

Doc. no.	P0024R1
Date:	2016-02-12
Project:	Programming Language C++
Reply to:	Jared Hoberock < jhoberock@nvidia.com >

The Parallelism TS Should be Standardized

Abstract

We survey implementation experience with the recently published C++ Technical Specification for Extensions for Parallelism and conclude that ample experience with its functionality exists to justify standardization in C++17. This paper describes various existing and pre-existing implementations of the TS's content and describes the additions to be made to the current C++ working paper (N4527) to integrate execution policies and parallel algorithms into the C++ Standard Library.

Implementation Experience

Although the first version of the Parallelism TS has been published only recently (N4354), we believe practioners have suitable experience with both existing and pre-existing implementations of its functionality to allow prompt standardization. Both implementors and users are ready for parallel algorithms in C++.

Development History

The technical content of the Parallelism TS was developed by domain experts in parallelism over the course of a few years. In 2012, representatives from NVIDIA (N3408) as well as representatives from Microsoft and Intel (N3429) independently proposed library approaches to parallelism within the C++ Standard Library. At the suggestion of SG1 the authors of these proposals submitted a design in a joint proposal (N3554) to parallelize the existing standard algorithms library. This proposal was refined into the Parallelism TS over the course of two years. During that refinement process, the authors of the Parallelism TS incorporated feedback from experimental implementations into the final design which was published in 2015. In total, the C++ Standardization Committee has three years of experience with the TS's design.

Existing implementations

Several different implementations of the Parallelism TS emerged during its preparation. We are aware of the following publically documented implementations.

- Microsoft: <http://parallelstl.codeplex.com>
- HPX: <http://stellar-group.github.io/hpx/docs/html/hpx/manual/parallel.html>
- Codeplay: <http://github.com/KhronosGroup/SyclParallelSTL>
- HSA: <http://www.hsafoundation.com/hsa-for-math-science>
- Thibaut Lutz: <http://github.com/t-lutz/ParallelSTL>
- NVIDIA: <http://github.com/n3554/n3554>

These implementations implement the functionality of the Parallelism TS to varying degrees and in different ways. For example, Microsoft's implementation appears complete and is implemented via Windows-specific tasking facilities. Thibaut Lutz' version also appears complete and is implemented by manipulating `std::thread` in a standard way. NVIDIA's implementation is partial and is implemented as a thin wrapper around Thrust, a pre-existing library similar in content to the Parallelism TS. This variety of implementation approaches exists by design: the abstractions of the Parallelism TS are intended to maximize flexibility of implementation.

Pre-existing implementations

The design of the Parallelism TS's functionality was inspired by several pre-existing libraries. Each of the following parallel algorithms libraries expose an iterator-based algorithm interface based on the conventions of the original Standard Template Library. We believe these libraries are a reasonable proxy for the content of the Parallelism TS.

- Thrust: <http://thrust.github.io>
- Boost.Compute: <http://github.com/boostorg/compute>
- Bolt: <http://github.com/HSA-Libraries/Bolt>
- libstdc++ Parallel Mode: http://gcc.gnu.org/onlinedocs/libstdc++/manual/parallel_mode.html
- AMP Algorithms Library: <http://ampalgorithms.codeplex.com>

These libraries have existed for several years, and some are widely deployed in production. Accordingly, we believe the features of the Parallelism TS are proven abstractions that represent standard practice and solve real challenges faced by real C++ programmers. These challenges exist because parallel architectures are so pervasive, and programming them correctly with existing low-level standard components is difficult. As a remedy, we believe that the high-level abstractions of the Parallelism TS must be standardized as soon as possible. C++ programmers should not have to wait beyond 2017 for standard parallel algorithms.

Future Support

The parallel algorithms and execution policies of the Parallelism TS are only a starting point. Already we anticipate opportunities for extending the Parallelism TS's functionality to increase programmer flexibility and expressivity. A fully-realized executors feature (N4414, N4406) will yield new, flexible ways of creating execution, including the execution of parallel algorithms. For example, executors will provide a programmatic means of specifying *where* execution is allowed to occur during parallel algorithm execution and will open the door for user-defined execution policies in addition to the Parallelism TS's closed set of standard policies. If the first version of the Parallelism TS is standardized in 2017, such additional features for parallelism will be well-positioned for 2020.

Summary of proposed changes

We propose to standardize the functionality of the Parallelism TS as specified. In summary:

- Add `exception_list` as a new subclause to Clause 19.
- Add execution policies as a new subclause to Clause 20.
- Introduce general wording for parallel algorithms as a new subclause to Clause 25.
- Add the new algorithms introduced by the Parallelism TS to Clauses 20, 25, and 26.

The locations of these changes merely represent our best guess and we expect them to change at the Library Working Group's discretion.

Acknowledgements

Thanks to Jens Maurer for wording suggestions.

References

1. N4308 - [Parallelizing the Standard Algorithms Library](#), J. Hoberock, M. Garland, O. Giroux, V. Grover, U. Kapasi, and J. Marathe. 2012.
2. N4329 - [A Library Solution to Parallelism](#), A. Laksberg, H. Sutter, A. Robison, and S. Mithani. 2012.
3. N3554 - [A Parallel Algorithms Library](#), J. Hoberock, J. Marathe, M. Garland, O. Giroux, V. Grover, A. Laksberg, H. Sutter, and A. Robison. 2013.
4. N4354 - [Programming Languages - Technical Specification for C++ Extensions for Parallelism](#), International Standards Organization. 2015.
5. N4414 - [Executors and schedulers, revision 5](#), C. Mysen. 2015.
6. N4406 - [Parallel Algorithms Need Executors](#), J. Hoberock et al. 2015.

Changelog

- P0024R0
 - Initial document.
- P0024R1
 - Eliminate greenlining.
 - Eliminate anachronistic use of TS-specific namespaces.
 - Eliminate NaN-numbered sections and number sequentially.
 - Introduce Acknowledgments section.
 - Eliminate unwarranted explanations in `for_each`'s notes.
 - Use the term "unspecified" in `reduce`'s note.
 - Replace use of "sequential order" with "indeterminately sequenced" in S 25.4p2.
 - Turn redundant restatement of 1.9 in S 25.4p3 into a note.

Contents

20 [General utilities library](#)

20.15 [Execution policies](#)

20.15.1 [In general](#)

20.15.2 [Header `<execution\_policy>` synopsis](#)

20.15.3 [Execution policy type trait](#)

20.15.4 [Sequential execution policy](#)

20.15.5 [Parallel execution policy](#)

20.15.6 [Parallel+Vector execution policy](#)

20.15.7 [Dynamic execution policy](#)

20.15.7.1 [execution_policy construct/assign](#)

20.15.7.2 [execution_policy object access](#)

20.15.8 [Execution policy objects](#)

19 [Diagnostics library](#)

19.6 [Class exception_list](#)

25 [Algorithms library](#)

25.2 [Parallel algorithms](#)

25.2.1 [Terms and definitions](#)

25.3 [Requirements on user-provided function objects](#)

25.4 [Effect of execution policies on algorithm execution](#)

25.5 [Parallel algorithm exceptions](#)

25.6 [ExecutionPolicy algorithm overloads](#)

25 [Algorithms library](#)

25.2 [Non-modifying sequence operations](#)

25.2.4 [For each](#)

26 [Numerics library](#)

26.2 [Definitions](#)

26 [Numerics library](#)

26.7 [Generalized numeric operations](#)

26.7.1 [Reduce](#)

26.7.2 [Exclusive scan](#)

26.7.3 [Inclusive scan](#)

26.7.4 [Transform reduce](#)

26.7.5 [Transform exclusive scan](#)

26.7.6 [Transform inclusive scan](#)

Introducing Execution Policies

Add the following entry to Table 44:

20.15	Execution policies	<code><execution_policy></code>
-------	--------------------	---------------------------------------

Add a new subclause to Clause 20:

This subclause describes classes that are *execution policy* types. An object of an execution policy type indicates the kinds of parallelism allowed in the execution of an algorithm and expresses the consequent requirements on the element access functions.

[*Example*:

```
std::vector<int> v = ...

// standard sequential sort
std::sort(v.begin(), v.end());

using namespace std::experimental::parallel;

// explicitly sequential sort
sort(seq, v.begin(), v.end());

// permitting parallel execution
sort(par, v.begin(), v.end());

// permitting vectorization as well
sort(par_vec, v.begin(), v.end());

// sort with dynamically-selected execution
size_t threshold = ...
execution_policy exec = seq;
if (v.size() > threshold)
{
    exec = par;
}

sort(exec, v.begin(), v.end());
```

— *end example*]

[*Note*: Because different parallel architectures may require idiosyncratic parameters for efficient execution, implementations of the Standard Library may provide additional execution policies to those described in this standard as extensions. — *end note*]

```
namespace std {
    // 20.15.3, Execution policy type trait
    template<class T> struct is_execution_policy;
    template<class T> constexpr bool is_execution_policy_v = is_execution_policy<T>::value;

    // 20.15.4, Sequential execution policy
    class sequential_execution_policy;

    // 20.15.5, Parallel execution policy
    class parallel_execution_policy;

    // 20.15.6, Parallel+Vector execution policy
    class parallel_vector_execution_policy;

    // 20.15.7, Dynamic execution policy
    class execution_policy;
}
```

```
template<class T> struct is_execution_policy { see below };
```

`is_execution_policy` can be used to detect parallel execution policies for the purpose of excluding function signatures from otherwise ambiguous overload resolution participation.

`is_execution_policy<T>` shall be a `UnaryTypeTrait` with a `BaseCharacteristic` of `true_type` if `T` is the type of a standard or implementation-defined execution policy, otherwise `false_type`.

[*Note*: This provision reserves the privilege of creating non-standard execution policies to the library implementation. — *end note*]

The behavior of a program that adds specializations for `is_execution_policy` is undefined.

```
class sequential_execution_policy{ unspecified };
```

The class `sequential_execution_policy` is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and require that a parallel algorithm's execution may not be parallelized.

20.15.5 Parallel execution policy

[execpol.par]

```
class parallel_execution_policy{ unspecified };
```

The class `parallel_execution_policy` is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be parallelized.

20.15.6 Parallel+Vector execution policy

[execpol.vec]

```
class parallel_vector_execution_policy{ unspecified };
```

The class `parallel_vector_execution_policy` is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be vectorized and parallelized.

20.15.7 Dynamic execution policy

[execpol.dynamic]

```
class execution_policy
{
public:
    // 20.15.7.1, execution_policy construct/assign
    template<class T> execution_policy(const T& exec);
    template<class T> execution_policy& operator=(const T& exec);

    // 20.15.7.2, execution_policy object access
    const type_info& type() const noexcept;
    template<class T> T* get() noexcept;
    template<class T> const T* get() const noexcept;
};
```

The class `execution_policy` is a container for execution policy objects. `execution_policy` allows dynamic control over standard algorithm execution.

[Example:

```
std::vector<float> sort_me = ...

using namespace std;experimental::parallel;
execution_policy exec = seq;

if(sort_me.size() > threshold)
{
    exec = std::par;
}

std::sort(exec, std::begin(sort_me), std::end(sort_me));
```

— end example]

Objects of type `execution_policy` shall be constructible and assignable from objects of type `T` for which `is_execution_policy<T>::value` is true.

20.15.7.1 execution_policy construct/assign

[parallel.execpol.con]

```
template<class T> execution_policy(const T& exec);
```

- ² *Effects:* Constructs an `execution_policy` object with a copy of `exec`'s state.
- ³ *Remarks:* This constructor shall not participate in overload resolution unless `is_execution_policy<T>::value` is true.

```
template<class T> execution_policy& operator=(const T& exec);
```

- ⁵ *Effects:* Assigns a copy of `exec`'s state to `*this`.
- ⁶ *Returns:* `*this`.

20.15.7.2 execution_policy object access

[parallel.execpol.access]

```
const type_info& type() const noexcept;
```

- ² *Returns:* `typeid(T)`, such that `T` is the type of the execution policy object contained by `*this`.

```
template<class T> T* get() noexcept;
template<class T> const T* get() const noexcept;
```

- ⁴ *Returns:* If `target_type() == typeid(T)`, a pointer to the stored execution policy object; otherwise a null pointer.
- ⁵ *Requires:* `is_execution_policy<T>::value` is true.

20.15.8 Execution policy objects

[parallel.execpol.objects]

```
constexpr sequential_execution_policy    seq{};
constexpr parallel_execution_policy      par{};
constexpr parallel_vector_execution_policy par_vec{};
```

The header `<experimental/execution_policy>` declares a global object associated with each type of execution policy defined by this Technical Specification.

Introducing exception_list

Add the following entry to Table 41:

19.6	Exception list	<exception_list>
------	----------------	------------------

Add a new subclause to Clause 19:

```
namespace std {  
namespace experimental {  
namespace parallel {  
inline namespace v1 {  
  
    class exception_list : public exception  
    {  
    public:  
        typedef unspecified iterator;  
  
        size_t size() const noexcept;  
        iterator begin() const noexcept;  
        iterator end() const noexcept;  
  
        const char* what() const noexcept override;  
    };  
}  
}  
}  
}
```

The class `exception_list` owns a sequence of `exception_ptr` objects. The parallel algorithms may use the `exception_list` to communicate uncaught exceptions encountered during parallel execution to the caller of the algorithm.

The type `exception_list::iterator` shall fulfill the requirements of `ForwardIterator`.

`size_t size() const noexcept;`

⁴ *Returns:* The number of `exception_ptr` objects contained within the `exception_list`.

⁵ *Complexity:* Constant time.

`iterator begin() const noexcept;`

⁷ *Returns:* An iterator referring to the first `exception_ptr` object contained within the `exception_list`.

`iterator end() const noexcept;`

⁹ *Returns:* An iterator that is past the end of the owned sequence.

`const char* what() const noexcept override;`

¹¹ *Returns:* An implementation-defined NTBS.

Introducing general parallel algorithms content

Add the following subclause to Clause 25:

25.2 Parallel algorithms

[algorithms.parallel]

This clause describes components that C++ programs may use to perform operations on containers and other sequences in parallel.

25.2.1 Terms and definitions

[algorithms.parallel.defns]

A *parallel algorithm* is a function template ~~described by this Technical Specification declared in namespace `std::experimental::parallel::v1`~~ listed in Table 1 with a formal template parameter named `ExecutionPolicy`.

Parallel algorithms access objects indirectly accessible via their arguments by invoking the following functions:

- All operations of the categories of the iterators that the algorithm is instantiated with.
- Functions on those sequence elements that are required by its specification.
- User-provided function objects to be applied during the execution of the algorithm, if required by the specification.
- Operations on those function objects required by the specification. [*Note:* See clause 25.1 of C++ *Standard Algorithms Library*. — *end note*]

These functions are herein called *element access functions*.

[*Example:* The `sort` function may invoke the following element access functions:

- Methods of the random-access iterator of the actual template argument, as per 24.2.7, as implied by the name of the template parameters `RandomAccessIterator`.
- The `swap` function on the elements of the sequence (as per 25.4.1.1 [sort]/2).
- The user-provided `Compare` function object.

— *end example*]

25.3 Requirements on user-provided function objects

[algorithms.parallel.user]

Function objects passed into parallel algorithms as objects of type `BinaryPredicate`, `Compare`, and `BinaryOperation` shall not directly or indirectly modify objects via their arguments.

25.4 Effect of execution policies on algorithm execution

[algorithms.parallel.exec]

Parallel algorithms have template parameters named `ExecutionPolicy` which describe the manner in which the execution of these algorithms may be parallelized and the manner in which they apply the element access functions.

The invocations of element access functions in parallel algorithms invoked with an execution policy object of type `sequential_execution_policy` ~~execute in sequential order in~~ are indeterminately sequenced (1.9) the calling thread.

The invocations of element access functions in parallel algorithms invoked with an execution policy object of type `parallel_execution_policy` are permitted to execute in an unordered fashion in either the invoking thread or in a thread implicitly created by the library to support parallel algorithm execution. ~~Any such invocations executing in the same thread are indeterminately sequenced with respect to each other.~~ [*Note:* Any such invocations executing in the same thread are indeterminately sequenced with respect to each other (1.9). It is the caller's responsibility to ensure correctness, for example that the invocation does not introduce data races or deadlocks. — *end note*]

[*Example:*

```
using namespace std::experimental::parallel;
int a[] = {0,1};
std::vector<int> v;
std::for_each(std::par, std::begin(a), std::end(a), [&](int i) {
    v.push_back(i*2+1);
});
```

The program above has a data race because of the unsynchronized access to the container `v`. — *end example*]

[*Example:*

```
using namespace std::experimental::parallel;
std::atomic<int> x = 0;
int a[] = {1,2};
std::for_each(std::par, std::begin(a), std::end(a), [&](int n) {
    x.fetch_add(1, std::memory_order_relaxed);
    // spin wait for another iteration to change the value of x
    while (x.load(std::memory_order_relaxed) == 1) { }
});
```

The above example depends on the order of execution of the iterations, and is therefore undefined (may deadlock). — *end example*]

[*Example:*

```
using namespace std::experimental::parallel;
int x=0;
std::mutex m;
int a[] = {1,2};
std::for_each(std::par, std::begin(a), std::end(a), [&](int) {
    m.lock();
    ++x;
    m.unlock();
});
```

The above example synchronizes access to object `x` ensuring that it is incremented correctly. — *end example*]

The invocations of element access functions in parallel algorithms invoked with an execution policy of type `parallel_vector_execution_policy` are permitted to execute in an unordered fashion in unspecified threads, and unsequenced with respect to one another within each thread. [*Note:* This means that multiple

function object invocations may be interleaved on a single thread. — *end note*]

[*Note:* This overrides the usual guarantee from the C++ standard, Section 1.9 [intro.execution] that function executions do not interleave with one another. — *end note*]

Since `parallel_vector_execution_policy` allows the execution of element access functions to be interleaved on a single thread, synchronization, including the use of mutexes, risks deadlock. Thus the synchronization with `parallel_vector_execution_policy` is restricted as follows:

A standard library function is *vectorization-unsafe* if it is specified to synchronize with another function invocation, or another function invocation is specified to synchronize with it, and if it is not a memory allocation or deallocation function. Vectorization-unsafe standard library functions may not be invoked by user code called from `parallel_vector_execution_policy` algorithms.

[*Note:* Implementations must ensure that internal synchronization inside standard library routines does not induce deadlock. — *end note*]

[*Example:*

```
using namespace std::experimental::parallel;
int x=0;
std::mutex m;
int a[] = {1,2};
std::for_each(std::par_vec, std::begin(a), std::end(a), [&](int) {
    m.lock();
    ++x;
    m.unlock();
});
```

The above program is invalid because the applications of the function object are not guaranteed to run on different threads. — *end example*]

[*Note:* The application of the function object may result in two consecutive calls to `m.lock` on the same thread, which may deadlock. — *end note*]

[*Note:* The semantics of the `parallel_execution_policy` or the `parallel_vector_execution_policy` invocation allow the implementation to fall back to sequential execution if the system cannot parallelize an algorithm invocation due to lack of resources. — *end note*]

Algorithms invoked with an execution policy object of type `execution_policy` execute internally as if invoked with the contained execution policy object.

The semantics of parallel algorithms invoked with an execution policy object of implementation-defined type are implementation-defined.

25.5 Parallel algorithm exceptions

[algorithms.parallel.exceptions]

During the execution of a standard parallel algorithm, if temporary memory resources are required and none are available, the algorithm throws a `std::bad_alloc` exception.

During the execution of a standard parallel algorithm, if the invocation of an element access function exits via an uncaught exception, the behavior of the program is determined by the type of execution policy used to invoke the algorithm:

- If the execution policy object is of type `class parallel_vector_execution_policy`, `std::terminate` shall be called.
- If the execution policy object is of type `sequential_execution_policy` or `parallel_execution_policy`, the execution of the algorithm exits via an exception. The exception shall be an `exception_list` containing all uncaught exceptions thrown during the invocations of element access functions, or optionally the uncaught exception if there was only one.

[*Note:* For example, when `for_each` is executed sequentially, if an invocation of the user-provided function object throws an exception, `for_each` can exit via the uncaught exception, or throw an `exception_list` containing the original exception. — *end note*]

[*Note:* These guarantees imply that, unless the algorithm has failed to allocate memory and exits via `std::bad_alloc`, all exceptions thrown during the execution of the algorithm are communicated to the caller. It is unspecified whether an algorithm implementation will "forge ahead" after encountering and capturing a user exception. — *end note*]

[*Note:* The algorithm may exit via the `std::bad_alloc` exception even if one or more user-provided function objects have exited via an exception. For example, this can happen when an algorithm fails to allocate memory while creating or adding elements to the `exception_list` object. — *end note*]
- If the execution policy object is of any other type, the behavior is implementation-defined.

25.6 ExecutionPolicy algorithm overloads

[algorithms.parallel.overloads]

The Parallel Algorithms Library provides overloads for each of the algorithms named in Table 1, corresponding to the algorithms with the same name in the C++ Standard Algorithms Library. For each algorithm in Table 1, if there are overloads for corresponding algorithms with the same name in the C++ Standard Algorithms Library, the overloads shall have an additional template type parameter named `ExecutionPolicy`, which shall be the first template parameter. In addition, each such overload shall have the new function parameter as the first function parameter of type `ExecutionPolicy&&`.

Unless otherwise specified, the semantics of `ExecutionPolicy` algorithm overloads are identical to their overloads without.

Parallel algorithms shall not participate in overload resolution unless `is_execution_policy<decay_t<ExecutionPolicy>>::value` is true.

Table 1 — Table of parallel algorithms			
<code>adjacent_difference</code>	<code>adjacent_find</code>	<code>all_of</code>	<code>any_of</code>
<code>copy</code>	<code>copy_if</code>	<code>copy_n</code>	<code>count</code>
<code>count_if</code>	<code>equal</code>	<code>exclusive_scan</code>	<code>fill</code>
<code>fill_n</code>	<code>find_end</code>	<code>find_end</code>	<code>find_first_of</code>
<code>find_if</code>	<code>find_if_not</code>	<code>for_each</code>	<code>for_each_n</code>
<code>generate</code>	<code>generate_n</code>	<code>includes</code>	<code>inclusive_scan</code>
<code>inner_product</code>	<code>inplace_merge</code>	<code>is_heap</code>	<code>is_heap_until</code>
<code>is_partitioned</code>	<code>is_sorted</code>	<code>is_sorted_until</code>	<code>lexicographical_compare</code>
<code>max_element</code>	<code>merge</code>	<code>min_element</code>	<code>minmax_element</code>
<code>mismatch</code>	<code>move</code>	<code>none_of</code>	<code>nth_element</code>
<code>partial_sort</code>	<code>partial_sort_copy</code>	<code>partition</code>	<code>partition_copy</code>
<code>reduce</code>	<code>remove</code>	<code>remove_copy</code>	<code>remove_copy_if</code>
<code>remove_if</code>	<code>replace</code>	<code>replace_copy</code>	<code>replace_copy_if</code>
<code>replace_if</code>	<code>reverse</code>	<code>reverse_copy</code>	<code>rotate</code>
<code>rotate_copy</code>	<code>search</code>	<code>search_n</code>	<code>set_difference</code>
<code>set_intersection</code>	<code>set_symmetric_difference</code>	<code>set_union</code>	<code>sort</code>

stable_partition	stable_sort	swap_ranges	transform
transform_exclusive_scan	transform_inclusive_scan	transform_reduce	uninitialized_copy
uninitialized_copy_n	uninitialized_fill	uninitialized_fill_n	unique
unique_copy			

[*Note:* Not all algorithms in the Standard Library have counterparts in [Table 1](#). — *end note*]

Modifications to header synopses

For each algorithm listed in [Table 1](#), add the signature of a parallel algorithm overload to the corresponding synopsis in Clause 20, Clause 25, or Clause 26.

Specify new algorithms

Add `for_each` with `ExecutionPolicy`, sequential `for_each_n`, and `for_each_n` with `ExecutionPolicy` to subclause 4.3.2:

```
template<class ExecutionPolicy,
         class InputIterator, class Function>
void for_each(ExecutionPolicy&& exec,
              InputIterator first, InputIterator last,
              Function f);
```

² *Effects:* Applies `f` to the result of dereferencing every iterator in the range `[first,last)`. [*Note:* If the type of `first` satisfies the requirements of a mutable iterator, `f` may apply nonconstant functions through the dereferenced iterator. — *end note*]

³ *Complexity:* Applies `f` exactly `last - first` times.

⁴ *Remarks:* If `f` returns a result, the result is ignored.

⁵ *Notes:* ~~Unlike its sequential form,~~^tThe parallel overload of `for_each` does not return a copy of its `Function` parameter, since parallelization may not permit efficient state accumulation.

⁶ *Requires:* ~~Unlike its sequential form,~~^tThe parallel overload of `for_each` requires `Function` to meet the requirements of `CopyConstructible`.

```
template<class InputIterator, class Size, class Function>
InputIterator for_each_n(InputIterator first, Size n,
                        Function f);
```

⁸ *Requires:* `Function` shall meet the requirements of `MoveConstructible` [*Note:* `Function` need not meet the requirements of `CopyConstructible`. — *end note*]

⁹ *Effects:* Applies `f` to the result of dereferencing every iterator in the range `[first,first + n)`, starting from `first` and proceeding to `first + n - 1`. [*Note:* If the type of `first` satisfies the requirements of a mutable iterator, `f` may apply nonconstant functions through the dereferenced iterator. — *end note*]

¹⁰ *Returns:* `first + n` for non-negative values of `n` and `first` for negative values.

¹¹ *Remarks:* If `f` returns a result, the result is ignored.

```
template<class ExecutionPolicy,
         class InputIterator, class Size, class Function>
InputIterator for_each_n(ExecutionPolicy && exec,
                        InputIterator first, Size n,
                        Function f);
```

¹³ *Effects:* Applies `f` to the result of dereferencing every iterator in the range `[first,first + n)`, starting from `first` and proceeding to `first + n - 1`. [*Note:* If the type of `first` satisfies the requirements of a mutable iterator, `f` may apply nonconstant functions through the dereferenced iterator. — *end note*]

¹⁴ *Returns:* `first + n` for non-negative values of `n` and `first` for negative values.

¹⁵ *Remarks:* If `f` returns a result, the result is ignored.

¹⁶ *Notes:* Unlike its sequential form, the parallel overload of `for_each_n` requires `Function` to meet the requirements of `CopyConstructible`.

Introducing numeric parallel algorithms definitions

Insert the following entry to Table 113:

26.2	Definitions
------	-------------

Insert the following subclause to Clause 26:

Define *GENERALIZED_SUM*(op, a1, ..., aN) as follows:

- a1 when N is 1
- op(*GENERALIZED_SUM*(op, b1, ..., bK), *GENERALIZED_SUM*(op, bM, ..., bN)) where
 - b1, ..., bN may be any permutation of a1, ..., aN and
 - 1 < K+1 = M ≤ N.

Define *GENERALIZED_NONCOMMUTATIVE_SUM*(op, a1, ..., aN) as follows:

- a1 when N is 1
- op(*GENERALIZED_NONCOMMUTATIVE_SUM*(op, a1, ..., aK), *GENERALIZED_NONCOMMUTATIVE_SUM*(op, aM, ..., aN) where 1 < K+1 = M ≤ N.

Specify new numeric algorithms

Add reduce, exclusive_scan, inclusive_scan, transform_reduce, transform_exclusive_scan, and transform_inclusive_scan to Clause 26.7. ~~The innermost section number is given as NaN to allow editorial discretion.~~

```
template<class InputIterator>
    typename iterator_traits<InputIterator>::value_type
        reduce(InputIterator first, InputIterator last);

2 Effects: Same as reduce(first, last, typename iterator_traits<InputIterator>::value_type{}).
```

```
template<class InputIterator, class T>
    T reduce(InputIterator first, InputIterator last, T init);

4 Effects: Same as reduce(first, last, init, plus<>()).
```

```
template<class InputIterator, class T, class BinaryOperation>
    T reduce(InputIterator first, InputIterator last, T init,
        BinaryOperation binary_op);

6 Returns: GENERALIZED_SUM(binary_op, init, *first, ..., *(first + (last - first) - 1)).

7 Requires: binary_op shall not invalidate iterators or subranges, nor modify elements in the range [first,last).

8 Complexity: O(last - first) applications of binary_op.

9 Notes: The primary difference between reduce and accumulate is that the behavior of reduce may be non-deterministic for non-associative or non-commutative binary_op. The difference between reduce and accumulate is that reduce applies binary_op in an unspecified order, which yields a non-deterministic result for non-associative or non-commutative binary_op such as floating-point addition.
```

```
template<class InputIterator, class OutputIterator, class T>
    OutputIterator exclusive_scan(InputIterator first, InputIterator last,
                                OutputIterator result,
                                T init);

2 Effects: Same as exclusive_scan(first, last, result, init, plus<>()).
```

```
template<class InputIterator, class OutputIterator, class T, class BinaryOperation>
    OutputIterator exclusive_scan(InputIterator first, InputIterator last,
                                OutputIterator result,
                                T init, BinaryOperation binary_op);

4 Effects: Assigns through each iterator i in [result,result + (last - first)) the value of
    GENERALIZED_NONCOMMUTATIVE_SUM(binary_op, init, *first, ..., *(first + (i - result) - 1)).

5 Returns: The end of the resulting range beginning at result.

6 Requires: binary_op shall not invalidate iterators or subranges, nor modify elements in the ranges [first,last) or [result,result + (last - first)).

7 Complexity: O(last - first) applications of binary_op.

8 Notes: The difference between exclusive_scan and inclusive_scan is that exclusive_scan excludes the ith input element from the ith sum. If binary_op is not mathematically associative, the behavior of exclusive_scan may be non-deterministic.
```

```
template<class InputIterator, class OutputIterator>
    OutputIterator inclusive_scan(InputIterator first, InputIterator last,
                                OutputIterator result);

2 Effects: Same as inclusive_scan(first, last, result, plus<>()).
```

```
template<class InputIterator, class OutputIterator, class BinaryOperation>
    OutputIterator inclusive_scan(InputIterator first, InputIterator last,
                                OutputIterator result,
                                BinaryOperation binary_op);

template<class InputIterator, class OutputIterator, class BinaryOperation>
    OutputIterator inclusive_scan(InputIterator first, InputIterator last,
                                OutputIterator result,
                                BinaryOperation binary_op, T init);

4 Effects: Assigns through each iterator i in [result,result + (last - first)) the value of
    GENERALIZED_NONCOMMUTATIVE_SUM(binary_op, *first, ..., *(first + (i - result))) OR
    GENERALIZED_NONCOMMUTATIVE_SUM(binary_op, init, *first, ..., *(first + (i - result))) if init is provided.

5 Returns: The end of the resulting range beginning at result.

6 Requires: binary_op shall not invalidate iterators or subranges, nor modify elements in the ranges [first,last) or [result,result + (last - first)).

7 Complexity: O(last - first) applications of binary_op.

8 Notes: The difference between exclusive_scan and inclusive_scan is that inclusive_scan includes the ith input element in the ith sum. If binary_op is not mathematically associative, the behavior of inclusive_scan may be non-deterministic.
```

```
template<class InputIterator, class UnaryFunction, class T, class BinaryOperation>
    T transform_reduce(InputIterator first, InputIterator last,
                      UnaryOperation unary_op, T init, BinaryOperation binary_op);
```

² *Returns:* *GENERALIZED_SUM*(binary_op, init, unary_op(*first), ..., unary_op(*(first + (last - first) - 1))).

³ *Requires:* Neither unary_op nor binary_op shall invalidate subranges, or modify elements in the range [first,last)

⁴ *Complexity:* O(last - first) applications each of unary_op and binary_op.

⁵ *Notes:* transform_reduce does not apply unary_op to init.

26.7.5 Transform exclusive scan

[\[transform.exclusive.scan\]](#)

```
template<class InputIterator, class OutputIterator,
        class UnaryOperation,
        class T, class BinaryOperation>
    OutputIterator transform_exclusive_scan(InputIterator first, InputIterator last,
                                          OutputIterator result,
                                          UnaryOperation unary_op,
                                          T init, BinaryOperation binary_op);
```

² *Effects:* Assigns through each iterator i in [result,result + (last - first)) the value of *GENERALIZED_NONCOMMUTATIVE_SUM*(binary_op, init, unary_op(*first), ..., unary_op(*(first + (i - result) - 1))).

³ *Returns:* The end of the resulting range beginning at result.

⁴ *Requires:* Neither unary_op nor binary_op shall invalidate iterators or subranges, or modify elements in the ranges [first,last) or [result,result + (last - first)).

⁵ *Complexity:* O(last - first) applications each of unary_op and binary_op.

⁶ *Notes:* The difference between transform_exclusive_scan and transform_inclusive_scan is that transform_exclusive_scan excludes the ith input element from the ith sum. If binary_op is not mathematically associative, the behavior of transform_exclusive_scan may be non-deterministic. transform_exclusive_scan does not apply unary_op to init.

26.7.6 Transform inclusive scan

[\[transform.inclusive.scan\]](#)

```
template<class InputIterator, class OutputIterator,
        class UnaryOperation,
        class BinaryOperation>
    OutputIterator transform_inclusive_scan(InputIterator first, InputIterator last,
                                          OutputIterator result,
                                          UnaryOperation unary_op,
                                          BinaryOperation binary_op);
```

```
template<class InputIterator, class OutputIterator,
        class UnaryOperation,
        class BinaryOperation, class T>
    OutputIterator transform_inclusive_scan(InputIterator first, InputIterator last,
                                          OutputIterator result,
                                          UnaryOperation unary_op,
                                          BinaryOperation binary_op, T init);
```

² *Effects:* Assigns through each iterator i in [result,result + (last - first)) the value of *GENERALIZED_NONCOMMUTATIVE_SUM*(binary_op, unary_op(*first), ..., unary_op(*(first + (i - result)))) or *GENERALIZED_NONCOMMUTATIVE_SUM*(binary_op, init, unary_op(*first), ..., unary_op(*(first + (i - result)))) if init is provided.

³ *Returns:* The end of the resulting range beginning at result.

⁴ *Requires:* Neither unary_op nor binary_op shall invalidate iterators or subranges, or modify elements in the ranges [first,last) or [result,result + (last - first)).

⁵ *Complexity:* O(last - first) applications each of unary_op and binary_op.

⁶ *Notes:* The difference between transform_exclusive_scan and transform_inclusive_scan is that transform_inclusive_scan includes the ith input element from the ith sum. If binary_op is not mathematically associative, the behavior of transform_inclusive_scan may be non-deterministic. transform_inclusive_scan does not apply unary_op to init.