

Text Formatting

1. [Revision History](#)
 1. [Changes since R2](#)
 2. [Changes since R1](#)
 3. [Changes since R0](#)
2. [Introduction](#)
3. [Design](#)
 1. [Format string syntax](#)
 2. [Extensibility](#)
 3. [Safety](#)
 4. [Locale support](#)
 5. [Positional arguments](#)
 6. [Performance](#)
 7. [Binary footprint](#)
 8. [Compile-time processing of format strings](#)
 9. [Argument visitation](#)
 10. [Impact on existing code](#)
4. [Proposed wording](#)
 1. [Formatting utilities \[format\]](#)
 1. [Header <format> synopsis \[format.syn\]](#)
 2. [Class format_error \[format.error\]](#)
 3. [Arguments \[format.arguments\]](#)
 1. [Class template basic_format_arg](#)
 2. [Argument visitation](#)
 3. [Class template basic_format_args](#)
 4. [Function template make_format_args](#)
 5. [Function template make_wformat_args](#)
 4. [Formatting functions \[format.functions\]](#)
 5. [Formatter \[format.formatter\]](#)
 1. [Class template basic_parse_context](#)
 2. [Class template basic_format_context](#)
 3. [Class template formatter](#)
 4. [Formatter requirements](#)
 2. [Header <format> synopsis \[format.syn\]](#)
 3. [Class format_error \[format.error\]](#)
 4. [Arguments \[format.arguments\]](#)
 1. [Class template basic_format_arg](#)
 2. [Argument visitation](#)
 3. [Class template basic_format_args](#)
 4. [Function template make_format_args](#)
 5. [Function template make_wformat_args](#)
 5. [Formatting functions \[format.functions\]](#)
 6. [Formatter \[format.formatter\]](#)
 1. [Class template basic_parse_context](#)
 2. [Class template basic_format_context](#)
 3. [Class template formatter](#)
 4. [Formatter requirements](#)
5. [Related work](#)
6. [Implementation](#)
7. [Acknowledgements](#)
8. [References](#)
9. [Appendix A: Benchmarks](#)
10. [Appendix B: Binary code comparison](#)

Revision History

Changes since [R2](#)

- Add section [Argument visitation](#) with an example of how to use the visitation API to implement dynamic format specifiers.
- Rename visit to visit_format_arg to distinguish from variant's visit.
- Merge [format.syntax] into [format.functions].
- Remove the restriction that '\0' cannot be used as a fill character.
- Replace Postconditions with Ensures.

Changes since [R1](#)

- Add the `format_to_n` function taking an output iterator and a size.
- Add a note that compile-time processing of format strings applies to user-defined types with custom parsers to section [Compile-time processing of format strings](#).
- Rename `count` to `formatted_size`.
- Drop nested namespace `fmt`.
- Add `format` prefix or infix to class and function names to avoid potential name collision after removing the nested namespace.
- Improve wording.
- Replace the requirement of implementing a formatter via ostream insertion operator if the latter is provided with a note that it can be implemented in such way.
- Expand the Acknowledgements section and remove "Reply to".

Changes since [R0](#)

- Add section [Compile-time processing of format strings](#).
- Separate parsing and formatting in the extension API replacing `format_value` function template with class template `formatter` to allow compile-time processing of format strings.
- Change return type of `format_to` and `vformat_to` to `OutputIterator` in synopsis.
- Remove sections Null-terminated string view and Format string, and replace `basic_cstring_view` with `basic_string_view`.
- Add a link to the implementation in [Introduction](#).
- Add a note regarding time formatting and compatibility with D0355 "Extending <chrono> to Calendars and Time Zones" [\[16\]](#) to section [Extensibility](#).
- Rename `basic_args` to `basic_format_args`.
- Rename `is_numeric` to `is_arithmetic`.
- Add the `count` function that counts the number of characters and use it to define output ranges.
- Remove `basic_buffer` and section Formatting buffer and replace buffers with output iterators.
- Add [Appendix A: Benchmarks](#).
- Explain the purpose of the type-erased API in more details in the [Binary footprint](#) section.
- Add [Appendix B: Binary code comparison](#).
- Add formatting function overloads for `wchar_t` strings.

Introduction

Even with proliferation of graphical and voice user interfaces, text remains one of the main ways for humans to interact with computer programs and programming languages provide a variety of methods to perform text formatting. The first thing we do when learning a new programming language is often write a "Hello, World!" program that performs simple formatted output.

C++ has not one but two standard APIs for producing formatted output, the `printf` family of functions inherited from C and the I/O streams library (`iostreams`). While `iostreams` are usually the recommended way of producing formatted output in C++ for safety and extensibility reasons, `printf` offers some advantages, such as an arguably more natural function call API, the separation of formatted message and arguments, possibly with argument reordering as a POSIX extension, and often more compact source and binary code.

This paper proposes a new text formatting library that can be used as a safe and extensible alternative to the `printf` family of functions. It is intended to complement the existing C++ I/O streams library and reuse some of its infrastructure such as overloaded insertion operators for user-defined types.

Example:

```
string message = format("The answer is {}.", 42);
```

A full implementation of this proposal is available at <https://github.com/fmtlib/fmt/tree/std>.

Design

Format string syntax

Variations of the `printf` format string syntax are arguably the most popular among the programming languages and C++ itself inherits `printf` from C [\[1\]](#). The advantage of the `printf` syntax is that many programmers are familiar with it. However, in its current form it has a number of issues:

- Many format specifiers like `hh`, `h`, `l`, `j`, etc. are used only to convey type information. They are redundant in type-safe formatting and would unnecessarily complicate specification and parsing.
- There is no standard way to extend the syntax for user-defined types.
- There are subtle differences between different implementations. For example, POSIX positional arguments [\[2\]](#) are not supported on some systems [\[6\]](#).
- Using `'%'` in a custom format specifier poses difficulties, e.g. for `put_time`-like time formatting.

Although it is possible to address these issues while maintaining resemblance to the original `printf` format, this will still break compatibility and can potentially be more confusing to users than introducing a different syntax.

Therefore we propose a new syntax based on the ones used in Python [\[3\]](#), the .NET family of languages [\[4\]](#), and Rust [\[5\]](#). This syntax employs `'{'` and `'}'` as replacement field delimiters instead of `'%'` and it is described in detail in [\[format.functions\]](#). Some advantages of the proposal are:

- A consistent and easy to parse mini-language focused on formatting rather than conveying type information
- Extensibility and support for custom format strings for user-defined types
- Positional arguments
- Support for both locale-specific and locale-independent formatting (see [Locale support](#))
- Formatting improvements such as better alignment control, fill character, and binary representation

The syntax is expressive enough to enable translation, possibly automated, of most `printf` format strings. The correspondence between `printf` and the new syntax is given in the following table:

printf	new
-	<
+	+
<i>space</i>	<i>space</i>
#	#
0	0
hh	unused
h	unused
l	unused
ll	unused
j	unused
z	unused
t	unused
L	unused
c	c (optional)
s	s (optional)
d	d (optional)
i	d (optional)
o	o
x	x
X	X
u	d (optional)
f	f
F	F
e	e
E	E
a	a
A	A
g	g (optional)
G	G
n	unused
p	p (optional)

Width and precision are represented similarly in `printf` and the proposed syntax with the only

difference that runtime value is specified by '*' in the former and '{}' in the latter, possibly with the index of the argument inside the braces:

```
printf("%*s", 10, "foo");
format("{:{}s}", "foo", 10);
```

As can be seen from the table above, most of the specifiers remain the same which simplifies migration from `printf`. A notable difference is in the alignment specification. The proposed syntax allows left, center, and right alignment represented by '<', '^', and '>' respectively which is more expressive than the corresponding `printf` syntax. The latter only supports left and right alignment.

The following example uses center alignment and '*' as a fill character:

```
format("{:^30s}", "centered");
```

resulting in "*****centered*****". The same formatting cannot be easily achieved with `printf`.

In addition to positional arguments, the grammar can be easily extended to support named arguments.

Extensibility

Both the format string syntax and the API are designed with extensibility in mind. The mini-language can be extended for user-defined types and users can provide functions that implement parsing, possibly at compile time, and formatting for such types.

The general syntax of a replacement field in a format string is

```
replacement-field ::= '{' [arg-id] [':' format-spec] '}'
```

where `format-spec` is predefined for built-in types, but can be customized for user-defined types. For example, the syntax can be extended for `put_time`-like date and time formatting

```
time_t t = time(nullptr);
string date = format("The date is {0:%Y-%m-%d}.", *localtime(&t));
```

by providing a specialization of `formatter` for `tm`:

```
template<>
struct formatter<tm> {
    constexpr parse_context::iterator parse(parse_context& ctx);

    template<class FormatContext>
    typename FormatContext::iterator format(const tm& tm, FormatContext& ctx);
};
```

The `formatter<tm>::parse` function parses the `format-spec` portion of the format string corresponding to the current argument and `formatter<tm>::format` formats the value and writes the output via the iterator `ctx.begin()`.

Note that date and time formatting is not covered by this proposal but formatting facilities provided by D0355 "Extending `<chrono>` to Calendars and Time Zones" [\[16\]](#) can be easily implemented using this extension API.

An implementation of `formatter<T>::format` can use ostream insertion operator<< for user-defined type `T` if available.

The extension API is based on specialization instead of the argument-dependent lookup, because the `parse` function doesn't take the object to be formatted as an argument and therefore some other way of parameterizing it on the argument type `T` such as introducing a dummy argument has to be used, e.g.

```
constexpr auto parse(type<T>, parse_context& ctx);
```

Safety

Formatting functions rely on variadic templates instead of the mechanism provided by `<cstdarg>`. The type information is captured automatically and passed to formatters guaranteeing type safety and making many of the `printf` specifiers redundant (see [Format String Syntax](#)). Memory management is automatic to prevent buffer overflow errors common to `printf`.

Locale support

As pointed out in P0067 "Elementary string conversions"[\[17\]](#) there is a number of use cases that do not require internationalization support, but do require high throughput when produced by a server. These include various text-based interchange formats such as JSON or XML. The need for locale-independent functions for conversions between integers and strings and between floating-point numbers and strings has also been highlighted in [N4412: Shortcomings of iostreams](#). Therefore a user should be able to easily control whether to use locales or not during formatting.

We follow Python's approach [\[3\]](#) and designate a separate format specifier 'n' for locale-aware numeric formatting. It applies to all integral and floating-point types. All other specifiers produce output unaffected by locale settings. This can also have positive effect on performance because locale-independent formatting can be implemented more efficiently.

Positional arguments

An important feature for localization is the ability to rearrange formatting arguments as the word order may vary in different languages [\[7\]](#). For example:

```
printf("String '%s' has %d characters\n", string, length(string));
```

A possible German translation of the format string might be:

```
"%2$d Zeichen lang ist die Zeichenkette '%1$s'\n"
```

using POSIX positional arguments [\[2\]](#). Unfortunately these positional specifiers are not portable [\[6\]](#). The C++ I/O streams don't support such rearranging of arguments by design because they are interleaved with the portions of the literal string:

```
cout << "String '" << string << "' has " << length(string) << " characters\n";
```

The current proposal allows both positional and automatically numbered arguments, for example:

```
format("String '{}'" has {} characters\n", string, length(string));
```

with the German translation of the format string being:

```
"{1} Zeichen lang ist die Zeichenkette '{0}'\n"
```

Performance

The formatting library has been designed with performance in mind. It tries to minimize the number of virtual function calls and dynamic memory allocations done per formatting operation. In particular, if formatting output can fit into a fixed-size array allocated on stack, it should be possible to avoid both of them altogether by using a suitable API.

The `format_to` function takes an arbitrary output iterator and, for performance reasons, can be specialized for random-access and contiguous iterators as shown in the reference implementation [\[14\]](#).

The locale-independent formatting can also be implemented more efficiently than the locale-aware one. However, the main goal for the former is to support specific use cases (see [Locale support](#)) rather than to improve performance.

See [Appendix A: Benchmarks](#) for a small performance comparison of the reference implementation of this proposal versus the standard formatting facilities.

Binary footprint

In order to minimize binary code size each formatting function that uses variadic templates can be implemented as a small inline wrapper around its non-variadic counterpart. This wrapper creates a `basic_format_args` object, representing an array of type-erased argument references, with `make_format_args` and calls the non-variadic function to do the actual work. For example, the `format` variadic function calls `vformat`:

```
string vformat(string_view fmt, format_args args);

template<class... Args>
inline string format(string_view fmt, const Args&... args) {
    return vformat(fmt, make_format_args(args...));
}
```

```
}
```

`basic_format_args` can be implemented as an array of tagged unions. If the number of arguments is small then the tags that indicate the arguments types can be combined and passed into a formatting function as a single integer. This single integer representing all argument types is computed at compile time and only needs to be stored, resulting in smaller binary code.

Given a reasonable optimizing compiler, this will result in a compact per-call binary code, effectively consisting of placing argument pointers (or, possibly, copies for primitive types) and packed tags on the stack and calling a formatting function. See [Appendix B: Binary code comparison](#) for a specific example.

Exposing the type-erased API rather than making it an implementation detail and only providing variadic functions allows applying the same technique to the user code. For example, consider a simple logging function that writes a formatted error message to `clog`:

```
template<class... Args>
void log_error(int code, string_view fmt, const Args&... args) {
    clog << "Error " << code << ": " << format(fmt, args...);
}
```

The code for this function will be generated for every combination of argument types which can be undesirable. However, if we use the type-erased API, there will be only one instance of the logging function while the wrapper function can be trivially inlined:

```
void vlog_error(int code, string_view fmt, format_args args) {
    clog << "Error " << code << ": " << vformat(fmt, args);
}

template<class... Args>
inline void log_error(int code, string_view fmt, const Args&... args) {
    vlog_error(code, fmt, make_format_args(args...));
}
```

The current design allows users to easily switch between the two approaches.

Compile-time processing of format strings

It is possible to parse format strings at compile time with `constexpr` functions which has been demonstrated in the reference implementation [\[14\]](#) and in [\[18\]](#). Unfortunately a satisfactory API cannot be provided using existing C++17 features. Ideally we would like to use a function call API similar to the one proposed in this paper:

```
template<class String, class... Args>
string format(String fmt, const Args&... args);
```

where `String` is some type representing a compile-time format string, for example

```
struct MyString {
    static constexpr string_view value() { return string_view("{} ", 2); }
};
```

However, requiring a user to create a new type either manually or via a macro for every format string is not acceptable. P0424R1 "Reconsidering literal operator templates for strings" [\[15\]](#) provides a solution to this problem based on user-defined literal operators. If this or other similar proposal for compile-time strings is accepted into the standard, it should be easy to provide additional formatting APIs that make use of this feature. Obviously runtime checks will still be needed in cases where the format string is not known at compile time, but as shown in [Appendix A: Benchmarks](#) even with runtime parsing performance can be on par with or better than that of existing methods.

Compile-time processing of format strings [can work with a user-defined type `T`](#) if the `formatter<T>::parse` function is `constexpr`.

Argument visitation

The argument visitation API can be used to implement dynamic format specifiers, for example, width and precision passed as additional formatting arguments as opposed to being encoded in the format string itself:

```
int width = 10;
int precision = 3;
```

```
auto s = format("{0:{1}.{2}f}", 12.345678, width, precision);
// s == "      12.346"
```

An example of a user-defined formatter with dynamic width:

```
struct Answer {};

template<>
struct formatter<Answer> {
    int width_arg_index = 0;

    // Parses dynamic width in the format "<digit>".
    auto parse(parse_context& parse_ctx) {
        auto iter = parse_ctx.begin();
        auto get_char = [&]() { return iter != parse_ctx.end() ? *iter : 0; };

        if (get_char() != '{')
            return iter;
        ++iter;
        char c = get_char();
        if (!std::isdigit(c) || (++iter, get_char()) != '}')
            throw format_error("invalid format");
        width_arg_index = c - '0';
        return ++iter;
    }

    auto format(Answer, format_context& format_ctx) {
        auto arg = format_ctx.args().get(width_arg_index);
        int width = visit_format_arg([](auto value) -> int {
            if constexpr (!std::is_integral_v<decltype(value)>)
                throw format_error("width is not integral");
            else if (value < 0 || value > std::numeric_limits<int>::max())
                throw format_error("invalid width");
            else
                return value;
        }, arg);
        return format_to(format_ctx.out(), "{:{}", 42, width);
    }
};

std::string s = format("{0:{1}}", Answer(), 10);
// s == "      42"
```

In many cases formatting can be delegated to standard formatters which makes manual handling of dynamic specifiers unnecessary, but the latter is still important for more complex cases.

This API can also be used to implement a custom formatting engine, such as the one compatible with the printf syntax, to provide some of the benefits of the current proposal to the legacy code.

Impact on existing code

The proposed formatting API is defined in the new header `<format>` and should have no impact on existing code.

Proposed wording

Add a new section in 19 [utilities].

Formatting utilities [format]

Header `<format>` synopsis [format.syn]

```
namespace std {
    // [format.error], class format_error
    class format_error;

    // [format.formatter], formatter
    template<class charT> class basic_parse_context;
    using parse_context = basic_parse_context<char>;
    using wparse_context = basic_parse_context<wchar_t>;

    template<class OutputIterator, class charT> class basic_format_context;
    using format_context = basic_format_context<unspecified, char>;
    using wformat_context = basic_format_context<unspecified, wchar_t>;
```

```

template<class T, class charT = char> struct formatter;

// \[format.arguments\], arguments
template<class Context> class basic_format_arg;

template<class Visitor, class Context>
    see below visit_format_arg(Visitor&& vis, basic_format_arg<Context> arg);

template<class Context, class... Args> using format_arg_store = unspecified;

template<class Context> class basic_format_args;
using format_args = basic_format_args<format_context>;
using wformat_args = basic_format_args<wformat_context>;

template<class OutputIterator, class charT>
    using format_args_t = basic_format_args<basic_format_context<OutputIterator, charT>>;

template<class Context, class... Args>
    format_arg_store<Context, Args...> make_format_args(const Args&... args);
template<class... Args>
    format_arg_store<format_context, Args...> make_format_args(const Args&... args);
template<class... Args>
    format_arg_store<wformat_context, Args...> make_wformat_args(const Args&... args);

// \[format.functions\], formatting functions
template<class... Args>
    string format(string_view fmt, const Args&... args);
template<class... Args>
    wstring format(wstring_view fmt, const Args&... args);

string vformat(string_view fmt, format_args args);
wstring vformat(wstring_view fmt, wformat_args args);

template<class OutputIterator, class... Args>
    OutputIterator format_to(OutputIterator out, string_view fmt,
                             const Args&... args);
template<class OutputIterator, class... Args>
    OutputIterator format_to(OutputIterator out, wstring_view fmt,
                             const Args&... args);

template<class OutputIterator>
    OutputIterator vformat_to(OutputIterator out, string_view fmt,
                             format_args_t<OutputIterator, char> args);
template<class OutputIterator>
    OutputIterator vformat_to(OutputIterator out, wstring_view fmt,
                             format_args_t<OutputIterator, wchar_t> args);

template<class OutputIterator, class Size>
    struct format_to_n_result {
        OutputIterator out;
        Size size;
    };

template<class OutputIterator, class Size, class... Args>
    format_to_n_result<OutputIterator, Size>
        format_to_n(OutputIterator out, Size n, string_view fmt,
                     const Args&... args);
template<class OutputIterator, class Size, class... Args>
    format_to_n_result<OutputIterator, Size>
        format_to_n(OutputIterator out, Size n, wstring_view fmt,
                     const Args&... args);

template<class... Args>
    size_t formatted_size(string_view fmt, const Args&... args);
template<class... Args>
    size_t formatted_size(wstring_view fmt, const Args&... args);
}

```

Class `format_error` [\[format.error\]](#)

```

namespace std {
    class format_error : public runtime_error {
    public:
        explicit format_error(const string& what_arg);
        explicit format_error(const char* what_arg);
    };
}

```


The class `format_error` defines the type of objects thrown as exceptions to report errors from the formatting library.

```
format_error(const string& what_arg);
```

Effects: Constructs an object of class `format_error`.

Ensures: `strcmp(what(), what_arg.c_str()) == 0`.

```
format_error(const char* what_arg);
```

Effects: Constructs an object of class `format_error`.

Ensures: `strcmp(what(), what_arg) == 0`.

Arguments [format.arguments]

Class template `basic_format_arg`

```
namespace std {
    template<class Context>
    class basic_format_arg {
    public:
        class handle;

        basic_format_arg();

        explicit operator bool() const noexcept;

        bool is_arithmetic() const;
        bool is_integral() const;
    };
}
```

The class template `basic_format_arg` describes an object that can refer to a formatting function argument. It is parameterized on a formatting context type (see [\[format.formatter\]](#)) and can hold a value of one of the following types:

- `bool`
- `charT`
- any integral type other than `bool` and `charT`
- any floating-point type
- `const charT*`
- `basic_string_view<charT>`
- `const void*`
- `basic_format_arg::handle` referring to an argument of a user-defined type `T` with an enabled specialization of `formatter<T>`
- monostate representing an empty state, i.e. when the object does not refer to an argument

where `charT` is `typename Context::char_type`. The value can be accessed via the visitation interface defined in the next section. `basic_format_arg` must preserve argument values, but does not have to distinguish between integral types other than `charT` and `bool`, e.g. an argument of type `short` can be stored as `int` and observed as `int` via the visitation interface. Similarly, `basic_format_arg` does not have to distinguish between floating-point types.

```
basic_format_arg();
```

Effects: Constructs a `basic_format_arg` object that doesn't refer to an argument.

Ensures: `!(*this)`.

```
explicit operator bool() const noexcept;
```

Returns: `true` if and only if `*this` refers to an argument.

```
bool is_arithmetic() const;
```

Returns: `true` if and only if `*this` refers to an argument of an arithmetic type.

```
bool is_integral() const;
```

Returns: true if and only if `*this` refers to an argument of an integral type.

```
namespace std {
    template<class Context>
    class basic_format_arg<Context>::handle {
    public:
        void format(Context& ctx) const;
    };
}
```

The class `handle` describes an object that can refer to a formatting function argument of a user-defined type.

```
void format(Context& ctx) const;
```

Effects: Constructs a local variable of type `typename Context::template formatter_type<T>` (named `f` for exposition purposes), where `T` is the type of the object `obj` referred to by this handle. Executes `ctx.parse_context().advance_to(f.parse(ctx.parse_context()))`. Then calls `ctx.advance_to(f.format(obj, ctx))`.

Argument visitation

```
template<class Visitor, class Context>
    see below visit_format_arg(Visitor&& vis, basic_format_arg<Context> arg);
```

Requires: The expression in the Effects section shall be a valid expression of the same type, for all alternative value types of a formatting argument. Otherwise, the program is ill-formed.

Effects: Let `value` be the value stored in `arg`. Returns `INVOKE(std::forward(vis), value)`.

Remarks: The return type is the common type of all possible `INVOKE` expressions in the Effects section. Since exact value types are implementation-defined, visitors should use type traits to handle multiple types.

[*Example:*

```
    auto uint_value = visit_format_arg([](auto value) -> unsigned {
        if constexpr (std::is_unsigned_v<decltype(value)>)
            return value;
        return 0;
    }, arg);
```

— *end example*]

Class template `basic_format_args`

```
namespace std {
    template<class Context>
    class basic_format_args {
    public:
        using size_type = size_t;

        basic_format_args() noexcept;

        template<class... Args>
            basic_format_args(const format_arg_store<Context, Args...>& store);

        basic_format_arg<Context> get(size_type i) const;
    };
}
```

An object of type `basic_format_args` provides access to formatting arguments. Copying a `basic_format_args` object does not copy the arguments.

```
basic_format_args() noexcept;
```

Effects: Constructs an empty `basic_format_args` object.

Ensures: `!get(0)`.

```
template<class... Args>
    basic_format_args(const format_arg_store<Context, Args...>& store);
```

Effects: Constructs a `basic_format_args` object that provides access to the arguments captured in

store.

```
basic_format_arg<Context> get(size_type i) const;
```

Returns: A `basic_format_arg` object that represents an argument at index `i` if `i < the number of arguments`. Otherwise, returns an empty `basic_format_arg` object.

Function template `make_format_args`

```
template<class Context, class... Args>
    format_arg_store<Context, Args...> make_format_args(const Args&... args);
template<class... Args>
    format_arg_store<format_context, Args...> make_format_args(const Args&... args);
```

Effects: The function returns a `format_arg_store` object that refers to formatting arguments `args`.

Function template `make_wformat_args`

```
template<class... Args>
    format_arg_store<wformat_context, Args...> make_wformat_args(const Args&... args);
```

Effects: The function returns a `format_arg_store` object that refers to formatting arguments `args`.

Formatting functions [format.functions]

```
template<class... Args>
    string format(string_view fmt, const Args&... args);
```

Returns: `vformat(fmt, make_format_args(args...))`.

```
template<class... Args>
    wstring format(wstring_view fmt, const Args&... args);
```

Returns: `vformat(fmt, make_wformat_args(args...))`.

```
string vformat(string_view fmt, format_args args);
wstring vformat(wstring_view fmt, wformat_args args);
```

Effects: The function returns a `string` object constructed from the format string argument `fmt` with each replacement field substituted with the character representation of the argument it refers to, formatted according to the specification given in the field.

Returns: The formatted string.

Throws: `format_error` if `fmt` is not a valid format string.

```
template<class OutputIterator, class... Args>
    OutputIterator format_to(OutputIterator out, string_view fmt,
                           const Args&... args);
```

Returns: `vformat_to(out, fmt, make_format_args(args...))`.

```
template<class OutputIterator, class... Args>
    OutputIterator format_to(OutputIterator out, wstring_view fmt,
                           const Args&... args);
```

Returns: `vformat_to(out, fmt, make_wformat_args(args...))`.

```
template<class OutputIterator, class... Args>
    OutputIterator vformat_to(OutputIterator out, string_view fmt,
                           format_args_t<OutputIterator, char> args);
template<class OutputIterator, class... Args>
    OutputIterator vformat_to(OutputIterator out, wstring_view fmt,
                           format_args_t<OutputIterator, wchar_t> args);
```

Effects: The function writes to the range `[out, out + size)`, where `size` is the output size, the format string `fmt` with each replacement field substituted with the character representation of the argument it refers to, formatted according to the specification given in the field.

Returns: The end of the output range.

Throws: `format_error` if `fmt` is not a valid format string.

```
template<class OutputIterator, class Size, class... Args>
```

```
format_to_n_result<OutputIterator, Size>
format_to_n(OutputIterator out, Size n, string_view fmt,
            const Args&... args);
template<class OutputIterator, class Size, class... Args>
format_to_n_result<OutputIterator, Size>
format_to_n(OutputIterator out, Size n, wstring_view fmt,
            const Args&... args);
```

Effects: The function writes to the range `[out, out + n)` the format string `fmt` with each replacement field substituted with the character representation of the argument it refers to, formatted according to the specification given in the field.

Returns: `format_to_n_result` containing the end of the output range and the total (not truncated) output size.

Throws: `format_error` if `fmt` is not a valid format string.

```
template<class... Args>
size_t formatted_size(string_view fmt, const Args&... args);
template<class... Args>
size_t formatted_size(wstring_view fmt, const Args&... args);
```

Effects: The function returns the number of characters in the output of `format(fmt, args...)` without constructing the formatted string.

Returns: The number of characters in the output of `format(fmt, args...)`.

Throws: `format_error` if `fmt` is not a valid format string.

The `fmt` string consists of zero or more *replacement fields*, *escape sequences*, and ordinary multibyte characters. All ordinary multibyte characters are copied unchanged to the output. An escape sequence is one of `{ {` or `}}`; it is replaced with `{` or `}` respectively in the output. The syntax of replacement fields is as follows:

```
replacement-field ::= '{' [arg-id] [':' format-spec] '}'
arg-id            ::= integer
integer           ::= digit [integer]
digit             ::= '0'...'9'
```

The `arg-id` field specifies the index of the argument in `args` whose value is to be formatted and inserted into the output instead of the replacement field. The optional `format-spec` field specifies a non-default format for the replacement value.

If the numeric `arg-ids` in a format string are 0, 1, 2, ... in sequence, they can all be omitted (not just some) and the numbers 0, 1, 2, ... will be automatically used in that order. Mixing automatic and manual indexing is not allowed.

[*Example:*

```
string s0 = format("From {} to {}", "a", "b"); // OK: automatic indexing
string s1 = format("From {1} to {0}", "a", "b"); // OK: manual indexing
string s2 = format("From {0} to {}", "a", "b"); // Error: mixing automatic and manual indexing
string s3 = format("From {} to {1}", "a", "b"); // Error: mixing automatic and manual indexing
```

— *end example*]

The `format-spec` field contains *format specifications* that define how the value should be presented, including such details as field width, alignment, padding, and decimal precision. Each type can define its own *formatting mini-language* or interpretation of the `format-spec` field. The syntax of format specifications is as follows:

```
format-spec      ::= std-format-spec | custom-format-spec
std-format-spec ::= [[fill] align] [sign] ['#'] ['0'] [width] ['.' precision] [type]
fill            ::= <a character other than '{' or '}'>
align           ::= '<' | '>' | '=' | '^'
sign            ::= '+' | '-' | ' '
width           ::= integer | '{' arg-id '}'
precision       ::= integer | '{' arg-id '}'
type            ::= 'a' | 'A' | 'b' | 'B' | 'c' | 'd' | 'e' | 'E' | 'f' | 'F' |
                  'g' | 'G' | 'n' | 'o' | 'p' | 's' | 'x' | 'X'
```

where `std-format-spec` defines a common formatting mini-language supported by built-in types and strings, while `custom-format-spec` is a placeholder for user-defined mini-languages. Some of the

formatting options are only supported by arithmetic types.

The `fill` character can be any character other than `'{'` or `'}'`. The presence of a fill character is signaled by the character following it, which must be one of the alignment options. If the second character of `format-spec` is not a valid alignment option, then it is assumed that both the fill character and the alignment option are absent.

The meaning of the various alignment options is as follows:

Option	Meaning
'<'	Forces the field to be left-aligned within the available space. This is the default for non-arithmetic types.
'>'	Forces the field to be right-aligned within the available space. This is the default for arithmetic types.
'='	Forces the padding to be placed after the sign (if any) but before the digits. This is used for printing fields in the form <code>+000000120</code> . This alignment option is only valid for arithmetic types.
'^'	Forces the field to be centered within the available space.

Note that unless a minimum field width is defined, the field width will be determined by the width of the content, meaning that the alignment option has no effect.

The `sign` option is only valid for arithmetic types, and can be one of the following:

Option	Meaning
'+'	Indicates that a sign should be used for both positive as well as negative numbers.
'-'	Indicates that a sign should be used only for negative numbers (this is the default behavior).
space	Indicates that a leading space should be used for positive numbers, and a minus sign for negative numbers.

The `'#'` option causes the *alternate form* to be used for the conversion. This option is only valid for integer and floating-point types. For integers, when binary, octal, or hexadecimal output is used, this option adds the respective prefix `"0b"` (`"0B"`), `"0"`, or `"0x"` (`"0X"`) to the output value. Whether the prefix is lower-case or upper-case is determined by the case of the type format specifier. For floating-point numbers the alternate form causes the result of the conversion to always contain a decimal-point character, even if no digits follow it. Normally, a decimal-point character appears in the result of these conversions only if a digit follows it. In addition, for `'g'` and `'G'` conversions, trailing zeros are not removed from the result.

`width` is a decimal integer defining the minimum field width. If not specified, then the field width will be determined by the content.

Preceding the `width` field by a zero (`'0'`) character enables sign-aware zero-padding for arithmetic types. This is equivalent to a `fill` character of `'0'` with an alignment type of `'='`.

The `precision` is a decimal number indicating how many digits should be displayed after the decimal point for a floating-point value formatted with `'f'` and `'F'`, or before and after the decimal point for a floating-point value formatted with `'g'` or `'G'`. For non-arithmetic types the field indicates the maximum field size - in other words, how many characters will be used from the field content. The `precision` is not allowed for integer, character, boolean, and pointer values.

Finally, the type determines how the data should be presented.

The available string presentation types are:

Type	Meaning
's'	String format. This is the default type for strings and may be omitted.
none	The same as <code>'s'</code> .

The available character presentation types are:

Type	Meaning
'c'	Character format. This is the default type for characters and may be omitted.
none	The same as <code>'c'</code> .

The available integer presentation types are:

Type Meaning

'b'	Binary format. Outputs the number in base 2. Using the '#' option with this type adds the prefix "0b" to the output value.
'B'	Binary format. Outputs the number in base 2. Using the '#' option with this type adds the prefix "0B" to the output value.
'd'	Decimal integer. Outputs the number in base 10.
'o'	Octal format. Outputs the number in base 8. Using the '#' option with this type adds the prefix "0" to the output value.
'x'	Hex format. Outputs the number in base 16, using lower-case letters for the digits above 9. Using the '#' option with this type adds the prefix "0x" to the output value.
'X'	Hex format. Outputs the number in base 16, using upper-case letters for the digits above 9. Using the '#' option with this type adds the prefix "0X" to the output value.
'n'	Number. This is the same as 'd' except that it uses the current locale to insert the appropriate number separator characters.
none	The same as 'd'.

Integer presentation types can also be used with character and boolean values. Boolean values are formatted using textual representation, either `true` or `false`, if the presentation type is not specified.

The available presentation types for floating-point values are:

Type Meaning

'a'	Hexadecimal floating point format. Prints the number in base 16 with prefix "0x" and lower-case letters for digits above 9. Uses 'p' to indicate the exponent.
'A'	Same as 'a' except it uses upper-case letters for the prefix, digits above 9 and to indicate the exponent.
'e'	Exponent notation. Prints the number in scientific notation using the letter 'e' to indicate the exponent.
'E'	Exponent notation. Same as 'e' except it uses an upper-case 'E' as the separator character.
'f'	Fixed point. Displays the number as a fixed-point number.
'F'	Fixed point. Same as 'f', but converts <code>nan</code> to <code>NAN</code> and <code>inf</code> to <code>INF</code> .
'g'	General format. For a given precision <code>p >= 1</code> , this rounds the number to <code>p</code> significant digits and then formats the result in either fixed-point format or in scientific notation, depending on its magnitude. A precision of 0 is treated as equivalent to a precision of 1.
'G'	General format. Same as 'g' except switches to 'E' if the number gets too large. The representations of infinity and NaN are uppercased, too.
'n'	Number. This is the same as 'g', except that it uses the locale to insert the appropriate number separator characters.
none	The same as 'g'.

The available presentation types for pointers are:

Type Meaning

'p'	Pointer format. This is the default type for pointers and may be omitted.
none	The same as 'p'.

Formatter [format.formatter]**Class template basic_parse_context**

```
namespace std {
    template<class charT>
    class basic_parse_context {
    public:
        using char_type = charT;
        using const_iterator = typename basic_string_view<charT>::const_iterator;
        using iterator = const_iterator;

        explicit constexpr basic_parse_context(basic_string_view<charT> fmt);

        constexpr const_iterator begin() const noexcept;
        constexpr const_iterator end() const noexcept;
        constexpr void advance_to(const_iterator it);
    };
}
```

```

    constexpr size_t next_arg_id();
    constexpr void check_arg_id(size_t id);
};
}

```

The class template `basic_parse_context` provides access the format string range being parsed and the argument counter for automatic indexing.

```
explicit constexpr basic_parse_context(basic_string_view<charT> fmt);
```

Effects: Constructs a parsing context from the format string `fmt`.

Ensures: `begin() == fmt.begin(), end() == fmt.end()`.

```
constexpr const_iterator begin() const noexcept;
```

Returns: An iterator referring to the first character in the format string range being parsed.

```
constexpr const_iterator end() const noexcept;
```

Returns: An iterator referring to the position one past the last character in the format string range being parsed.

```
void advance_to(iterator it);
```

Effects: Advances the beginning of the parsed format string range to `it`.

Requires: `end()` shall be reachable from `it`.

Ensures: `begin() == it`.

```
constexpr size_t next_arg_id();
```

Returns: The next argument identifier starting from 0. Each subsequent call to `next_arg_id` gives an identifier which is 1 greater than the one returned from the previous call for the same context object.

Throws: `format_error` if `check_arg_id` has been called earlier which indicates mixing of automatic and manual argument indexing.

```
constexpr void check_arg_id(size_t id);
```

Throws: `format_error` if `next_arg_id` has been called earlier which indicates mixing of automatic and manual argument indexing.

Class template `basic_format_context`

```

namespace std {
    template<class OutputIterator, class charT>
    class basic_format_context {
    public:
        using iterator = OutputIterator;
        using char_type = charT;
        using args_type = basic_format_args<basic_format_context>;

        template<class T>
            using formatter_type = formatter<T>;

        basic_parse_context<charT>& parse_context();

        iterator out();
        void advance_to(iterator it);

        args_type args() const;
    };
}

```

The class template `basic_format_context` provides access to format string parsing context, output iterator and formatting arguments.

```
basic_parse_context<charT>& parse_context();
```

Returns: A reference to the format string parsing context.

iterator out();

Returns: An iterator referring to the current position in the output range.

void advance_to(iterator it);

Effects: Advances the current position in the output range to it.

Ensures: out() == it.

args_type args() const;

Returns: A copy of the args_type object that provides access to formatting arguments.

Class template `formatter`

The functions defined in [\[format.functions\]](#) use specializations of the class template `formatter` to format individual arguments.

Each specialization of `formatter` is either enabled or disabled, as described below. [*Note:* Enabled specializations meet the `Formatter` requirements, and disabled specializations do not. — *end note*] Each header that declares the template `formatter` provides enabled specializations `formatter<const char*, char>`, `formatter<const wchar_t*, wchar_t>` as well as specializations of `formatter` for `nullptr_t`, `const void*`, and all cv-unqualified arithmetic types. For any type τ for which neither the library nor the user provides an explicit or partial specialization of the class template `formatter`, `formatter<T>` is disabled.

If the library provides an explicit or partial specialization of `formatter<T>`, that specialization is enabled except as noted otherwise.

If F is a disabled specialization of `formatter`, these values are false: `is_default_constructible_v<F>`, `is_copy_constructible_v<F>`, `is_move_constructible_v<F>`, `is_copy_assignable_v<F>`, `is_move_assignable_v<F>`.

An enabled specialization `formatter<T>` will satisfy the `Formatter` requirements (Table 1), with τ as the formatted argument type, the `DefaultConstructible` requirements, and the `CopyAssignable` requirements.

Formatter requirements

A type F meets the `Formatter` requirements if:

- it satisfies the `CopyConstructible` and `Destructible` requirements, and
- the expressions shown in Table 1 are valid and have the indicated semantics.

Given character type `charT`, output iterator type `OutputIterator`, and formatted argument type τ , in Table 1 f is a value of type F , u is an lvalue of type τ , t is a value of a type convertible to (possibly `const`) τ , pc is an lvalue of type `basic_parse_context<charT>` (denoted by PC), and fc is an lvalue of type `basic_format_context<OutputIterator, charT>` (denoted by FC). `pc.begin()` points to the beginning of the `format-spec` ([\[format.syntax\]](#)) portion of the format string. If `format-spec` is empty then either `pc.begin() == pc.end()` or `*pc.begin() == '}'`.

Table 1 — Formatter requirements

Expression	Return type	Requirement
<code>f.parse(pc)</code>	<code>PC::iterator</code>	Shall parse <code>format-spec</code> for type τ , store the parsed specifiers in <code>*this</code> , and return an iterator past the end of the parsed range.
<code>f.format(t, fc)</code>	<code>FC::iterator</code>	Shall format t according to the specifiers stored in <code>*this</code> , write the output to <code>fc.out()</code> and return an iterator past the end of the output range.
<code>f.format(u, fc)</code>	<code>FC::iterator</code>	Shall not modify u .

Related work

The Boost Format library [\[8\]](#) is an established formatting library that uses `printf`-like format string syntax with extensions. The main differences between this library and the current proposal are:

- Syntax: for the reasons described in section [Format String Syntax](#) this proposal uses a new syntax instead of extending the `printf` one. This allows much simpler and easier to parse grammar, not burdened by legacy specifiers used to convey type information. For example, Boost Format has two ways to refer to an argument by index and allows but ignores some format specifiers.
- API: Boost Format uses `operator%` to pass formatting arguments while this proposal uses variadic function templates.
- Performance: the implementation of this proposal is several times faster than the implementation of Boost Format on `tinyformat` benchmarks [\[9\]](#), generates smaller binary code and is faster to compile.

A `printf`-like Interface for the Streams Library [\[10\]](#) is similar to the Boost Format library but uses variadic templates instead of `operator%`. Unfortunately it hasn't been updated since 2013 and the same arguments about format string syntax apply to it.

The FastFormat library [\[11\]](#) is another well-known formatting library. Similarly to this proposal, FastFormat uses brace-delimited format specifiers, but otherwise the format string syntax is different and the library has significant limitations [\[12\]](#):

Three features that have no hope of being accommodated within the current design are:

- Leading zeros (or any other non-space padding)
- Octal/hexadecimal encoding
- Runtime width/alignment specification

Formatting facilities of the Folly library [\[13\]](#) are the closest to the current proposal. Folly also uses Python-like format string syntax nearly identical to the one described here. However, the API details are quite different. The current proposal tries to address performance and code bloat issues that are largely ignored by Folly Format. For instance formatting functions in Folly Format are parameterized on all argument types while in this proposal, only the inlined wrapper functions are, which results in much smaller binary code and better compile times.

Implementation

An implementation of this proposal is available in the `std` branch of the open-source `fmt` library [\[14\]](#).

Acknowledgements

Thanks to Beman Dawes, Bengt Gustafsson, Eric Niebler, Jason McKesson, Jeffrey Yasskin, Joël Lamotte, Lee Howes, Louis Dionne, Matt Clarke, Michael Park, Sergey Ignatchenko, Thiago Macieira, Zach Laine, Zhihao Yuan and participants of the Library Evolution Working Group for their feedback, support, constructive criticism and contributions to the proposal. Special thanks to Howard Hinnant who encouraged me to write the proposal and gave useful early advice on how to go about it.

The format string syntax is based on the Python documentation [\[3\]](#).

References

- [1] *The `fprintf` function*. ISO/IEC 9899:2011. 7.21.6.1.
- [2] [`fprintf`, `printf`, `snprintf`, `sprintf` - print formatted output](#). The Open Group Base Specifications Issue 6 IEEE Std 1003.1, 2004 Edition.
- [3] [6.1.3. Format String Syntax](#). Python 3.5.2 documentation.
- [4] [String.Format Method](#). .NET Framework Class Library.
- [5] [Module `std::fmt`](#). The Rust Standard Library.
- [6] [Format Specification Syntax: `printf` and `wprintf` Functions](#). C++ Language and Standard Libraries.
- [7] [10.4.2 Rearranging `printf` Arguments](#). The GNU Awk User's Guide.
- [8] [Boost Format library](#). Boost 1.63 documentation.
- [9] [Speed Test](#). The `fmt` library repository.
- [10] [A `printf`-like Interface for the Streams Library \(revision 1\)](#).
- [11] [The FastFormat library website](#).
- [12] [An Introduction to Fast Format \(Part 1\): The State of the Art](#). Overload Journal #89 - February 2009
- [13] [The folly library repository](#).
- [14] [The `fmt` library repository](#).
- [15] [P0424: Reconsidering literal operator templates for strings](#).
- [16] [D0355: Extending `<chrono>` to Calendars and Time Zones](#).

- [17] [P0067: Elementary string conversions.](#)
- [18] [MPark.Format: Compile-time Checked, Type-Safe Formatting in C++14.](#)
- [19] [Google Benchmark: A microbenchmark support library.](#)

Appendix A: Benchmarks

To demonstrate that the formatting functions described in this paper can be implemented efficiently, we compare the reference implementation [14] of `format` and `format_to` to `sprintf`, `ostringstream` and `to_string` on the following benchmark. This benchmark generates a set of integers with random numbers of digits, applies each method to convert each integer into a string (either `std::string` or a char buffer depending on the API) and uses the Google Benchmark library [19] to measure timings:

```
#include <algorithm>
#include <cmath>
#include <cstdio>
#include <limits>
#include <sstream>
#include <string>
#include <utility>
#include <vector>

#include <benchmark/benchmark.h>
#include <fmt/format.h>

// Returns a pair with the smallest and the largest value of integral type T
// with the given number of digits.
template<typename T>
std::pair<T, T> range(int num_digits) {
    T first = std::pow(T(10), num_digits - 1);
    int max_digits = std::numeric_limits<T>::digits10 + 1;
    T last = num_digits < max_digits ? first * 10 - 1 :
        std::numeric_limits<T>::max();
    return {num_digits > 1 ? first : 0, last};
}

// Generates values of integral type T with random number of digits.
template<typename T>
std::vector<T> generate_random_data(int numbers_per_digit) {
    int max_digits = std::numeric_limits<T>::digits10 + 1;
    std::vector<T> data;
    data.reserve(max_digits * numbers_per_digit);
    for (int i = 1; i <= max_digits; ++i) {
        auto r = range<T>(i);
        auto value = r.first;
        std::generate_n(std::back_inserter(data), numbers_per_digit, [=]() mutable {
            T result = value;
            value = value < r.second ? value + 1 : r.first;
            return result;
        });
    }
    std::random_shuffle(data.begin(), data.end());
    return data;
}

auto data = generate_random_data<int>(1000);

void sprintf(benchmark::State &s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            char buffer[12];
            result += std::sprintf(buffer, "%d", i);
        }
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(sprintf);

void ostringstream(benchmark::State &s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            std::ostringstream ss;
            ss << i;
            result += ss.str().size();
        }
    }
}
```

```

    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(ostringstream);

void to_string(benchmark::State &s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data)
            result += std::to_string(i).size();
        }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(to_string);

void format(benchmark::State &s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data)
            result += fmt::format("{} ", i).size();
        }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(format);

void format_to(benchmark::State &s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            char buffer[12];
            result += fmt::format_to(buffer, "{} ", i) - buffer;
        }
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(format_to);

BENCHMARK_MAIN();

```

The benchmark was compiled with clang (Apple LLVM version 9.0.0 clang-900.0.39.2) with `-O3 -DNDEBUG` and run on a macOS system. Below are the results:

```

Run on (4 X 3100 MHz CPU s)
2018-01-27 07:12:00
Benchmark      Time          CPU Iterations
-----
sprintf        882311 ns      881076 ns       781
ostringstream  2892035 ns     2888975 ns       242
to_string      1167422 ns     1166831 ns       610
format         675636 ns      674382 ns      1045
format_to      499376 ns      498996 ns      1263

```

The `format` and `format_to` functions show much better performance than the other methods. The `format` function that constructs `std::string` is even 30% faster than the system's version of `sprintf` that uses stack-allocated `char` buffer. `format_to` with a stack-allocated buffer is ~60% faster than `sprintf`.

Appendix B: Binary code comparison

In this section we compare per-call binary code size between the reference implementation that uses techniques described in section [Binary footprint](#) and standard formatting facilities. All the code snippets are compiled with clang (Apple LLVM version 9.0.0 clang-900.0.39.2) with `-O3 -DNDEBUG -c -std=c++14` and the resulted binaries are disassembled with `objdump -S`:

```

void consume(const char*);

void sprintf_test() {
    char buffer[100];
    sprintf(buffer, "The answer is %d.", 42);
    consume(buffer);
}

__Z12sprintf_testv:
    0:      55      pushq   %rbp
    1:      48 89 e5      movq    %rsp, %rbp
    4:      53      pushq   %rbx
    5:      48 83 ec 78      subq    $120, %rsp

```

```

9:      48 8b 05 00 00 00 00 movq    (%rip), %rax
10:     48 8b 00              movq    (%rax), %rax
13:     48 89 45 f0          movq    %rax, -16(%rbp)
17:     48 8d 35 37 00 00 00 leaq    55(%rip), %rsi
1e:     48 8d 5d 80          leaq    -128(%rbp), %rbx
22:     ba 2a 00 00 00 00 movl    $42, %edx
27:     31 c0              xorl    %eax, %eax
29:     48 89 df              movq    %rbx, %rdi
2c:     e8 00 00 00 00 00 callq   0 <__Z12sprintf_testv+0x31>
31:     48 89 df              movq    %rbx, %rdi
34:     e8 00 00 00 00 00 callq   0 <__Z12sprintf_testv+0x39>
39:     48 8b 05 00 00 00 00 movq    (%rip), %rax
40:     48 8b 00              movq    (%rax), %rax
43:     48 3b 45 f0          cmpq    -16(%rbp), %rax
47:     75 07              jne     7 <__Z12sprintf_testv+0x50>
49:     48 83 c4 78          addq    $120, %rsp
4d:     5b                  popq    %rbx
4e:     5d                  popq    %rbp
4f:     c3                  retq
50:     e8 00 00 00 00 00 callq   0 <__Z12sprintf_testv+0x55>

```

```

void format_test() {
    consume(format("The answer is {}.", 42).c_str());
}

```

```

__Z11format_testv:
0:      55              pushq    %rbp
1:      48 89 e5          movq    %rsp, %rbp
4:      53              pushq    %rbx
5:      48 83 ec 28        subq    $40, %rsp
9:      48 c7 45 d0 2a 00 00 00 movq    $42, -48(%rbp)
11:     48 8d 35 f4 83 01 00 leaq    99316(%rip), %rsi
18:     48 8d 7d e0          leaq    -32(%rbp), %rdi
1c:     4c 8d 45 d0          leaq    -48(%rbp), %r8
20:     ba 11 00 00 00 00 movl    $17, %edx
25:     b9 02 00 00 00 00 movl    $2, %ecx
2a:     e8 00 00 00 00 00 callq   0 <__Z11format_testv+0x2F>
2f:     f6 45 e0 01          testb   $1, -32(%rbp)
33:     48 8d 7d e1          leaq    -31(%rbp), %rdi
37:     48 0f 45 7d f0        cmovneq -16(%rbp), %rdi
3c:     e8 00 00 00 00 00 callq   0 <__Z11format_testv+0x41>
41:     f6 45 e0 01          testb   $1, -32(%rbp)
45:     74 09              je      9 <__Z11format_testv+0x50>
47:     48 8b 7d f0          movq    -16(%rbp), %rdi
4b:     e8 00 00 00 00 00 callq   0 <__Z11format_testv+0x50>
50:     48 83 c4 28          addq    $40, %rsp
54:     5b                  popq    %rbx
55:     5d                  popq    %rbp
56:     c3                  retq
57:     48 89 c3              movq    %rax, %rbx
5a:     f6 45 e0 01          testb   $1, -32(%rbp)
5e:     74 09              je      9 <__Z11format_testv+0x69>
60:     48 8b 7d f0          movq    -16(%rbp), %rdi
64:     e8 00 00 00 00 00 callq   0 <__Z11format_testv+0x69>
69:     48 89 df              movq    %rbx, %rdi
6c:     e8 00 00 00 00 00 callq   0 <__Z11format_testv+0x71>
71:     66 66 66 66 66 66 2e 0f 1f 84 00 00 00 00 00 00 nopw    %cs:(%rax,%rax)

```

```

void ostream_test() {
    std::ostream ss;
    ss << "The answer is " << 42 << ".";
    consume(ss.str().c_str());
}

```

```

__Z18ostream_testv:
0:      55              pushq    %rbp
1:      48 89 e5          movq    %rsp, %rbp
4:      41 57              pushq    %r15
6:      41 56              pushq    %r14
8:      41 55              pushq    %r13
a:      41 54              pushq    %r12
c:      53              pushq    %rbx
d:      48 81 ec 38 01 00 00 subq    $312, %rsp
14:     4c 8d b5 18 ff ff ff leaq    -232(%rbp), %r14
1b:     4c 8d a5 b0 fe ff ff leaq    -336(%rbp), %r12
22:     48 8b 05 00 00 00 00 movq    (%rip), %rax
29:     48 8d 48 18          leaq    24(%rax), %rcx
2d:     48 89 8d a8 fe ff ff movq    %rcx, -344(%rbp)
34:     48 83 c0 40          addq    $64, %rax

```

```

38: 48 89 85 18 ff ff ff movq    %rax, -232(%rbp)
3f: 4c 89 f7          movq    %r14, %rdi
42: 4c 89 e6          movq    %r12, %rsi
45: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x4A>
4a: 48 c7 45 a0 00 00 00 movq    $0, -96(%rbp)
52: c7 45 a8 ff ff ff ff movl    $4294967295, -88(%rbp)
59: 48 8b 1d 00 00 00 00 movq    (%rip), %rbx
60: 4c 8d 6b 18      leaq    24(%rbx), %r13
64: 4c 89 ad a8 fe ff ff movq    %r13, -344(%rbp)
6b: 48 83 c3 40      addq    $64, %rbx
6f: 48 89 9d 18 ff ff ff movq    %rbx, -232(%rbp)
76: 4c 89 e7          movq    %r12, %rdi
79: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x7E>
7e: 4c 8b 3d 00 00 00 00 movq    (%rip), %r15
85: 49 83 c7 10      addq    $16, %r15
89: 4c 89 bd b0 fe ff ff movq    %r15, -336(%rbp)
90: 48 c7 85 08 ff ff ff 00 00 00 00 movq    $0, -248(%rbp)
9b: 48 c7 85 00 ff ff ff 00 00 00 00 movq    $0, -256(%rbp)
a6: 48 c7 85 f8 fe ff ff 00 00 00 00 movq    $0, -264(%rbp)
b1: 48 c7 85 f0 fe ff ff 00 00 00 00 movq    $0, -272(%rbp)
bc: c7 85 10 ff ff ff 10 00 00 00 movl    $16, -240(%rbp)
c6: 0f 57 c0        xorps   %xmm0, %xmm0
c9: 0f 29 45 b0      movaps  %xmm0, -80(%rbp)
cd: 48 c7 45 c0 00 00 00 00 movq    $0, -64(%rbp)
d5: 48 8d 75 b0      leaq    -80(%rbp), %rsi
d9: 4c 89 e7          movq    %r12, %rdi
dc: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0xE1>
e1: f6 45 b0 01      testb   $1, -80(%rbp)
e5: 74 09          je      9 <__Z18ostream_testv+0xF0>
e7: 48 8b 7d c0      movq    -64(%rbp), %rdi
eb: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0xF0>
f0: 48 8d 35 dd 10 00 00 leaq    4317(%rip), %rsi
f7: 48 8d bd a8 fe ff ff leaq    -344(%rbp), %rdi
fe: ba 0e 00 00 00 movl    $14, %edx
103: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x108>
108: be 2a 00 00 00 movl    $42, %esi
10d: 48 89 c7          movq    %rax, %rdi
110: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x115>
115: 48 8d 35 c7 10 00 00 leaq    4295(%rip), %rsi
11c: ba 01 00 00 00 movl    $1, %edx
121: 48 89 c7          movq    %rax, %rdi
124: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x129>
129: 48 8d 7d b0      leaq    -80(%rbp), %rdi
12d: 4c 89 e6          movq    %r12, %rsi
130: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x135>
135: f6 45 b0 01      testb   $1, -80(%rbp)
139: 48 8d 7d b1      leaq    -79(%rbp), %rdi
13d: 48 0f 45 7d c0   cmovneq -64(%rbp), %rdi
142: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x147>
147: f6 45 b0 01      testb   $1, -80(%rbp)
14b: 74 09          je      9 <__Z18ostream_testv+0x156>
14d: 48 8b 7d c0      movq    -64(%rbp), %rdi
151: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x156>
156: 4c 89 ad a8 fe ff ff movq    %r13, -344(%rbp)
15d: 48 89 9d 18 ff ff ff movq    %rbx, -232(%rbp)
164: 4c 89 bd b0 fe ff ff movq    %r15, -336(%rbp)
16b: f6 85 f0 fe ff ff 01 testb   $1, -272(%rbp)
172: 74 0c          je      12 <__Z18ostream_testv+0x180>
174: 48 8b bd 00 ff ff ff movq    -256(%rbp), %rdi
17b: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x180>
180: 4c 89 e7          movq    %r12, %rdi
183: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x188>
188: 48 8b 35 00 00 00 00 movq    (%rip), %rsi
18f: 48 83 c6 08      addq    $8, %rsi
193: 48 8d bd a8 fe ff ff leaq    -344(%rbp), %rdi
19a: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x19F>
19f: 4c 89 f7          movq    %r14, %rdi
1a2: e8 00 00 00 00 callq   0 <__Z18ostream_testv+0x1A7>
1a7: 48 81 c4 38 01 00 00 addq    $312, %rsp
1ae: 5b              popq    %rbx
1af: 41 5c          popq    %r12
1b1: 41 5d          popq    %r13
1b3: 41 5e          popq    %r14
1b5: 41 5f          popq    %r15
1b7: 5d              popq    %rbp
1b8: c3              retq
1b9: 48 89 45 d0      movq    %rax, -48(%rbp)
1bd: f6 45 b0 01      testb   $1, -80(%rbp)
1c1: 74 3b          je      59 <__Z18ostream_testv+0x1FE>

```

```

1c3: 48 8b 7d c0    movq    -64(%rbp), %rdi
1c7: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x1CC>
1cc: eb 30         jmp     48 <__Z18ostream_testv+0x1FE>
1ce: eb 2a         jmp     42 <__Z18ostream_testv+0x1FA>
1d0: 48 89 45 d0    movq    %rax, -48(%rbp)
1d4: f6 45 b0 01    testb   $1, -80(%rbp)
1d8: 74 39         je      57 <__Z18ostream_testv+0x213>
1da: 48 8b 7d c0    movq    -64(%rbp), %rdi
1de: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x1E3>
1e3: f6 85 f0 fe ff ff 01 testb   $1, -272(%rbp)
1ea: 75 30         jne     48 <__Z18ostream_testv+0x21C>
1ec: eb 3a         jmp     58 <__Z18ostream_testv+0x228>
1ee: 48 89 45 d0    movq    %rax, -48(%rbp)
1f2: eb 3c         jmp     60 <__Z18ostream_testv+0x230>
1f4: 48 89 45 d0    movq    %rax, -48(%rbp)
1f8: eb 4d         jmp     77 <__Z18ostream_testv+0x247>
1fa: 48 89 45 d0    movq    %rax, -48(%rbp)
1fe: 4c 89 ad a8 fe ff ff movq    %r13, -344(%rbp)
205: 48 89 9d 18 ff ff ff movq    %rbx, -232(%rbp)
20c: 4c 89 bd b0 fe ff ff movq    %r15, -336(%rbp)
213: f6 85 f0 fe ff ff 01 testb   $1, -272(%rbp)
21a: 74 0c         je      12 <__Z18ostream_testv+0x228>
21c: 48 8b bd 00 ff ff ff movq    -256(%rbp), %rdi
223: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x228>
228: 4c 89 e7       movq    %r12, %rdi
22b: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x230>
230: 48 8b 35 00 00 00 00 movq    (%rip), %rsi
237: 48 83 c6 08    addq    $8, %rsi
23b: 48 8d bd a8 fe ff ff leaq     -344(%rbp), %rdi
242: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x247>
247: 4c 89 f7       movq    %r14, %rdi
24a: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x24F>
24f: 48 8b 7d d0    movq    -48(%rbp), %rdi
253: e8 00 00 00 00 callq    0 <__Z18ostream_testv+0x258>
258: 0f 1f 84 00 00 00 00 nopl     (%rax,%rax)

```

The code generated for the `format_test` function that uses the reference implementation of the `format` function described in this proposal is several times smaller than the `ostream` code and only 40% larger than the one generated for `sprintf` which is a moderate price to pay for full type and memory safety.

The following factors contribute to the difference in binary code size between `format` and `sprintf`:

- Passing format string as `string_view` instead of `const char*`.
- Using `string` instead of a `char` buffer.
- Preparing the array of formatting arguments.

We can exclude the first two factors from the experiment by mimicking parts of the `sprintf` API:

```
int vraw_format(char* buffer, const char* format, format_args args);
```

```
template<typename... Args>
```

```
inline int raw_format(char* buffer, const char* format, const Args&... args) {
    return vraw_format(buffer, format, make_format_args(args...));
}
```

```
void raw_format_test() {
    char buffer[100];
    raw_format(buffer, "The answer is {}.", 42);
}
```

```
__Z15raw_format_testv:
```

```

0: 55            pushq   %rbp
1: 48 89 e5      movq    %rsp, %rbp
4: 48 81 ec 80 00 00 00 subq    $128, %rsp
b: 48 8b 05 00 00 00 00 movq    (%rip), %rax
12: 48 8b 00      movq    (%rax), %rax
15: 48 89 45 f8    movq    %rax, -8(%rbp)
19: 48 c7 45 80 2a 00 00 00 movq    $42, -128(%rbp)
21: 48 8d 35 24 12 00 00 leaq     4644(%rip), %rsi
28: 48 8d 7d 90    leaq     -112(%rbp), %rdi
2c: 48 8d 4d 80    leaq     -128(%rbp), %rcx
30: ba 02 00 00 00 movl     $2, %edx
35: e8 00 00 00 00 callq    0 <__Z15raw_format_testv+0x3A>
3a: 48 8b 05 00 00 00 00 movq    (%rip), %rax
41: 48 8b 00      movq    (%rax), %rax
44: 48 3b 45 f8    cmpq    -8(%rbp), %rax

```

```
48:      75 09    jne     9 <__Z15raw_format_testv+0x53>
4a:      48 81 c4 80 00 00 00    addq    $128, %rsp
51:      5d      popq    %rbp
52:      c3      retq
53:      e8 00 00 00 00    callq   0 <__Z15raw_format_testv+0x58>
58:      0f 1f 84 00 00 00 00 00    nopl     (%rax,%rax)
```

This shows that passing formatting arguments adds very little overhead and is comparable with `printf`.