

Text Formatting

☐ Show only wording and recent revision history

1. [Revision History](#)
 1. [Changes since R8](#)
 2. [Changes since R7](#)
 3. [Changes since R6](#)
 4. [Changes since R5](#)
 5. [Changes since R4](#)
 6. [Changes since R3](#)
 7. [Changes since R2](#)
 8. [Changes since R1](#)
 9. [Changes since R0](#)
2. [Introduction](#)
3. [Design](#)
 1. [Format string syntax](#)
 2. [Extensibility](#)
 3. [Safety](#)
 4. [Locale support](#)
 5. [Positional arguments](#)
 6. [Performance](#)
 7. [Binary footprint](#)
 8. [Format strings and char traits](#)
 9. [Compile-time processing of format strings](#)
 10. [Argument visitation](#)
 11. [Impact on existing code](#)
4. [Proposed wording](#)
 1. [Editing notes](#)
 2. [19.❖ Formatting \[format\]](#)
 1. [19.❖.1 Header <format> synopsis \[format.syn\]](#)
 2. [19.❖.2 Format string \[format.string\]](#)
 3. [19.❖.3 Formatting functions \[format.functions\]](#)
 4. [19.❖.4 Formatter \[format.formatter\]](#)
 1. [19.❖.4.1 Formatter requirements \[formatter.requirements\]](#)
 2. [19.❖.4.2 Class template basic_format_parse_context \[format.parse_context\]](#)
 3. [19.❖.4.3 Class template basic_format_context \[format.context\]](#)
 5. [19.❖.5 Arguments \[format.arguments\]](#)
 1. [19.❖.5.1 Class template basic_format_arg \[format.arg\]](#)
 2. [19.❖.5.2 Argument visitation \[format.visit\]](#)
 3. [19.❖.5.3 Class template format_arg_store \[format.arg_store\]](#)
 4. [19.❖.5.4 Function template make_format_args \[format.make_args\]](#)
 5. [19.❖.5.5 Function template make_wformat_args \[format.make_wargs\]](#)
 6. [19.❖.5.6 Class template basic_format_args \[format.basic_args\]](#)
 6. [19.❖.6 Class format_error \[format.error\]](#)
5. [Related work](#)
6. [Implementation](#)
7. [Acknowledgements](#)
8. [References](#)
9. [Appendix A: Benchmarks](#)
10. [Appendix B: Binary code comparison](#)

Revision History

Changes since R8

- Handle the case of undefined `uintptr_t` in [\[format.string\]](#).
- Clarify that digit separator characters include the decimal radix separator in [\[format.string\]](#).
- Clarify that a format string is a possibly empty sequence in [\[format.string\]](#).
- Add the default value for `num_args` in [basic_format_parse_context](#) to match the declaration in synopsis.
- Add a *Remarks* clause to [basic_format_parse_context::check_arg_id](#).
- Add missing forward in [\[format.visit\]](#).
- Fix namespace qualification in example in [\[format.context\]](#).
- Replace *Requires* with *Constraints*, *Mandates*, and *Expects* as appropriate in [vformat_to](#), [format_to_n](#), [formatted_size](#), and [make_format_args](#) per [p1369](#).
- Replace "satisfies the *X* requirements" with "meets the *X* requirements" when referring to old-style requirements throughout the paper.
- Replace the undocumented `Formattable` requirement with the requirement that `formatter<T, charT>` is enabled in [format_to_n](#), [formatted_size](#), and [make_format_args](#).
- Replace "A format string consists" with "A format string is a sequence" in [\[format.string\]](#).
- Replace "A format string is invalid if it contains a mix of automatic and manual indexing." with "A format string does not contain a mixture of automatic and manual indexing." in [\[format.string\]](#).
- Replace "Places the character representation" with "Places the first *M* characters of the character representation" in the specifications of [format_to_n](#).
- Add missing `std` namespace qualifications in the `formatter` specialization example in [\[format.formatter\]](#).
- Add a colon after "provides enabled specializations" in [\[format.formatter\]](#).
- Add a missing `charT` template argument to `formatter` in the definition of `basic_format_context::formatter_type` in [\[format.context\]](#).
- Replace "Parses the range `[pc.begin(), pc.end())` according to format specifications for type `T` until the first unmatched character" with "Parses `format-spec` ([\[format.string\]](#)) for type `T` in the range `[pc.begin(), pc.end())` until the first unmatched character" in [Table 4 — Formatter requirements](#).
- Remove "format string" from "format string range `[pc.begin(), pc.end())`" in [Table 4 — Formatter requirements](#).
- Replace "Does not modify `u`" with "As above but does not modify `u`" in [Table 4 — Formatter requirements](#).
- Clarify "`id` is a valid argument `id`" in the *Expects* clause in [basic_format_parse_context::check_arg_id](#) using a new exposition-only `num_args_member`.
- Fix a typo: "back_insert_iterato" -> "back_insert_iterator" in [\[format.context\]](#).
- Rephrase the note in [\[format.context\]](#) to clarify that it applies to types other than `char` too.
- Replace "*Requires*: `Out` models `OutputIterator<const charT&>`" with "`Out` shall model `OutputIterator<const charT&>`" in [\[format.context\]](#).
- Change the type of `width_arg_id` from `int` to `size_t` in the example in [\[format.context\]](#).
- Make exposition-only code private and add exposition-only friends in [\[format.arg\]](#) for the code to remain well-formed.
- Disallow specializations of `basic_format_arg` ([\[format.arg\]](#)).
- Merge integral and unconstrained overloads of the `basic_format_arg` constructor to support specializations of `formatter` for extended integer types larger than `long long int`.
- Replace "a standard (un)signed integer type or an extended (un)signed integer type" with "a(n) (un)signed integer type" in [basic_format_arg\(const T&\)](#).
- Add "otherwise" to bullet points in [basic_format_arg\(const T&\)](#).
- Replace "the program is ill-formed" bullet point with "*Constraints*: `typename Context::template formatter_type<T>` is enabled" in [basic_format_arg\(const T&\)](#).
- Move the description of class `basic_format_arg::handle` before the definition.
- Replace `&val` with `addressof(val)` in the `basic_format_arg::handle` constructor.
- Simplify `make_wformat_args` by specifying it in terms of `make_format_args`.
- Move sections [\[format.make_args\]](#) and [\[format.make_wargs\]](#) next to section [\[format.arg_store\]](#).

Changes since R7

- Suggest placing 19.4 Formatting as a subclause after 19.19 Primitive numeric conversions in the ["General utilities library summary" table](#).
- Replace concept-related annotations with *Requires* clauses per [P1369 Guidelines for Formulating Library Semantics Specifications](#).
- Rename `o` to `out` in template parameters to make them distinguishable from `o`.

- Replace "Returns:" with "Effects: Equivalent to:" in [format](#) and [format_to](#) specifications to propagate *Throws*.
- Merge the two [format_to](#) specifications into one.
- Add commas around "formatted according to specifications given in *fmt*" in the specifications of [vformat_to](#) and [format_to_n](#).
- Replace "*N* is/be the formatted output size" with $N = \text{formatted_size}(\text{fmt}, \text{args} \dots)$ in the specifications of [vformat_to](#) and [format_to_n](#).
- Denote $\min(\max(n, 0), N)$ by *M* in the specification of [format_to_n](#) since *N* is already taken.
- Require that each τ_i in *Args* satisfies the *Formattable* requirements in the specifications of [format_to_n](#), [formatted_size](#), [make_format_args](#), and [make_wformat_args](#). Other functions get this transitively via `make_(w)format_args`.
- Move the format string specification to the new section [\[format.string\]](#).
- Remove the word "valid" when referring to a format string.
- Replace "The *fmt* string" with "A format string" in [\[format.string\]](#).
- Replace "by" with "for" in "Some of the formatting options are only supported by arithmetic types."
- Clarify that padding is placed after the prefix if any when using the numeric alignment and add an example.
- Rename "number separator" to "digit separator".
- Added a note that pointer presentation types also apply to `nullptr_t`.
- Replace "where *charT* is `char` or `wchar_t`" with "Let *charT* be either `char` or `wchar_t`" in [\[format.formatter\]](#).
- Clarify that infinity and NaN are formatted as "inf" and "nan" for non-uppercase specifiers in [\[format.string\]](#).
- Reword the definitions of *PC*, *FC*, *pc*, and *fc* in [\[formatter.requirements\]](#).
- Remove several occurrences of "shall" in [Table 4 — Formatter requirements](#).
- Replace *Returns* and *Remarks* with *Effects* in [\[format.visit\]](#).
- Use kebab case for exposition-only types renaming `format_arg_store` to *format-arg-store* and `basic_format_arg::char_type` to `basic_format_arg::char-type`.
- Add a description of *format-arg-store* in [\[format.arg.store\]](#).
- Clarify what can be used as an iterator in `format_context` in [\[format.context\]](#).
- Remove unnecessary `std` qualifications.
- Remove `noexcept` from [basic_format_arg\(const char_type* s\)](#) because it has a precondition.

Changes since R6

- Rebase wording onto the pre-Kona C++ working draft [N4800](#).
- Use placeholder for section and table numbers and drop "utilities" from the section name.
- Clarify the parse wording in [Table 4 — Formatter requirements](#).
- Add an entry to the "General utilities library summary" table.
- Replace overload resolution *Remarks* with *Constraints* per [p1369](#).
- Remove "shall" from *Expects* wording.
- Add Kona LEWG poll results for R6.
- Minor improvements to the wording and editing notes.

Changes since R5

- Separate parsing context from formatting context and remove `basic_format_context::parse_context()`.
- Require that `id` in `basic_format_parse_context::check_arg_id` is a valid argument id.
- Replace *Requires* with *Expects* per [p1369](#).

[LEWG polls \(R6\):](#)

Remove `basic_format_context<>::parse_context()`.

SF	F	N	A	SA
7	2	0	0	0

Unanimous consent to all other design changes presented.

Forward back to LWG as written for C++20.

SF	F	N	A	SA
8	4	0	0	0

Changes since [R4](#)

- Clarify that `s` in `basic_format_arg(const char_type* s)` shall point to a NTCTS.
- Remove `is_arithmetic` and `is_integral`.
- Clarify that the formatter output should be deterministic in [Table 4 — Formatter requirements](#).
- Minor wording corrections.

Changes since [R3](#)

- Use `OutputIterator` concept in formatting functions that take output iterators and replace the `size` template parameter with `iter_difference_t`.
- Rename `*parse_context` to `*format_parse_context`.
- Disallow consecutive zeros in `arg-id` and `width`.
- Specify formatting of arithmetic and pointer types in terms of `to_chars`.
- Clarify what is the current locale.
- Use the shortest round-trip format as the default floating-point formatting.
- Add missing specializations of `formatter`.
- Clarify that the `sign` option applies to floating-point infinity and NaN.
- Propose adding a new entry `<format>` to table 18 in section 15.5.1.2 [headers].
- Add section numbers in Proposed wording.
- Add a feature test macro.
- Remove tautological `Ensures` clauses.
- Merge two `make_format_args` overloads.
- Replace `args()` with `arg(size_t)` in `basic_format_context`.
- Make `format_arg_store` exposition-only.

Changes since [R2](#)

- Add section [Argument visitation](#) with an example of how to use the visitation API to implement dynamic format specifiers.
- Rename `visit` to `visit_format_arg` to distinguish from variant's `visit`.
- Merge `[format.syntax]` into `[format.functions]`.
- Remove the restriction that `'\0'` cannot be used as a fill character.
- Replace `Postconditions` with `Ensures`.

Changes since [R1](#)

- Add the `format_to_n` function taking an output iterator and a size.
- Add a note that compile-time processing of format strings applies to user-defined types with custom parsers to section [Compile-time processing of format strings](#).
- Rename `count` to `formatted_size`.
- Drop nested namespace `fmt`.
- Add `format` prefix or infix to class and function names to avoid potential name collision after removing the nested namespace.
- Improve wording.
- Replace the requirement of implementing a formatter via ostream insertion operator if the latter is provided with a note that it can be implemented in such way.
- Expand the Acknowledgements section and remove "Reply to".

Changes since [R0](#)

- Add section [Compile-time processing of format strings](#).
- Separate parsing and formatting in the extension API replacing `format_value` function template with class template `formatter` to allow compile-time processing of format strings.
- Change return type of `format_to` and `vformat_to` to `OutputIterator` in synopsis.
- Remove sections Null-terminated string view and Format string, and replace `basic_cstring_view` with `basic_string_view`.
- Add a link to the implementation in [Introduction](#).
- Add a note regarding time formatting and compatibility with D0355 "Extending `<chrono>` to Calendars and Time Zones" [\[16\]](#) to section [Extensibility](#).
- Rename `basic_args` to `basic_format_args`.
- Rename `is_numeric` to `is_arithmetic`.

- Add the `count` function that counts the number of characters and use it to define output ranges.
- Remove `basic_buffer` and section Formatting buffer and replace buffers with output iterators.
- Add [Appendix A: Benchmarks](#).
- Explain the purpose of the type-erased API in more details in the [Binary footprint](#) section.
- Add [Appendix B: Binary code comparison](#).
- Add formatting function overloads for `wchar_t` strings.

Introduction

Even with proliferation of graphical and voice user interfaces, text remains one of the main ways for humans to interact with computer programs and programming languages provide a variety of methods to perform text formatting. The first thing we do when learning a new programming language is often write a "Hello, World!" program that performs simple formatted output.

C++ has not one but two standard APIs for producing formatted output, the `printf` family of functions inherited from C and the I/O streams library (`iostreams`). While `iostreams` are usually the recommended way of producing formatted output in C++ for safety and extensibility reasons, `printf` offers some advantages, such as an arguably more natural function call API, the separation of formatted message and arguments, possibly with argument reordering as a POSIX extension, and often more compact source and binary code.

This paper proposes a new text formatting library that can be used as a safe and extensible alternative to the `printf` family of functions. It is intended to complement the existing C++ I/O streams library and reuse some of its infrastructure such as overloaded insertion operators for user-defined types.

Example:

```
string message = format("The answer is {}. ", 42);
```

A full implementation of this proposal is available at <https://github.com/fmtlib/fmt/tree/std>.

Design

Format string syntax

Variations of the `printf` format string syntax are arguably the most popular among the programming languages and C++ itself inherits `printf` from C [1]. The advantage of the `printf` syntax is that many programmers are familiar with it. However, in its current form it has a number of issues:

- Many format specifiers like `hh`, `h`, `l`, `j`, etc. are used only to convey type information. They are redundant in type-safe formatting and would unnecessarily complicate specification and parsing.
- There is no standard way to extend the syntax for user-defined types.
- There are subtle differences between different implementations. For example, POSIX positional arguments [2] are not supported on some systems [6].
- Using `'%'` in a custom format specifier poses difficulties, e.g. for `put_time`-like time formatting.

Although it is possible to address these issues while maintaining resemblance to the original `printf` format, this will still break compatibility and can potentially be more confusing to users than introducing a different syntax.

Therefore we propose a new syntax based on the ones used in Python [3], the .NET family of languages [4], and Rust [5]. This syntax employs `'{'` and `'}'` as replacement field delimiters instead of `'%'` and it is described in detail in [\[format.functions\]](#). Some advantages of the proposal are:

- A consistent and easy to parse mini-language focused on formatting rather than conveying type information
- Extensibility and support for custom format strings for user-defined types
- Positional arguments
- Support for both locale-specific and locale-independent formatting (see [Locale support](#))
- Formatting improvements such as better alignment control, fill character, and binary representation

The syntax is expressive enough to enable translation, possibly automated, of most `printf` format strings. The correspondence between `printf` and the new syntax is given in the following table:

printf	new
-	<
+	+
<i>space</i>	<i>space</i>
#	#
0	0
hh	unused
h	unused
l	unused
ll	unused
j	unused
z	unused
t	unused
L	unused
c	c (optional)
s	s (optional)
d	d (optional)
i	d (optional)
o	o
x	x
X	X
u	d (optional)
f	f
F	F
e	e
E	E
a	a
A	A
g	g (optional)
G	G
n	unused
p	p (optional)

Width and precision are represented similarly in `printf` and the proposed syntax with the only difference that runtime value is specified by '*' in the former and '{ }' in the latter, possibly with the index of the argument inside the braces:

```
printf("%*s", 10, "foo");
format("{:{} }", "foo", 10);
```

As can be seen from the table above, most of the specifiers remain the same which simplifies migration from `printf`. A notable difference is in the alignment specification. The proposed syntax allows left, center, and right alignment represented by '<', '^', and '>' respectively which is more expressive than the corresponding `printf` syntax. The latter only supports left and right alignment.

The following example uses center alignment and '*' as a fill character:

```
format("{:^30}", "centered");
```

resulting in "*****centered*****". The same formatting cannot be easily achieved with `printf`.

In addition to positional arguments, the grammar can be easily extended to support named arguments.

Extensibility

Both the format string syntax and the API are designed with extensibility in mind. The mini-language can be extended for user-defined types and users can provide functions that implement parsing, possibly at compile time, and formatting for such types.

The general syntax of a replacement field in a format string is

```
replacement-field ::= '{' [arg-id] [':' format-spec] '}'
```

where `format-spec` is predefined for built-in types, but can be customized for user-defined types. For example, the syntax can be extended for `put_time`-like date and time formatting

```
time_t t = time(nullptr);
string date = format("The date is {0:%Y-%m-%d}.", *localtime(&t));
```

by providing a specialization of `formatter` for `tm`:

```
template<>
struct formatter<tm> {
    constexpr format_parse_context::iterator parse(format_parse_context& ctx);

    template<class FormatContext>
    typename FormatContext::iterator format(const tm& tm, FormatContext& ctx);
};
```

The `formatter<tm>::parse` function parses the `format-spec` portion of the format string corresponding to the current argument and `formatter<tm>::format` formats the value and writes the output via the iterator `ctx.begin()`.

Note that date and time formatting is not covered by this proposal but formatting facilities provided by D0355 "Extending `<chrono>` to Calendars and Time Zones" [\[16\]](#) can be easily implemented using this extension API.

An implementation of `formatter<T>::format` can use ostream insertion `operator<<` for user-defined type `T` if available.

The extension API is based on specialization instead of the argument-dependent lookup (ADL), because the `parse` function doesn't take the object to be formatted as an argument and therefore some other way of parameterizing it on the argument type `T` such as introducing a dummy argument has to be used, e.g.

```
constexpr auto parse(type<T>, format_parse_context& ctx);
```

Another problem with ADL-based approach is compile-time performance as pointed out in [\[20\]](#):

Overload resolution on `operator<<` tends to get expensive for larger projects with hundreds or thousands of candidates in the overload set. This seems hard to resolve, since choosing a different name for `operator<<` simply shifts the expense of overload resolution to a differently-named function.

Safety

Formatting functions rely on variadic templates instead of the mechanism provided by `<csdarg>`. The type information is captured automatically and passed to formatters guaranteeing type safety and making many of the `printf` specifiers redundant (see [Format String Syntax](#)). Memory management is automatic to prevent buffer overflow errors common to `printf`.

Locale support

As pointed out in P0067 "Elementary string conversions" [\[17\]](#) there is a number of use cases that do not require internationalization support, but do require high throughput when produced by a server. These include various text-based interchange formats such as JSON or XML. The need for locale-independent functions for conversions between integers and strings and between floating-point numbers and strings has also been highlighted in [\[20\]](#). Therefore a user should be able to easily control whether to use locales or not during formatting.

We follow Python's approach [\[3\]](#) and designate a separate format specifier `'n'` for locale-aware numeric formatting. It applies to all integral and floating-point types. All other specifiers produce output unaffected by locale settings. This can also have positive effect on performance because locale-independent formatting can be implemented more efficiently.

Positional arguments

An important feature for localization is the ability to rearrange formatting arguments as the word order may vary in different languages [\[7\]](#). For example:

```
printf("String '%s' has %d characters\n", string, length(string));
```


A possible German translation of the format string might be:

```
"%2$d Zeichen lang ist die Zeichenkette '%1$s'\n"
```

using POSIX positional arguments [2]. Unfortunately these positional specifiers are not portable [6]. The C++ I/O streams don't support such rearranging of arguments by design because they are interleaved with the portions of the literal string:

```
cout << "String `" << string << "' has " << length(string) << " characters\n";
```

The current proposal allows both positional and automatically numbered arguments, for example:

```
format("String `{}`' has {} characters\n", string, length(string));
```

with the German translation of the format string being:

```
"{1} Zeichen lang ist die Zeichenkette `{0}'\n"
```

Performance

The formatting library has been designed with performance in mind. It tries to minimize the number of virtual function calls and dynamic memory allocations done per formatting operation. In particular, if formatting output can fit into a fixed-size array allocated on stack, it should be possible to avoid both of them altogether by using a suitable API.

The `format_to` function takes an arbitrary output iterator and, for performance reasons, can be specialized for random-access and contiguous iterators as shown in the reference implementation [14].

The locale-independent formatting can also be implemented more efficiently than the locale-aware one. However, the main goal for the former is to support specific use cases (see [Locale support](#)) rather than to improve performance.

See [Appendix A: Benchmarks](#) for a small performance comparison of the reference implementation of this proposal versus the standard formatting facilities.

Binary footprint

In order to minimize binary code size each formatting function that uses variadic templates can be implemented as a small inline wrapper around its non-variadic counterpart. This wrapper creates a `basic_format_args` object, representing an array of type-erased argument references, with `make_format_args` and calls the non-variadic function to do the actual work. For example, the `format` variadic function calls `vformat`:

```
string vformat(string_view fmt, format_args args);

template<class... Args>
inline string format(string_view fmt, const Args&... args) {
    return vformat(fmt, make_format_args(args...));
}
```

`basic_format_args` can be implemented as an array of tagged unions. If the number of arguments is small then the tags that indicate the arguments types can be combined and passed into a formatting function as a single integer. This single integer representing all argument types is computed at compile time and only needs to be stored, resulting in smaller binary code.

Given a reasonable optimizing compiler, this will result in a compact per-call binary code, effectively consisting of placing argument pointers (or, possibly, copies for primitive types) and packed tags on the stack and calling a formatting function. See [Appendix B: Binary code comparison](#) for a specific example.

Exposing the type-erased API rather than making it an implementation detail and only providing variadic functions allows applying the same technique to the user code. For example, consider a simple logging function that writes a formatted error message to `clog`:

```
template<class... Args>
void log_error(int code, string_view fmt, const Args&... args) {
    clog << "Error " << code << ": " << format(fmt, args...);
}
```

The code for this function will be generated for every combination of argument types which can be undesirable. However, if we use the type-erased API, there will be only one instance of the logging

function while the wrapper function can be trivially inlined:

```
void vlog_error(int code, string_view fmt, format_args args) {
    clog << "Error " << code << ": " << vformat(fmt, args);
}

template<class... Args>
inline void log_error(int code, string_view fmt, const Args&... args) {
    vlog_error(code, fmt, make_format_args(args...));
}
```

The current design allows users to easily switch between the two approaches.

For similar reasons formatting functions take format strings as instances of `basic_string_view` instead of being parameterized on a string type. This allows implicit conversions from null-terminated character strings and `basic_string` without affecting code size.

Format strings and `char_traits`

Format strings intentionally use the default character traits in the formatting functions API because making traits customizable creates potential problems with handling format specifications.

Consider the following example:

```
using ci_string_view = std::basic_string_view<char, ci_char_traits>;
auto s = format(ci_string_view("{:X}"), 0xCAFE);
```

where `ci_char_traits` is a character trait with case-insensitive comparison.

Since the format specification grammar is case-sensitive, the behavior of such code is no longer intuitive. There are several options:

1. ill-formed (current proposal),
2. well-formed and `s == "cafe"`,
3. well-formed and `s == "CAFE"`.

Option 2 is a potential pessimization because comparisons will have to go through traits while option 3 effectively ignores traits.

Existing standard functions that take format strings (`printf`, `put_time`, `strftime`) do not allow passing character traits.

Compile-time processing of format strings

It is possible to parse format strings at compile time with `constexpr` functions which has been demonstrated in the reference implementation [\[14\]](#) and in [\[18\]](#). Unfortunately a satisfactory API cannot be provided using existing C++17 features. Ideally we would like to use a function call API similar to the one proposed in this paper:

```
template<class String, class... Args>
string format(String fmt, const Args&... args);
```

where `String` is some type representing a compile-time format string, for example

```
struct MyString {
    static constexpr string_view value() { return string_view("{}", 2); }
};
```

However, requiring a user to create a new type either manually or via a macro for every format string is not acceptable. P0424R1 "Reconsidering literal operator templates for strings" [\[15\]](#) provides a solution to this problem based on user-defined literal operators. If this or other similar proposal for compile-time strings is accepted into the standard, it should be easy to provide additional formatting APIs that make use of this feature. Obviously runtime checks will still be needed in cases where the format string is not known at compile time, but as shown in [Appendix A: Benchmarks](#) even with runtime parsing performance can be on par with or better than that of existing methods.

Compile-time processing of format strings [can work with a user-defined type T](#) if the `formatter<T>::parse` function is `constexpr`.

Argument visitation

The argument visitation API can be used to implement dynamic format specifiers, for example, width and precision passed as additional formatting arguments as opposed to being encoded in the format string itself:

```
int width = 10;
int precision = 3;
auto s = format("{0:{1}.{2}f}", 12.345678, width, precision);
// s == "      12.346"
```

An example of a user-defined formatter with dynamic width:

```
struct Answer {};
```

```
template<>
struct formatter<Answer> {
    int width_arg_index = 0;

    // Parses dynamic width in the format "<digit>".
    auto parse(format_parse_context& parse_ctx) {
        auto iter = parse_ctx.begin();
        auto get_char = [&]() { return iter != parse_ctx.end() ? *iter : 0; };

        if (get_char() != '{')
            return iter;
        ++iter;
        char c = get_char();
        if (!std::isdigit(c) || (++iter, get_char()) != '}')
            throw format_error("invalid format");
        width_arg_index = c - '0';
        return ++iter;
    }

    auto format(Answer, format_context& format_ctx) {
        auto arg = format_ctx.arg(width_arg_index);
        int width = visit_format_arg([&](auto value) -> int {
            if constexpr (!std::is_integral_v<decltype(value)>)
                throw format_error("width is not integral");
            else if (value < 0 || value > std::numeric_limits<int>::max())
                throw format_error("invalid width");
            else
                return value;
        }, arg);
        return format_to(format_ctx.out(), "{:{}", 42, width);
    }
};
```

```
std::string s = format("{0:{1}}", Answer(), 10);
// s == "      42"
```

In many cases formatting can be delegated to standard formatters which makes manual handling of dynamic specifiers unnecessary, but the latter is still important for more complex cases.

This API can also be used to implement a custom formatting engine, such as the one compatible with the `printf` syntax, to provide some of the benefits of the current proposal to the legacy code.

Impact on existing code

The proposed formatting API is defined in the new header `<format>` and should have no impact on existing code.

Proposed wording

Editing notes

All wording is relative to the pre-Kona C++ working draft [N4800](#).

The  character is used to denote a placeholder section number, table number, or paragraph number which the editor shall determine.

Add the header `<format>` to the "C++ library headers" table in [\[headers\]](#), in a place that respects the table's current alphabetic order.

Add an entry for `__cpp_lib_format` to the "Standard library feature-test macros" table in

[\[support.limits.general\]](#), in a place that respects the table's current alphabetic order:

Macro name	Value	Header(s)
...	...	...
_cpp_lib_format	201902L **placeholder**	<format>
...	...	...

Add a new entry to the "General utilities library summary" table in [\[utilities.general\]](#):

Subclause	Header(s)
...	...
19.19 Primitive numeric conversions	<charconv>
19.❖ Formatting	<format>

Add a new section in 19 [\[utilities\]](#).

19.❖ Formatting [\[format\]](#)

19.❖.1 Header [<format>](#) synopsis [\[format.syn\]](#)

```
namespace std {
    // \[format.error\], class format_error
    class format_error;

    // \[format.formatter\], formatter
    template<class charT> class basic_format_parse_context;
    using format_parse_context = basic_format_parse_context<char>;
    using wformat_parse_context = basic_format_parse_context<wchar_t>;

    template<class Out, class charT> class basic_format_context;
    using format_context = basic_format_context<unspecified, char>;
    using wformat_context = basic_format_context<unspecified, wchar_t>;

    template<class T, class charT = char> struct formatter;

    // \[format.arguments\], arguments
    template<class Context> class basic_format_arg;

    template<class Visitor, class Context>
        see below visit_format_arg(Visitor&& vis, basic_format_arg<Context> arg);

    template<class Context, class... Args> struct format_arg_store; // exposition only

    template<class Context> class basic_format_args;
    using format_args = basic_format_args<format_context>;
    using wformat_args = basic_format_args<wformat_context>;

    template<class Out, class charT>
        using format_args_t = basic_format_args<basic_format_context<Out, charT>>;

    template<class Context = format_context, class... Args>
        format_arg_store<Context, Args...>
        make_format_args(const Args&... args);
    template<class... Args>
        format_arg_store<wformat_context, Args...>
        make_wformat_args(const Args&... args);

    // \[format.functions\], formatting functions
    template<class... Args>
        string format(string_view fmt, const Args&... args);
    template<class... Args>
        wstring format(wstring_view fmt, const Args&... args);

    string vformat(string_view fmt, format_args args);
    wstring vformat(wstring_view fmt, wformat_args args);

    template<class Out, class... Args>
        Out format_to(Out out, string_view fmt, const Args&... args);
    template<class Out, class... Args>
        Out format_to(Out out, wstring_view fmt, const Args&... args);

    template<class Out>
        Out vformat_to(Out out, string_view fmt, format_args_t<Out, char> args);
    template<class Out>
```

```

    Out vformat_to(Out out, wstring_view fmt, format_args_t<Out, wchar_t> args);

template<class Out>
    struct format_to_n_result {
        Out out;
        iter_difference_t<Out> size;
    };

template<class Out, class... Args>
    format_to_n_result<Out> format_to_n(Out out, iter_difference_t<Out> n,
                                         string_view fmt, const Args&... args);

template<class Out, class... Args>
    format_to_n_result<Out> format_to_n(Out out, iter_difference_t<Out> n,
                                         wstring_view fmt, const Args&... args);

template<class... Args>
    size_t formatted_size(string_view fmt, const Args&... args);
template<class... Args>
    size_t formatted_size(wstring_view fmt, const Args&... args);
}

```

19.2 Format string [format.string]

A format string is a possibly empty sequence of zero or more *replacement fields*, *escape sequences*, and other characters. Each character that is not part of a replacement field or an escape sequence is copied unchanged to the output. An escape sequence is one of { or }. It is replaced with { or } respectively in the output. The syntax of replacement fields is as follows:

```

replacement-field ::= '{' [arg-id] [':' format-spec] '}'
arg-id           ::= '0' | nonzero-digit [integer]
integer          ::= digit [integer]
nonzero-digit    ::= '1'...'9'
digit            ::= '0'...'9'

```

The `arg-id` field specifies the index of the argument in `args` whose value is to be formatted and inserted into the output instead of the replacement field. The optional `format-spec` field specifies a non-default format for the replacement value.

[*Example*:

```
string s = format("{0}-{1}", 8); // s == "8-{"
```

— *end example*]

If the numeric `arg-ids` in a format string are 0, 1, 2, ... in sequence, they can all be omitted (not just some) and the numbers 0, 1, 2, ... will be automatically used in that order. A format string does not contain a mixture of automatic and manual indexing.

[*Example*:

```

string s0 = format("{} to {}", "a", "b"); // OK: automatic indexing
string s1 = format("{1} to {0}", "a", "b"); // OK: manual indexing
string s2 = format("{0} to {}", "a", "b"); // Error: mixing automatic and manual indexing
string s3 = format("{} to {1}", "a", "b"); // Error: mixing automatic and manual indexing

```

— *end example*]

The `format-spec` field contains *format specifications* that define how the value should be presented, including such details as field width, alignment, padding, and decimal precision. Each type can define its own *formatting mini-language* or interpretation of the `format-spec` field. The syntax of format specifications is as follows:

```

format-spec      ::= std-format-spec | custom-format-spec
std-format-spec  ::= [[fill] align] [sign] ['#'] ['0'] [width] ['.' precision] [type]
fill             ::= <a character other than '{' or '>'>
align            ::= '<' | '>' | '=' | '^'
sign             ::= '+' | '-' | ' '
width            ::= nonzero-digit [integer] | '{' arg-id '}'
precision        ::= integer | '{' arg-id '}'
type             ::= 'a' | 'A' | 'b' | 'B' | 'c' | 'd' | 'e' | 'E' | 'f' | 'F' |
                    'g' | 'G' | 'n' | 'o' | 'p' | 's' | 'x' | 'X'

```

where `std-format-spec` defines a common formatting mini-language supported by fundamental and string types, while `custom-format-spec` is a placeholder for user-defined mini-languages. Some of the

formatting options are only supported for arithmetic types.

The `fill` character can be any character other than `'{'` or `'}'`. The presence of a fill character is signaled by the character following it, which must be one of the alignment options. If the second character of `std::format-spec` is not a valid alignment option, then it is assumed that both the fill character and the alignment option are absent.

Let `charT` be `decltype(fmt)::value_type`.

The meaning of the various alignment options is as follows:

Option	Meaning
'<'	Forces the field to be left-aligned within the available space. This is the default for non-arithmetic types, <code>charT</code> , and <code>bool</code> , unless an integer presentation type is specified.
'>'	Forces the field to be right-aligned within the available space. This is the default for arithmetic types other than <code>charT</code> and <code>bool</code> or when an integer presentation type is specified.
'='	Forces the padding to be placed after the sign or prefix (if any) but before the digits. This is used for printing fields in the form <code>+000000120</code> . This alignment option is only valid for arithmetic types other than <code>charT</code> and <code>bool</code> or when an integer presentation type is specified.
'^'	Forces the field to be centered within the available space by inserting $N / 2$ and $N - N / 2$ fill characters before and after the value respectively, where N is the total number of fill characters to insert.

[Example:

```
char c = 120;
string s0 = format("{:6}", 42);      // s0 == "    42"
string s1 = format("{:6}", 'x');     // s1 == "x     "
string s2 = format("{:<6}", 'x');    // s2 == "x*****"
string s3 = format("{:>6}", 'x');    // s3 == "*****x"
string s4 = format("{:^6}", 'x');    // s4 == "***x***"
string s5 = format("{:=6}", 'x');    // Error: '=' with charT and no integer presentation type
string s6 = format("{:6d}", c);      // s6 == "    120"
string s7 = format("{:=+06d}", c);   // s7 == "+00120"
string s8 = format("{:0=#6x}", 0xa); // s8 == "0x000a"
string s9 = format("{:6}", true);    // s9 == "true  "
```

— end example]

Unless a minimum field width is defined, the field width will be determined by the size of the content and the alignment option will have no effect.

The `sign` option is only valid for arithmetic types other than `charT` and `bool` or when an integer presentation type is specified. It can be one of the following:

Option	Meaning
'+'	Indicates that a sign should be used for both non-negative and negative numbers.
'-'	Indicates that a sign should be used only for negative numbers (this is the default behavior).
space	Indicates that a leading space should be used for non-negative numbers, and a minus sign for negative numbers.

The `sign` option applies to floating-point infinity and NaN.

[Example:

```
double inf = numeric_limits<double>::infinity();
double nan = numeric_limits<double>::quiet_NaN();
string s0 = format("{0:} {0:+} {0:-} {0: }", 1); // s0 == "1 +1 1 1"
string s1 = format("{0:} {0:+} {0:-} {0: }", -1); // s1 == "-1 -1 -1 -1"
string s2 = format("{0:} {0:+} {0:-} {0: }", inf); // s2 == "inf +inf inf inf"
string s3 = format("{0:} {0:+} {0:-} {0: }", nan); // s3 == "nan +nan nan nan"
```

— end example]

The `'#'` option causes the *alternate form* to be used for the conversion. This option is only valid for arithmetic types other than `charT` and `bool` or when an integer presentation type is specified. For integers, when binary, octal, or hexadecimal output is used, this option adds the respective prefix `"0b"` (`"0B"`), `"0o"`, or `"0x"` (`"0X"`) to the output value. Whether the prefix is lower-case or upper-case is determined by the case of the type format specifier. For floating-point numbers the alternate form

causes the result of the conversion to always contain a decimal-point character, even if no digits follow it. Normally, a decimal-point character appears in the result of these conversions only if a digit follows it. In addition, for 'g' and 'G' conversions, trailing zeros are not removed from the result.

width is a decimal integer defining the minimum field width. If not specified, then the field width will be determined by the content.

Preceding the width field by a zero ('0') character enables sign-aware zero-padding for arithmetic types. This is equivalent to a fill character of '0' with an alignment type of '='.

The precision field is a decimal integer defining the precision or maximum field size. It can only be used with floating-point and string types. For floating-point types this field specifies the formatting precision. For string types it specifies how many characters will be used from the string.

Finally, the type determines how the data should be presented.

The available string presentation types are:

Type	Meaning
's'	Copies the string to the output.
none	The same as 's'.

The available charT presentation types are:

Type	Meaning
'c'	Copies the character to the output.
none	The same as 'c'.

Formatting of objects of arithmetic types and const void* is done as if by calling to_chars and copying the output through the output iterator of the format context with additional padding and adjustments as per format specifiers.

Let [first, last) be a range large enough to hold the to_chars output and value be the formatting argument value.

The available integer presentation types and their mapping to to_chars are:

Type	Meaning
'b'	to_chars(first, last, value, 2); using the '#' option with this type adds the prefix "0b" to the output.
'B'	The same as 'b', except that the '#' option adds the prefix "0B" to the output.
'd'	to_chars(first, last, value).
'o'	to_chars(first, last, value, 8); using the '#' option with this type adds the prefix "0" to the output.
'x'	to_chars(first, last, value, 16); using the '#' option with this type adds the prefix "0x" to the output.
'X'	The same as 'x', except that it uses uppercase letters for digits above 9 and the '#' option adds the prefix "0X" to the output.
'n'	The same as 'd', except that it uses the current global locale to insert the appropriate digit separator characters including the decimal radix separator.
none	The same as 'd' if the formatting argument type is not charT or bool.

Integer presentation types can also be used with charT and bool values. Values of type bool are formatted using textual representation, either "true" or "false", if the presentation type is not specified.

[Example:

```
string s0 = format!("{}", 42); // s0 == "42"
string s1 = format("{0:b} {0:d} {0:o} {0:x}", 42); // s1 == "101010 42 52 2a"
string s2 = format("{0:#x} {0:#X}", 42); // s2 == "0x2a 0X2A"
string s3 = format("{:n}", 1234); // s3 == "1,234" (depends on the locale)
```

— end example]

The available floating-point presentation types and their mapping to to_chars are:

--

Type	Meaning
'a'	to_chars(first, last, value, chars_format::hex, precision) if precision is specified, to_chars(first, last, value, chars_format::hex) otherwise.
'A'	The same as 'a', except that it uses uppercase letters for digits above 9, "P" to indicate the exponent, "INF" for infinity, and "NaN" for NaN.
'e'	to_chars(first, last, value, chars_format::scientific, precision); precision defaults to 6 if not specified.
'E'	The same as 'e', except that it uses "E" to indicate exponent, "INF" for infinity, and "NaN" for NaN.
'f'	to_chars(first, last, value, chars_format::fixed, precision); precision defaults to 6 if not specified.
'F'	The same as 'f', except that it uses "INF" for infinity, and "NaN" for NaN.
'g'	to_chars(first, last, value, chars_format::general, precision); precision defaults to 6 if not specified.
'G'	The same as 'g', except that it uses "E" to indicate exponent, "INF" for infinity, and "NaN" for NaN.
'n'	The same as 'g', except that it uses the current global locale to insert the appropriate digit separator characters including the decimal radix separator.
none	to_chars(first, last, value, chars_format::general, precision) if precision is specified, to_chars(first, last, value) otherwise.

Infinity and NaN are formatted as "inf" and "nan" respectively unless specified otherwise.

The available pointer presentation types and their mapping to to_chars are:

Type	Meaning
'p'	to_chars(first, last, reinterpret_cast<uintptr_t>(value), 16) with the prefix "0x" added to the output if uintptr_t is defined, implementation-defined otherwise.
none	The same as 'p'.

[Note: Pointer presentation types also apply to nullptr_t. — end note]

19.3 Formatting functions [format.functions]

```
template<class... Args>
string format(string_view fmt, const Args&... args);
```

Effects: Equivalent to: return vformat(fmt, make_format_args(args...));

```
template<class... Args>
wstring format(wstring_view fmt, const Args&... args);
```

Effects: Equivalent to: return vformat(fmt, make_wformat_args(args...));

```
string vformat(string_view fmt, format_args args);
wstring vformat(wstring_view fmt, wformat_args args);
```

Returns: A string object holding the character representation of formatting arguments provided by args formatted according to specifications given in fmt.

Throws: format_error if fmt is not a format string.

```
template<class Out, class... Args>
Out format_to(Out out, string_view fmt, const Args&... args);
template<class Out, class... Args>
Out format_to(Out out, wstring_view fmt, const Args&... args);
```

Effects: Equivalent to:

```
using context = basic_format_context<Out, decltype(fmt)::value_type>;
return vformat_to(out, fmt, {make_format_args<context>(args...)});

template<class Out>
Out vformat_to(Out out, string_view fmt, format_args_t<Out, char> args);
template<class Out>
```


Out vformat_to(Out out, wstring_view fmt, format_args_t<Out, wchar_t> args);

Let charT be `decltype(fmt)::value_type`.

Constraints: Out satisfies the `OutputIterator<const charT&>` requirements.

Expects: Out models `OutputIterator<const charT&>`.

Effects: Places the character representation of formatting arguments provided by `args`, formatted according to specifications given in `fmt`, into the range `[out, out + N)`, where $N = \text{formatted_size}(\text{fmt}, \text{args} \dots)$.

Returns: `out + N`.

Throws: `format_error` if `fmt` is not a format string.

```
template<class Out, class... Args>
    format_to_n_result<Out> format_to_n(Out out, iter_difference_t<Out> n,
                                         string_view fmt, const Args&... args);

template<class Out, class... Args>
    format_to_n_result<Out> format_to_n(Out out, iter_difference_t<Out> n,
                                         wstring_view fmt, const Args&... args);
```

Let charT be `decltype(fmt)::value_type`, $N = \text{formatted_size}(\text{fmt}, \text{args} \dots)$, and $M = \min(\max(n, 0), N)$.

Constraints: Out satisfies the `OutputIterator<const charT&>` requirements.

Mandates: `formatter<Ti, charT>` is enabled for each T_i in `Args`.

Expects: Out models `OutputIterator<const charT&>`.

Effects: Places the first M characters of the character representation of formatting arguments provided by `args`, formatted according to specifications given in `fmt`, into the range `[out, out + M)`.

Returns: `{out + M, N}`.

Throws: `format_error` if `fmt` is not a format string.

```
template<class... Args>
    size_t formatted_size(string_view fmt, const Args&... args);

template<class... Args>
    size_t formatted_size(wstring_view fmt, const Args&... args);
```

Mandates: `formatter<Ti, charT>` is enabled for each T_i in `Args`.

Returns: The number of characters in the character representation of formatting arguments `args` formatted according to specifications given in `fmt`.

Throws: `format_error` if `fmt` is not a format string.

19.4 Formatter [format.formatter]

The functions defined in [\[format.functions\]](#) use specializations of the class template `formatter` to format individual arguments.

Let charT be either `char` or `wchar_t`. Each specialization of `formatter` is either enabled or disabled, as described below. [*Note:* Enabled specializations meet the *Formatter* requirements, and disabled specializations do not. — *end note*] Each header that declares the template `formatter` provides enabled specializations:

- `template<> struct formatter<charT, charT>;`
- `template<> struct formatter<char, wchar_t>;`
- `template<> struct formatter<charT*, charT>;`
- `template<> struct formatter<const charT*, charT>;`
- `template<size_t N> struct formatter<const charT[N], charT>;`

- `template<class traits, class Allocator>`
`struct formatter<basic_string<charT, traits, Allocator>, charT>;`
- `template<class traits>`
`struct formatter<basic_string_view<charT, traits>, charT>;`
- specializations of `formatter` for `nullptr_t`, `void*`, `const void*`, `bool`, and all cv-unqualified standard integer types, extended integer types, and floating-point types as the first template argument and `charT` as the second template argument,

[*Note:* Specializations such as `formatter<wchar_t, char>` and `formatter<const char*, wchar_t>` that require implicit multibyte / wide string or character conversion are intentionally disabled. — *end note*]

For any types `T` and `charT` for which neither the library nor the user provides an explicit or partial specialization of the class template `formatter`, `formatter<T, charT>` is disabled.

If the library provides an explicit or partial specialization of `formatter<T, charT>`, that specialization is enabled except as noted otherwise.

If `F` is a disabled specialization of `formatter`, these values are false: `is_default_constructible_v<F>`, `is_copy_constructible_v<F>`, `is_move_constructible_v<F>`, `is_copy_assignable_v<F>`, `is_move_assignable_v<F>`.

An enabled specialization `formatter<T, charT>` will meet the *Formatter* requirements.

[*Example:*

```
#include <format>

enum color { red, green, blue };

const char* color_names[] = { "red", "green", "blue" };

template<> struct std::formatter<color> : std::formatter<const char*> {
    auto format(color c, format_context& ctx) {
        return formatter<const char*>::format(color_names[c], ctx);
    }
};

struct err {};
```

```
std::string s0 = std::format("{} ", 42);    // OK: library-provided formatter
std::string s1 = std::format("{} ", L"foo"); // Ill-formed: disabled formatter
std::string s2 = std::format("{} ", red);   // OK: user-provided formatter
std::string s3 = std::format("{} ", err{}); // Ill-formed: disabled formatter
```

— *end example*]

19.4.1 *Formatter* requirements [formatter.requirements]

A type `F` meets the *Formatter* requirements if:

- it meets the *Cpp17DefaultConstructible* (Table 25), *Cpp17CopyConstructible* (Table 27), *Cpp17CopyAssignable* (Table 29), and *Cpp17Destructible* (Table 30) requirements,
- it is swappable ([swappable.requirements]) for lvalues, and
- the expressions shown in Table 4 are valid and have the indicated semantics.

Given character type `charT`, output iterator type `Out`, and formatting argument type `T`, in Table 4 `f` is a value of type `F`, `u` is an lvalue of type `T`, `t` is a value of a type convertible to (possibly `const`) `T`, `pc` is `basic_format_parse_context<charT>`, `fc` is `basic_format_context<Out, charT>`, `pc` is an lvalue of type `PC`, and `fc` is an lvalue of type `FC`. `pc.begin()` points to the beginning of the `format-spec` ([format.syntax]) portion of the format string. If `format-spec` is empty then either `pc.begin() == pc.end()` OR `*pc.begin() == '}'`.

Table 4 — *Formatter* requirements

Expression	Return type	Requirement
<code>f.parse(pc)</code>	<code>PC::iterator</code>	Parses <code>format-spec</code> ([format.string]) for type <code>T</code> in the range <code>[pc.begin(), pc.end())</code> until the first unmatched character. Throws <code>format_error</code> unless the whole range is parsed or the unmatched character is <code>'}'</code> . [<i>Note:</i> This allows formatters to emit meaningful error messages. — <i>end note</i>]

	Stores the parsed format specifiers in <code>*this</code> and returns an iterator past the end of the parsed range.
<code>f.format(t, fc) FC::iterator</code>	Formats <code>t</code> according to the specifiers stored in <code>*this</code> , writes the output to <code>fc.out()</code> and returns an iterator past the end of the output range. The output shall only depend on <code>t</code> , the current global locale, and the range <code>[pc.begin(), pc.end())</code> from the last call to <code>f.parse(pc)</code> .
<code>f.format(u, fc) FC::iterator</code>	As above but does not modify <code>u</code> .

19.4.2 Class template `basic_format_parse_context` [`format.parse_context`]

```
namespace std {
    template<class charT>
    class basic_format_parse_context {
    public:
        using char_type = charT;
        using const_iterator = typename basic_string_view<charT>::const_iterator;
        using iterator = const_iterator;

    private:
        iterator begin_;                // exposition only
        iterator end_;                  // exposition only
        enum indexing { unknown, manual, automatic }; // exposition only
        indexing indexing_;             // exposition only
        size_t next_arg_id_;            // exposition only
        size_t num_args_;               // exposition only

    public:
        explicit constexpr basic_format_parse_context(basic_string_view<charT> fmt,
                                                       size_t num_args = 0) noexcept;
        basic_format_parse_context(const basic_format_parse_context&) = delete;
        basic_format_parse_context& operator=(const basic_format_parse_context&) = delete;

        constexpr const_iterator begin() const noexcept;
        constexpr const_iterator end() const noexcept;
        constexpr void advance_to(const_iterator it);

        constexpr size_t next_arg_id();
        constexpr void check_arg_id(size_t id);
    };
}
```

An instance of `basic_format_parse_context` holds the format string parsing state consisting of the format string range being parsed and the argument counter for automatic indexing.

```
explicit constexpr basic_format_parse_context(basic_string_view<charT> fmt,
                                              size_t num_args = 0) noexcept;
```

Effects: Initializes `begin_` with `fmt.begin()`, `end_` with `fmt.end()`, `indexing_` with `unknown`, `next_arg_id_` with `0`, and `num_args_` with `num_args`.

```
constexpr const_iterator begin() const noexcept;
```

Returns: `begin_`.

```
constexpr const_iterator end() const noexcept;
```

Returns: `end_`.

```
constexpr void advance_to(const_iterator it);
```

Expects: `end()` is reachable from `it`.

Effects: Equivalent to: `begin_ = it;`

```
constexpr size_t next_arg_id();
```

Effects: If `indexing_ != manual`, equivalent to:

```
if (indexing_ == unknown)
    indexing_ = automatic;
return next_arg_id++;
```

Throws: `format_error` if `indexing_ == manual` which indicates mixing of automatic and manual

argument indexing.

```
constexpr void check_arg_id(size_t id);
```

Expects: !is_constant_evaluated() || id < num_args_.

Effects: If indexing_ != automatic, equivalent to:

```
if (indexing_ == unknown)
    indexing_ = manual;
```

Throws: format_error if indexing_ == automatic which indicates mixing of automatic and manual argument indexing.

Remarks: Call expressions that violate the precondition in the *Expects* element are not core constant expressions ([\[expr.const\]](#)).

19.4.3 Class template basic_format_context [format.context]

```
namespace std {
    template<class Out, class charT>
    class basic_format_context {
        basic_format_args<basic_format_context> args_; // exposition only
        Out out_; // exposition only

    public:
        using iterator = Out;
        using char_type = charT;
        template<class T> using formatter_type = formatter<T, charT>;

        basic_format_arg<basic_format_context> arg(size_t id) const;

        iterator out();
        void advance_to(iterator it);
    };
}
```

An instance of basic_format_context holds formatting state consisting of the formatting arguments and the output iterator.

Out shall model OutputIterator<const charT>.

format_context is an alias for the instantiation of basic_format_context with an output iterator that appends to string such as back_insert_iterator<string>.

[*Note:* For a given type charT, implementations are encouraged to provide a single instantiation of basic_format_context for appending to basic_string<charT>, vector<charT>, or any other container with contiguous storage by wrapping those in temporary objects with uniform interface (like span) and polymorphic reallocation. — *end note*]

```
basic_format_arg<basic_format_context> arg(size_t id) const;
```

Returns: args_.get(id).

```
iterator out();
```

Returns: out_.

```
void advance_to(iterator it);
```

Effects: Equivalent to: out_ = it;

[*Example:*

```
struct S {
    int value;
};

template<> struct std::formatter<S> {
    size_t width_arg_id = 0;

    // Parses a width argument id in the format { <digit> }.
    constexpr auto parse(format_parse_context& ctx) {
        auto iter = ctx.begin();
```

```

    auto get_char = [&]() { return iter != ctx.end() ? *iter : 0; };
    if (get_char() != '{}')
        return iter;
    ++iter;
    char c = get_char();
    if (!isdigit(c) || (++iter, get_char()) != '{}')
        throw format_error("invalid format");
    width_arg_id = c - '0';
    ctx.check_arg_id(width_arg_id);
    return ++iter;
}

// Formats S with width given by the argument width_arg_id.
auto format(S s, format_context& ctx) {
    int width = visit_format_arg([&](auto value) -> int {
        if constexpr (!is_integral_v<decltype(value)>)
            throw format_error("width is not integral");
        else if (value < 0 || value > numeric_limits<int>::max())
            throw format_error("invalid width");
        else
            return value;
    }, ctx.arg(width_arg_id));
    return format_to(ctx.out(), "{0:{1}}", s.value, width);
}
};

std::string s = std::format("{0:{1}}", S{42}, 10); // s == "      42"

```

— end example]

19.5 Arguments [format.arguments]

19.5.1 Class template `basic_format_arg` [format.arg]

```

namespace std {
    template<class Context>
    class basic_format_arg {
    public:
        class handle;

    private:
        using char_type = typename Context::char_type;

        variant<monostate, bool, char_type,
            int, unsigned int, long long int, unsigned long long int,
            double, long double,
            const char_type*, basic_string_view<char_type>,
            const void*, handle> value;

        template<typename T> explicit basic_format_arg(const T& v) noexcept; // exposition only
        explicit basic_format_arg(float n) noexcept; // exposition only
        explicit basic_format_arg(double n) noexcept; // exposition only
        explicit basic_format_arg(long double n) noexcept; // exposition only
        explicit basic_format_arg(const char_type* s); // exposition only

        template<class traits>
        explicit basic_format_arg(
            basic_string_view<char_type, traits> s) noexcept; // exposition only

        template<class traits, class Allocator>
        explicit basic_format_arg(
            const basic_string<char_type, traits, Allocator>& s) noexcept; // exposition only

        explicit basic_format_arg(nullptr_t) noexcept; // exposition only

        template<class T>
        explicit basic_format_arg(const T* p) noexcept; // exposition only

        template<class Visitor, class Ctx>
        friend auto visit_format_arg(Visitor&& vis,
            basic_format_arg<Ctx> arg); // exposition only

        template<class Ctx, class... Args>
        friend format_arg_store<Ctx, Args...>
            make_format_args(const Args&... args); // exposition only

    public:
        basic_format_arg() noexcept;

```

```

    explicit operator bool() const noexcept;
};
}

```

An instance of `basic_format_arg` provides access to a formatting argument for user-defined formatters.

The behavior of a program that adds specializations of `basic_format_arg` is undefined.

```
basic_format_arg() noexcept;
```

Ensures: `!(*this)`.

```
template<class T> explicit basic_format_arg(const T& v) noexcept;
```

Constraints: `typename Context::template formatter_type<T>` is enabled.

Effects:

- if `T` is `bool` or *char-type*, initializes value with `v`; otherwise,
- if `T` is `char` and *char-type* is `wchar_t`, initializes value with `static_cast<wchar_t>(v)`; otherwise,
- if `T` is a signed integer type ([\(basic.fundamental\)](#)) and `sizeof(T) <= sizeof(int)`, initializes value with `static_cast<int>(v)`; otherwise,
- if `T` is an unsigned integer type and `sizeof(T) <= sizeof(unsigned int)`, initializes value with `static_cast<unsigned int>(v)`; otherwise,
- if `T` is a signed integer type and `sizeof(T) <= sizeof(long long int)`, initializes value with `static_cast<long long int>(v)`; otherwise,
- if `T` is an unsigned integer type and `sizeof(T) <= sizeof(unsigned long long int)`, initializes value with `static_cast<unsigned long long int>(v)`; otherwise,
- initializes value with `handle(v)`.

```
explicit basic_format_arg(float n) noexcept;
```

Effects: Initializes value with `static_cast<double>(n)`.

```
explicit basic_format_arg(double n) noexcept;
```

```
explicit basic_format_arg(long double n) noexcept;
```

Effects: Initializes value with `n`.

```
explicit basic_format_arg(const char-type* s);
```

Expects: `s` points to a NTCTS ([\(defs.ntcts\)](#)).

Effects: Initializes value with `s`.

```
template<class traits>
    explicit basic_format_arg(basic_string_view<char-type, traits> s) noexcept;
```

Effects: Initializes value with `s`.

```
template<class traits, class Allocator>
    explicit basic_format_arg(
        const basic_string<char-type, traits, Allocator>& s) noexcept;
```

Effects: Initializes value with `basic_string_view<char-type>(s.data(), s.size())`.

```
explicit basic_format_arg(nullptr_t) noexcept;
```

Effects: Initializes value with `static_cast<const void*>(nullptr)`.

```
template<class T> explicit basic_format_arg(T* p) noexcept;
```

Constraints: `is_void_v<T>` is true.

Effects: Initializes value with `p`.

[*Note*: Constructing `basic_format_arg` from a pointer to a member is ill-formed unless the user provides an enabled specialization of `formatter` for this pointer to member type. — *end note*]

explicit operator bool() const noexcept;

Returns: `!holds_alternative<monostate>(value)`.

The class `handle` allows formatting an object of a user-defined type.

```
namespace std {
    template<class Context>
    class basic_format_arg<Context>::handle {
        const void* ptr_; // exposition only
        void (*format_)(basic_format_parse_context<char-type>&,
                        Context&, const void*); // exposition only

        template<class T> explicit handle(const T& val) noexcept; // exposition only

        friend class basic_format_arg<Context>; // exposition only

    public:
        void format(basic_format_parse_context<char-type>&, Context& ctx) const;
    };
}
```

template<class T> explicit handle(const T& val) noexcept;

Effects: Initializes `ptr_` with `addressof(val)` and `format_` with

```
{
    [(basic_format_parse_context<char-type>& parse_ctx, Context& format_ctx, const void* ptr) {
        typename Context::template formatter_type<T> f;
        parse_ctx.advance_to(f.parse(parse_ctx));
        format_ctx.advance_to(f.format(*static_cast<const T*>(ptr), format_ctx));
    }
}
```

void format(basic_format_parse_context<char-type>& parse_ctx, Context& format_ctx) const;

Effects: Equivalent to: `format_(parse_ctx, format_ctx, ptr_)`;

19.5.2 Argument visitation [format.visit]

```
template<class Visitor, class Context>
    see below visit_format_arg(Visitor&& vis, basic_format_arg<Context> arg);
```

Effects: Equivalent to: `return visit(forward<Visitor>(vis), arg.value)`;

19.5.3 Class template *format-arg-store* [format.arg_store]

```
namespace std {
    template<class Context, class... Args>
    struct format_arg_store { // exposition only
        array<basic_format_arg<Context>, sizeof...(Args)> args;
    };
}
```

An instance of *format-arg-store* stores formatting arguments.

19.5.4 Function template `make_format_args` [format.make_args]

```
template<class Context = format_context, class... Args>
    format_arg_store<Context, Args...> make_format_args(const Args&... args);
```

Mandates: `typename Context::template formatter_type<Ti>` is enabled for each `Ti` in `Args`.

Returns: `{basic_format_arg<Context>(args) ...}`.

19.5.5 Function template `make_wformat_args` [format.make_wargs]

```
template<class... Args>
    format_arg_store<wformat_context, Args...> make_wformat_args(const Args&... args);
```

Effects: Equivalent to: `return make_format_args<wformat_context>(args...)`;

19.5.6 Class `template basic_format_args` [`format.basic_args`]

```
namespace std {
    template<class Context>
    class basic_format_args {
        size_t size_; // exposition only
        const basic_format_arg<Context>* data_; // exposition only

    public:
        basic_format_args() noexcept;

        template<class... Args>
        basic_format_args(const format-arg-store<Context, Args...>& store) noexcept;

        basic_format_arg<Context> get(size_t i) const noexcept;
    };
}
```

An instance of `basic_format_args` provides access to formatting arguments.

```
basic_format_args() noexcept;
```

Effects: Initializes `size_` with 0.

```
template<class... Args>
    basic_format_args(const format-arg-store<Context, Args...>& store) noexcept;
```

Effects: Initializes `size_` with `sizeof...(Args)` and `data_` with `store.args.data()`.

```
basic_format_arg<Context> get(size_t i) const noexcept;
```

Returns: `i < size_ ? data_[i] : basic_format_arg<Context>()`.

[*Note:* Implementations are encouraged to optimize the representation of `basic_format_args` for small number of formatting arguments by storing indices of type alternatives separately from values and packing the former. — *end note*]

19.6 Class `format_error` [`format.error`]

```
namespace std {
    class format_error : public runtime_error {
    public:
        explicit format_error(const string& what_arg);
        explicit format_error(const char* what_arg);
    };
}
```

The class `format_error` defines the type of objects thrown as exceptions to report errors from the formatting library.

```
format_error(const string& what_arg);
```

Ensures: `strcmp(what(), what_arg.c_str()) == 0`.

```
format_error(const char* what_arg);
```

Ensures: `strcmp(what(), what_arg) == 0`.

Related work

The Boost Format library [8] is an established formatting library that uses `printf`-like format string syntax with extensions. The main differences between this library and the current proposal are:

- Syntax: for the reasons described in section [Format String Syntax](#) this proposal uses a new syntax instead of extending the `printf` one. This allows much simpler and easier to parse grammar, not burdened by legacy specifiers used to convey type information. For example, Boost Format has two ways to refer to an argument by index and allows but ignores some format specifiers.
- API: Boost Format uses `operator%` to pass formatting arguments while this proposal uses variadic function templates.
- Performance: the implementation of this proposal is several times faster than the implementation of Boost Format on `tinyformat` benchmarks [9], generates smaller binary code and is faster to compile.

A `printf`-like Interface for the Streams Library [10] is similar to the Boost Format library but uses variadic templates instead of `operator%`. Unfortunately it hasn't been updated since 2013 and the same arguments about format string syntax apply to it.

The FastFormat library [11] is another well-known formatting library. Similarly to this proposal, FastFormat uses brace-delimited format specifiers, but otherwise the format string syntax is different and the library has significant limitations [12]:

Three features that have no hope of being accommodated within the current design are:

- Leading zeros (or any other non-space padding)
- Octal/hexadecimal encoding
- Runtime width/alignment specification

Formatting facilities of the Folly library [13] are the closest to the current proposal. Folly also uses Python-like format string syntax nearly identical to the one described here. However, the API details are quite different. The current proposal tries to address performance and code bloat issues that are largely ignored by Folly Format. For instance formatting functions in Folly Format are parameterized on all argument types while in this proposal, only the inlined wrapper functions are, which results in much smaller binary code and better compile times.

Implementation

An implementation of this proposal is available in the `std` branch of the open-source `fmt` library [14].

Acknowledgements

Thanks to Alberto Barbati, Antony Polukhin, Beman Dawes, Bengt Gustafsson, Daniel Krügler, Eric Niebler, Jason McKesson, Jeffrey Yasskin, Joël Lamotte, Lars Gullik Bjønnes, Lee Howes, Louis Dionne, Matt Clarke, Michael Park, Sergey Ignatchenko, Thiago Macieira, Tomasz Kamiński, Zach Laine, Zhihao Yuan and participants of the Library Evolution Working Group and the Library Working Group for their feedback, support, constructive criticism and contributions to the proposal. Special thanks to Howard Hinnant who encouraged me to write the proposal and gave useful early advice on how to go about it.

The format string syntax is based on the Python documentation [3].

References

- [1] *The `fprintf` function*. ISO/IEC 9899:2011. 7.21.6.1.
- [2] [*`fprintf`, `printf`, `snprintf`, `sprintf` - print formatted output*](#). The Open Group Base Specifications Issue 6 IEEE Std 1003.1, 2004 Edition.
- [3] [*6.1.3. Format String Syntax*](#). Python 3.5.2 documentation.
- [4] [*String.Format Method*](#). .NET Framework Class Library.
- [5] [*Module `std::fmt`*](#). The Rust Standard Library.
- [6] [*Format Specification Syntax: `printf` and `wprintf` Functions*](#). C++ Language and Standard Libraries.
- [7] [*10.4.2 Rearranging `printf` Arguments*](#). The GNU Awk User's Guide.
- [8] [*Boost Format library*](#). Boost 1.63 documentation.
- [9] [*Speed Test*](#). The `fmt` library repository.
- [10] [*A `printf`-like Interface for the Streams Library \(revision 1\)*](#).
- [11] [*The FastFormat library website*](#).
- [12] [*An Introduction to Fast Format \(Part 1\): The State of the Art*](#). Overload Journal #89 - February 2009
- [13] [*The folly library repository*](#).
- [14] [*The fmt library repository*](#).
- [15] [*P0424: Reconsidering literal operator templates for strings*](#).
- [16] [*D0355: Extending `<chrono>` to Calendars and Time Zones*](#).
- [17] [*P0067: Elementary string conversions*](#).
- [18] [*MPark.Format: Compile-time Checked, Type-Safe Formatting in C++14*](#).
- [19] [*Google Benchmark: A microbenchmark support library*](#).
- [20] [*N4412: Shortcomings of `iostreams`*](#).

Appendix A: Benchmarks

To demonstrate that the formatting functions described in this paper can be implemented efficiently, we compare the reference implementation [14] of `format` and `format_to` to `sprintf`, `ostringstream` and `to_string` on the following benchmark. This benchmark generates a set of integers with random numbers of digits, applies each method to convert each integer into a string (either `std::string` or a char buffer depending on the API) and uses the Google Benchmark library [19] to measure timings:

```
#include <algorithm>
#include <cmath>
#include <cstdio>
#include <limits>
#include <sstream>
#include <string>
#include <utility>
#include <vector>

#include <benchmark/benchmark.h>
#include <fmt/format.h>

// Returns a pair with the smallest and the largest value of integral type T
// with the given number of digits.
template<typename T>
std::pair<T, T> range(int num_digits) {
    T first = std::pow(T(10), num_digits - 1);
    int max_digits = std::numeric_limits<T>::digits10 + 1;
    T last = num_digits < max_digits ? first * 10 - 1 :
        std::numeric_limits<T>::max();
    return {num_digits > 1 ? first : 0, last};
}

// Generates values of integral type T with random number of digits.
template<typename T>
std::vector<T> generate_random_data(int numbers_per_digit) {
    int max_digits = std::numeric_limits<T>::digits10 + 1;
    std::vector<T> data;
    data.reserve(max_digits * numbers_per_digit);
    for (int i = 1; i <= max_digits; ++i) {
        auto r = range<T>(i);
        auto value = r.first;
        std::generate_n(std::back_inserter(data), numbers_per_digit, [=]() mutable {
            T result = value;
            value = value < r.second ? value + 1 : r.first;
            return result;
        });
    }
    std::random_shuffle(data.begin(), data.end());
    return data;
}

auto data = generate_random_data<int>(1000);

void sprintf(benchmark::State& s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            char buffer[12];
            result += std::sprintf(buffer, "%d", i);
        }
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(sprintf);

void ostream(benchmark::State& s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            std::ostringstream ss;
            ss << i;
            result += ss.str().size();
        }
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(ostream);

void to_string(benchmark::State& s) {
    size_t result = 0;
    while (s.KeepRunning()) {
```

```

        for (auto i: data)
            result += std::to_string(i).size();
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(to_string);

void format(benchmark::State& s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data)
            result += fmt::format("{}", i).size();
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(format);

void format_to(benchmark::State& s) {
    size_t result = 0;
    while (s.KeepRunning()) {
        for (auto i: data) {
            char buffer[12];
            result += fmt::format_to(buffer, "{}", i) - buffer;
        }
    }
    benchmark::DoNotOptimize(result);
}
BENCHMARK(format_to);

BENCHMARK_MAIN();

```

The benchmark was compiled with clang (Apple LLVM version 9.0.0 clang-900.0.39.2) with `-O3 -DNDEBUG` and run on a macOS system. Below are the results:

```

Run on (4 X 3100 MHz CPU s)
2018-01-27 07:12:00
Benchmark                Time          CPU Iterations
-----
sprintf                 882311 ns      881076 ns       781
ostringstream          2892035 ns     2888975 ns       242
to_string               1167422 ns     1166831 ns       610
format                  675636 ns      674382 ns      1045
format_to               499376 ns      498996 ns      1263

```

The `format` and `format_to` functions show much better performance than the other methods. The `format` function that constructs `std::string` is even 30% faster than the system's version of `sprintf` that uses stack-allocated `char` buffer. `format_to` with a stack-allocated buffer is ~60% faster than `sprintf`.

Appendix B: Binary code comparison

In this section we compare per-call binary code size between the reference implementation that uses techniques described in section [Binary footprint](#) and standard formatting facilities. All the code snippets are compiled with clang (Apple LLVM version 9.0.0 clang-900.0.39.2) with `-O3 -DNDEBUG -c -std=c++14` and the resulted binaries are disassembled with `objdump -S`:

```

void consume(const char*);

void sprintf_test() {
    char buffer[100];
    sprintf(buffer, "The answer is %d.", 42);
    consume(buffer);
}

__Z12sprintf_testv:
    0:      55      pushq   %rbp
    1:      48 89 e5      movq    %rsp, %rbp
    4:      53      pushq   %rbx
    5:      48 83 ec 78    subq    $120, %rsp
    9:      48 8b 05 00 00 00 00 movq    (%rip), %rax
   10:      48 8b 00      movq    (%rax), %rax
   13:      48 89 45 f0      movq    %rax, -16(%rbp)
   17:      48 8d 35 37 00 00 00 leaq    55(%rip), %rsi
   1e:      48 8d 5d 80      leaq    -128(%rbp), %rbx
   22:      ba 2a 00 00 00 00 movl    $42, %edx
   27:      31 c0      xorl    %eax, %eax
   29:      48 89 df      movq    %rbx, %rdi

```

```

2c:    e8 00 00 00 00 callq 0 <__Z12sprintf_testv+0x31>
31:    48 89 df        movq    %rbx, %rdi
34:    e8 00 00 00 00 callq 0 <__Z12sprintf_testv+0x39>
39:    48 8b 05 00 00 00 00 movq    (%rip), %rax
40:    48 8b 00        movq    (%rax), %rax
43:    48 3b 45 f0     cmpq    -16(%rbp), %rax
47:    75 07 jne      7 <__Z12sprintf_testv+0x50>
49:    48 83 c4 78     addq    $120, %rsp
4d:    5b             popq    %rbx
4e:    5d             popq    %rbp
4f:    c3             retq
50:    e8 00 00 00 00 callq 0 <__Z12sprintf_testv+0x55>

```

```

void format_test() {
    consume(format("The answer is {}.", 42).c_str());
}

```

```

__Z11format_testv:
0:    55             pushq   %rbp
1:    48 89 e5       movq    %rsp, %rbp
4:    53             pushq   %rbx
5:    48 83 ec 28     subq    $40, %rsp
9:    48 c7 45 d0 2a 00 00 00 movq    $42, -48(%rbp)
11:   48 8d 35 f4 83 01 00 leaq    99316(%rip), %rsi
18:   48 8d 7d e0     leaq    -32(%rbp), %rdi
1c:   4c 8d 45 d0     leaq    -48(%rbp), %r8
20:   ba 11 00 00 00 movl    $17, %edx
25:   b9 02 00 00 00 movl    $2, %ecx
2a:   e8 00 00 00 00 callq 0 <__Z11format_testv+0x2F>
2f:   f6 45 e0 01     testb   $1, -32(%rbp)
33:   48 8d 7d e1     leaq    -31(%rbp), %rdi
37:   48 0f 45 7d f0 cmovneq -16(%rbp), %rdi
3c:   e8 00 00 00 00 callq 0 <__Z11format_testv+0x41>
41:   f6 45 e0 01     testb   $1, -32(%rbp)
45:   74 09 je        9 <__Z11format_testv+0x50>
47:   48 8b 7d f0     movq    -16(%rbp), %rdi
4b:   e8 00 00 00 00 callq 0 <__Z11format_testv+0x50>
50:   48 83 c4 28     addq    $40, %rsp
54:   5b             popq    %rbx
55:   5d             popq    %rbp
56:   c3             retq
57:   48 89 c3       movq    %rax, %rbx
5a:   f6 45 e0 01     testb   $1, -32(%rbp)
5e:   74 09 je        9 <__Z11format_testv+0x69>
60:   48 8b 7d f0     movq    -16(%rbp), %rdi
64:   e8 00 00 00 00 callq 0 <__Z11format_testv+0x69>
69:   48 89 df       movq    %rbx, %rdi
6c:   e8 00 00 00 00 callq 0 <__Z11format_testv+0x71>
71:   66 66 66 66 66 66 2e 0f 1f 84 00 00 00 00 00 00 nopw    %cs:(%rax,%rax)

```

```

void ostringstream_test() {
    std::ostringstream ss;
    ss << "The answer is " << 42 << ".";
    consume(ss.str().c_str());
}

```

```

__Z18ostringstream_testv:
0:    55             pushq   %rbp
1:    48 89 e5       movq    %rsp, %rbp
4:    41 57          pushq   %r15
6:    41 56          pushq   %r14
8:    41 55          pushq   %r13
a:    41 54          pushq   %r12
c:    53             pushq   %rbx
d:    48 81 ec 38 01 00 00 subq    $312, %rsp
14:   4c 8d b5 18 ff ff ff leaq    -232(%rbp), %r14
1b:   4c 8d a5 b0 fe ff ff leaq    -336(%rbp), %r12
22:   48 8b 05 00 00 00 00 movq    (%rip), %rax
29:   48 8d 48 18     leaq    24(%rax), %rcx
2d:   48 89 8d a8 fe ff ff movq    %rcx, -344(%rbp)
34:   48 83 c0 40     addq    $64, %rax
38:   48 89 85 18 ff ff ff movq    %rax, -232(%rbp)
3f:   4c 89 f7       movq    %r14, %rdi
42:   4c 89 e6       movq    %r12, %rsi
45:   e8 00 00 00 00 callq 0 <__Z18ostringstream_testv+0x4A>
4a:   c7 c7 45 a0 00 00 00 00 movq    $0, -96(%rbp)
52:   c7 45 a8 ff ff ff ff movl    $4294967295, -88(%rbp)
59:   48 8b 1d 00 00 00 00 movq    (%rip), %rbx
60:   4c 8d 6b 18     leaq    24(%rbx), %r13

```

```

64: 4c 89 ad a8 fe ff ff    movq    %r13, -344(%rbp)
6b: 48 83 c3 40            addq    $64, %rbx
6f: 48 89 9d 18 ff ff ff    movq    %rbx, -232(%rbp)
76: 4c 89 e7              movq    %r12, %rdi
79: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x7E>
7e: 4c 8b 3d 00 00 00 00    movq    (%rip), %r15
85: 49 83 c7 10            addq    $16, %r15
89: 4c 89 bd b0 fe ff ff    movq    %r15, -336(%rbp)
90: 48 c7 85 08 ff ff ff 00 00 00 00    movq    $0, -248(%rbp)
9b: 48 c7 85 00 ff ff ff 00 00 00 00    movq    $0, -256(%rbp)
a6: 48 c7 85 f8 fe ff ff 00 00 00 00    movq    $0, -264(%rbp)
b1: 48 c7 85 f0 fe ff ff 00 00 00 00    movq    $0, -272(%rbp)
bc: c7 85 10 ff ff ff 10 00 00 00    movl    $16, -240(%rbp)
c6: 0f 57 c0              xorps   %xmm0, %xmm0
c9: 0f 29 45 b0          movaps  %xmm0, -80(%rbp)
cd: 48 c7 45 c0 00 00 00 00    movq    $0, -64(%rbp)
d5: 48 8d 75 b0          leaq    -80(%rbp), %rsi
d9: 4c 89 e7              movq    %r12, %rdi
dc: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0xE1>
e1: f6 45 b0 01          testb   $1, -80(%rbp)
e5: 74 09                je      9 <__Z18ostream_testv+0xF0>
e7: 48 8b 7d c0          movq    -64(%rbp), %rdi
eb: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0xF0>
f0: 48 8d 35 dd 10 00 00    leaq    4317(%rip), %rsi
f7: 48 8d bd a8 fe ff ff    leaq    -344(%rbp), %rdi
fe: ba 0e 00 00 00 00    movl    $14, %edx
103: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x108>
108: be 2a 00 00 00 00    movl    $42, %esi
10d: 48 89 c7              movq    %rax, %rdi
110: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x115>
115: 48 8d 35 c7 10 00 00    leaq    4295(%rip), %rsi
11c: ba 01 00 00 00 00    movl    $1, %edx
121: 48 89 c7              movq    %rax, %rdi
124: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x129>
129: 48 8d 7d b0          leaq    -80(%rbp), %rdi
12d: 4c 89 e6              movq    %r12, %rsi
130: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x135>
135: f6 45 b0 01          testb   $1, -80(%rbp)
139: 48 8d 7d b1          leaq    -79(%rbp), %rdi
13d: 48 0f 45 7d c0        cmovneq -64(%rbp), %rdi
142: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x147>
147: f6 45 b0 01          testb   $1, -80(%rbp)
14b: 74 09                je      9 <__Z18ostream_testv+0x156>
14d: 48 8b 7d c0          movq    -64(%rbp), %rdi
151: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x156>
156: 4c 89 ad a8 fe ff ff    movq    %r13, -344(%rbp)
15d: 48 89 9d 18 ff ff ff    movq    %rbx, -232(%rbp)
164: 4c 89 bd b0 fe ff ff    movq    %r15, -336(%rbp)
16b: f6 85 f0 fe ff ff 01    testb   $1, -272(%rbp)
172: 74 0c                je      12 <__Z18ostream_testv+0x180>
174: 48 8b bd 00 ff ff ff    movq    -256(%rbp), %rdi
17b: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x180>
180: 4c 89 e7              movq    %r12, %rdi
183: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x188>
188: 48 8b 35 00 00 00 00    movq    (%rip), %rsi
18f: 48 83 c6 08          addq    $8, %rsi
193: 48 8d bd a8 fe ff ff    leaq    -344(%rbp), %rdi
19a: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x19F>
19f: 4c 89 f7              movq    %r14, %rdi
1a2: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x1A7>
1a7: 48 81 c4 38 01 00 00    addq    $312, %rsp
1ae: 5b                  popq    %rbx
1af: 41 5c                popq    %r12
1b1: 41 5d                popq    %r13
1b3: 41 5e                popq    %r14
1b5: 41 5f                popq    %r15
1b7: 5d                  popq    %rbp
1b8: c3                  retq
1b9: 48 89 45 d0          movq    %rax, -48(%rbp)
1bd: f6 45 b0 01          testb   $1, -80(%rbp)
1c1: 74 3b                je      59 <__Z18ostream_testv+0x1FE>
1c3: 48 8b 7d c0          movq    -64(%rbp), %rdi
1c7: e8 00 00 00 00 00    callq  0 <__Z18ostream_testv+0x1CC>
1cc: eb 30                jmp     48 <__Z18ostream_testv+0x1FE>
1ce: eb 2a                jmp     42 <__Z18ostream_testv+0x1FA>
1d0: 48 89 45 d0          movq    %rax, -48(%rbp)
1d4: f6 45 b0 01          testb   $1, -80(%rbp)
1d8: 74 39                je      57 <__Z18ostream_testv+0x213>
1da: 48 8b 7d c0          movq    -64(%rbp), %rdi

```

```

1de: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x1E3>
1e3: f6 85 f0 fe ff ff 01 testb $1, -272(%rbp)
1ea: 75 30 jne 48 <__Z18ostream_testv+0x21C>
1ec: eb 3a jmp 58 <__Z18ostream_testv+0x228>
1ee: 48 89 45 d0 movq %rax, -48(%rbp)
1f2: eb 3c jmp 60 <__Z18ostream_testv+0x230>
1f4: 48 89 45 d0 movq %rax, -48(%rbp)
1f8: eb 4d jmp 77 <__Z18ostream_testv+0x247>
1fa: 48 89 45 d0 movq %rax, -48(%rbp)
1fe: 4c 89 ad a8 fe ff ff movq %r13, -344(%rbp)
205: 48 89 9d 18 ff ff ff movq %rbx, -232(%rbp)
20c: 4c 89 bd b0 fe ff ff movq %r15, -336(%rbp)
213: f6 85 f0 fe ff ff 01 testb $1, -272(%rbp)
21a: 74 0c je 12 <__Z18ostream_testv+0x228>
21c: 48 8b bd 00 ff ff ff movq -256(%rbp), %rdi
223: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x228>
228: 4c 89 e7 movq %r12, %rdi
22b: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x230>
230: 48 8b 35 00 00 00 00 movq (%rip), %rsi
237: 48 83 c6 08 addq $8, %rsi
23b: 48 8d bd a8 fe ff ff leaq -344(%rbp), %rdi
242: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x247>
247: 4c 89 f7 movq %r14, %rdi
24a: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x24F>
24f: 48 8b 7d d0 movq -48(%rbp), %rdi
253: e8 00 00 00 00 callq 0 <__Z18ostream_testv+0x258>
258: 0f 1f 84 00 00 00 00 nopl (%rax,%rax)

```

The code generated for the `format_test` function that uses the reference implementation of the `format` function described in this proposal is several times smaller than the `ostream` code and only 40% larger than the one generated for `sprintf` which is a moderate price to pay for full type and memory safety.

The following factors contribute to the difference in binary code size between `format` and `sprintf`:

- Passing format string as `string_view` instead of `const char*`.
- Using `string` instead of a `char` buffer.
- Preparing the array of formatting arguments.

We can exclude the first two factors from the experiment by mimicking parts of the `sprintf` API:

```
int vraw_format(char* buffer, const char* format, format_args args);
```

```

template<typename... Args>
inline int raw_format(char* buffer, const char* format, const Args&... args) {
    return vraw_format(buffer, format, make_format_args(args...));
}

```

```

void raw_format_test() {
    char buffer[100];
    raw_format(buffer, "The answer is {}.", 42);
}

```

```

__Z15raw_format_testv:
0: 55 pushq %rbp
1: 48 89 e5 movq %rsp, %rbp
4: 48 81 ec 80 00 00 00 subq $128, %rsp
b: 48 8b 05 00 00 00 00 movq (%rip), %rax
12: 48 8b 00 movq (%rax), %rax
15: 48 89 45 f8 movq %rax, -8(%rbp)
19: 48 c7 45 80 2a 00 00 00 movq $42, -128(%rbp)
21: 48 8d 35 24 12 00 00 leaq 4644(%rip), %rsi
28: 48 8d 7d 90 leaq -112(%rbp), %rdi
2c: 48 8d 4d 80 leaq -128(%rbp), %rcx
30: ba 02 00 00 00 movl $2, %edx
35: e8 00 00 00 00 callq 0 <__Z15raw_format_testv+0x3A>
3a: 48 8b 05 00 00 00 00 movq (%rip), %rax
41: 48 8b 00 movq (%rax), %rax
44: 48 3b 45 f8 cmpq -8(%rbp), %rax
48: 75 09 jne 9 <__Z15raw_format_testv+0x53>
4a: 48 81 c4 80 00 00 00 addq $128, %rsp
51: 5d popq %rbp
52: c3 retq
53: e8 00 00 00 00 callq 0 <__Z15raw_format_testv+0x58>
58: 0f 1f 84 00 00 00 00 nopl (%rax,%rax)

```

This shows that passing formatting arguments adds very little overhead and is comparable with `sprintf`.