

Document number: P1063R2

Date: 2019-01-16

Reply To:

Geoff Romer (gromer@google.com)

James Dennett (jdennett@google.com)

Chandler Carruth (chandlerc@google.com)

Audience: WG21 Evolution Working Group

Core Coroutines

Making coroutines simpler, faster, and more general

[Introduction](#)

[Non-goals](#)

[Proposed Design](#)

[Unwrap operator syntax](#)

[Return syntax](#)

[Coroutine semantics](#)

[Coroutine lambdas](#)

[Expository implementation](#)

[Coroutine lambda size and layout](#)

[Coroutine functions](#)

[Tail calls](#)

[constexpr](#)

[Comparison](#)

[Conclusion](#)

[Acknowledgements](#)

[Revision History](#)

[Appendix: Examples](#)

[Futures](#)

[Simple generator](#)

[Zero-allocation generator](#)

[Parser Combinators](#)

Introduction

“C++ is a general-purpose programming language providing a direct and efficient model of hardware combined with facilities for defining lightweight and efficient abstractions.

Or terser:

C++ is a language for developing and using elegant and efficient abstractions.”

— Bjarne Stroustrup, *The C++ Programming Language (4th Edition)*

The Coroutines TS provides users with an elegant and efficient abstraction for writing asynchronous code. We mean that as both sincere praise, and as a critique: the Coroutines TS provides an abstraction, but it does not provide programmers with the facilities they need to define their own elegant and efficient abstractions. Furthermore, the TS’s abstraction prioritizes the asynchronous use case in a variety of ways that prevent it from being general-purpose. It gives programmers ways of extending and reusing the asynchrony abstraction, but they remain locked into many of the design tradeoffs motivated by the original use case.

Fundamentally, the Coroutines TS does not provide a direct and efficient model of hardware¹: the primitive objects and operations that are used to implement coroutines are hidden behind an abstraction boundary.

Nearly all of the serious issues we identified in [P0973R0](#) are reflections of this problem:

- Programmers cannot reliably prevent coroutine-based code from allocating memory, even if they know the allocation is unnecessary, because the allocation takes place behind the abstraction boundary.
- Programmers cannot control variable-capture semantics, and can all too easily overlook them entirely, because the capture is hidden behind an interface that presents itself as a function call.
- The library bindings are extremely complex because different abstractions require the underlying primitives to be composed in different ways. The TS supports this by providing APIs for the programmer to *configure how they are composed*, rather than permitting the programmer to *write code that composes them*.
- The `co_await` keyword is an overt manifestation of the TS’s preference for the asynchronous use case.

¹ We suggest as a friendly amendment that the quote should say “a direct and efficient model of *the platform*”. For example, C++ templates provide a direct and efficient map of the compiler’s code generation facilities, rather than of any hardware feature.

In this paper, we propose exposing a minimal set of coroutine primitives that map directly to the underlying implementation. This results in a design for coroutines that is substantially simpler and yet can efficiently support a broader range of uses.

Non-goals

This proposal is solely concerned with “stackless” coroutines, and does not address the kinds of problems that are solved by “stackful” coroutines. We fully support the committee’s decision to pursue stackless and stackful coroutines independently.

This proposal does not attempt to extend coroutines to be a fully general monad facility. For programmers who wish to adopt a monadic approach, both our proposal and the TS are limited to supporting *linear monads*, because they do not support copying a suspended coroutine frame, and consequently do not support nondeterministically resuming from the same state with multiple inputs.

Proposed Design

Unwrap operator syntax

We propose replacing the `co_await` keyword with an operator-like token, which we tentatively suggest spelling `[<-]` (we are very open to committee feedback on the spelling):

```
optional<string> f();  
string s = [<-] f();
```

```
future<string> g();  
string s = [<-] g();
```

```
expected<string> h();  
string s = [<-] h();
```

Our proposed spelling is intended to suggest *unwrapping*, which we regard as the most central meaning of these expressions. Correspondingly, we propose to refer to them as *unwrap expressions* rather than “await expressions”, and we refer to the operand of an unwrap expression as a *wrapper* (and its type as a *wrapper type*).

An operator-like token has two major advantages over an English-derived keyword:

- An operator can more easily avoid tying itself to a particular use case, as `co_await` is tied to asynchrony.

- An operator need not choose between colliding with existing identifiers in user code, or being so awkwardly spelled that no existing code uses it. It must avoid colliding with existing C++ syntax, but that's a far more manageable problem.

Option: We could also introduce a binary operator analogous to `->`, such that `x op y` is equivalent to `([<-]x).y`. This would make it easier to chain applications of the unwrap operator.

Alternative: The unwrap token could be a suffix, rather than a prefix. This has the advantage of naturally supporting chaining:

```
optional_struct[->].optional_sub_struct[->].field
```

However, this would depart from C++ convention (unary operators are generally prefixes), and could reduce readability by making the token less prominent.

Alternative: A different keyword spelling could be less use-case-specific than `co_await`. However, any keyword will still suffer from the need to avoid collisions with identifiers, and it is doubtful if any keyword can fully capture the breadth of possible use cases. Our best suggestion along these lines would be something like `co_unwrap`, but unwrapping is not the sole meaning of this operation; just the most central one. For example, it's a poor fit for unidirectional generators:

```
co_unwrap std::yield(foo); // Huh?
```

Return syntax

We propose to allow `return` statements in coroutines, and in fact we do not propose any syntax other than `return` (e.g. `co_return`) for returning from a coroutine. Under our proposal coroutines can always be distinguished from ordinary functions without inspecting the function body, so `co_return` is never needed for disambiguation. While `co_return` arguably acts as a signal that the return semantics of coroutines are different, those differences are observable only *outside* the coroutine body; within the body, `co_return` plays exactly the same role as `return` does in an ordinary function. Thus, forbidding `return` would be an example of what [N4074](#) memorably described as

“... the (arguably most) hated category of C++ compiler diagnostics: ‘I know exactly what you meant. My error message even tells you exactly what you must type. But I will make *you* type it.’ ”

Coroutine semantics

Consider how an expression `[<-] x` (or `co_await x` in the TS) is evaluated: the state of the enclosing coroutine is reified as an object, and passed to an algorithm that is controlled by the library associated with `x`. That algorithm may eventually do two things:

1. return a value to the coroutine's caller, and

2. resume the coroutine, specifying the value of the `co_await` expression.

The first is mandatory and synchronous, whereas the second is optional and may be synchronous or asynchronous.

We propose to allow the library to implement that entire algorithm directly in C++ code. Specifically, the library will define an `operator[<-]` overload which takes a coroutine object and returns the value that is returned to the coroutine's caller, while arranging for the coroutine to be resumed in whatever manner is appropriate to the library. This approach is motivated by the observation that C++ code is a far simpler and more general way to specify an algorithm than overloading the ~15 extension points of a fixed algorithm specified by the standard.

In addition to specifying the semantics of unwrapping, a wrapper type must also specify the semantics of returning; we propose allowing it to do so by defining an overload of `operator return`.

Both functions take as input an object we will refer to as the *shared state*. This an object associated with the coroutine's return type that lives as long as the coroutine², can store long-lived state associated with the wrapper, and acts as a "hook" for name lookup to find the aforementioned functions. Thus, it corresponds closely to the concept of a "promise" from the Coroutines TS, but we have chosen not to adopt that term because it is too closely tied to the asynchronous use case.

Here's an example of how these functions might be used to define the semantics of a coroutine that returns `expected<T,E>`:

```
template <typename T, typename E>
struct expected_shared_state {
    template <template<typename> Continuation, typename U>
    expected<T, E> operator[<-] (
        const expected<U,E>& e, Continuation<const T>&& continuation) {
        if (e.has_value()) {
            tail return continuation([&] { return *e; });
        } else {
            return unexpected(e.error());
        }
    }
};

expected<T, E> operator return(const T& value) {
    return value;
};
```

² In a subsequent revision of this paper, we intend to propose that the shared state be a sub-object of the coroutine, which will greatly simplify the library code that manages the shared state (among many other advantages).

(See [below](#) for a discussion of `tail return`)

As shown in this example, these functions can be members of the shared state (which in this case is stateless, and is being used only as a name lookup hook), but they can also be free functions that take the shared state as the first parameter (the two are treated as equivalent for purposes of name lookup and overload resolution, as with the operators that can be overloaded as members or non-members). Both overloads also take parameters representing the operands of the corresponding operation, although `operator return` can also be overloaded with no additional parameters, to handle `return` statements with no operand³.

In addition, `operator[<-]` takes a parameter representing the *continuation*⁴, a compiler-generated function object which can be called to resume the coroutine (this plays a role akin to `coroutine_handle`, but without the mandatory type erasure). The continuation takes an argument specifying the value of the unwrap expression, and its type takes the form of a specialization of a template template parameter, whose template argument specifies the type of the unwrap expression. It must be represented this way so that the type of the unwrap expression can be determined statically⁵.

The value of the unwrap expression is passed inside a callback, rather than as an explicit value, so that it never needs to be represented as a temporary in the stack frame of the caller, where it could inhibit [tail call elimination](#). This also gives wrapper type authors a means of resuming a coroutine with an exception, as though the unwrap expression had thrown it. Note that we could save the caller the trouble of embedding the value in a callback if this parameter were a lazy parameter, as proposed by [P0927R0](#).

The return type of the continuation is called the *suspension type*, because its primary role is to convey information about the suspended state of the coroutine to the library code that invoked the continuation. The suspension type is determined by the return types of every `operator[<-]` and `operator return` call generated for the coroutine, which must be identical other than cv-qualification in order to facilitate [tail call elimination](#). Note that in this simple example, the suspension type happens to also be the return type of the coroutine, but that is not the case in general.

³ "Falling off the end" of a function behaves like an explicit `return` statement with no operand (or an operand that is a call to a function with return type `void`), except that if no suitable `operator return` overload is found, the behavior is undefined (whereas an explicit `return` statement would be ill-formed). This corresponds exactly to the semantics of falling off the end of an ordinary function, with the existence of a no-operand `operator return` overload taking the place of a `void` return type.

⁴ We no longer propose a separate parameter representing the suspension point (as in R0), because we no longer believe the performance benefits of that approach (if any) will be significant enough to justify the extra complexity.

⁵ We no longer propose a separate overload to convey this information (as in R1), because that made it inordinately difficult for the library to pass information through from suspension to resumption. It also raised difficult questions about how to spell the two overloads without obscuring their connection to the `[<-]` operation.

Thus, the example above says that if `e` holds a `T` value, the coroutine is resumed, with `*e` as the value of the unwrap expression. This happens synchronously as part of the `[<-]` operation, so even though the coroutine returns when the `[<-]` operation returns, the effect is as if the `[<-]` operation simply returned control to the coroutine, which then returns normally. On the other hand, if `e` holds an error, that error is returned immediately to the coroutine's caller, and consequently the coroutine returns immediately, and the remainder of the coroutine is never executed.

Note that when evaluating an unwrap expression, control leaves the enclosing coroutine (without exiting any scopes) before `operator[<-]` is invoked. Consequently, if `operator[<-]` throws an exception, the coroutine will not be found during stack unwinding. An unwrap expression can only throw if its operand throws, or if the callback passed to the continuation throws.

Note that for simplicity, the example above glosses over the issue of qualifiers on the `expected<T,E>` object: like `*e`, `[<-] e` should be mutable if and only if `e` is mutable, and should be an rvalue if and only if `e` is an rvalue. This can be accomplished via a set of four overloads (with the unwrapped return type qualified to match the parameter), and/or perfect forwarding (with the unwrapped return type computed via a metafunction such as [P0847R0](#)'s `like_t`).

Coroutine lambdas

A fundamental distinction between a C++ coroutine and an ordinary function is that the state of a running coroutine (i.e. the coroutine frame) is effectively an object: it has storage and a lifetime, and provides operations that the program can invoke. However, the Coroutines TS does not expose coroutines as objects; instead, it creates them implicitly (via another core-language algorithm with its own extension points) and exposes only a type-erased handle.

We propose instead to make coroutines fully-fledged objects, which can be created and managed in the same way as any other object. Lambda syntax has exactly the properties we need for this purpose: it lets us define a callable object from a function body, and allows us to explicitly specify capture semantics. A coroutine lambda is distinguished from an ordinary lambda by the fact that it specifies a wrapper type rather than a return type, using the `[->]` syntax:

```
future<string> foo();
future<int> bar();
...
auto my_coroutine = [] [->] future<int> {
    int i = ([<-]foo()).size();
    return i + [<-]bar();
};
```

A coroutine lambda is much like an ordinary lambda, except that its state includes not only its captures, but also any local variables that must be preserved across suspensions. Similarly, it exposes not only a call operator for the initial invocation, but also one for resuming execution. As discussed earlier, the return type of both call operators is the suspension type, which is determined from the `operator[<-]` and `operator return` overloads invoked by the coroutine. The call operator that begins execution takes the shared state as a parameter (with a type determined by the `shared_state_type` member of the wrapper type, e.g. `future<int>::shared_state_type` in the example above), essentially as a workaround for the lack of a viable constructor syntax for lambdas. We have already described the semantics of the call operator for resumption, because the continuation passed to `operator[<-]` is just a coroutine lambda: it takes a callback representing the value of the unwrap expression the coroutine is suspended at.

The call operators are never `const` (in effect, coroutine lambdas are always implicitly `mutable`); in principle we could allow the user to specify or omit `mutable` as with ordinary lambdas, but in practice `mutable` would just be boilerplate, since it would only be correct and safe to omit it in cases where the coroutine has effectively no mutable stack variables, which we expect to be rare and marginal.

Coroutine lambdas cannot have user-defined parameters. This is for reasons of safety: the code in a coroutine may continue executing after the initial function call has returned (from the caller's point of view), so if any temporary values were passed to pointer or reference parameters of the coroutine, they would be left dangling. The inputs to a coroutine lambda are instead expressed via the capture group. See [below](#) for how coroutine lambdas can be used to define ordinary functions with parameters, etc.

Modeling coroutines as lambdas rather than functions has two major benefits: first, it enables library code to control the creation, usage, and destruction of coroutine frames in exactly the same way as any other object (and in particular, allows the creation of coroutine libraries that are allocation-free by construction, rather than at the whim of the optimizer). Second, the capture syntax gives programmers explicit control over capture semantics (in the Coroutines TS, capture semantics are controlled by the parameter types, but parameter types are API-visible, and so API owners are not always at liberty to change them). Use of capture syntax also leverages programmers' existing intuitions: reference and pointer inputs to a coroutine are potentially hazardous in the same way, and for the same reasons, as the reference and pointer captures of an ordinary lambda.

Note that exceptions have no special semantics inside a coroutine: any exception that isn't caught in the body of the coroutine will propagate to the caller that resumed the coroutine (which will typically be library code associated with the coroutine, so this shouldn't have any major functional effects).

Alternative: we could specify that exceptions that escape the coroutine are caught and forwarded to an extension point comparable to `operator return`. This would provide some minor benefits (primarily, greater consistency in how exceptions and ordinary returns are propagated), but also some minor drawbacks: it complicates the API, and we would not be able to handle exceptions thrown from tail calls (i.e. unwrap expressions and return statements), which may be surprising, and may limit the consistency benefits.

Expository implementation

The following example illustrates how a compiler might generate equivalent C++17 code for a given coroutine. Of course, this is not how we expect coroutine compilation to actually work, but it can serve as a “reference implementation” to understand the API and behavior of coroutine objects.

Consider the following code:

```
expected<string, Err> foo(const string& s);
expected<int, Err> bar();

void f(const string& s) {
    auto coroutine = [&s] [->] expected<int, Err> {
        int i = ([<-]foo(s)).size();
        return i + [<-]bar();
    };
}
```

The compiler could implement that by generating the following code:

```
// Convenience helper shared by all coroutine implementations
template <typename T>
class __manual_lifetime {
    std::aligned_storage_t<sizeof(T), alignof(T)> storage_;

public:
    template <typename... Args>
    void emplace(Args&&... args) {
        new (&storage_) (std::forward<Args>(args)...);
    }

    T& get() { return *reinterpret_cast<T*>(&storage_); }

    void destroy() {
        get().~T();
    }
};
```

```

// The generated type of the coroutine lambda
class _f_1 {
    using wrapped_return_type = expected<int, Err>;
    using suspension_type = expected<int, Err>;
    using shared_state_type = typename wrapped_return_type::shared_state_type;
public:
    _f_1(const _f_1&) = delete;
    _f_1(_f_1&&) = delete;
    _f_1& operator=(const _f_1&) = delete;
    _f_1& operator=(_f_1&&) = delete;

    // Beginning of execution
    suspension_type operator()(shared_state_type shared_state) {
        __shared_state.emplace(std::move(shared_state));
        __tmp_1.emplace(foo(s));
        __suspend_point = 1;
        tail return __shared_state.get().operator[<-](__tmp_1.get(), *this);
    }

    template <typename F>
    suspension_type operator()(F&& value_callback) {
        switch (__suspend_point) {
            case 1:
                i.emplace(value_callback());
                __tmp_1.destroy();
                __tmp_2.emplace(bar());
                __suspend_point = 2;
                tail return __shared_state.get().operator[<-](__tmp_2.get(), *this);

            case 2:
                int __result = i.get() + value_callback();
                __tmp_2.destroy();
                __suspend_point = 3;
                i.destroy();
                return __shared_state.get().operator return(__result);
        }
    }
}

~_f_1() {
    switch (__suspend_point) {
        case 1:
            __tmp_1.destroy();
            break;
        case 2:
            __tmp_2.destroy();
            i.destroy();
            break;
    }
}

```

```

        case 0:
        case 3:
            break;
    }
    __shared_state.destroy();
}

private:
    // Implicitly invoked via lambda capture syntax
    _f_1(const string& s) : s(s), __suspend_point(0) {}

    __manual_lifetime<shared_state_type> __shared_state;

    size_t __suspend_point;

    // Captures
    const string& s;

    // Stack variables
    //
    // The layout of these members is purely illustrative; in practice we expect
    // the compiler to lay out this class using the same algorithms it uses to
    // lay out ordinary stack frames.
    __manual_lifetime<int> i;

    union {
        __manual_lifetime<expected<string, Err>> __tmp_1;
        __manual_lifetime<expected<int, Err>> __tmp_2;
    };
};

void f(const string& s) {
    auto coroutine = _f_1(s);
}

```

Some notes on how we present this implementation:

- The generated code must also ensure that the local-variable members are suitably destroyed, and the object is left in a destructible state, if an exception escapes the coroutine body. This logic is omitted from the above code for clarity and simplicity.
- See [below](#) for a discussion of `tail return`, and note in particular that the example code above presumes that the functions being invoked are defined in the current translation unit; otherwise ordinary `return` would be used.

Note that the coroutine transformation does not affect the existing rules for the sequenced-before relation; if an `unwrap` expression and some other operation are

unsequenced, the latter operation may be evaluated before the unwrap expression, or it may not be evaluated until the continuation is resumed (which, of course, may never happen).

Coroutine lambda size and layout

It is vitally important that the compiler be free to apply the same transformations to a coroutine as it would to an ordinary function, which means that it must be able to apply the same transformations to the coroutine lambda object layout that it would apply to a stack frame.

However, we see two closely-related obstacles to that goal:

- Within a single translation unit, earlier stages of compilation may need to know the size of a coroutine lambda type, e.g. to determine the result of `sizeof` or to determine the layout of an object containing a coroutine lambda. The optimizer then would not be able to apply transformations (e.g. inlining) if they would cause the coroutine frame to exceed that size.
- When compiling and linking multiple translation units that refer to the same coroutine lambda type, those translation units may need to agree on the size and layout of that type (e.g. if a coroutine lambda is created in one TU and passed to a function defined in another TU, or if it appears in an inline function or template instantiation). Such agreement is typically achieved by having the platform ABI fix the exact size and layout of the type, which would leave the optimizer no flexibility to change those properties.

The coroutines TS avoids these problems by ensuring that coroutine frames are not objects and do not have types, and providing a layer of indirection when it's necessary for them to have object-like behavior. For example, opaque `coroutine_handle` objects permit you to refer to specific coroutine frames, manage their lifetime, and even take their address (or rather, obtain a `void*` value that uniquely identifies a coroutine frame in some unspecified manner, via a method with the oddly suggestive name `address`), and special rules not only ensure that coroutine frames are allocated via `operator new` (despite not being objects), but even ensure that their allocation can be customized via the type system (despite not having types).

We propose instead to take a more minimalist approach, treating coroutine lambdas as genuine objects with genuine types, and departing from the C++ type system and object model only to the extent needed to support optimization.

Specifically, we propose defining a new category of types, tentatively called *hidden-layout types*, which includes coroutine lambda types as well as the types of any objects that have sub-objects of hidden-layout type. A hidden-layout type `T` is subject to the following restrictions:

- `T` cannot be a base class.
- `T` cannot be the element type of an array.
- `T` cannot be the type of a non-static data member, unless it is the last non-static data member of the class.
- If a `sizeof` expression has operand type `T`, it is not a constant expression.

These restrictions are not negligible, but neither are they prohibitive, and they appear to be sufficient to address the issues that arise within a single translation unit. It is worth noting that these are substantially the same restrictions that apply to most implementations of C-style variable-length arrays, and indeed we expect coroutine lambdas to be laid out in much the same way (namely, via tail-padded allocation).

The issues that arise across translation units are more challenging; we believe they are tractable, but do not yet have a solution in hand. One possible approach is to generate code such that all operations that depend on the layout of a hidden-layout type are expressed as function calls, which are resolved by the linker to a single TU that's canonical for that type. However, that would pose its own challenges for the optimizer because those calls couldn't normally be inlined except in the canonical TU (and then only if the optimizer can tell that the current TU will be chosen as canonical).

Another possible approach is to allow separate TUs to have separate sizes and layouts for the hidden-layout type. This seems particularly promising for cases such as inline functions and template instantiations, which are typically implemented by generating the code independently for each TU. Although some use cases preclude that, those cases are relatively unusual (e.g. static local variables, which must have a consistent address in all TUs), and they can often be worked around.

These approaches complement each other well: the first provides a flexible but potentially costly mechanism for different TUs to agree on the layout of a hidden-layout type, and the second focuses on minimizing the situations where such agreement is needed in the first place. The primary open question here is whether there are any usages where agreement is necessary, but unacceptably costly, and whether those situations are rare/avoidable enough that we can feasibly make them ill-formed.

Coroutine functions

We expect that in most use cases, coroutine lambdas will not be part of public APIs; instead, they will be hidden implementation details of ordinary functions, which wrap the coroutine lambdas to handle issues such as parameter passing/capture, lifetime management, and whether to defer initial invocation of the lambda. We propose a sugar syntax for defining such functions, in order to mitigate the associated boilerplate.

As a motivating example, consider this asynchronous function:

```
// Consumes all bytes from `connection`, and returns the number
// of bytes consumed. `connection` must remain live until the returned
// future is ready.
auto count_bytes(Connection& connection) {
    return future<int>([&] {
```

```

return [&connection] [->] future<int> {
    int bytes_read = 0;
    vector<char> buffer(1024);
    while(!connection.done()) {
        bytes_read += [<-]connection.Read(buffer.data(), buffer.size());
    }
    return bytes_read;
}
}));
}

```

The intent here is for `count_bytes` to construct and return a `future<int>` representing the result of executing the coroutine lambda's body. We expect this to be a very general pattern for all or nearly all coroutine functions (notice that one consequence of this design is that the suspension type becomes purely an implementation detail of the library).

However, the meaning of this code is obscured by some troublesome boilerplate:

- The coroutine lambda must be wrapped in an ordinary lambda, so that the constructor can control how the coroutine lambda object is allocated. This problem could instead be addressed by having `future`'s constructor take a lazy parameter as proposed by [P0927](#).
- The wrapper type `future<int>` must be named in two places, the explicit constructor invocation and the wrapper type of the coroutine. Neither can realistically be deduced from the other (although the explicit constructor invocation might be able to omit the template argument, if `future` provides suitable deduction guides).
- The coroutine body proper is nested inside three levels of braces.

We propose a sugar syntax for defining such a coroutine function, in which the function body (including braces) is replaced by a coroutine lambda expression:

```

auto count_bytes(Connection& connection) [&connection] [->] future<int> {
    int bytes_read = 0;
    vector<char> buffer(1024);
    while(!connection.done()) {
        bytes_read += [<-]connection.Read(buffer.data(), buffer.size());
    }
    return bytes_read;
}

```

This syntax can be specified as a pure rewrite to the form shown above: if a coroutine lambda expression `L` with wrapper type `R` appears in place of a function body, it behaves as if the function body were `{ return R([&] { return L; }); }`.

We contend that this syntax contains almost no boilerplate other than a smattering of punctuation. The additional syntactic elements not present in the Coroutines TS all have important, user-facing functional roles:

- The capture group specifies the capture semantics of the coroutine object.
- `[->]` acts as an introducer, specifying that the following block is a coroutine.
- `future<int>` specifies what kind of coroutine function this is, including its return type.

In all three cases, making these properties syntactically explicit has important advantages:

- The programmer has explicit control over capture behavior, so that for example an argument can be captured by value (for safety) even if the API is obliged to pass by reference. Symmetrically, the capture behavior is explicitly visible in the code, cueing the reader (and programmer) to possible safety or performance concerns.
- The explicit introducer enables both the reader and the compiler to immediately and reliably recognize coroutine code. This eliminates the need for a separate `co_return` syntax to cue the compiler that it's processing a coroutine.
- The programmer has explicit, local control over what kind of coroutine is being defined, even if they do not control the function signature, e.g. because they must match an existing API (In the Coroutines TS, this can be controlled only via a trait parameterized by the parameter and return types). Symmetrically, the reader can easily tell what kind of coroutine they are reading.

In order to support this pattern, most wrapper types will need a constructor taking a single generic argument, which is interpreted as a callback that returns a coroutine. It will often be important to constrain such a constructor to avoid matching arbitrary single arguments, so we propose introducing a type trait `std::is_coroutine`, which is true only for coroutine lambda types, to facilitate such constraints.

Although we do not propose it here, this design could plausibly be extended to support deduction of the explicit template argument `int` from the coroutine body. This is somewhat non-trivial because the semantics of the coroutine body technically depend on the concrete wrapper type of the coroutine, which would make ordinary return type deduction circular. However, we believe that for non-pathological wrapper types, the deduced return type will not have a logical dependency on the template argument, and so we can perform ordinary return type deduction using an arbitrarily-chosen type argument as a placeholder, and then replace it with the result of the deduction.

It is less clear whether we can deduce the *entire* wrapper type (not just the template argument `int`, but the template `future`). We will not go into that issue here except to say that the Coroutines TS currently does not support such deduction either, and the proposed solutions that we are aware of can be applied equally well to either proposal.

Tail calls

Consider a coroutine like the following:

```
[&connection] [->] expected<int, Err> {  
    int bytes_read = 0;  
    vector<char> buffer(1024);  
    while(!connection.done()) {  
        bytes_read += [<-]connection.Read(buffer.data(), buffer.size());  
    }  
    return bytes_read;  
}
```

With the design described above, the unbounded iteration in this code will be transformed into an unbounded *recursion*, raising obvious concerns about stack size. However, the mutual recursion between `operator[<-]` and the generated coroutine code is actually all tail recursion, because every mutually recursive call is actually the final operation before the enclosing function returns. Consequently, the compiler should be able to apply tail call elimination (hereinafter “TCE”) to avoid growing the stack.

For this approach to be viable, programmers will need to have complete confidence that TCE will in fact be applied (even in non-optimizing build modes). Consequently, we propose to standardize TCE as a C++ feature. Although obviously motivated by the coroutines use case, this proposal is completely independent of coroutines.

There are two reasons TCE is difficult to achieve in C++:

- There is currently no way for the standard to specify that TCE will take place, because the standard has no explicit concept of stack storage as a finite resource.
- It’s not as easy as it seems to determine whether a call is eligible for TCE in the first place. For example, a statement of the form `return f(...);` is nevertheless ineligible if there are any local variables with nontrivial destructors still live at that point (because then the function call is not actually the last operation before the return), or if the `f()` call takes a pointer or reference to any local variable. This is not an issue for `operator[<-]` calls inside the coroutine generated code (because the compiler can ensure that it’s able to apply TCE to the code it generates), but it is an issue when `operator[<-]` invokes the continuation synchronously.

To address the first issue, we propose adding standard wording such as the following:

“If this International Standard specifies that a function invocation is a *tail call*, then before entering that invocation, the implementation must disregard the invoking function call for purposes of enforcing any implementation-defined limits concerning the number of simultaneously active function calls, or the number or size of simultaneously-live variables with automatic storage duration. [Note: The effect of this requirement is that on

implementations with a bounded stack, a tail call must reuse the stack frame of the calling function. — *end note*]

To address the second issue, we propose introducing a new syntax `tail return`, which requires its operand to be a tail call (`tail` is a contextual keyword, with a special meaning only when followed by `return`, so this should not break any existing code). This would be both a constraint on the operand (to make it eligible for TCE) and a requirement on the implementation (to apply TCE). The standard wording would be something like the following:

If a `return` statement is preceded by `tail`, then evaluation of its operand will be a tail call, and the program is ill-formed if:

- the statement is within a *try-block* or *function-try-block*,
- any live object with automatic storage duration within the scope of the enclosing function has a non-trivial destructor, or its address is taken or it is bound to a reference (including the implicit object parameter of a member function) anywhere within the function body,
- the operand is not a function call expression,
- the operand is a function call expression whose *postfix-expression* has a function pointer type,
- the operand is a call to a virtual function that is not named by a *qualified-id*, or
- the function designated by the function call expression is not defined in the current translation unit, or has a return type that is not the same as the return type of the calling function (ignoring cv-qualifiers), or has a *parameter-declaration-clause* that terminates with an ellipsis.

We believe that the above conditions are minimally sufficient to permit TCE in Clang, and probably in any other reasonable C++ implementation (of course, we particularly welcome implementer feedback on this point). Note that the generated code for a coroutine lambda can easily ensure that all these conditions hold for its invocations of `operator[<-]`, except that it cannot guarantee that the operator is defined in the current translation unit. We will therefore specify that invocation of `operator[<-]` by a coroutine lambda is always a tail call *if* the selected overload is defined in the current translation unit.

Note that the above rules do not permit a function invoked via operator syntax (other than an `operator()` overload) to be a tail call. This is for reasons of readability: a statement like `tail return *foo();` is apt to mislead the reader into thinking that `foo()` is the tail call, rather than `operator*()`.

Possible extension: we could loosen the above rules somewhat to permit taking addresses of and forming references to local variables, but specify that the lifetime of local variables ends when the tail call begins (since we forbid nontrivial destructors, the effect of this is just that it's UB to access them after that point). However, that would make this construct less safe, since

changing `return` to `tail return` could break code in ways that can't be detected at compile time.

Alternative: we could achieve the same behavior via an attribute, e.g. `[[tail_call]]`. This would be more conceptually lightweight than a new contextual keyword, correctly signalling to programmers that they can disregard this feature unless they have a specific need for it. However, an attribute might not allow us to normatively mandate TCE, which we believe is necessary.

Alternative: rather than allow users to force TCE, we could make it inherent in the API for unwrap expressions. Specifically, we could allow `operator[<-]` to return either the suspension type of the coroutine, or a nullary callback returning the same type as `operator[<-]`. The generated code would then apply a "trampoline" technique, repeatedly checking if the result is a callback, and if so invoking it to obtain a new result, until it obtains an instance of the suspension type. However, this would substantially complicate the `operator[<-]` API (notice for example that it makes the return type of `operator[<-]` self-referential), and would not have the benefit of allowing TCE in non-coroutine contexts. Note also that the library may be able to implement this technique under our existing proposal.

constexpr

We have not worked through this issue in detail, but we see no obstacles to allowing coroutines to be `constexpr` (and uses of them to be core constant expressions) on the same terms as ordinary functions. The sample implementation given above cannot be `constexpr` because of its use of `reinterpret_cast`, but that is only as an expository way of depicting the compiler's management of the stack frame, which we know it can do in `constexpr` code because it already does.

Comparison

The following chart summarizes what we see as the key functional differences between the Coroutines TS status quo and our proposal:

	Coroutines TS	Core coroutines
Library customization points	15: <code>await_transform</code> <code>operator co_await</code> <code>await_ready</code> <code>await_suspend</code> <code>await_resume</code> <code>yield_value</code>	3: <code>operator[<-]</code> <code>operator return</code> <code>shared_state_type</code>

	return_value return_void promise_type get_return_object get_return_object_on_allocation_failure coroutine_traits initial_suspend final_suspend unhandled_exception	
Coroutine object representation	Type-erased as <code>coroutine_handle</code>	Concrete object with anonymous type
Coroutine allocation (normative)	All coroutine objects are heap-allocated by default. This can be disabled by explicit collaboration between library and user code.	Coroutine objects are allocated by explicit code, just like all other objects. Allocation will normally be a hidden detail of the library.
Coroutine allocation (QoI)	Optimizers have demonstrated ability to elide coroutine allocations in many common cases. Techniques sufficient to reliably elide allocation for specific types are on the drawing board. Unclear whether optimizations will apply in all build modes.	Allocation elision applies equally to all kinds of objects, including coroutines.
(N)RVO in coroutines	No	No NRVO if there's a suspend point between construction and return.
Programmer control of capture	No	Yes
<code>return</code> in coroutines	Forbidden	Allowed without restriction (<code>co_return</code> is unnecessary)
User-facing syntax	Keyword, concurrency-specific	Operator token, general-purpose

Conclusion

C++ is a language that enables programmers to build powerful and efficient abstractions by composing simple primitives that are efficiently supported by the platform. This is a defining property of C++, and a cornerstone of its success, so we should not abandon it (or even postpone it) without extremely compelling reasons.

The current design of the Coroutines TS is not consistent with that principle, because it does not provide simple, composable primitives, but only a complex abstraction that is tuned for a particular kind of use case. Shipping the current design as part of a C++ IS would be either an outright rejection of that principle or, at best, a wholly unjustified gamble that we'll be able to add the necessary primitives as a non-breaking extension, and still end up with a coherent design.

We believe that C++ can still be a vital language 50 years from now, and the language should be designed with that goal in mind. In 50 years nobody will even remember whether coroutines shipped in C++20 or C++23, but if we lock ourselves into a coroutines design that lacks such an essential ingredient of C++'s success, the consequences could easily last that long.

We have shown that a revised design that accords with that principle is well within reach, and that the resulting facility will be simpler, more general, and more efficient. We therefore urge the committee not to merge the Coroutines TS into the IS in its current form, and instead to allow sufficient time for this design to be fleshed out and validated.

Acknowledgements

Many thanks to Richard Smith, Gor Nishanov, Roman Perepelitsa, Jeffrey Yasskin, Bryce Lebach, Michael Spencer, Davide Italiano and Gabriel Kerneis for their valuable design discussions and feedback.

Revision History

Changes since P1063R1:

- Reverted to overloading `operator[<-]` rather than a pair of functions `coroutine_suspend` and `coroutine_resume`. Instead, we require a special structure for the template and function parameters in order to convey the information formerly conveyed by the `coroutine_resume` return type.
- Correspondingly, renamed `coroutine_return` to `operator_return`.
- Added discussion of the size and layout of coroutine lambdas.
- Dropped discussion of incremental fixes to the TS, which are no longer timely.

Changes since P1063R0:

- Replaced `operator[<-]` with `coroutine_suspend` and `coroutine_resume`, thereby eliminating the novel return-type syntax, and correspondingly revised the coroutine lambda to not take the unwrapped value as an argument on resumption.
- Eliminated explicit passing of the suspend point; the coroutine generated code now tracks it internally.
- Renamed the "final return functor" to "shared state", replaced its call operator with a named function `coroutine_return`, dropped the requirement to pass it into the

coroutine on every resumption, and required the wrapper to specify the shared state type rather than vice-versa.

- Revised the coroutine lambda syntax, dropping the `do` keyword and adding a mandatory syntax for specifying the wrapper type.
- Proposed a coroutine-specific sugar syntax rather than one built out of general-purpose extensions.
- Dropped dependencies on [P0927R0](#).
- Dropped `raise()` operation on coroutine lambdas, which is superseded by `coroutine_resume`
- Fleshed out the specification of `tail return` based on implementation experience.
- Removed `no_alloc` member from alternative proposal; programmers can force allocation to be a build failure by deleting `operator new`.
- Added a parser combinator example.
- Miscellaneous copyediting, clarification, and improved exposition.

Appendix: Examples

Caveat: unless otherwise indicated, these examples are completely untested.

Futures

The following is a (very) rough implementation of a future library that supports coroutines. All types other than `promise` and `future` are hidden implementation details. This implementation leaks all shared states, in order to avoid a lot of distracting reference-counting machinery:

```
// Interface of all future shared states. This API should be sufficient
// to implement future<T>.
template <typename T>
class future_shared_state {
public:
    virtual bool is_ready() const = 0;
    virtual T& get() const = 0;

    virtual concrete_shared_state<T>& get_concrete() = 0;

    virtual ~future_shared_state() = 0;
};

// Interface of all promises.
template <typename T>
class promise_interface {
public:
    virtual void set_value(const T& value) = 0;
    virtual void set_exception(std::exception_ptr ptr) = 0;
    virtual ~promise_interface() = 0;
};

// Tag type representing a shared state that is not yet ready.
struct not_ready{};

// A shared state implementation for ordinary promise/future patterns.
template <typename T>
class concrete_shared_state
    : public promise_interface<T>, future_shared_state<T> {

    std::variant<not_ready, T, std::exception_ptr> state_{not_ready{}};

    std::function<void(void)> continuation_ = nullptr;
```

```

mutable std::mutex mu_;
mutable std::condition_variable done_;

public:
concrete_shared_state() = default;

bool is_ready() const override {
    std::lock_guard guard(mu_);
    return !std::holds_alternative<not_ready>(state_);
}

T& get() const override {
    std::lock_guard guard(mu_);
    done_.wait(guard, [&] {
        return !std::holds_alternative<not_ready>(state_);
    });
    return std::visit(overloaded(
        [] (not_ready) -> T& { std::abort(); },
        [] (T& t) -> T& { return t; },
        [] (std::exception_ptr ptr) -> T& { std::rethrow_exception(ptr); }),
        state_);
}

concrete_shared_state<T>& get_concrete() override {
    return *this;
}

template <typename U, typename Coroutine>
void fuse_to(concrete_shared_state<U>& continuation_shared_state,
             Coroutine& continuation) {
    {
        std::lock_guard guard(mu_);
        assert(continuation_ == nullptr);
        if (std::holds_alternative<not_ready>(state_)) {
            continuation_ = [&coroutine, this] {
                try {
                    continuation([this] { return get(); });
                } catch (...) {
                    continuation_shared_state.set_exception(std::current_exception());
                }
            };
            return;
        }
    }

    tail return coroutine();
}

```

```

void set_value(const T& value) override {
    std::function<void(void)> continuation = nullptr;
    {
        std::lock_guard guard(mu_);
        if (continuation_ == nullptr) {
            state_.template emplace<T>(value);
            done_.notify_all();
            return;
        }
        std::swap(continuation_, continuation);
    }
    continuation();
}

void set_exception(std::exception_ptr ptr) override {
    std::function<void(void)> continuation = nullptr;
    {
        std::lock_guard guard(mu_);
        if (!continuation_) {
            state_.template emplace<std::exception_ptr>(ptr);
            done_.notify_all();
            return;
        }
        std::swap(continuation_, continuation);
    }
    continuation();
}
};

template <typename T>
class coroutine_shared_state_base : public future_shared_state<T> {
    concrete_shared_state<T> shared_state_;
public:
    coroutine_shared_state_base() = default;

    template <template<typename> Continuation, typename U>
    void operator[<->](future<U>& f, Continuation<U>&& continuation) {
        auto& concrete = f.state_->get_concrete();
        tail return concrete.fuse_to(*this, continuation);
    }

    void operator return(const T& value) {
        shared_state_.set_value(value);
    }

    bool is_ready() const override { return shared_state_.is_ready(); }

```



```

T& get() const override { return shared_state_.get(); }
concrete_shared_state<T>& get_concrete() override {
    return shared_state_;
}

private:
    void set_exception(std::exception_ptr e) { shared_state_.set_exception(e); }
};

// A shared state co-allocated with a coroutine. Does not implement
// promise_interface, because the value is determined by running the coroutine.
//
// We catch all exceptions when starting the coroutine, and when resuming it
// asynchronously. We do not catch exceptions when resuming synchronously,
// because that would prevent tail call elimination; instead the exception
// propagates back to the start, and/or asynchronous resumption.
template <typename T, typename Coroutine>
class coroutine_shared_state : public coroutine_shared_state_base<T> {
    Coroutine coroutine_;
public:
    template <typename F>
    coroutine_shared_state(F&& coroutine_callback)
        : coroutine_(coroutine_callback()) {}

    void run() {
        // Begin execution of the coroutine, and return the first time it
        // blocks.
        try {
            coroutine_(*this);
        } catch (...) {
            this->set_exception(std::current_exception());
        }
    }
};

template <typename T>
class promise {
    // Invariant: if two promise objects have equal shared_state_ values, they are
    // both null.
    promise_interface<T>* shared_state_;

public:
    promise()
        : shared_state_(nullptr) {}

    promise(promise_interface<T>* shared_state)
        : shared_state_(shared_state) {}

```

```

promise(promise&& other)
    : shared_state_(other.shared_state_) {
    other.shared_state_ = nullptr;
}
promise& operator=(promise&& rhs) {
    shared_state_ = rhs.shared_state_;
    rhs.shared_state_ = nullptr;
    return *this;
}

explicit operator bool() { return shared_state_ != nullptr; }

void set_value(const T& value) {
    shared_state_->set_value(value);
}

void set_exception(std::exception_ptr ptr) {
    shared_state_->set_exception(ptr);
}
};

template <typename T>
class future {
    // Invariant: if two future objects have equal state_ values, they are both null
    future_shared_state<T>* state_;
public:
    using shared_state_type = coroutine_shared_state_base<T>&;

    future(const future&) = delete;
    future& operator=(const future&) = delete;

    template <typename F>
    future(F&& coroutine_callback) {
        auto state = new coroutine_shared_state<
            T, std::decay_t<decltype(coroutine_callback())>>(coroutine_callback);
        state->run();
        state_ = state;
    }

    // Public API left as exercise for reader
};

```

A typical usage, as shown earlier, could look like:

```

auto count_bytes(Connection& connection) [&connection] [->] future<int> {
    int bytes_read = 0;

```

```

vector<char> buffer(1024);
while (!connection.done()) {
    bytes_read += [<-]connection.Read(buffer.data(), buffer.size());
}
return bytes_read;
}

```

Simple generator

This example prints the contents of a binary tree in order, using a generator:

```

struct BstNode {
    BstNode* left, right;
    string value;
};

auto Traverse(BstNode<int>* node) [node] [->] generator<string> {
    if (node == nullptr) {
        return;
    }
    [<-] Traverse(node->left);
    [<-] std::yield(node->value);
    [<-] Traverse(node->right);
}

void PrintBst(BstNode* root) {
    generator<string> g = Traverse(root);
    while (g) {
        cout << *g << endl;
        g.next();
    }
}

```

And here's the implementation that supports it:

```

namespace std {
// yield_handle represents the result of a `yield` call. It has no semantics
// of its own; semantics are provided by the overloads for specific
// generators. Thus, all generators can use the same `yield` function.
template <typename T>
struct yield_handle {
    T& value;
};

template <typename T>
yield_handle<T> yield(T& value) {
    return {value};
}

```

```

}

template <typename T>
yield_handle<const T> yield(const T& value) {
    return {value};
}
} // namespace std

// The current state of a generator<T,P>. This is a hidden implementation
// detail, but it must be a namespace-scope template in order to facilitate
// deduction of T and P.
template <typename T, typename P>
struct generator_state {
    // The code to execute to resume this generator. Null if this generator
    // is done.
    std::function<generator_state(P&)> continuation = nullptr;

    // Pointer to the currently yielded value. Null if this generator is done.
    T* value = nullptr;
}

// generator<T, P> represents a bidirectional generator, i.e. that not only
// yields values of type T, but takes arguments of type P (which become values
// of the yield expression). Yielded values are accessed by dereferencing,
// and the generator is advanced to the next yielded value by calling next().
// Like an iterator, a generator has a special past-the-end state, signifying
// the end of the generated sequence, which cannot be dereferenced or advanced.
//
// The generator<T, void> specialization (which represents a traditional
// unidirectional generator) is omitted for brevity; the differences are
// mostly obvious, but note that it could easily implement MoveIterator
// (see P0902R0).
template <typename T, typename P = void>
class generator {
    generator_state<T,P> state_;

    // Manages lifetime of the coroutine lambda. Is not accessed otherwise.
    std::unique_ptr<void, void(*)(&void*)> coroutine_;

    class shared_state_type;
    std::unique_ptr<shared_state_type> shared_state_;

public:
    struct shared_state_type {
        // This overload defines the semantics of yielding from a generator<T,P>
        template<template Continuation<typename>>
        generator_state<T,P> operator[<->](

```

```

        yield_handle<T>& handle, Continuation<P>&& continuation) {
    tail return { [&continuation](P& yield_result) {
        continuation([&]() -> P& { return yield_result; })
    },
        &handle.value};
}

// Unwrapping a generator object behaves like python's `yield from`: next()
// operations on the outer generator are delegated to the inner generator
// until it is done, and then the outer generator's coroutine is resumed.
// Consequently, it does not return a value, even for bidirectional generators.
template <template<typename> Continuation, typename U, typename Q>
generator_state<T, P> operator<[-]>(
    generator<U, Q>&& inner_generator, Continuation<void>&& continuation) {
    if (inner_generator) {
        return { [&] (P& yield_result) {
            inner_generator.next(yield_result);
            tail return operator<[-]>(std::move(inner_generator), continuation);
        },
            &*inner_generator};
    } else {
        tail return continuation();
    }
}

generator_state<T,P> operator return() { return {}; }
};

friend void swap(generator& lhs, generator& rhs) {
    using std::swap;
    swap(lhs.state_, rhs.state_);
    swap(lhs.coroutine_, rhs.coroutine_);
}

// Move only
generator(generator&& rhs) { swap(*this, rhs); }
generator& operator=(generator&& rhs) { swap(*this, rhs); }

// Constructs a generator which exposes the values yielded by
// coroutine_callback().
template <typename F>
generator(F& coroutine_callback)
    : shared_state_(std::make_unique<shared_state_type>()) {
    using Coroutine = decltype(coroutine_callback());
    unique_ptr<Coroutine, void(*) (void*)> coroutine(
        new Coroutine(coroutine_callback()),
        +[] (void* ptr) { delete static_cast<Coroutine*>(ptr); });
}

```

```

    state_ = (*coroutine)(*shared_state_);
    coroutine_ = std::move(coroutine);
}

// Returns whether the generator is dereferenceable. False indicates
// the end of the generated sequence.
explicit operator bool() const { return state_.continuation != nullptr; }

// Accessors for the currently yielded value. static_cast<bool>(*this) must
// be true. Valid only until the following `next()` call.
T& operator*() { return *state_.value; }
T* operator->() { return state_.value; }

// Advance to the next yielded value. static_cast<bool>(*this) must be true.
void next(P& p) {
    auto continuation = std::move(state_.continuation);
    state_ = continuation(p);
}
};

```

Zero-allocation generator

The above generator is comparable to generators as proposed by the Coroutines TS; in particular, it allocates every coroutine frame on the heap, which is extremely inefficient in many cases. The following example shows a generator that always stores its state on the stack, which isn't possible with the Coroutines TS (without changing the signatures of generator functions). As a consequence of storing its state on the stack, generator functions defined this way cannot recurse (i.e. the maximum generator stack depth must be statically known).

It should be possible to use similar techniques to define a generator library that supports recursion by using a side stack (i.e. at most one more allocation than the corresponding non-generator-based recursive code), but the API design of the side stack abstraction raises issues beyond the scope of this paper.

Please note that this example is known to have some bugs, which will be corrected in a future revision. In the meantime, it is presented only as a suggestive illustration.

First, a usage example:

```

// Returns a generator whose output consists of the concatenated
// outputs of each generator produced by `generators`.
template <typename T, typename P>
auto flatten(stack_generator_base<stack_generator_base<T,P>>&& generators)
    [&] [->] stack_generator<T,P> {
    while (generators) {

```

```

    [<-]*generators;
    generators.next();
}
}

// Returns a generator that iterates over the given range.
template <typename Range>
auto traverser(const Range& range)
    [&] [->] stack_generator<decltype(*begin(range)), void> {
    for (auto& element: range) {
        [<-] std::yield(element);
    }
}

// Returns a generator that yields `f(x)`, for each `x` yielded by `g`.
template <typename T, typename F>
auto transform_generator(stack_generator_base<T>&& g, F f)
    [&g, f] [->] stack_generator<decltype(f(*g)), void> {
    while (g) {
        [<-] std::yield(f(*g));
        g.next();
    }
}

// Toy example: turn a nested vector into nested generators, and then
// flatten them.
//
// Caveat: this code contains a dangling-reference bug that we did not have
// time to fix before publication.
void f(const std::vector<std::vector<int>>& vectors) {
    stack_generator<int> gen = flatten(transform_generator(
        traverser(vectors),
        [] (const std::vector<int>& vec) { return traverser(vec); }));

    while (gen) {
        // Do stuff with *gen
        gen.next();
    }
}

```

And the underlying implementation:

```

// The internal state of a stack_generator<T,P,Coroutine>
template <typename T, typename P>
struct stack_generator_state {
    // Pointer to the currently yielded value
    T* value = nullptr;

```

```

    // The generator we have recursed into, if any
    stack_generator_base<T, P>* nested_generator = nullptr;
};

// Base class of all stack_generators that take P and yield T.
// Allows us to type-erase the coroutine.
template <typename T, typename P>
class stack_generator_base {
    stack_generator_state<T,P> state_;
public:
    stack_generator_base(stack_generator_base&&) = delete;
    stack_generator_base& operator=(stack_generator_base&&) = delete;

    void next(P& p) {
        (void) next_impl(p);
    }

    operator bool() const {
        return state_.value != nullptr;
    }
    T& operator*() { return *state_.value; }
    T* operator->() { return state_.value; }

    template<template<typename> Continuation>
    stack_generator_state<T,P> operator[<-](
        std::yield_handle<T>& handle, Continuation<P>&&) {
        return {handle.value, nullptr};
    }

    template <template<typename> Continuation, typename U, typename Q>
    stack_generator_state<T,P> operator[<-](
        stack_generator_base<U,Q>&& inner_generator, Continuation<void>&) {
        return {&*inner_generator, &inner_generator};
    }

    stack_generator_state<T,P> operator return() {
        return {};
    }

private:
    template <typename T2, typename P2>
    friend class stack_generator_base;

    // Resumes execution of the generator, and returns the new state
    virtual stack_generator_state<T,P> resume(P& p) = 0;

```



```

T* next_impl(P& p) {
    if (state_.nested_generator != nullptr) {
        T* value = state_.nested_generator.next_impl(p);
        if (value != nullptr) {
            return value;
        } else {
            state_.nested_generator = nullptr;
        }
    }
    assert(state_.nested_generator == nullptr);

    state_ = resume(p);
    return state_.value;
}
};

template <typename T, typename P, typename Coroutine>
class stack_generator : public stack_generator_base<T,P> {
public:
    using shared_state_type = stack_generator_base<T,P>;

    template <typename F>
    stack_generator(F& coroutine_callback)
        : coroutine_(coroutine_callback()) {}

private:
    Coroutine coroutine_;

    stack_generator_state<T,P> resume(P& p) override {
        return coroutine_([&]() -> P& { return p; });
    }
};

```

Parser Combinators

Coroutine syntax can be used to produce clear and elegant recursive-descent parsers. Here's an example usage, which parses a toy arithmetic expression language, computing the expression value on the fly:

```

// expr ::= expr addop factor | factor
// addop ::= '+' | '-'
// factor ::= number | '(' expr ')'

enum class Sign { Plus, Minus };

```

```

Parser<Sign> AddOp() {
    return FirstMatch<Sign>(
        Map(Consume("+"), Sign::Plus),
        Map(Consume("-"), Sign::Minus));
}

auto Number() [] [->] Parser<int> {
    string_view input = [<-] Peek();
    int value;
    auto result = std::from_chars(input.begin(), input.end(), value);
    if (result.ec != std::errc{}) {
        [<-] Fail(result.ec);
    }
    [<-] AdvanceInput(result.ptr - input.begin());
    return value;
}

Parser<int> Factor() {
    return FirstMatch<int>(
        Number(),
        [] [->] Parser<int> {
            [<-] Consume("(");
            int result = [<-] Expression();
            [<-] Consume(")");
            return result;
        });
}

auto Expression() [] [->] Parser<int> {
    int result = [<-] Factor();
    while (!([<-] Peek()).empty()) {
        Sign sign = [<-] AddOp();
        int next = [<-] Factor();
        switch (sign) {
            case Sign::Plus:
                result += next;
                break;
            case Sign::Minus:
                result -= next;
                break;
        }
    }
    return result;
}

```

Here's the underlying `Parser` type (note that for exposition purposes, we err on the side of simplicity rather than efficiency):

```
// An intermediate state of a Parser. `value` represents the value just parsed
// (nullopt if the parse failed), and `tail` represents the unparsed suffix of
// the input.
template <typename T>
struct ParseState {
    optional<T> value;
    string_view tail;
};

// A Parser<T> is essentially a function that takes a string_view, parses
// it to produce a T, and returns the result.
template <typename T>
class Parser {
    std::function<ParseState<T>(string_view)> parse_;

public:
    class shared_state_type {
        template <template<typename> Continuation, typename U>
        Parser<T> operator[<->](Parser<U>& parser, Continuation<U>& continuation) {
            return [&] (string_view input) -> ParseState<T> {
                ParseState<U> state = parser.parse(input);
                if (!state.value) {
                    return {nullopt, state.tail};
                }
                return continuation([&]() { return *state.value; }).parse(state.tail);
            }
        }

        Parser<T> operator return(const T& value) {
            return Return(value);
        }
    };

    template <typename F>
    Parser(F& coroutine_callback)
        requires std::is_coroutine_v<decltype(coroutine_callback())>
        : parse_([&] { return coroutine_callback(); }) (string_view input)
            -> ParseState<T> {
            return coroutine_callback();
        }) {}

    template <typename F>
    Parser(F& parse)
```

```

    requires std::is_convertible_v<
               decltype(parse(std::declval<string_view>)), ParseState<T>>
    : parse_(parse) {}

    ParseState<T> parse(string_view str) const {
        return parse_(str);
    }
};

```

And here are the reusable low-level parsing operations used in the above example:

```

// Parser which returns the given value, without consuming any input
template <typename T>
Parser<T> Return(const T& val) {
    return [=] (string_view input) { return val; };
}

// Parser which produces the entire remaining input as a string_view,
// without consuming any of it.
Parser<string_view> Peek() {
    return [] (string_view input) {
        return ParseState<string_view>{ input, input };
    };
}

// Parser which consumes n characters of input.
Parser<void> AdvanceInput(size_t n) {
    return [] (string_view input) {
        input.remove_prefix(n);
        return ParseState<void>{input};
    }
}

// Parser which fails without consuming any input.
Parser<void> Fail() {
    return [] (string_view input) {
        return {nullopt, input};
    }
}

namespace internal {
    ParseState<T> FirstMatchImpl(string_view input) {
        return {nullopt, input};
    }
}

```

```

template <typename T, typename... Parser_Ts>
ParseState<T> FirstMatchImpl(string_view input, Parser<T> parser,
                             Parser_Ts... parsers) {
    ParseState<T> result = parser.parse(input);
    if (result.value) {
        return result;
    } else {
        return FirstMatchImpl(input, std::move(parsers)...);
    }
}

}

} // namespace internal

// Parser which parses a T value from the first of parsers... which
// succeeds. parsers... must all be Parser<T> objects.
template <typename T, typename... Parser_Ts>
Parser<T> FirstMatch(Parser_Ts... parsers) {
    return [=] (string_view input) {
        return internal::FirstMatchImpl(input, std::move(parsers)...);
    }
}

// Parser which consumes the given string value from the beginning of
// the input, or fails if it is not present.
Parser<void> Consume(string expected) {
    return [=] (string_view input) {
        if (input.starts_with(expected)) {
            input.remove_prefix(expected.size());
            return ParseState<void>{input};
        } else {
            return Fail();
        }
    };
}

// Parser which parses the same inputs as `parser`, and produces `value`.
template <typename T>
Parser<T> Map(Parser<void> parser, T value) [=] [->] Parser<T> {
    [<-] parser;
    return std::move(value);
}

```