

P1072R6

basic_string::resize_and_overwrite

Published Proposal, 2020-12-14

This version:

<http://wg21.link/P1072R6>

Authors:

[Chris Kennelly](#) (Google)

[Mark Zeren](#) (VMware)

Audience:

LEWG, LWG, SG16

Project:

ISO/IEC JTC1/SC22/WG21 14882: Programming Language — C++

Abstract

Allow access to default-initialized elements of `basic_string`.

Table of Contents

- 1 Motivation**
- 2 Proposal**
- 3 Implementation**
- 4 Examples**
 - 4.1 Stamping a Pattern
 - 4.2 Interacting with C
- 5 Design Considerations**
 - 5.1 Method vs. Alternatives

5.2 Allocator Interaction

6 Proposal with Exceptions

6.1 Bikeshedding

7 Alternatives Considered

7.1 Tag Type

7.2 Non-Lambda Approach

7.3 Standalone Type: `storage_buffer`

7.4 Externally-Allocated Buffer Injection

8 Related Work

8.1 Google

8.2 MongoDB

8.3 VMware

8.4 Discussion on std-proposals

8.5 DynamicBuffer

8.6 P1020R1

8.7 Boost

8.8 Thrust

9 Wording

9.1 `[basic.string]`

9.2 `[container.requirements.general]`

10 History

10.1 R5 → R6

10.2 R4 → R5

10.3 R3 → R4

10.4 R2 → R3

10.5 R1 → R2

10.6 R0 → R1

11 Acknowledgments

References

Informative References

§ 1. Motivation

Performance sensitive code is impacted by the cost of initializing and manipulating strings. When streaming data into a `basic_string`, a programmer is faced with an unhappy choice:

- **Pay for extra initialization** — `resize`, which zero initializes, followed by copy.
- **Pay for extra copies** — Populate a temporary buffer, copy it to the string.
- **Pay for extra "bookkeeping"** — `reserve` followed by small appends, each of which checks capacity and null terminates.

The effects of these unhappy choices can all be measured at scale, yet C++'s hallmark is to leave no room between the language (or in this case the library) and another language.

LEWGI polled on this at the [\[SAN\]](#) Meeting:

We should promise more committee time to [\[P1072R1\]](#), knowing that our time is scarce and this will leave less time for other work?

Unanimous consent

LEWG at the [\[SAN\]](#) Meeting:

We want to solve this problem

SF	F	N	A	SA
17	5	0	0	0

Willing to solve this for string without solving it for vector

SF	F	N	A	SA
6	9	2	1	0

§ 2. Proposal

This proposal addresses the problem by adding `basic_string::resize_and_overwrite`:

```
template<typename Operation>
constexpr void resize_and_overwrite(size_type n, Operation op);
```

8. *Effects*: Let $o = \text{size}()$. Alters the value of `*this` as follows:

1. — If $n \leq \text{size}()$, erases the last $\text{size}() - n$ elements.
2. — If $n > \text{size}()$, appends $n - \text{size}()$ default-initialized elements.

Invokes `erase(begin() + op(data(), n), end())`.

9. *Remarks*: If `op` throws the behavior is undefined.

Let $m = \text{op}(\text{data}(), n)$.

$m \leq n$ otherwise the behavior is undefined.

If $m > o$, `op` shall replace ([**expr.ass**]) the values stored in the character array [`data()+o`, `data()+m`). Until replaced, the values may be indeterminate [**basic.indet**] (6.6.4). [*Note*: `*(data()+o)` may not be `CharT()`. - *end note*]

When `op` is called `*this` is in an unspecified state. `op` shall not bind or in any way access `*this`

`op` shall not allow its first argument to be accessible after it returns.

In order to enable `resize_and_overwrite`, this proposal makes it implementation-defined whether `basic_string` uses `allocator_traits::construct` and `allocator_traits::destroy` to construct and destroy the "char-like objects" that it controls. See [§ 5.2 Allocator Interaction](#) and [§ 9 Wording](#) below for more details.

§ 3. Implementation

`libc++` includes a private implementation of a prior version of this proposal (based on an earlier revision) and uses it to avoid a dynamic allocation in `std::filesystem` [LIBC++].

§ 4. Examples

§ 4.1. Stamping a Pattern

Consider writing a pattern several times into a string:

```
std::string GeneratePattern(const std::string& pattern, size_t count) {
    std::string ret;

    ret.reserve(pattern.size() * count);
    for (size_t i = 0; i < count; i++) {
        // SUB-OPTIMAL:
        // * Writes 'count' nulls
        // * Updates size and checks for potential resize 'count' times
        ret.append(pattern);
    }

    return ret;
}
```

Alternatively, we could adjust the output string's size to its final size, avoiding the bookkeeping in `append` at the cost of extra initialization:

```
std::string GeneratePattern(const std::string& pattern, size_t count) {
    std::string ret;

    const auto step = pattern.size();
    // SUB-OPTIMAL: We memset step*count bytes only to overwrite them.
    ret.resize(step * count);
    for (size_t i = 0; i < count; i++) {
        // GOOD: No bookkeeping
        memcpy(ret.data() + i * step, pattern.data(), step);
    }

    return ret;
}
```

With this proposal:

```
std::string GeneratePattern(const std::string& pattern, size_t count) {
    std::string ret;

    const auto step = pattern.size();
    // GOOD: No initialization
    ret.resize_and_overwrite(step * count, [&](char* buf, size_t n) {
        for (size_t i = 0; i < count; i++) {
            // GOOD: No bookkeeping
            memcpy(buf + i * step, pattern.data(), step);
        }
        return step * count;
    });

    return ret;
}
```

§ 4.2. Interacting with C

Consider wrapping a C API while working in terms of C++'s `basic_string` vocabulary. We *anticipate* over-allocating, as computation of the *length* of the data is done simultaneously with the computation of the *contents*.

```
extern "C" {
    int compress(void* out, size_t* out_size, const void* in, size_t in_size)
    size_t compressBound(size_t bound);
}

std::string CompressWrapper(std::string_view input) {
    std::string compressed;
    // Compute upper limit of compressed input.
    size_t size = compressBound(input.size());

    // SUB-OPTIMAL: Extra initialization
    compressed.resize(size);
    int ret = compress(&*compressed.begin(), &size, input.data(), input.size());
    if (ret != OK) {
        throw ...some error...
    }
}
```

```
// Set actual size.  
compressed.erase(size);  
return compressed;  
}
```

With this proposal:

```
extern "C" {  
    int compress(void* out, size_t* out_size, const void* in, size_t in_size)  
    size_t compressBound(size_t bound);  
}  
  
std::string CompressWrapper(std::string_view input) {  
    std::string compressed;  
    // Compute upper limit of compressed input.  
    size_t size = compressBound(input.size());  
  
    // GOOD: No initialization  
    compressed.resize_and_overwrite(size, [&](char* buf, size_t n) {  
        int ret = compress(buf, &size, input.data(), input.size());  
        if (ret != OK) {  
            throw ...some error...  
        }  
        return n;  
    })  
    return compressed;  
}
```

§ 5. Design Considerations

§ 5.1. Method vs. Alternatives

`resize_and_overwrite` exposes users to UB if they read indeterminate ([**basic.indet**] (6.6.4)) values in the string buffer passed to `op`. Despite this foot-gun, `resize_and_overwrite` is appealing because:

- Uninitialized buffers are not uncommon in C++. See for example `new T[20]` or the recently adopted `make_unique_for_overwrite`.
- Dynamic analyzers like Memory Sanitizer [\[MSan\]](#) and Valgrind [\[Valgrind\]](#) can catch uninitialized reads.

During the [\[SAN\]](#) Meeting LEWG expressed a preference for implementing this functionality as a new method on `basic_string` (as proposed in [\[P1072R0\]](#)) rather than a standalone "storage buffer" type (one option in [\[P1072R1\]](#)):

Method on string vs storage_buffer type:

Strong Method	Method	Neutral	Type	Strong Type
9	2	3	2	2

During the [\[KOA\]](#) Meeting, LEWG expressed a preference for not weakening `basic_string`'s invariants after `resize_and_overwrite` returns. Under this revision:

- A default-initialized buffer is *only* accessible to `op`.
- Under penalty of undefined behavior, `op` must assign (replace) the values in resulting buffer. The default initialized portion (`[data()+m, data()+n)`) above is erased and inaccessible to the program without further violation of class invariants.
- This proposal follows the "lambda" option discussed in [\[SAN\]](#).

Several other alternatives were discussed at [\[SAN\]](#), [\[KOA\]](#), and on the reflector. Please see [§ 7 Alternatives Considered](#) below for more details.

§ 5.2. Allocator Interaction

Unlocking the desired optimizations requires *some* change to `basic_string`'s interaction with allocators. This proposal does what we think is the simplest possible change: remove the requirement that `basic_string` call `allocator_traits::construct` or `allocator_traits::destroy`.

This restriction should be acceptable because `basic_string` is defined to only hold "non-array trivial standard-layout" types. [\[strings.general\]](#) p1:

This Clause describes components for manipulating sequences of any non-array trivial standard-layout (6.7) type. Such types are called *char-like types*, and objects of *char-like types* are called *char-like objects* or simply *characters*.

Removing calls to `construct` and `destroy` is compatible with `pmr::basic_string` as long as `uses_allocator_v<charT>` is `false`, which should be the case in practice.

Along the way, this proposal clarifies ambiguous language in `[string.require]` and `[container.requirements.general]` by:

- Stating explicitly that `basic_string` is allocator-aware.
- Introducing the term "allocator-aware" earlier in `[container.requirements.general]`.
- Not attempting to mention other "components" in `[container.requirements.general]`. The other allocator-aware containers (just `basic_string` and `regex::match_results`?) can independently state that they are "allocator-aware" in their own clauses.

libstdc++ and msvc allow strings of non-trivial type. That might force those libraries to continue to support `construct` and `destroy` as "extensions". On the other hand, libc++ disallows non-trivial `charTs`, so any such extensions are non-portable. See godbolt.org.

§ 6. Proposal with Exceptions

The version of this proposal seen by LEWG in [\[Cologne\]](#) throws `std::range_error` when `op` returns a value greater than `n`. The authors now recommend and LEWG confirmed via [\[LEWGNovTelecon\]](#) to use UB. For the `op` to have written greater than `n` bytes, it would have had to overrun the provided buffer (already UB).

§ 6.1. Bikeshedding

What do we want to call this method?

- `resize_and_overwrite` (as proposed). This is consistent with `make_unique_for_overwrite` [\[P1973R1\]](#). Unlike the `make_unique` version, this concludes with the buffer being overwritten (as far as `op` specifies), rather than requiring further action be taken.
- `resize_default_init` (as proposed). This is consistent with `make_unique_default_init` [\[P1020R1\]](#), adopted in San Diego. This name was removed from the working paper subsequently and renamed `make_unique_for_overwrite`.
- `resize_uninitialized` (as proposed in R0). "Uninitialized" is different from "default-initialized", which is what we want to specify. Also, `uninitialized` is already used elsewhere

in the Library to mean "not constructed at all", as in `uninitialized_copy`; so this would be an overloading of the term.

§ 7. Alternatives Considered

§ 7.1. Tag Type

At the [\[SAN\]](#) Meeting, LEWG showed support for a tag argument type:

Approval (vote for as many as you find acceptable)

13	Go back to <code>resize_uninitialized</code>
15	Do tag type (<code>default_initialize</code>) for c'tor / <code>resize()</code>
12	Continue with <code>storage_buffer</code> (as per R2 of this paper)
7	Crack open with a lambda
7	RAII separate type

For example:

```
std::string GeneratePattern(const std::string& pattern, size_t count) {
    const auto step = pattern.size();

    // GOOD: No initialization
    std::string ret(step * count, std::string::default_init);
    for (size_t i = 0; i < count; i++) {
        // GOOD: No bookkeeping
        memcpy(ret.data() + i * step, pattern.data(), step);
    }

    return ret;
}
```

Benefits:

- A constructor that takes a tag type simplifies some use cases, like the example above.
- Matches existing practice in Boost. See [§ 8.7 Boost](#).

Drawbacks:

- Feature creep / complexity — A tag type invites generalizing for all allocator-aware container types. We agree that this is desirable and even has implementation experience in Boost. However default initialization is not enough. There is also "implicit construction" (see [\[P1010R1\]](#), [\[P0593R2\]](#)) and "relocation" (see [\[P1144R0\]](#), [\[P1029R1\]](#)). Neither of these features are yet in the language. It is too early to generalize. Note that the second poll quoted in [§ 1 Motivation](#) shows support for solving this problem for `[basic_]string` but not `vector`.
- In reflector discussion of `make_unique_default_init [Zero]`, there was a preference for avoiding tag types. The standard library has `copy_backward`, not `copy` with a tag, and `count_if`, rather than `count` with a predicate.

Conclusion:

LEWG should explore tags for allocator-aware containers, but that work should not block near-term enablement of efficient `[std::]string` builders.

§ 7.2. Non-Lambda Approach

In [\[P1072R0\]](#) and [\[P1072R3\]](#) of this proposal, the authors considered a method `resize_default_init` / `resize_uninitialized` which left a default-initialized buffer accessible to users of the `basic_string` instance after the method returned. This method was rejected in [\[KOA\]](#), due to the weakened class invariants.

For illustration:

```
extern "C" {
    int compress(void* out, size_t* out_size, const void* in, size_t in_size)
    size_t compressBound(size_t bound);
}

std::string CompressWrapper(std::string_view input) {
    std::string compressed;
    // Compute upper limit of compressed input.
    size_t size = compressBound(input.size());

    // GOOD: No initialization
    compressed.resize_default_init(size);
    int ret = compress(&*compressed.begin(), &size, input.data(), input.size());
    if (ret != OK) {
        throw ...some error...
    }
}
```

```

    // Suppose size is the value of size before the call to compress and si
    // is the value of size after the call to compress.
    //
    // If size' < size, then:
    //   std::cout << compressed[size' + 1]
    // ...triggers a read of uninitialized data.

    // Set actual size.
    compressed.erase(size);
    return compressed;
}

```

§ 7.3. Standalone Type: `storage_buffer`

In [\[P1072R1\]](#), we considered `storage_buffer`, a standalone type providing a `prepare` method (similar to the `resize_uninitialized` method proposed here) and `commit` (to promise, under penalty of UB, that the buffer had been initialized).

At the [\[SAN\]](#) Meeting, this approach received support from LEWGI in light of the [\[post-Rapperswil\]](#) email review indicating support for a distinct type. This approach was rejected by the larger LEWG room in San Diego Meeting Diego.

The proposed type would be move-only.

```

std::string GeneratePattern(const std::string& pattern, size_t count) {
    std::storage_buffer<char> tmp;

    const auto step = pattern.size();
    // GOOD: No initialization
    tmp.prepare(step * count + 1);
    for (size_t i = 0; i < count; i++) {
        // GOOD: No bookkeeping
        memcpy(tmp.data() + i * step, pattern.data(), step);
    }

    tmp.commit(step * count);
    return std::string(std::move(tmp));
}

```

For purposes of the container API, `size()` corresponds to the *committed* portion of the buffer. This leads to more consistency when working with (and explicitly copying to) other containers via iterators, for example:

```
std::storage_buffer<char> buf;
buf.prepare(100);
*fill in data*
buf.commit(50);

std::string a(buf.begin(), buf.end());
std::string b(std::move(buf));

assert(a == b);
```

Benefits:

- By being move-only, we would not have the risk of copying types with trap representations (thereby triggering UB).
- Uninitialized data is only accessible from the `storage_buffer` type. For an API working with `basic_string`, no invariants are weakened. Crossing an API boundary with a `storage_buffer` is much more obvious than a "`basic_string` with possibly uninitialized data." Uninitialized bytes (the promise made by `commit`) never escape into the valid range of `basic_string`.

Drawbacks:

- `storage_buffer` requires introducing an extra type to the standard library, even though its novel functionality (from `string` and `vector`) is limited to the initialization abilities.
- `basic_string` is often implemented with a short-string optimization (SSO) and an extra type would need to implement that (likely by additional checks when moving to/from the `storage_buffer`) that are often unneeded.

§ 7.4. Externally-Allocated Buffer Injection

In [P1072R1], we considered that `basic_string` could "adopt" an externally `allocator::allocate`'d buffer. At the [SAN] Meeting, we concluded that this was:

- **Not critical**
- **Not constrained in the future**

- **Overly constraining to implementers.** Allowing users to provide their own buffers runs into the "offset problem". Consider an implementation that stores its `size` and `capacity` inline with its data, so `sizeof(container) == sizeof(void*)`.

```
class container {  
    struct Rep {  
        size_t size;  
        size_t capacity;  
    };  
  
    Rep* rep_;  
};
```

If using a `Rep`-style implementation, the mismatch in offsets requires an $O(N)$ move to shift the contents into place and trigger a possible reallocation.

- **Introducing new pitfalls.** It would be easy to mix `new[]` and `allocator::allocate` inadvertently.

§ 8. Related Work

§ 8.1. Google

Google has a local extension to `basic_string` called `resize_uninitialized` which is wrapped as `STLStringResizeUninitialized`.

- [\[Abseil\]](#) uses this to avoid bookkeeping overheads in `StrAppend` and `StrCat`.
- [\[Snappy\]](#)
 - In [decompression](#), the final size of the output buffer is known before the contents are ready.
 - During [compression](#), an upperbound on the final compressed size is known, allowing data to be efficiently added to the output buffer (eliding `append`'s checks) and the string to be shrunk to its final, correct size.
- [\[Protobuf\]](#) avoids extraneous copies or initialization when the size is known before data is available (especially during parsing or serialization).

§ 8.2. MongoDB

MongoDB has a string builder that could have been implemented in terms of `basic_string` as a return value. However, as explained by Mathias Stearn, zero initialization was measured and was too costly. Instead a custom string builder type is used:

E.g.: <https://github.com/mongodb/mongo/blob/master/src/mongo/db/fts/unicode/string.h>

```
/**
 * Strips diacritics and case-folds the utf8 input string, as needed to sup
 * options.
 *
 * The options field specifies what operations to *skip*, so kCaseSensitive
 * means to skip case folding and kDiacriticSensitive means to skip diacrit
 * striping. If both flags are specified, the input utf8 StringData is retu
 * directly without any processing or copying.
 *
 * If processing is performed, the returned StringData will be placed in
 * buffer. buffer's contents (if any) will be replaced. Since we may retur
 * the input unmodified the returned StringData's lifetime is the shorter o
 * the input utf8 and the next modification to buffer. The input utf8 must
 * point into buffer.
 */
static StringData caseFoldAndStripDiacritics(StackBufBuilder* buffer,
                                             StringData utf8,
                                             SubstrMatchOptions options,
                                             CaseFoldMode mode);
```

(Comments re-wrapped.)

§ 8.3. VMware

VMware has an internal string builder implementation that avoids `std::string` due, in part, to `reserve`'s zero-writing behavior. This is similar in spirit to the MongoDB example above.

§ 8.4. Discussion on std-proposals

This topic was discussed in 2013 on std-proposals in a thread titled "Add `basic_string::resize_uninitialized` (or a similar mechanism)":

§ 8.5. DynamicBuffer

The [\[N4734\]](#) (the Networking TS) has *dynamic buffer* types.

§ 8.6. P1020R1

See also [\[P1020R1\]](#) "Smart pointer creation functions for default initialization". Adopted in San Diego.

§ 8.7. Boost

Boost provides a related optimization for vector-like containers, introduced in [\[SVN r85964\]](#) by Ion Gaztañaga.

E.g.: [boost/container/vector.hpp](#):

```
//! <b>Effects</b>: Constructs a vector that will use a copy of allocator a
//!    and inserts n default initialized values.
//!
//! <b>Throws</b>: If allocator_type's allocation
//!    throws or T's default initialization throws.
//!
//! <b>Complexity</b>: Linear to n.
//!
//! <b>Note</b>: Non-standard extension
vector(size_type n, default_init_t);
vector(size_type n, default_init_t, const allocator_type &a)
...
void resize(size_type new_size, default_init_t);
...
```

These optimizations are also supported in Boost Container's [small_vector](#), [static_vector](#), [deque](#), [stable_vector](#), and [string](#).

§ 8.8. Thrust

The Thrust library has "a RAII-type `thrust::detail::temporary_array` which has a vector-like interface and a constructor with a tag parameter that indicates its elements should not be initialized." - [Bryce Adelstein Leibach].

E.g. [thrust/thrust/detail/temporary_array.inl](#):

```
template<typename T, typename TemporaryArray, typename Size>
__host__ __device__
typename thrust::detail::disable_if<
    avoid_initialization<T>::value
>::type
    construct_values(TemporaryArray &a,
                    Size n)
{
    a.default_construct_n(a.begin(), n);
} // end construct_values()
```

§ 9. Wording

Wording is relative to [\[N4830\]](#).

Motivation for some of these edits can be found in [§ 5.2 Allocator Interaction](#).

§ 9.1. [basic.string]

In [basic.string] [21.3.2], in the synopsis, add `resize_and_overwrite`:

```
...
// 21.3.2.4, capacity
constexpr size_type size() const noexcept;
constexpr size_type length() const noexcept;
constexpr size_type max_size() const noexcept;
constexpr void resize(size_type n, charT c);
constexpr void resize(size_type n);
template<typename Operation>
void resize_and_overwrite(size_type n, Operation op);
constexpr size_type capacity() const noexcept;
constexpr void reserve(size_type res_arg);
constexpr void shrink_to_fit();
constexpr void clear() noexcept;
[[nodiscard]] constexpr bool empty() const noexcept;
```

In [string.require] [21.3.2.1] clarify that `basic_string` is an allocator-aware container, and then add an exception for `construct` and `destroy`.

3. In every specialization `basic_string<charT, traits, Allocator>`, the type `allocator_traits<Allocator>::value_type` shall name the same type as `charT`.
~~Every object of type `basic_string<charT, traits, Allocator>` uses an object of type `Allocator` to allocate and free storage for the contained `charT` objects as needed. The `Allocator` object used is obtained as described in [container.requirements.general]. In every specialization `basic_string<charT, traits, Allocator>`, and the type `traits` shall satisfy the character traits requirements ([char.traits]). [Note: The program is ill-formed if `traits::char_type` is not the same type as `charT`. — end note]~~
4. `basic_string` is an allocator-aware container as described in [container.requirements.general], except that it is implementation-defined whether `allocator_traits::construct` and `allocator_traits::destroy` are used to construct and destroy the contained `charT` objects.
5. References, pointers, and iterators referring to the elements of a `basic_string` sequence may be invalidated by the following uses of that `basic_string` object:
...

In `[string.capacity]` [21.3.2.4]:

```
constexpr void resize(size_type n, charT c);
```

6. *Effects*: Alters the value of `*this` as follows:

1. — If `n <= size()`, erases the last `size() - n` elements.
2. — If `n > size()`, appends `n - size()` copies of `c`.

```
constexpr void resize(size_type n);
```

7. *Effects*: Equivalent to `resize(n, charT())`.

```
template<typename Operation>
```

```
void resize_and_overwrite(size_type n, Operation op);
```

8. *Effects*: Let `o = size()`. Alters the value of `*this` as follows:

1. — If `n <= size()`, erases the last `size() - n` elements.
2. — If `n > size()`, appends `n - size()` default-initialized elements.

Invokes `erase(begin() + op(data(), n), end())`.

9. *Remarks*: If `op` throws the behavior is undefined.

Let `m = op(data(), n)`.

`m <= n` otherwise the behavior is undefined.

If `m > o`, `op` shall replace ([**expr.ass**]) the values stored in the character array [`data() + o`, `data() + m`). Until replaced, the values may be indeterminate [**basic.indet**] (6.6.4). [*Note*: `*(data() + o)` may not be `CharT()`. - end note]

When `op` is called `*this` is in an unspecified state. `op` shall not bind or in any way access `*this`

`op` shall not allow its first argument to be accessible after it returns.

```
constexpr size_type capacity() const noexcept;
```

...

§ 9.2. [container.requirements.general]

In [container.requirements.general] [22.2.1] clarify the ambiguous "components affected by this subclause" terminology in p3. Just say "allocator-aware containers".

3. ~~All of the containers defined in this Clause except array are allocator-aware containers. For the components affected by this subclause that declare an allocator_type Objects objects stored in these components~~ allocator-aware containers, unless otherwise specified, shall be constructed using the function
allocator_traits<allocator_type>::rebind_traits<U>::construct and
destroyed using the function
allocator_traits<allocator_type>::rebind_traits<U>::destroy (20.10.9.2),
where U is either allocator_type::value_type or an internal type used by the container.
...

We can then simplify p15:

15. ~~All of the containers defined in this Clause and in 21.3.2 except array meet the additional requirements of an allocator-aware container, as~~ Allocator-aware containers meet the additional requirements described in Table 75.

§ 10. History

§ 10.1. R5 → R6

LEWG reviewed R5 via [\[LEWGNovTelecon\]](#).

POLL: `resize_default_init` should be able to work in `constexpr` context.

No objection to unanimous consent. Attendance: 23

POLL: We prefer UB over throwing an exception if the operation reports it wrote more bytes than it has access to.

No objection to unanimous consent. Attendance: 21

POLL: We want to rename `resize_default_init` for consistency, in a mailing list review.

SF	F	N	A	SA
3	12	0	1	0

Modifications since R5:

- Indicated poll outcomes for UB when `op` reports an invalid value.
- Renamed to `resize_and_overwrite`. This is based on naming discussions for `std::make_unique_default_init` being ultimately renamed to `std::make_unique_for_overwrite`, but differs slightly because the memory accessible has been overwritten. Overwriting the buffer is an integral part of the operation performed.
- Applied `constexpr`

§ 10.2. R4 → R5

A draft of this revision was presented to LEWG at the [\[Cologne\]](#) meeting.

Forward to LWG for C++23

SF	F	N	A	SA
10	5	1	0	0

- Rebased on [\[N4830\]](#).
- Proposed UB instead of `throwing std::range_error`.

§ 10.3. R3 → R4

- Applied feedback from the [\[KOA\]](#) meeting.
- Moved to using a callback routine, rather than leaving invariants weakened from the time `resize_default_init` returned until values in the buffer were assigned.

§ 10.4. R2 → R3

- Applied Jonathan Wakely's editorial comments on [§ 9 Wording](#).
- Rebased on [\[N4791\]](#).

- Editorial changes to [§ 1 Motivation](#) and [§ 7.1 Tag Type](#) to support our design choice.
- Added the reference to [\[LIBC++\]](#) in [§ 3 Implementation](#).

§ 10.5. R1 → R2

Applied feedback from [\[SAN\]](#) Meeting reviews.

- Reverted design to "option A" proposed in [\[P1072R0\]](#).
- Switched from `resize_uninitialized` to `resize_default_init`.
- Added discussion of alternatives considered.
- Specified allocator interaction.
- Added wording.

§ 10.6. R0 → R1

Applied feedback from LEWG [\[post-Rapperswil\]](#) Email Review:

- Shifted to a new vocabulary types: `storage_buffer` / `storage_node`
 - Added presentation of `storage_buffer` as a new container type
 - Added presentation of `storage_node` as a node-like type
- Added discussion of design and implementation considerations.
- Most of [\[P1010R1\]](#) Was merged into this paper.

§ 11. Acknowledgments

Thanks go to **Arthur O'Dwyer** for help with wording and proof reading, to **Jonathan Wakely** for hunting down the language that makes `basic_string` allocator-aware, and to **Glen Fernandes**, **Corentin Jabot**, **Billy O'Neal**, and **Mathias Stearn** for design discussion. Special thanks to **Eric Fiselier** for providing the implementation.

§ References

§ Informative References

[Abseil]

Abseil. 2018-09-22. URL: <https://github.com/abseil/abseil-cpp>

[Cologne]

Cologne Meeting Minutes. 2019-07-19. URL:
<http://wiki.edg.com/bin/view/Wg21cologne2019/P1072>

[KOA]

Kona Meeting Minutes. 2019-02-19. URL: <http://wiki.edg.com/bin/view/Wg21kona2019/P1072>

[LEWGNovTelecon]

LEWG November 2020 Minutes. 2020-11-16. URL:
<https://wiki.edg.com/bin/view/Wg21summer2020/P1072>

[LIBC++]

Add basic_string::__resize_default_init (from P1072). 2018-11-26. URL:
<http://llvm.org/viewvc/llvm-project?view=revision&revision=347589>

[MSan]

Memory Sanitizer. URL: <https://clang.llvm.org/docs/MemorySanitizer.html>

[N4734]

Jonathan Wakely. Working Draft, C ++ Extensions for Networking. 4 April 2018. URL:
<https://wg21.link/n4734>

[N4791]

Richard Smith. Working Draft, Standard for Programming Language C++. 26 November 2018.
URL: <https://wg21.link/n4791>

[N4830]

Working Draft, Standard for Programming Language C++. 2019-08-15. URL: <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2019/n4830.pdf>

[P0593R2]

Richard Smith. Implicit creation of objects for low-level object manipulation. 11 February 2018.
URL: <https://wg21.link/p0593r2>

[P1010R1]

Mark Zeren, Chris Kennelly. Container support for implicit lifetime types. 8 October 2018. URL:
<https://wg21.link/p1010r1>

[P1020R1]

Glen Joseph Fernandes, Peter Dimov. Smart pointer creation with default initialization. 6 November 2018. URL: <https://wg21.link/p1020r1>

[P1029R1]

Niall Douglas. [\[\[move_relocates\]\]](https://wg21.link/p1029r1). 7 August 2018. URL: <https://wg21.link/p1029r1>

[P1072R0]

Chris Kennelly, Mark Zeren. [Default Initialization for basic_string](https://wg21.link/p1072r0). 4 May 2018. URL: <https://wg21.link/p1072r0>

[P1072R1]

Chris Kennelly, Mark Zeren. [Optimized Initialization for basic_string and vector](https://wg21.link/p1072r1). 7 October 2018. URL: <https://wg21.link/p1072r1>

[P1072R3]

Chris Kennelly, Mark Zeren. [basic_string::resize_default_init](https://wg21.link/p1072r3). 21 January 2019. URL: <https://wg21.link/p1072r3>

[P1144R0]

Arthur O'Dwyer, Mingxin Wang. [Object relocation in terms of move plus destroy](https://wg21.link/p1144r0). 4 October 2018. URL: <https://wg21.link/p1144r0>

[P1973R1]

[Rename _default_init Functions](http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2020/p1973r1.pdf). 2020-02-12. URL: <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2020/p1973r1.pdf>

[post-Rapperswil]

[LEWG Weekly - P1072](http://lists.isocpp.org/lib-ext/2018/08/8299.php). 2018-08-26. URL: <http://lists.isocpp.org/lib-ext/2018/08/8299.php>

[Protobuf]

[Protocol Buffers](https://github.com/protocolbuffers/protobuf). 2018-09-22. URL: <https://github.com/protocolbuffers/protobuf>

[SAN]

[San Diego Meeting Minutes](http://wiki.edg.com/bin/view/Wg21sandiego2018/P1072). 2018-11-07. URL: <http://wiki.edg.com/bin/view/Wg21sandiego2018/P1072>

[Snappy]

[Snappy](https://github.com/google/snappy). 2018-09-21. URL: <https://github.com/google/snappy>

[Valgrind]

[Valgrind](http://www.valgrind.org). URL: <http://www.valgrind.org>

[Zero]

[Listening to our customers: Zero-initialization issue](http://lists.isocpp.org/lib-ext/2018/04/6712.php). 2018-04-13. URL: <http://lists.isocpp.org/lib-ext/2018/04/6712.php>