

P1609R2

C++ Should Support Just-in-Time Compilation

Published Proposal, 2019-10-07

This version:

<http://wg21.link/p1609>

Author:

[Hal Finkel](#) (Argonne National Laboratory)

Audience:

SG17

Project:

ISO/IEC JTC1/SC22/WG21 14882: Programming Language — C++

Abstract

C++ has been specifically designed to support the creation of performance-sensitive systems. C++ templates have become a critical piece of this support, allowing the creation of algorithmic specializations to which the application can dispatch based on runtime conditions. However, many applications are pushing this capability to its limit, generating so many specializations using template instantiation that compile times are interfering with programmer productivity. In other cases, programmers simply must settle for suboptimal performance because the ahead-of-time generation of specializations with better performance is not practical given the size of the relevant parameter space. To address these challenges, this paper proposes to add to C++ the capability to support programmer-directed just-in-time (JIT) compilation.

Table of Contents

- 1 Introduction
- 2 Changelog

3	Updated Proposal
4	Implementation
5	Proposal
5.1	Attribute
5.1.1	Example
5.1.2	Example
5.2	Restrictions
6	A Design Question
7	On State
8	ABI and Error Handling
9	A Benchmark
9.1	Compile Time
9.2	Runtime Performance
9.3	Overheads
10	Other C++ JIT Approaches
11	Acknowledgments

§ 1. Introduction

C++ is the go-to language for programming performance-sensitive systems. In part, this is due to C++'s design philosophy of "leav[ing] no room for a lower-level language below C++ (except assembler)" (see Appendix A in p0559r0 and Bjarne's 2007 HoPL paper). As C++ programmers, however, we are faced with a fundamental dilemma:

1. A compiler can often generate more-efficient code for algorithms in which some important parameters have (compile-time) known values compared to code for the same algorithms for which those same parameters are known only during program execution. These parameters are sometimes values (e.g., the number of rows in the matrix is three) and sometimes types. Types here include both types being operated upon and types representing behavioral composition.

2. In some cases, we can template the algorithms on these relevant parameters and instantiate the templates for a large set of the parameters likely to be relevant during program execution. The program can then dispatch to a relevant instantiation during program execution, perhaps falling back to a generic implementation. However, this can have a large compile-time cost. In fact, practical limits on compile time often limit the size of the parameter space which can be covered by instantiated templates.
3. In other cases, the relevant parameter space is so large that a relevant set of instantiations cannot be feasibly selected. In such cases, we're forced to settle for a generic implementation, even if the actual sets of parameters that occur during any particular program execution is not large.

If we had the ability to instantiate templates during program execution, using parameters not known until program execution, we could provide the higher performance of fixed-parameter implementations while simultaneously providing improved compile times. This paper explores the following question: Can we naturally integrate this kind of just-in-time (JIT) compilation capability, including the ability to construct novel template instantiations during program execution, into C++?

Clearly, this kind of JIT compilation capability might not be something that all C++ implementation can provide. In some cases, this is because of technical limitation of the underlying platforms (e.g., in certain kinds of embedded systems). In other cases, this is because the additional risk factors introduced by dynamic compilation are unacceptable in certain environments. Regardless, it seems natural that capabilities in this space would fall into a conditionally-supported category.

§ 2. Changelog

- r2: Updated based on feedback from the Cologne meeting.
- r1: Updated to discuss overheads in more detail, ABI considerations, and other issues raised on the reflector.

§ 3. Updated Proposal

After the evening-session discussion in Cologne, and review by EWGI, what follows is an updated proposal for the interface for this functionality:

First, we need some way to represent templates to enable runtime instantiation without referring to types using strings. We need, essentially, some runtime representation of a dependent template id, where the dependence can now be a runtime dependence.

Here, I'll propose to extend `typeid` so that it can be applied to template ids. Thus, we can write:

```
template <int i>
struct x {};

...

auto a = typeid(x);
```

We'll also add a new standard-library function:

```
template <typename... Args>
std::type_info std::compose_type_info(const std::type_info &base_type, const
```

where each argument must be a const type_info reference, where the template takes a type argument, and for a non-type template argument, a const reference to the corresponding type.

It is often true in the context of JIT that we want some template arguments to be provided at compile time and others to be provided during program execution. We can enable this by using template aliases. For example, if we have:

```
template <typename T, int A>
struct x {};
```

then we can define template aliases such as:

```
template <int d>
using i = x<int, d>;
```

and then apply typeid to that template alias.

There's some potential relationship here to future reflection features, and that's not yet addressed here.

The interface to instantiation is also changed to be an expression using the keyword:

```
dynamic_function_instantiation
```

which takes the function template as its first argument followed by arguments for each function-template argument. For type arguments, a const type_info reference must be provided.

The result of this expression is an implementation-defined type that has a `()` operator which allows invoking the newly-instantiated function, a `bool` conversion operator (which returns `true` if calling the `()` operator is permitted), and two additional member functions: `errors()` and `warnings()`. These functions return `const` references to implementation-defined types which are valid as range expressions. Each element in the range is an implementation-defined type which has the members: `message()`, which returns a `std::string`, `note()`, which also returns a `std::string`, and `loc()` which returns a `const source_location` reference.

Thus, we can do the following:

```

template <typename T, int I>
void foo(int a) { }

...

template <int J>
struct A { };

...

auto A_tid = typeid(A);

...

auto A5 = std::compose_type_info(A_tid, 5);

...

int i = ..., j = ...;
auto foo_TI = dynamic_function_instantiation(foo, A5, i);

for (auto &W : foo_TI.warnings()) {
    std::cerr << "warning: " << W.message() << " @ " << W.loc() << "\n";
}

if (foo_TI) {
    foo_TI(j);
} else {
    std::cerr << "compilation failed!\n";

    for (auto &W : foo_TI.errors()) {
        std::cerr << "error: " << W.message() << " @ " << W.loc() << "\n";
    }
}

```

§ 4. Implementation

The proposal below has been implemented, and that implementation is available from the repository: [github:hfinkel/llvm-project-cxxjit](https://github.com/hfinkel/llvm-project-cxxjit). There is some additional documentation on the wiki associated with

the repository: [github:hfinkel/llvm-project-cxxjit/wiki](https://github.com/hfinkel/llvm-project-cxxjit/wiki). The data presented below was generated using this implementation.

§ 5. Proposal

This proposal is designed to be minimal, and proposes a JIT capability with an interface limited to function templates. JIT compilation engines, as a practical matter, are generally designed to provide new functions (e.g., to provide a function pointer to some newly-compiled function), making this a natural fit to the underlying technology. Moreover, this allows us to elide questions regarding partial specialization.

§ 5.1. Attribute

Function templates can be tagged for just-in-time compilation by using the attribute:

```
[[jit]]
```

[*Note:*

In the aforementioned implementation of this proposal, the attribute is named `[[clang::jit]]`.

-- end note]

The attributed function template provides for additional features and restrictions. Features:

1. Instantiations of this function template will not be constructed at compile time, but rather calling a specialization of the template, or taking the address of a specialization of the template, will trigger the instantiation and compilation of the template during program execution. Note that this property is non-normative (i.e., not observable within the abstract machine) in some cases (e.g., when all template parameters are known at compile time and there are no observable order-of-instantiation dependences (see the discussion of state below)).
2. Non-constant expressions may be provided for the non-type template parameters, and these values will be used during program execution to construct the type of the requested instantiation. For const array references, the data in the array will be treated as an initializer of a constexpr variable.

§ 5.1.1. Example

```
#include <iostream>
#include <cstdlib>

template <int x>
[[jit]] void run() {
    std::cout << "I was compiled at runtime, x = " << x << "\n";
}

int main(int argc, char *argv[]) {
    int a = std::atoi(argv[1]);
    run<a>();
}
```

3. Type arguments to the template can be provided as strings. If the argument is implicitly convertible to a `const char *`, then that conversion is performed, and the result is used to identify the requested type. Otherwise, if an object is provided, and that object has a member function named `c_str()`, and the result of that function can be converted to a `const char *`, then the call and conversion (if necessary) are performed in order to get a string used to identify the type. The string is parsed and analyzed to identify the type in the declaration context of the parent to the function triggering the instantiation. Whether types defined after the point in the source code that triggers the instantiation are available is not specified.

§ 5.1.2. Example

```
#include <iostream>

struct F {
    int i;
    double d;
};

template <typename T, int S>
struct G {
    T arr[S];
};

template <typename T>
[[jit]] void run() {
    std::cout << "I was compiled at runtime, sizeof(T) = " << sizeof(T) << "\n";
}

int main(int argc, char *argv[]) {
    std::string t(argv[1]);
    run<t>();
}
```

§ 5.2. Restrictions

1. Because the body of the template is not instantiated at compile time, `decltype(auto)` and any other type-deduction mechanisms depending on the body of the function are not available.
2. Because the template specializations are not compiled until during program execution, they're not available at compile time for use as non-type template arguments, etc.

Explicit specializations of a JIT function template are not just-in-time compiled.

Note: A JIT template with a pointer/reference non-type template parameter which is provided with a runtime pointer value will generate a different instantiation for each pointer value. If the pointer provided points to a global object, no attempt is made to map that pointer value back to the name of the global object when constructing the new type.

Note: In general, pointer/reference-type non-type template arguments are not permitted to point to subobjects. This restriction still applies formally to the templates instantiated at runtime using runtime-provided pointer values. This has important optimization benefits: pointers that can be traced back to distinct underlying objects are known not to alias, and these template parameters appear to the optimizer to have this unique-object property.

§ 6. A Design Question

The proposal above reflects what has been implemented, but the use of a attribute might be suboptimal from a language-design perspective. Some points against this attribute approach:

1. The template is not otherwise special, it's the point of instantiation that's special (and that might not compile at all without the JIT support).
2. The uses of the template look vaguely normal, and so places where you might invoke the JIT compiler will be difficult to spot during code review.
3. The current mechanism provides no place to get out an error or provide a fall-back execution path - except that having the runtime throw an exception might work.
4. Should the mechanism be synchronous (as in this proposal), or asynchronous (e.g., something that returns a promise).

Thus, it might be better to use, e.g., a keyword near the point of instantiation. Perhaps something like this:

```
jit_this_template foo<argc>();
```

or maybe:

```
foo jit_this_template <argc>();
```

where `jit_this_template` is a new keyword.

The disadvantage of tying the JIT use to the point of instantiation instead of to the template itself, is that we need to decide that happens if the same instantiation is created both at compile time in the usual manner and also requested to be delayed until program execution. This might be okay or it might be some kind of ODR violation.

A library-like syntax is also possible (suggestion by Connor Waters):

```
std::jit<foo, a, _, b, c, _, _>(x, y, z)
```

§ 7. On State

Because the instantiation of some templates can affect the instantiation of other templates (e.g., because friend injection can affect later overload resolution), as currently proposed, the implementation of the JIT-compilation functionality cannot be "stateless." While not clearly desirable in general, although it enables some interesting counting tricks (e.g., [Non-Constant Constant-Expressions in C++](#)), this might be particularly undesirable in the JIT-compilation context: It likely makes it more difficult for the JIT-compilation engine to discard its in-memory state even in cases where that might be otherwise advantageous (e.g., to reduce its memory footprint if resource pressure becomes an issue). It is already a non-trivial question as to whether, and under what conditions, the runtime system might be able to determine that a particular instantiation might become unused, thus enabling the freeing of associated resources.

As the JIT-compilation process operates on semantically-processed inputs, and so there is no preprocessor state relevant to the JIT-compilation process. This imposes some restriction on the use of the feature: the JIT-compiled code cannot use differences in the preprocessor state (e.g., whether `__AVX2__` is defined) to customize its behavior (and/or implementation strategy) based on the capabilities of the system on which the program is running. This restriction is likely undesirable, although it is not clear how best to lift the restriction.

§ 8. ABI and Error Handling

It is important that the JIT-compilation facility provide robust error handling capabilities. In this light, although observable within the abstract machine, a production-quality implementation might choose to segregate the JIT-compilation engine into a separate process. JIT compilation can fail because, in addition to compiler bugs, the compilation engine might lack some necessary resources, or the code might otherwise trigger some implementation limit. In addition, compilation might fail because an invalid type string was provided or the provided type or value triggered some semantic error (including triggering a `static_assert`). It might be correct to say that the latter set of conditions are simply undefined behavior, but regardless, providing some ability for the application to handle such errors at the point of instantiation seems useful. No specific mechanism is proposed in this paper, but feedback is certainly desired.

This is related to questions regarding the ABI: If an application is compiled with all of the necessary metadata embedded with in to compile the JIT-attributed templates, does that metadata format and the

underlying interface to the JIT-compilation engine that uses it become part of the ABI that the system must support in order to run the application? The answer to this question seems likely to be yes. In practice, this represents another reason why compilation might fail: ABI incompatibility (e.g., running an application on a system with a JIT-compilation engine incompatible with the one required by the application - one that is too old, for example). At the cost of some (perhaps significant) binary-size overhead, this issue can be avoided by linking the necessary parts of the compiler into the application itself.

Also related to the question of the ABI is that of inter-compiler compatibility: If different translation units that comprise an application's code base are compiled with different compilers, and some of these different translation units use JIT-compilation features, does this necessarily cause problems? Implementation experience is certainly needed here, but it seems likely that an appropriate system ABI could be defined that would allow this kind of combination to work. In the case where the use of different compilers implied the corresponding use of different JIT engines, and we want to ensure properties such as global uniqueness of static local variables in inline functions generated by just-in-time compilation, the system would need provide some way for each JIT engine to register with the others which symbols (functions, global variables, etc.) it had generated. It is unknown what else, if anything, might be required. This seems conceptually similar to the notification functionality that already exists in order to notify a debugger about the newly-generated (JIT-compiled) symbols.

§ 9. A Benchmark

As a benchmark to illustrate the feature, we'll adapt a benchmark from the [Eigen](https://github.com/eigenteam/eigen-git-mirror/blob/master/bench/benchmark.cpp) library. Specifically, this one: <https://github.com/eigenteam/eigen-git-mirror/blob/master/bench/benchmark.cpp>

We want to look at two aspects: Compile time and runtime performance. Eigen provides a matrix type which can either have compile-time-specific or runtime-specified sizes (i.e., the number of rows and columns).

```
#include <iostream>
#include <string>
#include <chrono>
#include <cstdlib>

#include <Eigen/Core>

using namespace std;
using namespace Eigen;
```

If we wish to support a variant of this benchmark supporting float, double, and long double, and supporting any size at runtime, we can adapt the code as:

```
template <typename T>
void test_aot(int size, int repeat) {
    Matrix<T,Dynamic,Dynamic> I = Matrix<T,Dynamic,Dynamic>::Ones(size, size);
    Matrix<T,Dynamic,Dynamic> m(size, size);
    for(int i = 0; i < size; i++)
        for(int j = 0; j < size; j++) {
            m(i,j) = (i+size*j);
        }

    auto start = chrono::system_clock::now();

    for (int r = 0; r < repeat; ++r) {
        m = Matrix<T,Dynamic,Dynamic>::Ones(size, size) + T(0.00005) * (m + (m^3));
    }

    auto end = chrono::system_clock::now();
    cout << "AoT: " << chrono::duration<double>(end - start).count() << " s\r\n";
}

void test_aot(std::string &type, int size, int repeat) {
    if (type == "float")
        test_aot<float>(size, repeat);
    else if (type == "double")
        test_aot<double>(size, repeat);
    else if (type == "long double")
        test_aot<long double>(size, repeat);
    else
        cout << type << "not supported for AoT\n";
}
```

To do the same thing with the JIT feature, we can write:

```

template <typename T, int size>
[[jit]] void test_jit_sz(int repeat) {
    Matrix<T,size,size> I = Matrix<T,size,size>::Ones();
    Matrix<T,size,size> m;
    for(int i = 0; i < size; i++)
        for(int j = 0; j < size; j++) {
            m(i,j) = (i+size*j);
        }

    auto start = chrono::system_clock::now();

    for (int r = 0; r < repeat; ++r) {
        m = Matrix<T,size,size>::Ones() + T(0.00005) * (m + (m*m));
    }

    auto end = chrono::system_clock::now();
    cout << "JIT: " << chrono::duration<double>(end - start).count() << " s\r\n";
}

void test_jit(std::string &type, int size, int repeat) {
    return test_jit_sz<type, size>(repeat);
}

```

And we can use very-similar code to construct explicit instantiations at compile time, but of course, then we're limited to support for the explicit sizes we have selected.

§ 9.1. Compile Time

Compiling using the implementation linked above on an Intel Xeon E5-2699 using the flags `-march=native -ffast-math -O3`, and measuring compile time using "user" time from the Linux `time` command.

	Time	Time over Base
JIT Only	3.5s	0.92s
(AoT) Single Specialization (double, size = 16)	4.95s	2.37s
(AoT) Single Specialization (double, size = 7)	3.3s	0.72s
(AoT) Single Specialization (double, size = 3)	3.2s	0.62s
(AoT) Single Specialization (double, size = 1)	2.95s	0.37s
(AoT) Two Specializations (double, size = 16) and (double, 7)	5.7s	3.12s

	Time	Time over Base
Generic AoT Only (three floating-point types with dispatch)	9.7s	7.12s
Generic AoT Only (double only)	5.3s	2.72s
Nothing (just the includes and a main function)	2.58s	-

As you can see, the time for generating each specific specialization is essentially additive, and they get more expensive as the fixed matrix sizes get longer. Generating the code for the JIT has a compile-time cost, but it's not even as expensive as a single non-fixed-size implementation.

§ 9.2. Runtime Performance

For (double, size = 3); a repeat count of 40000000. Times as reported by the code (excludes JIT compilation time).

	Time
JIT	1.0s
Single Specialization	1.01s
AoT	8.05s

For (double, size = 7)

	Time
JIT	8.34s
Single Specialization	8.45s
AoT	20s

For (double, size = 16)

	Time
JIT	35.3s
Single Specialization	35.1s
AoT	36.2s

A few trends to notice:

The JIT-generated code is **significantly** faster than the ahead-of-time-generated code for small matrix sizes. The advantage becomes less significant as the matrix sizes become larger. At a high level, this is easy to understand: as the matrix sizes become larger, the value to the optimizer of knowing the exact size decreases.

Thus, using the JIT gives the performance advantages of using many ahead-of-time specializations, and is sometimes even better, with very low compile-time cost.

For more information, see: [arXiv:1904.08555](https://arxiv.org/abs/1904.08555).

§ 9.3. Overheads

There is no overhead to the presence of the JIT-compilation feature in the language when it is not used (although if the feature is enabled, and that causes the compiler to link with the JIT-compilation engine, that can cause both an increase in link time and binary size if the implementation is not aggressive about pruning those dependencies). When the JIT feature is used, however, there are additional costs. Clang's serialized AST can be large, and the current implementation makes no attempt to prune it before embedding it into each object file. For example, in this benchmark, the AoT-only executable is 135 KB in size. The executable making use of the JIT-compilation feature is 12 MB in size, and this does not include the size of the compiled Clang and LLVM libraries. If these libraries are linked to the application statically, instead of dynamically, the size of the binary increases to 75 MB (again, no attempt has yet been made to exclude unnecessary parts of the compiler in the current implementation). In the current implementation, the memory overhead of the JIT-compilation engine is approximately 127 MB, although efforts to prune the serialized AST might be able to reduce that significantly.

Regarding runtime overheads, on an Intel Haswell processor, for the simplest lookup involving a single template argument, the instantiation-cache lookup takes approximately 350 cycles (approximately 140 ns). Should the cache lookup fail, but solely because some types were specified using different names, resolving the instantiation request to the previously-compiled function takes approximately 160 thousand cycles (approximately 65 us). Compiling new instantiations takes, at the very least, tens of millions of cycles (a few milliseconds). The compilation time of a new instantiation, of course, depends heavily on the amount and complexity of the code being compiled (for the example above, this is essentially the time over baseline for each (AoT) single specialization).

§ 10. Other C++ JIT Approaches

Combining C++ with JIT compilation is not new. See:

- Facebook's C++ JIT [talk](#) / [poster](#)
- CERN's [Cling](#)
- From the RuntimeCompiledCPlusPlus project, [A large list](#)

An alternate method for providing an C++ JIT interface is to provide a mechanism, some kind of `std::eval`, which takes a string representing C++ code and evaluates it. Such an interface raises a number of additional questions compared to the interface proposed here, including: Should preprocessor macros be expanded inside of the string, and if so, with what preprocessor state? Do different calls to `std::eval` have any effect on each other (i.e., can one call to `std::eval` define a function called by code in another call to `std::eval`)? If so, are these state changes scoped in any way? If not, would such an interface likely be inefficient to use because it would require reprocessing large amounts of C++ code every time it is invoked? What facilities would need to be provided so that the code designed to construct the to-be-dynamically-executed C++ code could be maintainable?

The flavor of JIT-compilation proposed here aims to provide this capability in keeping with the general C++ philosophy of minimal overhead and explicit programmer control. This makes it a different kind of JIT-compilation than that provided in many of JIT-compiled languages (e.g., Java or Javascript) - in these languages all of the code is JIT-compiled (or interpreted), and sophisticated multi-tiered JIT-compilation, using techniques such as OSR (on-stack replacement) and deoptimization, is used to minimize compilation time spent on all code not measured to be hot (i.e., profiling information is not just used by the JIT-compilation engine itself, but is used to figure out what the JIT-compilation engine should compile in the first place). This proposal focuses on providing a user-directed JIT compilation of code that the programmer already knows is worth compiling during program execution. Future extensions to this proposal might allow for providing some range of possible values/types for template instantiation for use by a JIT-compilation engine capable of performing some online autotuning using that information.

§ 11. Acknowledgments

I'd like to thank David Poliakoff for a lot of testing and feedback on the implementation, and I'd like to thank the many committee members who provided me with feedback on this idea during the Kona meeting and on the reflector (including, but certainly not limited to, Herb Sutter, Chandler Carruth, Olivier Giroux, Vinod Grover, Matthias Kretz, and JF Bastien). I'd also like to thank Michał Dominiak, Nir Friedman, and Connor Waters for providing early feedback online.

This research was supported by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of two U.S. Department of Energy organizations (Office of Science and the National Nuclear Security Administration) responsible for the planning and preparation of a capable exascale ecosystem, including software, applications, hardware, advanced system engineering, and early testbed platforms, in support of the nation's exascale computing imperative. Additionally, this research used resources of the Argonne Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC02-06CH11357.

