

Conversion to literal encoding should not lead to loss of meaning

Document #:	D1854R2
Date:	2021-11-23
Programming Language C++	
Audience:	SG-16
Reply-to:	Corentin Jabot < corentin.jabot@gmail.com >

It's just semantic! - Kevlin Henney

Target

C++23

Abstract

This paper makes non-encodable characters (characters that cannot be represented in the literal encoding) in character and string literals ill-formed. We also restrict the valid characters in a multicharacter literal to avoid visual ambiguity caused by graphemes constituted of multiple codepoints.

Revisions

R2

- Modify the wording to mandate that each element of a multicharacter literal is representable as a single code unit, rather than restricting them to members of the basic character set.
- Expand motivation to clarify that multicharacter literals are changed only to avoid visual ambiguities.

R1

- Rebase the wording
- Propose to make non-basic characters in multi-characters ill-formed
- Add more motivation.

Non-encodable character-literals

Implementation defined behaviors related to conversion to literal encoding reduce the portability of C++ programs and lead to silently incorrect programs as implementations are allowed to substitute the characters they cannot represent. Strings are text which carries intent and meaning. We believe an implementation should not be able to alter that meaning.

A program should either conserve the same sequence of abstract character as in the source or be ill-formed.

This constitutes a breaking change in the wording, as well as some implementations(MSVC), and matches other existing implementations' behavior.

This is a follow-up to [P2362R3](#) [1] and a realization of the plan outlined in [P2178R1](#) [2].

Impact on the standard and implementations

- Clang always use UTF-8 to encode narrow literals
- GCC emits a diagnostics

converting to execution character set: Invalid or incomplete multibyte or wide character

- MSVC uses ? as a replacement character. For example, the string "こんにちは" becomes ' ?? ?? ', 00H. MSVC does emits warning for this scenario (enabled by default).

```
<source>(4): warning C4566: character represented by universal-character-name
'\u3053' cannot be represented in the current code page (20127)
<source>(4): warning C4566: character represented by universal-character-name
'\u3093' cannot be represented in the current code page (20127)
<source>(4): warning C4566: character represented by universal-character-name
'\u306B' cannot be represented in the current code page (20127)
<source>(4): warning C4566: character represented by universal-character-name
'\u3061' cannot be represented in the current code page (20127)
<source>(4): warning C4566: character represented by universal-character-name
'\u306F' cannot be represented in the current code page (20127)
```

A demonstration of existing behavior is available on [Compiler Explorer](#).

We argue that the code which breaks never matches the developer's intent.

Are we removing a capability?

- The exact nature of the literal encoding can be observed by a dedicated API [P1885R7](#) [3], and in general, the relying on non-encodable characters to detect the literal encoding is non-portable as it can only work on windows. It is also very difficult to use such clever tricks in a way that has no false positives or false negatives.

- `?` can be inserted in string and character literals.
- `u8` strings can be used portably.
- If the code's author does not care about the content of a string being preserved, then presumably that character can be removed.

Impact on C

This makes a behavior that is implementation-defined in C ill-formed in C++. GCC exposes the same behavior (the one proposed by this paper) in all language modes.

Multi character literals

Narrow multicharacter literals such as `'int1'` are widely used. We are not proposing to remove them. However, grapheme clusters - for example, `'é'` (e, ACUTE ACCENT) - read as single characters. A multicharacter literal can be visually indistinguishable from a character literal, leading to the accidental creation of multicharacter literals. This is the same issue described for wide literals in [P2362R3](#) [1].

Unlike what we proposed for wide-literals, we can't make all the multi-characters literals ill-formed. Instead, we propose that multi-characters literals can only contain characters representable as a single code unit.

This has two benefits

- It excludes all combining characters or characters that do not constitute a full grapheme. That takes care of the visual ambiguity.
- It makes multi-characters literals slightly less confusing as it is difficult to imagine how multiple codepoints over 0x80 could be stuffed into an `int` in any sensible way.

And indeed, the documentation of GCC shows that codepoints that do not fit in a single byte are not preserved but instead truncated.

The compiler evaluates a multicharacter character constant a character at a time, shifting the previous value left by the number of bits per target character and then or-ing in the bit-pattern of the new character truncated to the width of a target character. The final bit-pattern is given type `int`, and is therefore signed, regardless of whether single characters are signed or not. If there are more characters in the constant than would fit in the target `int` the compiler issues a warning, and the excess leading characters are ignored. For example, `'ab'` for a target with an 8-bit `char` would be interpreted as `(int) ((unsigned char) 'a' * 256 + (unsigned char) 'b')`, and `'\234a'` as `(int) ((unsigned char) '\234' * 256 + (unsigned char) 'a')`.

Because this proposal only cares about visual ambiguities between character literal and multicharacter literals, we do not propose to make any escape sequences in multicharacter

literal ill-formed (the number of escape sequences in a multicharacter or character literal is visually identifiable).

It is likely that escape sequences in multicharacter literals are in the general case non-sensible but given the implementation-defined nature of multicharacter literals, we do not think there is value in adding further restriction; our only goal is for users to accidentally introduce them.

Alternatives considered

We considered

- Restricting elements of multi-characters literals to elements of the basic translation set. However, that would exclude \$ and @.
- Restricting elements of multi-characters literals to U+0000-U+007F.

These solutions are virtually isomorphic (given that characters in these ranges are visually distinguishable). However, restricting encodability by a single unit is arguably the more direct expression of intent.

Impact on the standard and implementations

GCC, Clang ICC(EDG) emit a warning for any multi-characters literals in general. They also emit a warning when the computed value exceeds the size of `int`.

No compiler emits a warning for non-encodable characters in multicharacter literals. Because this feature cannot produce a sensible result, we do not think its removal would affect users.

Feature macro

No feature macro is proposed because the transformation to characters literals and string literals is not observable by the program.

Proposed wording

[Editor's note: The Magenta text with blue squiggly lines correspond to wording previously removed by P2362R3 [1]. This wording should stay removed.]

❖ Character literals

[lex.ccon]

A non-encodable character literal is a character-literal whose c-char-sequence consists of a single c-char that is not a numeric-escape-sequence and that specifies a character that either lacks representation in the literal's associated character encoding or that cannot be encoded as a single code unit. *A multicharacter literal is a character-literal whose c-char-sequence consists of more than one c-char.* *A multicharacter literal shall not have an encoding-prefix. If a*

multicharacter literal contains a *basic-c-char* representing a codepoint that is not encodable as a single code unit in the ordinary literal encoding, the program is ill-formed.

The *encoding-prefix* of a non-encodable character literal or a multicharacter literal shall be absent or L.

Such ~~character-literals~~ Multicharacter literals are conditionally-supported.

The kind of a *character-literal*, its type, and its associated character encoding are determined by its *encoding-prefix* and its *c-char-sequence* as defined by [lex.ccon.literal]. ~~The special cases for non-encodable character literals and multicharacter literals take precedence over their respective base kinds.~~

[*Note:* The associated character encoding for ordinary and wide character literals determines encodability, but does not determine the value of non-encodable ordinary or wide character literals or ordinary or wide multicharacter literals. The examples in [lex.ccon.literal] for non-encodable ordinary and wide character literals assume that the specified character lacks representation in the execution character set or execution wide-character set, respectively, or that encoding it would require more than one code unit. — *end note*]

Table 1: Character literals

Encoding prefix	Kind	Type	Associated character encoding	Example
none	<i>ordinary character literal</i>	char	encoding of	'v'
	non-encodable ordinary character literal	int	the execution	'\U0001F525'
	ordinary multicharacter literal	int	character set	'abcd'
L	<i>wide character literal</i>	wchar_t	encoding of	L 'w'
	non-encodable wide character literal	wchar_t	the execution	L '\U0001F32A'
	<u>wide multicharacter literal</u>	<u>wchar_t</u>	wide-character set	<u>L 'abcd'</u>
u8	<i>UTF-8 character literal</i>	char8_t	UTF-8	u8 'x'
u	<i>UTF-16 character literal</i>	char16_t	UTF-16	u 'y'
U	<i>UTF-32 character literal</i>	char32_t	UTF-32	U 'z'

In translation phase 4, the value of a *character-literal* is determined using the range of representable values of the *character-literal's* type in translation phase 7. A ~~non-encodable character literal or a~~ multicharacter literal has an implementation-defined value. The value of any other kind of *character-literal* is determined as follows:

- A *character-literal* with a *c-char-sequence* consisting of a single *basic-c-char*, *simple-escape-sequence*, or *universal-character-name* is the code unit value of the specified character as encoded in the literal's associated character encoding. [*Note:* If the specified character lacks representation in the literal's associated character encoding or if it cannot be encoded as a single code unit, then the literal is ~~a non-encodable character literal ill-formed~~. —*end note*]
- A *character-literal* with a *c-char-sequence* consisting of a single *numeric-escape-sequence* that specifies an integer value *v* has a value as follows:

- If v does not exceed the range of representable values of the *character-literal's* type, then the value is v .
 - Otherwise, if the *character-literal's encoding-prefix* is absent or L, and v does not exceed the range of representable values of the corresponding unsigned type for the underlying type of the *character-literal's* type, then the value is the unique value of the *character-literal's* type T that is congruent to v modulo 2^N , where N is the width of T .
 - Otherwise, the *character-literal* is ill-formed.
- A *character-literal* with a *c-char-sequence* consisting of a single *conditional-escape-sequence* is conditionally-supported and has an implementation-defined value.



String literals

[lex.string]

String literal objects are initialized with the sequence of code unit values corresponding to the *string-literal's* sequence of *s-char* s (for a non-raw string literal) and *r-char* s (for a raw string literal) in order as follows:

- The sequence of characters denoted by each contiguous sequence of *basic-s-char* s , *r-char* s , *simple-escape-sequence* s , and *universal-character-name* s is encoded to a code unit sequence using the *string-literal's* associated character encoding. If a character lacks representation in the associated character encoding, then : the *string-literal* is ill-formed.
 - If the *string-literal's encoding-prefix* is absent or L, then the *string-literal* is conditionally-supported and an implementation-defined code unit sequence is encoded.
 - Otherwise, the *string-literal* is ill-formed.

When encoding a stateful character encoding, implementations should encode the first such sequence beginning with the initial encoding state and encode subsequent sequences beginning with the final encoding state of the prior sequence. [Note: The encoded code unit sequence can differ from the sequence of code units that would be obtained by encoding each character independently. — end note]

- Each *numeric-escape-sequence* that specifies an integer value v contributes a single code unit with a value as follows:
 - If v does not exceed the range of representable values of the *string-literal's* array element type, then the value is v .
 - Otherwise, if the *string-literal's encoding-prefix* is absent or L, and v does not exceed the range of representable values of the corresponding unsigned type for the underlying type of the *string-literal's* array element type, then the value is the unique value of the *string-literal's* array element type T that is congruent to v modulo 2^N , where N is the width of T .

- Otherwise, the *string-literal* is ill-formed.

When encoding a stateful character encoding, these sequences should have no effect on encoding state.

- Each *conditional-escape-sequence* contributes an implementation-defined code unit sequence. When encoding a stateful character encoding, it is implementation-defined what effect these sequences have on encoding state.

Acknowledgments

Many thanks to JeanHeyd Meneide, Peter Bindels, Zach Laine, Tom Honermann, and Steve Downey, who reviewed this paper and offered valuable feedback.

References

- [1] Peter Brett and Corentin Jabot. P2362R3: Remove non-encodable wide character literals and multicharacter wide character literals. <https://wg21.link/p2362r3>, 8 2021.
- [2] Corentin Jabot. P2178R1: Misc lexing and string handling improvements. <https://wg21.link/p2178r1>, 7 2020.
- [3] Corentin Jabot and Peter Brett. P1885R7: Naming text encodings to demystify them. <https://wg21.link/p1885r7>, 9 2021.
- [N4885] Thomas Köppe *Working Draft, Standard for Programming Language C++* <https://wg21.link/N4885>