

Document Number: P1928R1
Date: 2022-10-15
Reply-to: Matthias Kretz <m.kretz@gsi.de>
Audience: SG1, LEWG
Target: C++26

MERGE DATA-PARALLEL TYPES FROM THE PARALLELISM TS 2

ABSTRACT

After the Parallelism TS 2 was published in 2018, data-parallel types (`simd<T>`) have been implemented and used. Now there is sufficient feedback to improve and merge Section 9 of the Parallelism TS 2 into the IS working draft.

CONTENTS

1	CHANGELOG	1
1.1	CHANGES FROM REVISION 0	1
1.2	CHANGES FROM REVISION 1	2
2	INTRODUCTION	2
2.1	RELATED PAPERS	2
3	CHANGES AFTER TS FEEDBACK	3
3.1	IMPROVE ABI TAGS	3
3.2	SIMPLIFY/GENERALIZE CASTS	6
3.3	ADD <code>simd_mask</code> GENERATOR CONSTRUCTOR	10
3.4	COMPLETE CASTS FOR <code>simd_mask</code>	10
3.5	<code>element_reference</code> IS OVERSPECIFIED	10
3.6	DEFAULT LOAD/STORE FLAGS TO <code>element_aligned</code>	11
3.7	CLEAN UP MATH FUNCTION OVERLOADS	11

4	OPEN QUESTIONS	12
4.1	REMOVE SUPPORT FOR SAFE PARAMETER PASSING OVER ABI BOUNDARIES	12
4.2	CASTS	12
4.3	INTEGRATION WITH RANGES	13
4.4	CORRECT PLACE FOR <code>simd</code> IN THE IS?	13
5	WORDING	14
5.1	ADD SECTION 9 OF N4808 WITH MODIFICATIONS	14
5.2	DIFF AGAINST PARALLELISM TS 2 (N4808)	46
A	BIBLIOGRAPHY	79

1

CHANGELOG

1.1

CHANGES FROM REVISION 0

Previous revision: P1928R0

- Target C++26, addressing SG1 and LEWG.
- Call for a merge of the (improved & adjusted) TS specification to the IS.
- Discuss changes to the ABI tags as consequence of TS experience; calls for polls to change the status quo.
- Add template parameter `T` to `simd_abi::fixed_size`.
- Remove `simd_abi::compatible`.
- Add (but ask for removal) `simd_abi::abi_stable`.
- Mention TS implementation in GCC releases.
- Add more references to related papers.
- Adjust the clause number for [numbers] to latest draft.
- Add open question: what is the correct clause for [simd]?
- Add open question: integration with ranges.
- Add `simd_mask` generator constructor.
- Consistently add `simd` and `simd_mask` to headings.
- Remove `experimental` and `parallelism_v2` namespaces.
- Present the wording twice: with and without diff against N4808 (Parallelism TS 2).
- Default load/store flags to `element_aligned`.
- Generalize casts: conditionally `explicit` converting constructors.
- Remove named cast functions.

1.2

CHANGES FROM REVISION 1

Previous revision: P1928R1

- ✎ Apply the new library specification style from P0788R3.
- ✎ Add `numeric_limits` / numeric traits specializations since behavior of e.g. `simd<float>` and `float` may differ for reasonable implementations.

2

INTRODUCTION

[P0214R9] introduced `simd<T>` and related types and functions into the Parallelism TS 2 Section 9. The TS was published in 2018. An incomplete and non-conforming (because P0214 evolved) implementation existed for the whole time P0214 progressed through the committee. Shortly after the GCC 9 release, a complete implementation of Section 9 of the TS was made available. Since GCC 11 a complete `simd` implementation of the TS is part of its standard library.

Note: The first revision of this paper did not proposing a merge. In the meantime the TS feedback progressed to a point where a merge should happen ASAP.

2.1

RELATED PAPERS

P0350 Before publication of the TS, SG1 approved [P0350R0] which did not progress in time in LEWG to make it into the TS. P0350 is moving forward independently.

P0918 After publication of the TS, SG1 approved [P0918R2] which adds `shuffle`, `interleave`, `sum_to`, `multiply_sum_to`, and `saturated_simd_cast`. P0918 will move forward independently.

P1068 R3 of the paper removed discussion/proposal of a `simd` based API because it was targeting C++23 with the understanding of `simd` not being ready for C++23. This is unfortunate as the presence of `simd` in the IS might lead to a considerably different assessment of the iterator/range-based API proposed in P1068.

P0917 The ability to write code that is generic wrt. arithmetic types and `simd` types is considered to be of high value (TS feedback). Conditional expressions via the `where` function were not all too well received. Conditional expressions via the conditional operator would provide a solution deemed perfect by those giving feedback (myself included).

DRAFT ON NON-MEMBER OPERATOR[] TODO

P2600 The fix for ADL is important to ensure the above two papers do not break existing code.

P0543 The paper proposing functions for saturation arithmetic expects `simd` overloads as soon as `simd` is merged to the IS.

P0553 The bit operations that are part of C++20 expects `simd` overloads as soon as `simd` is merged to the IS.

P2638 Intel's response to P1915R0 for `std::simd`

P2663 `std::simd<std::complex<T>>`.

P2664 Permutations for `simd`.

The papers P0350, P0918, P2663, P2664, and the `simd`-based P1068 fork currently have no shipping vehicle and are basically blocked on this paper.

3

CHANGES AFTER TS FEEDBACK

This section is mostly a stub. [P1915R0] (Expected Feedback from `simd` in the Parallelism TS 2) was just published and asks for specific feedback. After gathering feedback, the relevant changes will be added to a new revision of this paper.

3.1

IMPROVE ABI TAGS

I received consistent feedback that `simd_abi::compatible<T>` is the wrong default and it should rather be `simd_abi::native<T>` instead. All my tutorial material instructed users to use `std::experimental::native_simd<T>`. There really is little use for `simd_abi::compatible<T>`. The preferred approach should be the use of `simd_abi::native<T>` together with compiler flags that limit the available registers and instructions to whatever the user deems “compatible”. Consequently, there is no reason to keep `simd_abi::compatible<T>` in its current form.

Another common question was about a “fixed size” ABI tag, similar to `std::experimental::simd_abi::fixed_size<N>` but without the ABI compatibility cost.¹ Basically, the ABI footgun should be as dangerous as `std::experimental::simd_abi::native<T>`. The answer to that FAQ is to use `std::experimental::simd_abi::deduce_t<T, N>` as ABI tag. This will provide you with a high-performance footgun, if supported, but might also fall back to `std::experimental::simd_abi::fixed_size<N>`. With `std::experimental::simd_abi::deduce_t<T, N>` turning out to be used potentially more often than `std::experimental::simd_abi::fixed_size<N>` the aliases and names should be revisited. My proposal:

¹ Implementations of the TS are encouraged to make passing `fixed_size` objects ABI compatible between different hardware generations and/or even different architectures.

A0 = simd_abi::native<T>

`simd<T, A0>` abstracts a SIMD register (or similar) with highest performance on the target system (typically widest available register, but that's a QoI choice). Consequently, the number of elements is chosen by the implementation and may differ for different `T` and different compiler flags. `simd_abi::native<T>` is an alias for an unspecified type. `simd_abi::native<T>` can be an alias for `simd_abi::scalar`. If `sizeof(simd<T, A0>)` or `alignof(simd<T, A0>)` in TU1 differ from the same expressions in TU2, then the types `A0` in TU1 and TU2 have a different name.

A1 = simd_abi::fixed_size<T, N>

`simd<T, A1>` abstracts one or more registers storing `N` values. The actual hardware resources might store more values; but instructions are generated to make it appear as if there are exactly `N` values stored and manipulated.

Parameter passing may be ABI incompatible between different TUs when compiled with different compiler flags. Therefore, if `sizeof(simd<T, A0>)`² or `alignof(simd<T, A0>)` in TU1 differ from the same expressions in TU2, then the types `A1` in TU1 and TU2 have a different name. This new requirement (wrt. the TS) is the reason for the additional `T` parameter. This allows an implementation to define the `fixed_size` alias as e.g. `template <typename T, int N> using fixed_size = _Fixed<N, native<T>>;`. `A1` and `A0` are always different types, i.e. even if `simd_size_v<T, A0> == N`.

A major difference to `std::experimental::simd_abi::fixed_size<N>` in the TS is about `simd_mask`. In order to support ABI stability the `simd_mask` implementation must choose one form of storage for all possible targets:

- full SIMD registers with all bits set to 1 or 0 per corresponding element
- bitmasks
- an array of `bool` or similar

The new intent for the `fixed_size<T, N>` ABI tag would be to allow `simd_mask<T, A1>` to use either choice depending only on compiler flags.

A2 = simd_abi::abi_stable<N>

`simd<T, A2>` abstracts the same hardware resources as `simd<T, A1>`, however with the intent of ABI compatibility between differently compiled TUs. This can't be a normative requirement; and thus `A2` and `A1` may be equivalent types (QoI). (Note: There is no use for a `simd_abi::abi_stable<T>` type since the parameter passing differences are tied to the number of elements in a SIMD vector. Thus, either `simd_abi::native<T>` would work

² `A0` is not a typo; this depends on `simd_abi::native<T>`

just fine or `simd_abi::abi_stable<T>` would lead to incompatibilities, because types with the same name had different `sizeof` and `alignof`.)

A3 = `simd_abi::scalar`

No change.

At this point the `simd_abi::deduce` facility seems to be obsolete. However, it is still a useful tool for implementing the `rebind_simd` and `resize_simd` traits. Without more compelling reason for removal, it should be merged as is.

3.1.1

NAMING DISCUSSION

For context on naming, consider the use-cases that the ABI tags serve:

`simd_abi::native<T>` The equivalent to `T`: a direct abstraction of available hardware resources in terms of registers and instructions.

`simd_abi::fixed_size<T, N>` Higher abstraction level than `native<T>`: the user/algorithm dictates the number of elements to be processed in parallel. The objects might not be direct mappings to hardware resources, but they use the best that is available on the given target system.

`simd_abi::abi_stable<N>` The type to use as function parameter type on ABI boundaries. This type allows TUs translated with different machine-related compiler flags to interact correctly.

`simd_abi::scalar` The actual type of `native<T>` if the target hardware has no support for parallel processing of elements of `T`.³ In addition, `simd<T, simd_abi::scalar>` can be a useful debugging tool.

For reference, the above names are my preference over the following alternatives.

- alternatives for `fixed_size`:
 - `simd_abi::fixed_native<N>` with `simd` alias `fixed_native_simd<T, N>`
 - `simd_abi::fixed<N>` with `simd` alias `fixed_simd<T, N>`
 - `simd_abi::sized<N>` with `simd` alias `sized_simd<T, N>`
- alternatives for `abi_stable`:
 - `simd_abi::fixed_compatible<N>` with `simd` alias `fixed_compatible_simd<T, N>`

³ A typical example is `simd_abi::native<long double>`.

- `simd_abi::compatible<N>` with `simd` alias `compatible_simd<T, N>`
- `simd_abi::abi_compatible<N>` with `simd` alias `abi_compatible_simd<T, N>`
- `simd_abi::parameter<N>` with `simd` alias `parameter_simd<T, N>` or `simd_parameter<T, N>`

3.1.2

ALTERNATIVE: NO ABI BOUNDARY HELPER

See Section 4.1.

3.2

SIMPLIFY/GENERALIZE CASTS

The change to the ABI tags requires a reconsideration of cast functions and implicit and explicit casts between data-parallel types of different ABI tags. This is in addition to TS feedback on casts being too strict or cumbersome to use.

3.2.1

INTERCONVERTIBLE `fixed_size` AND `abi_stable`

Conversions between `fixed_size_simd<T, N>` and `abi_stable_simd<T, N>` should be implicit, because the only difference is in how the parameter is passed to a function. Thus, a function on an ABI boundary might want to declare function arguments with `abi_stable<N>` and then continue with `fixed_size<T, N>` in its implementation. However, conversions between `fixed_size_simd_mask<T, N>` and `abi_stable_simd_mask<T, N>` might require costly conversion between different mask representations. It would certainly be inconsistent to allow implicit conversions for `simd` but not `simd_mask`. Also this potential mask conversion cost exists for today's `std::experimental::simd_abi::fixed_size_simd_mask<T, N>` for every mask operation, so reducing it to implicit conversions on ABI boundaries would already be a huge improvement.

3.2.2

MORE (EXPLICIT) CONVERTING CONSTRUCTORS

The TS allows implicit casts between `fixed_size<N>` types that only differ in element type and where the values are preserved ("every possible value of `U` can be represented with type `value_type`").

However, from experience with the TS, it is better to also enable implicit conversions between any `simd` specializations with equal element count, even if such a conversion might be non-portable between targets with different native SIMD widths. The expectation is, that users set up their types according to a pattern similar to Listing 1. Thus, users will work with a set of types that have equal number of elements by construction. Some of the types may use the `fixed_size` ABI tag and some may use an extended ABI tag. This detail should not stop the user from being able to cast between a compiler-flag dependent subset of these types.

```

1 using floatv = std::simd<float>;
2 using doublev = std::rebind_simd_t<double, floatv>;
3 using int32v = std::rebind_simd_t<std::int32_t, floatv>;
4 using uint32v = std::rebind_simd_t<std::uint32_t, floatv>;
5 using int16v = std::rebind_simd_t<std::int16_t, floatv>;
6 using uint16v = std::rebind_simd_t<std::uint16_t, floatv>;
7 // ...

```

Listing 1: Recommended setup of simd types

Besides a constraint on the number of elements being equal, the converting constructor should be conditionally `explicit`: Implicit casts are only allowed if the element type conversion is value-preserving (same wording as in the TS).

This resolves major inconveniences when working with mixed-precision operations (cf. Tony Table 1). Type conversions for `simd` are still less error-prone than builtin types, because conversions

before	with P1928R1
<pre> namespace stdx = std::experimental; template <class T> void f(T a, int b) { using I = std::conditional_t< stdx::is_simd_v<T>, stdx::rebind_simd_t<int, T>, int>; I c; if constexpr (stdx::is_simd_v<T>) { c = stdx::static_simd_cast<I>(a) + b; } else { c = static_cast<int>(a) + b; } g(c); } </pre>	<pre> // assuming simd in namespace std template <class T> void f(T a, int b) { using I = std::conditional_t< std::is_simd_v<T>, std::rebind_simd_t<int, T>, int>; I c = static_cast<I>(a) + b; g(c); } </pre>

Tony Table 1: Improved generic code after adding converting constructors

that might lose information require an explicit cast. Also, an unintended widening of the SIMD register size can happen but is not completely invisible as shown in Listing 2.

3.2.3

REMOVE NAMED CAST FUNCTIONS

From the cast functions `std::experimental::to_fixed_size`, `std::experimental::to_native`, and `std::experimental::to_compatible` only the conversions from `simd_abi::fixed_size<T, N>` to `simd_abi::native<T>` and back may still benefit from a named cast function. Most importantly, the conversion from `native` to its `fixed_size` counterpart benefits from a cast expression that does not require spelling out the destination type. If only (explicit) converting constructors were provided by the standard library, it would be simple enough for users

```

1 void f(int32v a, doublev b, floatv c)
2 {
3     doublev x = a * b + c; // OK: implicit (value-preserving) conversion from int and float
4     // to double. Requires twice the register space, but there's no way around it and the
5     // result type requires it anyway.
6     int32v y = a * b; // ERROR: implicit conversion from double to int not value-preserving
7     int32v z1 = static_cast<int32v>(a * b); // OK: cast hints at implicit register widening
8     int32v z2 = a * static_cast<int32v>(b); // OK
9 }

```

Listing 2: Mixed precision code using the types from Listing 1, ensuring equal element count

to define their own `to_fixed_size` function if they want one (e.g. Listing 3). The reverse cast can

```

1 template <typename T>
2 constexpr auto to_fixed_size(std::simd<T> x)
3 {
4     using R = std::fixed_size_simd<T, std::simd<T>::size()>;
5     return static_cast<R>(x);
6 }

```

Listing 3: Example of a user-defined `to_fixed_size` implementation if explicit casts are provided

trivially be spelled out as `static_cast<simd<T>>(y)` in program code. The only motivation for adding a `to_native` function would be the provision of a counterpart for the `to_fixed_size` cast function.

If there is consensus to allow implicit conversions, then I believe the named cast functions should be removed altogether.

3.2.4

REMOVE `simd_cast` AND `static_simd_cast`

There are two cast function templates in the TS: `simd_cast` and `static_simd_cast`. The former is equivalent to the latter except that only value-preserving conversions are allowed. The template parameter can either be a `simd` specialization or a vectorizable type `T`. In the latter case, the cast function determines the return type as `fixed_size_simd<T, input.size()>`.

If we allow all conversions covered by `std::experimental::simd_cast` and `std::experimental::static_simd_cast` via `std::simd` constructors, then the cast functions can be removed altogether. The lost feature (cast via element type) can be replaced using `rebind_simd` as shown in Tony Table 2.

3.2.5

MASK CASTS

`simd_mask` casts should work when `simd` casts work. I.e. if `simd<T0, A0>` is implicitly convertible to `simd<T1, A1>` then `simd_mask<T0, A0>` is implicitly convertible to `simd_mask<T1,`

before	with P1928R1
<pre> template <typename V> void f(V x) { const auto y = std::static_simd_cast<double>(x); // ... } </pre>	<pre> template <typename V> void f(V x) { const auto y = std::rebind_simd_t<double, V>(x); // ... } </pre>

Tony Table 2: Casting without specifying the target ABI tag

A1>. The reverse (if `simd_mask` is convertible then `simd` is convertible) does not have to be true. Specifically, the TS allows all `fixed_size<N>` mask to be interconvertible, irrespective of the element type. For the IS merge, the proposal is to make this more consistent with `simd` while also preserving most of the convenience: Allow implicit conversions if the `sizeof` of the element types are equal, otherwise the conversion must be explicit.

Conversions with different element count are not possible via a constructor (consistent with `simd`). This would require a different function, such as the `resize<N>(simd)` function proposed by Towner et al. [P2638R0].

3.2.6

SUMMARY OF CASTS

1. `simd<T0, A0>` is convertible to `simd<T1, A1>` if `simd_size_v<T0, A0> == simd_size_v<T1, A1>`.
2. `simd<T0, A0>` is implicitly convertible to `simd<T1, A1>` if, additionally, the conversion `T0` to `T1` is value-preserving.
3. `simd_mask<T0, A0>` is convertible to `simd_mask<T1, A1>` if `simd_size_v<T0, A0> == simd_size_v<T1, A1>`.
4. `simd_mask<T0, A0>` is implicitly convertible to `simd_mask<T1, A1>` if, additionally,
5. `sizeof(T0) == sizeof(T1)`.
6. `simd<T0, A0>` can be `bit_casted` to `simd<T1, A1>` if `sizeof(simd<T0, A0>) == sizeof(simd<T1, A1>)` and neither `A0` nor `A1` are specializations of `simd_abi::abi_stable`.
7. `simd_mask<T0, A0>` can be `bit_casted` to `simd_mask<T1, A1>` if `sizeof(simd_mask<T0, A0>) == sizeof(simd_mask<T1, A1>)` and neither `A0` nor `A1` are specializations of `simd_abi::abi_stable`.

3.3

ADD `simd_mask` GENERATOR CONSTRUCTOR

The `simd` generator constructor is very useful for initializing objects from scalars in a portable (i.e. different `simd::size()`) fashion. The need for a similar constructor for `simd_mask` is less frequent, but, even if only for consistency, there should be one. Besides consistency, it is also useful, of course. Consider a predicate function that is given without `simd` interface (e.g. from a library). How do you construct a `simd_mask` from it? With a generator constructor it is easy:

```
simd<T> f(simd<T> x, Predicate p) {
    const simd_mask<T> k([&](auto i) { return p(x[i]); });
    where(k, x) = 0;
    return x;
}
```

Without the generator constructor one has to write e.g.:

```
simd<T> f(simd<T> x, Predicate p) {
    simd_mask<T> k;
    for (size_t i = 0; i < simd<T>::size(); ++i) {
        k[i] = p(x[i]);
    }
    where(k, x) = 0;
    return x;
}
```

The latter solution makes it hard to initialize the `simd_mask` as `const`, is more verbose, is harder to optimize, and cannot use the sequencing properties the generator constructor allows.

Therefore add:

```
template<class G> simd_mask(G&& gen) noexcept;
```

3.4

COMPLETE CASTS FOR `simd_mask`

The `simd_cast` and `static_simd_cast` overloads for `simd_mask` were forgotten for the TS. Without those casts (and no casts via constructors) mixing different arithmetic types is painful. There is no motivation for forbidding casts on `simd_mask`.

However, since the proposed change for casts removes `simd_cast` and `static_simd_cast`, the problem is already solved.

3.5

element_reference IS OVERSPECIFIED

`element_reference` is spelled out in a lot of detail. It may be better to define its requirements in a table instead.

This change is not reflected in the wording, pending encouragement from WG21 (mostly LWG).

3.6

DEFAULT LOAD/STORE FLAGS TO `element_aligned`

Consider:

```
1 std::simd<float> v(addr, std::vector_aligned);
2 v.copy_from(addr + 1, std::element_aligned);
3 v.copy_to(dest, std::element_aligned);
```

Line 1 supplies an optimization hint to the load operation. Line 2 says what really? “Please don’t crash. I know this is not a vector aligned access⁴.” Line 3 says: “I don’t know whether it’s vector aligned or not. Compiler, if you know more, please optimize, otherwise just don’t make it crash.” (To clarify, the difference between lines 2 and 3 is what line 1 says about the alignment of `addr`.) In both cases of `element_aligned` access, the developer requested a behavior we take as given in all other situations. Why does the TS force to spell it out in this case?

Since C++20, we also have another option:

```
1 std::simd<float> v(std::assume_aligned<std::memory_alignment_v<std::simd<float>>>>(addr));
2 v.copy_from(addr + 1);
3 v.copy_to(dest);
```

This seems to compose well, except that line 1 is rather long for a common pattern in this interface. Also, this removes implementation freedom because the library cannot statically determine the alignment properties of the pointer.

Consequently, as a minimal improvement to the TS keep the load/store flags as is, but default them to `element_aligned`. I.e.:

```
1 std::simd<float> v(addr, std::vector_aligned);
2 v.copy_from(addr + 1);
3 v.copy_to(dest);
```

3.7

CLEAN UP MATH FUNCTION OVERLOADS

The wording that produces `simd` overloads misses a few cases and leaves room for ambiguity. There is also no explicit mention of integral overloads that are supported in `<cmath>` (e.g. `std::cos(1)` calling `std::cos(double)`). At the very least, `std::abs(simd<signed-integral>)` should be specified.

Also, from implementation experience “undefined behavior” for domain, pole, or range error is unnecessary. It could either be an unspecified result or even match the expected result of the function according to Annex F in the C standard. The latter could also be only a recommendation, i.e. QoI.

This needs more work and is not reflected in the wording at this point.

⁴ Of course, vector aligned is equivalent to element aligned if `simd<float>::size() == 1`

4

OPEN QUESTIONS

4.1

REMOVE SUPPORT FOR SAFE PARAMETER PASSING OVER ABI BOUNDARIES

The feature, that `std::experimental::simd_abi::fixed_size<N>` acts as a type for safely passing over ABI boundaries was a request/idea from SG1 and was not part of the original design. My preferred approach would be that users or compilers solve this problem, and not `simd`. Users could pass data via `std::array<T, N>` instead of `std::experimental::fixed_size_simd<T, N>`; in some situations this might be less efficient, though. Alternatively, the `simd` object can be wrapped in a `struct` with non-trivial copy constructor (at least for the Itanium ABI).

A more general solution would require a language extension. An implementation could provide attributes for adjusting the calling convention for certain functions or function parameters. This would make it trivial to request less efficient copying at the place where it is needed without extra performance cost in other places.

Removing the `std::experimental::simd_abi::fixed_size<N>/std::simd_abi::abi_stable<N>` ABI would resolve the problem that one of the `simd` specializations is not trivially copyable and thus not bit-castable. Furthermore, no type in the C++ standard can *guarantee* compatibility on ABI boundaries (i.e. TUs compiled with different compiler flags), since there can be no normative wording to such an effect. Thus, the type is less useful in this respect to users than it should be.

Consequently, this is my preferred set of standard ABI tags:

- `simd_abi::native<T>`
- `simd_abi::fixed_size<T, N>`
- `simd_abi::scalar`

I.e. compared to the TS, remove `compatible<T>` and remove the ABI suggestion for `fixed_size<T, N>`. However, that goes against a decision / discussion in SG1, so I would prefer for (at least) SG1 to weigh in on this issue.

4.1.1

SUGGESTED POLL

Poll: Remove provision of `simd` type (ABI tag) that *might* help passing `simd` objects correctly over ABI boundaries

SF	F	N	A	SA

4.2

CASTS

See Section 3.2 for a discussion.

4.2.1

SUGGESTED POLLS

Poll: Allow implicit conversions between `simd<T>` and `fixed_size_simd<T, N>` (both directions, `N == simd<T>::size()`)⁵

SF	F	N	A	SA

Poll: Remove named cast functions (`to_fixed_size`, `to_native`)⁶

SF	F	N	A	SA

4.3

INTEGRATION WITH RANGES

`simd` itself is not a range. The value of a data-parallel object is not an array of elements but rather needs to be understood as a single opaque value that happens to have means for reading and writing element values. I.e. `simd<int> x = {};` does not start the lifetime of `int` objects. The `element_reference` for `operator[]` is bad enough as is. Let's not exacerbate the problem by adding iterators to `simd`.

Instead, a `simd` can be converted into an array (e.g. Listing 4). Conversely, a `std::span` with static extent can be converted into a `std::simd`.

```

1 template <class T, class A>
2 std::array<T, std::simd_size_v<T, A>> to_array(std::simd<T, A> x)
3 {
4     std::array<T, std::simd_size_v<T, A>> r;
5     x.copy_to(r.data());
6     return r;
7 }
```

Listing 4: `simd` to array conversion

I plan to pursue conversions to array and from span with static extent in a follow-up paper. I believe it is not necessary to resolve this question before merging `simd` from the TS.

4.4

CORRECT PLACE FOR `simd` IN THE IS?

While `simd` is certainly very important for numerics and therefore fits into the “Numerics library” clause, it is also more than that. E.g. `simd` can be used for vectorization of text processing. In principle `simd` should be understood similar to fundamental types. Is the “General utilities library”

⁵ This is a change of the TS status quo. The paper argues that it is an improvement.

⁶ This is a change of the TS status quo. The paper argues that if the above poll passes, the functions are unnecessary.

clause a better place? Or rename “Concurrency support library” to “Parallelism and concurrency support library” and put it there? Alternatively, add a new library clause?

The author is seeking feedback before making a recommendation.

5

WORDING

The following section presents the wording to be applied against the C++ working draft. The subsequent section, Section 5.2, reproduces the same wording as a diff against the Parallelism TS 2.

There is still work to be done. In this revision consider this section as a preview how this paper is going to present the wording.

5.1

ADD SECTION 9 OF N4808 WITH MODIFICATIONS

Add a new subclause after §28.8 [numbers]

(5.1.1) 28.9 Data-Parallel Types [simd]

(5.1.1.1) 28.9.1 General [simd.general]

- 1 The data-parallel library consists of data-parallel types and operations on these types. A data-parallel type consists of elements of an underlying arithmetic type, called the *element type*. The number of elements is a constant for each data-parallel type and called the *width* of that type.
- 2 Throughout this Clause, the term *data-parallel type* refers to all *supported* (28.9.6.1) specializations of the `simd` and `simd_mask` class templates. A *data-parallel object* is an object of *data-parallel type*.
- 3 An *element-wise operation* applies a specified operation to the elements of one or more data-parallel objects. Each such application is unsequenced with respect to the others. A *unary element-wise operation* is an element-wise operation that applies a unary operation to each element of a data-parallel object. A *binary element-wise operation* is an element-wise operation that applies a binary operation to corresponding elements of two data-parallel objects.
- 4 Throughout this Clause, the set of *vectorizable types* for a data-parallel type comprises all cv-unqualified arithmetic types other than `bool`.
- 5 [*Note*: The intent is to support acceleration through data-parallel execution resources, such as SIMD registers and instructions or execution units driven by a common instruction decoder. If such execution resources are unavailable, the interfaces support a transparent fallback to sequential execution. — *end note*]

(5.1.1.2) 28.9.2 Header `<simd>` synopsis [simd.synopsis]

```
namespace std {
  namespace simd_abi {
    using scalar = see below;
    template<class T, int N> using fixed_size = see below;
    template<class T> inline constexpr int max_fixed_size = implementation-defined;
    template<class T> using native = implementation-defined;
    template<int N> using abi_stable = see below;
```

```

template<class T, size_t N, class... Abis> struct deduce { using type = see below; };
template<class T, size_t N, class... Abis> using deduce_t =
    typename deduce<T, N, Abis...>::type;
}

```

```

struct element_aligned_tag {};
struct vector_aligned_tag {};
template<size_t> struct overaligned_tag {};
inline constexpr element_aligned_tag element_aligned{};
inline constexpr vector_aligned_tag vector_aligned{};
template<size_t N> inline constexpr overaligned_tag<N> overaligned{};

```

// 28.9.4, *simd type traits*

```

template<class T> struct is_abi_tag;
template<class T> inline constexpr bool is_abi_tag_v = is_abi_tag<T>::value;

template<class T> struct is_simd;
template<class T> inline constexpr bool is_simd_v = is_simd<T>::value;

template<class T> struct is_simd_mask;
template<class T> inline constexpr bool is_simd_mask_v = is_simd_mask<T>::value;

template<class T> struct is_simd_flag_type;
template<class T> inline constexpr bool is_simd_flag_type_v =
    is_simd_flag_type<T>::value;

template<class T, class Abi = simd_abi::native<T>> struct simd_size;
template<class T, class Abi = simd_abi::native<T>>
    inline constexpr size_t simd_size_v = simd_size<T, Abi>::value;

template<class T, class U = typename T::value_type> struct memory_alignment;
template<class T, class U = typename T::value_type>
    inline constexpr size_t memory_alignment_v = memory_alignment<T, U>::value;

template<class T, class V> struct rebind_simd { using type = see below; };
template<class T, class V> using rebind_simd_t = typename rebind_simd<T, V>::type;
template<int N, class V> struct resize_simd { using type = see below; };
template<int N, class V> using resize_simd_t = typename resize_simd<N, V>::type;

```

// 28.9.6, *Class template simd*

```

template<class T, class Abi = simd_abi::native<T>> class simd;
template<class T, int N> using fixed_size_simd = simd<T, simd_abi::fixed_size<T, N>>;
template<class T, int N> using abi_stable_simd = simd<T, simd_abi::abi_stable<N>>;

```

// 28.9.8, *Class template simd_mask*

```

template<class T, class Abi = simd_abi::native<T>> class simd_mask;

```

```
template<class T, int N> using fixed_size_simd_mask = simd_mask<T, simd_abi::fixed_size<T, N>>;
template<class T, int N> using abi_stable_simd_mask = simd_mask<T, simd_abi::abi_stable<N>>;
```

```
template<size_t... Sizes, class T, class Abi>
    tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd<T, Abi>&) noexcept;
template<size_t... Sizes, class T, class Abi>
    tuple<simd_mask<T, simd_mask_abi::deduce_t<T, Sizes>>...>
        split(const simd_mask<T, Abi>&) noexcept;
template<class V, class Abi>
    array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
        split(const simd<typename V::value_type, Abi>&) noexcept;
template<class V, class Abi>
    array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
        split(const simd_mask<typename V::simd_type::value_type, Abi>&) noexcept;

template<size_t N, class T, class A>
    array<resize_simd_t<simd_size_v<T, A> / N, simd<T, A>>, N>
        split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
    array<resize_simd_t<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
        split_by(const simd_mask<T, A>& x) noexcept;

template<class T, class... Abis>
    simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...)>>
        concat(const simd<T, Abis>&...) noexcept;
template<class T, class... Abis>
    simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...)>>
        concat(const simd_mask<T, Abis>&...) noexcept;

template<class T, class Abi, size_t N>
    resize_simd_t<simd_size_v<T, Abi> * N, simd<T, Abi>>
        concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
    resize_simd_t<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
        concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;
```

// 28.9.9.4, simd_mask *reductions*

```
template<class T, class Abi> bool all_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool any_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool none_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool some_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> int popcount(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> int find_first_set(const simd_mask<T, Abi>&);
template<class T, class Abi> int find_last_set(const simd_mask<T, Abi>&);
```

```

bool all_of(T) noexcept;
bool any_of(T) noexcept;
bool none_of(T) noexcept;
bool some_of(T) noexcept;
int popcount(T) noexcept;
int find_first_set(T);
int find_last_set(T);

```

// 28.9.5, *Where expression class templates*

```

template<class M, class T> class const_where_expression;
template<class M, class T> class where_expression;

```

// 28.9.9.5, *Where functions*

```

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, simd<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, const simd<T, Abi>&) noexcept;

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, simd_mask<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, const simd_mask<T, Abi>&) noexcept;

template<class T>
    where_expression<bool, T>
        where(see below k, T& d) noexcept;

template<class T>
    const_where_expression<bool, T>
        where(see below k, const T& d) noexcept;

```

// 28.9.7.4, *simd reductions*

```

template<class T, class Abi, class BinaryOperation = plus<>>
    T reduce(const simd<T, Abi>&,
        BinaryOperation = {});

template<class M, class V, class BinaryOperation>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        typename V::value_type identity_element,

```

```

        BinaryOperation binary_op);

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        plus<> binary_op = {}) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        multiplies<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_and<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_or<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_xor<> binary_op) noexcept;

template<class T, class Abi>
    T hmin(const simd<T, Abi>&) noexcept;
template<class M, class V>
    typename V::value_type hmin(const const_where_expression<M, V>&) noexcept;
template<class T, class Abi>
    T hmax(const simd<T, Abi>&) noexcept;
template<class M, class V>
    typename V::value_type hmax(const const_where_expression<M, V>&) noexcept;

// 28.9.7.6, Algorithms
template<class T, class Abi>
    simd<T, Abi>
        min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    simd<T, Abi>
        max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    pair<simd<T, Abi>, simd<T, Abi>>
        minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    simd<T, Abi>
        clamp(const simd<T, Abi>& v,
            const simd<T, Abi>& lo,
            const simd<T, Abi>& hi);
}

```

- ¹ The header `<simd>` defines class templates, tag types, trait types, and function templates for element-wise operations on data-parallel objects.

```

namespace simd_abi {
    using scalar = see below;
    template<class T, int N> using fixed_size = see below;
    template<class T> inline constexpr int max_fixed_size = implementation-defined;
    template<class T> using native = implementation-defined;
    template<int N> using abi_stable = see below;
}

```

- 1 An *ABI tag* is a type in the `std::simd_abi` namespace that indicates a choice of size and binary representation for objects of data-parallel type. [*Note:* The intent is for the size and binary representation to depend on the target architecture. — *end note*] The ABI tag, together with a given element type implies a number of elements. ABI tag types are used as the second template argument to `simd` and `simd_mask`.
- 2 [*Note:* The ABI tag is orthogonal to selecting the machine instruction set. The selected machine instruction set limits the usable ABI tag types, though (see 28.9.6.1). The ABI tags enable users to safely pass objects of data-parallel type between translation unit boundaries (e.g. function calls or I/O). — *end note*]
- 3 `scalar` is an alias for an unspecified ABI tag that is different from `fixed_size<1>` and `abi_stable<1>`. Use of the `scalar` tag type requires data-parallel types to store a single element (i.e., `simd<T, simd_abi::scalar>::size()` returns 1).
- 4 The value of `max_fixed_size<T>` is at least 32.
- 5 `fixed_size<N>` is an alias for an unspecified ABI tag. `fixed_size` does not introduce a non-deduced context. Use of the `simd_abi::fixed_size<N>` tag type requires data-parallel types to store `N` elements (i.e. `simd_size_v<T, simd_abi::fixed_size<N>>` is `N`). `simd<T, fixed_size<N>>` and `simd_mask<T, fixed_size<N>>` with `N > 0` and `N <= max_fixed_size<T>` shall be supported. Additionally, for every supported `simd<T, Abi>` (see 28.9.6.1), where `Abi` is an ABI tag that is not a specialization of `simd_abi::fixed_size`, `N == simd<T, Abi>::size()` shall be supported.
- 6 [*Note:* It is unspecified whether `simd<T, fixed_size<N>>` with `N > max_fixed_size<T>` is supported. The value of `max_fixed_size<T>` can depend on compiler flags and can change between different compiler versions. — *end note*]
- 7 `abi_stable<N>` is an alias for an unspecified ABI tag. `abi_stable` does not introduce a non-deduced context. Use of the `simd_abi::abi_stable<N>` tag type requires data-parallel types to store `N` elements (i.e. `simd_size_v<T, simd_abi::abi_stable<N>>` is `N`). `abi_stable_simd<T, N>` and `abi_stable_simd_mask<T, N>` are supported iff `fixed_size_simd<T, N>` and `fixed_size_simd_mask<T, N>` are supported.
- 8 The type of `fixed_size<T, N>` in TU1 differs from the type of `fixed_size<T, N>` in TU2 iff the type of `native<T>` in TU1 differs from the type of `native<T>` in TU2. The type of `abi_stable<N>` is the same in TU1 and TU2.
- 9 [*Note:* An implementation should ensure ABI compatibility between differently compiled translation units for `simd` and `simd_mask` specializations using the same `simd_abi::abi_stable<N>` tag. — *end note*]
- 10 An implementation may define additional *extended ABI tag* types in the `std::simd_abi` namespace, to support other forms of data-parallel computation.
- 11 `native<T>` is an implementation-defined alias for an ABI tag. [*Note:* The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that is supported on the currently targeted system. For target architectures with ISA extensions, compiler flags may change the type of the `native<T>` alias. — *end note*] [*Example:* Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where hardware support for `__simd256` only exists for floating-point types. The implementation therefore defines `native<T>` as an alias for

- `__simd256` if `T` is a floating-point type, and
- `__simd128` otherwise.

— *end example*]

```
template<T, size_t N, class... Abis> struct deduce { using type = see below; };
```

- 12 **The member type shall be present if and only if**
- T is a vectorizable type, and
 - `simd_abi::fixed_size<N>` is supported (see 28.9.3), and
 - every type in the Abis pack is an ABI tag.
- 13 **Where present, the member typedef type shall name an ABI tag type that satisfies**
- `simd_size<T, type> == N`, and
 - `simd<T, type>` is default constructible (see 28.9.6.1).
- If N is 1, the member typedef type is `simd_abi::scalar`. [*Note:* Implementations can base the choice on Abis, but can also ignore the Abis arguments. — *end note*]**
- 14 **The behavior of a program that adds specializations for `deduce` is undefined.**

(5.1.1.4) 28.9.4 `simd` type traits

[`simd.traits`]

```
template<class T> struct is_abi_tag { see below };
```

- 1 **The type `is_abi_tag<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if T is a standard or extended ABI tag, and `false_type` otherwise.**
- 2 **The behavior of a program that adds specializations for `is_abi_tag` is undefined.**

```
template<class T> struct is_simd { see below };
```

- 3 **The type `is_simd<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if T is a specialization of the `simd` class template, and `false_type` otherwise.**
- 4 **The behavior of a program that adds specializations for `is_simd` is undefined.**

```
template<class T> struct is_simd_mask { see below };
```

- 5 **The type `is_simd_mask<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if T is a specialization of the `simd_mask` class template, and `false_type` otherwise.**
- 6 **The behavior of a program that adds specializations for `is_simd_mask` is undefined.**

```
template<class T> struct is_simd_flag_type { see below };
```

- 7 **The type `is_simd_flag_type<class T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if T is one of**
- `element_aligned_tag`, or
 - `vector_aligned_tag`, or
 - `overaligned_tag<N>` with $N > 0$ and N an integral power of two,
- and `false_type` otherwise.**
- 8 **The behavior of a program that adds specializations for `is_simd_flag_type` is undefined.**

```
template<class T, class Abi = simd_abi::native<T>> struct simd_size { see below };
```

9 `simd_size<T, Abi>` shall have a member value if and only if

- `T` is a vectorizable type, and
- `is_abi_tag_v<Abi>` is true.

[*Note*: The rules are different from those in (28.9.6.1). — *end note*]

10 If value is present, the type `simd_size<T, Abi>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` with `N` equal to the number of elements in a `simd<T, Abi>` object. [*Note*: If `simd<T, Abi>` is not supported for the currently targeted system, `simd_size<T, Abi>::value` produces the value `simd<T, Abi>::size()` would return if it were supported. — *end note*]

11 The behavior of a program that adds specializations for `simd_size` is undefined.

```
template<class T, class U = typename T::value_type> struct memory_alignment { see below };
```

12 `memory_alignment<T, U>` shall have a member value if and only if

- `is_simd_mask_v<T>` is true and `U` is `bool`, or
- `is_simd_v<T>` is true and `U` is a vectorizable type.

13 If value is present, the type `memory_alignment<T, U>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` for some implementation-defined `N` (see 28.9.6.5 and 28.9.8.4). [*Note*: value identifies the alignment restrictions on pointers used for (converting) loads and stores for the give type `T` on arrays of type `U`. — *end note*]

14 The behavior of a program that adds specializations for `memory_alignment` is undefined.

```
template<class T, class V> struct rebind_simd { using type = see below; };
```

15 The member type is present if and only if

- `V` is either `simd<U, Abi0>` or `simd_mask<U, Abi0>`, where `U` and `Abi0` are deduced from `V`, and
- `T` is a vectorizable type, and
- `simd_abi::deduce<T, simd_size_v<U, Abi0>, Abi0>` has a member type type.

16 Let `Abi1` denote the type `deduce_t<T, simd_size_v<U, Abi0>, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if `V` is `simd<U, Abi0>` or `simd_mask<T, Abi1>` if `V` is `simd_mask<U, Abi0>`.

```
template<int N, class V> struct resize_simd { using type = see below; };
```

17 The member type is present if and only if

- `V` is either `simd<T, Abi0>` or `simd_mask<T, Abi0>`, where `T` and `Abi0` are deduced from `V`, and
- `simd_abi::deduce<T, N, Abi0>` has a member type type.

18 Let `Abi1` denote the type `deduce_t<T, N, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if `V` is `simd<T, Abi0>` or `simd_mask<T, Abi1>` if `V` is `simd_mask<T, Abi0>`.

```

template<class M, class T> class const_where_expression {
    const M mask;      // exposition only
    T& data;           // exposition only

public:
    const_where_expression(const const_where_expression&) = delete;
    const_where_expression& operator=(const const_where_expression&) = delete;

    T operator-() const && noexcept;
    T operator+() const && noexcept;
    T operator~() const && noexcept;

    template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const &&
};

template<class M, class T>
class where_expression : public const_where_expression<M, T> {
public:
    template<class U> void operator=(U&& x) && noexcept;
    template<class U> void operator+=(U&& x) && noexcept;
    template<class U> void operator-=(U&& x) && noexcept;
    template<class U> void operator*=(U&& x) && noexcept;
    template<class U> void operator/=(U&& x) && noexcept;
    template<class U> void operator%=(U&& x) && noexcept;
    template<class U> void operator&=(U&& x) && noexcept;
    template<class U> void operator|=(U&& x) && noexcept;
    template<class U> void operator^=(U&& x) && noexcept;
    template<class U> void operator<<=(U&& x) && noexcept;
    template<class U> void operator>>=(U&& x) && noexcept;

    void operator++() && noexcept;
    void operator++(int) && noexcept;
    void operator--() && noexcept;
    void operator--(int) && noexcept;

    template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
};

```

- 1 The class templates `const_where_expression` and `where_expression` abstract the notion of selecting elements of a given object of arithmetic or data-parallel type.
- 2 The first templates argument `M` shall be cv-unqualified `bool` or a cv-unqualified `simd_mask` specialization.
- 3 If `M` is `bool`, `T` shall be a cv-unqualified arithmetic type. Otherwise, `T` shall either be `M` or `typename M::simd_type`.
- 4 In this subclause, if `M` is `bool`, `data[0]` is used interchangeably for `data`, `mask[0]` is used interchangeably for `mask`, and `M::size()` is used interchangeably for 1.
- 5 The *selected indices* signify the integers $i \in \{j \in \mathbb{N} | j < M::size() \wedge \text{mask}[j]\}$. The *selected elements* signify the elements `data[i]` for all selected indices i .

6 In this subclause, the type `value_type` is an alias for `T` if `M` is `bool`, or an alias for `typename T::value_type` if `is_simd_mask_v<M>` is true.

7 [*Note:* The `where` functions 28.9.9.5 initialize `mask` with the first argument to `where` and `data` with the second argument to `where`. — *end note*]

```
T operator-() const && noexcept;
T operator+() const && noexcept;
T operator~() const && noexcept;
```

8 **Returns:** A copy of `data` with the indicated unary operator applied to all selected elements.

```
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const &&
```

9 **Requires:**

- If `M` is not `bool`, the largest selected index is less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

10 **Effects:** Copies the selected elements as if `mem[i] = static_cast<U>(data[i])` for all selected indices *i*.

11 **Throws:** Nothing.

12 **Remarks:** This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
 - `U` is `bool` and `value_type` is `bool`, or
 - `U` is a vectorizable type and `value_type` is not `bool`.

```
template<class U> void operator=(U&& x) && noexcept;
```

13 **Effects:** Replaces `data[i]` with `static_cast<T>(std::forward<U>(x))[i]` for all selected indices *i*.

14 **Remarks:** This operator shall not participate in overload resolution unless `U` is convertible to `T`.

```
template<class U> void operator+=(U&& x) && noexcept;
template<class U> void operator-=(U&& x) && noexcept;
template<class U> void operator*=(U&& x) && noexcept;
template<class U> void operator/=(U&& x) && noexcept;
template<class U> void operator%=(U&& x) && noexcept;
template<class U> void operator&=(U&& x) && noexcept;
template<class U> void operator|=(U&& x) && noexcept;
template<class U> void operator^=(U&& x) && noexcept;
template<class U> void operator<<=(U&& x) && noexcept;
template<class U> void operator>>=(U&& x) && noexcept;
```

15 *Effects:* Replaces `data[i]` with `static_cast<T>(data @ std::forward<U>(x))[i]` (where @ denotes the indicated operator) for all selected indices *i*.

16 *Remarks:* Each of these operators shall not participate in overload resolution unless the return type of `data @ std::forward<U>(x)` is convertible to `T`. It is unspecified whether the binary operator, implied by the compound assignment operator, is executed on all elements or only on the selected elements.

```
void operator++() && noexcept;
void operator++(int) && noexcept;
void operator--() && noexcept;
void operator--(int) && noexcept;
```

17 *Effects:* Applies the indicated operator to the selected elements.

18 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `T`.

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
```

19 *Requires:*

- If `is_simd_flag_type_v<U>` is true, for all selected indices *i*, *i* shall be less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

20 *Effects:* Replaces the selected elements as if `data[i] = static_cast<value_type>(mem[i])` for all selected indices *i*.

21 *Throws:* Nothing.

22 *Remarks:* This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
 - `U` is `bool` and `value_type` is `bool`, or
 - `U` is a vectorizable type and `value_type` is not `bool`.

(5.1.1.6) 28.9.6 Class template `simd` [simd.class]

(5.1.1.6.1) 28.9.6.1 Class template `simd` overview [simd.overview]

```
template<class T, class Abi> class simd {
public:
    using value_type = T;
    using reference = see below;
    using mask_type = simd_mask<T, Abi>;
```

```

using abi_type = Abi;

static constexpr size_t size() noexcept;

constexpr simd() noexcept = default;

// 28.9.6.4, simd constructors
template<class U> constexpr simd(U&& value) noexcept;
template<class U, class UAbi> explicit(see below) constexpr simd(const simd<U, UAbi>&) noexcept;
template<class G> explicit constexpr simd(G&& gen) noexcept;
template<class U, class Flags = element_aligned_tag> constexpr simd(const U* mem, Flags f = {});

// 28.9.6.5, simd copy functions
template<class U, class Flags = element_aligned_tag> constexpr void copy_from(const U* mem, Flags f
= {});
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const;

// 28.9.6.6, simd subscript operators
constexpr reference operator[](size_t);
constexpr value_type operator[](size_t) const;

// 28.9.6.7, simd unary operators
constexpr simd& operator++() noexcept;
constexpr simd operator++(int) noexcept;
constexpr simd& operator--() noexcept;
constexpr simd operator--(int) noexcept;
constexpr mask_type operator!() const noexcept;
constexpr simd operator~() const noexcept;
constexpr simd operator+() const noexcept;
constexpr simd operator-() const noexcept;

// 28.9.7.1, simd binary operators
constexpr friend simd operator+(const simd&, const simd&) noexcept;
constexpr friend simd operator-(const simd&, const simd&) noexcept;
constexpr friend simd operator*(const simd&, const simd&) noexcept;
constexpr friend simd operator/(const simd&, const simd&) noexcept;
constexpr friend simd operator%(const simd&, const simd&) noexcept;
constexpr friend simd operator&(const simd&, const simd&) noexcept;
constexpr friend simd operator|(const simd&, const simd&) noexcept;
constexpr friend simd operator^(const simd&, const simd&) noexcept;
constexpr friend simd operator<<(const simd&, const simd&) noexcept;
constexpr friend simd operator>>(const simd&, const simd&) noexcept;
constexpr friend simd operator<<(const simd&, int) noexcept;
constexpr friend simd operator>>(const simd&, int) noexcept;

// 28.9.7.2, simd compound assignment

```

```
constexpr friend simd& operator+=(simd&, const simd&) noexcept;
constexpr friend simd& operator-=(simd&, const simd&) noexcept;
constexpr friend simd& operator*=(simd&, const simd&) noexcept;
constexpr friend simd& operator/=(simd&, const simd&) noexcept;
constexpr friend simd& operator%=(simd&, const simd&) noexcept;
constexpr friend simd& operator&=(simd&, const simd&) noexcept;
constexpr friend simd& operator|=(simd&, const simd&) noexcept;
constexpr friend simd& operator^=(simd&, const simd&) noexcept;
constexpr friend simd& operator<=(simd&, const simd&) noexcept;
constexpr friend simd& operator>=(simd&, const simd&) noexcept;
constexpr friend simd& operator<=(simd&, int) noexcept;
constexpr friend simd& operator>=(simd&, int) noexcept;
```

// 28.9.7.3, *simd compare operators*

```
constexpr friend mask_type operator==(const simd&, const simd&) noexcept;
constexpr friend mask_type operator!=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator>=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator<=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator>(const simd&, const simd&) noexcept;
constexpr friend mask_type operator<(const simd&, const simd&) noexcept;
};
```

- 1 The class template `simd` is a data-parallel type. The width of a given `simd` specialization is a constant expression, determined by the template parameters.
- 2 Every specialization of `simd` is a complete type. The specialization `simd<T, Abi>` is supported if `T` is a vectorizable type and

- `Abi` is `simd_abi::scalar`, or
- `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in 28.9.3.

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd<T, Abi>` is supported. [Note: The intent is for implementations to decide on the basis of the currently targeted system. — end note]

If `simd<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd<T, Abi>>`.

[Example: Consider an implementation that defines the extended ABI tags `__simd_x` and `__gpu_y`. When the compiler is invoked to translate to a machine that has support for the `__simd_x` ABI tag for all arithmetic types other than `long double` and no support for the `__gpu_y` ABI tag, then:

- `simd<T, simd_abi::__gpu_y>` is not supported for any `T` and has a deleted constructor.
- `simd<long double, simd_abi::__simd_x>` is not supported and has a deleted constructor.
- `simd<double, simd_abi::__simd_x>` is supported.
- `simd<long double, simd_abi::scalar>` is supported.

— *end example*]

- 3 Default initialization performs no initialization of the elements; value-initialization initializes each element with `T()`. [*Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- 4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd`:

```
explicit constexpr operator implementation-defined() const;
explicit constexpr simd(const implementation-defined& init);
```

[*Example:* Consider an implementation that supports the type `__vec4f` and the function `__vec4f _vec4f_addsub(__vec4f, __vec4f)` for the currently targeted system. A user may require the use of `_vec4f_addsub` for maximum performance and thus writes:

```
using V = simd<float, simd_abi::__simd128>;
V addsub(V a, V b) {
    return static_cast<V>(_vec4f_addsub(static_cast<__vec4f>(a), static_cast<__vec4f>(b)));
}
```

— *end example*]

(5.1.1.6.2) 28.9.6.2 `simd` width [`simd.width`]

```
static constexpr size_t size() noexcept;
```

- 1 *Returns:* The width of `simd<T, Abi>`.

(5.1.1.6.3) 28.9.6.3 Element references [`simd.reference`]

- 1 A reference is an object that refers to an element in a `simd` or `simd_mask` object. `reference::value_type` is the same type as `simd::value_type` or `simd_mask::value_type`, respectively.
- 2 Class `reference` is for exposition only. An implementation is permitted to provide equivalent functionality without providing a class with this name.

```
class reference // exposition only
{
public:
    reference() = delete;
    reference(const reference&) = delete;

    operator value_type() const noexcept;

    template<class U> reference operator=(U&& x) && noexcept;

    template<class U> reference operator+=(U&& x) && noexcept;
    template<class U> reference operator-=(U&& x) && noexcept;
    template<class U> reference operator*=(U&& x) && noexcept;
    template<class U> reference operator/=(U&& x) && noexcept;
    template<class U> reference operator%=(U&& x) && noexcept;
    template<class U> reference operator|=(U&& x) && noexcept;
    template<class U> reference operator&=(U&& x) && noexcept;
```

```
template<class U> reference operator^=(U&& x) && noexcept;
template<class U> reference operator<<=(U&& x) && noexcept;
template<class U> reference operator>=(U&& x) && noexcept;
```

```
reference operator++() && noexcept;
value_type operator++(int) && noexcept;
reference operator--() && noexcept;
value_type operator--(int) && noexcept;
```

```
friend void swap(reference&& a, reference&& b) noexcept;
friend void swap(value_type& a, reference&& b) noexcept;
friend void swap(reference&& a, value_type& b) noexcept;
};
```

```
operator value_type() const noexcept;
```

3 **Returns:** The value of the element referred to by *this.

```
template<class U> reference operator=(U&& x) && noexcept;
```

4 **Effects:** Replaces the referred to element in `simd` or `simd_mask` with `static_cast<value_type>(std::forward<U>(x))`.

5 **Returns:** A copy of *this.

6 **Remarks:** This function shall not participate in overload resolution unless `declval<value_type&>() = std::forward<U>(x)` is well-formed.

```
template<class U> reference operator+=(U&& x) && noexcept;
template<class U> reference operator-=(U&& x) && noexcept;
template<class U> reference operator*=(U&& x) && noexcept;
template<class U> reference operator/=(U&& x) && noexcept;
template<class U> reference operator%=(U&& x) && noexcept;
template<class U> reference operator|=(U&& x) && noexcept;
template<class U> reference operator&=(U&& x) && noexcept;
template<class U> reference operator^=(U&& x) && noexcept;
template<class U> reference operator<<=(U&& x) && noexcept;
template<class U> reference operator>=(U&& x) && noexcept;
```

7 **Effects:** Applies the indicated compound operator to the referred to element in `simd` or `simd_mask` and `std::forward<U>(x)`.

8 **Returns:** A copy of *this.

9 **Remarks:** This function shall not participate in overload resolution unless `declval<value_type&>() @= std::forward<U>(x)` (where @= denotes the indicated compound assignment operator) is well-formed.

```
reference operator++() && noexcept;
reference operator--() && noexcept;
```

- 10 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.
 11 *Returns:* A copy of `*this`.
 12 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
value_type operator++(int) && noexcept;
value_type operator--(int) && noexcept;
```

- 13 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.
 14 *Returns:* A copy of the referred to element before applying the indicated operator.
 15 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend void swap(reference&& a, reference&& b) noexcept;
friend void swap(value_type& a, reference&& b) noexcept;
friend void swap(reference&& a, value_type& b) noexcept;
```

- 16 *Effects:* Exchanges the values `a` and `b` refer to.

(5.1.1.6.4) 28.9.6.4 `simd` constructors

[`simd.ctor`]

```
template<class U> simd(U&&) noexcept;
```

- 1 *Effects:* Constructs an object with each element initialized to the value of the argument after conversion to `value_type`.
 2 *Remarks:* Let `From` denote the type `remove_cv_t<remove_reference_t<U>>`. This constructor shall not participate in overload resolution unless:
- `From` is a vectorizable type and every possibly value of `From` can be represented with type `value_type`, or
 - `From` is not an arithmetic type and is implicitly convertible to `value_type`, or
 - `From` is `int`, or
 - `From` is `unsigned int` and `value_type` is an unsigned integral type.

```
template<class U, class UAbi> explicit(see below) constexpr simd(const simd<U, UAbi>& x) noexcept;
```

- 3 *Effects:* Constructs an object where the i^{th} element equals `static_cast<T>(x[i])` for all i in the range of `[0, size())`.
 4 *Remarks:* This constructor shall not participate in overload resolution unless `simd_size_v<U, UAbi> == size()`.

- 5 The constructor is `explicit` iff

- not every possible value of `U` can be represented with type `value_type`, or
- if both `U` and `value_type` are integral, the integer conversion rank (??) of `U` is greater than the integer conversion rank of `value_type`.

```
template<class G> simd(G&& gen) noexcept;
```

- 6 *Effects:* Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.
- 7 *Remarks:* This constructor shall not participate in overload resolution unless `simd(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.
- 8 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<class U, class Flags = element_aligned_tag> simd(const U* mem, Flags = {});
```

- 9 *Requires:*
- `[mem, mem + size())` is a valid range.
 - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
 - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
 - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.
- 10 *Effects:* Constructs an object where the i^{th} element is initialized to `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.
- 11 *Remarks:* This constructor shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and
 - `U` is a vectorizable type.

(5.1.1.6.5) 28.9.6.5 `simd` copy functions

[`simd.copy`]

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {});
```

- 1 *Requires:*
- `[mem, mem + size())` is a valid range.
 - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
 - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
 - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.
- 2 *Effects:* Replaces the elements of the `simd` object such that the i^{th} element is assigned with `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.
- 3 *Remarks:* This function shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and

- `U` is a vectorizable type.

```
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const;
```

4 **Requires:**

- `[mem, mem + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

5 **Effects:** Copies all `simd` elements as if `mem[i] = static_cast<U>(operator[])(i)` for all i in the range of `[0, size())`.

6 **Remarks:** This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- `U` is a vectorizable type.

(5.1.1.6.6) 28.9.6.6 `simd` subscript operators

[simd.subscr]

```
reference operator[](size_t i);
```

1 **Requires:** $i < \text{size}()$.

2 **Returns:** A reference (see 28.9.6.3) referring to the i^{th} element.

3 **Throws:** Nothing.

```
value_type operator[](size_t i) const;
```

4 **Requires:** $i < \text{size}()$.

5 **Returns:** The value of the i^{th} element.

6 **Throws:** Nothing.

(5.1.1.6.7) 28.9.6.7 `simd` unary operators

[simd.unary]

1 Effects in this subclause are applied as unary element-wise operations.

```
simd& operator++() noexcept;
```

2 **Effects:** Increments every element by one.

3 **Returns:** `*this`.

```
simd operator++(int) noexcept;
```

4 *Effects:* Increments every element by one.

5 *Returns:* A copy of `*this` before incrementing.

```
simd& operator--() noexcept;
```

6 *Effects:* Decrements every element by one.

7 *Returns:* `*this`.

```
simd operator--(int) noexcept;
```

8 *Effects:* Decrements every element by one.

9 *Returns:* A copy of `*this` before decrementing.

```
mask_type operator!() const noexcept;
```

10 *Returns:* A `simd_mask` object with the i^{th} element set to `!operator[](i)` for all i in the range of `[0, size())`.

```
simd operator~() const noexcept;
```

11 *Returns:* A `simd` object where each bit is the inverse of the corresponding bit in `*this`.

12 *Remarks:* This operator shall not participate in overload resolution unless `T` is an integral type.

```
simd operator+() const noexcept;
```

13 *Returns:* `*this`.

```
simd operator-() const noexcept;
```

14 *Returns:* A `simd` object where the i^{th} element is initialized to `-operator[](i)` for all i in the range of `[0, size())`.

(5.1.1.7) 28.9.7 `simd` non-member operations

[`simd.nonmembers`]

(5.1.1.7.1) 28.9.7.1 `simd` binary operators

[`simd.binary`]

```
friend simd operator+(const simd& lhs, const simd& rhs) noexcept;
friend simd operator-(const simd& lhs, const simd& rhs) noexcept;
friend simd operator*(const simd& lhs, const simd& rhs) noexcept;
friend simd operator/(const simd& lhs, const simd& rhs) noexcept;
friend simd operator%(const simd& lhs, const simd& rhs) noexcept;
friend simd operator&(const simd& lhs, const simd& rhs) noexcept;
friend simd operator|(const simd& lhs, const simd& rhs) noexcept;
friend simd operator^(const simd& lhs, const simd& rhs) noexcept;
friend simd operator<<(const simd& lhs, const simd& rhs) noexcept;
friend simd operator>>(const simd& lhs, const simd& rhs) noexcept;
```

1 **Returns:** A `simd` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

2 **Remarks:** Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend simd operator<<(const simd& v, int n) noexcept;
friend simd operator>>(const simd& v, int n) noexcept;
```

3 **Returns:** A `simd` object where the i^{th} element is initialized to the result of applying the indicated operator to `v[i]` and `n` for all i in the range of $[0, \text{size}())$.

4 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

(5.1.1.7.2) 28.9.7.2 `simd` compound assignment

[`simd.cassign`]

```
friend simd& operator+=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator-=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator*=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator/=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator%=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator&=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator|=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator^=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator<<=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator>>=(simd& lhs, const simd& rhs) noexcept;
```

1 **Effects:** These operators apply the indicated operator to `lhs` and `rhs` as an element-wise operation.

2 **Returns:** `lhs`.

3 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend simd& operator<<=(simd& lhs, int n) noexcept;
friend simd& operator>>=(simd& lhs, int n) noexcept;
```

4 **Effects:** Equivalent to: `return operator@=(lhs, simd(n));`

5 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

(5.1.1.7.3) 28.9.7.3 `simd` compare operators

[`simd.comparison`]

```
friend mask_type operator==(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator!=(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator>=(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator<=(const simd& lhs, const simd& rhs) noexcept;
```

```
friend mask_type operator>(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator<(const simd& lhs, const simd& rhs) noexcept;
```

- 1 **Returns:** A `simd_mask` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

(5.1.1.7.4) 28.9.7.4 `simd` reductions

[`simd.reductions`]

- 1 In this subclause, `BinaryOperation` shall be a binary element-wise operation.

```
template<class T, class Abi, class BinaryOperation = plus<>>
    T reduce(const simd<T, Abi>& x, BinaryOperation binary_op = {});
```

- 2 **Requires:** `binary_op` shall be callable with two arguments of type `T` returning `T`, or callable with two arguments of type `simd<T, A1>` returning `simd<T, A1>` for every `A1` that is an ABI tag type.
- 3 **Returns:** `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all `i` in the range of `[0, size())`(??).
- 4 **Throws:** Any exception thrown from `binary_op`.

```
template<class M, class V, class BinaryOperation>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
                                   typename V::value_type identity_element,
                                   BinaryOperation binary_op = {});
```

- 5 **Requires:** `binary_op` shall be callable with two arguments of type `T` returning `T`, or callable with two arguments of type `simd<T, A1>` returning `simd<T, A1>` for every `A1` that is an ABI tag type. The results of `binary_op(identity_element, x)` and `binary_op(x, identity_element)` shall be equal to `x` for all finite values `x` representable by `V::value_type`.
- 6 **Returns:** If `none_of(x.mask)`, returns `identity_element`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.
- 7 **Throws:** Any exception thrown from `binary_op`.

```
template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x, plus<> binary_op) noexcept;
```

- 8 **Returns:** If `none_of(x.mask)`, returns 0. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x, multiplies<> binary_op) noexcept;
```

- 9 **Returns:** If `none_of(x.mask)`, returns 1. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x, bit_and<> binary_op) noexcept;
```

- 10 *Requires:* `is_integral_v<V::value_type>` is true.
- 11 *Returns:* If `none_of(x.mask)`, returns `~V::value_type()`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices *i*.

```
template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x, bit_or<> binary_op) noexcept;
template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x, bit_xor<> binary_op) noexcept;
```

- 12 *Requires:* `is_integral_v<V::value_type>` is true.
- 13 *Returns:* If `none_of(x.mask)`, returns 0. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices *i*.

```
template<class T, class Abi> T hmin(const simd<T, Abi>& x) noexcept;
```

- 14 *Returns:* The value of an element `x[j]` for which `x[j] <= x[i]` for all *i* in the range of `[0, size())`.

```
template<class M, class V> typename V::value_type hmin(const const_where_expression<M, V>& x) noexcept;
```

- 15 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::max()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] <= x.data[i]` for all selected indices *i*.

```
template<class T, class Abi> T hmax(const simd<T, Abi>& x) noexcept;
```

- 16 *Returns:* The value of an element `x[j]` for which `x[j] >= x[i]` for all *i* in the range of `[0, size())`.

```
template<class M, class V> typename V::value_type hmax(const const_where_expression<M, V>& x) noexcept;
```

- 17 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::lowest()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] >= x.data[i]` for all selected indices *i*.

(5.1.1.7.5) 28.9.7.5 Casts

[simd.casts]

```
template<class T, class U, class Abi> see below simd_cast(const simd<U, Abi>& x) noexcept;
```

- 1 Let `To` denote `T::value_type` if `is_simd_v<T>` is true, or `T` otherwise.
- 2 *Returns:* A `simd` object with the *i*th element initialized to `static_cast<To>(x[i])` for all *i* in the range of `[0, size())`.
- 3 *Remarks:* The function shall not participate in overload resolution unless
- every possible value of type `U` can be represented with type `To`, and
 - either
 - `is_simd_v<T>` is false, or

– `T::size() == simd<U, Abi>::size()` is true.

4 The return type is

- `T` if `is_simd_v<T>` is true;
- otherwise, `simd<T, Abi>` if `U` is the same type as `T`;
- otherwise, `simd<T, simd_abi::fixed_size<simd<U, Abi>::size()>>`

```
template<class T, class U, class Abi> see below static_simd_cast(const simd<U, Abi>& x) noexcept;
```

5 Let `To` denote `T::value_type` if `is_simd_v<T>` is true or `T` otherwise.

6 **Returns:** A `simd` object with the i^{th} element initialized to `static_cast<To>(x[i])` for all i in the range of `[0, size())`.

7 **Remarks:** The function shall not participate in overload resolution unless either

- `is_simd_v<T>` is false, or
- `T::size() == simd<U, Abi>::size()` is true.

8 The return type is

- `T` if `is_simd_v<T>` is true;
- otherwise, `simd<T, Abi>` if either `U` is the same type as `T` or `make_signed_t<U>` is the same type as `make_signed_t<T>`;
- otherwise, `simd<T, simd_abi::fixed_size<simd<U, Abi>::size()>>`.

```
template<class T, class U, class Abi> see below simd_cast(const simd_mask<U, Abi>& x) noexcept;
```

```
template<class T, class U, class Abi> see below static_simd_cast(const simd_mask<U, Abi>& x) noexcept;
```

9 **Returns:** A `simd_mask` object with the i^{th} element initialized to `x[i]` for all i in the range of `[0, size())`.

10 **Remarks:** The functions shall not participate in overload resolution unless either

- `is_simd_mask_v<T>` is false, or
- `T::size() == simd_size_v<U, Abi>` is true.

11 The return type is

- `T` if `is_simd_mask_v<T>` is true;
- otherwise, `simd_mask<T, Abi>` if either `U` is the same type as `T` or `make_signed_t<U>` is the same type as `make_signed_t<T>`;
- otherwise, `simd_mask<T, simd_abi::fixed_size<simd_size_v<U, Abi>>>`

```
template<class T, class Abi>
```

```
    fixed_size_simd<T, simd_size_v<T, Abi>> to_fixed_size(const simd<T, Abi>& x) noexcept;
```

```
template<class T, class Abi>
```

```
    fixed_size_simd_mask<T, simd_size_v<T, Abi>> to_fixed_size(const simd_mask<T, Abi>& x) noexcept;
```

12 **Returns:** A data-parallel object with the i^{th} element initialized to `x[i]` for all i in the range of `[0, size())`.

```
template<class T, int N> native_simd<T> to_native(const fixed_size_simd<T, N>& x) noexcept;
```

```
template<class T, int N> native_simd_mask<T> to_native(const fixed_size_simd_mask<T, N>& x) noexcept;
```

- 13 **Returns:** A data-parallel object with the i^{th} element initialized to $x[i]$ for all i in the range of $[0, \text{size}())$.
- 14 **Remarks:** These functions shall not participate in overload resolution unless `simd_size_v<T, simd_abi::native<T>> == N` is true.

```
template<class T, int N> simd<T> to_compatible(const fixed_size_simd<T, N>& x) noexcept;
template<class T, int N> simd_mask<T> to_compatible(const fixed_size_simd_mask<T, N>& x) noexcept;
```

- 15 **Returns:** A data-parallel object with the i^{th} element initialized to $x[i]$ for all i in the range of $[0, \text{size}())$.
- 16 **Remarks:** These functions shall not participate in overload resolution unless `simd_size_v<T, simd_abi::compatible<T>> == N` is true.

```
template<size_t... Sizes, class T, class Abi>
    tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd<T, Abi>& x) noexcept;
template<size_t... Sizes, class T, class Abi>
    tuple<simd_mask<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd_mask<T, Abi>& x) noexcept;
```

- 17 **Returns:** A tuple of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} tuple element initialized to the value of the element x with index $i + \text{sum of the first } j \text{ values in the Sizes pack}$.
- 18 **Remarks:** These functions shall not participate in overload resolution unless the sum of all values in the `Sizes` pack is equal to `simd_size_v<T, Abi>`.

```
template<class V, class Abi>
    array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
        split(const simd<typename V::value_type, Abi>& x) noexcept;
template<class V, class Abi>
    array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
        split(const simd_mask<typename V::simd_type::value_type, Abi>& x) noexcept;
```

- 19 **Returns:** An array of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} array element initialized to the value of the element in x with index $i + j * V::size()$.
- 20 **Remarks:** These functions shall not participate in overload resolution unless either:
- `is_simd_v<V>` is true and `simd_size_v<typename V::value_type, Abi>` is an integral multiple of `V::size()`, or
 - `is_simd_mask_v<V>` is true and `simd_size_v<typename V::simd_type::value_type, Abi>` is an integral multiple of `V::size()`.

```
template<size_t N, class T, class A>
    array<resize_simd<simd_size_v<T, A> / N, simd<T, A>>, N>
        split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
    array<resize_simd<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
        split_by(const simd_mask<T, A>& x) noexcept;
```

- 21 **Returns:** An array `arr`, where `arr[i][j]` is initialized by `x[i * (simd_size_v<T, A> / N) + j]`.
- 22 **Remarks:** The functions shall not participate in overload resolution unless `simd_size_v<T, A>` is an integral multiple of `N`.

```
template<class T, class... Abis>
    simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
        const simd<T, Abis>&... xs) noexcept;
template<class T, class... Abis>
    simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
        const simd_mask<T, Abis>&... xs) noexcept;
```

- 23 **Returns:** A data-parallel object initialized with the concatenated values in the `xs` pack of data-parallel objects: The i^{th} `simd/simd_mask` element of the j^{th} parameter in the `xs` pack is copied to the return value's element with index i + the sum of the width of the first j parameters in the `xs` pack.

```
template<class T, class Abi, size_t N>
    resize_simd<simd_size_v<T, Abi> * N, simd<T, Abi>>
        concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
    resize_simd<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
        concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;
```

- 24 **Returns:** A data-parallel object, the i^{th} element of which is initialized by `arr[i / simd_size_v<T, Abi>][i % simd_size_v<T, Abi>]`.

(5.1.1.7.6) 28.9.7.6 Algorithms

[simd.alg]

```
template<class T, class Abi> simd<T, Abi> min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 1 **Returns:** The result of the element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())`.

```
template<class T, class Abi> simd<T, Abi> max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 2 **Returns:** The result of the element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())`.

```
template<class T, class Abi>
    pair<simd<T, Abi>, simd<T, Abi>> minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 3 **Returns:** A pair initialized with
- the result of element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())` in the first member, and
 - the result of element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())` in the second member.

```
template<class T, class Abi> simd<T, Abi>
clamp(const simd<T, Abi>& v, const simd<T, Abi>& lo, const simd<T, Abi>& hi);
```

- 4 *Requires:* No element in `lo` shall be greater than the corresponding element in `hi`.
- 5 *Returns:* The result of element-wise application of `std::clamp(v[i], lo[i], hi[i])` for all i in the range of `[0, size())`.

(5.1.1.7.7) 28.9.7.7 `simd` math library [simd.math]

- 1 For each set of overloaded functions within `<cmath>`, there shall be additional overloads sufficient to ensure that if any argument corresponding to a double parameter has type `simd<T, Abi>`, where `is_floating_point_v<T>` is true, then:
- All arguments corresponding to double parameters shall be convertible to `simd<T, Abi>`.
 - All arguments corresponding to double* parameters shall be of type `simd<T, Abi>*`.
 - All arguments corresponding to parameters of integral type `U` shall be convertible to `fixed_size_simd<U, simd_size_v<T, Abi>>`.
 - All arguments corresponding to `U*`, where `U` is integral, shall be of type `fixed_size_simd<U, simd_size_v<T, Abi>>*`.
 - If the corresponding return type is `double`, the return type of the additional overloads is `simd<T, Abi>`. Otherwise, if the corresponding return type is `bool`, the return type of the additional overload is `simd_mask<T, Abi>`. Otherwise, the return type is `fixed_size_simd<R, simd_size_v<T, Abi>>`, with `R` denoting the corresponding return type.

It is unspecified whether a call to these overloads with arguments that are all convertible to `simd<T, Abi>` but are not of type `simd<T, Abi>` is well-formed.

- 2 Each function overload produced by the above rules applies the indicated `<cmath>` function element-wise. For the mathematical functions, the results per element only need to be approximately equal to the application of the function which is overloaded for the element type.
- 3 The result is unspecified if a domain, pole, or range error occurs when the input argument(s) are applied to the indicated `<cmath>` function. [*Note:* Implementations are encouraged to follow the C specification (especially Annex F). — *end note*]
- 4 TODO: Allow `abs(simd<signed-integral>)`.
- 5 If `abs` is called with an argument of type `simd<X, Abi>` for which `is_unsigned_v<X>` is true, the program is ill-formed.

(5.1.1.8) 28.9.8 Class template `simd_mask` [simd.mask.class]

(5.1.1.8.1) 28.9.8.1 Class template `simd_mask` overview [simd.mask.overview]

```
template<class T, class Abi> class simd_mask {
public:
    using value_type = bool;
    using reference = see below;
    using simd_type = simd<T, Abi>;
```

```

using abi_type = Abi;

static constexpr size_t size() noexcept;

simd_mask() noexcept = default;

// 28.9.8.3, simd_mask constructors
explicit simd_mask(value_type) noexcept;
template<class U>
    simd_mask(const simd_mask<U, simd_abi::fixed_size<size()>>&) noexcept;
template<class G> explicit simd_mask(G&& gen) noexcept;
template<class Flags = element_aligned_tag> simd_mask(const value_type* mem, Flags = {});

// 28.9.8.4, simd_mask copy functions
template<class Flags = element_aligned_tag> void copy_from(const value_type* mem, Flags = {});
template<class Flags = element_aligned_tag> void copy_to(value_type* mem, Flags = {});

// 28.9.8.5, simd_mask subscript operators
reference operator[](size_t);
value_type operator[](size_t) const;

// 28.9.8.6, simd_mask unary operators
simd_mask operator!() const noexcept;

// 28.9.9.1, simd_mask binary operators
friend simd_mask operator&&(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator||(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator&(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator|(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator^(const simd_mask&, const simd_mask&) noexcept;

// 28.9.9.2, simd_mask compound assignment
friend simd_mask& operator&=(simd_mask&, const simd_mask&) noexcept;
friend simd_mask& operator|=(simd_mask&, const simd_mask&) noexcept;
friend simd_mask& operator^=(simd_mask&, const simd_mask&) noexcept;

// 28.9.9.3, simd_mask comparisons
friend simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
};

```

- ¹ The class template `simd_mask` is a data-parallel type with the element type `bool`. The width of a given `simd_mask` specialization is a constant expression, determined by the template parameters. Specifically, `simd_mask<T, Abi>::size() == simd<T, Abi>::size()`.
- ² Every specialization of `simd_mask` is a complete type. The specialization `simd_mask<T, Abi>` is supported if `T` is a vectorizable type and
 - `Abi` is `simd_abi::scalar`, or

- `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in (28.9.3).

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd_mask<T, Abi>` is supported. [*Note:* The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd_mask<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd_mask<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd_mask<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd_mask<T, Abi>>`.

3 Default initialization performs no initialization of the elements; value-initialization initializes each element with false. [*Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note*]

4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd_mask`:

```
explicit operator implementation-defined() const;
explicit simd_mask(const implementation-defined& init) const;
```

5 The member type reference has the same interface as `simd<T, Abi>::reference`, except its `value_type` is `bool`. (28.9.6.3)

(5.1.1.8.2) 28.9.8.2 `simd_mask` width [simd.mask.width]

```
static constexpr size_t size() noexcept;
```

1 **Returns:** The width of `simd<T, Abi>`.

(5.1.1.8.3) 28.9.8.3 `simd_mask` constructors [simd.mask.ctor]

```
explicit simd_mask(value_type x) noexcept;
```

1 **Effects:** Constructs an object with each element initialized to `x`.

```
template<class U> simd_mask(const simd_mask<U, simd_abi::fixed_size<size()>>& x) noexcept;
```

2 **Effects:** Constructs an object of type `simd_mask` where the i^{th} element equals `x[i]` for all i in the range of `[0, size())`.

3 **Remarks:** This constructor shall not participate in overload resolution unless `abi_type` is `simd_abi::fixed_size<size()>`.

```
template<class G> simd_mask(G&& gen) noexcept;
```

4 **Effects:** Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.

5 **Remarks:** This constructor shall not participate in overload resolution unless `static_cast<bool>(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.

6 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<class Flags = element_aligned_tag> simd_mask(const value_type* mem, Flags = {});
```

- 7 **Requires:**
- [mem, mem + size()) is a valid range.
 - If the template parameter Flags is vector_aligned_tag, mem shall point to storage aligned by memory_alignment_v<simd_mask>.
 - If the template parameter Flags is overaligned_tag<N>, mem shall point to storage aligned by N.
 - If the template parameter Flags is element_aligned_tag, mem shall point to storage aligned by alignof(value_type).
- 8 **Effects:** Constructs an object where the i^{th} element is initialized to mem[i] for all i in the range of [0, size()).
- 9 **Throws:** Nothing.
- 10 **Remarks:** This constructor shall not participate in overload resolution unless is_simd_flag_type_v<Flags> is true.

(5.1.1.8.4) 28.9.8.4 simd_mask copy functions

[simd.mask.copy]

```
template<class Flags = element_aligned_tag> void copy_from(const value_type* mem, Flags = {});
```

- 1 **Requires:**
- [mem, mem + size()) is a valid range.
 - If the template parameter Flags is vector_aligned_tag, mem shall point to storage aligned by memory_alignment_v<simd_mask>.
 - If the template parameter Flags is overaligned_tag<N>, mem shall point to storage aligned by N.
 - If the template parameter Flags is element_aligned_tag, mem shall point to storage aligned by alignof(value_type).
- 2 **Effects:** Replaces the elements of the simd_mask object such that the i^{th} element is replaced with mem[i] for all i in the range of [0, size()).
- 3 **Throws:** Nothing.
- 4 **Remarks:** This function shall not participate in overload resolution unless is_simd_flag_type_v<Flags> is true.

```
template<class Flags = element_aligned_tag> void copy_to(value_type* mem, Flags = {});
```

- 5 **Requires:**
- [mem, mem + size()) is a valid range.
 - If the template parameter Flags is vector_aligned_tag, mem shall point to storage aligned by memory_alignment_v<simd_mask>.
 - If the template parameter Flags is overaligned_tag<N>, mem shall point to storage aligned by N.
 - If the template parameter Flags is element_aligned_tag, mem shall point to storage aligned by alignof(value_type).

- 6 *Effects:* Copies all `simd_mask` elements as if `mem[i] = operator[](i)` for all *i* in the range of `[0, size())`.
 7 *Throws:* Nothing.
 8 *Remarks:* This function shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is `true`.

(5.1.1.8.5) 28.9.8.5 `simd_mask` subscript operators [`simd.mask.subscr`]

```
reference operator[](size_t i);
```

- 1 *Requires:* `i < size()`.
 2 *Returns:* A reference (see 28.9.6.3) referring to the *i*th element.
 3 *Throws:* Nothing.

```
value_type operator[](size_t i) const;
```

- 4 *Requires:* `i < size()`.
 5 *Returns:* The value of the *i*th element.
 6 *Throws:* Nothing.

(5.1.1.8.6) 28.9.8.6 `simd_mask` unary operators [`simd.mask.unary`]

```
simd_mask operator!() const noexcept;
```

- 1 *Returns:* The result of the element-wise application of `operator!`.

(5.1.1.9) 28.9.9 Non-member operations [`simd.mask.nonmembers`]

(5.1.1.9.1) 28.9.9.1 `simd_mask` binary operators [`simd.mask.binary`]

```
friend simd_mask operator&&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator||(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator|(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator^(const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

- 1 *Returns:* A `simd_mask` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

(5.1.1.9.2) 28.9.9.2 `simd_mask` compound assignment [`simd.mask.cassign`]

```
friend simd_mask& operator&=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask& operator|=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask& operator^=(simd_mask& lhs, const simd_mask& rhs) noexcept;
```

- 1 *Effects:* These operators apply the indicated operator to lhs and rhs as a binary element-wise operation.
 2 *Returns:* lhs.

(5.1.1.9.3) 28.9.9.3 simd_mask comparisons [simd.mask.comparison]

```
friend simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
```

- 1 *Returns:* A simd_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(5.1.1.9.4) 28.9.9.4 simd_mask reductions [simd.mask.reductions]

```
template<class T, class Abi> bool all_of(const simd_mask<T, Abi>& k) noexcept;
```

- 1 *Returns:* true if all boolean elements in k are true, false otherwise.

```
template<class T, class Abi> bool any_of(const simd_mask<T, Abi>& k) noexcept;
```

- 2 *Returns:* true if at least one boolean element in k is true, false otherwise.

```
template<class T, class Abi> bool none_of(const simd_mask<T, Abi>& k) noexcept;
```

- 3 *Returns:* true if none of the one boolean elements in k is true, false otherwise.

```
template<class T, class Abi> bool some_of(const simd_mask<T, Abi>& k) noexcept;
```

- 4 *Returns:* true if at least one of the one boolean elements in k is true and at least one of the boolean elements in k is false, false otherwise.

```
template<class T, class Abi> int popcount(const simd_mask<T, Abi>& k) noexcept;
```

- 5 *Returns:* The number of boolean elements in k that are true.

```
template<class T, class Abi> int find_first_set(const simd_mask<T, Abi>& k);
```

- 6 *Requires:* any_of(k) returns true.

- 7 *Returns:* The lowest element index i where $k[i]$ is true.

- 8 *Throws:* Nothing.

```
template<class T, class Abi> int find_last_set(const simd_mask<T, Abi>& k);
```

- 9 *Requires:* any_of(k) returns true.

- 10 *Returns:* The greatest element index i where $k[i]$ is true.

- 11 *Throws:* Nothing.

```

bool all_of(T) noexcept;
bool any_of(T) noexcept;
bool none_of(T) noexcept;
bool some_of(T) noexcept;
int popcount(T) noexcept;

```

12 *Returns:* `all_of` and `any_of` return their arguments; `none_of` returns the negation of its argument; `some_of` returns false; `popcount` returns the integral representation of its argument.

13 *Remarks:* The parameter type `T` is an unspecified type that is only constructible via implicit conversion from `bool`.

```

int find_first_set(T);
int find_last_set(T);

```

14 *Requires:* The value of the argument is true.

15 *Returns:* 0.

16 *Throws:* Nothing.

17 *Remarks:* The parameter type `T` is an unspecified type that is only constructible via implicit conversion from `bool`.

(5.1.1.9.5) 28.9.9.5 `where` functions

[`simd.mask.where`]

```

template<class T, class Abi>
  where_expression<simd_mask<T, Abi>, simd<T, Abi>>
    where(const typename simd<T, Abi>::mask_type& k, simd<T, Abi>& v) noexcept;
template<class T, class Abi>
  const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
    where(const typename simd<T, Abi>::mask_type& k, const simd<T, Abi>& v) noexcept;
template<class T, class Abi>
  where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
    where(const type_identity_t<simd_mask<T, Abi>>& k, simd_mask<T, Abi>& v) noexcept;
template<class T, class Abi>
  const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
    where(const type_identity_t<simd_mask<T, Abi>>& k, const simd_mask<T, Abi>& v) noexcept;

```

1 *Returns:* An object (28.9.5) with `mask` and `data` initialized with `k` and `v` respectively.

```

template<class T>
  where_expression<bool T>
    where(see below k, T& v) noexcept;
template<class T>
  const_where_expression<bool, T>
    where(see below k, const T& v) noexcept;

```

2 *Remarks:* The functions shall not participate in overload resolution unless

- `T` is neither a `simd` nor a `simd_mask` specialization, and
- the first argument is of type `bool`.

3 *Returns:* An object (28.9.5) with `mask` and `data` initialized with `k` and `v` respectively.

5.2

DIFF AGAINST PARALLELISM TS 2 (N4808)

In the following, the wording from Section 5.1 is repeated with additional indications of differences with regard to N4808. Changes relative to N4808, which contains editorial changes after the publication of the TS, are marked using color for **additions** and **removals**.

(5.2.1) 28.9 Data-Parallel Types [simd]

(5.2.1.1) 28.9.1 General [simd.general]

- 1 The data-parallel library consists of data-parallel types and operations on these types. A data-parallel type consists of elements of an underlying arithmetic type, called the *element type*. The number of elements is a constant for each data-parallel type and called the *width* of that type.
- 2 Throughout this Clause, the term *data-parallel type* refers to all *supported* (28.9.6.1) specializations of the `simd` and `simd_mask` class templates. A *data-parallel object* is an object of *data-parallel type*.
- 3 An *element-wise operation* applies a specified operation to the elements of one or more data-parallel objects. Each such application is unsequenced with respect to the others. A *unary element-wise operation* is an element-wise operation that applies a unary operation to each element of a data-parallel object. A *binary element-wise operation* is an element-wise operation that applies a binary operation to corresponding elements of two data-parallel objects.
- 4 Throughout this Clause, the set of *vectorizable types* for a data-parallel type comprises all cv-unqualified arithmetic types other than `bool`.
- 5 [*Note:* The intent is to support acceleration through data-parallel execution resources, such as SIMD registers and instructions or execution units driven by a common instruction decoder. If such execution resources are unavailable, the interfaces support a transparent fallback to sequential execution. — *end note*]

(5.2.1.2) 28.9.2 Header ~~<experimental/simd>~~ synopsis [simd.synopsis]

```
namespace std::experimental { inline namespace parallelism_v2 {
    namespace simd_abi {
        using scalar = see below;
        template<class T, int N> using fixed_size = see below;
        template<class T> inline constexpr int max_fixed_size = implementation-defined;
        template<class T> using compatible = implementation-defined;
        template<class T> using native = implementation-defined;
        template<int N> using abi_stable = see below;

        template<class T, size_t N, class... Abis> struct deduce { using type = see below; };
    }
}
```

```

    template<class T, size_t N, class... Abis> using deduce_t =
        typename deduce<T, N, Abis...>::type;
}

struct element_aligned_tag {};
struct vector_aligned_tag {};
template<size_t> struct overaligned_tag {};
inline constexpr element_aligned_tag element_aligned{};
inline constexpr vector_aligned_tag vector_aligned{};
template<size_t N> inline constexpr overaligned_tag<N> overaligned{};

// 28.9.4, simd type traits
template<class T> struct is_abi_tag;
template<class T> inline constexpr bool is_abi_tag_v = is_abi_tag<T>::value;

template<class T> struct is_simd;
template<class T> inline constexpr bool is_simd_v = is_simd<T>::value;

template<class T> struct is_simd_mask;
template<class T> inline constexpr bool is_simd_mask_v = is_simd_mask<T>::value;

template<class T> struct is_simd_flag_type;
template<class T> inline constexpr bool is_simd_flag_type_v =
    is_simd_flag_type<T>::value;

template<class T, class Abi = simd_abi::compatible_native<T>> struct simd_size;
template<class T, class Abi = simd_abi::compatible_native<T>>
    inline constexpr size_t simd_size_v = simd_size<T, Abi>::value;

template<class T, class U = typename T::value_type> struct memory_alignment;
template<class T, class U = typename T::value_type>
    inline constexpr size_t memory_alignment_v = memory_alignment<T, U>::value;

template<class T, class V> struct rebind_simd { using type = see below; };
template<class T, class V> using rebind_simd_t = typename rebind_simd<T, V>::type;
template<int N, class V> struct resize_simd { using type = see below; };
template<int N, class V> using resize_simd_t = typename resize_simd<N, V>::type;

// 28.9.6, Class template simd
template<class T, class Abi = simd_abi::compatible_native<T>> class simd;
template<class T> using native_simd = simd<T, simd_abi::native<T>>
template<class T, int N> using fixed_size_simd = simd<T, simd_abi::fixed_size<T, N>>;
template<class T, int N> using abi_stable_simd = simd<T, simd_abi::abi_stable<N>>

// 28.9.8, Class template simd_mask

```

```

template<class T, class Abi = simd_abi::compatiblenative<T>> class simd_mask;
template<class T> using native_simd_mask = simd_mask<T, simd_abi::native<T>>
template<class T, int N> using fixed_size_simd_mask = simd_mask<T, simd_abi::fixed_size<T, N>>;
template<class T, int N> using abi_stable_simd_mask = simd_mask<T, simd_abi::abi_stable<N>>

// 28.9.7.5, Casts
template<class T, class U, class Abi> see below simd_cast(const simd<U, Abi>&) noexcept;
template<class T, class U, class Abi> see below static_simd_cast(const simd<U, Abi>&) noexcept;

template<class T, class Abi>
    fixed_size_simd<T, simd_size_v<T, Abi>
        to_fixed_size(const simd<T, Abi>&) noexcept;
template<class T, class Abi>
    fixed_size_simd_mask<T, simd_size_v<T, Abi>
        to_fixed_size(const simd_mask<T, Abi>&) noexcept;
template<class T, int N>
    native_simd<T> to_native(const fixed_size_simd<T, N>&) noexcept;
template<class T, int N>
    native_simd_mask<T> to_native(const fixed_size_simd_mask<T, N>&) noexcept;
template<class T, int N>
    simd<T> to_compatible(const fixed_size_simd<T, N>&) noexcept;
template<class T, int N>
    simd_mask<T> to_compatible(const fixed_size_simd_mask<T, N>&) noexcept;

template<size_t... Sizes, class T, class Abi>
    tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd<T, Abi>&) noexcept;
template<size_t... Sizes, class T, class Abi>
    tuple<simd_mask<T, simd_mask_abi::deduce_t<T, Sizes>>...>
        split(const simd_mask<T, Abi>&) noexcept;
template<class V, class Abi>
    array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
        split(const simd<typename V::value_type, Abi>&) noexcept;
template<class V, class Abi>
    array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
        split(const simd_mask<typename V::simd_type::value_type, Abi>&) noexcept;

template<size_t N, class T, class A>
    array<resize_simd_t<simd_size_v<T, A> / N, simd<T, A>>, N>
        split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
    array<resize_simd_t<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
        split_by(const simd_mask<T, A>& x) noexcept;

template<class T, class... Abis>
    simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...)>>

```

```

        concat(const simd<T, Abi>&...) noexcept;
template<class T, class... Abi>
    simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abi> + ...) >>
        concat(const simd_mask<T, Abi>&...) noexcept;

template<class T, class Abi, size_t N>
    resize_simd_t<simd_size_v<T, Abi> * N, simd<T, Abi>>
        concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
    resize_simd_t<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
        concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;

```

// 28.9.9.4, *Rsimd_mask reductions*

```

template<class T, class Abi> bool all_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool any_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool none_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> bool some_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> int popcount(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> int find_first_set(const simd_mask<T, Abi>&);
template<class T, class Abi> int find_last_set(const simd_mask<T, Abi>&);

bool all_of(T) noexcept;
bool any_of(T) noexcept;
bool none_of(T) noexcept;
bool some_of(T) noexcept;
int popcount(T) noexcept;
int find_first_set(T);
int find_last_set(T);

```

// 28.9.5, *Where expression class templates*

```

template<class M, class T> class const_where_expression;
template<class M, class T> class where_expression;

```

// 28.9.9.5, *Where functions*

```

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, simd<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, const simd<T, Abi>&) noexcept;

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, simd_mask<T, Abi>&) noexcept;

```

```

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, const simd_mask<T, Abi>&) noexcept;

template<class T>
    where_expression<bool, T>
        where(see below k, T& d) noexcept;

template<class T>
    const_where_expression<bool, T>
        where(see below k, const T& d) noexcept;

// 28.9.7.4, Rsimd reductions
template<class T, class Abi, class BinaryOperation = plus<>>
    T reduce(const simd<T, Abi>&,
        BinaryOperation = {});

template<class M, class V, class BinaryOperation>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        typename V::value_type identity_element,
        BinaryOperation binary_op);

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        plus<> binary_op = {}) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        multiplies<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_and<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_or<> binary_op) noexcept;

template<class M, class V>
    typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_xor<> binary_op) noexcept;

template<class T, class Abi>
    T hmin(const simd<T, abi>&) noexcept;
template<class M, class V>
    typename V::value_type hmin(const const_where_expression<M, V>&) noexcept;
template<class T, class Abi>
    T hmax(const simd<T, abi>&) noexcept;
template<class M, class V>
    typename V::value_type hmax(const const_where_expression<M, V>&) noexcept;

```

// 28.9.7.6, Algorithms

```

template<class T, class Abi>
    simd<T, Abi>
        min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    simd<T, Abi>
        max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    pair<simd<T, Abi>, simd<T, Abi>>
        minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    simd<T, Abi>
        clamp(const simd<T, Abi>& v,
              const simd<T, Abi>& lo,
              const simd<T, Abi>& hi);
}

```

- 1 The header `<experimental/simd>` defines class templates, tag types, trait types, and function templates for element-wise operations on data-parallel objects.

(5.2.1.3) 28.9.3 simd ABI tags

[simd.abi]

```

namespace simd_abi {
    using scalar = see below;
    template<class T, int N> using fixed_size = see below;
    template<class T> inline constexpr int max_fixed_size = implementation-defined;
    template<class T> using compatible = implementation-defined;
    template<class T> using native = implementation-defined;
    template<int N> using abi_stable = see below;
}

```

- 1 An ABI tag is a type in the `std::experimental::parallelism_v2::simd_abi` namespace that indicates a choice of size and binary representation for objects of data-parallel type. [*Note*: The intent is for the size and binary representation to depend on the target architecture. — *end note*] The ABI tag, together with a given element type implies a number of elements. ABI tag types are used as the second template argument to `simd` and `simd_mask`.
- 2 [*Note*: The ABI tag is orthogonal to selecting the machine instruction set. The selected machine instruction set limits the usable ABI tag types, though (see 28.9.6.1). The ABI tags enable users to safely pass objects of data-parallel type between translation unit boundaries (e.g. function calls or I/O). — *end note*]
- 3 `scalar` is an alias for an unspecified ABI tag that is different from `fixed_size<1>` and `abi_stable<1>`. Use of the `scalar` tag type requires data-parallel types to store a single element (i.e., `simd<T, simd_abi::scalar>::size()` returns 1).
- 4 The value of `max_fixed_size<T>` is at least 32.
- 5 `fixed_size<N>` is an alias for an unspecified ABI tag. `fixed_size` does not introduce a non-deduced context. Use of the `simd_abi::fixed_size<N>` tag type requires data-parallel types to store `N` elements (i.e. ~~`simd<T, simd_abi::fixed_size<N>>::size(simd_size_v<T, simd_abi::fixed_size<N>)`~~ is `N`). `simd<T, fixed_size<N>>` and `simd_mask<T, fixed_size<N>>` with `N > 0` and `N <= max_fixed_size<T>` shall be supported. Additionally, for every supported `simd<T, Abi>` (see 28.9.6.1), where `Abi` is an ABI tag that is not a specialization of `simd_abi::fixed_size`, `N == simd<T, Abi>::size()` shall be supported.

- 6 [*Note:* It is unspecified whether `simd<T, fixed_size<N>>` with `N > max_fixed_size<T>` is supported. The value of `max_fixed_size<T>` can depend on compiler flags and can change between different compiler versions. — *end note*]
- 7 `abi_stable<N>` is an alias for an unspecified ABI tag. `abi_stable` does not introduce a non-deduced context. Use of the `simd_abi::abi_stable<N>` tag type requires data-parallel types to store `N` elements (i.e. `simd_size_v<T, simd_abi::abi_stable<N>` is `N`). `abi_stable_simd<T, N>` and `abi_stable_simd_mask<T, N>` are supported iff `fixed_size_simd<T, N>` and `fixed_size_simd_mask<T, N>` are supported.
- 8 The type of `fixed_size<T, N>` in TU1 differs from the type of `fixed_size<T, N>` in TU2 iff the type of `native<T>` in TU1 differs from the type of `native<T>` in TU2. The type of `abi_stable<N>` is the same in TU1 and TU2.
- 9 [*Note:* An implementation ~~can forego~~should ensure ABI compatibility between differently compiled translation units for `simd` and `simd_mask` specializations using the same `simd_abi::fixed_sizeabi_stable<N>` tag. ~~Otherwise, the efficiency of `simd<T, Abi>` is likely to be better than for `simd<T, fixed_size<simd_size_v<T, Abi>>` (with `Abi` not a specialization of `simd_abi::fixed_size`).~~ — *end note*]
- 10 An implementation may define additional *extended ABI tag* types in the `std::experimental::parallelism_v2::simd_abi` namespace, to support other forms of data-parallel computation.
- 11 ~~`compatible<T>` is an implementation-defined alias for an ABI tag. [*Note:* The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that ensures ABI compatibility between translation units on the target architecture. — *end note*] [*Example:* Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where the `__simd256` type requires an optional ISA extension on said architecture. Also, the target architecture does not support `long double` with either ABI tag. The implementation therefore defines `compatible<T>` as an alias for:~~
- ~~• scalar if `T` is the same type as `long double`, and~~
 - ~~• `__simd128` otherwise.~~
- ~~— *end example*]~~
- 12 `native<T>` is an implementation-defined alias for an ABI tag. [*Note:* The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that is supported on the currently targeted system. ~~For target architectures without ISA extensions, the `native<T>` and `compatible<T>` aliases will likely be the same.~~ For target architectures with ISA extensions, compiler flags may ~~influence~~change the type of the `native<T>` alias ~~while `compatible<T>` will be the same independent of such flags.~~ — *end note*] [*Example:* Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where hardware support for `__simd256` only exists for floating-point types. The implementation therefore defines `native<T>` as an alias for
- `__simd256` if `T` is a floating-point type, and
 - `__simd128` otherwise.
- *end example*]
- ```
template<T, size_t N, class... Abis> struct deduce { using type = see below; };
```
- 13 The member type shall be present if and only if
- `T` is a vectorizable type, and
  - `simd_abi::fixed_size<N>` is supported (see 28.9.3), and
  - every type in the `Abis` pack is an ABI tag.
- 14 Where present, the member typedef `type` shall name an ABI tag type that satisfies
- `simd_size<T, type> == N`, and

- `simd<T, type>` is default constructible (see 28.9.6.1).

If `N` is 1, the member typedef `type` is `simd_abi::scalar`. ~~Otherwise, if there are multiple ABI tag types that satisfy the constraints, the member typedef `type` is implementation-defined.~~ [ *Note: It is expected that extended ABI tags can produce better optimizations and thus are preferred over `simd_abi::fixed_size<N>`.* Implementations can base the choice on `Abis`, but can also ignore the `Abis` arguments. — end note ]

15 The behavior of a program that adds specializations for `deduce` is undefined.

#### (5.2.1.4) 28.9.4 `simd` type traits

[`simd.traits`]

```
template<class T> struct is_abi_tag { see below };
```

- 1 The type `is_abi_tag<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a standard or extended ABI tag, and `false_type` otherwise.
- 2 The behavior of a program that adds specializations for `is_abi_tag` is undefined.

```
template<class T> struct is_simd { see below };
```

- 3 The type `is_simd<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd` class template, and `false_type` otherwise.
- 4 The behavior of a program that adds specializations for `is_simd` is undefined.

```
template<class T> struct is_simd_mask { see below };
```

- 5 The type `is_simd_mask<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd_mask` class template, and `false_type` otherwise.
- 6 The behavior of a program that adds specializations for `is_simd_mask` is undefined.

```
template<class T> struct is_simd_flag_type { see below };
```

- 7 The type `is_simd_flag_type<class T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is one of
  - `element_aligned_tag`, or
  - `vector_aligned_tag`, or
  - `overaligned_tag<N>` with `N > 0` and `N` an integral power of two,
 and `false_type` otherwise.
- 8 The behavior of a program that adds specializations for `is_simd_flag_type` is undefined.

```
template<class T, class Abi = simd_abi::compatible_native<T>> struct simd_size { see below };
```

- 9 `simd_size<T, Abi>` shall have a member `value` if and only if
  - `T` is a vectorizable type, and
  - `is_abi_tag_v<Abi>` is true.

[ *Note*: The rules are different from those in (28.9.6.1). — *end note* ]

- 10 If value is present, the type `simd_size<T, Abi>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` with `N` equal to the number of elements in a `simd<T, Abi>` object. [ *Note*: If `simd<T, Abi>` is not supported for the currently targeted system, `simd_size<T, Abi>::value` produces the value `simd<T, Abi>::size()` would return if it were supported. — *end note* ]
- 11 The behavior of a program that adds specializations for `simd_size` is undefined.

```
template<class T, class U = typename T::value_type> struct memory_alignment { see below };
```

- 12 `memory_alignment<T, U>` shall have a member `value` if and only if
- `is_simd_mask_v<T>` is true and `U` is `bool`, or
  - `is_simd_v<T>` is true and `U` is a vectorizable type.
- 13 If value is present, the type `memory_alignment<T, U>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` for some implementation-defined `N` (see 28.9.6.5 and 28.9.8.4). [ *Note*: `value` identifies the alignment restrictions on pointers used for (converting) loads and stores for the give type `T` on arrays of type `U`. — *end note* ]
- 14 The behavior of a program that adds specializations for `memory_alignment` is undefined.

```
template<class T, class V> struct rebind_simd { using type = see below; };
```

- 15 The member type is present if and only if
- `V` is either `simd<U, Abi0>` or `simd_mask<U, Abi0>`, where `U` and `Abi0` are deduced from `V`, and
  - `T` is a vectorizable type, and
  - `simd_abi::deduce<T, simd_size_v<U, Abi0>, Abi0>` has a member type type.
- 16 Let `Abi1` denote the type `deduce_t<T, simd_size_v<U, Abi0>, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if `V` is `simd<U, Abi0>` or `simd_mask<T, Abi1>` if `V` is `simd_mask<U, Abi0>`.

```
template<int N, class V> struct resize_simd { using type = see below; };
```

- 17 The member type is present if and only if
- `V` is either `simd<T, Abi0>` or `simd_mask<T, Abi0>`, where `T` and `Abi0` are deduced from `V`, and
  - `simd_abi::deduce<T, N, Abi0>` has a member type type.
- 18 Let `Abi1` denote the type `deduce_t<T, N, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if `V` is `simd<T, Abi0>` or `simd_mask<T, Abi1>` if `V` is `simd_mask<T, Abi0>`.

#### (5.2.1.5) 28.9.5 Where expression class templates

[`simd.whereexpr`]

```
template<class M, class T> class const_where_expression {
 const M mask; // exposition only
 T& data; // exposition only

public:
 const_where_expression(const const_where_expression&) = delete;
```

```

const_where_expression& operator=(const const_where_expression&) = delete;

T operator-() const && noexcept;
T operator+() const && noexcept;
T operator~() const && noexcept;

template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const &&
};

template<class M, class T>
class where_expression : public const_where_expression<M, T> {
public:
 template<class U> void operator=(U&& x) && noexcept;
 template<class U> void operator+=(U&& x) && noexcept;
 template<class U> void operator-=(U&& x) && noexcept;
 template<class U> void operator*=(U&& x) && noexcept;
 template<class U> void operator/=(U&& x) && noexcept;
 template<class U> void operator%=(U&& x) && noexcept;
 template<class U> void operator&=(U&& x) && noexcept;
 template<class U> void operator|=(U&& x) && noexcept;
 template<class U> void operator^=(U&& x) && noexcept;
 template<class U> void operator<<=(U&& x) && noexcept;
 template<class U> void operator>>=(U&& x) && noexcept;

 void operator++() && noexcept;
 void operator++(int) && noexcept;
 void operator--() && noexcept;
 void operator--(int) && noexcept;

 template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
};

```

- 1 The class templates `const_where_expression` and `where_expression` abstract the notion of selecting elements of a given object of arithmetic or data-parallel type.
- 2 The first templates argument `M` shall be cv-unqualified `bool` or a cv-unqualified `simd_mask` specialization.
- 3 If `M` is `bool`, `T` shall be a cv-unqualified arithmetic type. Otherwise, `T` shall either be `M` or typename `M::simd_type`.
- 4 In this subclause, if `M` is `bool`, `data[0]` is used interchangeably for `data`, `mask[0]` is used interchangeably for `mask`, and `M::size()` is used interchangeably for 1.
- 5 The *selected indices* signify the integers  $i \in \{j \in \mathbb{N} | j < M::size() \wedge \text{mask}[j]\}$ . The *selected elements* signify the elements `data[i]` for all selected indices  $i$ .
- 6 In this subclause, the type `value_type` is an alias for `T` if `M` is `bool`, or an alias for typename `T::value_type` if `is_simd_mask_v<M>` is true.
- 7 [Note: The `where` functions 28.9.9.5 initialize `mask` with the first argument to `where` and `data` with the second argument to `where`. — end note ]

```

T operator-() const && noexcept;
T operator+() const && noexcept;

```

```
T operator~() const && noexcept;
```

8 **Returns:** A copy of data with the indicated unary operator applied to all selected elements.

```
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const &&;
```

9 **Requires:**

- If M is not bool, the largest selected index is less than the number of values pointed to by mem.
- If the template parameter Flags is vector\_aligned\_tag, mem shall point to storage aligned by memory\_alignment\_v<T, U>.
- If the template parameter Flags is overaligned\_tag<N>, mem shall point to storage aligned by N.
- If the template parameter Flags is element\_aligned\_tag, mem shall point to storage aligned by alignof(U).

10 **Effects:** Copies the selected elements as if `mem[i] = static_cast<U>(data[i])` for all selected indices *i*.

11 **Throws:** Nothing.

12 **Remarks:** This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
  - U is bool and value\_type is bool, or
  - U is a vectorizable type and value\_type is not bool.

```
template<class U> void operator=(U&& x) && noexcept;
```

13 **Effects:** Replaces `data[i]` with `static_cast<T>(std::forward<U>(x))[i]` for all selected indices *i*.

14 **Remarks:** This operator shall not participate in overload resolution unless U is convertible to T.

```
template<class U> void operator+=(U&& x) && noexcept;
template<class U> void operator-=(U&& x) && noexcept;
template<class U> void operator*=(U&& x) && noexcept;
template<class U> void operator/=(U&& x) && noexcept;
template<class U> void operator%=(U&& x) && noexcept;
template<class U> void operator&=(U&& x) && noexcept;
template<class U> void operator|=(U&& x) && noexcept;
template<class U> void operator^=(U&& x) && noexcept;
template<class U> void operator<=(U&& x) && noexcept;
template<class U> void operator>=(U&& x) && noexcept;
```

15 **Effects:** Replaces `data[i]` with `static_cast<T>(data @ std::forward<U>(x))[i]` (where @ denotes the indicated operator) for all selected indices *i*.

16 **Remarks:** Each of these operators shall not participate in overload resolution unless the return type of `data @ std::forward<U>(x)` is convertible to T. It is unspecified whether the binary operator, implied by the compound assignment operator, is executed on all elements or only on the selected elements.

```
void operator++() && noexcept;
void operator++(int) && noexcept;
void operator--() && noexcept;
void operator--(int) && noexcept;
```

17 *Effects:* Applies the indicated operator to the selected elements.

18 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type  $T$ .

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&;
```

19 *Requires:*

- If `is_simd_flag_type_v<U>` is true, for all selected indices  $i$ ,  $i$  shall be less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by  $N$ .
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

20 *Effects:* Replaces the selected elements as if `data[i] = static_cast<value_type>(mem[i])` for all selected indices  $i$ .

21 *Throws:* Nothing.

22 *Remarks:* This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
  - $U$  is `bool` and `value_type` is `bool`, or
  - $U$  is a vectorizable type and `value_type` is not `bool`.

(5.2.1.6) 28.9.6 Class template `simd`

[simd.class]

(5.2.1.6.1) 28.9.6.1 Class template `simd` overview

[simd.overview]

```
template<class T, class Abi> class simd {
public:
 using value_type = T;
 using reference = see below;
 using mask_type = simd_mask<T, Abi>;
 using abi_type = Abi;

 static constexpr size_t size() noexcept;

 constexpr simd() noexcept = default;
```

*// 28.9.6.4, simd constructors*

```
template<class U> constexpr simd(U&& value) noexcept;
template<class U, class UAbi>
 explicit(see below) constexpr simd(const simd<U, simd_abi::fixed_size<size()>UAbi>&) noexcept;
template<class G> explicit constexpr simd(G&& gen) noexcept;
template<class U, class Flags = element_aligned_tag> constexpr simd(const U* mem, Flags f = {});
```

*// 28.9.6.5, simd copy functions*

```
template<class U, class Flags = element_aligned_tag> constexpr void copy_from(const U* mem, Flags f
= {});
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const;
```

*// 28.9.6.6, simd subscript operators*

```
constexpr reference operator[](size_t);
constexpr value_type operator[](size_t) const;
```

*// 28.9.6.7, simd unary operators*

```
constexpr simd& operator++() noexcept;
constexpr simd operator++(int) noexcept;
constexpr simd& operator--() noexcept;
constexpr simd operator--(int) noexcept;
constexpr mask_type operator!() const noexcept;
constexpr simd operator~() const noexcept;
constexpr simd operator+() const noexcept;
constexpr simd operator-() const noexcept;
```

*// 28.9.7.1, simd binary operators*

```
constexpr friend simd operator+(const simd&, const simd&) noexcept;
constexpr friend simd operator-(const simd&, const simd&) noexcept;
constexpr friend simd operator*(const simd&, const simd&) noexcept;
constexpr friend simd operator/(const simd&, const simd&) noexcept;
constexpr friend simd operator%(const simd&, const simd&) noexcept;
constexpr friend simd operator&(const simd&, const simd&) noexcept;
constexpr friend simd operator|(const simd&, const simd&) noexcept;
constexpr friend simd operator^(const simd&, const simd&) noexcept;
constexpr friend simd operator<<(const simd&, const simd&) noexcept;
constexpr friend simd operator>>(const simd&, const simd&) noexcept;
constexpr friend simd operator<<(const simd&, int) noexcept;
constexpr friend simd operator>>(const simd&, int) noexcept;
```

*// 28.9.7.2, simd compound assignment*

```
constexpr friend simd& operator+=(simd&, const simd&) noexcept;
constexpr friend simd& operator-=(simd&, const simd&) noexcept;
constexpr friend simd& operator*=(simd&, const simd&) noexcept;
constexpr friend simd& operator/=(simd&, const simd&) noexcept;
constexpr friend simd& operator%=(simd&, const simd&) noexcept;
```

```
constexpr friend simd& operator&=(simd&, const simd&) noexcept;
constexpr friend simd& operator|=(simd&, const simd&) noexcept;
constexpr friend simd& operator^=(simd&, const simd&) noexcept;
constexpr friend simd& operator<<=(simd&, const simd&) noexcept;
constexpr friend simd& operator>>=(simd&, const simd&) noexcept;
constexpr friend simd& operator<<=(simd&, int) noexcept;
constexpr friend simd& operator>>=(simd&, int) noexcept;
```

// 28.9.7.3, simd *compare operators*

```
constexpr friend mask_type operator==(const simd&, const simd&) noexcept;
constexpr friend mask_type operator!=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator>=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator<=(const simd&, const simd&) noexcept;
constexpr friend mask_type operator>(const simd&, const simd&) noexcept;
constexpr friend mask_type operator<(const simd&, const simd&) noexcept;
```

```
};
```

- 1 The class template `simd` is a data-parallel type. The width of a given `simd` specialization is a constant expression, determined by the template parameters.
- 2 Every specialization of `simd` ~~shall be~~ is a complete type. The specialization `simd<T, Abi>` is supported if `T` is a vectorizable type and

- `Abi` is `simd_abi::scalar`, or
- `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in 28.9.3.

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd<T, Abi>` is supported. [Note: The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd<T, Abi>>`.

[Example: Consider an implementation that defines the extended ABI tags `__simd_x` and `__gpu_y`. When the compiler is invoked to translate to a machine that has support for the `__simd_x` ABI tag for all arithmetic types other than `long double` and no support for the `__gpu_y` ABI tag, then:

- `simd<T, simd_abi::__gpu_y>` is not supported for any `T` and has a deleted constructor.
- `simd<long double, simd_abi::__simd_x>` is not supported and has a deleted constructor.
- `simd<double, simd_abi::__simd_x>` is supported.
- `simd<long double, simd_abi::scalar>` is supported.

— *end example* ]

- 3 Default initialization performs no initialization of the elements; value-initialization initializes each element with `T()`. [Note: Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- 4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd`:

```
explicit constexpr operator implementation-defined() const;
explicit constexpr simd(const implementation-defined& init);
```

[ *Example:* Consider an implementation that supports the type `__vec4f` and the function `__vec4f __vec4f_addsub(__vec4f, __vec4f)` for the currently targeted system. A user may require the use of `_vec4f_addsub` for maximum performance and thus writes:

```
using V = simd<float, simd_abi::__simd128>;
V addsub(V a, V b) {
 return static_cast<V>(_vec4f_addsub(static_cast<__vec4f>(a), static_cast<__vec4f>(b)));
}
```

— *end example* ]

(5.2.1.6.2) 28.9.6.2 `simd` width [`simd.width`]

```
static constexpr size_t size() noexcept;
```

1 *Returns:* The width of `simd<T, Abi>`.

(5.2.1.6.3) 28.9.6.3 Element references [`simd.reference`]

- 1 A reference is an object that refers to an element in a `simd` or `simd_mask` object. `reference::value_type` is the same type as `simd::value_type` or `simd_mask::value_type`, respectively.
- 2 Class reference is for exposition only. An implementation is permitted to provide equivalent functionality without providing a class with this name.

```
class reference // exposition only
{
public:
 reference() = delete;
 reference(const reference&) = delete;

 operator value_type() const noexcept;

 template<class U> reference operator=(U&& x) && noexcept;

 template<class U> reference operator+=(U&& x) && noexcept;
 template<class U> reference operator-=(U&& x) && noexcept;
 template<class U> reference operator*=(U&& x) && noexcept;
 template<class U> reference operator/=(U&& x) && noexcept;
 template<class U> reference operator%=(U&& x) && noexcept;
 template<class U> reference operator|=(U&& x) && noexcept;
 template<class U> reference operator&=(U&& x) && noexcept;
 template<class U> reference operator^=(U&& x) && noexcept;
 template<class U> reference operator<=(U&& x) && noexcept;
 template<class U> reference operator>=(U&& x) && noexcept;

 reference operator++() && noexcept;
```

```

value_type operator++(int) && noexcept;
reference operator--() && noexcept;
value_type operator--(int) && noexcept;

friend void swap(reference&& a, reference&& b) noexcept;
friend void swap(value_type& a, reference&& b) noexcept;
friend void swap(reference&& a, value_type& b) noexcept;
};

```

```
operator value_type() const noexcept;
```

3       **Returns:** The value of the element referred to by \*this.

```
template<class U> reference operator=(U&& x) && noexcept;
```

4       **Effects:** Replaces the referred to element in simd or simd\_mask with static\_cast<value\_type>(std::forward<U>(x)).

5       **Returns:** A copy of \*this.

6       **Remarks:** This function shall not participate in overload resolution unless declval<value\_type>() = std::forward<U>(x) is well-formed.

```

template<class U> reference operator+=(U&& x) && noexcept;
template<class U> reference operator-=(U&& x) && noexcept;
template<class U> reference operator*=(U&& x) && noexcept;
template<class U> reference operator/=(U&& x) && noexcept;
template<class U> reference operator%=(U&& x) && noexcept;
template<class U> reference operator|=(U&& x) && noexcept;
template<class U> reference operator&=(U&& x) && noexcept;
template<class U> reference operator^=(U&& x) && noexcept;
template<class U> reference operator<=(U&& x) && noexcept;
template<class U> reference operator>=(U&& x) && noexcept;

```

7       **Effects:** Applies the indicated compound operator to the referred to element in simd or simd\_mask and std::forward<U>(x).

8       **Returns:** A copy of \*this.

9       **Remarks:** This function shall not participate in overload resolution unless declval<value\_type>() @= std::forward<U>(x) (where @= denotes the indicated compound assignment operator) is well-formed.

```

reference operator++() && noexcept;
reference operator--() && noexcept;

```

10       **Effects:** Applies the indicated operator to the referred to element in simd or simd\_mask.

11       **Returns:** A copy of \*this.

12       **Remarks:** This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type value\_type.

```
value_type operator++(int) && noexcept;
value_type operator--(int) && noexcept;
```

- 13 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.  
 14 *Returns:* A copy of the referred to element before applying the indicated operator.  
 15 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend void swap(reference&& a, reference&& b) noexcept;
friend void swap(value_type& a, reference&& b) noexcept;
friend void swap(reference&& a, value_type& b) noexcept;
```

- 16 *Effects:* Exchanges the values `a` and `b` refer to.

#### (5.2.1.6.4) 28.9.6.4 `simd` constructors

[`simd.ctor`]

```
template<class U> simd(U&&) noexcept;
```

- 1 *Effects:* Constructs an object with each element initialized to the value of the argument after conversion to `value_type`.  
 2 *Remarks:* Let `From` denote the type `remove_cv_t<remove_reference_t<U>>`. This constructor shall not participate in overload resolution unless:
- `From` is a vectorizable type and every possibly value of `From` can be represented with type `value_type`, or
  - `From` is not an arithmetic type and is implicitly convertible to `value_type`, or
  - `From` is `int`, or
  - `From` is `unsigned int` and `value_type` is an unsigned integral type.

```
template<class U, class UAbi>
explicit(see below) constexpr simd(const simd<U, simd_abi::fixed_size<size()->UAbi>& x) noexcept;
```

- 3 *Effects:* Constructs an object where the  $i^{\text{th}}$  element equals `static_cast<T>(x[i])` for all  $i$  in the range of `[0, size())`.  
 4 *Remarks:* This constructor shall not participate in overload resolution unless `simd_size_v<U, UAbi> == size()`.
- ~~• `abi_type` is `simd_abi::fixed_size<size()->`, and~~
  - ~~• every possible value of `U` can be represented with type `value_type`, and~~
  - ~~• if both `U` and `value_type` are integral, the integer conversion rank (??) of `value_type` is greater than the integer conversion rank of `U`.~~
- 5 The constructor is explicit iff
- not every possible value of `U` can be represented with type `value_type`, or
  - if both `U` and `value_type` are integral, the integer conversion rank (??) of `U` is greater than the integer conversion rank of `value_type`.

```
template<class G> simd(G&& gen) noexcept;
```

- 6 *Effects:* Constructs an object where the  $i^{\text{th}}$  element is initialized to `gen(integral_constant<size_t, i>())`.
- 7 *Remarks:* This constructor shall not participate in overload resolution unless `simd(gen(integral_constant<size_t, i>()))` is well-formed for all  $i$  in the range of `[0, size())`.
- 8 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<class U, class Flags = element_aligned_tag> simd(const U* mem, Flags = {});
```

- 9 *Requires:*
- `[mem, mem + size())` is a valid range.
  - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
  - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
  - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.
- 10 *Effects:* Constructs an object where the  $i^{\text{th}}$  element is initialized to `static_cast<T>(mem[i])` for all  $i$  in the range of `[0, size())`.
- 11 *Remarks:* This constructor shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and
  - `U` is a vectorizable type.

(5.2.1.6.5) 28.9.6.5 `simd` copy functions

[`simd.copy`]

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {});
```

- 1 *Requires:*
- `[mem, mem + size())` is a valid range.
  - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
  - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
  - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.
- 2 *Effects:* Replaces the elements of the `simd` object such that the  $i^{\text{th}}$  element is assigned with `static_cast<T>(mem[i])` for all  $i$  in the range of `[0, size())`.
- 3 *Remarks:* This function shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and
  - `U` is a vectorizable type.

```
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const;
```

- 4 **Requires:**
- `[mem, mem + size())` is a valid range.
  - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
  - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
  - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.
- 5 **Effects:** Copies all `simd` elements as if `mem[i] = static_cast<U>(operator[](i))` for all *i* in the range of `[0, size())`.
- 6 **Remarks:** This function shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and
  - `U` is a vectorizable type.

(5.2.1.6.6) 28.9.6.6 `simd` subscript operators

[`simd.subscr`]

```
reference operator[](size_t i);
```

- 1 **Requires:** `i < size()`.
- 2 **Returns:** A reference (see 28.9.6.3) referring to the *i*<sup>th</sup> element.
- 3 **Throws:** Nothing.

```
value_type operator[](size_t i) const;
```

- 4 **Requires:** `i < size()`.
- 5 **Returns:** The value of the *i*<sup>th</sup> element.
- 6 **Throws:** Nothing.

(5.2.1.6.7) 28.9.6.7 `simd` unary operators

[`simd.unary`]

- 1 Effects in this subclause are applied as unary element-wise operations.

```
simd& operator++() noexcept;
```

- 2 **Effects:** Increments every element by one.
- 3 **Returns:** `*this`.

```
simd operator++(int) noexcept;
```

- 4 **Effects:** Increments every element by one.
- 5 **Returns:** A copy of `*this` before incrementing.

```
simd& operator--() noexcept;
```

6       *Effects:* Decrements every element by one.

7       *Returns:* \*this.

```
simd operator--(int) noexcept;
```

8       *Effects:* Decrements every element by one.

9       *Returns:* A copy of \*this before decrementing.

```
mask_type operator!() const noexcept;
```

10       *Returns:* A simd\_mask object with the  $i^{\text{th}}$  element set to !operator[]( $i$ ) for all  $i$  in the range of  $[0, \text{size}())$ .

```
simd operator~() const noexcept;
```

11       *Returns:* A simd object where each bit is the inverse of the corresponding bit in \*this.

12       *Remarks:* This operator shall not participate in overload resolution unless  $\tau$  is an integral type.

```
simd operator+() const noexcept;
```

13       *Returns:* \*this.

```
simd operator-() const noexcept;
```

14       *Returns:* A simd object where the  $i^{\text{th}}$  element is initialized to -operator[]( $i$ ) for all  $i$  in the range of  $[0, \text{size}())$ .

#### (5.2.1.7) 28.9.7 simd non-member operations

[simd.nonmembers]

##### (5.2.1.7.1) 28.9.7.1 simd binary operators

[simd.binary]

```
friend simd operator+(const simd& lhs, const simd& rhs) noexcept;
friend simd operator-(const simd& lhs, const simd& rhs) noexcept;
friend simd operator*(const simd& lhs, const simd& rhs) noexcept;
friend simd operator/(const simd& lhs, const simd& rhs) noexcept;
friend simd operator%(const simd& lhs, const simd& rhs) noexcept;
friend simd operator&(const simd& lhs, const simd& rhs) noexcept;
friend simd operator|(const simd& lhs, const simd& rhs) noexcept;
friend simd operator^(const simd& lhs, const simd& rhs) noexcept;
friend simd operator<<(const simd& lhs, const simd& rhs) noexcept;
friend simd operator>>(const simd& lhs, const simd& rhs) noexcept;
```

1       *Returns:* A simd object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

2       *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type value\_type.

```
friend simd operator<<(const simd& v, int n) noexcept;
friend simd operator>>(const simd& v, int n) noexcept;
```

- 3 **Returns:** A simd object where the  $i^{\text{th}}$  element is initialized to the result of applying the indicated operator to  $v[i]$  and  $n$  for all  $i$  in the range of  $[0, \text{size}())$ .
- 4 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

#### (5.2.1.7.2) 28.9.7.2 simd compound assignment

[simd.cassign]

```
friend simd& operator+=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator-=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator*=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator/=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator%=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator&=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator|=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator^=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator<<=(simd& lhs, const simd& rhs) noexcept;
friend simd& operator>>=(simd& lhs, const simd& rhs) noexcept;
```

- 1 **Effects:** These operators apply the indicated operator to `lhs` and `rhs` as an element-wise operation.
- 2 **Returns:** `lhs`.
- 3 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend simd& operator<<=(simd& lhs, int n) noexcept;
friend simd& operator>>=(simd& lhs, int n) noexcept;
```

- 4 **Effects:** Equivalent to: `return operator@=(lhs, simd(n));`
- 5 **Remarks:** These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

#### (5.2.1.7.3) 28.9.7.3 simd compare operators

[simd.comparison]

```
friend mask_type operator==(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator!=(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator>=(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator<=(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator>(const simd& lhs, const simd& rhs) noexcept;
friend mask_type operator<(const simd& lhs, const simd& rhs) noexcept;
```

- 1 **Returns:** A `simd_mask` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

(5.2.1.7.4) 28.9.7.4 **R**simd reductions

[simd.reductions]

1 In this subclause, BinaryOperation shall be a binary element-wise operation.

```
template<class T, class Abi, class BinaryOperation = plus<>>
 T reduce(const simd<T, Abi>& x, BinaryOperation binary_op = {});
```

2 *Requires:* binary\_op shall be callable with two arguments of type T returning T, or callable with two arguments of type simd<T, A1> returning simd<T, A1> for every A1 that is an ABI tag type.

3 *Returns:* GENERALIZED\_SUM(binary\_op, x.data[i], ...) for all i in the range of [0, size())(??).

4 *Throws:* Any exception thrown from binary\_op.

```
template<class M, class V, class BinaryOperation>
 typename V::value_type reduce(const const_where_expression<M, V>& x,
 typename V::value_type identity_element,
 BinaryOperation binary_op = {});
```

5 *Requires:* binary\_op shall be callable with two arguments of type T returning T, or callable with two arguments of type simd<T, A1> returning simd<T, A1> for every A1 that is an ABI tag type. The results of binary\_op(identity\_element, x) and binary\_op(x, identity\_element) shall be equal to x for all finite values x representable by V::value\_type.

6 *Returns:* If none\_of(x.mask), returns identity\_element. Otherwise, returns GENERALIZED\_SUM(binary\_op, x.data[i], ...) for all selected indices i.

7 *Throws:* Any exception thrown from binary\_op.

```
template<class M, class V>
 typename V::value_type reduce(const const_where_expression<M, V>& x, plus<> binary_op) noexcept;
```

8 *Returns:* If none\_of(x.mask), returns 0. Otherwise, returns GENERALIZED\_SUM(binary\_op, x.data[i], ...) for all selected indices i.

```
template<class M, class V>
 typename V::value_type reduce(const const_where_expression<M, V>& x, multiplies<> binary_op) noexcept;
```

9 *Returns:* If none\_of(x.mask), returns 1. Otherwise, returns GENERALIZED\_SUM(binary\_op, x.data[i], ...) for all selected indices i.

```
template<class M, class V>
 typename V::value_type reduce(const const_where_expression<M, V>& x, bit_and<> binary_op) noexcept;
```

10 *Requires:* is\_integral\_v<V::value\_type> is true.

11 *Returns:* If none\_of(x.mask), returns ~V::value\_type(). Otherwise, returns GENERALIZED\_SUM(binary\_op, x.data[i], ...) for all selected indices i.

```
template<class M, class V>
 typename V::value_type reduce(const const_where_expression<M, V>& x, bit_or<> binary_op) noexcept;
template<class M, class V>
 typename V::value_type reduce(const const_where_expression<M, V>& x, bit_xor<> binary_op) noexcept;
```

- 12       **Requires:** `is_integral_v<V::value_type>` is true.
- 13       **Returns:** If `none_of(x.mask)`, returns 0. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices *i*.

```
template<class T, class Abi> T hmin(const simd<T, Abi>& x) noexcept;
```

- 14       **Returns:** The value of an element `x[j]` for which `x[j] <= x[i]` for all *i* in the range of `[0, size())`.

```
template<class M, class V> typename V::value_type hmin(const const_where_expression<M, V>& x) noexcept;
```

- 15       **Returns:** If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::max()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] <= x.data[i]` for all selected indices *i*.

```
template<class T, class Abi> T hmax(const simd<T, Abi>& x) noexcept;
```

- 16       **Returns:** The value of an element `x[j]` for which `x[j] >= x[i]` for all *i* in the range of `[0, size())`.

```
template<class M, class V> typename V::value_type hmax(const const_where_expression<M, V>& x) noexcept;
```

- 17       **Returns:** If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::lowest()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] >= x.data[i]` for all selected indices *i*.

#### (5.2.1.7.5) 28.9.7.5 Casts [simd.casts]

```
template<class T, class U, class Abi> see below simd_cast(const simd<U, Abi>& x) noexcept;
```

- 1       Let `To` denote `T::value_type` if `is_simd_v<T>` is true, or `T` otherwise.
- 2       **Returns:** A `simd` object with the *i*<sup>th</sup> element initialized to `static_cast<To>(x[i])` for all *i* in the range of `[0, size())`.
- 3       **Remarks:** The function shall not participate in overload resolution unless
- every possible value of type `U` can be represented with type `To`, and
  - either
    - `is_simd_v<T>` is false, or
    - `T::size() == simd<U, Abi>::size()` is true.
- 4       The return type is
- `T` if `is_simd_v<T>` is true;
  - otherwise, `simd<T, Abi>` if `U` is the same type as `T`;

- otherwise, `simd<T, simd_abi::fixed_size<simd<U, Abi>::size()>>`

```
template<class T, class U, class Abi> see below static_simd_cast(const simd<U, Abi>& x) noexcept;
```

- 5       **Let** `To` denote `T::value_type` if `is_simd_v<T>` is true or `T` otherwise.
- 6       **Returns:** A `simd` object with the  $i^{\text{th}}$  element initialized to `static_cast<To>(x[i])` for all  $i$  in the range of `[0, size())`.
- 7       **Remarks:** The function shall not participate in overload resolution unless either
- `is_simd_v<T>` is false, or
  - `T::size() == simd<U, Abi>::size()` is true.
- 8       **The return type is**
- `T` if `is_simd_v<T>` is true;
  - otherwise, `simd<T, Abi>` if either `U` is the same type as `T` or `make_signed_t<U>` is the same type as `make_signed_t<T>`;
  - otherwise, `simd<T, simd_abi::fixed_size<simd<U, Abi>::size()>>`.

```
template<class T, class U, class Abi> see below simd_cast(const simd_mask<U, Abi>& x) noexcept;
template<class T, class U, class Abi> see below static_simd_cast(const simd_mask<U, Abi>& x) noexcept;
```

- 9       **Returns:** A `simd_mask` object with the  $i^{\text{th}}$  element initialized to `x[i]` for all  $i$  in the range of `[0, size())`.
- 10      **Remarks:** The functions shall not participate in overload resolution unless either
- `is_simd_mask_v<T>` is false, or
  - `T::size() == simd_size_v<U, Abi>` is true.
- 11      **The return type is**
- `T` if `is_simd_mask_v<T>` is true;
  - otherwise, `simd_mask<T, Abi>` if either `U` is the same type as `T` or `make_signed_t<U>` is the same type as `make_signed_t<T>`;
  - otherwise, `simd_mask<T, simd_abi::fixed_size<simd_size_v<U, Abi>>>`

```
template<class T, class Abi>
 fixed_size_simd<T, simd_size_v<T, Abi>> to_fixed_size(const simd<T, Abi>& x) noexcept;
template<class T, class Abi>
 fixed_size_simd_mask<T, simd_size_v<T, Abi>> to_fixed_size(const simd_mask<T, Abi>& x) noexcept;
```

- 12      **Returns:** A data-parallel object with the  $i^{\text{th}}$  element initialized to `x[i]` for all  $i$  in the range of `[0, size())`.

```
template<class T, int N> native_simd<T> to_native(const fixed_size_simd<T, N>& x) noexcept;
template<class T, int N> native_simd_mask<T> to_native(const fixed_size_simd_mask<T, N>& x) noexcept;
```

- 13      **Returns:** A data-parallel object with the  $i^{\text{th}}$  element initialized to `x[i]` for all  $i$  in the range of `[0, size())`.
- 14      **Remarks:** These functions shall not participate in overload resolution unless `simd_size_v<T, simd_abi::native<T>> == N` is true.

```
template<class T, int N> simd<T> to_compatible(const fixed_size_simd<T, N>& x) noexcept;
template<class T, int N> simd_mask<T> to_compatible(const fixed_size_simd_mask<T, N>& x) noexcept;
```

15 **Returns:** A data-parallel object with the  $i^{\text{th}}$  element initialized to  $x[i]$  for all  $i$  in the range of  $[0, \text{size}())$ .

16 **Remarks:** These functions shall not participate in overload resolution unless `simd_size_v<T, simd_abi::compatible<T>> == N` is true.

```
template<size_t... Sizes, class T, class Abi>
tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
split(const simd<T, Abi>& x) noexcept;
template<size_t... Sizes, class T, class Abi>
tuple<simd_mask<T, simd_abi::deduce_t<T, Sizes>>...>
split(const simd_mask<T, Abi>& x) noexcept;
```

17 **Returns:** A tuple of data-parallel objects with the  $i^{\text{th}}$  `simd/simd_mask` element of the  $j^{\text{th}}$  tuple element initialized to the value of the element  $x$  with index  $i + \text{sum of the first } j \text{ values in the Sizes pack}$ .

18 **Remarks:** These functions shall not participate in overload resolution unless the sum of all values in the Sizes pack is equal to `simd_size_v<T, Abi>`.

```
template<class V, class Abi>
array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
split(const simd<typename V::value_type, Abi>& x) noexcept;
template<class V, class Abi>
array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
split(const simd_mask<typename V::simd_type::value_type, Abi>& x) noexcept;
```

19 **Returns:** An array of data-parallel objects with the  $i^{\text{th}}$  `simd/simd_mask` element of the  $j^{\text{th}}$  array element initialized to the value of the element in  $x$  with index  $i + j * V::size()$ .

20 **Remarks:** These functions shall not participate in overload resolution unless either:

- `is_simd_v<V>` is true and `simd_size_v<typename V::value_type, Abi>` is an integral multiple of `V::size()`, or
- `is_simd_mask_v<V>` is true and `simd_size_v<typename V::simd_type::value_type, Abi>` is an integral multiple of `V::size()`.

```
template<size_t N, class T, class A>
array<resize_simd<simd_size_v<T, A> / N, simd<T, A>>, N>
split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
array<resize_simd<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
split_by(const simd_mask<T, A>& x) noexcept;
```

21 **Returns:** An array `arr`, where `arr[i][j]` is initialized by `x[i * (simd_size_v<T, A> / N) + j]`.

22 **Remarks:** The functions shall not participate in overload resolution unless `simd_size_v<T, A>` is an integral multiple of `N`.

```
template<class T, class... Abis>
 simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
 const simd<T, Abis>&... xs) noexcept;
template<class T, class... Abis>
 simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
 const simd_mask<T, Abis>&... xs) noexcept;
```

- 23 **Returns:** A data-parallel object initialized with the concatenated values in the `xs` pack of data-parallel objects: The  $i^{\text{th}}$  `simd/simd_mask` element of the  $j^{\text{th}}$  parameter in the `xs` pack is copied to the return value's element with index  $i$  + the sum of the width of the first  $j$  parameters in the `xs` pack.

```
template<class T, class Abi, size_t N>
 resize_simd<simd_size_v<T, Abi> * N, simd<T, Abi>>
 concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
 resize_simd<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
 concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;
```

- 24 **Returns:** A data-parallel object, the  $i^{\text{th}}$  element of which is initialized by `arr[i / simd_size_v<T, Abi>][i % simd_size_v<T, Abi>]`.

#### (5.2.1.7.6) 28.9.7.6 Algorithms

[simd.alg]

```
template<class T, class Abi> simd<T, Abi> min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 1 **Returns:** The result of the element-wise application of `std::min(a[i], b[i])` for all  $i$  in the range of  $[0, \text{size}())$ .

```
template<class T, class Abi> simd<T, Abi> max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 2 **Returns:** The result of the element-wise application of `std::max(a[i], b[i])` for all  $i$  in the range of  $[0, \text{size}())$ .

```
template<class T, class Abi>
 pair<simd<T, Abi>, simd<T, Abi>> minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 3 **Returns:** A pair initialized with
- the result of element-wise application of `std::min(a[i], b[i])` for all  $i$  in the range of  $[0, \text{size}())$  in the first member, and
  - the result of element-wise application of `std::max(a[i], b[i])` for all  $i$  in the range of  $[0, \text{size}())$  in the second member.

```
template<class T, class Abi> simd<T, Abi>
 clamp(const simd<T, Abi>& v, const simd<T, Abi>& lo, const simd<T, Abi>& hi);
```

- 4 *Requires:* No element in `lo` shall be greater than the corresponding element in `hi`.
- 5 *Returns:* The result of element-wise application of `std::clamp(v[i], lo[i], hi[i])` for all `i` in the range of `[0, size())`.

(5.2.1.7.7) 28.9.7.7 ~~M~~simd math library [simd.math]

- 1 For each set of overloaded functions within `<cmath>`, there shall be additional overloads sufficient to ensure that if any argument corresponding to a double parameter has type `simd<T, Abi>`, where `is_floating_point_v<T>` is true, then:
- All arguments corresponding to double parameters shall be convertible to `simd<T, Abi>`.
  - All arguments corresponding to double\* parameters shall be of type `simd<T, Abi>*`.
  - All arguments corresponding to parameters of integral type `U` shall be convertible to `fixed_size_simd<U, simd_size_v<T, Abi>>`.
  - All arguments corresponding to `U*`, where `U` is integral, shall be of type `fixed_size_simd<U, simd_size_v<T, Abi>*>`.
  - If the corresponding return type is `double`, the return type of the additional overloads is `simd<T, Abi>`. Otherwise, if the corresponding return type is `bool`, the return type of the additional overload is `simd_mask<T, Abi>`. Otherwise, the return type is `fixed_size_simd<R, simd_size_v<T, Abi>>`, with `R` denoting the corresponding return type.

It is unspecified whether a call to these overloads with arguments that are all convertible to `simd<T, Abi>` but are not of type `simd<T, Abi>` is well-formed.

- 2 Each function overload produced by the above rules applies the indicated `<cmath>` function element-wise. For the mathematical functions, the results per element only need to be approximately equal to the application of the function which is overloaded for the element type.
- 3 The ~~behavior is undefined~~result is unspecified if a domain, pole, or range error occurs when the input argument(s) are applied to the indicated `<cmath>` function. [ Note: Implementations are encouraged to follow the C specification (especially Annex F). — end note ]
- 4 TODO: Allow `abs(simd<signed-integral>)`.
- 5 If `abs` is called with an argument of type `simd<X, Abi>` for which `is_unsigned_v<X>` is true, the program is ill-formed.

(5.2.1.8) 28.9.8 Class template `simd_mask` [simd.mask.class]

(5.2.1.8.1) 28.9.8.1 Class template `simd_mask` overview [simd.mask.overview]

```
template<class T, class Abi> class simd_mask {
public:
 using value_type = bool;
 using reference = see below;
 using simd_type = simd<T, Abi>;
 using abi_type = Abi;

 static constexpr size_t size() noexcept;
```

```

simd_mask() noexcept = default;

// 28.9.8.3, Csimd_mask constructors
explicit simd_mask(value_type) noexcept;
template<class U>
 simd_mask(const simd_mask<U, simd_abi::fixed_size<size()>>&) noexcept;
template<class G> explicit simd_mask(G&& gen) noexcept;
template<class Flags = element_aligned_tag> simd_mask(const value_type* mem, Flags = {});

// 28.9.8.4, Csimd_mask copy functions
template<class Flags = element_aligned_tag> void copy_from(const value_type* mem, Flags = {});
template<class Flags = element_aligned_tag> void copy_to(value_type* mem, Flags = {});

// 28.9.8.5, Ssimd_mask subscript operators
reference operator[](size_t);
value_type operator[](size_t) const;

// 28.9.8.6, Usimd_mask unary operators
simd_mask operator!() const noexcept;

// 28.9.9.1, Bsimd_mask binary operators
friend simd_mask operator&&(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator|(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator&(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator|(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator^(const simd_mask&, const simd_mask&) noexcept;

// 28.9.9.2, Csimd_mask compound assignment
friend simd_mask& operator&=(simd_mask&, const simd_mask&) noexcept;
friend simd_mask& operator|=(simd_mask&, const simd_mask&) noexcept;
friend simd_mask& operator^=(simd_mask&, const simd_mask&) noexcept;

// 28.9.9.3, Csimd_mask comparisons
friend simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
};

```

- 1 The class template `simd_mask` is a data-parallel type with the element type `bool`. The width of a given `simd_mask` specialization is a constant expression, determined by the template parameters. Specifically, `simd_mask<T, Abi>::size() == simd<T, Abi>::size()`.
- 2 Every specialization of `simd_mask` shall be a complete type. The specialization `simd_mask<T, Abi>` is supported if `T` is a vectorizable type and
  - `Abi` is `simd_abi::scalar`, or
  - `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in (28.9.3).

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd_mask<T, Abi>` is supported. [ *Note:* The intent is for implementations to decide on the basis of the currently targeted system. — *end note* ]

If `simd_mask<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd_mask<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd_mask<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd_mask<T, Abi>>`.

3 Default initialization performs no initialization of the elements; value-initialization initializes each element with false. [ *Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note* ]

4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd_mask`:

```
explicit operator implementation-defined() const;
explicit simd_mask(const implementation-defined& init) const;
```

5 The member type reference has the same interface as `simd<T, Abi>::reference`, except its `value_type` is `bool`. (28.9.6.3)

(5.2.1.8.2) 28.9.8.2 `simd_mask` width [`simd.mask.width`]

```
static constexpr size_t size() noexcept;
```

1 **Returns:** The width of `simd<T, Abi>`.

(5.2.1.8.3) 28.9.8.3 `simd_mask` constructors [`simd.mask.ctor`]

```
explicit simd_mask(value_type x) noexcept;
```

1 **Effects:** Constructs an object with each element initialized to `x`.

```
template<class U> simd_mask(const simd_mask<U, simd_abi::fixed_size<size()>>& x) noexcept;
```

2 **Effects:** Constructs an object of type `simd_mask` where the  $i^{\text{th}}$  element equals `x[i]` for all  $i$  in the range of `[0, size())`.

3 **Remarks:** This constructor shall not participate in overload resolution unless `abi_type` is `simd_abi::fixed_size<size()>`.

```
template<class G> simd_mask(G&& gen) noexcept;
```

4 **Effects:** Constructs an object where the  $i^{\text{th}}$  element is initialized to `gen(integral_constant<size_t, i>())`.

5 **Remarks:** This constructor shall not participate in overload resolution unless `static_cast<bool>(gen(integral_constant<size_t, i>()))` is well-formed for all  $i$  in the range of `[0, size())`.

6 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<class Flags = element_aligned_tag> simd_mask(const value_type* mem, Flags = {});
```

- 7       **Requires:**
- [mem, mem + size()) is a valid range.
  - If the template parameter Flags is vector\_aligned\_tag, mem shall point to storage aligned by memory\_alignment\_v<simd\_mask>.
  - If the template parameter Flags is overaligned\_tag<N>, mem shall point to storage aligned by N.
  - If the template parameter Flags is element\_aligned\_tag, mem shall point to storage aligned by alignof(value\_type).
- 8       **Effects:** Constructs an object where the  $i^{\text{th}}$  element is initialized to mem[ $i$ ] for all  $i$  in the range of [0, size()).
- 9       **Throws:** Nothing.
- 10      **Remarks:** This constructor shall not participate in overload resolution unless is\_simd\_flag\_type\_v<Flags> is true.

(5.2.1.8.4)   28.9.8.4 **✗** simd\_mask copy functions

[simd.mask.copy]

```
template<class Flags = element_aligned_tag> void copy_from(const value_type* mem, Flags = {});
```

- 1       **Requires:**
- [mem, mem + size()) is a valid range.
  - If the template parameter Flags is vector\_aligned\_tag, mem shall point to storage aligned by memory\_alignment\_v<simd\_mask>.
  - If the template parameter Flags is overaligned\_tag<N>, mem shall point to storage aligned by N.
  - If the template parameter Flags is element\_aligned\_tag, mem shall point to storage aligned by alignof(value\_type).
- 2       **Effects:** Replaces the elements of the simd\_mask object such that the  $i^{\text{th}}$  element is replaced with mem[ $i$ ] for all  $i$  in the range of [0, size()).
- 3       **Throws:** Nothing.
- 4       **Remarks:** This function shall not participate in overload resolution unless is\_simd\_flag\_type\_v<Flags> is true.

```
template<class Flags = element_aligned_tag> void copy_to(value_type* mem, Flags = {});
```

- 5       **Requires:**
- [mem, mem + size()) is a valid range.
  - If the template parameter Flags is vector\_aligned\_tag, mem shall point to storage aligned by memory\_alignment\_v<simd\_mask>.
  - If the template parameter Flags is overaligned\_tag<N>, mem shall point to storage aligned by N.
  - If the template parameter Flags is element\_aligned\_tag, mem shall point to storage aligned by alignof(value\_type).

- 6 *Effects:* Copies all `simd_mask` elements as if `mem[i] = operator[](i)` for all  $i$  in the range of  $[0, \text{size}())$ .  
 7 *Throws:* Nothing.  
 8 *Remarks:* This function shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is `true`.

(5.2.1.8.5) 28.9.8.5 **S**simd\_mask subscript operators [simd.mask.subscr]

```
reference operator[](size_t i);
```

- 1 *Requires:*  $i < \text{size}()$ .  
 2 *Returns:* A reference (see 28.9.6.3) referring to the  $i^{\text{th}}$  element.  
 3 *Throws:* Nothing.

```
value_type operator[](size_t i) const;
```

- 4 *Requires:*  $i < \text{size}()$ .  
 5 *Returns:* The value of the  $i^{\text{th}}$  element.  
 6 *Throws:* Nothing.

(5.2.1.8.6) 28.9.8.6 **U**simd\_mask unary operators [simd.mask.unary]

```
simd_mask operator!() const noexcept;
```

- 1 *Returns:* The result of the element-wise application of `operator!`.

(5.2.1.9) 28.9.9 Non-member operations [simd.mask.nonmembers]

(5.2.1.9.1) 28.9.9.1 **B**simd\_mask binary operators [simd.mask.binary]

```
friend simd_mask operator&&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator||(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator|(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask operator^(const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

- 1 *Returns:* A `simd_mask` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

(5.2.1.9.2) 28.9.9.2 **C**simd\_mask compound assignment [simd.mask.cassign]

```
friend simd_mask& operator&=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask& operator|=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend simd_mask& operator^=(simd_mask& lhs, const simd_mask& rhs) noexcept;
```

- 1 *Effects:* These operators apply the indicated operator to lhs and rhs as a binary element-wise operation.  
 2 *Returns:* lhs.

(5.2.1.9.3) 28.9.9.3 ~~C~~simd\_mask comparisons [simd.mask.comparison]

```
friend simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
```

- 1 *Returns:* A simd\_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(5.2.1.9.4) 28.9.9.4 ~~R~~simd\_mask reductions [simd.mask.reductions]

```
template<class T, class Abi> bool all_of(const simd_mask<T, Abi>& k) noexcept;
```

- 1 *Returns:* true if all boolean elements in k are true, false otherwise.

```
template<class T, class Abi> bool any_of(const simd_mask<T, Abi>& k) noexcept;
```

- 2 *Returns:* true if at least one boolean element in k is true, false otherwise.

```
template<class T, class Abi> bool none_of(const simd_mask<T, Abi>& k) noexcept;
```

- 3 *Returns:* true if none of the one boolean elements in k is true, false otherwise.

```
template<class T, class Abi> bool some_of(const simd_mask<T, Abi>& k) noexcept;
```

- 4 *Returns:* true if at least one of the one boolean elements in k is true and at least one of the boolean elements in k is false, false otherwise.

```
template<class T, class Abi> int popcount(const simd_mask<T, Abi>& k) noexcept;
```

- 5 *Returns:* The number of boolean elements in k that are true.

```
template<class T, class Abi> int find_first_set(const simd_mask<T, Abi>& k);
```

- 6 *Requires:* any\_of(k) returns true.

- 7 *Returns:* The lowest element index  $i$  where  $k[i]$  is true.

- 8 *Throws:* Nothing.

```
template<class T, class Abi> int find_last_set(const simd_mask<T, Abi>& k);
```

- 9 *Requires:* any\_of(k) returns true.

- 10 *Returns:* The greatest element index  $i$  where  $k[i]$  is true.

- 11 *Throws:* Nothing.

```

bool all_of(T) noexcept;
bool any_of(T) noexcept;
bool none_of(T) noexcept;
bool some_of(T) noexcept;
int popcount(T) noexcept;

```

12       *Returns:* `all_of` and `any_of` return their arguments; `none_of` returns the negation of its argument; `some_of` returns false; `popcount` returns the integral representation of its argument.

13       *Remarks:* The parameter type  $T$  is an unspecified type that is only constructible via implicit conversion from `bool`.

```

int find_first_set(T);
int find_last_set(T);

```

14       *Requires:* The value of the argument is `true`.

15       *Returns:* 0.

16       *Throws:* Nothing.

17       *Remarks:* The parameter type  $T$  is an unspecified type that is only constructible via implicit conversion from `bool`.

(5.2.1.9.5)   28.9.9.5 `where` functions

[`simd.mask.where`]

```

template<class T, class Abi>
 where_expression<simd_mask<T, Abi>, simd<T, Abi>>
 where(const typename simd<T, Abi>::mask_type& k, simd<T, Abi>& v) noexcept;
template<class T, class Abi>
 const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
 where(const typename simd<T, Abi>::mask_type& k, const simd<T, Abi>& v) noexcept;
template<class T, class Abi>
 where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
 where(const type_identity_t<simd_mask<T, Abi>>& k, simd_mask<T, Abi>& v) noexcept;
template<class T, class Abi>
 const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
 where(const type_identity_t<simd_mask<T, Abi>>& k, const simd_mask<T, Abi>& v) noexcept;

```

1       *Returns:* An object (28.9.5) with `mask` and `data` initialized with `k` and `v` respectively.

```

template<class T>
 where_expression<bool T>
 where(see below k, T& v) noexcept;
template<class T>
 const_where_expression<bool, T>
 where(see below k, const T& v) noexcept;

```

2       *Remarks:* The functions shall not participate in overload resolution unless

- `T` is neither a `simd` nor a `simd_mask` specialization, and
- the first argument is of type `bool`.

3 **Returns:** An object (28.9.5) with `mask` and `data` initialized with `k` and `v` respectively.

---

## A

## BIBLIOGRAPHY

- [P0214R9] Matthias Kretz. *P0214R9: Data-Parallel Vector Types & Operations*. ISO/IEC C++ Standards Committee Paper. 2018. URL: <https://wg21.link/p0214r9>.
- [P0350R0] Matthias Kretz. *P0350R0: Integrating datapar with parallel algorithms and executors*. ISO/IEC C++ Standards Committee Paper. 2016. URL: <https://wg21.link/p0350r0>.
- [P1915R0] Matthias Kretz. *P1915R0: Expected Feedback from `simd` in the Parallelism TS 2*. ISO/IEC C++ Standards Committee Paper. 2019. URL: <https://wg21.link/p1915r0>.
- [P0918R2] Tim Shen. *P0918R2: More `simd<>` Operations*. ISO/IEC C++ Standards Committee Paper. 2018. URL: <https://wg21.link/p0918r2>.
- [P2638R0] Daniel Towner et al. *P2638R0: Intel's response to P1915R0 for `std::simd` parallelism in TS 2*. ISO/IEC C++ Standards Committee Paper. 2022. URL: <https://wg21.link/p2638r0>.