

Document Number: P1928R3
Date: 2023-02-03
Reply-to: Matthias Kretz <m.kretz@gsi.de>
Audience: LEWG
Target: C++26

MERGE DATA-PARALLEL TYPES FROM THE PARALLELISM TS 2

ABSTRACT

After the Parallelism TS 2 was published in 2018, data-parallel types (`simd<T>`) have been implemented and used. Now there is sufficient feedback to improve and merge Section 9 of the Parallelism TS 2 into the IS working draft.

CONTENTS

1	CHANGELOG	1
1.1	CHANGES FROM REVISION 0	1
1.2	CHANGES FROM REVISION 1	2
1.3	CHANGES FROM REVISION 2	2
2	STRAW POLLS	3
2.1	SG1 AT KONA 2022	3
3	INTRODUCTION	3
3.1	RELATED PAPERS	4

4	CHANGES AFTER TS FEEDBACK	5
4.1	IMPROVE ABI TAGS	5
4.2	SIMPLIFY/GENERALIZE CASTS	7
4.3	ADD <code>simd_mask</code> GENERATOR CONSTRUCTOR	11
4.4	DEFAULT LOAD/STORE FLAGS TO <code>element_aligned</code>	11
4.5	CONTIGUOUS ITERATORS FOR LOADS AND STORES	12
4.6	<code>constexpr</code> EVERYTHING	12
4.7	SPECIFY <code>simd::size</code> AS <code>integral_constant</code>	13
4.8	REPLACE <code>where</code> FACILITIES	13
4.9	MAKE USE OF <code>int</code> AND <code>size_t</code> CONSISTENT	16
4.10	CLEAN UP MATH FUNCTION OVERLOADS	18
5	OPEN QUESTIONS	18
5.1	ALTERNATIVES TO <code>hmin</code> AND <code>hmax</code>	18
5.2	MAKE <code>simd_mask</code> REDUCTIONS CONSISTENT WITH <code><bit></code>	19
5.3	ARGUMENT ORDER AND NAMING OF MASKED OVERLOADS	22
5.4	CONVERTING LOADS & STORES CONSISTENCY	23
5.5	INTEGRATION WITH RANGES	25
5.6	CORRECT PLACE FOR <code>simd</code> IN THE IS?	25
5.7	<code>element_reference</code> IS OVERSPECIFIED	25
6	WORDING	25
6.1	ADD SECTION 9 OF N4808 WITH MODIFICATIONS	26
6.2	DIFF AGAINST PARALLELISM TS 2 (N4808)	56
A	BIBLIOGRAPHY	89

1

CHANGELOG

1.1

CHANGES FROM REVISION 0

Previous revision: P1928R0

- Target C++26, addressing SG1 and LEWG.
- Call for a merge of the (improved & adjusted) TS specification to the IS.
- Discuss changes to the ABI tags as consequence of TS experience; calls for polls to change the status quo.
- Add template parameter `T` to `simd_abi::fixed_size`.
- Remove `simd_abi::compatible`.
- Add (but ask for removal) `simd_abi::abi_stable`.
- Mention TS implementation in GCC releases.
- Add more references to related papers.
- Adjust the clause number for [numbers] to latest draft.
- Add open question: what is the correct clause for [simd]?
- Add open question: integration with ranges.
- Add `simd_mask` generator constructor.
- Consistently add `simd` and `simd_mask` to headings.
- Remove `experimental` and `parallelism_v2` namespaces.
- Present the wording twice: with and without diff against N4808 (Parallelism TS 2).
- Default load/store flags to `element_aligned`.
- Generalize casts: conditionally `explicit` converting constructors.
- Remove named cast functions.

1.2

CHANGES FROM REVISION 1

Previous revision: P1928R1

- Add floating-point conversion rank to condition of `explicit` for converting constructors.
- Call out different or equal semantics of the new ABI tags.
- Update introductory paragraph of Section 4; R1 incorrectly kept the text from R0.
- Define `simd::size` as a `constexpr` static data-member of type `integral_constant<size_t, N>`. This simplifies passing the size via function arguments and still be useable as a constant expression in the function body.
- Document addition of `constexpr` to the API.
- Add `constexpr` to the wording.
- Removed ABI tag for passing `simd` over ABI boundaries.
- Apply cast interface changes to the wording.
- Explain the plan: what this paper wants to merge vs. subsequent papers for additional features. With an aim of minimal removal/changes of wording after this paper.
- Document rationale and design intent for `where` replacement.

1.3

CHANGES FROM REVISION 2

Previous revision: P1928R2

- Propose alternative to `hmin` and `hmax`.
- Discuss `simd_mask` reductions wrt. consistency with `<bit>`. Propose better names to avoid ambiguity.
- Remove `some_of`.
- Add unary `~` to `simd_mask`.
- Discuss and ask for confirmation of masked “overloads” names and argument order.
- Resolve inconsistencies wrt. `int` and `size_t`: Change `fixed_size` and `resize_simd` NTPs from `int` to `size_t` (for consistency).
- Discuss conversions on loads and stores. (Section 5.4)

- Point to [P2509R0] as related paper.
 - Generalize load and store from pointer to `contiguous_iterator`. (Section 4.5)
 - Moved “`element_reference` is overspecified” to “Open questions”.
- ✎ Replace where wording.
 - ✎ Apply the new library specification style from P0788R3.
 - ✎ Add `numeric_limits` / numeric traits specializations since behavior of e.g. `simd<float>` and `float` may differ for reasonable implementations.
 - ✎ Consider adding a note that recommends implementations to let `simd` primary operations behave like operations of arithmetic types to never be function calls. (cf. GCC PR108030)

2

STRAW POLLS

2.1

SG1 AT KONA 2022

Poll: After significant experience with the TS, we recommend that the next version (the TS version with improvements) of `std::simd` target the IS (C++26)

SF	F	N	A	SA
10	8	0	0	0

Poll: We like all of the recommended changes to `std::simd` proposed in p1928r1 (Includes making all of `std::simd constexpr`, and dropping an ABI stable type)

→ unanimous consent

Poll: Future papers and future revisions of existing papers that target `std::simd` should go directly to LEWG. (We do not believe there are SG1 issues with `std::simd` today.)

SF	F	N	A	SA
9	8	0	0	0

3

INTRODUCTION

[P0214R9] introduced `simd<T>` and related types and functions into the Parallelism TS 2 Section 9. The TS was published in 2018. An incomplete and non-conforming (because P0214 evolved) implementation existed for the whole time P0214 progressed through the committee. Shortly after

the GCC 9 release, a complete implementation of Section 9 of the TS was made available. Since GCC 11 a complete `simd` implementation of the TS is part of its standard library.

In the meantime the TS feedback progressed to a point where a merge should happen ASAP. This paper proposes to merge only the feature-set that is present in the Parallelism TS 2. (Note: The first revision of this paper did not propose a merge.) If, due to feedback, any of these features require a change, then this paper (P1928) is the intended vehicle. If a new feature is basically an addition to the wording proposed here, then it will progress in its own paper.

3.1

RELATED PAPERS

P0350 Before publication of the TS, SG1 approved [P0350R0] which did not progress in time in LEWG to make it into the TS. P0350 is moving forward independently.

P0918 After publication of the TS, SG1 approved [P0918R2] which adds `shuffle`, `interleave`, `sum_to`, `multiply_sum_to`, and `saturated_simd_cast`. P0918 will move forward independently.

P1068 R3 of the paper removed discussion/proposal of a `simd` based API because it was targeting C++23 with the understanding of `simd` not being ready for C++23. This is unfortunate as the presence of `simd` in the IS might lead to a considerably different assessment of the iterator/range-based API proposed in P1068.

P0917 The ability to write code that is generic wrt. arithmetic types and `simd` types is considered to be of high value (TS feedback). Conditional expressions via the `where` function were not all too well received. Conditional expressions via the conditional operator would provide a solution deemed perfect by those giving feedback (myself included).

DRAFT ON NON-MEMBER OPERATOR[] TODO

P2600 The fix for ADL is important to ensure the above two papers do not break existing code.

P0543 The paper proposing functions for saturation arithmetic expects `simd` overloads as soon as `simd` is merged to the IS.

P0553 The bit operations that are part of C++20 expects `simd` overloads as soon as `simd` is merged to the IS.

P2638 Intel's response to P1915R0 for `std::simd`

P2663 `std::simd<std::complex<T>>`.

P2664 Permutations for `simd`.

P2509 D'Angelo [P2509R0] proposes a “type trait to detect conversions between arithmetic-like types that always preserve the numeric value of the source object”. This matches the *value-preserving* conversions the `simd` specification uses.

The papers P0350, P0918, P2663, P2664, and the `simd`-based P1068 fork currently have no shipping vehicle and are basically blocked on this paper.

4

CHANGES AFTER TS FEEDBACK

[P1915R0] (Expected Feedback from `simd` in the Parallelism TS 2) was published in 2019, asking for feedback to the TS. I received feedback on the TS via the GitHub issue tracker, e-mails, and personal conversations. There is also a lot of valuable feedback published in P2638 “Intel’s response to P1915R0 for `std::simd`”. This paper captures the major change requests but should still be considered a work-in-progress.

4.1

IMPROVE ABI TAGS

I received consistent feedback that `simd_abi::compatible<T>` is the wrong default and it should rather be `simd_abi::native<T>` instead. All my tutorial material instructed users to use `std::experimental::native_simd<T>`. There really is little use for `simd_abi::compatible<T>`. The preferred approach should be the use of `simd_abi::native<T>` together with compiler flags that limit the available registers and instructions to whatever the user deems “compatible”. Consequently, there is no reason to keep `simd_abi::compatible<T>` in its current form.

Another common question was about a “fixed size” ABI tag, similar to `std::experimental::simd_abi::fixed_size<N>` but without the ABI compatibility cost.¹ Basically, the ABI footgun should be as dangerous as `std::experimental::simd_abi::native<T>`. The answer to that FAQ is to use `std::experimental::simd_abi::deduce_t<T, N>` as ABI tag. This will provide you with a high-performance footgun, if supported, but might also fall back to `std::experimental::simd_abi::fixed_size<N>`. With `std::experimental::simd_abi::deduce_t<T, N>` turning out to be used potentially more often than `std::experimental::simd_abi::fixed_size<N>` the aliases and names should be revisited. My proposal:

A0 = `simd_abi::native<T>` (no change from the TS semantics)

`simd<T, A0>` abstracts a SIMD register (or similar) with highest performance on the target system (typically widest available register, but that’s a QoI choice). Consequently, the number of elements is chosen by the implementation and may differ for different `T` and different compiler flags. `simd_abi::native<T>` is an alias for an unspecified type. `simd_abi::native<T>` can be an alias for `simd_abi::scalar`. If `sizeof(simd<T, A0>)` or

¹ Implementations of the TS are encouraged to make passing `fixed_size` objects ABI compatible between different hardware generations and/or even different architectures.

`alignof(simd<T, A0>)` in TU1 differ from the same expressions in TU2, then the types `A0` in TU1 and TU2 have a different name.

A1 = `simd_abi::fixed_size<T, N>` *(different to the TS semantics)*

`simd<T, A1>` abstracts one or more registers storing `N` values. The actual hardware resources might store more values; but instructions are generated to make it appear as if there are exactly `N` values stored and manipulated.

Parameter passing may be ABI incompatible between different TUs when compiled with different compiler flags. Therefore, if `sizeof(simd<T, A0>)`² or `alignof(simd<T, A0>)` in TU1 differ from the same expressions in TU2, then the types `A1` in TU1 and TU2 have a different name. This new requirement (wrt. the TS) is the reason for the additional `T` parameter. This allows an implementation to define the `fixed_size` alias as e.g. `template <typename T, int N> using fixed_size = _Fixed<N, native<T>>;`. `A1` and `A0` are always different types, i.e. even if `simd_size_v<T, A0> == N`.

A major difference to `std::experimental::simd_abi::fixed_size<N>` in the TS is about `simd_mask`. In order to support ABI stability the `simd_mask` implementation must choose one form of storage for all possible targets:

- full SIMD registers with all bits set to 1 or 0 per corresponding element
- bitmasks
- an array of `bool` or similar

The new intent for the `fixed_size<T, N>` ABI tag would be to allow `simd_mask<T, A1>` to use either choice depending only on compiler flags.

A2 = `simd_abi::scalar`

No change.

At this point the `simd_abi::deduce` facility seems to be obsolete. However, it is still a useful tool for implementing the `rebind_simd` and `resize_simd` traits. Without more compelling reason for removal, it should be merged as is.

4.1.1

NAMING DISCUSSION

For context on naming, consider the use-cases that the ABI tags serve:

`simd_abi::native<T>` The equivalent to `T`: a direct abstraction of available hardware resources in terms of registers and instructions.

² `A0` is not a typo; this depends on `simd_abi::native<T>`

simd_abi::fixed_size<T, N> Higher abstraction level than **native<T>**: the user/algorithm dictates the number of elements to be processed in parallel. The objects might not be direct mappings to hardware resources, but they use the best that is available on the given target system.

simd_abi::scalar The actual type of **native<T>** if the target hardware has no support for parallel processing of elements of T.³ In addition, **simd<T, simd_abi::scalar>** can be a useful debugging tool.

For reference, the name **fixed_size** is my preference over the following alternatives:

- **simd_abi::fixed_native<N>** with **simd** alias **fixed_native_simd<T, N>**
- **simd_abi::fixed<N>** with **simd** alias **fixed_simd<T, N>**
- **simd_abi::sized<N>** with **simd** alias **sized_simd<T, N>**

4.2

SIMPLIFY/GENERALIZE CASTS

The change to the ABI tags requires a reconsideration of cast functions and implicit and explicit casts between data-parallel types of different ABI tags. This is in addition to TS feedback on casts being too strict or cumbersome to use.

4.2.1

MORE (EXPLICIT) CONVERTING CONSTRUCTORS

The TS allows implicit casts between **fixed_size<N>** types that only differ in element type and where the values are preserved (“every possible value of U can be represented with type **value_type**”).

However, from experience with the TS, it is better to also enable implicit conversions between any **simd** specializations with equal element count, even if such a conversion might be non-portable between targets with different native SIMD widths. The expectation is, that users set up their types according to a pattern similar to Listing 1. Thus, users will work with a set of types that have equal

```

1 using floatv = std::simd<float>;
2 using doublev = std::rebind_simd_t<double, floatv>;
3 using int32v = std::rebind_simd_t<std::int32_t, floatv>;
4 using uint32v = std::rebind_simd_t<std::uint32_t, floatv>;
5 using int16v = std::rebind_simd_t<std::int16_t, floatv>;
6 using uint16v = std::rebind_simd_t<std::uint16_t, floatv>;
7 // ...

```

Listing 1: Recommended setup of **simd** types

³ A typical example is **simd_abi::native<long double>**.

number of elements by construction. Some of the types may use the `fixed_size` ABI tag and some may use an extended ABI tag. This detail should not stop the user from being able to cast between a compiler-flag dependent subset of these types.

Besides a constraint on the number of elements being equal, the converting constructor should be conditionally `explicit`: Implicit casts are only allowed if the element type conversion is value-preserving (same wording as in the TS).

This resolves major inconveniences when working with mixed-precision operations (cf. Tony Table 1). Type conversions for `simd` are still less error-prone than builtin types, because conversions

Parallelism TS 2	with P1928R3
<pre> namespace stdx = std::experimental; template <class T> void f(T a, int b) { using I = std::conditional_t< stdx::is_simd_v<T>, stdx::rebind_simd_t<int, T>, int>; I c; if constexpr (stdx::is_simd_v<T>) { c = stdx::static_simd_cast<I>(a) + b; } else { c = static_cast<int>(a) + b; } g(c); } </pre>	<pre> // assuming simd in namespace std template <class T> void f(T a, int b) { using I = std::conditional_t< std::is_simd_v<T>, std::rebind_simd_t<int, T>, int>; I c = static_cast<I>(a) + b; g(c); } </pre>

Tony Table 1: Improved generic code after adding converting constructors

that might lose information require an explicit cast. Also, unintended widening of the SIMD register size can happen, but typically leads to the need for an explicit cast in the complete statement (cf. Listing 2).

```

1 void f(int32v a, doublev b, floatv c)
2 {
3     doublev x = a * b + c; // OK: implicit (value-preserving) conversion from int and float
4     // to double. Requires twice the register space, but there's no way around it and the
5     // result type requires it anyway.
6     int32v y = a * b; // ERROR: implicit conversion from double to int not value-preserving
7     int32v z1 = static_cast<int32v>(a * b); // OK: cast hints at implicit register widening
8     int32v z2 = a * static_cast<int32v>(b); // OK
9 }

```

Listing 2: Mixed precision code using the types from Listing 1, ensuring equal element count

4.2.2

REMOVE NAMED CAST FUNCTIONS

From the cast functions `std::experimental::to_fixed_size`, `std::experimental::to_native`, and `std::experimental::to_compatible` only the conversions from `simd_abi::fixed_size<T, N>` to `simd_abi::native<T>` and back may still benefit from a named cast function. Most importantly, the conversion from `native` to its `fixed_size` counterpart benefits from a cast expression that does not require spelling out the destination type. However, since converting constructors are provided by the standard library, it is simple for users to define their own `to_fixed_size` function if they want one (e.g. Listing 3). The reverse cast can trivially be spelled

```

1 template <typename T>
2 constexpr std::fixed_size_simd<T, std::simd_size_v<T>> to_fixed_size(std::simd<T> x)
3 {
4     return x;
5 }

```

Listing 3: Example of a user-defined `to_fixed_size` implementation if explicit casts are provided

out as `static_cast<simd<T>>(y)` in program code. The only motivation for adding a `to_native` function would be the provision of a counterpart for the `to_fixed_size` cast function.

Besides the functions only implementing trivial implicit casts, there is little to no need for these functions. The named cast functions are therefore removed altogether.

4.2.3

REMOVE `simd_cast` AND `static_simd_cast`

There are two cast function templates in the TS: `simd_cast` and `static_simd_cast`. The former is equivalent to the latter except that only value-preserving conversions are allowed. The template parameter can either be a `simd` specialization or a vectorizable type `T`. In the latter case, the cast function determines the return type as `fixed_size_simd<T, input.size()>`.

Since we allow all conversions covered by `std::experimental::simd_cast` and `std::experimental::static_simd_cast` via `std::simd` constructors, the cast functions can be removed altogether. The lost feature (cast via element type) can be replaced using `rebind_simd` as shown in Tony Table 2.

4.2.4

MASK CASTS

`simd_mask` casts should work when `simd` casts work. I.e. if `simd<T0, A0>` is implicitly convertible to `simd<T1, A1>` then `simd_mask<T0, A0>` is implicitly convertible to `simd_mask<T1, A1>`. The reverse (if `simd_mask` is convertible then `simd` is convertible) does not have to be true. Specifically, the TS allows all `fixed_size<N>` mask to be interconvertible, irrespective of the element type. For the IS merge, the proposal is to make this more consistent with `simd` while also

Parallelism TS 2	with P1928R3
<pre> template <typename V> void f(V x) { const auto y = std::static_simd_cast<double>(x); // ... } </pre>	<pre> template <typename V> void f(V x) { const auto y = std::rebind_simd_t<double, V>(x); // ... } </pre>

Tony Table 2: Casting without specifying the target ABI tag

preserving most of the convenience: Allow implicit conversions if the `sizeof` the element types are equal, otherwise the conversion must be explicit.

Conversions with different element count are not possible via a constructor (consistent with `simd`). This would require a different function, such as the `resize<N>(simd)` function proposed by Towner et al. [P2638R0].

4.2.5

COMPLETE CASTS FOR `simd_mask`

The `simd_cast` and `static_simd_cast` overloads for `simd_mask` were forgotten for the TS. Without those casts (and no casts via constructors) mixing different arithmetic types is painful. There is no motivation for forbidding casts on `simd_mask`.

The proposed changes for casts solve this issue.

4.2.6

SUMMARY OF CASTS

1. `simd<T0, A0>` is convertible to `simd<T1, A1>` if `simd_size_v<T0, A0> == simd_size_v<T1, A1>`.
2. `simd<T0, A0>` is implicitly convertible to `simd<T1, A1>` if, additionally, the conversion `T0` to `T1` is value-preserving.
3. `simd_mask<T0, A0>` is convertible to `simd_mask<T1, A1>` if `simd_size_v<T0, A0> == simd_size_v<T1, A1>`.
4. `simd_mask<T0, A0>` is implicitly convertible to `simd_mask<T1, A1>` if, additionally, `sizeof(T0) == sizeof(T1)`.
5. `simd<T0, A0>` can be `bit_casted` to `simd<T1, A1>` if `sizeof(simd<T0, A0>) == sizeof(simd<T1, A1>)`.
6. `simd_mask<T0, A0>` can be `bit_casted` to `simd_mask<T1, A1>` if `sizeof(simd_mask<T0, A0>) == sizeof(simd_mask<T1, A1>)`.

4.3

ADD `simd_mask` GENERATOR CONSTRUCTOR

The `simd` generator constructor is very useful for initializing objects from scalars in a portable (i.e. different `simd::size()`) fashion. The need for a similar constructor for `simd_mask` is less frequent, but, even if only for consistency, there should be one. Besides consistency, it is also useful, of course. Consider a predicate function that is given without `simd` interface (e.g. from a library). How do you construct a `simd_mask` from it? With a generator constructor it is easy:

```
simd<T> f(simd<T> x, Predicate p) {
    const simd_mask<T> k([&](auto i) { return p(x[i]); });
    where(k, x) = 0;
    return x;
}
```

Without the generator constructor one has to write e.g.:

```
simd<T> f(simd<T> x, Predicate p) {
    simd_mask<T> k;
    for (size_t i = 0; i < simd<T>::size(); ++i) {
        k[i] = p(x[i]);
    }
    where(k, x) = 0;
    return x;
}
```

The latter solution makes it hard to initialize the `simd_mask` as `const`, is more verbose, is harder to optimize, and cannot use the sequencing properties the generator constructor allows.

Therefore add:

```
template<class G> simd_mask(G&& gen) noexcept;
```

4.4

DEFAULT LOAD/STORE FLAGS TO `element_aligned`

Consider:

```
1 std::simd<float> v(addr, std::vector_aligned);
2 v.copy_from(addr + 1, std::element_aligned);
3 v.copy_to(dest, std::element_aligned);
```

Line 1 supplies an optimization hint to the load operation. Line 2 says what really? “Please don’t crash. I know this is not a vector aligned access⁴.” Line 3 says: “I don’t know whether it’s vector aligned or not. Compiler, if you know more, please optimize, otherwise just don’t make it crash.” (To clarify, the difference between lines 2 and 3 is what line 1 says about the alignment of `addr`.) In

⁴ Of course, vector aligned is equivalent to element aligned if `simd<float>::size() == 1`

both cases of `element_aligned` access, the developer requested a behavior we take as given in all other situations. Why does the TS force to spell it out in this case?

Since C++20, we also have another option:

```
1 std::simd<float> v(std::assume_aligned<std::memory_alignment_v<std::simd<float>>>(addr));
2 v.copy_from(addr + 1);
3 v.copy_to(dest);
```

This seems to compose well, except that line 1 is rather long for a common pattern in this interface. Also, this removes implementation freedom because the library cannot statically determine the alignment properties of the pointer.

Consequently, as a minimal improvement to the TS keep the load/store flags as is, but default them to `element_aligned`. I.e.:

```
1 std::simd<float> v(addr, std::vector_aligned);
2 v.copy_from(addr + 1);
3 v.copy_to(dest);
```

Section 5.4 discusses an option for additional flags.

4.5

CONTIGUOUS ITERATORS FOR LOADS AND STORES

After Ranges and Concepts introduced `std::contiguous_iterator`, the load/store interface for `simd` can easily be generalized from `U*` to `std::contiguous_iterator` with additional constraints for `input_iterator/output_iterator` and `iter_value_t`. This was not a possible design choice for the TS but does make a lot of sense to modernize with the merge. Therefore, the merge generalizes the load/store interfaces to look like Listing 4.

```
1 template <contiguous_iterator It, typename Flags = element_aligned_tag>
2 requires detail::vectorizable<iter_value_t<It>>
3 constexpr void copy_from(const It& first, Flags f = {});
4
5 template <contiguous_iterator It, typename Flags = element_aligned_tag>
6 requires output_iterator<It, Tp> && detail::vectorizable<iter_value_t<It>>
7 constexpr void copy_to(const It& first, Flags f = {}) const;
```

Listing 4: `copy_from` and `copy_to` declarations using `contiguous_iterator`

4.6

constexpr EVERYTHING

The libstdc++ implementation implements the complete TS API as `constexpr` as an optional extension. This is useful (e.g. for computing constants) and not a significant implementation burden. Users (as well as Towner et al. [P2638R0]) have called for `constexpr`. The merge consequently adds `constexpr` to all functions.

4.7

SPECIFY `simd::size` AS `integral_constant`

The TS specifies `simd::size` as a `static constexpr` function returning the number of elements of the `simd` specialization. Instead of a function, this paper uses a static data member of type `std::integral_constant<std::size_t, N>`, which is both convertible to `std::size_t` and callable. The upside of using a static data member is that it can be used as function parameter without conversion to integer and thus easily pass the size into a function as constant expression. See Listing 5 for an example.

```

1  template <std::ranges::contiguous_range R, std::size_t Size>
2  std::span<const std::ranges::range_value_t<R>, Size>
3  auto subscript(const R& r, std::size_t first, std::integral_constant<std::size_t, Size>) {
4      return std::span<const std::ranges::range_value_t<R>, Size>(
5          std::ranges::data(r) + first, Size());
6  }
7
8  void g(std::vector<float> data) {
9      std::simd<float> v;
10     for (std::size_t i = 0; i + v.size < data.size(); i += v.size) {
11         v = subscript(data, i, v.size); // simd::simd(span) to be proposed
12         // ...
13     }
14 }
```

Listing 5: Example: Pass `simd::size` as “constant expression function argument”

4.8

REPLACE `where` FACILITIES

The `where` functions and corresponding `where_expression` have been the most controversial part going into the TS. My interpretation of the feedback I received is that users can work with it but do not find it intuitive. Instead, many have asked for a `blend` / `select` / `conditional` operator instead. Whenever I asked users whether they would like to use the `?:` operator I got positive and often enthusiastic responses. An overloaded `operator?:` would open the door to generic and intuitive SIMD code.

A major motivation for the `where` function in the TS was its ability to express masked *operations* in addition to masked assignments. This enables library implementations to explicitly use masked operation intrinsics instead of resorting to an unmasked operation with subsequent masked assignment. The latter can be contracted to a masked operation by compilers, but obviously there's no guarantee. In any case, the topic is a QoI issue that doesn't have to dictate the API.

If `operator?:` had been overloadable when I designed `std::experimental::simd` then I would have propopsed `?:` overloads for `simd_mask` and `simd`. Consequently, `where` would likely not have existed. Sadly we still cannot overload `operator?:` even though there has been positive feedback in EWG-I. That work is currently blocked on [P2600R0].

4.8.1

PROPOSED REPLACEMENTS FOR `where`

This paper proposes the following replacements for `std::experimental::where`:

- Overloads for `simd::copy_from`, `simd::copy_to`, `simd_mask::copy_from`, `simd_mask::copy_to`, `reduce`, `hmin`, and `hmax` with additional `simd_mask` parameter.

There are still open questions on these functions, discussed in Section 5.3.

- hidden friend `operator?:`⁵ / `conditional_operator` functions in `simd` and `simd_mask`:

- `simd simd::operator?: (mask_type, simd, simd)`
- `template <class U1, class U2>`
`requires convertible_to<simd_mask, rebind_simd_t<common_type_t<U1, U2>, simd_mask>`
`friend constexpr rebind_simd_t<common_type_t<U1, U2>, simd_type>`
`simd_mask::operator?: (simd_mask, U1, U2)`
- `simd_mask simd_mask::operator?: (simd_mask, simd_mask, simd_mask)`
- `simd_mask simd_mask::operator?: (simd_mask<K, KAbi>, simd_mask, simd_mask<U, UAbi>)`
(for disambiguation of the above because `simd_masks` can be interconvertible)
- `simd_mask simd_mask::operator?: (simd_mask, bool, bool)`
(for consistency; it's not very useful)

- facilities for converting `simd_mask<T>` to `simd<T>` with values 0 or 1:
 - `simd_mask::operator simd_type` (not explicit, preferably with 4.1 of [P2600R0] adopted)
 - unary `simd_mask::operator+`; equivalent to `+(operator simd_type())`
 - unary `simd_mask::operator-`; equivalent to `-(operator simd_type())`
 - unary `simd_mask::operator~`; equivalent to `~(operator simd_type())`

(Wording for the above is still TBD.)

4.8.2

EXAMPLES

Listing 6 presents a few simple examples of working with a `simd_mask` result in the absence of `where`. Note that the compiler I used implements the ADL fix proposed in [P2600R0] and implements `operator?:` overloading as explored in [D0917].

⁵ as soon as EWG lifts the restriction...

```

1  auto f0(std::simd<int> x) { return x > 0 ? 2 * x : x; }
2  /*      vpxor    xmm1, xmm1, xmm1
3          vpcmpd   k1, zmm1, zmm0, 1
4          vpslld   zmm0{k1}, zmm0, 1
5          ret
6  */
7  auto f1(std::simd<int> x) { return x > 0 ? 1 : 0; }
8  /*      vmovdqa32  zmm1, zmm0
9          vpxor     xmm0, xmm0, xmm0
10         vpcmpd    k1, zmm0, zmm1, 1
11         mov       eax, 1
12         vpbroadcastd zmm0{k1}{z}, eax
13         ret
14 */
15 auto f2(std::simd<int> x) { return std::simd(x > 0); }
16 /*      vmovdqa32  zmm1, zmm0
17         vpxor     xmm0, xmm0, xmm0
18         vpcmpd    k1, zmm0, zmm1, 1
19         mov       eax, 1
20         vpbroadcastd zmm0{k1}{z}, eax
21         ret
22 */
23 auto f3(std::simd<int> x) { return -(x > 0); }
24 /*      vmovdqa32  zmm1, zmm0
25         vpxor     xmm0, xmm0, xmm0
26         vpcmpd    k0, zmm0, zmm1, 1
27         vpmovm2d   zmm0, k0
28         ret
29 */
30 auto f4(std::simd<int> x) { return x > 0 ? -1 : 0; }
31 /*      vmovdqa32  zmm1, zmm0
32         vpxor     xmm0, xmm0, xmm0
33         vpcmpd    k0, zmm0, zmm1, 1
34         vpmovm2d   zmm0, k0
35         ret
36 */
37 auto f5(std::simd<int> x) { return x > 0 ? true : false; }
38 /*      vmovdqa32  zmm1, zmm0
39         vpxor     xmm0, xmm0, xmm0
40         vpcmpd    k0, zmm0, zmm1, 1
41         kmovw     eax, k0
42         ret
43 */

```

Listing 6: `simd` conditionals without `where` and with [P2600R0] and [D0917], showing the corresponding assembly output (`gcc -O2 -std=c++23 -march=skylake-avx512`; personal GCC 12.1 branch with patches implementing [P2600R0] and [D0917])

- The function `f0` scales all positive values in `x` by 2. The compiler contracts the blending of `2 * x` and `x` with the multiply operation (a left shift by 1) to a masked left shift instruction.
- The functions `f1` and `f2` both return a `simd<int>` where all positive entries of `x` are replaced by 1 and the remaining entries are 0. I.e. converting the comparison result to `simd` works analogue to promotion of `bool` to `int`.
- The functions `f3` and `f4` both return a `simd<int>` where all positive entries of `x` are replaced by -1 and the remaining entries are 0. The ISA allows a more efficient translation and the compiler recognizes the pattern in both variants.
- Finally, to complete the set, `f5` shows how one could even blend `bool` arguments into a `simd_mask`. The compiler recognizes that the conditional operator is a no-op and simply returns the result of the comparison itself.

Tony Table 3 presents an algorithm for counting all positive `float` values in a `std::vector`. For simplicity, the code uses `vector<simd<float>>` and assumes the `vector` is not empty. If a `simd_mask` implicitly converts to a `simd` (as proposed and analogue to `bool`), the code is simplified significantly. However, at this point, the TS implementation compiles to a masked add instruction while the implementation for this paper does not. The difference is that the former executes an unmasked addition followed up by a masked assignment while the latter converts the mask into a `simd` of 1s and 0s followed up by an unmasked addition. The compiler needs to recognize this pattern in order to reach the same performance (QoI).

4.9

MAKE USE OF `int` AND `size_t` CONSISTENT

The TS uses `int` as NTTP for

- `std::experimental::simd_abi::fixed_size`,
- `std::experimental::fixed_size_simd`,
- `std::experimental::fixed_size_simd_mask`, and
- `std::experimental::resize_simd`.

The constant `std::experimental::simd_abi::max_fixed_size` is of type `int`. The TS uses `size_t` as NTTP for

- `split`, and
- `split_by`.

This paper uses `integral_constant<size_t,  $\mathcal{W}$ >` for

Parallelism TS 2

with P1928R3

```

namespace stdx = std::experimental;
int count_positive(
    const std::vector<stdx::native_simd<float>>& x)
{
    // simplify generated assembly:
    if (x.size() == 0) std::unreachable();
    using floatv = stdx::native_simd<float>;
    using intv = stdx::rebind_simd_t<int, floatv>;
    intv counter = {};
    for (stdx::simd v : x) {
        auto k = stdx::static_simd_cast<
            intv::mask_type>(v > 0);
        ++where(k, counter);
    }
    return reduce(counter);
}
/*      mov     edx, 1
      mov     rcx, QWORD PTR [rdi+8]
      mov     rax, QWORD PTR [rdi]
      vpxor   xmm0, xmm0, xmm0
      vxorps  xmm2, xmm2, xmm2
      vpbroadcastd  zmm1, edx
.L12:
      vcmpps  k1, zmm2, ZMMWORD PTR [rax], 1
      add     rax, 64

      vpaddd  zmm0{k1}, zmm0, zmm1
      cmp     rcx, rax
      jne     .L12
      vmovdqa ymm1, ymm0
      vextracti64x4  ymm0, zmm0, 0x1
      vpaddd  ymm0, ymm1, ymm0
      vmovdqa xmm1, xmm0
      vextracti64x2  xmm0, ymm0, 0x1
      vpaddd  xmm1, xmm1, xmm0
      vpshufd xmm0, xmm1, 27
      vpaddd  xmm0, xmm0, xmm1
      vpunpckhqdq  xmm1, xmm0, xmm0
      vpaddd  xmm0, xmm0, xmm1
      vmovd   eax, xmm0
      vzeroupper
      ret
*/

```

```

int count_positive(
    const std::vector<std::simd<float>>& x)
{
    // simplify generated assembly:
    if (x.size() == 0) std::unreachable();
    using floatv = std::simd<float>;
    using intv = std::rebind_simd_t<int, floatv>;
    intv counter = {};
    for (std::simd v : x) {
        counter += v > 0;
    }
    return reduce(counter);
}
/*      mov     edx, 1
      mov     rcx, QWORD PTR [rdi+8]
      mov     rax, QWORD PTR [rdi]
      vpxor   xmm0, xmm0, xmm0
      vxorps  xmm3, xmm3, xmm3
      vpbroadcastd  zmm2, edx
.L9:
      vcmpps  k1, zmm3, ZMMWORD PTR [rax], 1
      add     rax, 64
      vmovdqa32  zmm1{k1}{z}, zmm2
      vpaddd  zmm0, zmm0, zmm1
      cmp     rcx, rax
      jne     .L9
      vmovdqa ymm1, ymm0
      vextracti64x4  ymm0, zmm0, 0x1
      vpaddd  ymm0, ymm1, ymm0
      vmovdqa xmm1, xmm0
      vextracti64x2  xmm0, ymm0, 0x1
      vpaddd  xmm1, xmm1, xmm0
      vpshufd xmm0, xmm1, 27
      vpaddd  xmm0, xmm0, xmm1
      vpunpckhqdq  xmm1, xmm0, xmm0
      vpaddd  xmm0, xmm0, xmm1
      vmovd   eax, xmm0
      vzeroupper
      ret
*/

```

Tony Table 3: Counting positive values in a `std::vector`

- `simd_size<T, Abi>`,
- `simd<T, Abi>::size`, and
- `simd_mask<T, Abi>::size`.

Finally, `simd_size_v<T, Abi>` is of type `size_t`.

All of these integers denote a SIMD width. They should be consistent. Since the `size` member will never get consensus to use type `int`, the decision falls on `size_t` for all.

The merge proposal therefore uses `size_t` for `std::experimental::simd_abi::fixed_size`, `std::experimental::fixed_size_simd`, `std::experimental::fixed_size_simd_mask`, `std::experimental::resize_simd`, and `std::experimental::simd_abi::max_fixed_size`.

4.10

CLEAN UP MATH FUNCTION OVERLOADS

The wording that produces `simd` overloads misses a few cases and leaves room for ambiguity. There is also no explicit mention of integral overloads that are supported in `<cmath>` (e.g. `std::cos(1)` calling `std::cos(double)`). At the very least, `std::abs(simd<signed-integral>)` should be specified.

Also, from implementation experience, “undefined behavior” for domain, pole, or range error is unnecessary. It could either be an unspecified result or even match the expected result of the function according to Annex F in the C standard. The latter could possibly be a recommendation, i.e. QoI. The intent is to avoid `errno` altogether, while still supporting floating-point exceptions (possibly depending on compiler flags).

This needs more work and is not reflected in the wording at this point.

5

OPEN QUESTIONS

5.1

ALTERNATIVES TO `hmin` AND `hmax`

The functions `hmin(simd)` and `hmax(simd)` are basically specializations of `reduce(simd)`. I received feedback asking for better names.

With C++17, there was nothing equivalent to `std::plus` for minimum and maximum. Since the merge of `Ranges` (C++20), we have `std::ranges::min` and `std::ranges::max`. The `reduce(simd)` specification requires the `binary_op` to be callable with two `simd` arguments, though (split initial argument in half, call `binary_op`, split again, call `binary_op`, ...until only a scalar is left). This doesn't work with `std::ranges::min` (and `max`) because it requires an lvalue reference as return type. If we added another `operator()` to `std::ranges::min`, then their use with `reduce(simd)` would be slightly inconsistent:

```

simd<unsigned> v = ...;
auto a = reduce(v, std::bit_and<>); // must type <>
auto b = reduce(v, std::ranges::min); // must *not* type <>

```

However, if `simd` will be an `input_range` (see Section 5.5) then the `std::ranges::min(ranges::input_range auto&&, ...)` overload matches and `std::ranges::min(simd)` works out of the box. We could then leave it up to QoI to recognize the opportunity for a SIMD implementation of the reduction.

Alternatively (or in addition) we could rename the TS functions to `reduce_min(simd)` and `reduce_max(simd)`.

5.1.1

SUGGESTED POLLS

Poll: We want to do something about `hmin` and `hmax`; i.e. the TS status quo is not acceptable for the IS.

SF	F	N	A	SA

Poll: Rename to `reduce_min` and `reduce_max`.

SF	F	N	A	SA

Poll: Extend `std::ranges::min` and `max` to allow prvalue return types.

SF	F	N	A	SA

Poll: Remove `hmin` and `hmax` expecting `simd` to become a range.

SF	F	N	A	SA

5.2

MAKE `simd_mask` REDUCTIONS CONSISTENT WITH `<bit>`

The functions in `<bit>` were added in C++20 [P0553R4] and now `simd_mask` needs to adjust for consistency. Overloads for `simd<unsigned integer type>` will also have to be added. But not with this *merge* paper. I promise to write a follow-up paper.

The three relevant functions in the TS are

- `popcount(simd_mask)` and `popcount(bool)`:

The names match. However, currently `popcount(bool)` is ill-formed. This would be very

unfortunate for generic code that calls `popcount` on the result of a comparison. For `simd_mask<T>` this operation provides value, for `T` it simply casts the `bool` to an `int`. The TS merge would therefore add `bool` to the list of valid types for `popcount`.

- `find_first_set(simd_mask)` and `find_first_set(bool)`:
This matches either `countl_zero` or `countr_zero`, except that the `simd_mask` version has a precondition. Same issue for `bool` as discussed for `popcount`.
- `find_last_set(simd_mask)` and `find_last_set(bool)`:
This matches either `countl_zero` or `countr_zero`, except that the `simd_mask` version has a precondition. Same issue for `bool` as discussed for `popcount`.

I believe we should remove the precondition on the `simd_mask` reductions and allow masks where `none_of` returns `true`. Then behavior is identical to `countl_zero` and `countr_zero`. The reason for the precondition was efficiency on older ISAs. At this point, especially with the precedent set by the `count[lr]` functions, I don't believe the performance issue is worth the UB.

Renaming the `find_first/last_set` functions is not obvious though. Consider that a `simd_mask k` can be indexed via subscript operator. Then `k[0]` identifies the first element in `k`. For all of us who think left-to-right, it would therefore be intuitive to have `countl_zero` be the replacement for `find_first_set`. But if you were to convert `k` into a `std::bitset` and convert the `bitset` into an unsigned integer, then `countl_zero` would have the opposite meaning (starting from the most significant bit). Indexing a `simd_mask` and a `bitset` is consistent as is the confusion about left and right with regard to the integer representation of a `bitset`⁶.

I believe it was wrong to call a bit position left or right. That is always ambiguous. Assuming that we will not change the existing names, the consistent name to use for `std::experimental::find_first_set` is `std::countr_zero`. To put the new name into perspective take a look at Listing 7. My mental model is iterating the `string` from left to right, but on every SIMD chunk I

```

1 int find(char needle, const std::string& hay) {
2     using V = std::simd<char>;
3     for (std::size_t i = 0; i + V::size <= hay.size(); i += V::size) {
4         const V chunk(hay.data() + i);
5         if (any_of(chunk == needle)) {
6             return i + countr_zero(chunk == needle);
7             // was: i + find_first_set(chunk == needle) in the TS
8         }
9     }
10 }
```

Listing 7: SIMD algorithm to find a `char` in a `string`

need a twist and search right to left. I believe `countr_zero` is an indefensible name for returning

⁶ cf. <https://godbolt.org/z/ov78q1751>

the lowest index where the mask is `true`. On the other hand, using `countl_zero` would be too inconsistent.

My recommendation is to overload `count[lr]_zero/one` for `simd` but *not* for `simd_mask`. Unless ...

5.2.1

THE STORY OF ELEMENT-WISE OPERATIONS

Everything you can do with `T` you can do with `simd<T>`, only that the operations apply element-wise.

Since `simd_mask<T>` is the return type of `simd<T>` comparisons, the goal for `simd_mask` is:

Everything you can do with `bool` you can do with `simd_mask<T>`, only that the operations apply element-wise.

By that reasoning, `popcount(simd_mask)` should be element-wise `popcount(bool)`, and `countr_zero(simd_mask)` should be element-wise `countr_zero(bool)`. The `<bit>` functions would therefore not be mask reductions. If `<bit>` were extended to include `bool` overloads, there should also be `simd_mask` overloads, applying element-wise.

As a consequence, the mask reductions need a different name than the `<bit>` functions. See Table 1 for a suggestion. Note that if `simd_mask` becomes a `std::ranges::input_range`, then

TS names	alternative names
<pre>void f(stdx::native_simd_mask<int> k) { bool all = all_of(k); bool any = any_of(k); bool none = none_of(k); int count = popcount(k); int first = find_first_set(k); int last = find_last_set(k); }</pre>	<pre>void f(std::simd_mask<int> k) { bool all = reduce_and(k); bool any = reduce_or(k); bool none = reduce_nand(k); int count = reduce_count(k); // or reduce_add int first = reduce_min_index(k); int last = reduce_max_index(k); }</pre>

Table 1: Mask reductions

`std::ranges::all_of(simd_mask)` will be well-formed and agree with the result of `std::experimental::all_of`. (Same for `any_of` and `none_of`.) The naming precedent suggests that `all_of`, `any_of`, and `none_of` should not be renamed.

My recommendation is to merge `all_of`, `any_of`, and `none_of` without change and rename `popcount`, `find_first_set`, and `find_last_set` according to Table 2.

5.2.2

REMOVE `some_of`

While we're at the topic. Let's drop `std::experimental::some_of`. Nobody understands or wants the reduction, apparently. It can be added back later if really needed (unlikely).

TS name	IS name
<code>all_of</code>	<code>all_of</code>
<code>any_of</code>	<code>any_of</code>
<code>none_of</code>	<code>none_of</code>
<code>popcount</code>	<code>reduce_count</code>
<code>find_first_set</code>	<code>reduce_min_index</code>
<code>find_last_set</code>	<code>reduce_max_index</code>

Table 2: Recommended renaming of mask reductions

5.2.3

SUGGESTED POLLS

Poll: Mask reductions should use the `<bit>` functions.

SF	F	N	A	SA

Poll: Rename mask reductions as proposed in P1928R3 Table 2.

SF	F	N	A	SA

Poll: Remove UB of `find_first_set` / `find_last_set`.

SF	F	N	A	SA

Poll: Do not merge / remove `some_of`.

SF	F	N	A	SA

5.3

ARGUMENT ORDER AND NAMING OF MASKED OVERLOADS

In the TS, where-expressions made it possible to reuse existing function names and argument orders for masked operations. With the removal of where-expressions the mask must become a function argument. See Table 3 for a possible pattern to replace where-expressions.

There are more options, of course. E.g. possible replacements for `std::experimental::where_expression::copy_from`:

- `v.copy_from_if(v > 0, ptr)`

Here the condition directly follows the word “if”, which seems helpful. However:

Parallelism TS 2	possible replacement
<pre>stdx::native_simd<float> v = ...; where(v > 0, v).copy_from(ptr, stdx::element_aligned); where(v < 0, v).copy_to(ptr, stdx::element_aligned); float pos_sum = reduce(where(v > 0, v));</pre>	<pre>std::simd<float> v = ...; v.copy_from_if(ptr, v > 0); v.copy_to_if(ptr, v < 0); float pos_sum = reduce_if(v, v > 0);</pre>

Table 3: Possible replacement for where-expressions

- `v.copy_from_if(ptr, v > 0)`
This argument order follows precedent from algorithms, which always append the predicate to the list of arguments (e.g. `copy_if(first, last, d_first, pred)`). In addition, the argument order matches the order of words in the function name: “from” — `ptr`, “if” — `mask`.
- `v.copy_from(v > 0, ptr)`
- `v.copy_from(ptr, v > 0)`
- `v.copy_from_where(v > 0, ptr)`
- `v.copy_from_where(ptr, v > 0)`

The TS has no facility for masked load constructors. I did not receive feedback that such a constructor is needed/wanted, so this paper will not propose one.

I propose to append `_if` to the masked functions and append the mask argument (but before default arguments or additional arguments required for masked operations, such as the `identity_element` in `reduce`).

5.3.1

SUGGESTED POLLS

Poll: The names of masked “overloads” should include an `if` and follow the argument order proposed in P1928R3 Section 5.3

SF	F	N	A	SA

5.4

CONVERTING LOADS & STORES CONSISTENCY

For the TS, we allowed pointers to any *vectorizable* type as valid arguments to `copy_from` and `copy_to`. I.e. loads and stores can be converting operations without a clue in the code other than the type of the pointer. It can therefore happen that a conversion that is not value-preserving goes

unnoticed. The broadcast and `simd` conversion constructors guard against accidental use of such conversions.

I have not received feedback that users wrote buggy because of this liberal interface. However, in the TS process this question was never really considered. Therefore, I just wanted to show a suggestion for a stricter but just as powerful interface. Listing 8 presents converting broadcast and cast expressions, which are ill-formed because the type conversion is not value-preserving.

```

1  float fmem[std::simd_size_v<float>] = {};
2  double dmem[std::simd_size_v<float>] = {};
3  short smem[std::simd_size_v<float>] = {};
4
5  std::simd<float> a = 1.; // ERROR: double -> float conversion is not value-preserving
6  std::simd<float> b = std::rebind_simd_t<int, std::simd<float>>(1); // ERROR:
7                                // int -> float is not value-preserving
8
9  // TS:
10 std::simd<float> ts;
11 ts.copy_from(fmem, std::element_aligned); // OK
12 ts.copy_from(dmem, std::element_aligned); // OK
13 ts.copy_from(smem, std::element_aligned); // OK
14
15 // idea, not status quo of this paper:
16 std::simd<float> v;
17 v.copy_from(fmem); // OK
18 v.copy_from(dmem); // ERROR: converting load
19 v.copy_from(smem); // ERROR: converting load
20
21 // Option (a) - one flag only:
22 v.copy_from(dmem, std::simd_converting); // OK
23 v.copy_from(smem, std::simd_converting); // OK
24
25 // Option (b) - two flags:
26 v.copy_from(dmem, std::simd_safe_cvt); // ERROR: double -> float is not value-preserving
27 v.copy_from(dmem, std::simd_any_cvt); // OK
28 v.copy_from(smem, std::simd_safe_cvt); // OK
29 v.copy_from(smem, std::simd_any_cvt); // OK

```

Listing 8: Load-store flags as opt-in to converting loads and stores

The equivalent conversions on `copy_from` are well-formed, though. I believe it would be better for users to opt-in to conversions on load and store. There are value-preserving and non-value-preserving conversions, which could be combined into the same opt-in or we provide a separate spelling for value-preserving conversions (the safe kind of conversion).

If LEWG is interested, I would be thankful for naming suggestions. I do not believe that using “safe” is a good term here.

5.5

INTEGRATION WITH RANGES

`simd` itself is not a container [P0851R0]. The value of a data-parallel object is not an array of elements but rather needs to be understood as a single opaque value that happens to have means for reading and writing element values. I.e. `simd<int> x = {};` does not start the lifetime of `int` objects. This implies that `simd` cannot model a contiguous range but only a random-access range. `simd` can trivially model `input_range`. However, in order to model `output_range`, the iterator of every non-const `simd` would have to return an `element_reference` on dereference. Without the ability of `element_reference` to decay to the element type (similar to how arrays decay to pointers on deduction), I would prefer to simply make `simd` model only `input_range`.

I plan to pursue adding iterators and conversions to array and from random-access ranges, specifically `span` with static extent, in a follow-up paper. I believe it is not necessary to resolve this question before merging `simd` from the TS.

5.6

CORRECT PLACE FOR `simd` IN THE IS?

While `simd` is certainly very important for numerics and therefore fits into the “Numerics library” clause, it is also more than that. E.g. `simd` can be used for vectorization of text processing. In principle `simd` should be understood similar to fundamental types. Is the “General utilities library” clause a better place? Or rename “Concurrency support library” to “Parallelism and concurrency support library” and put it there? Alternatively, add a new library clause?

I am seeking feedback before making a recommendation.

5.7

`element_reference` IS OVERSPECIFIED

`element_reference` is spelled out in a lot of detail. It may be better to define its requirements in a list of requirements or a table instead.

This change is not reflected in the wording, pending encouragement from WG21 (mostly LWG).

6

WORDING

The following section presents the wording to be applied against the C++ working draft. The subsequent section, Section 6.2, reproduces the same wording as a diff against the Parallelism TS 2.

The wording still needs work:

- Replace `where` & `where_expression` wording with `conditional_operator` and masked overloads.
- Apply the new library specification style from P0788R3.

6.1

ADD SECTION 9 OF N4808 WITH MODIFICATIONS

 Add a new subclause after §28.8 [numbers]

(6.1.1) 28.9 Data-Parallel Types [simd]

(6.1.1.1) 28.9.1 General [simd.general]

- 1 The data-parallel library consists of data-parallel types and operations on these types. A data-parallel type consists of elements of an underlying arithmetic type, called the *element type*. The number of elements is a constant for each data-parallel type and called the *width* of that type.
- 2 Throughout this Clause, the term *data-parallel type* refers to all *supported* (28.9.6.1) specializations of the `simd` and `simd_mask` class templates. A *data-parallel object* is an object of *data-parallel type*.
- 3 An *element-wise operation* applies a specified operation to the elements of one or more data-parallel objects. Each such application is unsequenced with respect to the others. A *unary element-wise operation* is an element-wise operation that applies a unary operation to each element of a data-parallel object. A *binary element-wise operation* is an element-wise operation that applies a binary operation to corresponding elements of two data-parallel objects.
- 4 Throughout this Clause, the set of *vectorizable types* for a data-parallel type comprises all cv-unqualified arithmetic types other than `bool`.
- 5 [*Note*: The intent is to support acceleration through data-parallel execution resources, such as SIMD registers and instructions or execution units driven by a common instruction decoder. If such execution resources are unavailable, the interfaces support a transparent fallback to sequential execution. — *end note*]

(6.1.1.2) 28.9.2 Header <simd> synopsis [simd.synopsis]

```

namespace std {
    namespace simd_abi {
        using scalar = see below;
        template<class T, size_t N> using fixed_size = see below;
        template<class T> inline constexpr size_t max_fixed_size = implementation-defined;
        template<class T> using native = implementation-defined;

        template<class T, size_t N, class... Abis> struct deduce { using type = see below; };
        template<class T, size_t N, class... Abis> using deduce_t =
            typename deduce<T, N, Abis...>::type;
    }

    struct element_aligned_tag {};
    struct vector_aligned_tag {};
    template<size_t> struct overaligned_tag {};
    inline constexpr element_aligned_tag element_aligned{};
    inline constexpr vector_aligned_tag vector_aligned{};
    template<size_t N> inline constexpr overaligned_tag<N> overaligned{};

    // 28.9.4, simd type traits
    template<class T> struct is_abi_tag;

```

```

template<class T> inline constexpr bool is_abi_tag_v = is_abi_tag<T>::value;

template<class T> struct is_simd;
template<class T> inline constexpr bool is_simd_v = is_simd<T>::value;

template<class T> struct is_simd_mask;
template<class T> inline constexpr bool is_simd_mask_v = is_simd_mask<T>::value;

template<class T> struct is_simd_flag_type;
template<class T> inline constexpr bool is_simd_flag_type_v =
    is_simd_flag_type<T>::value;

template<class T, class Abi = simd_abi::native<T>> struct simd_size;
template<class T, class Abi = simd_abi::native<T>>
    inline constexpr size_t simd_size_v = simd_size<T, Abi>::value;

template<class T, class U = typename T::value_type> struct memory_alignment;
template<class T, class U = typename T::value_type>
    inline constexpr size_t memory_alignment_v = memory_alignment<T, U>::value;

template<class T, class V> struct rebind_simd { using type = see below; };
template<class T, class V> using rebind_simd_t = typename rebind_simd<T, V>::type;
template<size_t N, class V> struct resize_simd { using type = see below; };
template<size_t N, class V> using resize_simd_t = typename resize_simd<N, V>::type;

// 28.9.6, Class template simd
template<class T, class Abi = simd_abi::native<T>> class simd;
template<class T, size_t N> using fixed_size_simd = simd<T, simd_abi::fixed_size<T, N>>;

// 28.9.8, Class template simd_mask
template<class T, class Abi = simd_abi::native<T>> class simd_mask;
template<class T, size_t N> using fixed_size_simd_mask = simd_mask<T, simd_abi::fixed_size<T, N>>;

template<size_t... Sizes, class T, class Abi>
    constexpr tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd<T, Abi>&) noexcept;
template<size_t... Sizes, class T, class Abi>
    constexpr tuple<simd_mask<T, simd_mask_abi::deduce_t<T, Sizes>>...>
        split(const simd_mask<T, Abi>&) noexcept;
template<class V, class Abi>
    constexpr array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
        split(const simd<typename V::value_type, Abi>&) noexcept;
template<class V, class Abi>
    constexpr array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
        split(const simd_mask<typename V::simd_type::value_type, Abi>&) noexcept;

```

```

template<size_t N, class T, class A>
constexpr array<resize_simd_t<simd_size_v<T, A> / N, simd<T, A>>, N>
    split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
constexpr array<resize_simd_t<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
    split_by(const simd_mask<T, A>& x) noexcept;

template<class T, class... Abis>
constexpr simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >>
    concat(const simd<T, Abis>&...) noexcept;
template<class T, class... Abis>
constexpr simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >>
    concat(const simd_mask<T, Abis>&...) noexcept;

template<class T, class Abi, size_t N>
constexpr resize_simd_t<simd_size_v<T, Abi> * N, simd<T, Abi>>
    concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
constexpr resize_simd_t<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
    concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;

```

// 28.9.9.4, *simd_mask reductions*

```

template<class T, class Abi> constexpr bool all_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool any_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool none_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool some_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr int popcount(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr int find_first_set(const simd_mask<T, Abi>&);
template<class T, class Abi> constexpr int find_last_set(const simd_mask<T, Abi>&);

```

```

constexpr bool all_of(T) noexcept;
constexpr bool any_of(T) noexcept;
constexpr bool none_of(T) noexcept;
constexpr bool some_of(T) noexcept;
constexpr int popcount(T) noexcept;
constexpr int find_first_set(T);
constexpr int find_last_set(T);

```

// 28.9.5, *Where expression class templates*

```

template<class M, class T> class const_where_expression;
template<class M, class T> class where_expression;

```

// 28.9.9.5, *Where functions*

```

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd<T, Abi>>

```

```

        where(const typename simd<T, Abi>::mask_type&, simd<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, const simd<T, Abi>&) noexcept;

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, simd_mask<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, const simd_mask<T, Abi>&) noexcept;

template<class T>
    where_expression<bool, T>
        where(see below k, T& d) noexcept;

template<class T>
    const_where_expression<bool, T>
        where(see below k, const T& d) noexcept;

// 28.9.7.4, simd reductions
template<class T, class Abi, class BinaryOperation = plus<>>
    constexpr T reduce(const simd<T, Abi>&,
        BinaryOperation = {});

template<class M, class V, class BinaryOperation>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        typename V::value_type identity_element,
        BinaryOperation binary_op);

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        plus<> binary_op = {}) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        multiplies<> binary_op) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_and<> binary_op) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_or<> binary_op) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        bit_xor<> binary_op) noexcept;

```

```

template<class T, class Abi>
    constexpr T hmin(const simd<T, abi>&) noexcept;
template<class M, class V>
    constexpr typename V::value_type hmin(const const_where_expression<M, V>&) noexcept;
template<class T, class Abi>
    constexpr T hmax(const simd<T, abi>&) noexcept;
template<class M, class V>
    constexpr typename V::value_type hmax(const const_where_expression<M, V>&) noexcept;

```

// 28.9.7.6, Algorithms

```

template<class T, class Abi>
    constexpr simd<T, Abi>
        min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr simd<T, Abi>
        max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr pair<simd<T, Abi>, simd<T, Abi>>
        minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr simd<T, Abi>
        clamp(const simd<T, Abi>& v,
              const simd<T, Abi>& lo,
              const simd<T, Abi>& hi);
}

```

- ¹ The header <simd> defines class templates, tag types, trait types, and function templates for element-wise operations on data-parallel objects.

(6.1.1.3) 28.9.3 simd ABI tags

[simd.abi]

```

namespace simd_abi {
    using scalar = see below;
    template<class T, size_t N> using fixed_size = see below;
    template<class T> inline constexpr size_t max_fixed_size = implementation-defined;
    template<class T> using native = implementation-defined;
}

```

- ¹ An *ABI tag* is a type in the `std::simd_abi` namespace that indicates a choice of size and binary representation for objects of data-parallel type. [*Note*: The intent is for the size and binary representation to depend on the target architecture. — *end note*] The ABI tag, together with a given element type implies a number of elements. ABI tag types are used as the second template argument to `simd` and `simd_mask`.
- ² [*Note*: The ABI tag is orthogonal to selecting the machine instruction set. The selected machine instruction set limits the usable ABI tag types, though (see 28.9.6.1). The ABI tags enable users to safely pass objects of data-parallel type between translation unit boundaries (e.g. function calls or I/O). — *end note*]
- ³ `scalar` is an alias for an unspecified ABI tag that is different from `fixed_size<1>`. Use of the `scalar` tag type requires data-parallel types to store a single element (i.e., `simd_size_v<T, simd_abi::scalar>` equals 1).
- ⁴ The value of `max_fixed_size<T>` is at least 32.

- 5 `fixed_size<N>` is an alias for an unspecified ABI tag. `fixed_size` does not introduce a non-deduced context. Use of the `simd_abi::fixed_size<N>` tag type requires data-parallel types to store `N` elements (i.e. `simd_size_v<T, simd_abi::fixed_size<N>>` equals `N`). `simd<T, fixed_size<N>>` and `simd_mask<T, fixed_size<N>>` with `N > 0` and `N <= max_fixed_size<T>` shall be supported. Additionally, for every supported `simd<T, Abi>` (see 28.9.6.1), where `Abi` is an ABI tag that is not a specialization of `simd_abi::fixed_size, N == simd<T, Abi>::size()` shall be supported.
- 6 [*Note*: It is unspecified whether `simd<T, fixed_size<N>>` with `N > max_fixed_size<T>` is supported. The value of `max_fixed_size<T>` can depend on compiler flags and can change between different compiler versions. — *end note*]
- 7 The type of `fixed_size<T, N>` in TU1 differs from the type of `fixed_size<T, N>` in TU2 iff the type of `native<T>` in TU1 differs from the type of `native<T>` in TU2.
- 8 An implementation may define additional *extended ABI tag* types in the `std::simd_abi` namespace, to support other forms of data-parallel computation.
- 9 `native<T>` is an implementation-defined alias for an ABI tag. [*Note*: The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that is supported on the currently targeted system. For target architectures with ISA extensions, compiler flags may change the type of the `native<T>` alias. — *end note*] [*Example*: Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where hardware support for `__simd256` only exists for floating-point types. The implementation therefore defines `native<T>` as an alias for

- `__simd256` if `T` is a floating-point type, and
- `__simd128` otherwise.

— *end example*]

```
template<T, size_t N, class... Abis> struct deduce { using type = see below; };
```

- 10 The member type shall be present if and only if
- `T` is a vectorizable type, and
 - `simd_abi::fixed_size<N>` is supported (see 28.9.3), and
 - every type in the `Abis` pack is an ABI tag.
- 11 Where present, the member typedef type shall name an ABI tag type that satisfies
- `simd_size<T, type> == N`, and
 - `simd<T, type>` is default constructible (see 28.9.6.1).
- If `N` is 1, the member typedef type is `simd_abi::scalar`. [*Note*: Implementations can base the choice on `Abis`, but can also ignore the `Abis` arguments. — *end note*]
- 12 The behavior of a program that adds specializations for `deduce` is undefined.

(6.1.1.4) 28.9.4 `simd` type traits

[`simd.traits`]

```
template<class T> struct is_abi_tag { see below };
```

- 1 The type `is_abi_tag<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a standard or extended ABI tag, and `false_type` otherwise.
- 2 The behavior of a program that adds specializations for `is_abi_tag` is undefined.

```
template<class T> struct is_simd { see below };
```

3 The type `is_simd<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd` class template, and `false_type` otherwise.

4 The behavior of a program that adds specializations for `is_simd` is undefined.

```
template<class T> struct is_simd_mask { see below };
```

5 The type `is_simd_mask<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd_mask` class template, and `false_type` otherwise.

6 The behavior of a program that adds specializations for `is_simd_mask` is undefined.

```
template<class T> struct is_simd_flag_type { see below };
```

7 The type `is_simd_flag_type<class T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is one of

- `element_aligned_tag`, or
- `vector_aligned_tag`, or
- `overaligned_tag<N>` with $N > 0$ and N an integral power of two,

and `false_type` otherwise.

8 The behavior of a program that adds specializations for `is_simd_flag_type` is undefined.

```
template<class T, class Abi = simd_abi::native<T>> struct simd_size { see below };
```

9 `simd_size<T, Abi>` shall have a member value if and only if

- `T` is a vectorizable type, and
- `is_abi_tag_v<Abi>` is true.

[*Note*: The rules are different from those in (28.9.6.1): The member value is present even if `simd<T, Abi>` is not supported for the currently targeted system. — *end note*]

10 If value is present, the type `simd_size<T, Abi>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` with N equal to the number of elements in a `simd<T, Abi>` object.

11 The behavior of a program that adds specializations for `simd_size` is undefined.

```
template<class T, class U = typename T::value_type> struct memory_alignment { see below };
```

12 `memory_alignment<T, U>` shall have a member value if and only if

- `is_simd_mask_v<T>` is true and `U` is `bool`, or
- `is_simd_v<T>` is true and `U` is a vectorizable type.

13 If value is present, the type `memory_alignment<T, U>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` for some implementation-defined N (see 28.9.6.5 and 28.9.8.4). [*Note*: value identifies the alignment restrictions on pointers used for (converting) loads and stores for the give type `T` on arrays of type `U`. — *end note*]

14 The behavior of a program that adds specializations for `memory_alignment` is undefined.

```
template<class T, class V> struct rebind_simd { using type = see below; };
```

- 15 The member type is present if and only if
- V is either `simd<U, Abi0>` or `simd_mask<U, Abi0>`, where U and Abi0 are deduced from V, and
 - T is a vectorizable type, and
 - `simd_abi::deduce<T, simd_size_v<U, Abi0>, Abi0>` has a member type type.
- 16 Let Abi1 denote the type `deduce_t<T, simd_size_v<U, Abi0>, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if V is `simd<U, Abi0>` or `simd_mask<T, Abi1>` if V is `simd_mask<U, Abi0>`.

```
template<size_t N, class V> struct resize_simd { using type = see below; };
```

- 17 The member type is present if and only if
- V is either `simd<T, Abi0>` or `simd_mask<T, Abi0>`, where T and Abi0 are deduced from V, and
 - `simd_abi::deduce<T, N, Abi0>` has a member type type.
- 18 Let Abi1 denote the type `deduce_t<T, N, Abi0>`. Where present, the member typedef type names `simd<T, Abi1>` if V is `simd<T, Abi0>` or `simd_mask<T, Abi1>` if V is `simd_mask<T, Abi0>`.

(6.1.1.5) 28.9.5 Where expression class templates

[simd.whereexpr]

```
template<class M, class T> class const_where_expression {
    const M mask;      // exposition only
    T& data;           // exposition only

public:
    const_where_expression(const const_where_expression&) = delete;
    const_where_expression& operator=(const const_where_expression&) = delete;

    T operator-() const && noexcept;
    T operator+() const && noexcept;
    T operator~() const && noexcept;

    template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const &&
};

template<class M, class T>
class where_expression : public const_where_expression<M, T> {
public:
    template<class U> void operator=(U&& x) && noexcept;
    template<class U> void operator+=(U&& x) && noexcept;
    template<class U> void operator-=(U&& x) && noexcept;
    template<class U> void operator*=(U&& x) && noexcept;
    template<class U> void operator/=(U&& x) && noexcept;
    template<class U> void operator%=(U&& x) && noexcept;
    template<class U> void operator&=(U&& x) && noexcept;
    template<class U> void operator|=(U&& x) && noexcept;
```

```

template<class U> void operator^=(U&& x) && noexcept;
template<class U> void operator<<=(U&& x) && noexcept;
template<class U> void operator>>=(U&& x) && noexcept;

void operator++() && noexcept;
void operator++(int) && noexcept;
void operator--() && noexcept;
void operator--(int) && noexcept;

template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
};

```

- 1 The class templates `const_where_expression` and `where_expression` abstract the notion of selecting elements of a given object of arithmetic or data-parallel type.
- 2 The first templates argument `M` shall be cv-unqualified `bool` or a cv-unqualified `simd_mask` specialization.
- 3 If `M` is `bool`, `T` shall be a cv-unqualified arithmetic type. Otherwise, `T` shall either be `M` or `typename M::simd_type`.
- 4 In this subclause, if `M` is `bool`, `data[0]` is used interchangeably for `data`, `mask[0]` is used interchangeably for `mask`, and `M::size()` is used interchangeably for 1.
- 5 The *selected indices* signify the integers $i \in \{j \in \mathbb{N} | j < M::size() \wedge \text{mask}[j]\}$. The *selected elements* signify the elements `data[i]` for all selected indices i .
- 6 In this subclause, the type `value_type` is an alias for `T` if `M` is `bool`, or an alias for `typename T::value_type` if `is_simd_mask_v<M>` is true.
- 7 [*Note:* The `where` functions 28.9.9.5 initialize `mask` with the first argument to `where` and `data` with the second argument to `where`. — *end note*]

```

T operator-() const && noexcept;
T operator+() const && noexcept;
T operator~() const && noexcept;

```

- 8 **Returns:** A copy of `data` with the indicated unary operator applied to all selected elements.

```

template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const &&

```

- 9 **Requires:**

- If `M` is not `bool`, the largest selected index is less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

- 10 **Effects:** Copies the selected elements as if `mem[i] = static_cast<U>(data[i])` for all selected indices i .

- 11 **Throws:** Nothing.

- 12 **Remarks:** This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
 - `U` is `bool` and `value_type` is `bool`, or

- U is a vectorizable type and `value_type` is not `bool`.

```
template<class U> void operator=(U&& x) && noexcept;
```

13 *Effects:* Replaces `data[i]` with `static_cast<T>(std::forward<U>(x))[i]` for all selected indices i .

14 *Remarks:* This operator shall not participate in overload resolution unless U is convertible to T .

```
template<class U> void operator+=(U&& x) && noexcept;
template<class U> void operator-=(U&& x) && noexcept;
template<class U> void operator*=(U&& x) && noexcept;
template<class U> void operator/=(U&& x) && noexcept;
template<class U> void operator%=(U&& x) && noexcept;
template<class U> void operator&=(U&& x) && noexcept;
template<class U> void operator|=(U&& x) && noexcept;
template<class U> void operator^=(U&& x) && noexcept;
template<class U> void operator<=(U&& x) && noexcept;
template<class U> void operator>=(U&& x) && noexcept;
```

15 *Effects:* Replaces `data[i]` with `static_cast<T>(data @ std::forward<U>(x))[i]` (where `@` denotes the indicated operator) for all selected indices i .

16 *Remarks:* Each of these operators shall not participate in overload resolution unless the return type of `data @ std::forward<U>(x)` is convertible to T . It is unspecified whether the binary operator, implied by the compound assignment operator, is executed on all elements or only on the selected elements.

```
void operator++() && noexcept;
void operator++(int) && noexcept;
void operator--() && noexcept;
void operator--(int) && noexcept;
```

17 *Effects:* Applies the indicated operator to the selected elements.

18 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type T .

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
```

19 *Requires:*

- If `is_simd_flag_type_v<U>` is true, for all selected indices i , i shall be less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by N .
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

- 20 *Effects:* Replaces the selected elements as if `data[i] = static_cast<value_type>(mem[i])` for all selected indices *i*.
- 21 *Throws:* Nothing.
- 22 *Remarks:* This function shall not participate in overload resolution unless
- `is_simd_flag_type_v<Flags>` is true, and
 - either
 - `U` is `bool` and `value_type` is `bool`, or
 - `U` is a vectorizable type and `value_type` is not `bool`.

(6.1.1.6) 28.9.6 Class template `simd` [simd.class]

(6.1.1.6.1) 28.9.6.1 Class template `simd` overview [simd.overview]

```
template<class T, class Abi> class simd {
public:
    using value_type = T;
    using reference = see below;
    using mask_type = simd_mask<T, Abi>;
    using abi_type = Abi;

    static constexpr typename simd_size<T, Abi>::type size;

    constexpr simd() noexcept = default;

    // 28.9.6.4, simd constructors
    template<class U> constexpr simd(U&& value) noexcept;
    template<class U, class UAbi> constexpr explicit(see below) simd(const simd<U, UAbi>&) noexcept;
    template<class G> constexpr explicit simd(G&& gen) noexcept;
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        constexpr simd(const It& first, Flags = {});

    // 28.9.6.5, simd copy functions
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        constexpr void copy_from(const It& first, Flags f = {});
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        requires output_iterator<It, value_type>
        constexpr void copy_to(const It& first, Flags f = {}) const;

    // 28.9.6.6, simd subscript operators
    constexpr reference operator[](size_t);
    constexpr value_type operator[](size_t) const;

    // 28.9.6.7, simd unary operators
    constexpr simd& operator++() noexcept;
```

```
constexpr simd operator++(int) noexcept;
constexpr simd& operator--() noexcept;
constexpr simd operator--(int) noexcept;
constexpr mask_type operator!() const noexcept;
constexpr simd operator~() const noexcept;
constexpr simd operator+() const noexcept;
constexpr simd operator-() const noexcept;
```

// 28.9.7.1, simd *binary operators*

```
friend constexpr simd operator+(const simd&, const simd&) noexcept;
friend constexpr simd operator-(const simd&, const simd&) noexcept;
friend constexpr simd operator*(const simd&, const simd&) noexcept;
friend constexpr simd operator/(const simd&, const simd&) noexcept;
friend constexpr simd operator%(const simd&, const simd&) noexcept;
friend constexpr simd operator&(const simd&, const simd&) noexcept;
friend constexpr simd operator|(const simd&, const simd&) noexcept;
friend constexpr simd operator^(const simd&, const simd&) noexcept;
friend constexpr simd operator<<(const simd&, const simd&) noexcept;
friend constexpr simd operator>>(const simd&, const simd&) noexcept;
friend constexpr simd operator<<(const simd&, int) noexcept;
friend constexpr simd operator>>(const simd&, int) noexcept;
```

// 28.9.7.2, simd *compound assignment*

```
friend constexpr simd& operator+=(simd&, const simd&) noexcept;
friend constexpr simd& operator-=(simd&, const simd&) noexcept;
friend constexpr simd& operator*=(simd&, const simd&) noexcept;
friend constexpr simd& operator/=(simd&, const simd&) noexcept;
friend constexpr simd& operator%=(simd&, const simd&) noexcept;
friend constexpr simd& operator&=(simd&, const simd&) noexcept;
friend constexpr simd& operator|=(simd&, const simd&) noexcept;
friend constexpr simd& operator^=(simd&, const simd&) noexcept;
friend constexpr simd& operator<=<=(simd&, const simd&) noexcept;
friend constexpr simd& operator>=>=(simd&, const simd&) noexcept;
friend constexpr simd& operator<=<=(simd&, int) noexcept;
friend constexpr simd& operator>=>=(simd&, int) noexcept;
```

// 28.9.7.3, simd *compare operators*

```
friend constexpr mask_type operator==(const simd&, const simd&) noexcept;
friend constexpr mask_type operator!=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator>=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator<=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator>(const simd&, const simd&) noexcept;
friend constexpr mask_type operator<(const simd&, const simd&) noexcept;
};
```

¹ The class template `simd` is a data-parallel type. The width of a given `simd` specialization is a constant expression, determined by the template parameters.

- 2 Every specialization of `simd` is a complete type. The specialization `simd<T, Abi>` is supported if `T` is a vectorizable type and

- `Abi` is `simd_abi::scalar`, or
- `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in 28.9.3.

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd<T, Abi>` is supported. [*Note:* The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd<T, Abi>>`.

[*Example:* Consider an implementation that defines the extended ABI tags `__simd_x` and `__gpu_y`. When the compiler is invoked to translate to a machine that has support for the `__simd_x` ABI tag for all arithmetic types other than `long double` and no support for the `__gpu_y` ABI tag, then:

- `simd<T, simd_abi::__gpu_y>` is not supported for any `T` and has a deleted constructor.
- `simd<long double, simd_abi::__simd_x>` is not supported and has a deleted constructor.
- `simd<double, simd_abi::__simd_x>` is supported.
- `simd<long double, simd_abi::scalar>` is supported.

— *end example*]

- 3 Default initialization performs no initialization of the elements; value-initialization initializes each element with `T()`. [*Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- 4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd`:

```
constexpr explicit operator implementation-defined() const;
constexpr explicit simd(const implementation-defined& init);
```

[*Example:* Consider an implementation that supports the type `__vec4f` and the function `__vec4f _vec4f_addsub(__vec4f, __vec4f)` for the currently targeted system. A user may require the use of `_vec4f_addsub` for maximum performance and thus writes:

```
using V = simd<float, simd_abi::__simd128>;
V addsub(V a, V b) {
    return static_cast<V>(_vec4f_addsub(static_cast<__vec4f>(a), static_cast<__vec4f>(b)));
}
```

— *end example*]

(6.1.1.6.2) 28.9.6.2 `simd` width

[`simd.width`]

```
static constexpr typename simd_size<T, Abi>::type size;
```

- 1 *Returns:* The width of `simd<T, Abi>`.

(6.1.1.6.3) 28.9.6.3 Element references

[simd.reference]

- 1 A reference is an object that refers to an element in a `simd` or `simd_mask` object. `reference::value_type` is the same type as `simd::value_type` or `simd_mask::value_type`, respectively.
- 2 Class `reference` is for exposition only. An implementation is permitted to provide equivalent functionality without providing a class with this name.

```

class reference // exposition only
{
public:
    reference() = delete;
    reference(const reference&) = delete;

    constexpr operator value_type() const noexcept;

    template<class U> constexpr reference operator=(U&& x) && noexcept;

    template<class U> constexpr reference operator+=(U&& x) && noexcept;
    template<class U> constexpr reference operator-=(U&& x) && noexcept;
    template<class U> constexpr reference operator*=(U&& x) && noexcept;
    template<class U> constexpr reference operator/=(U&& x) && noexcept;
    template<class U> constexpr reference operator%=(U&& x) && noexcept;
    template<class U> constexpr reference operator|=(U&& x) && noexcept;
    template<class U> constexpr reference operator&=(U&& x) && noexcept;
    template<class U> constexpr reference operator^=(U&& x) && noexcept;
    template<class U> constexpr reference operator<<=(U&& x) && noexcept;
    template<class U> constexpr reference operator>>=(U&& x) && noexcept;

    constexpr reference operator++() && noexcept;
    constexpr value_type operator++(int) && noexcept;
    constexpr reference operator--() && noexcept;
    constexpr value_type operator--(int) && noexcept;

    friend constexpr void swap(reference&& a, reference&& b) noexcept;
    friend constexpr void swap(value_type& a, reference&& b) noexcept;
    friend constexpr void swap(reference&& a, value_type& b) noexcept;
};

constexpr operator value_type() const noexcept;

```

- 3 **Returns:** The value of the element referred to by `*this`.

```

template<class U> constexpr reference operator=(U&& x) && noexcept;

```

- 4 **Effects:** Replaces the referred to element in `simd` or `simd_mask` with `static_cast<value_type>(std::forward<U>(x))`.
- 5 **Returns:** A copy of `*this`.
- 6 **Remarks:** This function shall not participate in overload resolution unless `declval<value_type>() = std::forward<U>(x)` is well-formed.

```

template<class U> constexpr reference operator+=(U&& x) && noexcept;
template<class U> constexpr reference operator-=(U&& x) && noexcept;
template<class U> constexpr reference operator*=(U&& x) && noexcept;
template<class U> constexpr reference operator/=(U&& x) && noexcept;
template<class U> constexpr reference operator%=(U&& x) && noexcept;
template<class U> constexpr reference operator|=(U&& x) && noexcept;
template<class U> constexpr reference operator&=(U&& x) && noexcept;
template<class U> constexpr reference operator^=(U&& x) && noexcept;
template<class U> constexpr reference operator<=(U&& x) && noexcept;
template<class U> constexpr reference operator>=(U&& x) && noexcept;

```

7 *Effects:* Applies the indicated compound operator to the referred to element in `simd` or `simd_mask` and `std::forward<U>(x)`.

8 *Returns:* A copy of `*this`.

9 *Remarks:* This function shall not participate in overload resolution unless `declval<value_type&&>() @= std::forward<U>(x)` (where `@=` denotes the indicated compound assignment operator) is well-formed.

```

constexpr reference operator++() && noexcept;
constexpr reference operator--() && noexcept;

```

10 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.

11 *Returns:* A copy of `*this`.

12 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```

constexpr value_type operator++(int) && noexcept;
constexpr value_type operator--(int) && noexcept;

```

13 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.

14 *Returns:* A copy of the referred to element before applying the indicated operator.

15 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```

friend constexpr void swap(reference&& a, reference&& b) noexcept;
friend constexpr void swap(value_type& a, reference&& b) noexcept;
friend constexpr void swap(reference&& a, value_type& b) noexcept;

```

16 *Effects:* Exchanges the values `a` and `b` refer to.

(6.1.1.6.4) 28.9.6.4 `simd` constructors

[`simd.ctor`]

```

template<class U> constexpr simd(U&&) noexcept;

```

1 *Effects:* Constructs an object with each element initialized to the value of the argument after conversion to `value_type`.

2 *Remarks:* Let `From` denote the type `remove_cvref_t<U>`. This constructor shall not participate in overload resolution unless:

- `From` is a vectorizable type and every possibly value of `From` can be represented with type `value_type`, or
- `From` is not an arithmetic type and is implicitly convertible to `value_type`, or
- `From` is `int`, or
- `From` is `unsigned int` and `is_unsigned_v<value_type>` is true.

```
template<class U, class UAbi> constexpr explicit(see below) simd(const simd<U, UAbi>& x) noexcept;
```

3 *Effects:* Constructs an object where the i^{th} element equals `static_cast<T>(x[i])` for all i in the range of `[0, size())`.

4 *Remarks:* This constructor shall not participate in overload resolution unless `simd_size_v<U, UAbi> == size()`.

5 The constructor is `explicit` iff

- at least one possible value of `U` cannot be represented with type `value_type`, or
- if both `U` and `value_type` are integral types, the integer conversion rank (??) of `U` is greater than the integer conversion rank of `value_type`, or
- if both `U` and `value_type` are floating-point types, the floating-point conversion rank (??) of `U` is greater than the floating-point conversion rank of `value_type`.

```
template<class G> constexpr simd(G&& gen) noexcept;
```

6 *Effects:* Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.

7 *Remarks:* This constructor shall not participate in overload resolution unless `simd(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.

8 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr simd(const It& first, Flags = {});
```

9 *Constraints:* `iter_value_t<It>` is a vectorizable type.

10 *Requires:*

- `[first, first + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.

- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

11

Effects: Constructs an object where the i^{th} element is initialized to `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.

12

Remarks: This constructor shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true.

(6.1.1.6.5) 28.9.6.5 `simd` copy functions[`simd.copy`]

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr void copy_from(const It& first, Flags f = {});
```

1

Constraints: `iter_value_t<It>` is a vectorizable type.

2

Requires:

- `[mem, mem + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

3

Effects: Replaces the elements of the `simd` object such that the i^{th} element is assigned with `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.

4

Remarks: This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true.

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
requires output_iterator<It, value_type>
constexpr void copy_to(const It& first, Flags f = {}) const;
```

5

Constraints: `iter_value_t<It>` is a vectorizable type.

6

Requires:

- `[mem, mem + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

7

Effects: Copies all `simd` elements as if `mem[i] = static_cast<U>(operator[](i))` for all i in the range of `[0, size())`.

8

Remarks: This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true.

(6.1.1.6.6) 28.9.6.6 simd subscript operators

[simd.subscr]

constexpr reference operator[](size_t i);

1 *Requires:* $i < \text{size}()$.2 *Returns:* A reference (see 28.9.6.3) referring to the i^{th} element.3 *Throws:* Nothing.

constexpr value_type operator[](size_t i) const;

4 *Requires:* $i < \text{size}()$.5 *Returns:* The value of the i^{th} element.6 *Throws:* Nothing.

(6.1.1.6.7) 28.9.6.7 simd unary operators

[simd.unary]

1 Effects in this subclause are applied as unary element-wise operations.

constexpr simd& operator++() noexcept;

2 *Effects:* Increments every element by one.3 *Returns:* *this.

constexpr simd operator++(int) noexcept;

4 *Effects:* Increments every element by one.5 *Returns:* A copy of *this before incrementing.

constexpr simd& operator--() noexcept;

6 *Effects:* Decrements every element by one.7 *Returns:* *this.

constexpr simd operator--(int) noexcept;

8 *Effects:* Decrements every element by one.9 *Returns:* A copy of *this before decrementing.

constexpr mask_type operator!() const noexcept;

10 *Returns:* A simd_mask object with the i^{th} element set to !operator[](i) for all i in the range of $[0, \text{size}())$.

```
constexpr simd operator~() const noexcept;
```

11 *Returns:* A simd object where each bit is the inverse of the corresponding bit in *this.

12 *Remarks:* This operator shall not participate in overload resolution unless τ is an integral type.

```
constexpr simd operator+() const noexcept;
```

13 *Returns:* *this.

```
constexpr simd operator-() const noexcept;
```

14 *Returns:* A simd object where the i^{th} element is initialized to `-operator[]( $i$ )` for all i in the range of `[0, size())`.

(6.1.1.7) 28.9.7 simd non-member operations [simd.nonmembers]

(6.1.1.7.1) 28.9.7.1 simd binary operators [simd.binary]

```
friend constexpr simd operator+(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator-(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator*(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator/(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator%(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator&(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator|(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator^(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator<<(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator>>(const simd& lhs, const simd& rhs) noexcept;
```

1 *Returns:* A simd object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

2 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend constexpr simd operator<<(const simd& v, int n) noexcept;
```

```
friend constexpr simd operator>>(const simd& v, int n) noexcept;
```

3 *Returns:* A simd object where the i^{th} element is initialized to the result of applying the indicated operator to `v[ $i$ ]` and `n` for all i in the range of `[0, size())`.

4 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

(6.1.1.7.2) 28.9.7.2 simd compound assignment [simd.cassign]

```

friend constexpr simd& operator+=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator-=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator*=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator/=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator%=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator&=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator|=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator^=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator<=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator>=(simd& lhs, const simd& rhs) noexcept;

```

- 1 *Effects:* These operators apply the indicated operator to lhs and rhs as an element-wise operation.
- 2 *Returns:* lhs.
- 3 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```

friend constexpr simd& operator<=(simd& lhs, int n) noexcept;
friend constexpr simd& operator>=(simd& lhs, int n) noexcept;

```

- 4 *Effects:* Equivalent to: `return operator@=(lhs, simd(n));`
- 5 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

(6.1.1.7.3) 28.9.7.3 simd compare operators

[simd.comparison]

```

friend constexpr mask_type operator==(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator!=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator>=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator<=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator>(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator<(const simd& lhs, const simd& rhs) noexcept;

```

- 1 *Returns:* A `simd_mask` object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(6.1.1.7.4) 28.9.7.4 simd reductions

[simd.reductions]

- 1 In this subclause, `BinaryOperation` shall be a binary element-wise operation.

```

template<class T, class Abi, class BinaryOperation = plus<>>
constexpr T reduce(const simd<T, Abi>& x, BinaryOperation binary_op = {});

```

- 2 *Requires:* `binary_op` shall be callable with two arguments of type `T` returning `T`, or callable with two arguments of type `simd<T, A1>` returning `simd<T, A1>` for every `A1` that is an ABI tag type.
- 3 *Returns:* `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all `i` in the range of `[0, size())`(??).
- 4 *Throws:* Any exception thrown from `binary_op`.

```
template<class M, class V, class BinaryOperation>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
                                       typename V::value_type identity_element,
                                       BinaryOperation binary_op = {});
```

5 *Requires:* `binary_op` shall be callable with two arguments of type `T` returning `T`, or callable with two arguments of type `simd<T, A1>` returning `simd<T, A1>` for every `A1` that is an ABI tag type. The results of `binary_op(identity_element, x)` and `binary_op(x, identity_element)` shall be equal to `x` for all finite values `x` representable by `V::value_type`.

6 *Returns:* If `none_of(x.mask)`, returns `identity_element`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

7 *Throws:* Any exception thrown from `binary_op`.

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, plus<> binary_op) noexcept;
```

8 *Returns:* If `none_of(x.mask)`, returns `0`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, multiplies<> binary_op) noexcept;
```

9 *Returns:* If `none_of(x.mask)`, returns `1`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_and<> binary_op) noexcept;
```

10 *Requires:* `is_integral_v<V::value_type>` is true.

11 *Returns:* If `none_of(x.mask)`, returns `~V::value_type()`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_or<> binary_op) noexcept;
```

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_xor<> binary_op) noexcept;
```

12 *Requires:* `is_integral_v<V::value_type>` is true.

13 *Returns:* If `none_of(x.mask)`, returns `0`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices `i`.

```
template<class T, class Abi> constexpr T hmin(const simd<T, Abi>& x) noexcept;
```

14 *Returns:* The value of an element `x[j]` for which `x[j] <= x[i]` for all `i` in the range of `[0, size())`.

```
template<class M, class V> constexpr typename V::value_type hmin(const const_where_expression<M, V>& x) noexcept;
```

- 15 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::max()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] <= x.data[i]` for all selected indices `i`.

```
template<class T, class Abi> constexpr T hmax(const simd<T, Abi>& x) noexcept;
```

- 16 *Returns:* The value of an element `x[j]` for which `x[j] >= x[i]` for all `i` in the range of `[0, size())`.

```
template<class M, class V> constexpr typename V::value_type hmax(const const_where_expression<M, V>& x) noexcept;
```

- 17 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::lowest()`. Otherwise, returns the value of an element `x.data[j]` for which `x.mask[j] == true` and `x.data[j] >= x.data[i]` for all selected indices `i`.

(6.1.1.7.5) 28.9.7.5 Casts

[simd.casts]

```
template<size_t... Sizes, class T, class Abi>
constexpr tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
split(const simd<T, Abi>& x) noexcept;
template<size_t... Sizes, class T, class Abi>
constexpr tuple<simd_mask<T, simd_abi::deduce_t<T, Sizes>>...>
split(const simd_mask<T, Abi>& x) noexcept;
```

- 1 *Returns:* A tuple of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} tuple element initialized to the value of the element `x` with index $i + \text{sum of the first } j \text{ values in the Sizes pack}$.
- 2 *Remarks:* These functions shall not participate in overload resolution unless the sum of all values in the `Sizes` pack is equal to `simd_size_v<T, Abi>`.

```
template<class V, class Abi>
constexpr array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
split(const simd<typename V::value_type, Abi>& x) noexcept;
template<class V, class Abi>
constexpr array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
split(const simd_mask<typename V::simd_type::value_type, Abi>& x) noexcept;
```

- 3 *Returns:* An array of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} array element initialized to the value of the element in `x` with index $i + j * V::size()$.
- 4 *Remarks:* These functions shall not participate in overload resolution unless either:
- `is_simd_v<V>` is true and `simd_size_v<typename V::value_type, Abi>` is an integral multiple of `V::size()`, or
 - `is_simd_mask_v<V>` is true and `simd_size_v<typename V::simd_type::value_type, Abi>` is an integral multiple of `V::size()`.

```

template<size_t N, class T, class A>
constexpr array<resize_simd<simd_size_v<T, A> / N, simd<T, A>>, N>
    split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
constexpr array<resize_simd<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
    split_by(const simd_mask<T, A>& x) noexcept;

```

- 5 *Returns:* An array `arr`, where `arr[i][j]` is initialized by `x[i * (simd_size_v<T, A> / N) + j]`.
- 6 *Remarks:* The functions shall not participate in overload resolution unless `simd_size_v<T, A>` is an integral multiple of `N`.

```

template<class T, class... Abis>
constexpr simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
    const simd<T, Abis>&... xs) noexcept;
template<class T, class... Abis>
constexpr simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
    const simd_mask<T, Abis>&... xs) noexcept;

```

- 7 *Returns:* A data-parallel object initialized with the concatenated values in the `xs` pack of data-parallel objects: The i^{th} `simd/simd_mask` element of the j^{th} parameter in the `xs` pack is copied to the return value's element with index i + the sum of the width of the first j parameters in the `xs` pack.

```

template<class T, class Abi, size_t N>
constexpr resize_simd<simd_size_v<T, Abi> * N, simd<T, Abi>>
    concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
constexpr resize_simd<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
    concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;

```

- 8 *Returns:* A data-parallel object, the i^{th} element of which is initialized by `arr[i / simd_size_v<T, Abi>][i % simd_size_v<T, Abi>]`.

(6.1.1.7.6) 28.9.7.6 Algorithms

[simd.alg]

```
template<class T, class Abi> constexpr simd<T, Abi> min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 1 *Returns:* The result of the element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())`.

```
template<class T, class Abi> constexpr simd<T, Abi> max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
```

- 2 *Returns:* The result of the element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())`.

```

template<class T, class Abi>
constexpr pair<simd<T, Abi>, simd<T, Abi>> minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;

```

- 3 *Returns:* A pair initialized with

- the result of element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())` in the first member, and
- the result of element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())` in the second member.

```
template<class T, class Abi> simd<T, Abi>
constexpr clamp(const simd<T, Abi>& v, const simd<T, Abi>& lo, const simd<T, Abi>& hi);
```

4 *Requires:* No element in `lo` shall be greater than the corresponding element in `hi`.

5 *Returns:* The result of element-wise application of `std::clamp(v[i], lo[i], hi[i])` for all i in the range of `[0, size())`.

(6.1.1.7.7) 28.9.7.7 `simd` math library

[`simd.math`]

1 For each set of overloaded functions within `<cmath>`, there shall be additional overloads sufficient to ensure that if any argument corresponding to a double parameter has type `simd<T, Abi>`, where `is_floating_point_v<T>` is true, then:

- All arguments corresponding to double parameters shall be convertible to `simd<T, Abi>`.
- All arguments corresponding to `double*` parameters shall be of type `simd<T, Abi>*`.
- All arguments corresponding to parameters of integral type `U` shall be convertible to `fixed_size_simd<U, simd_size_v<T, Abi>>`.
- All arguments corresponding to `U*`, where `U` is integral, shall be of type `fixed_size_simd<U, simd_size_v<T, Abi>>*`.
- If the corresponding return type is `double`, the return type of the additional overloads is `simd<T, Abi>`. Otherwise, if the corresponding return type is `bool`, the return type of the additional overload is `simd_mask<T, Abi>`. Otherwise, the return type is `fixed_size_simd<R, simd_size_v<T, Abi>>`, with `R` denoting the corresponding return type.

It is unspecified whether a call to these overloads with arguments that are all convertible to `simd<T, Abi>` but are not of type `simd<T, Abi>` is well-formed.

2 Each function overload produced by the above rules applies the indicated `<cmath>` function element-wise. For the mathematical functions, the results per element only need to be approximately equal to the application of the function which is overloaded for the element type.

3 The result is unspecified if a domain, pole, or range error occurs when the input argument(s) are applied to the indicated `<cmath>` function. [*Note:* Implementations are encouraged to follow the C specification (especially Annex F). — *end note*]

4 TODO: Allow `abs(simd<signed-integral>)`.

5 If `abs` is called with an argument of type `simd<X, Abi>` for which `is_unsigned_v<X>` is true, the program is ill-formed.

(6.1.1.8) 28.9.8 Class template `simd_mask`

[`simd.mask.class`]

(6.1.1.8.1) 28.9.8.1 Class template `simd_mask` overview

[`simd.mask.overview`]

```

template<class T, class Abi> class simd_mask {
public:
    using value_type = bool;
    using reference = see below;
    using simd_type = simd<T, Abi>;
    using abi_type = Abi;

    static constexpr typename simd_size<T, Abi>::type size;

    constexpr simd_mask() noexcept = default;

// 28.9.8.3, simd_mask constructors
    constexpr explicit simd_mask(value_type) noexcept;
    template<class U, class UAbi>
        constexpr explicit(sizeof(U) != sizeof(T)) simd_mask(const simd_mask<U, UAbi>&) noexcept;
    template<class G> constexpr explicit simd_mask(G&& gen) noexcept;
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        constexpr simd_mask(const It& first, Flags = {});

// 28.9.8.4, simd_mask copy functions
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        constexpr void copy_from(const It& first, Flags = {});
    template<contiguous_iterator It, class Flags = element_aligned_tag>
        requires output_iterator<It, value_type>
        constexpr void copy_to(const It& first, Flags = {}) const;

// 28.9.8.5, simd_mask subscript operators
    constexpr reference operator[](size_t);
    constexpr value_type operator[](size_t) const;

// 28.9.8.6, simd_mask unary operators
    constexpr simd_mask operator!() const noexcept;

// 28.9.9.1, simd_mask binary operators
    friend constexpr simd_mask operator&&(const simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask operator||(const simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask operator&(const simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask operator|(const simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask operator^(const simd_mask&, const simd_mask&) noexcept;

// 28.9.9.2, simd_mask compound assignment
    friend constexpr simd_mask& operator&=(simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask& operator|=(simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask& operator^=(simd_mask&, const simd_mask&) noexcept;

// 28.9.9.3, simd_mask comparisons

```

```

    friend constexpr simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
    friend constexpr simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
};

```

- 1 The class template `simd_mask` is a data-parallel type with the element type `bool`. The width of a given `simd_mask` specialization is a constant expression, determined by the template parameters. Specifically, `simd_mask<T, Abi>::size() == simd<T, Abi>::size()`.
- 2 Every specialization of `simd_mask` is a complete type. The specialization `simd_mask<T, Abi>` is supported if `T` is a vectorizable type and

- `Abi` is `simd_abi::scalar`, or
- `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in (28.9.3).

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd_mask<T, Abi>` is supported. [*Note:* The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd_mask<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd_mask<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd_mask<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd_mask<T, Abi>>`.

- 3 Default initialization performs no initialization of the elements; value-initialization initializes each element with `false`. [*Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- 4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd_mask`:

```

    constexpr explicit operator implementation_defined() const;
    constexpr explicit simd_mask(const implementation_defined& init) const;

```

- 5 The member type reference has the same interface as `simd<T, Abi>::reference`, except its `value_type` is `bool`. (28.9.6.3)

(6.1.1.8.2) 28.9.8.2 `simd_mask` width [`simd.mask.width`]

```

static constexpr typename simd_size<T, Abi>::type size;

```

- 1 *Returns:* The width of `simd<T, Abi>`.

(6.1.1.8.3) 28.9.8.3 `simd_mask` constructors [`simd.mask.ctor`]

```

constexpr explicit simd_mask(value_type x) noexcept;

```

- 1 *Effects:* Constructs an object with each element initialized to `x`.

```

template<class U, class UAbi>
constexpr explicit(sizeof(U) != sizeof(T)) simd_mask(const simd_mask<U, UAbi>& x) noexcept;

```

- 2 *Effects:* Constructs an object of type `simd_mask` where the i^{th} element equals `x[i]` for all i in the range of `[0, size())`.
 3 *Remarks:* This constructor shall not participate in overload resolution unless `simd_size_v<U, UAbi> == size()`.

```
template<class G> constexpr explicit simd_mask(G&& gen) noexcept;
```

- 4 *Effects:* Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.
 5 *Remarks:* This constructor shall not participate in overload resolution unless `static_cast<bool>(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.
 6 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr simd_mask(const It& first, Flags = {});
```

- 7 *Constraints:* `iter_value_t<It>` is of type `bool`.
 8 *Requires:*
- `[mem, mem + size())` is a valid range.
 - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd_mask>`.
 - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
 - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(value_type)`.
- 9 *Effects:* Constructs an object where the i^{th} element is initialized to `mem[i]` for all i in the range of `[0, size())`.
 10 *Throws:* Nothing.
 11 *Remarks:* This constructor shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is true.

(6.1.1.8.4) 28.9.8.4 `simd_mask` copy functions

[simd.mask.copy]

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr void copy_from(const It& first, Flags = {});
```

- 1 *Constraints:* `iter_value_t<It>` is of type `bool`.
 2 *Requires:*
- `[mem, mem + size())` is a valid range.
 - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd_mask>`.
 - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
 - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(value_type)`.

- 3 *Effects:* Replaces the elements of the `simd_mask` object such that the i^{th} element is replaced with `mem[i]` for all i in the range of `[0, size())`.
- 4 *Throws:* Nothing.
- 5 *Remarks:* This function shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is true.

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
requires output_iterator<It, value_type>
constexpr void copy_to(const It& first, Flags = {}) const;
```

- 6 *Constraints:* `iter_value_t<It>` is of type `bool`.
- 7 *Requires:*
- `[mem, mem + size())` is a valid range.
 - If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd_mask>`.
 - If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
 - If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(value_type)`.
- 8 *Effects:* Copies all `simd_mask` elements as if `mem[i] = operator[](i)` for all i in the range of `[0, size())`.
- 9 *Throws:* Nothing.
- 10 *Remarks:* This function shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is true.

(6.1.1.8.5) 28.9.8.5 `simd_mask` subscript operators [simd.mask.subscr]

```
constexpr reference operator[](size_t i);
```

- 1 *Requires:* $i < \text{size}()$.
- 2 *Returns:* A reference (see 28.9.6.3) referring to the i^{th} element.
- 3 *Throws:* Nothing.

```
constexpr value_type operator[](size_t i) const;
```

- 4 *Requires:* $i < \text{size}()$.
- 5 *Returns:* The value of the i^{th} element.
- 6 *Throws:* Nothing.

(6.1.1.8.6) 28.9.8.6 `simd_mask` unary operators [simd.mask.unary]

```
constexpr simd_mask operator!() const noexcept;
```

- 1 *Returns:* The result of the element-wise application of `operator!`.

(6.1.1.9) 28.9.9 Non-member operations [simd.mask.nonmembers]

(6.1.1.9.1) 28.9.9.1 simd_mask binary operators [simd.mask.binary]

```
friend constexpr simd_mask operator&&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask operator||(const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask operator& (const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask operator| (const simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask operator^ (const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

1 *Returns:* A simd_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(6.1.1.9.2) 28.9.9.2 simd_mask compound assignment [simd.mask.cassign]

```
friend constexpr simd_mask& operator&=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask& operator|=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask& operator^=(simd_mask& lhs, const simd_mask& rhs) noexcept;
```

1 *Effects:* These operators apply the indicated operator to lhs and rhs as a binary element-wise operation.

2 *Returns:* lhs.

(6.1.1.9.3) 28.9.9.3 simd_mask comparisons [simd.mask.comparison]

```
friend constexpr simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
```

1 *Returns:* A simd_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(6.1.1.9.4) 28.9.9.4 simd_mask reductions [simd.mask.reductions]

```
template<class T, class Abi> constexpr bool all_of(const simd_mask<T, Abi>& k) noexcept;
```

1 *Returns:* true if all boolean elements in k are true, false otherwise.

```
template<class T, class Abi> constexpr bool any_of(const simd_mask<T, Abi>& k) noexcept;
```

2 *Returns:* true if at least one boolean element in k is true, false otherwise.

```
template<class T, class Abi> constexpr bool none_of(const simd_mask<T, Abi>& k) noexcept;
```

3 *Returns:* true if none of the one boolean elements in k is true, false otherwise.

```
template<class T, class Abi> constexpr bool some_of(const simd_mask<T, Abi>& k) noexcept;
```

- 4 *Returns:* true if at least one of the one boolean elements in *k* is true and at least one of the boolean elements in *k* is false, false otherwise.

```
template<class T, class Abi> constexpr int popcount(const simd_mask<T, Abi>& k) noexcept;
```

- 5 *Returns:* The number of boolean elements in *k* that are true.

```
template<class T, class Abi> constexpr int find_first_set(const simd_mask<T, Abi>& k);
```

- 6 *Requires:* any_of(*k*) returns true.
 7 *Returns:* The lowest element index *i* where *k*[*i*] is true.
 8 *Throws:* Nothing.

```
template<class T, class Abi> constexpr int find_last_set(const simd_mask<T, Abi>& k);
```

- 9 *Requires:* any_of(*k*) returns true.
 10 *Returns:* The greatest element index *i* where *k*[*i*] is true.
 11 *Throws:* Nothing.

```
constexpr bool all_of(T) noexcept;  
constexpr bool any_of(T) noexcept;  
constexpr bool none_of(T) noexcept;  
constexpr bool some_of(T) noexcept;  
constexpr int popcount(T) noexcept;
```

- 12 *Returns:* all_of and any_of return their arguments; none_of returns the negation of its argument; some_of returns false; popcount returns the integral representation of its argument.
 13 *Remarks:* The parameter type *T* is an unspecified type that is only constructible via implicit conversion from bool.

```
constexpr int find_first_set(T);  
constexpr int find_last_set(T);
```

- 14 *Requires:* The value of the argument is true.
 15 *Returns:* 0.
 16 *Throws:* Nothing.
 17 *Remarks:* The parameter type *T* is an unspecified type that is only constructible via implicit conversion from bool.

```

template<class T, class Abi>
  where_expression<simd_mask<T, Abi>, simd<T, Abi>>
    where(const typename simd<T, Abi>::mask_type& k, simd<T, Abi>& v) noexcept;
template<class T, class Abi>
  const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
    where(const typename simd<T, Abi>::mask_type& k, const simd<T, Abi>& v) noexcept;
template<class T, class Abi>
  where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
    where(const type_identity_t<simd_mask<T, Abi>>& k, simd_mask<T, Abi>& v) noexcept;
template<class T, class Abi>
  const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
    where(const type_identity_t<simd_mask<T, Abi>>& k, const simd_mask<T, Abi>& v) noexcept;

```

1 **Returns:** An object (28.9.5) with mask and data initialized with *k* and *v* respectively.

```

template<class T>
  where_expression<bool T>
    where(see below k, T& v) noexcept;
template<class T>
  const_where_expression<bool, T>
    where(see below k, const T& v) noexcept;

```

2 **Remarks:** The functions shall not participate in overload resolution unless

- *T* is neither a `simd` nor a `simd_mask` specialization, and
- the first argument is of type `bool`.

3 **Returns:** An object (28.9.5) with mask and data initialized with *k* and *v* respectively.

6.2

DIFF AGAINST PARALLELISM TS 2 (N4808)

In the following, the wording from Section 6.1 is repeated with additional indications of differences with regard to N4808. Changes relative to N4808, which contains editorial changes after the publication of the TS, are marked using color for **additions** and **removals**.

(6.2.1) 28.9 Data-Parallel Types [simd]

(6.2.1.1) 28.9.1 General [simd.general]

- 1 The data-parallel library consists of data-parallel types and operations on these types. A data-parallel type consists of elements of an underlying arithmetic type, called the *element type*. The number of elements is a constant for each data-parallel type and called the *width* of that type.
- 2 Throughout this Clause, the term *data-parallel type* refers to all *supported* (28.9.6.1) specializations of the `simd` and `simd_mask` class templates. A *data-parallel object* is an object of *data-parallel type*.

- 3 An *element-wise operation* applies a specified operation to the elements of one or more data-parallel objects. Each such application is unsequenced with respect to the others. A *unary element-wise operation* is an element-wise operation that applies a unary operation to each element of a data-parallel object. A *binary element-wise operation* is an element-wise operation that applies a binary operation to corresponding elements of two data-parallel objects.
- 4 Throughout this Clause, the set of *vectorizable types* for a data-parallel type comprises all cv-unqualified arithmetic types other than `bool`.
- 5 [*Note*: The intent is to support acceleration through data-parallel execution resources, such as SIMD registers and instructions or execution units driven by a common instruction decoder. If such execution resources are unavailable, the interfaces support a transparent fallback to sequential execution. — *end note*]

(6.2.1.2) 28.9.2 Header `<experimental/simd>` synopsis

[simd.synopsis]

```
namespace std::experimental { inline namespace parallelism_v2 {
    namespace simd_abi {
        using scalar = see below;
        template<class T, intsize_t N> using fixed_size = see below;
        template<class T> inline constexpr intsize_t max_fixed_size = implementation-defined;
        template<class T> using compatible = implementation-defined;
        template<class T> using native = implementation-defined;

        template<class T, size_t N, class... Abis> struct deduce { using type = see below; };
        template<class T, size_t N, class... Abis> using deduce_t =
            typename deduce<T, N, Abis...>::type;
    }

    struct element_aligned_tag {};
    struct vector_aligned_tag {};
    template<size_t> struct overaligned_tag {};
    inline constexpr element_aligned_tag element_aligned{};
    inline constexpr vector_aligned_tag vector_aligned{};
    template<size_t N> inline constexpr overaligned_tag<N> overaligned{};

    // 28.9.4, simd type traits
    template<class T> struct is_abi_tag;
    template<class T> inline constexpr bool is_abi_tag_v = is_abi_tag<T>::value;

    template<class T> struct is_simd;
    template<class T> inline constexpr bool is_simd_v = is_simd<T>::value;

    template<class T> struct is_simd_mask;
    template<class T> inline constexpr bool is_simd_mask_v = is_simd_mask<T>::value;

    template<class T> struct is_simd_flag_type;
    template<class T> inline constexpr bool is_simd_flag_type_v =
        is_simd_flag_type<T>::value;
}
```

```

template<class T, class Abi = simd_abi::compatible_native<T>> struct simd_size;
template<class T, class Abi = simd_abi::compatible_native<T>>
    inline constexpr size_t simd_size_v = simd_size<T,Abi>::value;

template<class T, class U = typename T::value_type> struct memory_alignment;
template<class T, class U = typename T::value_type>
    inline constexpr size_t memory_alignment_v = memory_alignment<T,U>::value;

template<class T, class V> struct rebind_simd { using type = see below; };
template<class T, class V> using rebind_simd_t = typename rebind_simd<T, V>::type;
template<int_size_t N, class V> struct resize_simd { using type = see below; };
template<int_size_t N, class V> using resize_simd_t = typename resize_simd<N, V>::type;

```

// 28.9.6, Class template simd

```

template<class T, class Abi = simd_abi::compatible_native<T>> class simd;
template<class T> using native_simd = simd<T, simd_abi::native<T>>
template<class T, int_size_t N> using fixed_size_simd = simd<T, simd_abi::fixed_size<T, N>>;

```

// 28.9.8, Class template simd_mask

```

template<class T, class Abi = simd_abi::compatible_native<T>> class simd_mask;
template<class T> using native_simd_mask = simd_mask<T, simd_abi::native<T>>
template<class T, int_size_t N> using fixed_size_simd_mask = simd_mask<T, simd_abi::fixed_size<T, N>>;

```

// 28.9.7.5, Casts

```

template<class T, class U, class Abi> see below simd_cast(const simd<U, Abi>&) noexcept;
template<class T, class U, class Abi> see below static_simd_cast(const simd<U, Abi>&) noexcept;

```

```

template<class T, class Abi>
    fixed_size_simd<T, simd_size_v<T, Abi>
        to_fixed_size(const simd<T, Abi>&) noexcept;
template<class T, class Abi>
    fixed_size_simd_mask<T, simd_size_v<T, Abi>
        to_fixed_size(const simd_mask<T, Abi>&) noexcept;
template<class T, int N>
    native_simd<T> to_native(const fixed_size_simd<T, N>&) noexcept;
template<class T, int N>
    native_simd_mask<T> to_native(const fixed_size_simd_mask<T, N>&) noexcept;
template<class T, int N>
    simd<T> to_compatible(const fixed_size_simd<T, N>&) noexcept;
template<class T, int N>
    simd_mask<T> to_compatible(const fixed_size_simd_mask<T, N>&) noexcept;

```

```

template<size_t... Sizes, class T, class Abi>
    constexpr tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
        split(const simd<T, Abi>&) noexcept;
template<size_t... Sizes, class T, class Abi>

```

```

    constexpr tuple<simd_mask<T, simd_mask_abi::deduce_t<T, Sizes>>...>
        split(const simd_mask<T, Abi>&) noexcept;
template<class V, class Abi>
    constexpr array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
        split(const simd<typename V::value_type, Abi>&) noexcept;
template<class V, class Abi>
    constexpr array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
        split(const simd_mask<typename V::simd_type::value_type, Abi>&) noexcept;

template<size_t N, class T, class A>
    constexpr array<resize_simd_t<simd_size_v<T, A> / N, simd<T, A>>, N>
        split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
    constexpr array<resize_simd_t<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
        split_by(const simd_mask<T, A>& x) noexcept;

template<class T, class... Abis>
    constexpr simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...)>>
        concat(const simd<T, Abis>&...) noexcept;
template<class T, class... Abis>
    constexpr simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...)>>
        concat(const simd_mask<T, Abis>&...) noexcept;

template<class T, class Abi, size_t N>
    constexpr resize_simd_t<simd_size_v<T, Abi> * N, simd<T, Abi>>
        concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
    constexpr resize_simd_t<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
        concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;

```

// 28.9.9.4, *R*simd_mask reductions

```

template<class T, class Abi> constexpr bool all_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool any_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool none_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr bool some_of(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr int popcount(const simd_mask<T, Abi>&) noexcept;
template<class T, class Abi> constexpr int find_first_set(const simd_mask<T, Abi>&);
template<class T, class Abi> constexpr int find_last_set(const simd_mask<T, Abi>&);

constexpr bool all_of(T) noexcept;
constexpr bool any_of(T) noexcept;
constexpr bool none_of(T) noexcept;
constexpr bool some_of(T) noexcept;
constexpr int popcount(T) noexcept;
constexpr int find_first_set(T);
constexpr int find_last_set(T);

```

// 28.9.5, Where expression class templates

```
template<class M, class T> class const_where_expression;
template<class M, class T> class where_expression;
```

// 28.9.9.5, Where functions

```
template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, simd<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>
        where(const typename simd<T, Abi>::mask_type&, const simd<T, Abi>&) noexcept;

template<class T, class Abi>
    where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, simd_mask<T, Abi>&) noexcept;

template<class T, class Abi>
    const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>
        where(const type_identity_t<simd_mask<T, Abi>>&, const simd_mask<T, Abi>&) noexcept;

template<class T>
    where_expression<bool, T>
        where(see below k, T& d) noexcept;

template<class T>
    const_where_expression<bool, T>
        where(see below k, const T& d) noexcept;
```

// 28.9.7.4, Rsimd reductions

```
template<class T, class Abi, class BinaryOperation = plus<>>
    constexpr T reduce(const simd<T, Abi>&,
        BinaryOperation = {});

template<class M, class V, class BinaryOperation>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        typename V::value_type identity_element,
        BinaryOperation binary_op);

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        plus<> binary_op = {}) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
        multiplies<> binary_op) noexcept;

template<class M, class V>
```

```

    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
                                              bit_and<> binary_op) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
                                              bit_or<> binary_op) noexcept;

template<class M, class V>
    constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
                                              bit_xor<> binary_op) noexcept;

template<class T, class Abi>
    constexpr T hmin(const simd<T, Abi>&) noexcept;
template<class M, class V>
    constexpr typename V::value_type hmin(const const_where_expression<M, V>&) noexcept;
template<class T, class Abi>
    constexpr T hmax(const simd<T, Abi>&) noexcept;
template<class M, class V>
    constexpr typename V::value_type hmax(const const_where_expression<M, V>&) noexcept;

// 28.9.7.6, Algorithms
template<class T, class Abi>
    constexpr simd<T, Abi>
        min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr simd<T, Abi>
        max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr pair<simd<T, Abi>, simd<T, Abi>>
        minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;
template<class T, class Abi>
    constexpr simd<T, Abi>
        clamp(const simd<T, Abi>& v,
              const simd<T, Abi>& lo,
              const simd<T, Abi>& hi);
}

```

- ¹ The header `<experimental/simd>` defines class templates, tag types, trait types, and function templates for element-wise operations on data-parallel objects.

(6.2.1.3) 28.9.3 simd ABI tags

[simd.abi]

```

namespace simd_abi {
    using scalar = see below;
    template<class T, intsize_t N> using fixed_size = see below;
    template<class T> inline constexpr intsize_t max_fixed_size = implementation-defined;
    template<class T> using compatible = implementation-defined;
    template<class T> using native = implementation-defined;
}

```

- 1 An *ABI tag* is a type in the `std::experimental::parallelism_v2::simd_abi` namespace that indicates a choice of size and binary representation for objects of data-parallel type. [*Note:* The intent is for the size and binary representation to depend on the target architecture. — *end note*] The ABI tag, together with a given element type implies a number of elements. ABI tag types are used as the second template argument to `simd` and `simd_mask`.
- 2 [*Note:* The ABI tag is orthogonal to selecting the machine instruction set. The selected machine instruction set limits the usable ABI tag types, though (see 28.9.6.1). The ABI tags enable users to safely pass objects of data-parallel type between translation unit boundaries (e.g. function calls or I/O). — *end note*]
- 3 `scalar` is an alias for an unspecified ABI tag that is different from `fixed_size<1>`. Use of the `scalar` tag type requires data-parallel types to store a single element (i.e., `simd<T, simd_abi::scalar>::size()` returns `simd_size_v<T, simd_abi::scalar> equals 1`).
- 4 The value of `max_fixed_size<T>` is at least 32.
- 5 `fixed_size<N>` is an alias for an unspecified ABI tag. `fixed_size` does not introduce a non-deduced context. Use of the `simd_abi::fixed_size<N>` tag type requires data-parallel types to store `N` elements (i.e. `simd<T, simd_abi::fixed_size<N>>::size()` is `simd_size_v<T, simd_abi::fixed_size<N> equals N`). `simd<T, fixed_size<N>>` and `simd_mask<T, fixed_size<N>>` with `N > 0` and `N <= max_fixed_size<T>` shall be supported. Additionally, for every supported `simd<T, Abi>` (see 28.9.6.1), where `Abi` is an ABI tag that is not a specialization of `simd_abi::fixed_size, N == simd<T, Abi>::size()` shall be supported.
- 6 [*Note:* It is unspecified whether `simd<T, fixed_size<N>>` with `N > max_fixed_size<T>` is supported. The value of `max_fixed_size<T>` can depend on compiler flags and can change between different compiler versions. — *end note*]
- 7 The type of `fixed_size<T, N>` in TU1 differs from the type of `fixed_size<T, N>` in TU2 iff the type of `native<T>` in TU1 differs from the type of `native<T>` in TU2.
- 8 ~~[*Note:* An implementation can forego ABI compatibility between differently compiled translation units for `simd` and `simd_mask` specializations using the same `simd_abi::fixed_size<N>` tag. Otherwise, the efficiency of `simd<T, Abi>` is likely to be better than for `simd<T, fixed_size<simd_size_v<T, Abi>>` (with `Abi` not a specialization of `simd_abi::fixed_size`). — *end note*]~~
- 9 An implementation may define additional *extended ABI tag* types in the `std::experimental::parallelism_v2::simd_abi` namespace, to support other forms of data-parallel computation.
- 10 ~~`compatible<T>` is an implementation-defined alias for an ABI tag. [*Note:* The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that ensures ABI compatibility between translation units on the target architecture. — *end note*] [*Example:* Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where the `__simd256` type requires an optional ISA extension on said architecture. Also, the target architecture does not support long double with either ABI tag. The implementation therefore defines `compatible<T>` as an alias for:~~
 - ~~• `scalar` if `T` is the same type as long double, and~~
 - ~~• `__simd128` otherwise.~~~~— *end example*]~~
- 11 `native<T>` is an implementation-defined alias for an ABI tag. [*Note:* The intent is to use the ABI tag producing the most efficient data-parallel execution for the element type `T` that is supported on the currently targeted system. ~~For target architectures without ISA extensions, the `native<T>` and `compatible<T>` aliases will likely be the same.~~ For target architectures with ISA extensions, compiler flags may influence change the type of the `native<T>` alias ~~while `compatible<T>` will be the same independent of such flags.~~ — *end note*] [*Example:* Consider a target architecture supporting the extended ABI tags `__simd128` and `__simd256`, where hardware support for `__simd256` only exists for floating-point types. The implementation therefore defines `native<T>` as an alias for
 - `__simd256` if `T` is a floating-point type, and

- `__simd128` otherwise.

— *end example*]

```
template<T, size_t N, class... Abis> struct deduce { using type = see below; };
```

- 12 The member type shall be present if and only if
- `T` is a vectorizable type, and
 - `simd_abi::fixed_size<N>` is supported (see 28.9.3), and
 - every type in the `Abis` pack is an ABI tag.
- 13 Where present, the member typedef `type` shall name an ABI tag type that satisfies
- `simd_size<T, type> == N`, and
 - `simd<T, type>` is default constructible (see 28.9.6.1).

If `N` is 1, the member typedef `type` is `simd_abi::scalar`. ~~Otherwise, if there are multiple ABI tag types that satisfy the constraints, the member typedef type is implementation-defined.~~ [*Note: It is expected that extended ABI tags can produce better optimizations and thus are preferred over `simd_abi::fixed_size<N>`.* Implementations can base the choice on `Abis`, but can also ignore the `Abis` arguments. — *end note*]

- 14 The behavior of a program that adds specializations for `deduce` is undefined.

(6.2.1.4) 28.9.4 `simd` type traits

[`simd.traits`]

```
template<class T> struct is_abi_tag { see below };
```

- 1 The type `is_abi_tag<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a standard or extended ABI tag, and `false_type` otherwise.
- 2 The behavior of a program that adds specializations for `is_abi_tag` is undefined.

```
template<class T> struct is_simd { see below };
```

- 3 The type `is_simd<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd` class template, and `false_type` otherwise.
- 4 The behavior of a program that adds specializations for `is_simd` is undefined.

```
template<class T> struct is_simd_mask { see below };
```

- 5 The type `is_simd_mask<T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is a specialization of the `simd_mask` class template, and `false_type` otherwise.
- 6 The behavior of a program that adds specializations for `is_simd_mask` is undefined.

```
template<class T> struct is_simd_flag_type { see below };
```

- 7 The type `is_simd_flag_type<class T>` is a `UnaryTypeTrait` with a base characteristic of `true_type` if `T` is one of
- `element_aligned_tag`, or

- `vector_aligned_tag`, or
- `overaligned_tag<N>` with $N > 0$ and N an integral power of two,

and `false_type` otherwise.

8 The behavior of a program that adds specializations for `is_simd_flag_type` is undefined.

```
template<class T, class Abi = simd_abi::compatible_native<T>> struct simd_size { see below };
```

9 `simd_size<T, Abi>` shall have a member `value` if and only if

- T is a vectorizable type, and
- `is_abi_tag_v<Abi>` is true.

[*Note:* The rules are different from those in (28.9.6.1): The member `value` is present even if `simd<T, Abi>` is not supported for the currently targeted system. — *end note*]

10 If `value` is present, the type `simd_size<T, Abi>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` with N equal to the number of elements in a `simd<T, Abi>` object. [*Note:* If `simd<T, Abi>` is not supported for the currently targeted system, `simd_size<T, Abi>::value` produces the value `simd<T, Abi>::size()` would return if it were supported. — *end note*]

11 The behavior of a program that adds specializations for `simd_size` is undefined.

```
template<class T, class U = typename T::value_type> struct memory_alignment { see below };
```

12 `memory_alignment<T, U>` shall have a member `value` if and only if

- `is_simd_mask_v<T>` is true and U is `bool`, or
- `is_simd_v<T>` is true and U is a vectorizable type.

13 If `value` is present, the type `memory_alignment<T, U>` is a `BinaryTypeTrait` with a base characteristic of `integral_constant<size_t, N>` for some implementation-defined N (see 28.9.6.5 and 28.9.8.4). [*Note:* `value` identifies the alignment restrictions on pointers used for (converting) loads and stores for the give type T on arrays of type U . — *end note*]

14 The behavior of a program that adds specializations for `memory_alignment` is undefined.

```
template<class T, class V> struct rebind_simd { using type = see below; };
```

15 The member type is present if and only if

- V is either `simd<U, Abi0>` or `simd_mask<U, Abi0>`, where U and $Abi0$ are deduced from V , and
- T is a vectorizable type, and
- `simd_abi::deduce<T, simd_size_v<U, Abi0>, Abi0>` has a member type type.

16 Let $Abi1$ denote the type `deduce_t<T, simd_size_v<U, Abi0>, Abi0>`. Where present, the member `typedef` type names `simd<T, Abi1>` if V is `simd<U, Abi0>` or `simd_mask<T, Abi1>` if V is `simd_mask<U, Abi0>`.

```
template<int size_t N, class V> struct resize_simd { using type = see below; };
```

17 The member type is present if and only if

- V is either `simd<T, Abi0>` or `simd_mask<T, Abi0>`, where T and $Abi0$ are deduced from V , and

- `simd_abi::deduce<T, N, Abi0>` has a member type `type`.

18 Let `Abi1` denote the type `deduce_t<T, N, Abi0>`. Where present, the member typedef `type` names `simd<T, Abi1>` if `V` is `simd<T, Abi0>` or `simd_mask<T, Abi1>` if `V` is `simd_mask<T, Abi0>`.

(6.2.1.5) 28.9.5 Where expression class templates

[simd.whereexpr]

```
template<class M, class T> class const_where_expression {
    const M mask;      // exposition only
    T& data;           // exposition only

public:
    const_where_expression(const const_where_expression&) = delete;
    const_where_expression& operator=(const const_where_expression&) = delete;

    T operator-() const && noexcept;
    T operator+() const && noexcept;
    T operator~() const && noexcept;

    template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags f = {}) const &&;
};

template<class M, class T>
class where_expression : public const_where_expression<M, T> {
public:
    template<class U> void operator=(U&& x) && noexcept;
    template<class U> void operator+=(U&& x) && noexcept;
    template<class U> void operator-=(U&& x) && noexcept;
    template<class U> void operator*=(U&& x) && noexcept;
    template<class U> void operator/=(U&& x) && noexcept;
    template<class U> void operator%=(U&& x) && noexcept;
    template<class U> void operator&=(U&& x) && noexcept;
    template<class U> void operator|=(U&& x) && noexcept;
    template<class U> void operator^=(U&& x) && noexcept;
    template<class U> void operator<<=(U&& x) && noexcept;
    template<class U> void operator>>=(U&& x) && noexcept;

    void operator++() && noexcept;
    void operator++(int) && noexcept;
    void operator--() && noexcept;
    void operator--(int) && noexcept;

    template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&;
};
```

- 1 The class templates `const_where_expression` and `where_expression` abstract the notion of selecting elements of a given object of arithmetic or data-parallel type.
- 2 The first templates argument `M` shall be cv-unqualified `bool` or a cv-unqualified `simd_mask` specialization.

- 3 If M is `bool`, T shall be a cv-unqualified arithmetic type. Otherwise, T shall either be M or typename `M::simd_type`.
- 4 In this subclause, if M is `bool`, `data[0]` is used interchangeably for `data`, `mask[0]` is used interchangeably for `mask`, and `M::size()` is used interchangeably for 1.
- 5 The *selected indices* signify the integers $i \in \{j \in \mathbb{N} | j < M::size() \wedge \text{mask}[j]\}$. The *selected elements* signify the elements `data[i]` for all selected indices i .
- 6 In this subclause, the type `value_type` is an alias for T if M is `bool`, or an alias for typename `T::value_type` if `is_simd_mask_v<M>` is true.
- 7 [*Note*: The `where` functions 28.9.9.5 initialize `mask` with the first argument to `where` and `data` with the second argument to `where`. — *end note*]

```
T operator-() const && noexcept;
T operator+() const && noexcept;
T operator~() const && noexcept;
```

- 8 **Returns**: A copy of `data` with the indicated unary operator applied to all selected elements.

```
template<class U, class Flags = element_aligned_tag> void copy_to(U* mem, Flags = {}) const &&
```

- 9 **Requires**:

- If M is not `bool`, the largest selected index is less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by N .
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

- 10 **Effects**: Copies the selected elements as if `mem[i] = static_cast<U>(data[i])` for all selected indices i .

- 11 **Throws**: Nothing.

- 12 **Remarks**: This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
 - U is `bool` and `value_type` is `bool`, or
 - U is a vectorizable type and `value_type` is not `bool`.

```
template<class U> void operator=(U&& x) && noexcept;
```

- 13 **Effects**: Replaces `data[i]` with `static_cast<T>(std::forward<U>(x))[i]` for all selected indices i .

- 14 **Remarks**: This operator shall not participate in overload resolution unless U is convertible to T .

```
template<class U> void operator+=(U&& x) && noexcept;
template<class U> void operator-=(U&& x) && noexcept;
template<class U> void operator*=(U&& x) && noexcept;
template<class U> void operator/=(U&& x) && noexcept;
template<class U> void operator%=(U&& x) && noexcept;
template<class U> void operator&=(U&& x) && noexcept;
```

```
template<class U> void operator|=(U&& x) && noexcept;
template<class U> void operator^=(U&& x) && noexcept;
template<class U> void operator<<=(U&& x) && noexcept;
template<class U> void operator>>=(U&& x) && noexcept;
```

15 *Effects:* Replaces `data[i]` with `static_cast<T>(data @ std::forward<U>(x))[i]` (where @ denotes the indicated operator) for all selected indices *i*.

16 *Remarks:* Each of these operators shall not participate in overload resolution unless the return type of `data @ std::forward<U>(x)` is convertible to `T`. It is unspecified whether the binary operator, implied by the compound assignment operator, is executed on all elements or only on the selected elements.

```
void operator++() && noexcept;
void operator++(int) && noexcept;
void operator--() && noexcept;
void operator--(int) && noexcept;
```

17 *Effects:* Applies the indicated operator to the selected elements.

18 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `T`.

```
template<class U, class Flags = element_aligned_tag> void copy_from(const U* mem, Flags = {}) &&
```

19 *Requires:*

- If `is_simd_flag_type_v<U>` is true, for all selected indices *i*, *i* shall be less than the number of values pointed to by `mem`.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<T, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

20 *Effects:* Replaces the selected elements as if `data[i] = static_cast<value_type>(mem[i])` for all selected indices *i*.

21 *Throws:* Nothing.

22 *Remarks:* This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, and
- either
 - `U` is `bool` and `value_type` is `bool`, or
 - `U` is a vectorizable type and `value_type` is not `bool`.

(6.2.1.6) 28.9.6 Class template `simd` [simd.class]

(6.2.1.6.1) 28.9.6.1 Class template `simd` overview [simd.overview]

```

template<class T, class Abi> class simd {
public:
    using value_type = T;
    using reference = see below;
    using mask_type = simd_mask<T, Abi>;
    using abi_type = Abi;

    static constexpr size_t size() noexcept typename simd_size<T, Abi>::type size;

    constexpr simd() noexcept = default;

// 28.9.6.4, simd constructors
template<class U> constexpr simd(U&& value) noexcept;
template<class U, class UAbi>
    constexpr explicit(see below) simd(const simd<U, simd_abi::fixed_size<size()>UAbi>&) noexcept;
template<class G> constexpr explicit simd(G&& gen) noexcept;
template<class Ucontiguous_iterator It, class Flags = element_aligned_tag>
    constexpr simd(const U*memIt& first, Flags = {});

// 28.9.6.5, simd copy functions
template<class Ucontiguous_iterator It, class Flags = element_aligned_tag>
    constexpr void copy_from(const U*memIt& first, Flags f = {});
template<class Ucontiguous_iterator It, class Flags = element_aligned_tag>
    requires output_iterator<It, value_type>
    constexpr void copy_to(U*memconst It& first, Flags f = {}) const;

// 28.9.6.6, simd subscript operators
constexpr reference operator[](size_t);
constexpr value_type operator[](size_t) const;

// 28.9.6.7, simd unary operators
constexpr simd& operator++() noexcept;
constexpr simd operator++(int) noexcept;
constexpr simd& operator--() noexcept;
constexpr simd operator--(int) noexcept;
constexpr mask_type operator!() const noexcept;
constexpr simd operator~() const noexcept;
constexpr simd operator+() const noexcept;
constexpr simd operator-() const noexcept;

// 28.9.7.1, simd binary operators
friend constexpr simd operator+(const simd&, const simd&) noexcept;
friend constexpr simd operator-(const simd&, const simd&) noexcept;
friend constexpr simd operator*(const simd&, const simd&) noexcept;
friend constexpr simd operator/(const simd&, const simd&) noexcept;
friend constexpr simd operator%(const simd&, const simd&) noexcept;

```

```

friend constexpr simd operator&(const simd&, const simd&) noexcept;
friend constexpr simd operator|(const simd&, const simd&) noexcept;
friend constexpr simd operator^(const simd&, const simd&) noexcept;
friend constexpr simd operator<<(const simd&, const simd&) noexcept;
friend constexpr simd operator>>(const simd&, const simd&) noexcept;
friend constexpr simd operator<<(const simd&, int) noexcept;
friend constexpr simd operator>>(const simd&, int) noexcept;

```

// 28.9.7.2, simd compound assignment

```

friend constexpr simd& operator+=(simd&, const simd&) noexcept;
friend constexpr simd& operator-=(simd&, const simd&) noexcept;
friend constexpr simd& operator*=(simd&, const simd&) noexcept;
friend constexpr simd& operator/=(simd&, const simd&) noexcept;
friend constexpr simd& operator%=(simd&, const simd&) noexcept;
friend constexpr simd& operator&=(simd&, const simd&) noexcept;
friend constexpr simd& operator|=(simd&, const simd&) noexcept;
friend constexpr simd& operator^=(simd&, const simd&) noexcept;
friend constexpr simd& operator<=<=(simd&, const simd&) noexcept;
friend constexpr simd& operator>>=(simd&, const simd&) noexcept;
friend constexpr simd& operator<=<=(simd&, int) noexcept;
friend constexpr simd& operator>>=(simd&, int) noexcept;

```

// 28.9.7.3, simd compare operators

```

friend constexpr mask_type operator==(const simd&, const simd&) noexcept;
friend constexpr mask_type operator!=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator>=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator<=(const simd&, const simd&) noexcept;
friend constexpr mask_type operator>(const simd&, const simd&) noexcept;
friend constexpr mask_type operator<(const simd&, const simd&) noexcept;
};

```

- 1 The class template `simd` is a data-parallel type. The width of a given `simd` specialization is a constant expression, determined by the template parameters.
- 2 Every specialization of `simd` ~~shall be~~ is a complete type. The specialization `simd<T, Abi>` is supported if `T` is a vectorizable type and
 - `Abi` is `simd_abi::scalar`, or
 - `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in 28.9.3.

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd<T, Abi>` is supported. [*Note*: The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd<T, Abi>>`.

[*Example:* Consider an implementation that defines the extended ABI tags `__simd_x` and `__gpu_y`. When the compiler is invoked to translate to a machine that has support for the `__simd_x` ABI tag for all arithmetic types other than `long double` and no support for the `__gpu_y` ABI tag, then:

- `simd<T, simd_abi::__gpu_y>` is not supported for any `T` and has a deleted constructor.
- `simd<long double, simd_abi::__simd_x>` is not supported and has a deleted constructor.
- `simd<double, simd_abi::__simd_x>` is supported.
- `simd<long double, simd_abi::scalar>` is supported.

— *end example*]

- 3 Default initialization performs no initialization of the elements; value-initialization initializes each element with `τ()`. [*Note:* Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- 4 Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd`:

```
constexpr explicit operator implementation-defined() const;
constexpr explicit simd(const implementation-defined& init);
```

[*Example:* Consider an implementation that supports the type `__vec4f` and the function `__vec4f _vec4f_addsub(__vec4f, __vec4f)` for the currently targeted system. A user may require the use of `_vec4f_addsub` for maximum performance and thus writes:

```
using V = simd<float, simd_abi::__simd128>;
V addsub(V a, V b) {
    return static_cast<V>(_vec4f_addsub(static_cast<__vec4f>(a), static_cast<__vec4f>(b)));
}
```

— *end example*]

(6.2.1.6.2) 28.9.6.2 `simd` width [`simd.width`]

```
static constexpr size_t size() noexcept typename simd_size<T, Abi>::type size;
```

- 1 *Returns:* The width of `simd<T, Abi>`.

(6.2.1.6.3) 28.9.6.3 Element references [`simd.reference`]

- 1 A reference is an object that refers to an element in a `simd` or `simd_mask` object. `reference::value_type` is the same type as `simd::value_type` or `simd_mask::value_type`, respectively.
- 2 Class `reference` is for exposition only. An implementation is permitted to provide equivalent functionality without providing a class with this name.

```
class reference // exposition only
{
public:
    reference() = delete;
    reference(const reference&) = delete;

    constexpr operator value_type() const noexcept;
```

```
template<class U> constexpr reference operator=(U&& x) && noexcept;

template<class U> constexpr reference operator+=(U&& x) && noexcept;
template<class U> constexpr reference operator-=(U&& x) && noexcept;
template<class U> constexpr reference operator*=(U&& x) && noexcept;
template<class U> constexpr reference operator/=(U&& x) && noexcept;
template<class U> constexpr reference operator%=(U&& x) && noexcept;
template<class U> constexpr reference operator|=(U&& x) && noexcept;
template<class U> constexpr reference operator&=(U&& x) && noexcept;
template<class U> constexpr reference operator^=(U&& x) && noexcept;
template<class U> constexpr reference operator<<=(U&& x) && noexcept;
template<class U> constexpr reference operator>>=(U&& x) && noexcept;
```

```
constexpr reference operator++() && noexcept;
constexpr value_type operator++(int) && noexcept;
constexpr reference operator--() && noexcept;
constexpr value_type operator--(int) && noexcept;
```

```
friend constexpr void swap(reference&& a, reference&& b) noexcept;
friend constexpr void swap(value_type& a, reference&& b) noexcept;
friend constexpr void swap(reference&& a, value_type& b) noexcept;
};
```

```
constexpr operator value_type() const noexcept;
```

- 3 *Returns:* The value of the element referred to by *this.

```
template<class U> constexpr reference operator=(U&& x) && noexcept;
```

- 4 *Effects:* Replaces the referred to element in `simd` or `simd_mask` with `static_cast<value_type>(std::forward<U>(x))`.

- 5 *Returns:* A copy of *this.

- 6 *Remarks:* This function shall not participate in overload resolution unless `declval<value_type>() = std::forward<U>(x)` is well-formed.

```
template<class U> constexpr reference operator+=(U&& x) && noexcept;
template<class U> constexpr reference operator-=(U&& x) && noexcept;
template<class U> constexpr reference operator*=(U&& x) && noexcept;
template<class U> constexpr reference operator/=(U&& x) && noexcept;
template<class U> constexpr reference operator%=(U&& x) && noexcept;
template<class U> constexpr reference operator|=(U&& x) && noexcept;
template<class U> constexpr reference operator&=(U&& x) && noexcept;
template<class U> constexpr reference operator^=(U&& x) && noexcept;
template<class U> constexpr reference operator<<=(U&& x) && noexcept;
template<class U> constexpr reference operator>>=(U&& x) && noexcept;
```

- 7 *Effects:* Applies the indicated compound operator to the referred to element in `simd` or `simd_mask` and `std::forward<U>(x)`.

- 8 *Returns:* A copy of *this.
- 9 *Remarks:* This function shall not participate in overload resolution unless `declval<value_type&>() @= std::forward<U>(x)` (where @= denotes the indicated compound assignment operator) is well-formed.

```
constexpr reference operator++() && noexcept;
constexpr reference operator--() && noexcept;
```

- 10 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.
- 11 *Returns:* A copy of *this.
- 12 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
constexpr value_type operator++(int) && noexcept;
constexpr value_type operator--(int) && noexcept;
```

- 13 *Effects:* Applies the indicated operator to the referred to element in `simd` or `simd_mask`.
- 14 *Returns:* A copy of the referred to element before applying the indicated operator.
- 15 *Remarks:* This function shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend constexpr void swap(reference&& a, reference&& b) noexcept;
friend constexpr void swap(value_type& a, reference&& b) noexcept;
friend constexpr void swap(reference&& a, value_type& b) noexcept;
```

- 16 *Effects:* Exchanges the values `a` and `b` refer to.

(6.2.1.6.4) 28.9.6.4 `simd` constructors

[`simd.ctor`]

```
template<class U> constexpr simd(U&&) noexcept;
```

- 1 *Effects:* Constructs an object with each element initialized to the value of the argument after conversion to `value_type`.
- 2 *Remarks:* Let `From` denote the type `remove_cv_t<remove_reference_t<U>>remove_cvref_t<U>`. This constructor shall not participate in overload resolution unless:
- `From` is a vectorizable type and every possibly value of `From` can be represented with type `value_type`, or
 - `From` is not an arithmetic type and is implicitly convertible to `value_type`, or
 - `From` is `int`, or
 - `From` is unsigned `int` and `value_type is an unsigned integral type` `is_unsigned_v<value_type> is true`.

```
template<class U, class UAbi>
constexpr explicit(see below) simd(const simd<U, simd_abi::fixed_size<size()>UAbi& x) noexcept;
```

3 *Effects:* Constructs an object where the i^{th} element equals `static_cast<T>(x[i])` for all i in the range of `[0, size())`.

4 *Remarks:* This constructor shall not participate in overload resolution unless `simd_size_v<U, UAbi> == size()`.

- ~~`abi_type` is `simd_abi::fixed_size<size()>`, and~~
- ~~every possible value of `U` can be represented with type `value_type`, and~~
- ~~if both `U` and `value_type` are integral, the integer conversion rank (??) of `value_type` is greater than the integer conversion rank of `U`.~~

5 The constructor is explicit iff

- at least one possible value of `U` cannot be represented with type `value_type`, or
- if both `U` and `value_type` are integral types, the integer conversion rank (??) of `U` is greater than the integer conversion rank of `value_type`, or
- if both `U` and `value_type` are floating-point types, the floating-point conversion rank (??) of `U` is greater than the floating-point conversion rank of `value_type`.

```
template<class G> constexpr simd(G&& gen) noexcept;
```

6 *Effects:* Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.

7 *Remarks:* This constructor shall not participate in overload resolution unless `simd(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.

8 The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (??).

```
template<class U, contiguous_iterator It, class Flags = element_aligned_tag>
constexpr simd(const U* mem, It& first, Flags = {});
```

9 *Constraints:* `iter_value_t<It>` is a vectorizable type.

10 *Requires:*

- `[first, first + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

11 *Effects:* Constructs an object where the i^{th} element is initialized to `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.

12 *Remarks:* This constructor shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, ~~and~~
- ~~`U` is a vectorizable type.~~

(6.2.1.6.5) 28.9.6.5 `simd` copy functions

[`simd.copy`]

```
template<class U, contiguous_iterator It, class Flags = element_aligned_tag>
constexpr void copy_from(const U* mem, It& first, Flags f = {});
```

1 Constraints: `iter_value_t<It>` is a vectorizable type.

2 *Requires:*

- `[mem, mem + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

3 *Effects:* Replaces the elements of the `simd` object such that the i^{th} element is assigned with `static_cast<T>(mem[i])` for all i in the range of `[0, size())`.

4 *Remarks:* This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, ~~and~~
- ~~`U` is a vectorizable type.~~

```
template<class U contiguous_iterator It, class Flags = element_aligned_tag>
requires output_iterator<It, value_type>
constexpr void copy_to(U*mem const It& first, Flags f = {}) const;
```

5 Constraints: `iter_value_t<It>` is a vectorizable type.

6 *Requires:*

- `[mem, mem + size())` is a valid range.
- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd, U>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(U)`.

7 *Effects:* Copies all `simd` elements as if `mem[i] = static_cast<U>(operator[](i))` for all i in the range of `[0, size())`.

8 *Remarks:* This function shall not participate in overload resolution unless

- `is_simd_flag_type_v<Flags>` is true, ~~and~~
- ~~`U` is a vectorizable type.~~

(6.2.1.6.6) 28.9.6.6 `simd` subscript operators

[simd.subscr]

```
constexpr reference operator[](size_t i);
```

1 *Requires:* $i < \text{size}()$.

2 *Returns:* A reference (see 28.9.6.3) referring to the i^{th} element.

3 *Throws:* Nothing.

```
constexpr value_type operator[](size_t i) const;
```

- 4 *Requires:* $i < \text{size}()$.
- 5 *Returns:* The value of the i^{th} element.
- 6 *Throws:* Nothing.

(6.2.1.6.7) 28.9.6.7 `simd` unary operators[`simd.unary`]

- 1 Effects in this subclause are applied as unary element-wise operations.

```
constexpr simd& operator++() noexcept;
```

- 2 *Effects:* Increments every element by one.
- 3 *Returns:* `*this`.

```
constexpr simd operator++(int) noexcept;
```

- 4 *Effects:* Increments every element by one.
- 5 *Returns:* A copy of `*this` before incrementing.

```
constexpr simd& operator--() noexcept;
```

- 6 *Effects:* Decrements every element by one.
- 7 *Returns:* `*this`.

```
constexpr simd operator--(int) noexcept;
```

- 8 *Effects:* Decrements every element by one.
- 9 *Returns:* A copy of `*this` before decrementing.

```
constexpr mask_type operator!() const noexcept;
```

- 10 *Returns:* A `simd_mask` object with the i^{th} element set to `!operator[](i)` for all i in the range of $[0, \text{size}())$.

```
constexpr simd operator~() const noexcept;
```

- 11 *Returns:* A `simd` object where each bit is the inverse of the corresponding bit in `*this`.
- 12 *Remarks:* This operator shall not participate in overload resolution unless `T` is an integral type.

```
constexpr simd operator+() const noexcept;
```

- 13 *Returns:* `*this`.

```
constexpr simd operator-() const noexcept;
```

- 14 *Returns:* A `simd` object where the i^{th} element is initialized to `-operator[](i)` for all i in the range of $[0, \text{size}())$.

(6.2.1.7) 28.9.7 `simd` non-member operations [simd.nonmembers]

(6.2.1.7.1) 28.9.7.1 `simd` binary operators [simd.binary]

```
friend constexpr simd operator+(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator-(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator*(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator/(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator%(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator&(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator|(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator^(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator<<(const simd& lhs, const simd& rhs) noexcept;
friend constexpr simd operator>>(const simd& lhs, const simd& rhs) noexcept;
```

1 *Returns:* A `simd` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

2 *Remarks:* Each of these operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend constexpr simd operator<<(const simd& v, int n) noexcept;
friend constexpr simd operator>>(const simd& v, int n) noexcept;
```

3 *Returns:* A `simd` object where the i^{th} element is initialized to the result of applying the indicated operator to `v[i]` and `n` for all i in the range of `[0, size())`.

4 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

(6.2.1.7.2) 28.9.7.2 `simd` compound assignment [simd.cassign]

```
friend constexpr simd& operator+=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator-=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator*=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator/=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator%=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator&=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator|=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator^=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator<<=(simd& lhs, const simd& rhs) noexcept;
friend constexpr simd& operator>>=(simd& lhs, const simd& rhs) noexcept;
```

1 *Effects:* These operators apply the indicated operator to `lhs` and `rhs` as an element-wise operation.

2 *Returns:* `lhs`.

3 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type `value_type`.

```
friend constexpr simd& operator<=(simd& lhs, int n) noexcept;
friend constexpr simd& operator>=(simd& lhs, int n) noexcept;
```

- 4 *Effects:* Equivalent to: return operator@=(lhs, simd(n));
- 5 *Remarks:* These operators shall not participate in overload resolution unless the indicated operator can be applied to objects of type value_type.

(6.2.1.7.3) 28.9.7.3 simd compare operators [simd.comparison]

```
friend constexpr mask_type operator==(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator!=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator>=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator<=(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator>(const simd& lhs, const simd& rhs) noexcept;
friend constexpr mask_type operator<(const simd& lhs, const simd& rhs) noexcept;
```

- 1 *Returns:* A simd_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(6.2.1.7.4) 28.9.7.4 ~~R~~simd reductions [simd.reductions]

- 1 In this subclause, BinaryOperation shall be a binary element-wise operation.

```
template<class T, class Abi, class BinaryOperation = plus<>>
constexpr T reduce(const simd<T, Abi>& x, BinaryOperation binary_op = {});
```

- 2 *Requires:* binary_op shall be callable with two arguments of type T returning T, or callable with two arguments of type simd<T, A1> returning simd<T, A1> for every A1 that is an ABI tag type.
- 3 *Returns:* GENERALIZED_SUM(binary_op, x.data[i], ...) for all i in the range of [0, size())(?).
- 4 *Throws:* Any exception thrown from binary_op.

```
template<class M, class V, class BinaryOperation>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x,
                                         typename V::value_type identity_element,
                                         BinaryOperation binary_op = {});
```

- 5 *Requires:* binary_op shall be callable with two arguments of type T returning T, or callable with two arguments of type simd<T, A1> returning simd<T, A1> for every A1 that is an ABI tag type. The results of binary_op(identity_element, x) and binary_op(x, identity_element) shall be equal to x for all finite values x representable by V::value_type.
- 6 *Returns:* If none_of(x.mask), returns identity_element. Otherwise, returns GENERALIZED_SUM(binary_op, x.data[i], ...) for all selected indices i.
- 7 *Throws:* Any exception thrown from binary_op.

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, plus<> binary_op) noexcept;
```

- 8 *Returns:* If `none_of(x.mask)`, returns 0. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices i .

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, multiplies<> binary_op) noexcept;
```

- 9 *Returns:* If `none_of(x.mask)`, returns 1. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices i .

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_and<> binary_op) noexcept;
```

- 10 *Requires:* `is_integral_v<V::value_type>` is true.

- 11 *Returns:* If `none_of(x.mask)`, returns `~V::value_type()`. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices i .

```
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_or<> binary_op) noexcept;
template<class M, class V>
constexpr typename V::value_type reduce(const const_where_expression<M, V>& x, bit_xor<> binary_op) noexcept;
```

- 12 *Requires:* `is_integral_v<V::value_type>` is true.

- 13 *Returns:* If `none_of(x.mask)`, returns 0. Otherwise, returns `GENERALIZED_SUM(binary_op, x.data[i], ...)` for all selected indices i .

```
template<class T, class Abi> constexpr T hmin(const simd<T, Abi>& x) noexcept;
```

- 14 *Returns:* The value of an element $x[j]$ for which $x[j] \leq x[i]$ for all i in the range of $[0, \text{size}())$.

```
template<class M, class V> constexpr typename V::value_type hmin(const const_where_expression<M, V>& x) noexcept;
```

- 15 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::max()`. Otherwise, returns the value of an element $x.data[j]$ for which $x.mask[j] == \text{true}$ and $x.data[j] \leq x.data[i]$ for all selected indices i .

```
template<class T, class Abi> constexpr T hmax(const simd<T, Abi>& x) noexcept;
```

- 16 *Returns:* The value of an element $x[j]$ for which $x[j] \geq x[i]$ for all i in the range of $[0, \text{size}())$.

```
template<class M, class V> constexpr typename V::value_type hmax(const const_where_expression<M, V>& x) noexcept;
```

- 17 *Returns:* If `none_of(x.mask)`, the return value is `numeric_limits<V::value_type>::lowest()`. Otherwise, returns the value of an element $x.data[j]$ for which $x.mask[j] == \text{true}$ and $x.data[j] \geq x.data[i]$ for all selected indices i .

(6.2.1.7.5) 28.9.7.5 Casts

[simd.casts]

```
template<class T, class U, class Abi> see below simd_cast(const simd<U, Abi>& x) noexcept;
```

- 1 ~~Let T_0 denote $T::\text{value_type}$ if $\text{is_simd_v}<T>$ is true, or T otherwise.~~
- 2 ~~Returns: A simd object with the i^{th} element initialized to $\text{static_cast}<T_0>(x[i])$ for all i in the range of $[0, \text{size}())$.~~
- 3 ~~Remarks: The function shall not participate in overload resolution unless~~
 - ~~• every possible value of type U can be represented with type T_0 , and~~
 - ~~• either~~
 - ~~• $\text{is_simd_v}<T>$ is false, or~~
 - ~~• $T::\text{size}() == \text{simd}<U, \text{Abi}>::\text{size}()$ is true.~~
- 4 ~~The return type is~~
 - ~~• T if $\text{is_simd_v}<T>$ is true;~~
 - ~~• otherwise, $\text{simd}<T, \text{Abi}>$ if U is the same type as T ;~~
 - ~~• otherwise, $\text{simd}<T, \text{simd_abi}::\text{fixed_size}<\text{simd}<U, \text{Abi}>::\text{size}()>$~~

```
template<class T, class U, class Abi> see below static_simd_cast(const simd<U, Abi>& x) noexcept;
```

- 5 ~~Let T_0 denote $T::\text{value_type}$ if $\text{is_simd_v}<T>$ is true or T otherwise.~~
- 6 ~~Returns: A simd object with the i^{th} element initialized to $\text{static_cast}<T_0>(x[i])$ for all i in the range of $[0, \text{size}())$.~~
- 7 ~~Remarks: The function shall not participate in overload resolution unless either~~
 - ~~• $\text{is_simd_v}<T>$ is false, or~~
 - ~~• $T::\text{size}() == \text{simd}<U, \text{Abi}>::\text{size}()$ is true.~~
- 8 ~~The return type is~~
 - ~~• T if $\text{is_simd_v}<T>$ is true;~~
 - ~~• otherwise, $\text{simd}<T, \text{Abi}>$ if either U is the same type as T or $\text{make_signed_t}<U>$ is the same type as $\text{make_signed_t}<T>$;~~
 - ~~• otherwise, $\text{simd}<T, \text{simd_abi}::\text{fixed_size}<\text{simd}<U, \text{Abi}>::\text{size}()>$.~~

```
template<class T, class Abi>
```

```
fixed_size_simd<T, simd_size_v<T, Abi> to_fixed_size(const simd<T, Abi>& x) noexcept;
```

```
template<class T, class Abi>
```

```
fixed_size_simd_mask<T, simd_size_v<T, Abi> to_fixed_size(const simd_mask<T, Abi>& x) noexcept;
```

- 9 ~~Returns: A data-parallel object with the i^{th} element initialized to $x[i]$ for all i in the range of $[0, \text{size}())$.~~

```
template<class T, int N> native_simd<T> to_native(const fixed_size_simd<T, N>& x) noexcept;
```

```
template<class T, int N> native_simd_mask<T> to_native(const fixed_size_simd_mask<T, N>& x) noexcept;
```

- 10 ~~Returns: A data-parallel object with the i^{th} element initialized to $x[i]$ for all i in the range of $[0, \text{size}())$.~~
- 11 ~~Remarks: These functions shall not participate in overload resolution unless $\text{simd_size_v}<T, \text{simd_abi}::\text{native}<T> == N$ is true.~~

```
template<class T, int N> simd<T> to_compatible(const fixed_size_simd<T, N>& x) noexcept;
template<class T, int N> simd_mask<T> to_compatible(const fixed_size_simd_mask<T, N>& x) noexcept;
```

- 12 *Returns:* A data-parallel object with the i^{th} element initialized to $x[i]$ for all i in the range of $[0, \text{size}())$.
- 13 *Remarks:* These functions shall not participate in overload resolution unless `simd_size_v<T, simd_abi::compatible<T> == N` is true.

```
template<size_t... Sizes, class T, class Abi>
constexpr tuple<simd<T, simd_abi::deduce_t<T, Sizes>>...>
    split(const simd<T, Abi>& x) noexcept;
template<size_t... Sizes, class T, class Abi>
constexpr tuple<simd_mask<T, simd_abi::deduce_t<T, Sizes>>...>
    split(const simd_mask<T, Abi>& x) noexcept;
```

- 14 *Returns:* A tuple of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} tuple element initialized to the value of the element x with index $i + \text{sum of the first } j \text{ values in the Sizes pack}$.
- 15 *Remarks:* These functions shall not participate in overload resolution unless the sum of all values in the `Sizes` pack is equal to `simd_size_v<T, Abi>`.

```
template<class V, class Abi>
constexpr array<V, simd_size_v<typename V::value_type, Abi> / V::size()>
    split(const simd<typename V::value_type, Abi>& x) noexcept;
template<class V, class Abi>
constexpr array<V, simd_size_v<typename V::simd_type::value_type, Abi> / V::size()>
    split(const simd_mask<typename V::simd_type::value_type, Abi>& x) noexcept;
```

- 16 *Returns:* An array of data-parallel objects with the i^{th} `simd/simd_mask` element of the j^{th} array element initialized to the value of the element in x with index $i + j * V::size()$.
- 17 *Remarks:* These functions shall not participate in overload resolution unless either:
- `is_simd_v<V>` is true and `simd_size_v<typename V::value_type, Abi>` is an integral multiple of `V::size()`, or
 - `is_simd_mask_v<V>` is true and `simd_size_v<typename V::simd_type::value_type, Abi>` is an integral multiple of `V::size()`.

```
template<size_t N, class T, class A>
constexpr array<resize_simd<simd_size_v<T, A> / N, simd<T, A>>, N>
    split_by(const simd<T, A>& x) noexcept;
template<size_t N, class T, class A>
constexpr array<resize_simd<simd_size_v<T, A> / N, simd_mask<T, A>>, N>
    split_by(const simd_mask<T, A>& x) noexcept;
```

- 18 *Returns:* An array `arr`, where `arr[i][j]` is initialized by `x[i * (simd_size_v<T, A> / N) + j]`.
- 19 *Remarks:* The functions shall not participate in overload resolution unless `simd_size_v<T, A>` is an integral multiple of `N`.

```

template<class T, class... Abis>
    constexpr simd<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
        const simd<T, Abis>&... xs) noexcept;
template<class T, class... Abis>
    constexpr simd_mask<T, simd_abi::deduce_t<T, (simd_size_v<T, Abis> + ...) >> concat(
        const simd_mask<T, Abis>&... xs) noexcept;

```

- 20 *Returns:* A data-parallel object initialized with the concatenated values in the `xs` pack of data-parallel objects: The i^{th} `simd/simd_mask` element of the j^{th} parameter in the `xs` pack is copied to the return value's element with index i + the sum of the width of the first j parameters in the `xs` pack.

```

template<class T, class Abi, size_t N>
    constexpr resize_simd<simd_size_v<T, Abi> * N, simd<T, Abi>>
        concat(const array<simd<T, Abi>, N>& arr) noexcept;
template<class T, class Abi, size_t N>
    constexpr resize_simd_mask<simd_size_v<T, Abi> * N, simd_mask<T, Abi>>
        concat(const array<simd_mask<T, Abi>, N>& arr) noexcept;

```

- 21 *Returns:* A data-parallel object, the i^{th} element of which is initialized by `arr[i / simd_size_v<T, Abi>][i % simd_size_v<T, Abi>]`.

(6.2.1.7.6) 28.9.7.6 Algorithms

[simd.alg]

```

template<class T, class Abi> constexpr simd<T, Abi> min(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;

```

- 1 *Returns:* The result of the element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())`.

```

template<class T, class Abi> constexpr simd<T, Abi> max(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;

```

- 2 *Returns:* The result of the element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())`.

```

template<class T, class Abi>
    constexpr pair<simd<T, Abi>, simd<T, Abi>> minmax(const simd<T, Abi>& a, const simd<T, Abi>& b) noexcept;

```

- 3 *Returns:* A pair initialized with
- the result of element-wise application of `std::min(a[i], b[i])` for all i in the range of `[0, size())` in the first member, and
 - the result of element-wise application of `std::max(a[i], b[i])` for all i in the range of `[0, size())` in the second member.

```

template<class T, class Abi> simd<T, Abi>
    constexpr clamp(const simd<T, Abi>& v, const simd<T, Abi>& lo, const simd<T, Abi>& hi);

```

- 4 *Requires:* No element in `lo` shall be greater than the corresponding element in `hi`.
- 5 *Returns:* The result of element-wise application of `std::clamp(v[i], lo[i], hi[i])` for all i in the range of `[0, size())`.

(6.2.1.7.7) 28.9.7.7 ~~M~~simd math library [simd.math]

¹ For each set of overloaded functions within `<cmath>`, there shall be additional overloads sufficient to ensure that if any argument corresponding to a double parameter has type `simd<T, Abi>`, where `is_floating_point_v<T>` is true, then:

- All arguments corresponding to double parameters shall be convertible to `simd<T, Abi>`.
- All arguments corresponding to double* parameters shall be of type `simd<T, Abi>*`.
- All arguments corresponding to parameters of integral type `U` shall be convertible to `fixed_size_simd<U, simd_size_v<T, Abi>>`.
- All arguments corresponding to `U*`, where `U` is integral, shall be of type `fixed_size_simd<U, simd_size_v<T, Abi>*>`.
- If the corresponding return type is `double`, the return type of the additional overloads is `simd<T, Abi>`. Otherwise, if the corresponding return type is `bool`, the return type of the additional overload is `simd_mask<T, Abi>`. Otherwise, the return type is `fixed_size_simd<R, simd_size_v<T, Abi>>`, with `R` denoting the corresponding return type.

It is unspecified whether a call to these overloads with arguments that are all convertible to `simd<T, Abi>` but are not of type `simd<T, Abi>` is well-formed.

² Each function overload produced by the above rules applies the indicated `<cmath>` function element-wise. For the mathematical functions, the results per element only need to be approximately equal to the application of the function which is overloaded for the element type.

³ The ~~behavior is undefined~~result is unspecified if a domain, pole, or range error occurs when the input argument(s) are applied to the indicated `<cmath>` function. [Note: Implementations are encouraged to follow the C specification (especially Annex F). — end note]

⁴ TODO: Allow `abs(simd<signed-integral>)`.

⁵ If `abs` is called with an argument of type `simd<X, Abi>` for which `is_unsigned_v<X>` is true, the program is ill-formed.

(6.2.1.8) 28.9.8 Class template `simd_mask` [simd.mask.class]

(6.2.1.8.1) 28.9.8.1 Class template `simd_mask` overview [simd.mask.overview]

```
template<class T, class Abi> class simd_mask {
public:
    using value_type = bool;
    using reference = see below;
    using simd_type = simd<T, Abi>;
    using abi_type = Abi;

    static constexpr size_tsize_t size()size() noexceptnoexcept typename simd_size<T, Abi>::type size;

    constexpr simd_mask() noexcept = default;

    // 28.9.8.3, Csimd_mask constructors
```

```
constexpr explicit simd_mask(value_type) noexcept;
template<class U, class UAbi>
    constexpr explicit(sizeof(U) != sizeof(T))
        simd_mask(const simd_mask<U, simd_abi::fixed_size<size()->UAbi>&) noexcept;
template<class G> constexpr explicit simd_mask(G&& gen) noexcept;
template<contiguous_iterator It, class Flags = element_aligned_tag>
    constexpr simd_mask(const value_type* memIt& first, Flags = {});
```

// 28.9.8.4, *simd_mask copy functions*

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
    constexpr void copy_from(const value_type* memIt& first, Flags = {});
template<contiguous_iterator It, class Flags = element_aligned_tag>
    requires output_iterator<It, value_type>
    constexpr void copy_to(value_type* mem const It& first, Flags = {}) const;
```

// 28.9.8.5, *simd_mask subscript operators*

```
constexpr reference operator[](size_t);
constexpr value_type operator[](size_t) const;
```

// 28.9.8.6, *simd_mask unary operators*

```
constexpr simd_mask operator!() const noexcept;
```

// 28.9.9.1, *simd_mask binary operators*

```
friend constexpr simd_mask operator&&(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator||(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator&(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator|(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator^(const simd_mask&, const simd_mask&) noexcept;
```

// 28.9.9.2, *simd_mask compound assignment*

```
friend constexpr simd_mask& operator&=(simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask& operator|=(simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask& operator^=(simd_mask&, const simd_mask&) noexcept;
```

// 28.9.9.3, *simd_mask comparisons*

```
friend constexpr simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
};
```

- ¹ The class template `simd_mask` is a data-parallel type with the element type `bool`. The width of a given `simd_mask` specialization is a constant expression, determined by the template parameters. Specifically, `simd_mask<T, Abi>::size() == simd<T, Abi>::size()`.
- ² Every specialization of `simd_mask` ~~shall be~~ is a complete type. The specialization `simd_mask<T, Abi>` is supported if `T` is a vectorizable type and
 - `Abi` is `simd_abi::scalar`, or
 - `Abi` is `simd_abi::fixed_size<N>`, with `N` constrained as defined in (28.9.3).

If `Abi` is an extended ABI tag, it is implementation-defined whether `simd_mask<T, Abi>` is supported. [*Note*: The intent is for implementations to decide on the basis of the currently targeted system. — *end note*]

If `simd_mask<T, Abi>` is not supported, the specialization shall have a deleted default constructor, deleted destructor, deleted copy constructor, and deleted copy assignment. Otherwise, the following are true:

- `is_nothrow_move_constructible_v<simd_mask<T, Abi>>`, and
- `is_nothrow_move_assignable_v<simd_mask<T, Abi>>`, and
- `is_nothrow_default_constructible_v<simd_mask<T, Abi>>`.

- Default initialization performs no initialization of the elements; value-initialization initializes each element with false. [*Note*: Thus, default initialization leaves the elements in an indeterminate state. — *end note*]
- Implementations should enable explicit conversion from and to implementation-defined types. This adds one or more of the following declarations to class `simd_mask`:

```
constexpr explicit operator implementation-defined() const;
constexpr explicit simd_mask(const implementation-defined& init) const;
```

- The member type reference has the same interface as `simd<T, Abi>::reference`, except its `value_type` is `bool`. (28.9.6.3)

(6.2.1.8.2) 28.9.8.2 `simd_mask` width [`simd.mask.width`]

```
static constexpr size_tsize()noexcept typename simd_size<T, Abi>::type size;
```

- Returns*: The width of `simd<T, Abi>`.

(6.2.1.8.3) 28.9.8.3 ~~C~~simd_mask constructors [`simd.mask.ctor`]

```
constexpr explicit simd_mask(value_type x) noexcept;
```

- Effects*: Constructs an object with each element initialized to `x`.

```
template<class U, class UAbi>
```

```
constexpr explicit(sizeof(U) != sizeof(T)) simd_mask(const simd_mask<U, simd_abi::fixed_size<size()UAbi>& x) noexcept;
```

- Effects*: Constructs an object of type `simd_mask` where the i^{th} element equals `x[i]` for all i in the range of `[0, size())`.

- Remarks*: This constructor shall not participate in overload resolution unless ~~`abi_type is simd_abi::fixed_size<size()`~~`simd_size_v<U, UAbi> == size()`.

```
template<class G> constexpr explicit simd_mask(G&& gen) noexcept;
```

- Effects*: Constructs an object where the i^{th} element is initialized to `gen(integral_constant<size_t, i>())`.

- Remarks*: This constructor shall not participate in overload resolution unless `static_cast<bool>(gen(integral_constant<size_t, i>()))` is well-formed for all i in the range of `[0, size())`.

- The calls to `gen` are unsequenced with respect to each other. Vectorization-unsafe standard library functions may not be invoked by `gen` (?).

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr simd_mask(const value_type* memIt& first, Flags = {});
```

7 Constraints: iter_value_t<It> is of type bool.

8 *Requires:*

- [mem, mem + size()) is a valid range.
- If the template parameter Flags is vector_aligned_tag, mem shall point to storage aligned by memory_alignment_v<simd_mask>.
- If the template parameter Flags is overaligned_tag<N>, mem shall point to storage aligned by N.
- If the template parameter Flags is element_aligned_tag, mem shall point to storage aligned by alignof(value_type).

9 *Effects:* Constructs an object where the i^{th} element is initialized to mem[i] for all i in the range of [0, size()).

10 *Throws:* Nothing.

11 *Remarks:* This constructor shall not participate in overload resolution unless is_simd_flag_type_v<Flags> is true.

(6.2.1.8.4) 28.9.8.4 simd_mask copy functions

[simd.mask.copy]

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
constexpr void copy_from(const value_type* memIt& first, Flags = {});
```

1 Constraints: iter_value_t<It> is of type bool.

2 *Requires:*

- [mem, mem + size()) is a valid range.
- If the template parameter Flags is vector_aligned_tag, mem shall point to storage aligned by memory_alignment_v<simd_mask>.
- If the template parameter Flags is overaligned_tag<N>, mem shall point to storage aligned by N.
- If the template parameter Flags is element_aligned_tag, mem shall point to storage aligned by alignof(value_type).

3 *Effects:* Replaces the elements of the simd_mask object such that the i^{th} element is replaced with mem[i] for all i in the range of [0, size()).

4 *Throws:* Nothing.

5 *Remarks:* This function shall not participate in overload resolution unless is_simd_flag_type_v<Flags> is true.

```
template<contiguous_iterator It, class Flags = element_aligned_tag>
requires output_iterator<It, value_type>
constexpr void copy_to(value_type* memconst It& first, Flags = {}) const;
```

6 Constraints: iter_value_t<It> is of type bool.

7 *Requires:*

- [mem, mem + size()) is a valid range.

- If the template parameter `Flags` is `vector_aligned_tag`, `mem` shall point to storage aligned by `memory_alignment_v<simd_mask>`.
- If the template parameter `Flags` is `overaligned_tag<N>`, `mem` shall point to storage aligned by `N`.
- If the template parameter `Flags` is `element_aligned_tag`, `mem` shall point to storage aligned by `alignof(value_type)`.

8 *Effects:* Copies all `simd_mask` elements as if `mem[i] = operator[](i)` for all *i* in the range of `[0, size())`.

9 *Throws:* Nothing.

10 *Remarks:* This function shall not participate in overload resolution unless `is_simd_flag_type_v<Flags>` is true.

(6.2.1.8.5) 28.9.8.5 **S**simd_mask subscript operators [simd.mask.subscr]

```
constexpr reference operator[](size_t i);
```

1 *Requires:* `i < size()`.

2 *Returns:* A reference (see 28.9.6.3) referring to the *i*th element.

3 *Throws:* Nothing.

```
constexpr value_type operator[](size_t i) const;
```

4 *Requires:* `i < size()`.

5 *Returns:* The value of the *i*th element.

6 *Throws:* Nothing.

(6.2.1.8.6) 28.9.8.6 **U**simd_mask unary operators [simd.mask.unary]

```
constexpr simd_mask operator!() const noexcept;
```

1 *Returns:* The result of the element-wise application of `operator!`.

(6.2.1.9) 28.9.9 Non-member operations [simd.mask.nonmembers]

(6.2.1.9.1) 28.9.9.1 **B**simd_mask binary operators [simd.mask.binary]

```
friend constexpr simd_mask operator&&(const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

```
friend constexpr simd_mask operator||(const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

```
friend constexpr simd_mask operator& (const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

```
friend constexpr simd_mask operator| (const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

```
friend constexpr simd_mask operator^ (const simd_mask& lhs, const simd_mask& rhs) noexcept;
```

1 *Returns:* A `simd_mask` object initialized with the results of applying the indicated operator to `lhs` and `rhs` as a binary element-wise operation.

(6.2.1.9.2) 28.9.9.2 **E**simd_mask compound assignment [simd.mask.cassign]

```
friend constexpr simd_mask& operator&=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask& operator|=(simd_mask& lhs, const simd_mask& rhs) noexcept;
friend constexpr simd_mask& operator^=(simd_mask& lhs, const simd_mask& rhs) noexcept;
```

- 1 *Effects:* These operators apply the indicated operator to lhs and rhs as a binary element-wise operation.
- 2 *Returns:* lhs.

(6.2.1.9.3) 28.9.9.3 **E**simd_mask comparisons [simd.mask.comparison]

```
friend constexpr simd_mask operator==(const simd_mask&, const simd_mask&) noexcept;
friend constexpr simd_mask operator!=(const simd_mask&, const simd_mask&) noexcept;
```

- 1 *Returns:* A simd_mask object initialized with the results of applying the indicated operator to lhs and rhs as a binary element-wise operation.

(6.2.1.9.4) 28.9.9.4 **R**simd_mask reductions [simd.mask.reductions]

```
template<class T, class Abi> constexpr bool all_of(const simd_mask<T, Abi>& k) noexcept;
```

- 1 *Returns:* true if all boolean elements in k are true, false otherwise.

```
template<class T, class Abi> constexpr bool any_of(const simd_mask<T, Abi>& k) noexcept;
```

- 2 *Returns:* true if at least one boolean element in k is true, false otherwise.

```
template<class T, class Abi> constexpr bool none_of(const simd_mask<T, Abi>& k) noexcept;
```

- 3 *Returns:* true if none of the one boolean elements in k is true, false otherwise.

```
template<class T, class Abi> constexpr bool some_of(const simd_mask<T, Abi>& k) noexcept;
```

- 4 *Returns:* true if at least one of the one boolean elements in k is true and at least one of the boolean elements in k is false, false otherwise.

```
template<class T, class Abi> constexpr int popcount(const simd_mask<T, Abi>& k) noexcept;
```

- 5 *Returns:* The number of boolean elements in k that are true.

```
template<class T, class Abi> constexpr int find_first_set(const simd_mask<T, Abi>& k);
```

- 6 *Requires:* any_of(k) returns true.
- 7 *Returns:* The lowest element index *i* where k[*i*] is true.
- 8 *Throws:* Nothing.

```
template<class T, class Abi> constexpr int find_last_set(const simd_mask<T, Abi>& k);
```

- 9 *Requires:* any_of(k) returns true.
 10 *Returns:* The greatest element index i where $k[i]$ is true.
 11 *Throws:* Nothing.

```
constexpr bool all_of(T) noexcept;  

constexpr bool any_of(T) noexcept;  

constexpr bool none_of(T) noexcept;  

constexpr bool some_of(T) noexcept;  

constexpr int popcount(T) noexcept;
```

- 12 *Returns:* all_of and any_of return their arguments; none_of returns the negation of its argument; some_of returns false; popcount returns the integral representation of its argument.
 13 *Remarks:* The parameter type T is an unspecified type that is only constructible via implicit conversion from bool.

```
constexpr int find_first_set(T);  

constexpr int find_last_set(T);
```

- 14 *Requires:* The value of the argument is true.
 15 *Returns:* 0.
 16 *Throws:* Nothing.
 17 *Remarks:* The parameter type T is an unspecified type that is only constructible via implicit conversion from bool.

(6.2.1.9.5) 28.9.9.5 where functions [simd.mask.where]

```
template<class T, class Abi>  

  where_expression<simd_mask<T, Abi>, simd<T, Abi>>  

  where(const typename simd<T, Abi>::mask_type& k, simd<T, Abi>& v) noexcept;  

template<class T, class Abi>  

  const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>  

  where(const typename simd<T, Abi>::mask_type& k, const simd<T, Abi>& v) noexcept;  

template<class T, class Abi>  

  where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>  

  where(const type_identity_t<simd_mask<T, Abi>>& k, simd_mask<T, Abi>& v) noexcept;  

template<class T, class Abi>  

  const_where_expression<simd_mask<T, Abi>, simd_mask<T, Abi>>  

  where(const type_identity_t<simd_mask<T, Abi>>& k, const simd_mask<T, Abi>& v) noexcept;
```

- 1 *Returns:* An object (28.9.5) with mask and data initialized with k and v respectively.

```
template<class T>
  where_expression<bool T>
    where(see below k, T& v) noexcept;
template<class T>
  const_where_expression<bool, T>
    where(see below k, const T& v) noexcept;
```

- 2 **Remarks:** The functions shall not participate in overload resolution unless
- `T` is neither a `simd` nor a `simd_mask` specialization, and
 - the first argument is of type `bool`.
- 3 **Returns:** An object (28.9.5) with `mask` and `data` initialized with `k` and `v` respectively.
-

A

BIBLIOGRAPHY

- [P2509R0] Giuseppe D'Angelo. *P2509R0: A proposal for a type trait to detect value-preserving conversions*. ISO/IEC C++ Standards Committee Paper. 2021. URL: <https://wg21.link/p2509r0>.
- [D0917] Matthias Kretz. *D0917: Making operator?: overloadable*. ISO/IEC C++ Standards Committee Paper. 2023. URL: <https://web-docs.gsi.de/~mkretz/D0917.pdf>.
- [P0214R9] Matthias Kretz. *P0214R9: Data-Parallel Vector Types & Operations*. ISO/IEC C++ Standards Committee Paper. 2018. URL: <https://wg21.link/p0214r9>.
- [P0350R0] Matthias Kretz. *P0350R0: Integrating datapar with parallel algorithms and executors*. ISO/IEC C++ Standards Committee Paper. 2016. URL: <https://wg21.link/p0350r0>.
- [P0851R0] Matthias Kretz. *P0851R0: `simd<T>` is neither a product type nor a container type*. ISO/IEC C++ Standards Committee Paper. 2017. URL: <https://wg21.link/p0851r0>.
- [P1915R0] Matthias Kretz. *P1915R0: Expected Feedback from `simd` in the Parallelism TS 2*. ISO/IEC C++ Standards Committee Paper. 2019. URL: <https://wg21.link/p1915r0>.
- [P2600R0] Matthias Kretz. *P2600R0: A minimal ADL restriction to avoid ill-formed template instantiation*. ISO/IEC C++ Standards Committee Paper. 2022. URL: <https://wg21.link/p2600r0>.
- [P0553R4] Jens Maurer. *P0553R4: Bit operations*. ISO/IEC C++ Standards Committee Paper. 2019. URL: <https://wg21.link/p0553r4>.
- [P0918R2] Tim Shen. *P0918R2: More `simd<>` Operations*. ISO/IEC C++ Standards Committee Paper. 2018. URL: <https://wg21.link/p0918r2>.

- [P2638R0] Daniel Towner et al. *P2638R0: Intel's response to P1915R0 for std::simd parallelism in TS 2*. ISO/IEC C++ Standards Committee Paper. 2022. URL: <https://wg21.link/p2638r0>.