

Process proposal: double-check evolutionary material via a Tentatively Ready status

Abstract

Here's the problem, or a couple of them:

Sometimes (NOT all the time), we send material onwards from EWG and LEWG and that material

- was discussed with key stakeholders elsewhere
- ends up with surprises about bakedness of the design
- ends up with surprises about implementability and CWG/LWG consistency matters

and sometimes the last bullet is something we realize when the design groups think the matter is design-wise settled and ready for reviewing technical details, in other words, off of LEWG/EWG's plate.

Raising concerns after a LEWG/EWG sign-off is awkward and has a high overhead; it's a schedule disruption and a distraction, albeit necessary. Sometimes such concerns are noticed after the sign-off, and this paper doesn't suggest that to change; rather, the author encourages CWG to continue sending design material back to EWG eagerly, and encourages LWG to do more of the same.

Here's a suggestion for a solution:

Let's make it our default process that both LEWG and EWG, when sending a proposal onwards, first make it Tentatively Ready, announce such transitions loudly and widely, give it a (hopefully) brief look at the next meeting, and if nothing particular concerns anyone, flush it forwards.

And as always, that would be a default process, and can be expedited if need be or if there exists a high level of confidence that some proposal doesn't need to wait.

Elaboration, and what we expect to happen in the double-checking phase for Tentatively Ready papers

So, for an evolutionary group, here's our new proposed default process:

1. Discuss a paper in the evolutionary group, find that it has consensus, and the evolutionary group thinks it doesn't see anything that would necessitate the paper returning to it.
2. Mark the paper Tentatively Ready and announce it has such status, putting especially the next group in the pipeline on notice that such material is incoming. The announcement is also expected to implicitly solicit review and feedback from the audiences that weren't present in the discussion. The most high-order bit, however, is that members of the next group in the pipeline are not-so-implicitly invited to review the incoming proposal, and they are given time to do it.

3. For the next meeting, look at whether there was feedback, and if not, send forward without further ado (this can be done by The Chair), and if there was, discuss the feedback on the reflectors and if need be, schedule a discussion for the feedback.

Here are examples of exceptions to the rule, situations where we might think we can go straight to the next group:

- a proposal has all the language/library designers that are not authors or co-authors of the proposal frantically nodding their head, saying it's the right thing, and language/library technicians and wordsmiths agree
- a proposal is a smaller technical change and there's early feedback from CWG/LWG saying it's a slam dunk

In other words, while there's wiggle room for uncommon sense decisions to not use the default process, let's do that when there are strong indications that it's the right thing to do.

Is this a "slow things down" process change?

No. First of all, material forwarded from an evolutionary group to the next group can't realistically be expected to be dealt with in the same meeting by the next group, except in sunshine scenarios. Furthermore, if there are no concerns raised, the material just moves without any face-to-face time overhead or scheduling overhead. In case concerns were raised, this process is arguably one with less overhead than having to punt things back from what is supposed to be a wording review, or otherwise disrupt a pipeline that was expected to be at a later stage.