

P2093R7

Formatted output

Published Proposal, 2021-07-03

Author:

[Victor Zverovich](#)

Audience:

LEWG

Project:

ISO/IEC JTC1/SC22/WG21 14882: Programming Language — C++

Table of Contents

1	Introduction
2	Revision history
3	SG16 polls (R6)
4	LEWG polls (R5)
5	SG16 polls (R3)
6	LEWG polls (R2)
7	LEWG polls (R1)
8	Motivating examples
9	API and naming
10	Unicode
11	Performance
12	Binary code
13	Impact on existing code
14	Implementation
15	Wording
	Appendix A: Unicode tests
16	Acknowledgements
	References
	Informative References

" $\|f\|_{\mathbb{C}} \|g\|_{\mathbb{H}} = \|fg\|_{\mathbb{H}}$, $\|f\|_{\mathbb{C}} \|g\|_{\mathbb{H}} = \|fg\|_{\mathbb{H}}$ "
— anonymous user

§ 1. Introduction

A new I/O-agnostic text formatting library was introduced in C++20 ([\[FORMAT\]](#)). This paper proposes integrating it with standard I/O facilities via a simple and intuitive API achieving the following goals:

- Usability
- Unicode support
- Good performance
- Small binary footprint

§ 2. Revision history

Changes since R6:

- Added new SG16 poll results.
- Rebased the wording onto the latest draft, most importantly adding compile-time checks introduced by [\[P2216\]](#).
- Added "If `out` contains invalid code units, the behavior is undefined and implementations are encouraged to diagnose it." to the definition of `vprint_unicode` per SG16 feedback.
- Replaced "invalid code points are substituted with U+FFFD  REPLACEMENT CHARACTER" with "implementations should substitute invalid code units with U+FFFD  REPLACEMENT CHARACTER per The Unicode® Standard Version 13.0 – Core Specification, Chapter 3.9" in [§ 15 Wording](#) per SG16 feedback.
- Added "The Unicode® Standard Version 13.0 – Core Specification" to Normative references in [§ 15 Wording](#).
- Clarified the behavior when mixing encodings in [§ 10 Unicode](#).

Changes since R5:

- Added new LEWG poll results.
- Added new SG16 poll results.
- Replaced `<io>` with `<print>` per LEWG feedback.
- Clarified the choice of U+FFFD  REPLACEMENT CHARACTER for transcoding errors in [§ 10 Unicode](#).
- Clarified the choice of literal encoding in [§ 10 Unicode](#).
- Clarified that ANSI escape codes for specifying coding systems are not considered a native system API that supports Unicode in [§ 10 Unicode](#).
- Added a reference to Rust's standard output facility that implements the same mojibake prevention mechanism to [§ 14 Implementation](#)

Changes since R4:

- Added SG16 Unicode poll results.
- Added a list of candidate headers formatted output functions can be added to.
- Moved the non-`ostream` overloads of formatted output functions to a separate header for cleaner separation of formatting and I/O facilities.

Changes since R3:

- Replaced `_isatty(_fileno(stream))` with `GetConsoleMode(_get_osfhandle(_fileno(stream)), ...)` in a note in `[format.functions]` because the former may return 1 for streams not referring to a terminal.

Changes since R2:

- Added better compatibility with other formatted I/O facilities as another advantage of using `stdout` in [§ 9 API and naming](#) per SG16 feedback.
- Clarified that [\[P1885\]](#) can be used for literal encoding detection per SG16 feedback.
- Added comparison of Unicode handling in various languages in [Appendix A: Unicode tests](#) and a summary in [§ 10 Unicode](#) per SG16 request.
- Removed incorrect "exposition-only" in [§ 10 Unicode](#).
- Replaced "both source and literal encodings are UTF-8, which is enabled by the `/utf-8` compiler flag" with "the literal (execution) encoding is UTF-8, which is enabled by the `/execution-charset:utf-8` compiler flag" since the source encoding is irrelevant there.
- Rephrased *Effects* of `vprint_unicode` for clarity.

Changes since R1:

- Added missing `println` overloads that take `FILE*` and `ostream&`.
- Moved the print functions that take `ostream&` to `<ostream>` since `<format>` shouldn't depend on `<ostream>`.
- Clarified why it is useful to provide `vprint*` functions in [§ 12 Binary code](#).
- Rebased the wording onto the latest working draft, N4861, in particular updating the *Throws* clauses to match existing wording.
- Replaced `std::system_error` with `system_error` in the wording.
- Added paragraph numbers to the wording.

Changes since R0:

- Clarified that adding `wchar_t` overloads will be addressed in a separate paper per the UK C++ panel feedback.

§ 3. SG16 polls (R6)

Poll: When `<print>` facilities must transcode their formatting results for display on a device, and during that process invalidly-encoded text is encountered, `print()` should replace the erroneously-encoded code units with `U+FFFD REPLACEMENT CHARACTER`.

SF	F	N	A	SA
3	3	1	2	0

Outcome: Consensus for the position

Poll: When `<print>` facilities need not transcode their formatting results for display on a device, and during that process invalidly-encoded text is encountered, `print()` should nevertheless replace the erroneously-encoded code units with `U+FFFD REPLACEMENT CHARACTER`.

SF	F	N	A	SA
1	0	2	2	3

Outcome: Consensus against the direction.

Poll: `<format>` and `<print>` facilities should have consistent behavior with respect to encoding expectations for the format string.

No objection to unanimous consent.

Poll: formatters should not be sensitive to whether they are being used with a `<format>` or `<print>` facility.

No objection to unanimous consent.

Poll: Regardless of format string encoding assumptions, `<format>` facilities may be used to format binary data.

SF	F	N	A	SA
5	1	1	0	0

Consensus: Strong consensus in favor.

Poll: Regardless of format string encoding assumptions, `<print>` facilities may be used to format binary data.

SF	F	N	A	SA
2	1	3	1	0

Consensus: Weak consensus in favor.

Poll: `<print>` facilities exhibit undefined behavior when an encoding expectation is present and a format string or formatter output does not match those expectations.

SF	F	N	A	SA
2	4	0	0	1

Consensus: Strong consensus in favor.

Poll: `<print>` facilities exhibit undefined behavior when an encoding expectation is present and a format string or formatter output does not match those expectations and output is directed to a device that has encoding expectations.

SF	F	N	A	SA
6	0	1	0	0

Consensus: Stronger consensus in favor relative to previous poll.

Poll: `<print>` facility implementors are encouraged to provide a run-time means for diagnosing format strings and formatter output that is not well-formed according to the expected encoding.

SF	F	N	A	SA
4	0	2	1	0

Consensus: Consensus in favor.

Poll: `<format>` and `<print>` facilities should have consistent behavior with respect to encoding expectations for the output of formatters.

SF	F	N	A	SA
0	1	1	5	1

Consensus: Strong consensus against.

Poll: `<print>` facility implementors are encouraged to substitute `U+FFFD` replacement characters following Unicode guidance when output is directed to a device and transcoding is necessary.

SF	F	N	A	SA
2	5	0	0	1

Consensus: Consensus in favor.

Poll: `<print>` facilities must provide an explicit program-controlled error handling mechanism for violations of encoding expectations.

SF	F	N	A	SA
0	0	3	3	3

Poll: Use of UTF-8 as the literal encoding is sufficient for `<print>` facilities to establish encoding expectations.

SF	F	N	A	SA
3	1	3	2	0

Consensus: Very weak consensus.

§ 4. LEWG polls (R5)

Poll: Block P2093 until we have a proposal for a lower level facility that can query tty/console and perform direct output to the console

SF	F	N	A	SA
2	1	5	5	6

Outcome: Consensus Against

Poll: We want `std::print` et al. non `ostream` overloads to go to:

Header	Approve	Disapprove
<code><ostream></code>	3	7
<code><io></code>	6	5
<code><print></code>	10	3
<code><format></code>	12	1
<code><utility></code>	1	13

Outcome: both `<print>` and `<format>` have consensus, author can choose one for the next revision.

§ 5. SG16 polls (R3)

Poll: Forward P2093R3 to LEWG.

SF	F	N	A	SA
4	2	2	0	1

Consensus? Yes

§ 6. LEWG polls (R2)

Poll: We want P2093R2 to revert moving the `print(ostream)` overloads to the `<ostream>` header.

SF	F	N	A	SA
0	0	6	6	1

Outcome: Keep the paper as is

§ 7. LEWG polls (R1)

Poll: We prefer `std::cout` as the default output target.

SF	F	N	A	SA
3	2	6	6	5

Poll: Add a member function on `ostream` instead of a `std::print(ostream&, ...)` free function overload.

SF	F	N	A	SA
0	2	5	17	3

Consensus against.

Poll: Remove `std::println` from the paper?

SF	F	N	A	SA
1	10	7	4	4

No consensus for change.

Poll: We are happy with the design with regards to UTF-8 output.

Unanimous consent.

Attendance: 35

§ 8. Motivating examples

Consider a common task of printing formatted text to `stdout`:

C++20	Proposed
<code>std::cout << std::format("Hello, {}!", name);</code>	<code>std::print("Hello, {}!", name);</code>

The proposed `std::print` function improves usability, avoids allocating a temporary `std::string` object and calling `operator<<` which performs formatted I/O on text that is already formatted. The number of function calls is reduced to one which, together with `std::vformat`-like type erasure, results in much smaller binary code (see [§ 12 Binary code](#)).

Existing alternatives in C++20:

Code	Comments
<code>std::cout << "Hello, " << name << "!";</code>	Requires even more formatted I/O function calls; message is interleaved with parameters; can result in interleaved output.
<code>std::printf("Hello, %s!", name);</code>	Only works if <code>name</code> is a null-terminated character string.

Constructs a temporary string; requires a call to `c_str()` and a separate I/O function call, although potentially cheaper than `operator<<`.

```
std::cout << "Привет, κόσμος!";
```

$$\|f\|_{\mathcal{H}} = \|g\|_{\mathcal{H}} = 1, \quad \|f\|_{\mathcal{H}} = \|g\|_{\mathcal{H}} = 1, \quad \|f\|_{\mathcal{H}} = \|g\|_{\mathcal{H}} = 1.$$

```
std::print("Привет, κόσμος!");
```

```
>>> print("Привет, κόσμος!")
Привет, κόσμος!
```

§ 9. API and naming

Language	Function(s)
C	<code>printf</code> [N2176]
C#/.NET	<code>Console.WriteLine</code> [DOTNET-WRITE]
COBOL	<code>DISPLAY</code> statement [N0147]
Fortran	<code>print</code> and <code>write</code> statements [N2162]

Go	<code>Printf</code> [GO-FMT]
Java	<code>PrintStream.format</code> , <code>PrintStream.print</code> , <code>PrintStream.printf</code> [JAVA-PRINT]
JavaScript	<code>console.log</code> [WHATWG-CONSOLE]
Perl	<code>printf</code> [PERL-PRINTF]
PHP	<code>printf</code> [PHP-PRINTF]
Python	<code>print</code> statement or function [PY-FUNC]
R	<code>print</code> [R-PRINT]
Ruby	<code>print</code> and <code>printf</code> [RUBY-PRINT]
Rust	<code>print!</code> [RUST-PRINT]
Swift	<code>print</code> [SWIFT-PRINT]

Variations of `print[f]` appear to be the most popular naming choice for this functionality. It is either provided as a free function (most common) or a member function (less common) together with a global object representing standard output stream. Notable exceptions are COBOL, Fortran, and Python 2 which have dedicated language statements and Rust where `print!` is a function-like macro.

We propose adding a free function called `print` with overloads for writing to the standard output (the default) and an explicitly passed output stream object. The default output stream can be either `stdout` or `std::cout`. We propose using `stdout` for the following reasons:

- `stdout` is considerably faster on at least two major implementations (see § 11 Performance).
- Better compatibility with other formatted I/O facilities compared to `std::cout` and its associated `std::streambuf` that suffer from private buffering, localization and conversion services that must be synchronized at a lower level.
- `print` won't use any formatted output functionality of `ostream`.

Since `stdout` doesn't have an associated locale we propose using the current global locale for locale-specific formatting which is consistent with `format`. With `cout` or another explicitly passed stream, the stream's locale will be used. In all cases the default formatting is locale-independent.

Another option is to make `print` a member function of `basic_ostream`. This would make usage somewhat more awkward:

```
std::cout.print("Hello, {}!", name);
```

A free function can also be overloaded to take `FILE*` to simplify migration (possibly automated) of code from `printf` to the new facility.

There are multiple approaches to appending a trailing newline:

- Don't append a newline automatically: `printf` in C and other languages.
- Append a newline but don't format arguments: `puts` in C (inconsistent with `fputs`).

- Have two formatting functions/macros, one that appends newline and another that doesn't: `print/println` in Java, `print!/println!` in Rust, `Printf/Println` in Go, `Write/WriteLine` in C#.NET.
- Let the user choose a terminating string defaulting to `"\n"` and do limited formatting: `print` in Python and Swift.

We propose not appending a newline automatically for consistency with `printf` and `iostreams`:

```
std::print("Hello, {}!", name);    // doesn't print a newline
std::print("Hello, {}!\n", name); // prints a newline
```

Additionally we can provide a function that appends a newline:

```
std::println("Hello, {}!", name); // prints a newline
```

Although `println` doesn't provide much usability improvement compared to `print` with explicit `'\n'`, it has been a frequently requested feature in the `fmt` library ([FMT]).

Another question is which header non-`ostream` overloads of formatted output functions should go to. Possible options:

- `<io>`
- `<print>`
- `<format>`
- `<ostream>`
- `<utility>`

Earlier versions of the paper proposed `<io>` analogous to `<cstdio>` so that the future I/O facilities that don't depend on `ostream` could be added there. This was changed to a more narrow-focused `<print>` but `<io>` can be added in the future once a symmetric input facility becomes available. Using `<ostream>` is undesirable because this header and its transitive dependencies are very big (42 thousand lines preprocessed on `libc++`):

```
% echo '#include <ostream>' | clang++ -E -x c++ - | wc -l
42491
```

It also pulls in a lot of unrelated symbols such as `ostream` insertion operators, global `cout`, `cerr`, `clog` variables and their `wchar_t` counterparts.

`ostream` overloads are added to the `<ostream>` header.

§ 10. Unicode

We can prevent mojibake in the Unicode example by detecting if the string literal encoding is UTF-8 and dispatching to a different function that correctly handles Unicode, for example:

```
constexpr bool is_utf8() {
    const unsigned char micro[] = "\u00B5";
    return sizeof(micro) == 3 && micro[0] == 0xC2 && micro[1] == 0xB5;
}

template <typename... Args>
void print(string_view fmt, const Args&... args) {
    if (is_utf8())
        vprint_unicode(fmt, make_format_args(args...));
    else
        vprint_nonunicode(fmt, make_format_args(args...));
}
```

where the `vprint_unicode` function formats and prints text in UTF-8 using the native system API that supports Unicode and `vprint_nonunicode` does the same for other encodings. The latter ensures that interoperability with code using legacy encodings is preserved even though `print` is a new API and it is not strictly necessary. If calling the system API requires transcoding we propose substituting invalid code units with U+FFFD REPLACEMENT CHARACTER which is consistent with the treatment of malformed UTF-8 in UTF-8-native terminals. For example

```
#include <stdio.h>

int main() {
    puts("\xc3\x28"); // Invalid 2 Octet Sequence
}
```

prints `␣(` in iTerm2 and `?(` in macOS Terminal. So whether transcoding is done or not in the UTF-8 case, you will normally get similar observed behavior.

In Visual C++ `is_utf8` will return `true` if the literal (execution) encoding is UTF-8, which is enabled by the `/execution-charset:utf-8` compiler flags or other means, and `false` otherwise. Literal encoding detection can be implemented in a more elegant way using [\[P1885\]](#).

Note that ANSI escape codes for specifying coding systems ([\[ISO2022\]](#)) are not considered a native system API that supports Unicode for the purposes of this proposal.

We propose using the literal encoding for the following reasons:

1. Consistency with the design of `std::format` which is locale-independent by default ([\[P0645\]](#)) and disallows implicitly mixing encodings e.g. passing a narrow string into a wide `std::format` is ill-formed.

2. Consistency with the encoding used for width estimation ([P1868]). The standard wording doesn't mention the literal encoding explicitly but the fact that the format strings are either literals or other compile-time strings ([P2216]) makes it the only conformant option.
3. Safety: the result of `formatted_size` does not depend on the global locale by default and a buffer allocated with this size can be passed safely to `format_to` even if the locale has been changed in the meantime, possibly from another thread.
4. Implementation and usage experience.
5. In the vast majority of cases format strings are literals. For example, analyzing a sample of 100 `printf` calls from [CODESEARCH] showed that 98 of them are string literals and 2 are string literals wrapped in the `_gettext` macro.
6. The active code page and the terminal encoding being unrelated on popular Windows localizations such as Russian where the former is CP1251 while the latter is CP866. Instead of assuming one encoding regardless of the string origin which would often result in mojibake, an explicit encoding indication can be done via the standard extension API, e.g. (exposition only)

```
print("Привет, {}!", locale_enc(string_in_locale_encoding));
```

This is already possible to implement by providing appropriate `std::formatter` specializations.

This approach has been implemented in the `fmt` library ([FMT]), successfully tested and used on a variety of platforms.

Users can sometimes restrict the set of used characters to the common subset among multiple encodings (often ASCII) in which case encoding becomes mostly irrelevant. Such "polyglot" strings are fully supported for legacy encodings and partially supported for UTF-8 by the current proposal even though mixing encodings in such a way is a clearly bad practice from a general software engineering point of view.

Here's an example output on Windows:

```
C:\projects\fmt>cl print-test.cc /I include /utf-8 /EHsc
Microsoft (R) C/C++ Optimizing Compiler Version 19.24.28316 for x86
Copyright (C) Microsoft Corporation. All rights reserved.

print-test.cc
Microsoft (R) Incremental Linker Version 14.24.28316.0
Copyright (C) Microsoft Corporation. All rights reserved.

/out:print-test.exe
print-test.obj

C:\projects\fmt>print-test
Привет, κόσμος!
```

At the same time interoperability with legacy code is preserved when literal encoding is not UTF-8. In particular, in case of EBCDIC, Shift JIS or a non-Unicode Windows code page, `print` will perform no transcoding and the text will be printed as is.

The following table summarizes the behavior of formatted output facilities in different programming languages:

Language	Linux		macOS		Windows	
	Terminal	Redirect	Terminal	Redirect	Terminal	Redirect
C	Correct	UTF-8	Correct	UTF-8	Wrong	UTF-8
Go	Correct	UTF-8	Correct	UTF-8	Correct	UTF-8
Java	Correct	UTF-8*	Correct	UTF-8*	Wrong	CP1251 (lossy)
JavaScript	Correct	UTF-8*	Correct	UTF-8*	Correct	UTF-8*
Python	Correct	UTF-8*	Correct	UTF-8*	Correct	Error
Rust	Correct	UTF-8	Correct	UTF-8	Correct	UTF-8

* - the output is transcoded from a different UTF representation.

Correct means that the test message "Привет, κόσμος!" was fully readable in the terminal output. None of the tested language facilities were able to produce readable output when piped through the standard `findstr` command on Windows. Java gave the worst results producing both mojibake and replacement characters in this case: "≡ЁЩтхЄ, ??????!". Most other languages produced valid UTF-8 when the output of `findstr` was redirected to a file.

The current paper proposes following C, Go, JavaScript and Rust and preserving the original encoding (modulo UTF conversion). The only difference compared to `printf` is that we fix the console output on Windows. Java's approach is problematic for the following reasons:

- There is a silent data loss for **valid** Unicode code points when the output is redirected to a file.
- It is more expensive because of transcoding.
- It may give an unusable result when piped through standard Windows commands like `findstr`.
- It transcodes into legacy encodings that are rarely used in practice nowadays. For example, usage of CP1251 dropped from 4.3% to 0.8% in the last 12+ years ([\[ENCODING-TRENDS\]](#)), including a 0.1% drop while the current paper was in review.

The full listings of test programs are given in [Appendix A: Unicode tests](#).

§ 11. Performance

All the performance benefits of `std::format` ([\[FORMAT\]](#)) automatically carry over to this proposal. In particular, locale-independence by default reduces global state and makes formatting more efficient compared to `stdio` and `iostreams`. There are fewer function calls (see [§ 12 Binary code](#)) and no shared formatting state compared to `iostreams`.

The following benchmark compares the reference implementation of `print` with `printf` and `ostream`. This benchmark formats a simple message and prints it to the output stream redirected to `/dev/null`. It uses the Google Benchmark library [\[GOOGLE-BENCH\]](#) to measure timings:

```
#include <cstdio>
#include <iostream>

#include <benchmark/benchmark.h>
#include <fmt/ostream.h>

void printf(benchmark::State& s) {
    while (s.KeepRunning())
        std::printf("The answer is %d.\n", 42);
}
BENCHMARK(printf);

void ostream(benchmark::State& s) {
    std::ios::sync_with_stdio(false);
    while (s.KeepRunning())
        std::cout << "The answer is " << 42 << ".\n";
}
BENCHMARK(ostream);

void print(benchmark::State& s) {
    while (s.KeepRunning())
        fmt::print("The answer is {}.\\n", 42);
}
BENCHMARK(print);

void print_cout(benchmark::State& s) {
    std::ios::sync_with_stdio(false);
    while (s.KeepRunning())
        fmt::print(std::cout, "The answer is {}.\\n", 42);
}
BENCHMARK(print_cout);

void print_cout_sync(benchmark::State& s) {
    std::ios::sync_with_stdio(true);
    while (s.KeepRunning())
        fmt::print(std::cout, "The answer is {}.\\n", 42);
}
BENCHMARK(print_cout_sync);

BENCHMARK_MAIN();
```

The benchmark was compiled with Apple clang version 11.0.0 (clang-1100.0.33.17) with `-O3 -DNDEBUG` and run on macOS 10.15.4. Below are the results:

Run on (8 X 2800 MHz CPU s)

CPU Caches:

L1 Data 32K (x4)

L1 Instruction 32K (x4)

L2 Unified 262K (x4)

L3 Unified 8388K (x1)

Load Average: 1.83, 1.88, 1.82

Benchmark	Time	CPU	Iterations
printf	87.0 ns	86.9 ns	7834009
ostream	255 ns	255 ns	2746434
print	78.4 ns	78.3 ns	9095989
print_cout	89.4 ns	89.4 ns	7702973
print_cout_sync	91.5 ns	91.4 ns	7903889

Both `print` and `printf` are ~3 times faster than `cout` even with synchronization to the standard C streams turned off. `print` is 14% faster when printing to `stdout` than to `cout`. For this reason and because `print` doesn't use formatting facilities of `ostream` we propose using `stdout` as the default output stream and providing an overload for writing to `ostream`.

On Windows 10 with Visual C++ 2019 the results are similar although the difference between `print` writing to `stdout` and `cout` is smaller with `stdout` being 7% faster:

Run on (1 X 2808 MHz CPU)

CPU Caches:

L1 Data 32K (x1)

L1 Instruction 32K (x1)

L2 Unified 262K (x1)

L3 Unified 8388K (x1)

Benchmark	Time	CPU	Iterations
printf	835 ns	816 ns	746667
ostream	2410 ns	2400 ns	280000
print	580 ns	572 ns	1120000
print_cout	623 ns	614 ns	1120000
print_cout_sync	615 ns	614 ns	1120000

§ 12. Binary code

We propose minimizing per-call binary code size by applying the type erasure mechanism from [P0645]. In this approach all the formatting and printing logic is implemented in a non-variadic function `vprint`. Inline variadic `print` function only constructs a `format_args` object, representing an array of type-erased argument references, and passes it to `vprint*`. Here is a simplified example:

```
void vprint(string_view fmt, format_args args);

template<class... Args>
inline void print(string_view fmt, const Args&... args) {
    return vprint(fmt, make_format_args(args...));
}
```

We provide `vprint*` overloads so that users can apply the same technique to their own code. For example:

```
void vlog(log_level level, string_view fmt, format_args args) {
    // Print the log level and use vprint* overloads to format and print the
    // message.
}

template<class... Args>
inline void log(log_level level, string_view fmt, const Args&... args) {
    return vlog(level, fmt, make_format_args(args...));
}
```

Here `vlog` that implements the logging logic is not parameterized on formatting argument types resulting in less code bloat compared to a naive templated version. As a real-world example, this technique has been applied in the Folly Logger ([FOLLY]) bringing ~5x binary size reduction per logging function call.

Below we compare the reference implementation of `print` to standard formatting facilities. All the code snippets are compiled with clang (Apple clang version 11.0.0 clang-1100.0.33.17) with `-O3 -DNDEBUG -c -std=c++17` and the resulting binaries are disassembled with `objdump -S`:

```
void printf_test(const char* name) {
    printf("Hello, %s!", name);
}
```

```
__Z11printf_testPKc:
    0:      55      pushq   %rbp
    1:      48 89 e5      movq    %rsp, %rbp
    4:      48 89 fe      movq    %rdi, %rsi
    7:      48 8d 3d 08 00 00 00      leaq    8(%rip), %rdi
```

```

e:      31 c0    xorl    %eax, %eax
10:     5d      popq    %rbp
11:     e9 00 00 00 00 jmp     0 <__Z11printf_testPKc+0x16>

```

```

void ostream_test(const char* name) {
    std::cout << "Hello, " << name << "!";
}

```

__Z12ostream_testPKc:

```

0:      55      pushq    %rbp
1:      48 89 e5      movq    %rsp, %rbp
4:      41 56      pushq    %r14
6:      53      pushq    %rbx
7:      48 89 fb      movq    %rdi, %rbx
a:      48 8b 3d 00 00 00 00 movq    (%rip), %rdi
11:     48 8d 35 6c 03 00 00 leaq    876(%rip), %rsi
18:     ba 07 00 00 00 00 movl    $7, %edx
1d:     e8 00 00 00 00 00 callq   0 <__Z12ostream_testPKc+0x22>
22:     49 89 c6      movq    %rax, %r14
25:     48 89 df      movq    %rbx, %rdi
28:     e8 00 00 00 00 00 callq   0 <__Z12ostream_testPKc+0x2d>
2d:     4c 89 f7      movq    %r14, %rdi
30:     48 89 de      movq    %rbx, %rsi
33:     48 89 c2      movq    %rax, %rdx
36:     e8 00 00 00 00 00 callq   0 <__Z12ostream_testPKc+0x3b>
3b:     48 8d 35 4a 03 00 00 leaq    842(%rip), %rsi
42:     ba 01 00 00 00 00 movl    $1, %edx
47:     48 89 c7      movq    %rax, %rdi
4a:     5b      popq    %rbx
4b:     41 5e      popq    %r14
4d:     5d      popq    %rbp
4e:     e9 00 00 00 00 00 jmp     0 <__Z12ostream_testPKc+0x53>
53:     66 2e 0f 1f 84 00 00 00 00 00 nopw    %cs:(%rax,%rax)
5d:     0f 1f 00      nopl    (%rax)

```

```

void print_test(const char* name) {
    print("Hello, {}!", name);
}

```

__Z10print_testPKc:

```

0:      55      pushq    %rbp
1:      48 89 e5      movq    %rsp, %rbp
4:      48 83 ec 10      subq    $16, %rsp
8:      48 89 7d f0      movq    %rdi, -16(%rbp)
c:      48 8d 3d 19 00 00 00 leaq    25(%rip), %rdi

```

```

13:      48 8d 4d f0      leaq    -16(%rbp), %rcx
17:      be 0a 00 00 00    movl    $10, %esi
1c:      ba 0d 00 00 00    movl    $13, %edx
21:      e8 00 00 00 00    callq  0 <__Z10print_testPKc+0x26>
26:      48 83 c4 10      addq    $16, %rsp
2a:      5d              popq    %rbp
2b:      c3              retq

```

The code generated for the `print_test` function that uses the reference implementation of `print` described in this proposal is more than 2x smaller than the `ostream` code and has one function call instead of three. The `printf` code is further 2x smaller but doesn't have any error handling. Adding error handling would make its code size closer to that of `print`.

The following factors contribute to the difference in binary code size between `print` and `printf`:

- Passing format string as `string_view` instead of `const char*`.
- Capturing and passing argument type information.
- Preparing the array of formatting arguments.

§ 13. Impact on existing code

The current proposal adds new functions to the headers `<print>` and `<ostream>` and should have no impact on existing code.

§ 14. Implementation

The proposed `print` function has been implemented in the open-source `fmt` library [\[FMT\]](#) and has been in use for about 6 years.

Rust's standard output facility uses essentially the same approach for preventing mojibake when printing to console on Windows ([\[RUST-STDIO\]](#)). The main difference is that invalid code units are reported as errors in Rust.

§ 15. Wording

Add an entry for `__cpp_lib_print` to section "Header `<version>` synopsis [\[version.syn\]](#)", in a place that respects the table's current alphabetic order:

```
#define __cpp_lib_print 202005L **placeholder** // also in <format>
```

Add the header `<print>` to the "C++ library headers" table in [headers], in a place that respects the table's current alphabetic order.

29.7.? Header `<print>` synopsis [print.syn].

```
namespace std {
    template<class... Args>
        void print(format-string<Args...> fmt, const Args&... args);
    template<class... Args>
        void print(FILE* stream, format-string<Args...> fmt, const Args&... args);

    template<class... Args>
        void println(format-string<Args...> fmt, const Args&... args);
    template<class... Args>
        void println(FILE* stream, format-string<Args...> fmt, const Args&... args);

    void vprint_unicode(string_view fmt, format_args args);
    void vprint_unicode(FILE* stream, string_view fmt, format_args args);

    void vprint_nonunicode(string_view fmt, format_args args);
    void vprint_nonunicode(FILE* stream, string_view fmt, format_args args);
}
```

Modify section "Header `<ostream>` synopsis [ostream.syn]":

```
...
template<class charT, class traits, class T>
basic_ostream<charT, traits>& operator<< (basic_ostream<charT, traits>&& os, const T

template<class... Args>
    void print(ostream& os, format-string<Args...> fmt, const Args&... args);
template<class... Args>
    void println(ostream& os, format-string<Args...> fmt, const Args&... args);

void vprint_unicode(ostream& os, string_view fmt, format_args args);
void vprint_nonunicode(ostream& os, string_view fmt, format_args args);
```

29.7.? Print functions [print.fun].

```
template<class... Args>
    void print(format-string<Args...> fmt, const Args&... args);
```

Effects: Equivalent to:

```
print(stdout, fmt, make_format_args(args...));
```

```
template<class... Args>  
void print(FILE* stream, format-string<Args...> fmt, const Args&... args);
```

27 *Effects:* If string literal encoding is UTF-8, equivalent to:

```
vprint_unicode(stream, fmt.str, make_format_args(args...));
```

Otherwise, equivalent to:

```
vprint_nonunicode(stream, fmt.str, make_format_args(args...));
```

```
template<class... Args>  
void println(format-string<Args...> fmt, const Args&... args);
```

28 *Effects:* Equivalent to:

```
print("{}\n", format(fmt, args...));
```

```
template<class... Args>  
void println(FILE* stream, format-string<Args...> fmt, const Args&... args);
```

29 *Effects:* Equivalent to:

```
print(stream, "{}\n", format(fmt, args...));
```

```
void vprint_unicode(string_view fmt, format_args args);
```

30 *Effects:* Equivalent to:

```
vprint_unicode(stdout, fmt, args);
```

```
void vprint_unicode(FILE* stream, string_view fmt, format_args args);
```

31 *Effects:* Let `out = vformat(fmt, args)`. If `stream` refers to a terminal [Note: On POSIX and Windows meaning that `isatty(fileno(stream))` and `GetConsoleMode(_get_osfhandle(_fileno(stream)), ...)` return nonzero respectively. — end note] capable of displaying Unicode, writes `out` to the terminal using

the native Unicode API. [Note: On Windows this API is `WriteConsoleW`. — end note] If `out` contains invalid code units, the behavior is undefined and implementations are encouraged to diagnose it. If this requires transcoding then implementations should substitute invalid code units with U+FFFD  REPLACEMENT CHARACTER per The Unicode® Standard Version 13.0 – Core Specification, Chapter 3.9. Otherwise writes `out` to `stream` unchanged.

Throws: As specified in `[format.err.report]` or `system_error` if a call by the implementation to an operating system or other underlying API results in an error that prevents the function from meeting its specifications.

```
void vprint_nonunicode(string_view fmt, format_args args);
```

32 *Effects:* Equivalent to:

```
vprint_nonunicode(stdout, fmt, args);
```

```
void vprint_nonunicode(FILE* stream, string_view fmt, format_args args);
```

33 *Effects:* Writes the result of `vformat(fmt, args)` to `stream`.

Throws: As specified in `[format.err.report]` or `system_error` if a call by the implementation to an operating system or other underlying API results in an error that prevents the function from meeting its specifications.

Add subsection "Print `[ostream.formatted.print]`" to "Formatted output functions `[ostream.formatted]`":

```
template<class... Args>
void print(ostream& os, format-string<Args...> fmt, const Args&... args);
```

1 *Effects:* If string literal encoding is UTF-8, equivalent to:

```
vprint_unicode(os, fmt, make_format_args(args...));
```

Otherwise, equivalent to:

```
vprint_nonunicode(os, fmt, make_format_args(args...));
```

```
void vprint_unicode(ostream& os, string_view fmt, format_args args);
```

2 *Effects:* Let `out = vformat(os.getloc(), fmt, args)`. If `os` is a file stream (its associated stream buffer is an instance of `basic_filebuf`) that refers to a terminal capable of displaying Unicode, writes `out` transcoded to the native system Unicode encoding to the terminal using the native API that preserves the encoding. Otherwise writes `out` to `stream` without transcoding.

Throws: As specified in [\[format.err.report\]](#) or `system_error` if a call by the implementation to an operating system or other underlying API results in an error that prevents the function from meeting its specifications.

```
void vprint_nonunicode(ostream& os, string_view fmt, format_args args);
```

3 *Effects:* Writes the result of `vformat(os.getloc(), fmt, args)` to `os`.

Throws: As specified in [\[format.err.report\]](#) or `system_error` if a call by the implementation to an operating system or other underlying API results in an error that prevents the function from meeting its specifications.

Add to Normative references [intro.refs]:

– The Unicode® Standard Version 13.0 – Core Specification

§ Appendix A: Unicode tests

This appendix gives full listings of programs for testing Unicode handling in various formatting facilities as well as test commands and their output on different platforms. The code contains additional sanity checks to ensure that the strings are encoded in some form of UTF as opposed to a legacy encoding.

C ([test.c](#)):

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    const char* message = "Привет, κόσμος!\n";
    if ((unsigned char)message[0] != 0xD0 && (unsigned char)message[1] != 0x9F)
        abort();
    printf(message);
}
```

Go ([test.go](#)):

```
package main

import "fmt"
import "log"

func main() {
    var message = "Привет, κόσμος!"
    if message[0] != 0xD0 && message[1] != 0x9F {
        log.Fatal("wrong encoding")
    }
}
```

```
}  
fmt.Println(message)  
}
```

Java (`Test.java`):

```
class Test {  
    public static void main(String[] args) {  
        String message = "Привет, κόσμος!\n";  
        if (message.charAt(0) != 0x41F) throw new RuntimeException();  
        System.out.print(message);  
    }  
}
```

JavaScript / Node.js (`test.js`):

```
message = "Привет, κόσμος!";  
if (message.charCodeAt(0) != 0x41F) throw "wrong encoding";  
console.log(message);
```

Python (`test.py`):

```
message = "Привет, κόσμος!"  
if ord(message[0]) != 0x41F:  
    raise Exception()  
print(message)
```

Rust (`test.rs`):

```
fn main() {  
    if "Привет, κόσμος!".chars().nth(0).unwrap() as u32 != 0x41F {  
        panic!();  
    }  
    println!("Привет, κόσμος!");  
}
```

Linux:

```
$ cc test.c -o c-test  
$ ./c-test  
Привет, κόσμος!  
$ ./c-test > out-c-linux.txt
```



```
$ go build -o go-test test.go
$ ./go-test
Привет, κόσμος!
$ ./go-test > out-go-linux.txt

$ java Test
Привет, κόσμος!
$ java Test > out-java-linux.txt

$ node test.js
Привет, κόσμος!
$ node test.js > out-js-linux.txt

$ python3 test.py
Привет, κόσμος!
$ python3 test.py > out-py-linux.txt

$ rustc test.rs -o rust-test
$ ./rust-test
Привет, κόσμος!
$ ./rust-test > out-rust-linux.txt
```

All output files are in UTF-8:

- [out-c-linux.txt](#)
- [out-go-linux.txt](#)
- [out-java-linux.txt](#)
- [out-js-linux.txt](#)
- [out-py-linux.txt](#)
- [out-rust-linux.txt](#)

Linux configuration:

- Ubuntu Focal 20.04 with the ru_RU.UTF-8 locale
- cc: gcc 9.3.0
- go: go1.13.8
- java: openjdk 11.0.9.1
- node: v14.5.0
- python3: 3.7.5
- rustc: 1.47.0

macOS:

```
% cc test.c -o c-test
% ./c-test
Привет, κόσμος!
% ./c-test > out-c-macos.txt

% go build -o test-go test.go
% ./test-go
Привет, κόσμος!
% ./test-go > out-go-macos.txt

% java Test
Привет, κόσμος!
% java Test > out-java-macos.txt

% node test.js
Привет, κόσμος!
% node test.js > out-js-macos.txt

% python3 test.py
Привет, κόσμος!
% python3 test.py > out-py-macos.txt

% rustc test.rs -o rust-test
% ./rust-test
Привет, κόσμος!
% ./rust-test > out-rust-macos.txt
```

All output files are in UTF-8:

- [out-c-macos.txt](#)
- [out-go-macos.txt](#)
- [out-java-macos.txt](#)
- [out-js-macos.txt](#)
- [out-py-macos.txt](#)
- [out-rust-macos.txt](#)

macOS configuration:

- macOS Catalina 10.15.7 with the ru_RU.UTF-8 locale which is the default for Russian
- cc: Apple clang version 12.0.0 (clang-1200.0.32.27)
- go: go1.15.5

- java: openjdk 14.0.1
- node: v14.5.0
- python3: 3.7.5
- rustc: 1.47.0

Windows:

```
>cl /Fe:c-test.exe test.c
...
>c-test
Привет, κόσμος!
>c-test > out-c-windows.txt
>c-test | findstr ,
Привет, κόσμος!

>go build -o go-test.exe test.go
>go-test
Привет, κόσμος!
>go-test > out-go-windows.txt
>go-test | findstr ,
Привет, κόσμος!

>java Test
Привет, ??????
>java Test > out-java-windows.txt
>java Test | findstr ,
Привет, ??????

>node test.js
Привет, κόσμος!
>node test.js > out-js-windows.txt
>node test.js | findstr ,
Привет, κόσμος!

>python test.py
Привет, κόσμος!
>python test.py > out-py-windows.txt
Traceback (most recent call last):
  File "...\\test.py", line 4, in <module>
    print(message)
  File "...\\Python39\\lib\\encodings\\cp1251.py", line 19, in encode
    return codecs.charmap_encode(input,self.errors,encoding_table)[0]
UnicodeEncodeError: 'charmap' codec can't encode characters in position 8-13: chara
>python test.py | findstr ,
Traceback (most recent call last):
  File "...\\test.py", line 4, in <module>
```

```
print(message)
File "...\\Python39\\lib\\encodings\\cp1251.py", line 19, in encode
    return codecs.charmap_encode(input,self.errors,encoding_table)[0]
UnicodeEncodeError: 'charmap' codec can't encode characters in position 8-13: charac
```



```
>rustc test.rs -o rust-test.exe
>rust-test
Привет, κόσμος!
>rust-test > out-rust-windows.txt
>rust-test | findstr ,
␣␣A␣␣␣␣␣␣␣␣␣B, ␣␣␣M␣␣␣␣␣␣␣␣B!
```

C, JavaScript (node.js), Rust and Go produced valid UTF-8 when the output was redirected to files. Java produced a file in the legacy CP1251 encoding with `?` for non-representable code points. Python failed on transcoding to CP1251. Output files:

- [out-c-windows.txt](#)
- [out-go-windows.txt](#)
- [out-java-windows.txt](#)
- [out-js-windows.txt](#)
- [out-py-windows.txt](#)
- [out-rust-windows.txt](#)

Windows configuration:

- Windows 10 10.0.19041 with Russian Region and Language settings
- cl: Microsoft (R) C/C++ Optimizing Compiler Version 19.28.29335
- go: go1.15.6
- java: Java HotSpot(TM) 64-Bit Server VM (build 15.0.1+9-18)
- node: v14.15.3
- python: 3.9.1
- rustc: v14.15.3

§ 16. Acknowledgements

Thanks to Corentin Jabot for his work on text encodings in C++ and in particular [\[P1885\]](#) that will simplify implementation of the current proposal.

Thanks to Roger Orr, Peter Brett, Hubert Tong, the BSI C++ panel and Tom Honermann for their feedback, support, constructive criticism and contributions to the proposal.

§ References

§ Informative References

[CODESEARCH]

Andrew Tomazos. [Code search engine website](https://codesearch.isocpp.org). URL: <https://codesearch.isocpp.org>

[DOTNET-WRITE]

.NET documentation, [Console.Write Method](https://docs.microsoft.com/en-us/dotnet/api/system.console.write). URL: <https://docs.microsoft.com/en-us/dotnet/api/system.console.write>

[ENCODING-TRENDS]

Historical yearly trends in the usage statistics of character encodings for websites. URL: https://w3techs.com/technologies/history_overview/character_encoding/ms/y

[FMT]

Victor Zverovich; et al. [The fmt library](https://github.com/fmtlib/fmt). URL: <https://github.com/fmtlib/fmt>

[FOLLY]

Folly: Facebook Open-source Library. URL: <https://github.com/facebook/folly>

[FORMAT]

Richard Smith. [Working Draft, Standard for Programming Language C++, Formatting \[format\]](https://wg21.link/n4861#section.20.20). URL: <https://wg21.link/n4861#section.20.20>

[GO-FMT]

Go Package Documentation, [Package fmt](https://golang.org/pkg/fmt/). URL: <https://golang.org/pkg/fmt/>

[GOOGLE-BENCH]

Google Benchmark: A microbenchmark support library. URL: <https://github.com/google/benchmark>

[ISO2022]

ISO/IEC 2022 Information technology—Character code structure and extension techniques. URL: <https://www.iso.org/standard/22747.html>

[JAVA-PRINT]

Java™ Platform, Standard Edition 7 API Specification, [Class PrintStream](https://docs.oracle.com/javase/7/docs/api/java/io/PrintStream.html). URL: <https://docs.oracle.com/javase/7/docs/api/java/io/PrintStream.html>

[MSVC-UTF8]

Visual C++ Documentation, [/utf-8 \(Set Source and Executable character sets to UTF-8\)](https://docs.microsoft.com/en-us/cpp/build/reference/utf-8-set-source-and-executable-character-sets-to-utf-8). URL: <https://docs.microsoft.com/en-us/cpp/build/reference/utf-8-set-source-and-executable-character-sets-to-utf-8>

[N0147]

ISO/IEC IS 1989:2001 – Programming language COBOL, 14.8.10 DISPLAY statement. URL: https://web.archive.org/web/20020124065139/http://www.ncits.org/tc_home/j4htm/cobolv200112.zip

[N2162]

ISO/IEC 1539-1:2018 Information technology — Programming languages — Fortran.

[N2176]

ISO/IEC 9899:2017 Programming languages — C, 7.21.6.3. The fprintf function.

[P0645]

Victor Zverovich. [Text Formatting](https://wg21.link/p0645). URL: <https://wg21.link/p0645>

[P1868]

Victor Zverovich; Zach Laine. 🦊 [width: clarifying units of width and precision in std::format](https://wg21.link/p1868). URL: <https://wg21.link/p1868>

[P1885]

Corentin Jabot. [Naming Text Encodings to Demystify Them](https://wg21.link/p1885). URL: <https://wg21.link/p1885>

[P2216]

Victor Zverovich. [std::format improvements](https://wg21.link/p2216). URL: <https://wg21.link/p2216>

[PERL-PRINTF]

Perl 5 version 30.0 documentation, Language reference, [printf](https://perldoc.perl.org/functions/printf.html). URL: <https://perldoc.perl.org/functions/printf.html>

[PHP-PRINTF]

PHP Manual, Function Reference, [printf](https://www.php.net/manual/en/function.printf.php). URL: <https://www.php.net/manual/en/function.printf.php>

[PY-FUNC]

The Python Standard Library, Built-in Functions. URL: <https://docs.python.org/3/library/functions.html>

[R-PRINT]

The R Core Team. [R: A Language and Environment for Statistical Computing, Reference Index, printf](https://cran.r-project.org/doc/manuals/r-release/fullrefman.pdf#page=457). URL: <https://cran.r-project.org/doc/manuals/r-release/fullrefman.pdf#page=457>

[RUBY-PRINT]

Documentation for Ruby, [print](https://docs.ruby-lang.org/en/2.7.0/ARGF.html#method-i-print). URL: <https://docs.ruby-lang.org/en/2.7.0/ARGF.html#method-i-print>

[RUST-PRINT]

The Rust Standard Library, Macro [std::print](https://doc.rust-lang.org/std/macro.print.html). URL: <https://doc.rust-lang.org/std/macro.print.html>

[RUST-STDIO]

The Rust Programming Language repository, [windows_stdio](https://github.com/rust-lang/rust/blob/db492ecd5ba6bd82205612cebb9034710653f0c2/library/std/src/sys/windows/stdio.rs). URL: <https://github.com/rust-lang/rust/blob/db492ecd5ba6bd82205612cebb9034710653f0c2/library/std/src/sys/windows/stdio.rs>

[SWIFT-PRINT]

Swift Standard Library, [print](https://developer.apple.com/documentation/swift/1541053-print). URL: <https://developer.apple.com/documentation/swift/1541053-print>

[WHATWG-CONSOLE]

WHATWG Standards, [Console](https://console.spec.whatwg.org/). URL: <https://console.spec.whatwg.org/>