

Atomic Refs Bound to Memory Orderings & Atomic Accessors

Document #: D2689R3
Date: 2024-04-29
Project: Programming Language C++
SG1 & LEWG
Reply-to: Christian Trott
<crtrott@sandia.gov>
Damien Lebrun-Grandie
<lebrungrandt@ornl.gov>
Mark Hoemmen
<mhoemmen@nvidia.com>
Daniel Sunderland
<dansunderland@gmail.com>
Nic Morales
<nmmoral@sandia.gov>
Nevin Liber
<nliber@anl.gov>

Contents

- 1 Revision History
 - 1.1 P2689R3 (LEWG Feedback)
 - 1.2 P2689R2 (SG1 Issaquah 2023 discussion)
 - 1.2.1 Issaquah 2023 SG1 Polls
 - 1.3 P2689R1 2023-01 (pre-Issaquah 2023) mailing
 - 1.4 Initial Version 2022-10 Mailing
 - 1.4.1 Kona 2022 SG1 polls
- 2 Rational
- 3 Design decisions
 - 3.1 Exposition Only Or Not
 - 3.2 Memory Order Conversions
- 4 Open questions
- 5 Wording
 - 5.1 Bound atomic ref
 - 5.1.1 Change the header synopsis for atomic in [atomics.syn]:
 - 5.1.2 Add the following just before [atomics.types.generic]:
 - 5.1.3 In [version.syn]:

- 5.2 Atomic Accessors
 - 5.2.1 Add to the `mdspan` header synopsis in `[mdspan.syn]`
 - 5.2.2 Put the following before `[mdspan.mdspan]`:
- 5.3 In `[version.syn]`
- 6 Acknowledgments

§ 1 Revision History

§ 1.1 P2689R3 (LEWG Feedback)

- Split definition of `atomic-ref-bound` into one for arithmetic and pointer, and one for other types
 - needed because the typedef `difference_type` should only appear for arithmetic and pointer types
- Added discussion for making `atomic-ref-bound` exposition only or not
- Added discussion arguing for not having converting constructors with respect to memory order
- Added missing `fetch_min` and `fetch_max` functions
- Use “west const”
- Added header synopsis changes

§ 1.2 P2689R2 (SG1 Issaquah 2023 discussion)

- Renamed `atomic-ref-bounded` to `atomic-ref-bound`
- Renamed `atomic-ref-unbounded` to `atomic-ref-unbound`
- Fixed the wording for `basic-atomic-accessor::offset`
- Fixed the wording for `basic-atomic-accessor::access`
- If P2616R3 (Making `std::atomic` notification/wait operations usable in more situations) is also approved, similar changes should be applied to `atomic-ref-bound` as well

§ 1.2.1 Issaquah 2023 SG1 Polls

§ 1.2.1.1 P2689 poll

Forward P2689R1 to LEWG, targeting C++26 Incl. wording changes:

- Atomic ref `[un]bounded` to `[un]bound`
- Basic atomic accessors `access()` effects should be `return reference(p[i]);`
- Basic atomic accessors `offset()` effects should be `return p+i;`

SF	F	N	A	SA
5	9	0	0	0

Unanimous consent

§ 1.2.1.2 P2616 poll

Apply changes in P2616R3 to the additions of P2689R1

SF	F	N	A	SA
----	---	---	---	----

No objection to unanimous consent

§ 1.3 P2689R1 2023-01 (pre-Issaquah 2023) mailing

- Added `atomic_ref_bounded_exposition_only` template for an `atomic_ref` like type bounded to a particular `memory_order`
- Added `atomic_ref_relaxed`, `atomic_ref_acq_rel` and `atomic_ref_seq_cst` alias templates
- Added `basic_atomic_accessor_exposition_only` template
- Added `atomic_accessor_relaxed`, `atomic_accessor_acq_rel` and `atomic_accessor_seq_cst` templates

§ 1.4 Initial Version 2022-10 Mailing

§ 1.4.1 Kona 2022 SG1 polls

§ 1.4.1.1 SG1 Poll #1

We like the bound `memory_order` versions of `atomic_ref` (This says nothing about whether these are aliases of general a template that takes a `memory_order` argument.)

`atomic_ref_relaxed`
`atomic_ref_acq_rel` (acquire for load, release for store)
`atomic_ref_seq_cst`

SF	F	N	A	SA
1	11	0	1	0

Consensus

§ 1.4.1.2 SG1 Poll #2

- X) The meaning of the bound `memory_order` is that is the default order
Y) The meaning of the bound `memory_order` is that is the only order

SX	X	N	Y	SY
1	1	3	4	2

Consensus for Y

SX : we have a variant of this type in Folly and it has been useful to be able to override this
SY : it's weird for `atomic_ref_relaxed` to have a `seq_cst` atomic done to it; since you can re-form the `_ref` you can still do it

§ 1.4.1.3 SG1 Poll #3

the next version of this proposal, with these changes, should target SG1 & LEWG

`atomic_accessor` -> `atomic_ref`
`atomic_accessor_relaxed` -> `atomic_ref_relaxed`

atomic_accessor_acq_rel -> atomic_ref_acq_rel
atomic_accessor_seq_cst -> atomic_ref_seq_cst

(Return to SG1 for final wording when this is in LWG.)

SF	F	N	A	SA
6	6	2	0	0

§ 2 Rational

This proposal adds three `atomic_ref` like types each bound to a particular `memory_order`. The API differs from `atomic_ref` in that the `memory_order` cannot be specified at run time; i.e., none of its member functions take a `memory_order` parameter. In the specific case of `atomic_ref_acq_rel`, loads are done via `memory_order_acquire` and stores are done via `memory_order_release`.

Atomic Ref	memory_order	Loads	Stores
<code>atomic_ref_relaxed</code>	<code>memory_order_relaxed</code>	<code>memory_order_relaxed</code>	<code>memory_order_relaxed</code>
<code>atomic_ref_acq_rel</code>	<code>memory_order_acq_rel</code>	<code>memory_order_acquire</code>	<code>memory_order_release</code>
<code>atomic_ref_seq_cst</code>	<code>memory_order_seq_cst</code>	<code>memory_order_seq_cst</code>	<code>memory_order_seq_cst</code>

This proposal also adds four atomic accessors to be used with `mdspan`. They only differ with their reference types and all operations are performed atomically by using that corresponding reference type.

Accessor	reference
<code>atomic_accessor</code>	<code>atomic_ref</code>
<code>atomic_accessor_relaxed</code>	<code>atomic_ref_relaxed</code>
<code>atomic_accessor_acq_rel</code>	<code>atomic_ref_acq_rel</code>
<code>atomic_accessor_seq_cst</code>	<code>atomic_ref_seq_cst</code>

`atomic_accessor` was part of the rational provided in P0009 for `mdspans` *accessor policy* template parameter.

One of the primary use cases for these accessors is the ability to write algorithms with somewhat generic `mdspan` outputs, which can be called in sequential and parallel contexts. When called in parallel contexts users would simply pass an `mdspan` with an atomic accessor - the algorithm implementation itself could be agnostic to the calling context.

A variation on this use case is an implementation of an algorithm taking an execution policy, which adds the atomic accessor to its output argument if called with a parallel policy, while using the `default_accessor` when called with a sequential policy. The following demonstrates this with a function computing a histogram:

```
template<class T, class Extents, class LayoutPolicy>
auto add_atomic_accessor_if_needed(
    std::execution::sequenced_policy,
    mdspan<T, Extents, LayoutPolicy> m) {
    return m;
}

template<class ExecutionPolicy, class T, class Extents, class LayoutPolicy>
auto add_atomic_accessor_if_needed(
```

```

    ExecutionPolicy,
    mdspan<T, Extents, LayoutPolicy> m) {
        return mdspan(m.data_handle(), m.mapping(), atomic_accessor<T>());
    }

template<class ExecT>
void compute_histogram(ExecT exec, float bin_size,
    mdspan<int, stdex::dextents<int,1>> output,
    mdspan<float, stdex::dextents<int,1>> data) {

    static_assert(is_execution_policy_v<ExecT>);

    auto accumulator = add_atomic_accessor_if_needed(exec, output);

    for_each(exec, data.data_handle(), data.data_handle()+data.extent(0), [=](float
        val) {
        int bin = std::abs(val)/bin_size;
        if(bin > output.extent(0)) bin = output.extent(0)-1;
        accumulator[bin]++;
    });
}

```

The above example is available on godbolt: <https://godbolt.org/z/cWa6MG5dj>

§ 3 Design decisions

Three options for atomic refs were discussed in SG1 in Kona 2022: add new types for `memory_order` bound atomic refs, add a new `memory_order` template parameter to the existing `atomic_ref`, or add a constructor to the existing `atomic_ref` that takes a `memory_order` and stores it. Given that the last two are ABI breaks, the first option was polled and chosen. It was also decided that the new bound atomic refs would not support overriding the specified `memory_order` at run time.

This proposal has chosen to make a general exposition-only template `atomic-ref-bound` that takes a `memory_order` as a template argument and alias templates for the three specific bound atomic refs. Also, the various member functions are constrained by integral types not including `bool`, floating point types and pointer types, as opposed to the different template specializations specified for `atomic_ref`. Other than not being able to specify the `memory_order` at run time, the intention is that the bound atomic ref types have the same functionality and API as `atomic_ref`.

Similarly for the atomic accessors, it was decided in SG1 in Kona 2022 to add four new types. This proposal has chosen to make a general exposition-only template `basic-atomic-accessor` which takes the reference type as a template parameter, and four alias templates for the specific atomic accessors.

Assuming both papers are approved, SG1 voted that similar changes to `atomic_ref` in P2616R3 (Making `std::atomic` notification/wait operations usable in more situations) should also be applied to `atomic-ref-bound`. They are not yet in the wording of either paper, as we do not know what order LWG will apply them to the working draft.

§ 3.1 Exposition Only Or Not

As mentioned above, during SG1 review we introduced explicitly named symbols instead of making `basic-atomic-accessor` and `atomic-ref-bound` public.

The primary reason for not making the generic versions public is that the memory order is generally dictated by algorithmic considerations.

The only generic thing one may want to decide is whether to use atomics at all (e.g. as the function of the execution policy).

Consequently, we could not find a use-case where the generic types are useful. That said we also don't have a strong reason to not make them public - other than that it constrains implementers.

§ 3.2 Memory Order Conversions

A question was raised whether to make the various bounded `atomic_ref` and `atomic_accessor` variants constructible/convertible from each other. We don't see many good reasons to enable that:

- `atomic_ref` is meant to be almost used in-place for individual operations
e.g. `atomic_ref(a[i])++` and not much of an interface type
- atomic accessors are something you are supposed to add in a fairly local scope too
 - we don't expect folks to pass `mdspan` with atomic accessors around

However we identified a possible reason not to permit it:

Writing functions which take `atomic_ref_MEM_ORDER` imply ordering behavior: allowing conversions would make it very easy to violate that expectation.

Consider the following:

```
void foo(atomic_ref_relaxed<int> counter);

void bar(int& counter) {
    atomic_ref_relaxed atomic_counter(counter);
    atomic_counter++;
    ...
    foo(atomic_counter);
    ...
    atomic_counter++;
    atomic_thread_fence(memory_order::acq_rel);
}
```

We believe it would be potentially unexpected for `foo` to do any operations other than relaxed atomics on `counter`. Likewise if `foo` were to take `atomic_ref_seq_cst` it would be surprising if it did relaxed atomic accesses.

§ 4 Open questions

This proposal uses alias templates to exposition-only types for `atomic_ref_relaxed`, `atomic_accessor`, etc. However, we do not want to prescribe a particular implementation. For instance, if an implementer wished to derive from an `atomic-ref-bound-like` type (to get a more user friendly name mangling, which is something not normally covered by the Standard), would they be allowed to do so? We believe an alias template to an exposition-only type is not observable from the point of view of the Standard and such an implementation would be allowed, but we request clarification on this. Potentially this is something we can leave to LWG review to suggest wording that would achieve that goal.

§ 5 Wording

The proposed changes are relative to [N4917](#):

§ 5.1 Bound atomic ref

§ 5.1.1 Change the header synopsis for atomic in [atomics.syn]:

```
// [atomics.ref.pointer], partial specialization for pointers
template<class T> struct atomic_ref<T*>; //
freestanding
// [atomics.ref.bounded], class template atomic-ref-bound
template<class T, class MemoryOrder> atomic-ref-bound; // exposition only
template<class T> using atomic_ref_relaxed = atomic-ref-bound<T,
memory_order_relaxed>; // freestanding
template<class T> using atomic_ref_acq_rel = atomic-ref-bound<T,
memory_order_acq_rel>; // freestanding
template<class T> using atomic_ref_seq_cst = atomic-ref-bound<T,
memory_order_seq_cst>; // freestanding

// [atomics.types.generic], class template atomic
template<class T> struct atomic<T>; //
freestanding
```

§ 5.1.2 Add the following just before [atomics.types.generic]:

Class template atomic-ref-bound [atomics.ref.bound]

Exposition only helper

```
template<class T>
constexpr bool atomic_ref_non_generic_type();
```

Returns: true if the class atomic_ref<T> has a member type definition difference_type, otherwise false.

[Note: This is true for a set of defined integral types, floating point types, and pointers.]

General [atomics.ref.bound.general]

```
template <class T, memory_order MemoryOrder>
struct atomic_ref_bound { // exposition only
private:
    using atomic_ref_unbound = atomic_ref<T>; // exposition only
    atomic_ref_unbound ref; // exposition only

    static constexpr memory_order store_ordering =
        MemoryOrder == memory_order_acq_rel ? memory_order_release
        : MemoryOrder; // exposition only

    static constexpr memory_order load_ordering =
        MemoryOrder == memory_order_acq_rel ? memory_order_acquire
        : MemoryOrder; // exposition only

    static constexpr bool is_integral_value =
        is_integral_v<T> && !is_same_v<T, bool>; // exposition only
    static constexpr bool is_floating_point_value =
        is_floating_point_v<T>; // exposition only
```

```

static constexpr bool is_pointer_value =
    is_pointer_v<T>; // exposition only

public:
    using value_type = T;
    static constexpr memory_order memory_ordering = MemoryOrder;
    static constexpr size_t required_alignment =
        atomic_ref_unbound::required_alignment;

    static constexpr bool is_always_lock_free =
        atomic_ref_unbound::is_always_lock_free;
    bool is_lock_free() const noexcept;

    explicit atomic_ref_bound(T&);
    atomic_ref_bound(const atomic_ref_bound&) noexcept;
    atomic_ref_bound& operator=(const atomic_ref_bound&) = delete;

    void store(T desired) const noexcept;
    T operator=(T desired) const noexcept;
    T load() const noexcept;
    operator T() const noexcept;

    T exchange(T desired) const noexcept;
    bool compare_exchange_weak(T& expected, T desired) const noexcept;
    bool compare_exchange_strong(T& expected, T desired) const noexcept;

    void wait(T old) const noexcept;
    void notify_one() const noexcept;
    void notify_all() const noexcept;
};

```

1 Class `atomic_ref_bound` is for exposition only.

2 *Mandates:*

(2.1) — `is_trivially_copyable_v<T>` is true.

(2.2) — `is_same_v<T, remove_cv_t<T>>` is true.

Operations [`atomics.ref.bound.ops`]

```
bool is_lock_free() const noexcept;
```

1 *Effects:* Equivalent to: `return ref.is_lock_free();`

```
explicit atomic_ref_bound(T& t);
```

2 *Preconditions:* The referenced object is aligned to `required_alignment`.

3 *Postconditions:* `ref` references the object referenced by `t`.

4 *Throws:* Nothing.

```
atomic_ref_bound(const atomic_ref_bound& a) noexcept;
```

5 *Postconditions:* `*this` references the object referenced by `a`.

```
void store(T desired) const noexcept;
```

6 *Effects:* Equivalent to: `ref.store(desired, store_ordering);`


```
T operator=(T desired) const noexcept;
```

7 *Effects:* Equivalent to: store(desired); return desired;

```
T load() const noexcept;
```

8 *Effects:* Equivalent to: ref.load(load_ordering);

```
operator T() const noexcept;
```

9 *Effects:* Equivalent to: return load();

```
T exchange(T desired) const noexcept;
```

10 *Effects:* Equivalent to: return ref.exchange(desired, memory_ordering);

```
bool compare_exchange_weak(T& expected, T desired) const noexcept;
```

11 *Effects:* Equivalent to: return ref.compare_exchange_weak(expected, desired, memory_ordering, load_ordering);

```
bool compare_exchange_strong(T& expected, T desired) const noexcept;
```

12 *Effects:* Equivalent to: return ref.compare_exchange_strong(expected, desired, memory_ordering, load_ordering);

```
void wait(T old) const noexcept;
```

13 *Effects:* Equivalent to: ref.wait(old, load_ordering);

```
void notify_one() const noexcept;
```

14 *Effects:* Equivalent to: ref.notify_one(); }

```
void notify_all() const noexcept;
```

15 *Effects:* Equivalent to: ref.notify_all(); }

General [atomics.ref.bound.arithmetic]

```
template <arithmetic-or-pointer T, memory_order MemoryOrder>
requires(atomic-ref-non-generic-type<T>())
struct atomic_ref_bound<T, MemoryOrder> { // exposition only
private:
    using atomic_ref_unbound = atomic_ref<T>; // exposition only
    atomic_ref_unbound ref; // exposition only

    static constexpr memory_order store_ordering =
        MemoryOrder == memory_order_acq_rel ? memory_order_release
        : MemoryOrder; // exposition only

    static constexpr memory_order load_ordering =
        MemoryOrder == memory_order_acq_rel ? memory_order_acquire
        : MemoryOrder; // exposition only

    static constexpr bool is_integral_value =
        is_integral_v<T> && !is_same_v<T, bool>; // exposition only
    static constexpr bool is_floating_point_value =
        is_floating_point_v<T>; // exposition only
    static constexpr bool is_pointer_value =
```

```

        is_pointer_v<T>; // exposition only

public:
    using value_type = T;
    using difference_type = atomic_ref_bound::difference_type;
    static constexpr memory_order memory_ordering = MemoryOrder;
    static constexpr size_t required_alignment =
        atomic_ref_unbound::required_alignment;

    static constexpr bool is_always_lock_free =
        atomic_ref_unbound::is_always_lock_free;
    bool is_lock_free() const noexcept;

    explicit atomic_ref_bound(T& t);
    atomic_ref_bound(const atomic_ref_bound&) noexcept;
    atomic_ref_bound& operator=(const atomic_ref_bound&) = delete;

    void store(T desired) const noexcept;
    T operator=(T desired) const noexcept;
    T load() const noexcept;
    operator T() const noexcept;

    T exchange(T desired) const noexcept;
    bool compare_exchange_weak(T& expected, T desired) const noexcept;
    bool compare_exchange_strong(T& expected, T desired) const noexcept;

    T fetch_add(T operand) const noexcept;
    T fetch_add(difference_type operand) const noexcept;
    T fetch_sub(T operand) const noexcept;
    T fetch_sub(difference_type operand) const noexcept;

    T fetch_and(T operand) const noexcept;
    T fetch_or(T operand) const noexcept;
    T fetch_xor(T operand) const noexcept;
    T fetch_max(T operand) const noexcept;
    T fetch_min(T operand) const noexcept;

    T operator++(int) const noexcept;
    T operator++() const noexcept;
    T operator--(int) const noexcept;
    T operator--() const noexcept;

    T operator+=(T operand) const noexcept;
    T operator+=(difference_type operand) const noexcept;
    T operator-=(T operand) const noexcept;
    T operator-=(difference_type operand) const noexcept;

    T operator&=(T operand) const noexcept;
    T operator|=(T operand) const noexcept;
    T operator^=(T operand) const noexcept;

    void wait(T old) const noexcept;
    void notify_one() const noexcept;
    void notify_all() const noexcept;
};

```

¹ Class atomic_ref_bound is for exposition only.

² Mandates:

(2.1) — `is_trivially_copyable_v<T>` is true.

(2.2) — `is_same_v<T, remove_cv_t<T>>` is true.

[Note: Unlike `atomic_ref`, the memory ordering for the arithmetic operators is `MemoryOrder`, which is not necessarily `memory_order_seq_cst`. — end note]

Operations [`atomics.ref.bound.ops`]

```
bool is_lock_free() const noexcept;
```

1 *Effects:* Equivalent to: `return ref.is_lock_free();`

```
explicit atomic-ref-bound(T& t);
```

2 *Preconditions:* The referenced object is aligned to `required_alignment`.

3 *Postconditions:* `ref` references the object referenced by `t`.

4 *Throws:* Nothing.

```
atomic-ref-bound(const atomic-ref-bound& a) noexcept;
```

5 *Postconditions:* `*this` references the object referenced by `a`.

```
void store(T desired) const noexcept;
```

6 *Effects:* Equivalent to: `ref.store(desired, store_ordering);`

```
T operator=(T desired) const noexcept;
```

7 *Effects:* Equivalent to: `store(desired); return desired;`

```
T load() const noexcept;
```

8 *Effects:* Equivalent to: `ref.load(load_ordering);`

```
operator T() const noexcept;
```

9 *Effects:* Equivalent to: `return load();`

```
T exchange(T desired) const noexcept;
```

10 *Effects:* Equivalent to: `return ref.exchange(desired, memory_ordering);`

```
bool compare_exchange_weak(T& expected, T desired) const noexcept;
```

11 *Effects:* Equivalent to: `return ref.compare_exchange_weak(expected, desired, memory_ordering, load_ordering);`

```
bool compare_exchange_strong(T& expected, T desired) const noexcept;
```

12 *Effects:* Equivalent to: `return ref.compare_exchange_strong(expected, desired, memory_ordering, load_ordering);`

```
T fetch_add(T operand) const noexcept;
```

13 *Constraints:* `is_integral_value || is_floating_point_value` is true.

14 *Effects:* Equivalent to: `return ref.fetch_add(operand, memory_ordering);`

```
T fetch_add(difference_type operand) const noexcept;
```

15 *Constraints:* `is_pointer_value` is true.

16 *Effects:* Equivalent to: `return ref.fetch_add(operand, memory_ordering);`

```
T fetch_sub(T operand) const noexcept;
```

17 *Constraints:* `is_integral_value || is_floating_point_value` is true.

18 *Effects:* Equivalent to: `return ref.fetch_sub(operand, memory_ordering);`

```
T fetch_sub(difference_type operand) const noexcept;
```

19 *Constraints:* `is_pointer_value` is true.

20 *Effects:* Equivalent to: `return ref.fetch_sub(operand, memory_ordering);`

```
T fetch_and(difference_type operand) const noexcept;
```

21 *Constraints:* `is_integral_value` is true.

22 *Effects:* Equivalent to: `return ref.fetch_and(operand, memory_ordering);`

```
T fetch_or(difference_type operand) const noexcept;
```

23 *Constraints:* `is_integral_value` is true.

24 *Effects:* Equivalent to: `return ref.fetch_or(operand, memory_ordering);`

```
T fetch_xor(difference_type operand) const noexcept;
```

25 *Constraints:* `is_integral_value` is true.

26 *Effects:* Equivalent to: `return ref.fetch_xor(operand, memory_ordering);`

```
T fetch_max(difference_type operand) const noexcept;
```

25 *Constraints:* `is_integral_value` is true.

26 *Effects:* Equivalent to: `return ref.fetch_max(operand, memory_ordering);`

```
T fetch_min(difference_type operand) const noexcept;
```

25 *Constraints:* `is_integral_value` is true.

26 *Effects:* Equivalent to: `return ref.fetch_min(operand, memory_ordering);`

```
T operator++(int) const noexcept;
```

27 *Constraints:* `is_integral_value || is_pointer_value` is true.

28 *Effects:* Equivalent to: `return fetch_add(1);`

```
T operator++() const noexcept;
```

29 *Constraints:* `is_integral_value || is_pointer_value` is true.

30 *Effects:* Equivalent to: `return fetch_add(1) + 1;`

```
T operator--(int) const noexcept;
```

31 *Constraints:* `is_integral_value || is_pointer_value` is true.

32 *Effects:* Equivalent to: `return fetch_sub(1);`

```
T operator--() const noexcept;
```

33 *Constraints:* `is_integral_value || is_pointer_value` is true.

34 *Effects:* Equivalent to: `return fetch_sub(1) - 1;`

```
T operator+=(T operand) const noexcept;
```

35 *Constraints:* `is_integral_value || is_floating_point_value` is true.

36 *Effects:* Equivalent to: `return fetch_add(operand) + operand;`

```
T operator+=(difference_type operand) const noexcept;
```

37 *Constraints:* `is_pointer_value` is true.

38 *Effects:* Equivalent to: `return fetch_add(operand) + operand;`

```
T operator+=(T operand) const noexcept;
```

39 *Constraints:* `is_integral_value || is_floating_point_value` is true.

40 *Effects:* Equivalent to: `return fetch_sub(operand) - operand;`

```
T operator-=(difference_type operand) const noexcept;
```

41 *Constraints:* `is_pointer_value` is true.

42 *Effects:* Equivalent to: `return fetch_sub(operand) - operand;`

```
T operator-=(T operand) const noexcept;
```

43 *Constraints:* `is_integral_value` is true.

44 *Effects:* Equivalent to: `return fetch_and(operand) & operand;`

```
T operator|=(T operand) const noexcept;
```

45 *Constraints:* `is_integral_value` is true.

46 *Effects:* Equivalent to: `return fetch_or(operand) | operand;`

```
T operator^=(T operand) const noexcept;
```

47 *Constraints:* `is_integral_value` is true.

48 *Effects:* Equivalent to: `return fetch_xor(operand) ^ operand;`

```
void wait(T old) const noexcept;
```

49 *Effects:* Equivalent to: `ref.wait(old, load_ordering);`

```
void notify_one() const noexcept;
```

50 *Effects:* Equivalent to: `ref.notify_one();` }

```
void notify_all() const noexcept;
```

51 *Effects:* Equivalent to: `ref.notify_all();` }

Memory Order Specific Atomic Refs [atomics.refs.bound.order]

```
// all freestanding
namespace std {
template<class T>
using atomic_ref_relaxed = atomic_ref_bound<T, memory_order_relaxed>;

template<class T>
using atomic_ref_acq_rel = atomic_ref_bound<T, memory_order_acq_rel>;

template<class T>
using atomic_ref_seq_cst = atomic_ref_bound<T, memory_order_seq_cst>;
}
```

§ 5.1.3 In [version.syn]:

Update the feature test macro `__cpp_lib_atomic_ref`.

§ 5.2 Atomic Accessors

§ 5.2.1 Add to the `mdspan` header synopsis in [mdspan.syn]

```
// [mdspan.accessor.default], class template default_accessor
template<class ElementType>
class default_accessor;
```

```
// [atomics.accessor.atomic], class template basic-atomic-accessor
template<class T, class ReferenceType> class basic-atomic-accessor; // exposition only
template<class T> using atomic_accessor = basic-atomic-accessor<T, atomic_ref<T>>;
template<class T> using atomic_accessor_relaxed = basic-atomic-accessor<T,
atomic_ref_relaxed<T>>;
template<class T> using atomic_accessor_acq_rel = basic-atomic-accessor<T,
atomic_ref_acq_rel<T>>;
template<class T> using atomic_accessor_seq_cst = basic-atomic-accessor<T,
atomic_ref_seq_cst<T>>;
```

```
// [mdspan.mdspan], class template mdspan
template<class ElementType, class Extents, class LayoutPolicy = layout_right,
class AccessorPolicy = default_accessor<ElementType>>
class mdspan;
```

§ 5.2.2 Put the following before [mdspan.mdspan]:

Class template `basic-atomic-accessor` [mdspan.accessor.atomic.basic]

General [mdspan.accessor.atomic.basic.overview]

```

template <class ElementType, class ReferenceType>
struct basic-atomic-accessor { // exposition only
    using offset_policy = basic-atomic-accessor;
    using element_type = ElementType;
    using reference = ReferenceType;
    using data_handle_type = ElementType*;

    constexpr basic-atomic-accessor() noexcept = default;

    template <class OtherElementType>
    constexpr basic-atomic-accessor(default_accessor<OtherElementType>) noexcept;

    template <class OtherElementType>
    constexpr basic-atomic-accessor(basic-atomic-accessor<OtherElementType,
        ReferenceType>) noexcept;

    constexpr reference access(data_handle_type p, size_t i) const noexcept;
    constexpr data_handle_type offset(data_handle_type p, size_t i) const noexcept;
};

```

- 1 Class basic-atomic-accessor is for exposition only.
- 2 basic-atomic-accessor meets the accessor policy requirements.
- 3 ElementType is required to be a complete object type that is neither an abstract class type nor an array type.
- 4 Each specialization of basic-atomic-accessor is a trivially copyable type that models semiregular.
- 5 $[0, n)$ is an accessible range for an object p of type data_handle_type and an object of type basic-atomic-accessor if and only if $[p, p+n)$ is a valid range.

Members [mdspan.accessor.atomic.basic.members]

```

template <class OtherElementType>
constexpr basic-atomic-accessor(default_accessor<OtherElementType>) noexcept {}

template <class OtherElementType>
constexpr basic-atomic-accessor(basic-atomic-accessor<OtherElementType,
    ReferenceType>) noexcept {}

```

- 1 *Constraints:* is_convertible_v<OtherElementType (*)[], element_type (*)[]> is true.

```
constexpr reference access(data_handle_type p, size_t i) const noexcept;
```

- 2 *Effects:* Equivalent to return reference(p[i]);

```
constexpr data_handle_type offset(data_handle_type p, size_t i) const noexcept;
```

- 3 *Effects:* Equivalent to return p + i;

Atomic accessors [mdspan.accessor.atomic.bounded]

```

namespace std {
    template <class ElementType>
    using atomic_accessor = basic-atomic-accessor<ElementType, atomic_ref<ElementType>>;

    template <class ElementType>
    using atomic_accessor_relaxed = basic-atomic-accessor<ElementType,
        atomic_ref_relaxed<ElementType>>;
}

```

```

template <class ElementType>
using atomic_accessor_acq_rel = basic-atomic-accessor<ElementType,
    atomic_ref_acq_rel<ElementType>>;

template <class ElementType>
using atomic_accessor_seq_cst = basic-atomic-accessor<ElementType,
    atomic_ref_seq_cst<ElementType>>;
}

```

§ 5.3 In [version.syn]

Add the following feature test macro:

```
#define __cpp_lib_atomic_accessors YYYYMMML // also in <mdspan>
```

§ 6 Acknowledgments

Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy’s National Nuclear Security Administration under Grant DE-NA-0003525.

This manuscript has been authored by UTBattelle, LLC, under Grant DE-AC05-00OR22725 with the U.S. Department of Energy (DOE).

This work was supported by Exascale Computing Project 17-SC-20-SC, a joint project of the U.S. Department of Energy’s Office of Science and National Nuclear Security Administration, responsible for delivering a capable exascale ecosystem, including software, applications, and hardware technology, to support the nation’s exascale computing imperative.

This research used resources of the Argonne Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC02-06CH11357.