

Trivial Relocatability For C++26

Proposal to safely relocate objects in memory

Document #: P2786R8
Date: 2024-10-15
Project: Programming Language C++
Audience: EWG, LEWG
Reply-to: Mungo Gill
<mgill83@bloomberg.net>
Alisdair Meredith
<ameredith1@bloomberg.net>
Joshua Berne
<jberne4@bloomberg.net>
Corentin Jabot
<corentinjabot@gmail.com>
Lori Hughes
<lori@lorihughes.com>

Contents

1 Abstract	4
2 Revision History	5
R8: October 2024 (pre-Wrocław mailing)	5
R7: September 2024 (midterm mailing)	5
R1–6: January 2023 – April 2024	5
3 Document Conventions	6
3.1 Typography	6
3.2 Definitions	6
3.3 Core-language additions	6
3.4 Library additions	6
4 Introduction	8
5 Basic Ideas	9
5.1 Trivial relocatability	9
5.2 Replaceability	9
5.3 Examples of trivially relocatable and replaceable types	10
5.4 Examples of trivially relocatable types that are <i>not</i> replaceable	10
5.5 Examples of replaceable types that are not <i>trivially</i> relocatable	10
5.6 Examples of Standard Library types that must defer to the implementation	10
5.7 Independent features	10
6 Technical Background	11
7 Basic Design Principles	12
7.1 Create a feature for users, not just the Standard Library	12
7.2 A formal specification is built around object lifetimes	12
7.3 Behavior must be reliable	12

7.4	Guard against accidental undefined behavior	12
8	Core Proposal	13
8.1	Trivial relocatability	13
8.2	Replaceability	17
8.3	Optimizing <code>std::swap</code>	18
9	Plans for Library Update	22
9.1	Immediate updates	22
9.2	ABI concerns	22
9.3	Specific follow-ups	23
10	Use Cases	24
10.1	Optimizing <code>std::vector</code> at run time	24
10.2	Optimizing <code>std::optional</code> to be trivially relocatable and replaceable	24
11	Implementation Experience	26
12	FAQ	27
12.1	Is <code>void</code> trivially relocatable?	27
12.2	Are reference types trivially relocatable?	27
12.3	Why can a class with a reference member be trivially relocatable?	27
12.4	Are <i>cv</i> -qualified types, notably <code>const</code> types, trivially relocatable?	27
12.5	Can <code>const</code> -qualified types be passed to <code>trivially_relocate</code> ?	27
12.6	Can types that are not implicit-lifetime types be trivially relocatable?	27
12.7	Why are virtual base classes not trivially relocatable?	27
12.8	What happens if a relocate operation throws?	28
12.9	Why do deleted special members inhibit implicit trivial relocatability?	28
12.10	Can the compiler transform argument passing with trivial relocation?	28
12.11	Can the Standard Library containers use this new feature internally?	28
12.12	Do implementations need to mark classes <code>memberwise_trivially_relocatable</code> to benefit?	30
12.13	Which Standard Library types are safe for the <code>swap_value_representations</code> function?	30
12.14	Can I mark as trivially relocatable a type that is not replaceable?	30
12.15	Can I mark a type as replaceable but not trivially relocatable?	30
12.16	What happened to the predicates for the contextual keywords?	30
12.17	Why is copy-replacement unsupported?	31
12.18	Why is marking a class that can <i>never</i> be trivially relocatable not ill-formed?	31
12.19	Why is there no <code>is_trivially_replaceable</code> trait?	31
12.20	Is it UB to mark a nonconforming type as trivially relocatable?	31
12.21	Is it UB to mark a nonconforming type as replaceable?	31
13	Illustrative Examples	33
13.1	Unconstrained vector	33
13.2	Standard vector	33
13.3	Conforming implementation of a trivially relocatable <code>std::optional</code>	33
13.4	C++26 implementation using internal array	37
14	Proposed Wording	41
14.1	Add new identifiers with a special meaning	41
14.2	Specify trivially relocatable types	41
14.3	Address trivial relocation of lambdas	41
14.4	Update grammar to support <code>memberwise</code> contextual keywords	42
14.5	Specification for trivially relocatable classes	43
14.6	Add feature macros	44
14.7	Library wording	44

14.8 Add new type traits	45
14.9 Specify the compiler-magic functions	45
14.10Require the optimization for <code>std::swap</code>	47
15 Acknowledgements	48
16 References	48

1 Abstract

Many types in C++ cannot be trivially moved or destroyed but do support trivially moving an object from one location to another by copying its bits — an operation known as *trivial relocation*. Some types even support bitwise swapping, which requires replacing the objects passed to the `swap` function, without violating any object invariants. Optimizing containers to take advantage of this property of a type is already in widespread use throughout the industry but is undefined behavior as far as the language is concerned. This paper provides a mechanism to annotate types as having the appropriate properties to be eligible for these optimizations, along with library interfaces to make use of them in a well-defined manner.

2 Revision History

R8: October 2024 (pre-Wrocław mailing)

- Extracted [Document Conventions](#) to above the Introduction
- Added [Basic Ideas](#), a higher-level introduction to the new features
- Listed example types for the different type categories defined by this paper
- Corrected implementation of several example functions
- Extended the FAQ
 - What happens if a relocate operation throws?
 - Why is there no `is_trivially_replaceable` trait?
 - Is it UB to mark a nonconforming type as trivially relocatable?
 - Is it UB to mark a nonconforming type as replaceable?
- Explained how to apply the new features to optimize vectors

R7: September 2024 (midterm mailing)

- Significant redrafting since [\[P2786R6\]](#) to better address EWG and LEWG concerns
- Simplified presentation and discussion of trivial relocatability (for detailed history, consult [\[P2786R6\]](#))
- Integrated discussion of `swap`; the present paper supersedes [\[P3239R0\]](#)
- Behavior changes since [\[P2786R6\]](#)
 - User-provided move assignment now prevents a type from being implicitly trivially relocatable
 - The contextual keyword
 - gets a new name, `memberwise_trivially_relocatable`, to better reflect revised semantics
 - is opt-in only
 - deduces relocatability on bases and members
 - No predicate follows the contextual keyword, so no mechanism for opting out is present
 - A new `relocate` function additionally supports nontrivial types and constant evaluation
- Behavior changes since [\[P3239R0\]](#)
 - Clarifies that trivial swappability is based on being replaceable and trivially relocatable
 - Proposes a new contextual keyword, `memberwise_replaceable`
 - Proposes optimizing `std::swap`, using the new properties, and the `swap_value_representations` function

R1–6: January 2023 – April 2024

Early versions of this paper were careful to include comparison and contrast with other papers in this space. That progress is archived by [\[P2786R6\]](#).

The evolution groups requested a clean draft that presents just our proposal and integrates our follow-up papers such that a single coherent design is presented, and this revision, R7, is the response to that request.

3 Document Conventions

3.1 Typography

Throughout this paper, a **bold typeface** will be used for terms defined herein and ***bold italicized typeface*** for terms of art defined herein; the proposed wording, however, will use the conventions of the Standard.

3.2 Definitions

We define a **relocation operation** of a source object as one that ends the lifetime of that object and starts the lifetime of a new object at a new location. Importantly, the destructor is not necessarily run by a **relocation operation**. For types in which move construction and destruction are supported, a relocation can be accomplished by constructing an object in the new location from an xvalue referring to the source object, followed by invoking the destructor of the source object.

We define a **trivial relocation operation** as a **relocation operation** accomplished by performing a bitwise copy of its *object representation* to a new memory location that ends the lifetime of that source object — just as if that (source) object’s storage were used by another object (6.7.3 [basic.life]p5) — and starts the life of a new object at the new location. Importantly, nothing else is done to the source object; in particular, *its destructor is not run*. This operation will typically be semantically equivalent to a nontrivial **relocation operation** performed via move construction and destruction (though exceptions, while not encouraged, are not expressly forbidden).

We define **replacement** of a target object by a source object as destroying the target object immediately followed by move construction into the location of the target object from the source object. For many types, this operation is semantically equivalent to a move-assignment operation from the source object to the target object.

3.3 Core-language additions

We propose a new Core-language definition for a ***trivially relocatable type***. This new definition is inspired by the recursive nature and handling of special member functions used in the definition of a *trivially copyable* type. A ***trivially relocatable type*** is a scalar type, a ***trivially relocatable class***, an array of such types, or a cv-qualified version of such a type. A class will be implicitly ***trivially relocatable*** if all its bases and members are ***trivially relocatable*** and none of its eligible special member functions are user provided; a contextual keyword will signify that a class may still be ***trivially relocatable*** even if it has user-provided special member functions.

We similarly propose a new Core-language definition for a ***replaceable type*** in which a class will be a ***replaceable class*** if all its bases and members are ***replaceable*** and none of its relevant special member functions are user provided; a contextual keyword will signify that a class may still be a ***replaceable class***, even if it has those user-provided special member functions.

Because **replaceability** is explicitly about the equivalence of assignment with destruction followed by construction, we do not decide that a type is *implicitly replaceable* when the special member functions selected for those operations are user provided.

3.4 Library additions

The Standard Library APIs to support **trivial relocation** comprise

- a type trait to detect **trivial relocatability**
- a function, `trivially_relocate`, that performs relocation on a range of objects by moving their bytes (similarly to `memmove`) while starting the lifetime of the destination objects and ending the lifetime of the source objects
- a user-facing function, `relocate`, that emulates relocation by using the move-and-destroy idiom for types that are not ***trivially relocatable*** and delegates to `trivially_relocate` for types that are ***trivially relocatable***

To support common use cases, both `trivially_relocate` and `relocate` are specified to support overlapping ranges.

Further, to take advantage of the ability to specify that a type is ***replaceable***, we propose the following additional changes to Standard Library APIs:

- A type trait to detect **replaceability**
- A new magic function, `swap_value_representations`, that swaps just the value representations of the members of objects without impacting the dynamic types of the complete objects, typically by copying a range of bytes that excludes the vtable pointers and tail padding bytes
- An update to the primary `std::swap` template to use **replacement** optimization when it is supported

Finally, we propose modifications to Standard Library wording to describe when Standard Library types are allowed and expected to have various properties, including **trivial relocatability** and **replaceability**.

4 Introduction

Containers in C++, in particular those like `std::vector` and `std::deque` that manage objects within a range of continuous storage, live and die by the efficiency with which they can move objects around. One of the most common fundamental steps in many of the operations these types perform is that of relocation — taking an element at one location in memory and creating a new element at a different location in memory with the same value and then destroying that original value.

Many frequently used libraries have long recognized that, for many types, the two nontrivial steps of move construction and destruction often combine into a single operation that can be accomplished by a simple bitwise copy followed by discarding the source object instead of evaluating its destructor. Much of the work a move constructor might do to the source object, such as setting pointers to owned data to `nullptr`, is done only to make sure the destructor that will eventually run knows that no data is present that it is still responsible for freeing. By taking advantage of the knowledge that certain types can be relocated by simply copying bits, complex operations that can involve the invocation of many user-provided special members functions can be replaced by single calls to `memcpy`, realizing huge benefits in performance.

The problem, of course, is that moving objects in this fashion that are not trivially copyable violates the C++ object model and is undefined behavior. In this paper, we propose a mechanism to fix that problem.

- Types with no user-provided special member functions are identified as *trivially relocatable* based on whether their bases and nonstatic data members are *trivially relocatable*.
- Users are able to further identify those types whose user-provided special member functions compound into a trivial operation by marking those types with a new class specifier, `memberwise_trivially_relocatable`.
- The Standard Library then provides tools to safely perform relocation operations, both trivial and non-trivial.

Beyond containers, algorithms in C++ also build on a related basic operation, `std::swap`. While a first glance might imply that `std::swap` could be optimized by simply performing a series of relocation operations, such an approach would have major semantic differences from the `std::swap` we have today. Relocation is, by definition, about creating one object while destroying another at the same time. `std::swap`, however, is an operation composed of construction of a temporary and *move assignment* to the function's operands and has no natural and semantically equivalent definition in terms of relocation.

To solve this problem with swap, we identify the property that types require to have a `std::swap` implementation with the same semantics yet be implemented in terms of just move construction and destruction. That property is **replaceability**; i.e., a type declares that assignment from a source object is semantically equivalent to destruction followed by move construction from that source object. We will show that the class of types that have both **trivial relocatability** and **replaceability** is exactly the class of types for which we can safely perform a swap operation by exchanging the bytes that make up the value representations of two objects.

To incorporate this concept of **replaceability** into the language, we propose further additional changes.

- Types are implicitly recursively *replaceable* if all their members and bases are.
- Through the use of the `memberwise_replaceable` specifier on a class, users can specify that their user-provided constructors, destructors, and assignment operators still satisfy the appropriate equivalence specified by being *replaceable* as long as all members and bases also have this property.
- A new utility function, `std::swap_value_representations`, is provided for swapping the bytes of a type that is both *trivially relocatable* and *replaceable*.
- The standard template `std::swap` is updated to make use of bitwise swapping when invoked with types that have both properties.

Put together, we hope this proposal provides a complete picture of how to incorporate into the C++ Standard, in an understandable and effective manner, bitwise operations that are already performed by many libraries in the industry.

5 Basic Ideas

This paper introduces two new complementary but independent notions into C++.

5.1 Trivial relocatability

Relocation is the act of moving an object from one memory location to another and is typically achieved by calling the move constructor to make a new object at the new location followed by the destructor on the original object.

A type is *trivially relocatable* if it can be relocated by copying the bytes of its *object representation* from the old location to the new and the lifetime of the original object can be ended *without* running its destructor. However, C++ object lifetimes do not currently permit types (with the exception of those few types that meet the strict requirements of trivially copyability) to be relocated by means of byte copying (**trivial relocation**).

If the object model were to allow it, most C++ types could safely be trivially relocated. The two known exceptions are types that maintain an internal pointer to a data member and types that register their presence in an external registry that must point back to the object.

This paper proposes adding

- a Core-language definition for *trivially relocatable types*
- a way to explicitly mark types *trivially relocatable* when that cannot be deduced by a compiler
- a type trait to report whether a type is *trivially relocatable*
- a “compiler-magic” function to perform **trivial relocation**, respecting object lifetimes
- a user-facing `relocate` function that can be safely used with types that are not *trivially relocatable*

All the features of **trivial relocation** have decades of experience in which code has relied on compilers not reacting to the use of undefined behavior when copying nontrivial types breaks the C++ object lifetime rules.

5.2 Replaceability

Replaceability is a semantic property of a type, where move assignment is isomorphic to destroy then move-construct. Just like in **trivial relocatability**, a compiler cannot deduce whether a type is *replaceable* if the user provides a move-assignment operator, move constructor, or destructor without extra guidance from the user.

In many cases, a library would like to require or assume **replaceability**, such as when moving elements around a `std::vector` when inserting or erasing elements.

One important place where knowledge of **replaceability** can lead to performance improvements is in the implementation of `std::swap`, which is one of the most widely used library functions and is an implementation detail of many standard algorithms.

If a type is both *replaceable* and *trivially relocatable*, then its swap operation could be simplified to swapping the *value representations* of two objects. Note that the *value representation* is generally a subset of the bytes in the *object representation*, and the difference matters when trying to call `swap` on potentially overlapping data members in a class, which is discussed in more detail later in this paper.

This paper proposes adding

- a Core-language definition for *replaceable* types
- a way to mark types as *replaceable* if the compiler cannot deduce that property
- a type trait to report whether a type is *replaceable*
- a “compiler-magic” function to swap the value representation of two *trivially relocatable*, *replaceable* objects

and then requiring `std::swap` functions to use this optimization for such *replaceable* types.

5.3 Examples of trivially relocatable and replaceable types

If we assume the default template arguments, we would expect the following Standard Library types to be both *trivially relocatable* and *replaceable*:

- `std::shared_ptr`
- `std::future`
- `std::vector`

We would expect the following types to be both *trivially relocatable* and *replaceable* if all their template arguments are both *trivially relocatable* and *replaceable*:

- `std::pair`
- `std::tuple`

5.4 Examples of trivially relocatable types that are *not* replaceable

A variety of types, while *trivially relocatable*, do not maintain the invariants of **replaceability**:

- `std::tuple<T &>`
- `std::pmr` containers
- Any class with a `const` data member

In many contexts, relocation of such types is desirable, especially in user-defined data structures beyond the reach of the Standard Library.

5.5 Examples of replaceable types that are not *trivially* relocatable

The main example in this category is Standard Library containers with debug iterators that track their container with a back-pointer or some other registry, although we can easily imagine user-supplied types with similar constraints. Note that some implementations of `std::basic_string` fall into this category, where the short string optimization maintains a pointer to its internal short buffer.

This category of types would meet preconditions for algorithms in which the semantics of **replaceability** are important, and they might be enforced by the equivalent of Mandates, Constraints, or Preconditions in users' libraries.

5.6 Examples of Standard Library types that must defer to the implementation

We would expect the quality of implementation would decide whether the following types are *trivially relocatable and replaceable* or are just **replaceable**:

- `std::basic_string`, depending on whether the short string optimization maintains an internal pointer
- `std::list`, depending on whether the sentinel node is a nonstatic data member

5.7 Independent features

From the variety of types and usage examples above, we see that while **trivial relocation** and **replacement** are often used together, each has important use cases and neither can be built on top of the other.

6 Technical Background

Some very specific uses of terminology from the C++ Standard are important to understand when reading this proposal and are quickly summarized here.

For decades, C++ developers have been optimizing low-level data structures, such as their own **vector**-like types, by byte-wise copying objects from one location to another, even though doing so is often UB; see earlier papers¹ for rationale.

Earlier revisions of this paper initially proposed language and library extensions, termed **trivial relocation**, to make writing such code well defined and was forwarded to Core where it received a strong review that challenged our assumptions about copying bytes. From the perspective of the C++ abstract machine, we should not be making assumptions about in-memory representations — that is the compiler’s job — and should limit ourselves to copying the *object representation*, leaving the compiler itself to optimize copying and moving the object representations to efficient memory copying operations.

The Core review proceeded in parallel to the LEWG review, which subsequently sent the proposal back to EWG, asking for a more complete handling of bitwise operations, notably optimizations for byte-wise swaps. While single swaps are unlikely to see a realistic impact on performance, the many algorithms in the Standard Library, such as `std::rotate`, that build on `swap` as a foundational operation increased interest in seeing `swap` benefit from these efforts.

Implementing `std::swap` as a series of **trivial relocation operations**, however, must contend with the fact that such operations will end the lifetimes of both parameters and create new objects in their place. In general, this is incorrect for any type for which such destruction and recreation would not create a new object with the same invariants, such as any type with a `const` or reference member, even though many such types could be *trivially relocatable*.

Even for types for which recreating new objects in place would not violate the invariants of the old objects, such as any *replaceable type*, our design must address complexities in the Core language. In general, turning a `swap` operation into a sequence of **trivial relocation operations** is not possible due to the Standard’s specification for *transparent replacement*, which handles the case of potentially overlapping members in a class, such as when the compiler optimizes layout for empty base classes or members annotated `[[no_unique_address]]`. Since this concern is not related to the type of the arguments but rather to the specific objects being passed to `std::swap` and which complete objects they are members of, we must handle these edge cases correctly lest we risk breaking much existing code at run time and without warning.

That is where we introduce our final term from the Standard: *value representation*. Where the object representation comprises all the bytes that an object occupies (including all padding bytes, vtables, and so on), the value representation denotes just the bytes that contribute to an object’s state. Therefore, an empty class has a nonzero size but a zero-length value representation. Thus, the parts of this paper that allow optimization of swap operations refer to value representations, while the parts that allow for **trivial relocation** optimizations refer to object representation. Bearing this distinction in mind will be important.

¹Much rationale related to **trivial relocation** can be found in [P2786R6], and early discussion of handling `swap` is covered in [P3239R0].

7 Basic Design Principles

7.1 Create a feature for users, not just the Standard Library

Efficient implementation of many data structures often needs a means to efficiently move and exchange objects that those data structures are managing, especially for data structures that manage their elements in contiguous storage or in some other location that is not a node at the end of a pointer. However, such object manipulation is also a sharp tool that is not expected to see much use other than optimizing the internal management of data structures.

Overall, our goal is to provide the needed — possibly sharp-edged — tools for use by container implementers, while users are given a usable API to manage when these optimizations are safe and correct to apply to their types.

7.2 A formal specification is built around object lifetimes

C++ has a well-specified object model that is important to optimizers, sanitizers, and analysis tools alike. Such tools must reason about object lifetimes and, importantly, minimize the doubt created for developers regarding that reasoning leading to false positives or false negatives when seeking to optimize or alert users.

7.3 Behavior must be reliable

No freedom for quality of implementation (QoI) in semantics is an important quality that builds on the well-specified object model.

7.4 Guard against accidental undefined behavior

The new Library APIs support only types that would produce well-defined behavior. The specification prefers *Mandates* clauses to *Constraints*: clauses since SFINAE behavior carries no expected benefit and is likely to produce error messages with less useful information.

8 Core Proposal

Our core proposal comprises two parts: **trivial relocation** and **replaceability**, each including the library primitives that are necessary for well-defined use. **Trivial relocation** is a technique already widely used in the industry, and **replaceability** is a more novel property that must be identified to apply bitwise copy optimizations to `std::swap`.

8.1 Trivial relocatability

To ensure that libraries taking advantage of the *trivially relocatable* semantic do not introduce undefined behavior, the model of lifetimes for objects must be extended to allow for relocation of *trivially relocatable types*. Since the compiler cannot know if a specific `memcpy` or `memmove` call is intended to duplicate (or to move) an object, we propose introducing a new function template, `std::trivially_relocate`, that is restricted to *trivially relocatable types*. The purpose of the new function template is to efficiently move the *object representation*, typically with a call to `memmove` that also signifies to the compiler (and other analysis tools) that the lifetime of the new object(s) has begun — similar to calling `start_lifetime_as` on the destination location(s) — and that the lifetime of the original object(s) has ended (without running destructors).

This design deliberately puts all compiler-magic and Core-language interaction dealing with the object lifetimes into a single place, rather than into a number of different `relocate`-related overloads. Note that users are not permitted to copy the bytes to perform a relocation themselves, unlike with trivial copyability, although byte copies would continue to work for trivially copyable types.

8.1.1 New type category: Trivially relocatable

To better integrate language support, we further propose that the language can detect types as *trivially relocatable* where all their bases and nonstatic data members are, in turn, *trivially relocatable*: The constructor selected for construction from a single rvalue of the same type is neither user provided nor deleted, the same applies for the assignment operator for rvalues, and their destructor is neither user provided nor deleted. Conceptually, this definition combines the rules we would follow if there was a new user-definable special member function for relocation *and* when that operation would be trivial.

Note that our notion of relocation relies on being semantically equivalent to move construction of the target followed by destruction of the source. Even though it is not involved in this definition, we still consider assignment operations when deciding if a type is *implicitly trivially relocatable* for the same reasons that we consider assignment when deciding if a type should have an implicitly declared move constructor; any existing type with a particular set of user-provided special member functions should not begin to have new operations considered valid for it if those operations might subvert expectations due to compiling with a new language Standard.

8.1.2 New keyword and explicit rule

Without an opt-in mechanism, the only types that would be implicitly *trivially relocatable* would be those that are already trivially copyable, an important yet relatively small subset of the full universe of types in C++. To enable **trivial relocatability** on the many more interesting types that have nontrivial special member functions, explicitly marking such types must be possible. This marking is needed for only user-defined class types (including unions); hence, we propose adding a new contextual keyword, `memberwise_trivially_relocatable`, as part of the class definition, similar to how `final` applies to classes:

```
struct X; // Forward declaration does not admit `final`.
struct X final {}; // Class definition admits `final`.
struct Y memberwise_trivially_relocatable {}; // New contextual keyword placed like `final`.
```

We propose one new contextual keyword, `memberwise_trivially_relocatable`, that can be placed in a class-head (on a class definition) to indicate that a type's special operations do nothing that would violate the implicit rule that would make a type *trivially relocatable*.

By means of the `memberwise_trivially_relocatable` specification, a class will be determined to be *trivially relocatable* if, according to the implicit rules for a *trivially relocatable class*, the class would be *trivially relocatable* if the presence of user-declared special member functions were ignored.

Users considering whether to apply this keyword to a given type that has user-provided special member functions must simply inspect their move constructor and destructor and decide if, when applied together as part of a relocate operation, they have no net effect. Common examples include many types.

- For a resource-owning type, such as `std::unique_ptr`, the newly constructed object will have the same bits as the source object, the source object will have its pointer member set to `nullptr`, and the source object destructor will do nothing because, by the time it runs, that member will be `nullptr`. Simply copying the bytes and discarding the source object achieves the same semantic effect.
- A reference-counting type, however, might increment a count by one when constructing the target object and then decrement that same count by one when destroying the source object. Combining these operations clarifies that the constructor and destructor negate one another.

8.1.3 New type trait: `is_trivially_relocatable`

To expose the **relocatability** property of a type to library functions seeking to provide appropriate optimizations, we propose a new trait, `std::is_trivially_relocatable<T>`, which enables the detection of **trivial relocatability**:

```
template< class T >
struct is_trivially_relocatable;

template< class T >
constexpr bool is_trivially_relocatable_v = is_trivially_relocatable<T>::value;
```

The `std::is_trivially_relocatable<T>` trait has a base characteristic of `std::true_type` if `T` is *trivially relocatable* and has `std::false_type` otherwise.

Note that the `std::is_trivially_relocatable` trait reflects the underlying property that a type has, and like all similar traits in the Standard Library, it must not be *user specializable*. Compilers themselves are expected to determine this property internally and should not introduce a library dependency such as by instantiating this type trait.

We expect that the `std::is_trivially_relocatable` trait shall be implemented through a compiler intrinsic, much like `std::is_trivially_copyable`, so the compiler can use that intrinsic when the language semantics require **trivial relocatability**, rather than requiring actual instantiation (and knowledge) of the Standard Library trait. The trait must always agree with the intrinsic since users do *not* have permission to specialize standard type traits (unless explicitly granted permission for a specific trait).

We see no particular need to separately detect whether a type has attempted to make itself *trivially relocatable* with `memberwise_trivially_relocatable`.

8.1.4 New relocation function: `trivially_relocate`

As stated in “[Library additions](#),” we are proposing a new function, `trivially_relocate`, which is the unique entry point into the core magic that tracks and manages object lifetimes in the abstract machine:

```
template <class T>
T* trivially_relocate(T* begin, T* end, T* new_location) noexcept;
{
    static_assert( is_trivially_relocatable_v<T> && !is_const_v<T> );
    // ... (platform-provided implementation)
}
```

This function template *mandates* that `is_trivially_relocatable_v<T> && !is_const_v<T>` is `true` and has *preconditions* that `end` is reachable from `begin`. Its *postcondition* is that new objects in the range `[new_location,`

`new_location + sizeof(T) * (end - begin))` have the same object representation as the objects originally in the range `[begin, end)` and that the objects originally in the range `[begin, end)` have ended their lifetime, all accomplished without running any destructors or other clean-up code. Overlapping ranges shall be supported.

On most platforms, this template is functionally equivalent to

```
memmove(new_location, begin, sizeof(T) * (end - begin));
```

However, unlike `memmove` on its own, this function template is restricted to *trivially relocatable types* rather than to implicit lifetime types.

Note that, consistent with its low-level purpose often tied to move semantics, this function is denoted with `noexcept` despite having a narrow contract regarding valid and reachable pointers.

In addition to performing `memmove`, the function also has the following two important effects that matter to the abstract machine but have no apparent physical effect (i.e., these effects do not change bits in memory), much like `std::launder`.

1. The `trivially_relocate` function ends the lifetime of the objects `*begin`, `*(begin+1)`, ..., through to `*(end-1)`. This ending of the objects' lifetimes means accessing these objects or attempting to run their destructors will be *undefined behavior*.
2. The `trivially_relocate` function begins the lifetime of the objects `*new_location`, `*(new_location+1)`, ..., through to `*(new_location+end-begin-1)`. If any of the objects or their subobjects are unions, they have the same active elements as the corresponding objects in the range `[begin, end)`.
3. These operations are a single action, and for any locations where an overlap occurs between the source and target ranges, an existing object will be destroyed and a new object will be created in its place.

The current library-level mechanism to start the lifetime of an object without invoking a constructor is `std::start_lifetime_as`, a function that works for only implicit lifetime types that must have trivial default constructors. *Trivially relocatable types*, however, include a much wider range of types, including many that establish and maintain invariants in their special member functions and thus cannot be implicit lifetime types.

A tool for ending lifetimes is similarly unavailable in the Standard Library today. This task can be accomplished by reusing the storage of an object, but that requires modifications of some sort.

The `trivially_relocate` function, therefore, is interacting with the abstract machine in ways that are not currently available. Importantly, for many of the types we are concerned with (e.g., `std::vector`, `std::unique_ptr`, and so on), the component steps of the **relocation operation** are decidedly not trivial, so we are compelled to make this single function responsible for the needed compiler magic.

To remove the need for a larger family of functions and avoid overly limiting cases in which **trivial relocation** might be applied, the `trivially_relocate` function is intended to support overlapping source and destination ranges, just like `memmove`. If the ranges are overlapping, the implementation must take care around the management of the lifetime of objects relocated out of or into the overlap.

Note that this initial proposal does not provide a single-object relocation function since our primary motivation is to optimize relocating objects in bulk, which is expected to be the common use case. Adding single-object `trivially_relocate` functions would be easy, but the effect can be achieved by calling the proposed function with a range of a single object.

8.1.5 New library function: `std::relocate`

The function `trivially_relocate` is a sharp tool that requires compiler magic to implement, and the user must write an alternative code path for types that are not *trivially relocatable*. General relocation that supports both trivial and nontrivial relocation is, however, a subtle and tedious function to implement correctly, and we do not want to force all users to reimplement this function.

Therefore, we propose an additional user-friendly, general-purpose relocation function, `std::relocate`, that will use `trivially_relocate` for *trivially relocatable types* and otherwise `relocate` elements by calling the move

constructor to move each object, followed by their destructor. This function must correctly order its moves to support overlapping ranges, just like `trivially_relocate`.

In addition, `std::relocate` is `constexpr` to support easy implementation of `constexpr` containers like `std::vector`. Adding such support means that in addition to checking whether a type is *trivially relocatable* before calling `trivially_relocate`, we must also have an `if constexpr` path that does *not* call `trivially_relocate` during constant evaluation:

```
template <class T>
constexpr
T* relocate(T* begin, T* end, T* new_location)
{
    static_assert(is_trivially_relocatable_v<T>
                  || is_nothrow_move_constructible_v<T>);

    // When relocating to the same location or an empty range, do nothing.
    if (begin == new_location) return end;
    if (begin == end) return new_location;

    // Then, if we are not evaluating at compile time and the type supports
    // trivial relocation, delegate to `trivially_relocate`.
    if ! constexpr {
        if constexpr (is_trivially_relocatable_v<T>) {
            return trivially_relocate(begin, end, new_location);
        }
    }

    if constexpr (is_move_constructible_v<T>) {
        // For nontrivial relocatable types or any time during constant
        // evaluation, we must detect overlapping ranges and act accordingly,
        // which can be done only if the type is movable. Note that trivially
        // relocatable types are allowed to have throwing move constructors, and
        // any throwing move that occurs in this branch will cause constant
        // evaluation to fail.

        if ! constexpr {
            // At run time, when there is no overlap, we can, using other Standard
            // Library algorithms, do all moves at once followed by all destructions.
            if (less{}(end, new_location) || less{}(new_location + (end - begin), begin)) {
                T* result = uninitialized_move(begin, end, new_location);
                destroy(begin, end);
                return result;
            }
        }

        if (less{}(new_location, begin) || less{}(end, new_location)) {
            // Any move to a lower address in memory or any nonoverlapping move can be
            // done by iterating forward through the range.
            T* next = begin;
            T* dest = new_location;
            while (next != end) {
                ::new(dest) T(move(*next));
                next->~T();
                ++next; ++dest;
            }
        }
    }
}
```



```

}
else {
    // When moving to a higher address that overlaps, we must go backward through
    // the range.
    T* next = end;
    T* dest = new_location + (end-begin);
    while (next != begin) {
        --next; --dest;
        ::new(dest) T(move(*next));
        next->~T();
    }
}

return new_location + (end-begin);
}

// The only way to reach this point is during constant evaluation where type `T`
// is trivially relocatable but not move constructible. Such cases are not supported,
// so we mark this branch as unreachable.
unreachable();
}

```

8.2 Replaceability

In addition to **trivial relocation**, we introduce the orthogonal notion of **replaceability**. An object of type *T* is **replaceable** by an object of type *U* if destroying the object of type *T* and reconstructing an object of type *T* in its place from an xvalue of type *U* is equivalent to assigning to the original object of type *T* with an xvalue of type *U*. Note that **replacement** updates an object's value, so **const**-qualified objects are never **replaceable**.

Replaceability is an important property when we want to transform **relocation** into assignment or vice versa. Containers such as `std::vector` already make a general assumption that all types are **replaceable**, but `std::swap` does not make such an assumption, so we provide a mechanism to identify those types with this new property.

8.2.1 New type category: Replaceable type

A type *T* is a **replaceable type** if every object of type *T* is **replaceable** by every other object of type *T*. Note that **replaceable types** must be object types; function types, reference types, and `void` are never **replaceable**.

A cv-unqualified type *T* will implicitly be a **replaceable type** if all its bases and nonstatic members are **replaceable types** and if it has no user-provided move constructor, move-assignment operator, nor destructor.

8.2.2 New keyword and explicit rule

To enable **replaceability** to be useful for classes with user-provided special member functions, explicitly marking class (including union) types as potentially **replaceable** must be possible (just like for **trivially relocatable types**). To that end, we propose adding a new contextual keyword, `memberwise_replaceable`, as part of the class definition (mirroring the design of `memberwise_trivially_relocatable`).

```

struct X;                                // Forward declaration does not admit `final`.
struct X final {};                       // Class definition admits `final`.
struct Y memberwise_trivially_relocatable {}; // New contextual keyword placed like `final`.
struct Z memberwise_replaceable {};       // New contextual keyword placed like `final`.

```

A class can be marked with both `memberwise_trivially_relocatable` and `memberwise_replaceable`; in fact, we expect many uses of `memberwise_replaceable` to also require `memberwise_trivially_relocatable`.

8.2.3 New type trait: `is_replaceable`

To expose the **replaceability** property of a type to library functions seeking to provide appropriate optimizations, we propose a new trait, `std::is_replaceable<T>`, that enables the detection of *replaceable types*:

```
template< class T >
struct is_replaceable;

template< class T >
constexpr bool is_replaceable_v = is_replaceable<T>::value;
```

The `std::is_replaceable<T>` trait has a base characteristic of `std::true_type` if `T` is *replaceable* and `std::false_type` otherwise.

Note that the `std::is_replaceable` trait reflects the underlying property that a type has, and like all similar traits in the Standard Library, it must not be *user specializable*. Compilers themselves are expected to determine this property internally and should not introduce a library dependency such as by instantiating this type trait.

Note that we expect that the `std::is_replaceable` trait shall be implemented through a compiler intrinsic, much like `std::is_trivially_copyable`, so the compiler can use that intrinsic when the language semantics require **replaceability**, rather than requiring actual instantiation (and knowledge) of the Standard Library trait. The trait must always agree with the intrinsic since users do *not* have permission to specialize standard type traits (unless explicitly granted permission for a specific trait).

8.3 Optimizing `std::swap`

`std::swap` differs significantly from **trivial relocation** in several ways; `std::swap` is an existing well-specified function with a wide contract, and it starts and ends with two valid objects and cannot end the lifetimes of either without vastly changing its current expected behavior. We discuss such constraints and how they led to our proposal of a new magic function based on notions of **trivial relocation** and **replacement**.

8.3.1 Swapping values rather than representations

`std::swap` exchanges *values* that are essentially defined by the move constructor and move-assignment operator. Using operations to directly swap the bytes would result in swapping the whole *object representations*, which are not always the same thing.

8.3.2 Object lifetime and invariants

Fundamental to the C++ object model is the ability for the structure of an object, including the types of its members and its special member functions, to enable developers to maintain object invariants through an object's entire lifetime. Passing an object to a function by modifiable lvalue reference is never expected to invalidate all such invariants. In addition, ending the lifetime of an object also invalidates pointers and references to that object, a process colloquially known as *pointer end-zap*.

In certain cases, an object can be destroyed and a new object can be created in place without encountering pointer end-zap or other issues, and those are cases in which an object can be *transparently replaced*. This scenario occurs when the objects have the same type, are not `const`-qualified, and are *complete* objects or not *potentially overlapping* subobjects.

Many of these situations are not, however, requirements for using `std::swap` today, so we must ensure that any changes to `std::swap` will still work for objects that do not meet these criteria.

Note that whether one object can be transparently replaced by another is not a property of the type but depends on the context of how that type is used, which is additional information that cannot be easily passed through a function call.

8.3.3 Object representation vs. value representation

The *object representation* of an object is the set of bytes the compiler uses to store that object. To allow for alignment, this set of bytes is often larger than simply summing all the sizes of that object's member subobjects and introduces padding bytes. The *value representation* of an object is just those nonpadding bytes that store information. For example, an *empty* class has an object representation of at least one byte, which is important for several reasons, one of which is to correctly represent arrays of empty types. On the other hand, the value representation of an empty type is zero bytes, i.e., there is no value representation.

8.3.4 Toward a well-defined function

If we specify a new library function to swap the value representations of the direct and indirect members (rather than the object representations) of two objects of the same type, we expect to preserve the properties that matter to the abstract machine that defines C++. In particular, for empty types, we will swap nothing, avoiding any chance of interfering with potentially overlapping objects, and for polymorphic types, we will not swap vtable pointers; thus we do not attempt to change the dynamic types of either parameter.

When swapping the value representations of corresponding members of the same complete type, however, there is no concern about swapping vtable pointers that are not the same, and all other relevant parts of the value will be exchanged. If a union member is present, we will swap any bytes that might contribute to the value representation of any active member of the union.

8.3.5 Untyped nested subobjects

A complete object can store a variety of nested subobjects, the obvious case being all its member subobjects, yet nested subobjects can be created in other ways too. For example, if a class has a nonstatic data member that is an array of `std::byte`, a nested subobject with dynamic storage duration can be created in that storage.

Any new function that swaps value representations must support swapping two complete objects that might have such nested subobjects created within their storage. The value representation of the array of `std::byte` is all those bytes, which includes the full object representation of any nested subobjects. The easiest specification would simply invoke pointer end-zap on all such nested subobjects, starting the lifetime of new nested subobjects after the complete objects' value representations have been exchanged. A stronger specification would not invoke pointer end-zap if both complete objects are storing nested subobjects of the same complete object type.

8.3.6 Uninitialized subobjects

By exchanging just the bytes of the value representation of members, we avoid touching bytes that have indeterminate values, which would be UB. However, if an object comprises any uninitialized member subobjects, then the same issue arises. We note that special dispensation is given to reading and writing raw memory through pointers to `unsigned char` or `std::byte`, and we intend to channel a similar dispensation, without addressing specific implementation details in the formal wording.

8.3.7 New compiler-magic function: `swap_value_representations`

We propose, for the Standard Library, a new function, `swap_value_representations`, that exchanges the value representations of all direct and indirect members of two objects of the same type that are passed by mutable lvalue reference. This is a magic library function that achieves the postcondition without affecting the lifetime of either object. Note that by swapping the value representation rather than the object representation, we intend to support potentially overlapping empty subobjects since empty types have no value representation, i.e., a value representation of zero bytes, regardless of the size of their object representation. Also note that no empty *types*, but only *objects* of zero bytes, are mentioned in the Core language.

This function swaps just the value-representation bytes of members, so it is safe in case of overlapping subobjects that would not be *transparently replaceable* if swapping the whole *object representation*. This function is safe to use with polymorphic types since it swaps only the values of members and not the part of an object's value representation that defines its dynamic type, such as the vtable pointer.

To maintain well-defined behavior in the C++ object model, we mandate that the type of the swapped objects must be both *replaceable* and *trivially relocatable*. To understand why we need both, let us consider the expected behavior of a simplified version of the definition of `std::swap`:

```
#include <utility>    // for `std::move`

template <class T>
void swap(T& a, T& b)
{
    T tmp(std::move(a));    // Move-construct a temporary.
    a = std::move(b);       // Move-assign to `a`.
    b = std::move(tmp);     // Move-assign to `b`.
    return;                // Destroy temporary at end of block.
}
```

On its surface, this combination of move construction, assignment, and destruction is clearly not an implementation in a completely, semantically identical way by **relocation operations** alone since those are equivalent to only a combination of a move constructor and a destructor.

On the other hand, **replacement** is exactly what we need to transform a move-assignment operation into a semantically equivalent series of move-construction and destruction operations. We can begin by rewriting our `swap` implementation above to one that uses a temporary with dynamic storage duration to make the destruction of the temporary an explicit (and not automatic) invocation:

```
#include <new>        // for placement `new`
#include <utility>    // for `std::move`

template <typename T>
void swap(T& a, T& b)
{
    alignas(T) char tmp_buffer[sizeof(T)];
    T* tmp = ::new(tmp_buffer) T(std::move(a));    // Move-construct a temporary.
    a = std::move(b);                               // Move-assign to `a`.
    b = std::move(*tmp);                             // Move-assign to `b`.
    tmp->~T();                                         // Destroy temporary at end of block.
}
```

Now, for a *replaceable type*, we can take advantage of the semantic equivalence of move assignment with destruction followed by move construction, once for the assignment to `a` and once for the assignment to `b`:

```
template <typename T>
void swap(T& a, T& b)
{
    alignas(T) char tmp_buffer[sizeof(T)];
    T* tmp = ::new(tmp_buffer) T(std::move(a));    // Move-construct a temporary.
    a.~T();
    new (&a) T(std::move(b));                       // Replace `a` by `b`.
    b.~T();
    new (&b) T(std::move(*tmp));                     // Replace `b` by `*tmp`.
    tmp->~T();                                         // Destroy temporary at end of block.
}
```

But now, for any *trivially relocatable type*, something magic has happened: We have six operations, which are three pairs of a move construction from a source object into uninitialized storage followed by destruction of that source object. Therefore, because by definition they are semantically equivalent, we can reimplement `swap` with the same semantics using `std::trivially_relocate`:

```

template <typename T>
void swap(T& a, T& b)
{
    alignas(T) char tmp_buffer[sizeof(T)];
    T* tmp = reinterpret_cast<T*>(tmp_buffer);
    std::trivially_relocate(&a, &a+1, tmp);
    std::trivially_relocate(&b, &b+1, &a);
    std::trivially_relocate(tmp, tmp+1, &b);
}

```

Now, this transformation is not going to do the right thing if `a` and `b` are not complete objects subject to being transparently replaced; we’ve done nothing to avoid overwriting padding bytes or vtable pointers in the footprints of `a` and `b`. For this reason, we instead need to use `swap_value_representations`, which performs the same operations (copying all the bytes) on the member data of `a` and `b` without violating the aspects of `a` and `b` that are unknowable in the type system within `swap`:

```

template <typename T>
void swap(T& a, T& b)
{
    if constexpr {
        if constexpr (std::is_trivially_relocatable_v<T> &&
            std::is_replaceable_v<T>) {
            std::swap_value_representations(a,b)
        }
    }
    T tmp(std::move(a)); // Move-construct a temporary.
    a = std::move(b);    // Move-assign to `a`.
    b = std::move(tmp);  // Move-assign to `b`.
}

```

Naturally, such a low-level facility that is close to the compiler should be supported in a freestanding implementation. Our proposed solution has no dependencies on dynamic memory, no exceptions, and no other concerns that a freestanding implementation might raise.

8.3.8 Changes to `std::swap`

Once we have the notions of **replaceability**, **trivial relocatability**, and the `swap_value_representations` function, we are in a position to optimize `std::swap` by using byte-wise operations for appropriate types. First, such types must be **replaceable** to preserve the semantic requirement that move assignment and move construction produce the same end-state. Secondly, we will require that such types are also **trivially relocatable** to ensure that copying bytes produces the expected valid state. Both properties are the necessary requirements to call `std::swap_value_representations`. To provide portable and reliable behavior, we will mandate this optimization rather than leave it as a QoI feature. In doing so, we will want to update every free-function overload of `std::swap` in the Standard Library that should similarly use this optimization; the ADL overload should be no less efficient than the primary template.

Finally, objects do not have a byte representation at compile time, and therefore, attempting to apply byte-wise operations generally does not make sense during constant evaluation. Therefore, we should exclude this optimized behavior during constant evaluation and can do so via a simple `if constexpr`.

We observe that the `swap` overload for arrays is specified as equivalent to calling `swap_ranges`, which in turn calls `std::swap` for each element. Therefore, the array `swap` is indirectly specified to use the byte-wise optimization. Given the **trivially relocatable** requirement for `swap_value_representations` and following the as-if rule, an ambitious implementation might further optimize swapping ranges by using the contiguous memory operations over ranges that come from the `trivially_relocate` specification but only when the type is also **replaceable** to thus assure the full equivalence of the operations.

9 Plans for Library Update

We plan to enable the adoption of these new features in follow-up papers targeting LEWG.

9.1 Immediate updates

In addition to specifying the type traits and library functions that enable the facilities, we should update the library frontmatter to indicate whether and how the Library is allowed to use these features to enhance their QoI.

Clearly, under the as-if rule, the Library immediately gets permission to optimize algorithms and functions for *trivially relocatable* and *replaceable* types where such optimizations are not observable. For example, `std::vector` could optimize many of its operations for such types, given a suitable allocator, such as the default `std::allocator`. No updates to the Library specification are needed for these optimizations, and follow-up papers that suggest changing specifications to allow such optimizations that *would* be observable should be properly directed to LEWG.

The other category of interest is whether Library types themselves can — or should — be *trivially relocatable* or *replaceable*. For example, any implementation of `std::vector` should be able to satisfy the requirements to be both *trivially relocatable* and *replaceable* for any element type as long as its allocator has those properties; we might want to mandate that `std::vector` is relocatable and/or *replaceable* in such cases. Conversely, in the two common implementation strategies for `std::list`, the *sentinel node* is either dynamically allocated or stored directly in the footprint of the `list`. The dynamic node case is always *trivially relocatable* and *replaceable*, but the in-place representation is neither; however, the in-place representation is *nothrow-movable*, whereas the dynamic case must allocate a new node, which can potentially throw. In both cases, **relocation** will never throw, but different trade-offs must be considered when choosing an implementation strategy, and such cases are almost always better left for implementation QoI (especially since ABI concerns might require consideration).

When granting permission for implementations to use keywords that are in addition to those specified by the C++ Standard, we have taken two approaches that we will term the **noexcept** approach and the **constexpr** approach. In the **noexcept** approach, an implementation is granted permission to add **noexcept** specifications to functions as long as those specifications do not invalidate other aspects of the function contract; i.e., an exception specification cannot be added to a virtual function or to a function that is specified to throw exceptions. Conversely, the **constexpr** approach disallows adding **constexpr** to a function that is not declared as **constexpr** in the C++ Standard.

For the purposes of this paper, we believe the minimal necessary specification should use the **noexcept** approach, and we propose the appropriate wording to say so. That choice will allow implementations to experiment with the feature and then provide clear recommendations for specific cases as follow-up LEWG papers.

9.2 ABI concerns

We believe no ABI concerns exist for libraries applying these new features throughout the Standard Library, even as unspecified QoI improvements.

The name-mangling of a type should not depend on whether it is either *trivially relocatable* or *replaceable*. While these properties can be determined through type traits, by definition of being a new feature, no existing code will be SFINAE-enabled on these traits. Updating the internal layout of any Standard Library type to accommodate optimizations using these traits should be unnecessary.

The main concern might be adding constraints to implementation-specific functions used to dispatch to optimized algorithms, such as when growing a vector. In these cases, to avoid introducing new mangled names that would affect link compatibility, `if constexpr` within the dispatching function could be used to enable a fully link-compatible library.

9.3 Specific follow-ups

In a followup paper, we intended to propose adding a new specification element, *Class properties*, for any specification related to class properties 11.2 [\[class.prop\]](#). The Standard Library already makes some effort to specify whether a class must be trivially copyable, standard layout, and so on, and we believe tracking such specification would be more maintainable with a consistent presentation and using a consistent form.

Once we have a *Class properties* element, we can then review all library classes and decide whether to specify the **trivial-relocatability** behavior for that class, which might be conditional on its template arguments if it is a class template. We might also deliberately defer specifying behavior to allow for implementations making different choices, such as node-based containers allocating their end node vs. storing the pointers in the container's object representation.

Finally, once we have an easy way to document class properties, we might consider making stronger guarantees on existing library components where such specification would be useful, e.g., clarifying which types are implicit lifetime.

We would propose moving the specification for the following properties in this new element

- *Trivially Copyable*
- *Standard Layout*
- *Implicit Lifetime*
- *Structural*
- *Aggregate*
- *Empty*
- *Bitmask*

along with the two new properties specified in this paper

- ***Trivially Relocatable***
- ***Replaceable***

The following clauses in the Standard Library specification would then include additional notes regarding this new element and updated specification:

- 16.3.2.4 [\[structure.specifications\]](#) — class properties as well as invariants
- 16.3.3.3.5 [\[customization.point.object\]](#) — may be mildly reformulated with the new specification element
- 16.3.3.5 [\[objects.within.classes\]](#) — may be constraining which members may be added

10 Use Cases

10.1 Optimizing `std::vector` at run time

`std::vector` can optimize moving elements into a new buffer by relying strictly on **trivial relocation** when the allocator does not implement **construct** and **destroy**. A library paper targeting the broader issue of optimizing containers for allocators that use the **construct** and **destroy** customization points will follow since that is a concern for more than just **trivial relocation**.

We find that the current specification allows for **trivial relocation** on **insert** and **erase**, although that use might produce a change of semantics that implementations using assignment prefer to avoid. Hence, we will leave the choice to implementers and their interpretation of the specification.

We expect to provide a library-specific paper to address the semantics of inserting into and erasing from a `std::vector` that is independent of **trivial relocation** concerns.

10.2 Optimizing `std::optional` to be trivially relocatable and replaceable

If `std::optional` is implemented with a variant member (anonymous union) and a boolean flag to indicate if the `optional` is engaged, then memberwise determination of both **trivial relocatability** and **replaceability** will produce the correct property. Typical usage might be something like the following example, which clearly shows that any `optional` implementation is going to provide implementations of all the special member functions and thus require use of both contextual keywords.

Original	Optimized
<pre>template <class T> class optional { union { T d_object; }; bool d_engaged{false}; public: using value_type = T; ... };</pre>	<pre>template <class T> class optional <u>memberwise_trivially_relocatable</u> <u>memberwise_replaceable</u> { union { T d_object; }; bool d_engaged{false}; public: using value_type = T; ... };</pre>

Note that to support the `constexpr` operations required by the Standard, a union-based implementation is the only known way to conform. However, if we were not concerned about `constexpr` evaluations, then we might choose to store our active element in an array of bytes. Unfortunately, adding the `memberwise_trivially_relocatable` or `memberwise_replaceable` properties to the class definition will give our class that same property — *even when the array member is used as storage for a type without those properties* — since an array of `std::byte` is both **trivially relocatable** and **replaceable**.

This problem can be resolved in several ways, but the key is to include a data member that is *conditionally trivially relocatable* or *replaceable*. This resolution is most easily achieved by adding, to the class, an empty data member that ideally can preserve the object layout and ABI.

```
template <bool = true>
struct OptionallyRelocatable {};
```



```

template <>
struct OptionallyRelocatable<false> {
    ~OptionallyRelocatable(){}
};

static_assert( std::is_trivially_relocatable_v<OptionallyRelocatable<>>);
static_assert(!std::is_trivially_relocatable_v<OptionallyRelocatable<false>>);

static_assert( std::is_replaceable_v<OptionallyRelocatable<>>);
static_assert(!std::is_replaceable_v<OptionallyRelocatable<false>>);

```

Original	Optimized
<pre> template <class T> class optional { alignas (T) std::byte d_object[sizeof (T)]; bool d_engaged{false}; public: using value_type = T; ... }; </pre>	<pre> template <class T> class optional <u>memberwise_trivially_relocatable</u> <u>memberwise_replaceable</u> { alignas (T) std::byte d_object[sizeof (T)]; union { bool d_engaged{false}; <u>OptionallyRelocatable<</u> <u>std::is_trivially_relocatable_v< T></u> <u>&&std::is_replaceable_v< T>> _;</u> }; public: using value_type = T; ... }; </pre>

Note that in the above implementation, even though we have made a `union` to contain our empty conditionally relocatable object, the `d_engaged` member will always be active. A similar conditional *replaceable* object would have the same implementation and be simple to add as well.

11 Implementation Experience

An implementation of this proposal is available as a fork of Clang and can also be accessed on [Compiler Explorer](#).

In addition to the handling of the new keywords and class properties, the implementation relies on

- built-in type traits for `is_trivially_relocatable` and `is_replaceable`, which are not different than other type traits of the same nature
- built-ins to return the range of the value representation (excluding the vpointer and trailing padding bytes)

Our Clang implementation of `trivial_relocate` is implemented in terms of `memcpy`. We did not add the necessary machinery to end and start lifetimes since that task is unsupported by the Clang front end and the LLVM optimizer (a known deficiency of LLVM rather than with our implementation). In general, starting and ending lifetimes requires an implementation to add some optimization fences so that optimizers that perform type-based alias analysis are not overly eager and inappropriately prune all code that depends on the new object lifetimes. Either way, adding such fences to an implementation that supports `start_lifetime_as` would present no notable challenges. We have not explored whether sanitizers would need to be made aware of these function semantics.

`swap_value_representation` is implemented as a library function. The function, which is basically isomorphic to some hypothetical `memswap` function, can be implemented in different ways. For our implementation, we chose to perform all the offset computation at compile time such that `swap_value_representation` can benefit from auto-vectorization.

For small objects that are trivial, we found no benefit to using `swap_value_representation` in the implementation of `swap` since optimizers can already produce optimal assembly in these cases.

Note that Clang already supports the notion of *trivially relocatable types* in production, although with no opt-in mechanism. This property is used in the implementation of `std::vector` in libc++ (once again demonstrating an industry need for this feature, as well as deployment experience with very similar ideas).

Clang also offers a `[[clang::trivial_abi]]` type attribute that allows a type to be passed in registers when its destructor/constructor pair can be replaced by a `memcpy`. Types with that attribute can be passed in a register, which affects calling convention, and therefore ABI.

12 FAQ

12.1 Is `void` trivially relocatable?

No, nor is it trivially copyable.

12.2 Are reference types trivially relocatable?

No, nor are they trivially copyable.

Taking the address of a reference to pass it to `relocate` is not possible. How the compiler implements references is entirely unspecified and may not need physical storage if the reference never leaves a local scope. Asking about copying or relocating a naked reference, rather than the entity it refers to, is not meaningful, so these trivial properties are `false`.

12.3 Why can a class with a reference member be trivially relocatable?

A class with a reference member can be *trivially relocatable* for the same reason such a class can be trivially copyable. Strictly speaking, reference members are not nonstatic data members, and we cannot create a pointer-to-data-member to one; they deliberately escape the relevant wording by not appearing in the list of disallowed entities, despite not being trivially copyable or *trivially relocatable* as a distinct type in their own right. This wording is subtle and can entrap the unwary but has been standard practice for many years.

12.4 Are *cv*-qualified types, notably `const` types, trivially relocatable?

Yes, if the unqualified type is *trivially relocatable*.

12.5 Can `const`-qualified types be passed to `trivially_relocate`?

No. While `const`-qualified types are *trivially relocatable* and thus do not inhibit the **trivial relocatability** of a wrapping type, they are typically not safe to `relocate` due to leaving behind a dead object that cannot be replaced using well-defined behavior. Hence, the `trivially_relocate` function is constrained to exclude `const`-qualified types. This exclusion can be skirted using `const_cast` if doing so would not introduce undefined behavior.

12.6 Can types that are not implicit-lifetime types be trivially relocatable?

Yes, and our experience tells us to expect the majority of types, even those that own resources and have nontrivial move constructors and destructors, to still be *trivially relocatable*.

12.7 Why are virtual base classes not trivially relocatable?

Because they are not trivially copyable and because the implementation of virtual base classes on some platforms involves an internal pointer, virtual base classes are not *trivially relocatable*.

We believe that implementing virtual bases such that trivial copyability and relocatability would not be a concern is possible since all the needed data for indirection could be stored as offsets instead of direct pointers. However, whether all implementations could use such a layout or are able to switch to such a layout is unclear. Forcing this support might also require an ABI break.

In our opinion, this low-level behavior should be kept consistent across platforms, rather than left as an unspecified QoI concern, since our current experience has not yet turned up a usage of virtual base classes that would also benefit from this feature.

We would be happy to remove this restriction, but consistency must be maintained with the corresponding restriction on trivially copyable. If no current ABIs are affected, we might consider normatively allowing — or even encouraging — such an implementation (for both trivialities) as conditionally supported behavior on platforms that would not incur an ABI break.

Note that no issues occur with virtual functions since virtual function-table implementations do not take a pointer back into the class, so the vtable pointer can be safely relocated.

12.8 What happens if a relocate operation throws?

Relocation operations must be no-fail, so they do not permit exceptions; if a relocate operation were allowed to fail, whether the failed state had 0, 1, 2, or more valid objects would be unknowable, essentially leaving the program in an undefined state that cannot be cleaned up correctly, which is a significant problem with objects holding resources like a locked mutex.

Our proposal makes clear that `std::trivial_relocate` cannot fail, and the nontrivial implementation of `relocate` mandates that the object type is *nothrow* move constructible. Hence, neither of our operations can fail by throwing an exception.

12.9 Why do deleted special members inhibit implicit trivial relocatability?

Initially, we considered allowing **trivial relocation** of types with these special members functions deleted, based on a notion that we have been familiar with since C++17 when *mandatory copy elision* started propagating noncopyable and nonmovable return values. However, relocation is not the same as copy elision, so objections arose to the idea that, when a user deliberately removes an operation, we should not *silently* re-enable it via a backdoor method. Note that this inhibition changes only the default, preventing accidental relocation of noncopyable or nonmovable types for which relocatability was neither considered nor intended; if **trivial relocatability** is desired, such classes can be made explicitly *trivially relocatable* by means of the `memberwise_trivially_relocatable` keyword.

This design also follows that of the Core language for trivial copyability, which was changed by [CWG1734] to exclude types that deleted all copying operations and which landed in C++17.

12.10 Can the compiler transform argument passing with trivial relocation?

As currently specified, we do not yet enable such support. We believe that this could be accomplished with the appropriate allowances (which already exist for trivially copyable types), but significant work in platform ABIs would be needed to make this happen, similar to what is needed to support Clang's `[[trivial_abi]]` attribute.

To enable bitwise parameter passing, such as through registers, for *trivially relocatable types*, we would need to enable the compiler to freely create extra instances of our objects when passing arguments and return results from functions, which would then enable a compiler to pass the data itself via a register. Importantly and unlike for trivially copyable types (which have trivial destructors), major changes would be needed to ensure that the receiver of the final object is aware that it is now responsible for destruction of that object since currently the creator of parameters is responsible for their destruction on many ABIs.

A separate proposal for argument passing by relocation was offered in [P2839R0] but was not reviewed favorably on its initial presentation to EWG.

12.11 Can the Standard Library containers use this new feature internally?

Yes, where the current specification is permitted to use move construction to **relocate** an object (e.g., when growing or when moving objects within a `vector`), this feature can be used instead for *trivially relocatable types*.

A common misconception implies that `vector` is required to use assignment when inserting into or erasing from a `vector` (other than at the back). This requirement is not, however, explicitly specified in the Standard. The misunderstanding stems from a number of places, which are addressed individually in the subsections below.

12.11.1 Complexity constraints

The first source of this misunderstanding is that people incorrectly consider the requirement to be implied from (24.3.12.5 [vector.modifiers]p5), which states for `vector::erase`:

Complexity: The destructor of `T` is called the number of times equal to the number of the elements erased, but the assignment operator of `T` is called the number of times equal to the number of elements in the `vector` after the erased elements.

This complexity existed in C++98, and the only revision has been a change in C++11 where the text “assignment operator” was updated to “move assignment operator.” Note that `vector::insert` has no such complexity requirement; it is specified only for the `vector::erase` operation.

The misconception also comes from the following sentence in (16.3.2.4 [structure.specifications]p7):

Complexity requirements specified in the library clauses are upper bounds, and implementations that provide better complexity guarantees meet the requirements.

This statement is *not*, therefore, a mandate from the Standard that calls to `vector::erase` shall use the assignment operator as long as the implementation performs as well as or better than the specified complexity. Given that the `trivially_relocate` function as specified in this paper is guaranteed to perform a copy of bytes of the object representation, it must outperform the complexity requirement, and the Standard, therefore, permits implementations to use the `trivially_relocate` function for `vector::erase` operations.

12.11.2 Precondition specifications

The second source of this misunderstanding stems from phrases such as the following in (24.2.4 [sequence.reqmts]p29):

```
a.insert(p, rv)
```

Preconditions: `T` is *Cpp17MoveInsertable* into `X`. For `vector` and `deque`, `T` is also *Cpp17MoveAssignable*.

Effects: Inserts a copy of `rv` before `p`.

Although this specification requires that statements of the form `t = rv` be well-formed, it does *not* impose any limitations on implementations to use assignment when moving objects around internally.

Although the requirement that a type be *Cpp17CopyAssignable* or *Cpp17MoveAssignable* does impose semantic requirements on the assignment operator(s), the requirements are vague and specified in terms of a notion of “value” that is not defined in the Standard; see (16.4.4.2 [utility.arg.requirements]tab:cpp17.moveassignable). This requirement was added in C++11 and has not been revisited since then.

The above explanation refers to `vector::insert(p, rv)`, but the same argument applies to similar preconditions on other member functions. Observe that the postconditions are identical for all sequence containers, including those, such as `list`, that do not require *Cpp17MoveAssignable* as a precondition.

12.11.3 Conclusion

In other words, although most implementations of `vector::erase` and `vector::insert` currently use assignment, which is generally assumed the most efficient approach currently available, implementations are under no obligation whatsoever to do so. The various member functions of `vector` guarantee only that values will be moved around but grant implementations complete freedom as to how that action should be performed, whether by means of (move) assignment, (move) construction, or any other mechanism. Implementations will, therefore, be permitted to perform this move by means of `trivially_relocate` for types that are *trivially relocatable*.

In fact, some implementations avoid using assignment for some operations (for reasons of efficiency); see the linked examples for [GCC](#) and [LLVM](#).

Note that all the comments above apply equally to `deque` as well as to `vector`.

Note also that this lack of a clear requirement exposes an existing ambiguity for `vector::insert` and `vector::erase` operations where, for the contained type, move-assign plus destroy is not equivalent to destroy plus move-construct. That ambiguity is an issue that exists at the moment, and while we might address it with a future, orthogonal proposal, a solution is not required for **trivial relocation** as specified by this paper.

Similarly, we might choose to clarify the complexity and requirements clauses above at some point in the future, but that clarification is not required by this proposal and has been left for another time.

12.12 Do implementations need to mark classes `memberwise_trivially_relocatable` to benefit?

No, although some classes will need to be annotated to qualify as *trivially relocatable*. For example, the most common implementations of `std::array`, `std::pair`, and `std::tuple` will be implicitly *trivially relocatable* if all their members are *trivially relocatable*. `std::vector` can safely be marked as *trivially relocatable* if its allocator and pointer types are *trivially relocatable*. `std::list` might be marked as *trivially relocatable* if it allocates its tail node but not if the tail node is embedded in the object representation itself.

Once we establish a policy of how much we want to guarantee and how much we want to leave open to implementer choice, a follow-up paper will address desired guarantees for **trivial relocatability** in the Standard Library.

12.13 Which Standard Library types are safe for the `swap_value_representations` function?

The Standard Library specification says nothing — or as little as possible — about the implementation of the types that it specifies, so no guarantees are made in this paper regarding which types in the Standard Library users can safely pass to `swap_value_representations`.

We expect that something must be said, even if only to explicitly make all such concerns a choice deferred to implementers' QoI. We would expect a more significant follow-up paper to review the whole library and to make specific guarantees on a subset of types that do — or do not — offer such guarantees.

12.13.1 For what types can I swap values but not representations?

For types in which move assignment produces a different result than move construction, we can swap values but not representations. These types include, among other things, any type that has reference semantics, not value semantics.

12.13.2 For what types can I swap neither representations nor values?

For types with data members that cannot be replaced, including const-qualified nonstatic data members or nonstatic data members that are references, we can swap neither representations nor values. Note that a swap for values that does not change these members — such as `swap` for `tuple<T&>` that swaps the referenced elements — may still be defined, but such semantics strictly go beyond swapping values.

12.14 Can I mark as trivially relocatable a type that is not replaceable?

Yes! For example, this would be appropriate for types having data members that are references or using `std::pmr::polymorphic_allocator` or any other type that does not propagate on `swap`.

12.15 Can I mark a type as replaceable but not trivially relocatable?

Yes! This proposal does not offer any immediate advantages for doing so, but we expect to build on **replacement** to optimize other features, such as assignment, in the future.

12.16 What happened to the predicates for the contextual keywords?

An earlier version of this proposal included the option to add a predicate following each of the new contextual keywords to activate or inhibit their behavior. This feature was dropped for introducing too much complexity, including a new vexing parse to resolve, and having vague semantics when the predicate is **false** but the implicit specification would have been **true**. Given the rarity of such cases and the relative simplicity of the library-based workaround above, we chose to keep the core proposal as simple as possible, following EWG guidance.

12.17 Why is copy-replacement unsupported?

In practice, only **replaceability** of objects of type `T` from xvalues of type `T` seems relevant to the operations we are likely to optimize. We have, therefore, simplified the design to focus solely on such **replacement** (which could be termed move-replacement were we being pedantic) and not overcomplicate the language or users' lives by adding even more properties to consider.

12.18 Why is marking a class that can *never* be trivially relocatable not ill-formed?

A class with a virtual base class can never be *trivially relocatable*, so why is adding the `memberwise_trivially_relocatable` identifier to that class not ill-formed?

This case is still well-formed, but the class will indeed never be *trivially relocatable*, and the type trait will deterministically always return `false`. However, this type may also be used as a base class or data member when instantiating a class template, and we do not want to add complexity by considering such special-case instantiations as well-formed when the original case need not be marked as ill-formed.

However, the deterministic case of a direct virtual base class would make for a useful compiler warning. The more general case of a data member or nonvirtual base class not being relocatable (or *replaceable*) is deliberately not an error since we want to support different implementations of the same type that have different properties; e.g., different implementations of `std::list` choose different trade-offs on how to store the sentinel node marking the end of the list, yet some of those choices are *trivially relocatable* and some are not. We want to avoid the inconsistency of deterministically flagging an error when compiling a class with a `std::list` data member in some Standard Library implementations and not in others.

12.19 Why is there no `is_trivially_replaceable` trait?

A common use case is to require types that satisfy both `is_trivially_relocatable<T>` and `is_replaceable<T>`. We could consider whether this use case occurs frequently enough that adding another trait that is the logical conjunction of the two would be valuable.

We opted to omit this trait to our proposal since such a trait is not primitive to the Core-language design of this paper and could easily be added as an amendment in an LEWG follow-up paper well within the timeframe of C++26 if desired.

The lack of a core type category named *trivially replaceable* is another reason to defer to a follow-up paper, and we would be consuming that potential for future vocabulary for a pure library extension. Making that choice before advancing this paper is unnecessary.

Similarly, we deferred the idea of adding an `is_relocatable` trait that is the logical disjunction of `is_trivially_relocatable` and `is_nothrow_move_constructible`, which is the precondition to call `std::relocate`. We deliberately want to leave the vocabulary of nontrivial relocation available for any potential future work in EWG beyond C++26.

12.20 Is it UB to mark a nonconforming type as trivially relocatable?

First, the compiler has no way to validate that our class's constructors and destructor do not maintain an invariant that is not relocatable, so the compiler will trust us and enable the type trait. This in itself is not UB, but UB will likely follow when some library code makes a transformation that causes our invariant, such as an internal pointer, to no longer hold. Such UB will occur in the subsequent library call, not in the class definition.

12.21 Is it UB to mark a nonconforming type as replaceable?

Just as erroneously marking a type as *trivially relocatable* can lead to undefined behavior in library calls, so can marking a type as *replaceable*. However, where **replaceability** is used as a constraint without **trivial relocation**, there remain reasonable implementations that do not incur UB. For example, if operations are logged, then the act of writing to a log is typically an observable side effect. Library code that transforms between assignment and destroy-then-construct will have an observable change of behavior, such as the suggested

logging, but such changes do not in themselves constitute undefined behavior. The creator of the affected class must decide whether a change of such logging behavior would be problematic and then choose whether to mark their type as **replaceable**.

13 Illustrative Examples

13.1 Unconstrained vector

Let us consider the case of a user-written container, similar to `std::vector`. Since `std::relocate` is a nofail function that exploits **trivial relocation** where it is available, we have to consider only two kinds of elements:

- Types that are either *trivially relocatable* or no-throw move constructible
- All other types

An alternative summary of these two kinds are

- types that can be safely relocated
- types that cannot be safely relocated

The first case to optimize is relocating elements when the current capacity is exceeded by an insert operation. In this case, we clearly can simply relocate for those element types that are safely relocatable and must manually move-construct the second category, accounting for a possible thrown exception on move.

The next operation to consider is erasing an element. In this scenario, we will destroy the requested element(s) and then, for types that can be safely relocated, **relocate** the tail of the vector to the lower address since **relocate** is nofail and supports overlapping ranges. For the second category of types, we must perform the manual relocation and clear the remains of the tail if an exception is thrown.

The final operation to consider is an insertion in the middle of this vector. Here, the first thing we do, assuming capacity is not exceeded, is relocate all elements from the insertion point up by a distance to allow all the new elements to be inserted. Then we construct all the new elements, which is a potentially throwing operation. If an exception *is* thrown, we have several options for our custom vector. For the strong exception safety guarantee, we can destroy the newly inserted items and then safely relocate the original elements back in place since **relocate** is a nofail operation. Alternatively, we provide the basic guarantee by destroying the old tail — and potentially the newly inserted items — before adjusting the vector's size, or maybe we could even clear the whole vector.

Note that all these operations use only **trivial relocation** and never call for **replaceability**.

13.2 Standard vector

When we add the constraints that the Standard imposes on `std::vector`, we find that **replaceability** becomes a useful property. For both insertion and erasure, the Standard likes to assume that elements are *replaceable*, i.e., assignment is interchangeable with destroy-then-move-construct. Within that guarantee, the Standard Library vector can use relocation per our custom vector example, but for types that are relocatable but not *replaceable*, matters become more complicated. That topic will be the subject of a separate paper specific to vector, which is necessary regardless of whether we support relocation in C++26. Having the ability to detect *replaceable* types would be extremely helpful for that follow-up paper.

13.3 Conforming implementation of a trivially relocatable `std::optional`

The following implementation of `optional` satisfies the C++ Standard specification for the members that it implements and provides a minimal test driver. This implementation uses the new feature macro to ensure that the code compiles with both C++23 and C++26 and is *trivially relocatable* if and only if its element type is *trivially relocatable*.

To implement the `constexpr` members, the implementation is required to use a union to represent its internal state when engaged²:

```
#include <cassert>
#include <iostream>
#include <memory>
#include <new>
```

²This implementation can be seen compiling on Compiler Explorer here: [compiler-explorer](https://godbolt.org/z/9KdKz9).

```

#include <type_traits>
#include <utility>

template <class T>
class optional
    memberwise_trivially_relocatable
    memberwise_replaceable
{
    union {
        T d_object;
    };
    bool      d_engaged{false};

    constexpr T const * address() const noexcept
    { return ::std::addressof(d_object); };

    constexpr T      * address()      noexcept
    { return ::std::addressof(d_object); };

    template<class... Args>
    constexpr void do_emplace(Args&&... args) {
        ::new(address()) T(std::forward<Args>(args)...);
        d_engaged = true;
    }

public:
    using value_type = T;

    constexpr optional() noexcept {}

    constexpr optional(optional const & other)
    : d_engaged{other.d_engaged} {
        if (d_engaged) {
            ::new(address()) T( other.value() );
        }
    }

    constexpr optional(optional&& other)
        noexcept(std::is_nothrow_move_constructible_v<T>)
    : d_engaged{other.d_engaged}
    {
        if (d_engaged) {
            ::new(address()) T( std::move(other).value() );
        }
    }

    template<class U = T>
        requires (std::is_constructible_v<T, U>
            && !std::is_same_v<std::remove_cvref_t<U>, optional>)
    constexpr
    explicit(!std::is_convertible_v<U, T>)
    optional(U&& arg) {
        do_emplace( std::forward<U>(arg) );
    }
}

```

```

constexpr ~optional() {
    static_assert(std::is_replaceable_v< optional>== std::is_replaceable_v< T>);
    static_assert(
        std::is_trivially_relocatable_v< optional>== std::is_trivially_relocatable_v< T>);

    if (d_engaged) {
        d_object.~T();
    }
}

constexpr optional& operator=(optional const & rhs);

constexpr optional& operator=(optional && rhs)
    noexcept(std::is_nothrow_move_assignable_v<T>
        && std::is_nothrow_move_constructible_v<T>) {
    std::cout << "Assignment\n";
    if (!d_engaged) {
        if (rhs.d_engaged) {
            do_emplace( std::move(rhs.value()) );
        }
    }
    else if (!rhs.d_engaged) {
        d_object.~T();
        d_engaged = false;
    }
    else {
        value() = rhs.value();
    }
    return *this;
}

template<class U = T>
constexpr optional& operator=(U && arg) {
    std::cout << "Assignment\n";
    if (!d_engaged) {
        do_emplace( std::forward<U>(arg) );
    }
    else {
        d_object = std::forward<U>(arg);
    }
    return *this;
}

constexpr T const * operator->() const noexcept
{ assert(d_engaged); return address(); }

constexpr T      * operator->()      noexcept
{ assert(d_engaged); return address(); }

constexpr T const & operator*() const & noexcept
{ assert(d_engaged); return d_object; }
constexpr T      & operator*()      & noexcept
{ assert(d_engaged); return d_object; }

```

```

constexpr T      && operator*()      && noexcept
{ assert(d_engaged); return std::move(d_object); }
constexpr T const&& operator*() const&& noexcept
{ assert(d_engaged); return std::move(d_object); }

constexpr explicit operator bool() const noexcept
{ return d_engaged; }
constexpr bool      has_value() const noexcept
{ return d_engaged; }

constexpr T const & value() const &
{ assert(d_engaged); return d_object; }
constexpr T      & value()      &
{ assert(d_engaged); return d_object; }
constexpr T      && value()      &&
{ assert(d_engaged); return std::move(d_object); }
constexpr T const&& value() const&&
{ assert(d_engaged); return std::move(d_object); }
};

constexpr int number(int n) {
    optional<int> x{n};
    return x.value();
}

int a[number(5uz)];

int main() {
    optional<int> x;
    assert(!x);

    std::cout << "Assignments to x\n";
    x = 3;
    auto y = x;

    x = 4;
    std::cout << "swap x\n";
    std::swap(x, y);

    assert(3 == *x);
    assert(4 == *y);

    optional<std::shared_ptr<int>> p1;

    std::cout << "Assignments to p\n";
    p1 = std::make_shared<int>(3);
    auto p2 = p1;

    p2 = std::make_shared<int>(4);
    std::cout << "swap p\n";
    std::swap(p1, p2);
}

```

13.4 C++26 implementation using internal array

Using an internal array negates the ability to support `constexpr`, but this implementation strategy is used frequently for similar types in other libraries. Managing both *trivially relocatable* and *replaceable* properties with an empty member must be done with care since mistakenly disabling both properties is easy to do when intending to disable only one or the other.³

```
#include <cassert>
#include <cstddef>
#include <iostream>
#include <memory>
#include <new>
#include <type_traits>
#include <utility>

template <bool triviallyRelocatable,
          bool replaceable>
struct ConditionalProperties {};

template <>
struct ConditionalProperties<false,true> memberwise_replaceable {
    ~ConditionalProperties(){}
};
template <>
struct ConditionalProperties<true,false> memberwise_trivially_relocatable {
    ~ConditionalProperties(){}
};
template <>
struct ConditionalProperties<false,false> {
    ~ConditionalProperties(){}
};

static_assert( std::is_trivially_relocatable_v<ConditionalProperties<true,true>> );
static_assert( std::is_trivially_relocatable_v<ConditionalProperties<true,false>> );
static_assert( !std::is_trivially_relocatable_v<ConditionalProperties<false,true>> );
static_assert( !std::is_trivially_relocatable_v<ConditionalProperties<false,false>> );

static_assert( std::is_replaceable_v<ConditionalProperties<true,true>> );
static_assert( !std::is_replaceable_v<ConditionalProperties<true,false>> );
static_assert( std::is_replaceable_v<ConditionalProperties<false,true>> );
static_assert( !std::is_replaceable_v<ConditionalProperties<false,false>> );

template <class T>
class optional memberwise_trivially_relocatable memberwise_replaceable {
    alignas (T)
    std::byte d_object[sizeof (T)];
    union {
        bool d_engaged{false};
        ConditionalProperties<std::is_trivially_relocatable_v<T>,
                             std::is_replaceable_v<T>> enforce_properties;
    };

    constexpr T const * address() const noexcept
    { return reinterpret_cast<T const *>(d_object); };
};
```

³This implementation can be seen compiling on Compiler Explorer here: [compiler-explorer](https://godbolt.org/z/9vKvz9).

```

constexpr T      * address()      noexcept
{ return reinterpret_cast<T      *>(d_object); };

public:
    using value_type = T;

    // 22.5.3.2, constructors
    constexpr optional() noexcept = default;

    constexpr optional(optional const & other) : d_engaged{other.d_engaged} {
        if (d_engaged) {
            ::new(address()) T( other.value() );
        }
    }

    constexpr optional(optional&& other) noexcept(std::is_nothrow_move_constructible_v<T>)
    : d_engaged{other.d_engaged}
    {
        if (d_engaged) {
            ::new(address()) T( std::move(other).value() );
        }
    }

    template<class U = T>
        requires (std::is_constructible_v<T, U>
            && !std::is_same_v<std::remove_cvref_t<U>, optional>)
    constexpr
    explicit(!std::is_convertible_v<U, T>)
    optional(U&& arg) {
        ::new(address()) T( std::forward<U>(arg) );
        d_engaged = true;
    }

    // 22.5.3.3, destructor
    constexpr ~optional() {
        static_assert(std::is_trivially_relocatable_v<optional> ==
            std::is_trivially_relocatable_v<T>);
        static_assert(std::is_replaceable_v<optional> ==
            std::is_replaceable_v<T>);

        if (d_engaged) {
            address()->~T();
        }
    }

    // 22.5.3.4, assignment
    constexpr optional& operator=(optional const & rhs);

    constexpr optional& operator=(optional && rhs)
        noexcept(std::is_nothrow_move_assignable_v<T>
            && std::is_nothrow_move_constructible_v<T>)
    {
        std::cout << "Assignment\n";
    }

```

```

    if (!d_engaged) {
        if (rhs.d_engaged) {
            ::new(address()) T( std::move(rhs.value()) );
            rhs.d_engaged = false;
            d_engaged = true;
        }
    }
    else if (!rhs.d_engaged) {
        address()->~T();
        d_engaged = false;
    }
    else {
        value() = rhs.value();
    }
    return *this;
}

template<class U = T>
constexpr optional& operator=(U && arg) {
    std::cout << "Assignment\n";
    if (!d_engaged) {
        ::new(address()) T( std::forward<U>(arg) );
        d_engaged = true;
    }
    else {
        *address() = std::forward<U>(arg);
    }
    return *this;
}

// 22.5.3.7, observers
constexpr T const * operator->() const noexcept
{ assert(d_engaged); return address(); }
constexpr T      * operator->()      noexcept
{ assert(d_engaged); return address(); }

constexpr T const & operator*() const & noexcept
{ assert(d_engaged); return *address(); }
constexpr T      & operator*()      & noexcept
{ assert(d_engaged); return *address(); }
constexpr T      && operator*()      && noexcept
{ assert(d_engaged); return std::move(*address()); }
constexpr T const&& operator*() const&& noexcept
{ assert(d_engaged); return std::move(*address()); }

constexpr explicit operator bool() const noexcept
{ return d_engaged; }
constexpr bool      has_value() const noexcept
{ return d_engaged; }

constexpr T const & value() const &
{ assert(d_engaged); return *address(); }

```

```

constexpr T      & value()      &
{ assert(d_engaged); return *address(); }
constexpr T      && value()      &&
{ assert(d_engaged); return std::move(*address()); }
constexpr T const&& value() const&&
{ assert(d_engaged); return std::move(*address()); }
};

int main() {
    optional<int> x;
    assert(!x);

    std::cout << "Assignments to x\n";
    x = 3;
    auto y = x;

    x = 4;
    std::cout << "swap x\n";
    std::swap(x, y);

    assert(3 == *x);
    assert(4 == *y);

    optional<std::shared_ptr<int>> p1;

    std::cout << "Assignments to p\n";
    p1 = std::make_shared<int>(3);
    auto p2 = p1;

    p2 = std::make_shared<int>(4);
    std::cout << "swap p\n";
    std::swap(p1, p2);
}

```


14 Proposed Wording

Make the following changes to the C++ Working Draft. All wording is relative to [N4988], the latest draft at the time of writing.

14.1 Add new identifiers with a special meaning

5.10 [lex.name] Identifiers

Table 4: Identifiers with special meaning [tab:lex.name.special]

final	import	<u>memberwise_replaceable</u>	<u>memberwise_trivially_relocatable</u>	module	override
-------	--------	-------------------------------	---	--------	----------

14.2 Specify trivially relocatable types

Editorial note: *We have separated each sentence to improve clarity rather than trying to identify the definition of so many terms as a single paragraph.*

6.8.1 [basic.types.general] General

- ⁹ Arithmetic types (6.8.2 [basic.fundamental]), enumeration types, pointer types, pointer-to-member types (6.8.4 [basic.compound]), `std::nullptr_t`, and cv-qualified (6.8.5 [basic.type.qualifier]) versions of these types are collectively called *scalar types*.

Scalar types, trivially copyable class types (11.2 [class.prop]), arrays of such types, and cv-qualified versions of these types are collectively called *trivially copyable types*.

Scalar types, trivial class types (11.2 [class.prop]), arrays of such types, and cv-qualified versions of these types are collectively called *trivial types*.

Scalar types, trivially relocatable class types (11.2 [class.prop]), arrays of such types, and cv-qualified versions of these types are collectively called *trivially relocatable types*.

Scalar types, replaceable class types (11.2 [class.prop]), and arrays of such types are collectively called *replaceable types*.

Scalar types, standard-layout class types (11.2 [class.prop]), arrays of such types, and cv-qualified versions of these types are collectively called *standard-layout types*.

Scalar types, implicit-lifetime class types (11.2 [class.prop]), array types, and cv-qualified versions of these types are collectively called *implicit-lifetime types*.

14.3 Address trivial relocation of lambdas

7.5.6.2 [expr.prim.lambda.closure] Closure types

- ² The closure type is declared in the smallest block scope, class scope, or namespace scope that contains the corresponding *lambda-expression*.

[Note 1: This determines the set of namespaces and classes associated with the closure type (6.5.4 [basic.lookup.argdep]). The parameter types of a *lambda-declarator* do not affect these associated namespaces and classes. —end note]

- ³ The closure type is not an aggregate type (9.4.2 [dcl.init.aggr]); it is a structural type (13.2 [temp.param]) if and only if the lambda has no *lambda-capture*. An implementation may define the closure type differently from what is described below provided this does not alter the observable behavior of the program other than by changing:

(3.1) — the size and/or alignment of the closure type,

- (3.2) — whether the closure type is trivially copyable (11.2 [class.prop]), or
- (3.x) — whether the closure type is trivially relocatable (11.2 [class.prop]), or
- (3.y) — whether the closure type is replaceable (11.2 [class.prop]), or
- (3.3) — whether the closure type is a standard-layout class (11.2 [class.prop]).

An implementation shall not add members of rvalue reference type to the closure type.

14.4 Update grammar to support memberwise contextual keywords

11.1 [class.pre] Preamble

- ¹ A class is a type. Its name becomes a class-name (11.3 [class.name]) within its scope.

class-name :
identifier
simple-template-id

A *class-specifier* or an *elaborated-type-specifier* (9.2.9.5 [dcl.type.elab]) is used to make a *class-name*. An object of a class consists of a (possibly empty) sequence of members and base class objects.

class-specifier :
class-head { *member-specification*_{opt} }

class-head:
class-key *attribute-specifier-seq*_{opt} *class-head-name* ~~*class-virt-specifier*~~_{opt} *class-property-specifier-seq*_{opt}
*base-clause*_{opt}
class-key *attribute-specifier-seq*_{opt} *base-clause*_{opt}

class-head-name :
*nested-name-specifier*_{opt} *class-name*

class-property-specifier-seq:
class-property-specifier *class-property-specifier-seq*_{opt}

class-property-specifier:
final
memberwise_replaceable
memberwise_trivially_relocatable

~~*class-virt-specifier*~~:
~~**final**~~

class-key :
class
struct
union

A class declaration where the *class-name* in the *class-head-name* is a *simple-template-id* shall be ...

- ⁴ [Note 2: The *class-key* determines whether the class is a union (11.5 [class.union]) and whether access is public or private by default (11.8 [class.access]). A union holds the value of at most one data member at a time. —end note]
- ⁵ If a class is marked with the *class-virt-specifier* **final** and it appears as a *class-or-decltype* in a base-clause (11.7 [class.derived]), the program is ill-formed. Whenever a *class-key* is followed by a *class-head-name*, the identifier **final**, and a colon or left brace, **final** is interpreted as a *class-virt-specifier*.

- ⁵ The same *class-property-specifier* shall not appear multiple times within a single *class-property-specifier-seq*.

Whenever a *class-key* is followed by a *class-head-name*, one of the identifiers `final`, `memberwise_replaceable`, or `memberwise_trivially_relocatable`, and a colon or left brace, the identifier is interpreted as a *class-property-specifier*.

[Example 2:

```
struct A;
struct A final {}; // OK, definition of struct A,
                  // not value-initialization of variable final

struct X {
    struct C { constexpr operator int() { return 5; } };
    struct B finalmemberwise_trivially_relocatable : C{}; // OK, definition of nested class B,
                                                         // not declaration of a bit-field
                                                         // member finalmemberwise_trivially_relocatable
};
```

—end example]

- ^a If a class is marked with the ~~class-virt-specifier~~*class-property-specifier* `final` and ~~it~~that class appears as a *class-or-decltype* in a base-clause (11.7 [class.derived]), the program is ill-formed.
- ⁶ [Note 3: Complete objects of class type have nonzero size. Base class subobjects and members declared with the `no_unique_address` attribute (9.12.12 [dcl.attr.nouniqueaddr]) are not so constrained. —end note]

14.5 Specification for trivially relocatable classes

Design note:

Declaring a class as trivially relocatable is possible, by means of the `memberwise_trivially_relocatable` specifier, even if that class has user-provided special members. Note that such a declaration is not permitted to break the encapsulation of members or bases and allow for their trivial relocation when they, themselves, are not trivially relocatable.

11.2 [class.prop] Properties of classes

- ² A *trivial class* is a class that is trivially copyable and has one or more eligible default constructors (11.4.5.2 [class.default.ctor]), all of which are trivial.

[Note 1: In particular, a trivially copyable or trivial class does not have virtual functions or virtual base classes. —end note]

- ^a A class is *eligible for trivial relocation* unless it has
- any virtual base classes, or
 - a base class that is not a trivially relocatable class, or
 - a non-static data member of a non-reference type that is not of a trivially relocatable type.
- ^b A class *C* is *eligible for replacement* unless it has
- a base class that is not a replaceable class, or
 - a non-static data member that is not of a replaceable type,
 - no constructor that would be selected when an object of type *C* is direct-initialized from an xvalue of type *C*,
 - no assignment operator that would be selected when an object of type *C* is assigned from an xvalue of type *C*,
 - no destructor.

- ^c A class **C** is a *trivially relocatable class* if it is eligible for trivial relocation and
 - has a *class-trivially-relocatable-specifier*, or
 - is a union with no user-declared special member functions, or
 - satisfies all of the following:
 - when an object of type **C** is direct-initialized from an xvalue of type **C**, overload resolution would select a constructor that is neither user-provided nor deleted, and
 - when an xvalue of type **C** is assigned to an object of type **C**, overload resolution would select an assignment operator that is neither user-provided nor deleted, and
 - it has a destructor that is neither user-provided nor deleted.
 - ^d [Note 2: Accessibility of the special member functions is not considered when establishing trivial relocatability. —end note]
 - ^e [Note 3: A type with non-static data members that are const-qualified or are references can be trivially relocatable. —end note]
 - ^f [Note 4: Trivially copyable classes are trivially relocatable unless they have deleted special members. —end note]
 - ^g A class **C** is a *replaceable class* if it is eligible for replacement and
 - has a *class-replaceable-specifier*, or
 - is a union with no user-declared special member functions, or
 - satisfies all of the following:
 - when an object of type **C** is direct-initialized from an xvalue of type **C**, overload resolution would select a constructor that is neither user-provided nor deleted, and
 - when an xvalue of type **C** is assigned to an object of type **C**, overload resolution would select an assignment operator that is neither user-provided nor deleted, and
 - it has a destructor that is neither user-provided nor deleted.
 - ^h [Note 5: Accessibility of the special member functions is not relevant. —end note]
 - ⁱ [Note 6: Trivially copyable classes are replaceable unless they have deleted special members. —end note]
 - ³ A class **S** is a *standard-layout class* if it:
- (3.1) ...

14.6 Add feature macros

Add a `__cpp_trivial_relocatability` feature-test macro to the table in 15.11 [cpp.predefined], set to the date of adoption.

...

14.7 Library wording

Design note: The first paragraph explicitly captures the status quo that these class properties — the whole set specified in 11.2 [class.prop] — are deliberately left as a quality of implementation feature.

The second paragraph addresses permission to add the new annotation wherever an implementation might find it useful, without being constrained by its absence from the library specification, much like we grant permission to add `noexcept` specifications to functions of the implementation's choosing. The specification really needs only the second paragraph, but adding a section with the first paragraph gives us somewhere to hang the wording.

16.4.6.X Properties of library classes [library.class.props]

- ¹ Unless clearly stated, it is unspecified whether any class described in Clause 17 through Clause 34 and Annex D is a trivial class, a trivially copyable class, a trivially relocatable class, a standard-layout class, or an implicit-lifetime class (11.2 [class.prop]).
- ² An implementation may add the *class-property-specifier* `memberwise_trivially_relocatable` to any class whose implementation is eligible for trivial relocation.
- ³ An implementation may add the *class-property-specifier* `memberwise_replaceable` to any class whose implementation is eligible for replacement.

14.8 Add new type traits

21.3.3 [meta.type.synop] Header <type_traits> synopsis

```
template< class T >
struct is_replaceable;

template< class T >
struct is_trivially_relocatable;

template< class T >
inline constexpr bool is_replaceable_v = is_replaceable<T>::value;

template< class T >
inline constexpr bool is_trivially_relocatable_v = is_trivially_relocatable<T>::value;
```

21.3.5.4 [meta.unary.prop] Type properties

Template	Condition	Preconditions
template<class T> struct is_replaceable;	T is a replaceable type (6.8.1 [basic.types.general])	<code>remove_all_extents_t<T></code> shall be a complete type or <i>cv</i> void
template<class T> struct is_trivially_relocatable;	T is a trivially relocatable type (6.8.1 [basic.types.general])	<code>remove_all_extents_t<T></code> shall be a complete type or <i>cv</i> void

14.9 Specify the compiler-magic functions

Add to the <memory> header synopsis in 20.2.2 [memory.syn]p3.

20.2.2 [memory.syn] Header <memory> synopsis

```
// 20.2.6, explicit lifetime management
template<class T>
    T* start_lifetime_as(void* p) noexcept;                                // freestanding
template<class T>
    const T* start_lifetime_as(const void* p) noexcept;                    // freestanding
template<class T>
    volatile T* start_lifetime_as(volatile void* p) noexcept;              // freestanding
template<class T>
    const volatile T* start_lifetime_as(const volatile void* p) noexcept;  // freestanding
template<class T>
    T* start_lifetime_as_array(void* p, size_t n) noexcept;                // freestanding
```

```

template<class T>
    const T* start_lifetime_as_array(const void* p, size_t n) noexcept;           // freestanding
template<class T>
    volatile T* start_lifetime_as_array(volatile void* p, size_t n) noexcept;   // freestanding
template<class T>
    const volatile T* start_lifetime_as_array(const volatile void* p,
                                              size_t n) noexcept;                 // freestanding

template <class T>
    void swap_value_representations(T& a, T& b);                               // freestanding

template <class T>
    T* trivially_relocate(T* begin, T* end, T* new_location);                   // freestanding

template <class T>
    constexpr T* relocate(T* begin, T* end, T* new_location);                   // freestanding

```

20.2.6 [obj.lifetime] Explicit lifetime management

```

template <class T>
    void swap_value_representations(T& a, T& b);

```

- ¹ *Mandates:* T is a complete object type, and `is_trivially_relocatable_v<T> && is_replaceable_v<T>` is true.
- ³ *Postconditions:* a has the value representation that b had prior to this function call; b has the value representation that a had prior to this function call; both objects and all their base class subobjects retain their original dynamic types.
 - For every non-base object c within a with a corresponding object d within b, c has the value representation that d had prior to this function call and d has the value representation that c had prior to this function call. If c and d are unions, the active member of each is the member that was the active member of the other one prior to this function call.
 - For every trivially relocatable object x within a or b, create a new object y at the same relative location within the other enclosing object having the value x had prior to this function call and end the lifetime of x. If y is a union, its active member is the same member that x had prior to this function call.
 - End the lifetime of all other non-base objects within a and b.
- ⁴ *Throws:* Nothing.
- ^h *Remarks:* No constructors, destructors, or assignment operators are invoked.

```

template <class T>
    T* trivially_relocate(T* begin, T* end, T* new_location);

```

- ^a *Mandates:* T is a complete type, and `is_trivially_relocatable_v<T> && !is_const_v<T>` is true.
- ^c *Preconditions:*
 - (c.1) — `[begin, end)` is a valid range.
 - (c.2) — `[new_location, new_location + (end - begin))` denotes a region of storage that is a subset of the region of storage reachable through (6.8.4 [basic.compound]) `new_location` and suitably aligned for the type T.
- ^d *Postconditions:*

No effect if `new_location == begin`.

Otherwise, the range denoted by `[new_location, new_location + (end - begin))` contains objects (including subobjects) whose lifetime has begun and whose object representations are the original object representations of

the corresponding objects in the source range `[begin, end)`. If any of the aforementioned objects is a union, its active member is the same as that of the corresponding union in the source range. If any of the aforementioned objects has a non-static data member of reference type, that reference refers to the same entity as does the corresponding reference in the source range. The lifetime of the original objects in the source range has ended.

^e *Returns:* `new_location + (end - begin)`.

^f *Throws:* Nothing.

^g *Complexity:* Linear in the length of the source range.

^h *Remarks:* No constructors or destructors are invoked.

```
template <class T>
constexpr T* relocate(T* begin, T* end, T* new_location);
```

^w *Mandates:* `is_trivially_relocatable_v<T> || is_nothrow_move_constructible_v` is true

^x *Effects:* If not called during constant evaluation and `T` is trivially relocatable, then has effects equivalent to `trivially_relocate(begin, end, new_location)`; otherwise, for each element in `[begin, end)`, move constructs that object to the corresponding location in `[new_location, new_location + (end - begin))` and then runs that element's destructor.

^y *Remarks:* Overlapping ranges are supported.

^e *Returns:* `new_location + (end-begin)`.

^z *Throws:* Nothing.

14.10 Require the optimization for `std::swap`

22.2.2 [utility.swap] swap

```
template<class T>
constexpr void swap(T& a, T& b) noexcept(see below);
```

¹ *Constraints:* `is_move_constructible_v<T>` is true, and `is_move_assignable_v<T>` is true.

² *Preconditions:* Type `T` meets the *Cpp17MoveConstructible* (Table 31) and *Cpp17MoveAssignable* (Table 33) requirements.

³ *Effects:* Exchanges values stored in two locations.

If not called during constant evaluation and `is_trivially_relocatable_v<T> && is_replaceable_v<T>` is true, then this function has effects equivalent to `swap_value_representations(a, b)`.

⁴ *Remarks:* The exception specification is equivalent to:

```
is_nothrow_move_constructible_v<T> && is_nothrow_move_assignable_v<T>
```

14.10.1 Feature-test macro

Add a new `__cpp_lib_trivially_relocatable` feature-test macro in `[version.syn]`:

```
#define __cpp_lib_trivially_relocatable 20XXXXL // also in <memory>, <type_traits>
```

15 Acknowledgements

This document is written in Markdown and depends on the extensions in [pandoc](#) and [mpark/wg21](#), and we would like to thank the authors of those extensions and associated libraries.

The authors would also like to thank Brian Bi for his assistance in proofreading this paper, especially the proposed Core wording. Additional thanks to Jens Maurer who helped to greatly refine the wording in advance of its first Core review.

Also, this paper has been greatly improved by feedback from Arthur O'Dwyer, author of [\[P1144\]](#), who corrected many bad assumptions we made about his paper and helped bring the technical differences into focus. We also benefited from several examples he shared to help illustrate those differences and misunderstandings.

16 References

- [CWG1734] James Widman. 2013-08-09. Nontrivial deleted copy functions.
<https://wg21.link/cwg1734>
- [N4988] Thomas Köppe. 2024-08-05. Working Draft, Programming Languages — C++.
<https://wg21.link/n4988>
- [P1144] Arthur O'Dwyer. `std::is_trivially_relocatable`.
<https://wg21.link/p1144>
- [P2786R6] Mungo Gill, Alisdair Meredith. 2024-05-21. Trivial Relocatability For C++26.
<https://wg21.link/p2786r6>
- [P2839R0] Brian Bi, Joshua Berne. 2023-05-15. Nontrivial relocation via a new “owning reference” type.
<https://wg21.link/p2839r0>
- [P3239R0] Alisdair Meredith. 2024-05-22. A Relocating Swap.
<https://wg21.link/p3239r0>