

Standardized Constexpr Type Ordering

Document #: P2830R9
Date: 2025-01-10
Project: Programming Language C++
Audience: LWG
Reply-to: Nate Nichols
<natenichols@cox.net>
Gašper Ažman
<gasper.azman@gmail.com>

Contents

- 1 Abstract
- 2 Revision History
- 3 Motivation
 - 3.1 Motivating Examples
 - 3.1.1 Archetypal example: canonicalized typelist
 - 3.1.2 Canonicalizing policy-based libraries
 - 3.1.3 Canonicalized variant
 - 3.1.4 Canonicalized tuple
 - 3.1.5 API stability
- 4 Design discussion
 - 4.1 Desirable properties of `TYPE-ORDER(x, y)`
 - 4.1.1 Stability
 - 4.1.2 Free-standing
 - 4.1.3 Self-consistency
 - 4.1.4 Reflection compatibility
 - 4.1.5 non-goal: Consistency with `type_info::before()`
 - 4.2 Pros of implementation-defined
 - 4.2.1 ABI specifications already have to do this work
 - 4.2.2 The C++ standard doesn't own cross-compilation-unit semantics
 - 4.2.3 Any new proposal to C++ would have to consider the ordering

- 4.2.4 TYPE-ORDER(X, Y) already has public-facing API implications
- 4.3 Pros of completely defined order (not chosen)
 - 4.3.1 Static analyzers do not have a single backend
 - 4.3.2 consteval should be as portable as possible
 - 4.3.3 Producing and comparing two mangled names takes more memory to cache
- 5 Proposal
 - 5.1 Semantics
 - 5.2 Proposed Syntax
 - 5.2.1 Discussion
 - 5.2.2 Future extension
- 6 Proposed Wording
 - 6.1 Notes on wording
- 7 FAQ
 - 7.1 Why should this be standardized?
 - 7.2 Why not wait for reflection?
 - 7.3 But couldn't this be done faster with reflection?
- 8 Implementability
- 9 Acknowledgements
- 10 Appendix A: Discarded syntax options
 - 10.1 variable template `std::entity_ordering<X, Y>`
 - 10.2 reflection
 - 10.3 heterogeneous constexpr `std::type_identity::operator<=>` (bad)
 - 10.4 constexpr `std::__lift<arg>::operator<=>` (bad)
 - 10.5 Non-Option: constexpr `bool std::type_info::before()`
- 11 Appendix B: building `apply_canonicalized`
 - 11.1 Full code listing as tested and implemented
- 12 Appendix C: `__PRETTY_FUNCTION__` instability
- 13 Appendix D: pros/cons slide from EWG Wroclaw
- 14 Appendix ZZZ: This his how deep the rabbit hole goes
 - 14.1 Foreword
 - 14.2 Approach
 - 14.3 Structure of sort-key-tuples

- 14.4 decl-scopes
 - 14.4.1 Every declaration introduces a decl-scope.
- 14.5 What to do with undefined comparisons?
- 14.6 Modules
- 14.7 Namespaces
 - 14.7.1 The global namespace
 - 14.7.2 Named namespaces
 - 14.7.3 The anonymous namespace
 - 14.7.4 Namespace examples
- 14.8 cv-qualified types
 - 14.8.1 Qualifiers
 - 14.8.2 Scalar Types
 - 14.8.3 Array Types
 - 14.8.4 Expressions
 - 14.8.5 Lambdas
- 14.9 Old design that needs to be revamped
 - 14.9.1 Example 1: `class foo` is declared in `struct bar`:
 - 14.9.2 Example:
- 14.10 Atoms
- 14.11 identifiers
- 14.12 Kinds
 - 14.12.1 Identifiers
 - 14.12.2 Namespaces
 - 14.12.3 Types
 - 14.12.4 Ordering Class Types
- 15 References

§ 1 Abstract

As of C++23, `std::type_info` provides a stable but *unspecified* and non-constexpr order on types.

This paper explores a standardized, constexpr ordering of types in C++.

This paper is split into two parts:

1. the design of the exposition-only function `TYPE-ORDER(x, y)`, which deals with how such an order should be defined,
2. how to expose this capability to the programmer (syntax).

§ 2 Revision History

0. New Paper

1. Revision 1

- Introduce options to prevent changing `std::type_info::before`
- Anonymous namespaces can't be empty
- Add FAQ section
- Add motivating examples
- Add proposed syntax
- Add appendix

2. Revision 2

- Presented at the EWGi in Tokyo.
- Propose to make the `TYPE-ORDER(X, Y)` definition `constexpr` and *implementation-defined*, as suggested by EWGi and BSI reviewers.
- added wording
- added more motivating examples

3. Revision 3

- incorporates wording feedback from Jens Maurer, to be presented in an EWG telecon

4. Revision 4

- incorporated feedback from the EWG discussion on May 17th
- Added example of `__PRETTY_FUNCTION__` differences from Barry Revzin (thanks!)
- Incorporated the fact that EWG decided that the syntax is reasonable, pending LEWG review.
- Did a further pass on properly defining `SORT_KEY(x, y)`

5. Revision 5

- did a lot more work on what we would need to standardize if we wanted to completely define the ordering.
- presented to SG7 in Wroclaw, they said they want to fully define it

- presented to EWG in Wroclaw, they forwarded with strong consensus to leave it implementation defined
- wrote wording and relegated the sketch of the completely-defined order to a historical appendix

6. Revision 6

- added feature-test macro

7. Revision 7

- more wording fixes.

8. Revision 8

- Wording fixes from Jeff Garland, Jens Maurer, Tomasz Kamiński and Tim Song on the lwg mailing list.

9. Revision 9

- Editorial wording fixes from LWG review.

§ 3 Motivation

There is currently no way in portable C++ to sort types at compile-time.

Various libraries hack it together, mostly by utilizing `__PRETTY_FUNCTION__`, but all such methods are non-portable and error-prone (consider forward-declared enums - example in Annex). There are multiple stack-overflow questions on the topic.

[One such implementation](#) is part of the monadic functional library by Bronek Kozicki et al.

Fundamentally, we need a way to canonicalize sets and multisets of types. This is necessary when building any kind of compositional library in C++, from `std::execution` [P2300R7], mixin libraries, libraries of monadic components and operators, policy-based libraries, etc. The inability to sort and unique types leads to the instantiation of an combinatorial number of templates instead of just one per typeset, which leads to code-bloat and untenable compile-times. The problem is fundamentally that without typeset canonicalization, we generate different template specializations but equivalent behavior, and there is no way to avoid this.

The goal here is to provide a *flexible* language mechanism to let both `Foo<A, B, C>` and `Foo<C, B, A>` produce the same underlying `Foo_impl<A, B, C>`.

The reason we start with `TYPE-ORDER()` and not “just give me typesets” is that we need flexibility; consider

- a given library might want to deduplicate on a part and keep either last or first, or even make it ill-formed: `Foo<pair<A, X>, pair<B, Y>, pair<A, Z>>` might want

to be the same as `Foo_impl<pair<A, X>, pair<B, Y>>` or `Foo_impl<pair<A, Z>, pair<B, Y>>`

- a given library might actually just want canonicalized multisets: `Foo<A, B, A, A, C>` should perhaps be `Foo<A, A, A, B, C>`
- or treat the first one as special: `Matrix<float, policy1, policy2, policy3>` should only deduplicate policies.

We must provide `TYPE-ORDER` in order to sort and unique; they are required building blocks for any set primitive. Put another way, even if we standardized a set, we'd need to somehow canonicalize the order (due to mangling and debug info), leading us back here.

Without such canonicalization, it is infeasible to enumerate the set of function templates to instantiate in a separate compilation unit. For the same reason, it's utterly impossible to type-erase them in a fixed-size virtual function table.

§ 3.1 Motivating Examples

This section introduces the kind of code we would like to write, regardless of how this feature ends up being spelled. To not prejudice the reader in this section, please assume the existence of an exposition-only consteval macro `TYPE-ORDER(x, y) -> std::strong_ordering` whose arguments can be any cv-ref qualified types.

Note: while `TYPE-ORDER(x, y)` is defined on types, we can define a set of class templates that take an arbitrary template argument, and `TYPE-ORDER` those; this induces an order on all entities that can be template arguments.

Crucially, also consider the interactions with [\[P1985R3\]](#) and [\[P2841R1\]](#), which introduce concept and variable-template arguments.

§ 3.1.1 Archetypal example: canonicalized typelist

(Kept short because further examples give concrete needs).

```
struct A {};  
struct B {};  
  
// we want some way to enable  
static_assert(std::is_same_v<typeset<A, B>, typeset<B, A, A>>);
```

Furthermore, we need the order to be the same in different compilation units.

§ 3.1.2 Canonicalizing policy-based libraries

Consider the needs of a library for type-erasure; we'd like the user to specify the capabilities to erase. Fundamentally, these capabilities are a set.

Observe:

```
using T1 = lib::dyn<ostreamable, json_serializable, iterable<int>,
    scalar_multiplicable<int>>;
```

Different users of the library will likely provide these capabilities in different orders; this becomes especially problematic when writing functions:

```
using T2 = lib::dyn<json_serializable, ostreamable, iterable<int>,
    scalar_multiplicable<int>>;
~~~~~ flipped ~~~~~
void f(T2 const&);
```

We can solve this by having type aliases, but if these sets are *computed*, we are left without recourse:

```
int sum = 0;
auto pipeline1 =
    log(std::cerr) | sum_into(&sum) | imul(sum) | json_dump(std::cout);
//    dyn<ostreamable,    iterable<int>,    scalar_multiplicable<int>,
//        json_serializable>
auto pipeline2 =
    sum_into(&sum) | imul(sum) | json_dump(std::cout) | log(std::cerr);
//    dyn<iterable<int>,    scalar_multiplicable<int>,    json_serializable,
//        ostreamable>
```

The above pipelines need the same type-erased interface for its input, and neither is either T1 or T2.

§ 3.1.3 Canonicalized variant

The most obvious example is a canonicalized `std::variant`, that is, something like

```
template <typename... Ts>
using variant_for = apply_canonicalized<std::variant, Ts...>;
```

Building `apply_canonicalized` is not terribly difficult, as long as we have `TYPE-ORDER`. Please see the appendices on how to do it.

It *would* be nice if `apply_canonicalized` was a language built-in, but to do that, we need to first define `TYPE-ORDER(x, y)`. After we define `TYPE-ORDER`, putting `apply_canonicalized` into `type_traits` is a far simpler proposition.

Note: `apply_canonicalized` is roughly `mp11::mp_sort + mp11::unique` with the order derived from `TYPE-ORDER`.

Effectively any time a variant is type-indexed instead of integer-indexed, we would prefer the compiler generate just one type per set, instead of type-per-list.

Examples follow, but they are by no means exhaustive.

§ 3.1.3.1 Multiple error kinds and `std::expected`

Consider a `std::expected` for an algorithm that can fail in several ways:

```
auto read_packet(socket& sock)
-> expected<message, variant<io_error, decode_error>>;
```

Such an algorithm doesn't care about whether the error type is `variant<io_error, decode_error>` or `variant<decode_error, io_error>`. This is not as much of a problem in cases where the programmer writes the types by hand, but consider the generic situation where several errors must be aggregated generically:

```
template <typename Decoder>
auto read_packet(socket& sock)
-> expected<Decoder::message_type,
           variant</*uniqued and sorted io_error +
                  Decoder::error_types*/>>;
```

If we had member packs and pack aliases, and a nice built-in set, we could do the above as

```
template <typename... Ts>
...using set = __builtin_uniqued_sorted<Ts...>;

variant<...set<io_error, Decoder::...error_types...>...>;
```

§ 3.1.3.2 Suspension-point storage for asynchronous code

`std::execution` (and every other async library) needs a variant-type to store the current result at a suspension point. In general, since it models a “suspended” function call, it's analogous to some kind of

```
variant<
    closure_tuple<set_value, Arg1_1, Arg1_2, Arg1_3>, // first possible
                entry point
    closure_tuple<set_value, Arg2_1>,                // second possible
                entry point
    closure_tuple<set_error, ...>,
    ...
>
```

The venerable `rxcpp` library by Kirk Shoop calls this type a *notification*.

Consider the `transfer(scheduler)` algorithm; it receives all `set_value(PACK)`, `set_error(PACK)` and possibly `set_stopped()` parameter lists, *stores them in a* `variant<tuple<CHANNEL, PACK>...>`, and resumes with the values on a different scheduler once the compute resource becomes available.

The code using it looks like this:

```
some_computation | transfer(scheduler) | then([](auto&&...args)
                { /*...*/ })...
```

All the possible closure tuples are fundamentally unsorted, and generating the assembly, debug information, and everything else for this type more than once is a waste of compile-time and executable space (as well as instruction cache).

Furthermore, stamping out this type multiple times also duplicates the calling code, and given the different indices, COMDAT folding cannot deduplicate those code-paths.

§ 3.1.3.3 Compositional State machine models

The states of most state-machines are fundamentally unordered, as are the message types they receive. If a state machine is generated (such as in a compile-time regex library), canonicalizing identically-behaving components would lead to smaller state-machines and faster compile- and execution times.

§ 3.1.4 Canonicalized tuple

Similarly to a canonicalized variant, a canonicalized_tuple could be implemented as

```
template <typename...Ts>
using canonical_tuple = apply_canonicalized<tuple, Ts...>;
// or
using canonical_tuple = tuple<...set<Ts...>...>;
```

This is far less useful on its own than the variant, since initializing that type will be quite the chore. Instead, canonicalized tuples appear as a result of the environment monads.

std::execution mixes in an environment, which consists at least of a stoppable_token, a scheduler, and possibly other things like a domain, and, at least in the authors' implementation, arbitrary other things. An allocator is also optionally part of the environment.

The environment<pair<Tags, Ts>...> is a product type that is a map from Tag to a value of type T, say std::stop_token to never_stop_token{} and scheduler to thread_pool_scheduler.

One can push and pop things off the environment in different parts of the pipeline; if the pushes come in the wrong order, the environment is difficult to keep canonicalized.

```
auto ss = std::stop_source();
auto tp = std::static_thread_pool(5);
auto env = empty_env{}
    | push_env<stop_token>(ss.get_token())
    | push_env<scheduler_t>(tp.get_scheduler())
    | push_env<mylib::numa_domain_t>(2);
```

If you imagine those three calls to happen in different algorithm transformers, it's quite likely they'll be in a different order, manifesting a different type, which doesn't make much sense - the type behaves identically, this just leads to code bloat.

§ 3.1.5 API stability

The authors have observed a necessity to sometimes split parts of such pipeline compositions into different compilation units due to excessive compile times and repetitive compilation.

This has proven difficult due to the brittle nature of type ordering of the names of explicit template specializations involved. By far, the main culprit has been the lack of a sensible canonicalization of names, as the order changes far more frequently than the set of types.

Consider:

```
do_something()
| error_if_odd()           // there is no semantic change
| error_if_we_too_long()  // if we flip these lines
| ...
```

Flipping the lines will flip the two exit error types in the error list, however.

Having a defined, stable order of types proved invaluable (even if we did hack the type order manually).

§ 4 Design discussion

The first two revisions of this proposal tried to construct a full type ordering from first principles; EWGi then requested we change to an implementation-defined order (use the mangler), but EWG subsequently reversed that decision; SG7 affirmed this decision in Wroclaw, only to be reversed by EWG again to implementation-defined.

This paper proposes the implementation-defined version.

§ 4.1 Desirable properties of `TYPE-ORDER(x, y)`

§ 4.1.1 Stability

The order should be the same across compilation units. This is key for generating ABI-compatible vtables, for instance.

It also should be stable through time. An order based on an ABI-mangling scheme satisfies this notion, at least.

§ 4.1.2 Free-standing

The order should be available on freestanding implementations. One crucial use-case is replacing the exception mechanism on those platforms with return values of `std::expected` or similar, and compromising that is undesirable.

§ 4.1.3 Self-consistency

The ordering should be self-consistent, that is, for all possible template arguments `T`, `U`, and any unary template `some_template`:

`TYPE-ORDER(T, U) == TYPE-ORDER(some_template<T>, some_template<U>).`

§ 4.1.4 Reflection compatibility

Any operator<=>(std::meta::info, std::meta::info) should be consistent with this one.

While std::meta::info [P2320R0] objects can reflect more entities than just types and values (mainly: expressions), any ordering defined on them should be finer than this one, and specifically consistent with it. We should do something that reflection can subsume.

However, it doesn't seem like that proposal will define an ordering on info objects, and we need this solved sooner than that.

§ 4.1.5 non-goal: Consistency with type_info::before()

Ideally, whenever typeid(x).before(typeid(y)) == true, then TYPE-ORDER(x, y) == std::strong_ordering::less. The converse obviously cannot be true, since TYPE-ORDER(x, y) is finer than the order induced by type_info::before().

However, the standard currently says

The names, encoding rule, and collating sequence for types are all unspecified and may differ between programs.

Since this paper requires that the order not differ between programs, exposing this as a normative requirement is impossible without tightening this wording, which is outside the scope of this paper.

At least at present, some implementations decide this order at program startup time, so this is not implementable in general.

§ 4.2 Pros of implementation-defined

The overwhelming response of people working on implementations, as well as people sitting on the Itanium ABI standards committee, was that this is a fool's errand, and we need only look at the bugreports for the Itanium ABI to see why. This view won.

§ 4.2.1 ABI specifications already have to do this work

As a taster, consider the following conundrum:

```
using what_type = decltype(
    [](this auto&& self,
        decltype([](decltype(self)&){}) x = {}){ return x; }());
// ^^ outer ^^ inner
```

The inner lambda's "name" is clearly dependent on the outer's, but it also goes the other way! The ABI standards already have to deal with such mangling conundrums, and duplicating their work in the C++ standard itself seems highly counterproductive,

especially since the ABI standards already accomplish all the stability guarantees we would want.

§ 4.2.2 The C++ standard doesn't own cross-compilation-unit semantics

Another problem is that within the C++ standard, we cannot legislate what happens with orderings between compilation units. Granted, we do not need this for `constexpr`, but we do need it to be stable and make sense, so a certain amount of common sense will be expected from implementations *anyway*.

The recommendation was overwhelmingly to just let the order be implementation-defined, and let the implementations do the sensible thing.

§ 4.2.3 Any new proposal to C++ would have to consider the ordering

“punting it off to implementations” saves the C++ standardization process from having to figure this out for every new proposal that touches types; implementations, however, already have representation on the committee, and will veto any truly unimplementable things in this space.

§ 4.2.4 TYPE-ORDER(X, Y) already has public-facing API implications

The “normalized” versions of types will inevitably show up in function signatures; if different compilers on the same platform produced different orderings, compilation units from different compilers would refuse to link, despite both “signatures” being spelled identically.

Letting the compiler vendors synchronize based on the mangling scheme or something equivalently useful is the same forum as the current ABI discussion forums; it seems the appropriate venue to standardize the ordering anyway, without making WG21 duplicate the work.

§ 4.3 Pros of completely defined order (not chosen)

§ 4.3.1 Static analyzers do not have a single backend

This proved decisive. Since static analyzers and compilers that produce intermediate representations before choosing a backend are part of the ecosystem, it is important to be consistent and not rely on a particular mangler.

§ 4.3.2 `constexpr` should be as portable as possible

The abstract machine used at constant evaluation time has far fewer parameters than the runtime one. We should keep them to a minimum, and type ordering is not a necessary parameter. We should keep implementations consistent.

§ 4.3.3 Producing and comparing two mangled names takes more memory to cache

The key-tuples proposed in this paper are recursive; the parts of the type-tuple merely refer to the tuples of the constituent parts; thus, they can be independently cached and compared as-needed.

§ 5 Proposal

§ 5.1 Semantics

Let X and Y be (possibly cv-ref) qualified types, and $\text{TYPE-ORDER}(X, Y)$ be an exposition-only macro.

Then, $\text{TYPE-ORDER}(X, Y)$ denotes a constant expression of type `std::strong_ordering`.

- `std::same_as<X, Y> == true` if and only if $\text{TYPE-ORDER}(X, Y) == \text{std::strong_ordering::equal}$.
- Otherwise, $\text{TYPE-ORDER}(X, Y)$ is either `std::strong_ordering::less` or `std::strong_ordering::greater`. Which of those is implementation-defined, subject to the following semantic constraints.

Implementations must define $\text{TYPE-ORDER}(X, Y)$ such that it is an ordering, that is, it is transitive and antisymmetric; that is

- if $\text{TYPE-ORDER}(X, Y) == \text{std::strong_ordering::less}$, then $\text{TYPE-ORDER}(Y, X) == \text{std::strong_ordering::greater}$ and vice versa
- if $\text{TYPE-ORDER}(X, Y) == \text{std::strong_ordering::less}$ and $\text{TYPE-ORDER}(Y, Z) == \text{std::strong_ordering::less}$, then $\text{TYPE-ORDER}(X, Z)$ is also `std::strong_ordering::less`.

Implementations are encouraged, but not required to, make the order recursively-consistent. In other words:

For any class template `template <..., typename Pi, ...> class X`, where P_i is the i -th template argument, and any two types T and U : $\text{TYPE-ORDER}(T, U) == \text{TYPE-ORDER}(X<..., T, ...>, X<..., U, ...>)$ if only the i th template argument is varied.

If an implementation makes the order recursively consistent, it should document this fact.

§ 5.2 Proposed Syntax

EWG affirmed a library name to access the ordering.

```
// <compare>
template <typename T, typename U>
inline constexpr std::strong_ordering type_order_v = TYPE-ORDER(T, U); /*
    see below */
template <typename T, typename U>
struct type_order : integral_constant<strong_ordering, type_order_v<T, U>>
{
};
```

As a special provision for this specific metafunction, we allow the type arguments for it to be incomplete types.

§ 5.2.1 Discussion

This seems like a pretty good choice. It does not need a new keyword, only depends on `<compare>`, and the name seems relatively discoverable.

It's also freestanding, since it doesn't depend on `<typeinfo>`.

We do not want `<type_traits>` to depend on `<compare>`, so we put it into `<compare>` directly.

§ 5.2.2 Future extension

Once we have pack aliases, the authors will propose the following two metafunctions, to be implemented using compiler intrinsics:

```
// as a separate library proposal, once member packs make it
template <typename... Ts>
using ...typemultiset = /* pack of Ts, sorted by type_order_v */;
template <typename... Ts>
using ...typeset = /* unique'd ...typemultiset<Ts...>... */;
```

§ 6 Proposed Wording

In 17.11.1 [\[compare.syn\]](#), add

```
// [compare.type] type ordering
template <class T, class U>
struct type_order;
template <class T, class U>
constexpr strong_ordering type_order_v = type_order<T,
    U>::value;
```

At the end of 17.11 [\[cmp\]](#), just before 17.12 [\[support.coroutine\]](#), add:

17.11.7: Type Ordering [\[compare.type\]](#)

- 1 There is an implementation-defined total ordering of all types. For any (possibly incomplete) types X and Y , the expression $TYPE-ORDER(X, Y)$ is a constant expression (7.7 [\[expr.const\]](#)) of type `strong_ordering` (17.11.2.4 [\[cmp.strongord\]](#)). Its value is `strong_ordering::less` if X precedes Y in this implementation-defined total order, `strong_ordering::greater` if Y precedes X , and `strong_ordering::equal` if they are the same type.

[Note: `int`, `const int` and `int&` are different types. – end note]

```
template <class T, class U>
struct type_order;
```

- 2 The name `type_order` denotes a `Cpp17BinaryTypeTrait` (21.3.2 [meta.rqmts]) with a base characteristic of `integral_constant<strong_ordering, TYPE-ORDER(X, Y)>`.
- 3 *Recommended practice:* The order should be lexicographical on parameter-type-lists and template argument lists.

Add a feature-test macro into 17.3.2 [version.syn] in section 2

```
#define __cpp_lib_type_order 2025XXL // also in <compare>
```

§ 6.1 Notes on wording

- I’d like to thank Lewis Baker and Davis Herring for reminding me that I should probably make an introductory paragraph.
- Davis Herring suggested adding a note to remind readers that `int`, `const int` and `int&` are different types.
- Jeff Garland, Tim Song, Tomasz Kamiński and Jens Maurer for their wording fixes and suggestions.

Design note: the reason we do not state a recommended practice of what to do with function return types is to allow doing *whatever the mangling does*, which for the itanium ABI means *sort by return type first, and then by function parameter list*. It is feasible for there to be a mangling that would put the return type last.

§ 7 FAQ

§ 7.1 Why should this be standardized?

Because we have no way to reliably order types across compilation units at compile-time.

§ 7.2 Why not wait for reflection?

It’s a good question. However, reflection will do *nothing* for this problem by itself; the user will still have to implement ordering using `constexpr` functions, which have no hope of being as fast as a compiler-provided built-in.

User-programmed functions also won’t adapt to language evolution; this feature will.

Finally, sorting is arbitrary; having it be consistent throughout the software ecosystem is potentially a great enabler of interoperability.

§ 7.3 But couldn’t this be done faster with reflection?

No; Peter Dimov shares an interesting anecdote.

I have in Mp11 the algorithm `mp_unique`, which takes a list of types and removes the duplicates. In the course of writing the reflection papers, their authors occasionally took Mp11 code examples and tried to show how they are elegantly implemented using value-based reflection metaprogramming.

So, you take a `vector<info>` that contains types, and then you simply apply the existing algorithm `std::unique` to it, et voila... oh wait.

`std::unique` wants a sorted range, and you can't `std::sort` the `info` vector, because `info` objects aren't ordered, even when they refer to types.

§ 8 Implementability

The proposal has no questions on whether it *can* be implemented - the question is about the definition of the order.

The concerns raised by the implementers so far have been mostly around having to bring the name mangler from the backend and make it accessible to the compiler front-end, because the obvious implementation is to define the comparison result on the mangled type strings; this option has been rejected by EWG.

Making this order match up with `type_info::before` is a matter of bringing the name mangler for the correct platform to the frontend.

It is the opinion of the authors that this doesn't seem to be a layering violation, as the name mangling for a given platform is analogous to other platform properties, such as the size and alignment of pointers.

If a platform does not have a name mangling strategy, *any* name mangling scheme will still result in a standards-conforming implementation; however, after much discussion, it seems a better direction to completely define the order in the standard itself.

§ 9 Acknowledgements

Thanks to all of the following:

- Davis Herring for his suggestions on ordering non-type template parameters.
- Ville Voutilainen for his critique of examples, and providing a simple way of explaining the motivation
- Peter Dimov for a helpful anecdote, now in the FAQ.
- Erich Keane for pushing us back to the “implementation-defined” territory.
- Jens Maurer for his thorough review of the initial proposed wording and his guidance.

§ 10 Appendix A: Discarded syntax options

§ 10.1 variable template `std::entity_ordering<X, Y>`

Specifically:

```
template <universal template X, universal template Y>
inline constexpr strong_ordering entity_order_v = TYPE-ORDER(X, Y); /* see
    below */
template <universal template X, universal template Y>
struct entity_order : integral_constant<strong_ordering, entity_order_v<X,
    Y>> {};
```

This is a better option than Option 1 if we get universal template parameters, as we really want to also order class templates, not just types.

However, without universal template parameters, we really don't have much of a choice but to reach for Option 1.

The name `entity_order` is also slightly less obvious than `type_order`, but metaprogrammers shouldn't have trouble finding either.

§ 10.2 reflection

Specifically:

```
constexpr std::partial_ordering partial_order(std::meta::info x,
    std::meta::info y) {
    return __comparable(x, y) ? TYPE-ORDER(x, y) :
    std::partial_order::unordered;
}
```

We could standardize a type order as a function on `std::meta::info` objects in `std::meta`. However, once we're in `std::meta::info` space, it's more difficult to know which reflections are comparable and which aren't, so such a function would need to return a `std::partial_order`, which seems decidedly less desirable.

It also means we'd need to pass the correct kind of reflection into the ordering function, which is a bit less intuitive than just `decltype` or just the template parameter that we already have.

§ 10.3 heterogeneous `constexpr std::type_identity::operator<=>` (bad)

Specifically:

```
template <typename T, typename U>
constexpr std::strong_ordering operator<=>(std::type_identity<T>,
    std::type_identity<U>);
```

Pros:

- No new names.

Cons:

- Less discoverable than a new type-trait
- Requires template instantiations of `type_identity<T>`, `type_identity<U>`, as well as `operator<=>` overload resolution and substitution, which is quite expensive compared to a single `type_order_v<T, U>` direct substitution and lookup.
- An extra `operator<=>` overload in the `std` namespace is a drag on overload resolution
- adds `<compare>` to `<type_traits>`, since that's where `type_identity` is
- too cute?
- doesn't work for nontypes

§ 10.4 `constexpr std::__lift<arg>::operator<=>` (bad)

This option means we add `template <universal template> struct __lift {};` into `<type_traits>` and define `operator<=>` for it.

Pros:

- ... we'll need a lift sooner or later?

Cons:

- all the cons of `type_identity`
- even more nonobvious
- still needs a new name
- needs a tutorial to find

§ 10.5 **Non-Option:** `constexpr` `bool` **`std::type_info::before()`**

(Included because it's a common question.)

It would be nice, but alas, operates on *cv-unqualified* versions of the referenced type, so it's not sufficient.

`constexpr std::strong_order(std::type_info, std::type_info)` has similar issues.

§ 11 **Appendix** **B:** **building** **`apply_canonicalized`**

We will need a small metaprogramming library; a filter is difficult to do otherwise.

```
struct undefined;
template <typename... Ts> struct list {};

// apply<F, list<Ts...>> -> F<Ts...>
template <template <typename...> typename, typename> extern undefined
    _apply;
template <template <typename...> typename F, template <typename...>
    typename L,
    typename... Ts>
F<Ts...> _apply<F, L<Ts...>>;
template <template <typename...> typename F, typename List>
using apply = decltype(_apply<F, List>);

// concatenate<list<Ts...>, list<Us...>, list<Vs...>> -> list<Ts...,
    Us..., Vs...>
template <typename...> extern undefined _concatenate;
template <typename... Ts> list<Ts...> _concatenate<list<Ts...>>;
template <typename... Ts, typename... Us, typename... Lists>
decltype(_concatenate<list<Ts..., Us..., Lists...>())
    _concatenate<list<Ts...>, list<Us...>, Lists...>;
template <typename... Ts>
using concatenate = decltype(_concatenate<Ts...>);

// select: list<T> if true, list<> if false
template <bool v, typename T> extern list<> _select;
template <typename T> list<T> _select<true, T>;

template <bool v, typename T>
using select = decltype(_select<v, T>);
```

Canonicalization is now just a basic not-in-place quicksort-ish thing:

```
template <typename.../*empty*/> extern list<> _canon;
template <typename... Ts>
using canonicalized = decltype(_canon<Ts...>);

// a canonicalized T is just T
template <typename T>
list<T> _canon<T>;

template <typename T, typename... Ts>
concatenate<
    // canonicalized things less than T
    apply<canonicalized, concatenate<select<(TYPE-ORDER(Ts, T) < 0),
        Ts>...>>,
    list<T> /*T*/, // ~~~~~
    // canonicalized things greater than T
    apply<canonicalized, concatenate<select<(TYPE-ORDER(Ts, T) > 0),
        Ts>...>>
    > // ~~~~~
    _canon<T, Ts...>;
```

We now have canonicalized<Ts...> - but this still leaves list as a special type which we'd rather not expose to the user. Onto apply_canonicalized:

```
template <template <typename...> typename F, typename... Ts>
using apply_canonicalized = apply<F, canonicalized<Ts...>>;
```

§ 11.1 Full code listing as tested and implemented

Here for completeness, feel free to skip.

```
#include <compare>
#include <type_traits>

struct undefined;

#define TYPE-ORDER(x, y) type_order_v<x, y>

template <typename X, typename Y>
constexpr inline std::strong_ordering type_order_v;

template <template <typename...> typename, typename>
extern undefined _apply;

template <template <typename...> typename F, template <typename...>
    typename L,
    typename... Ts>
F<Ts...> _apply<F, L<Ts...>>;

template <template <typename...> typename F, typename List>
using apply = decltype(_apply<F, List>);

// some user-type
template <auto x>
struct value_t : std::integral_constant<decltype(x), x> {};
template <auto x>
inline constexpr value_t<x> value_v{};

// built-in
template <auto x, auto y>
constexpr inline std::strong_ordering type_order_v<value_t<x>, value_t<y>>
    =
    x <=> y;

template <typename... Ts>
struct list {};

template <typename...>
extern undefined _concatenate;
template <typename... Ts>
list<Ts...> _concatenate<list<Ts...>>;
template <typename... Ts, typename... Us, typename... Lists>
decltype(_concatenate<list<Ts..., Us..., Lists...>
    _concatenate<list<Ts...>, list<Us...>, Lists...>);

template <typename... Ts>
using concatenate = decltype(_concatenate<Ts...>);

template <bool v, typename T>
```

```

extern list<> _select;
template <typename T>
list<T> _select<true, T>;

template <bool v, typename T>
using select = decltype(_select<v, T>);

template <typename...>
extern list<> _canon;
template <typename... Ts>
using canonicalized = decltype(_canon<Ts...>);

template <typename T>
list<T> _canon<T>;

template <typename T, typename... Ts>
concatenate<
    apply<canonicalized, concatenate<select<(TYPE-ORDER(Ts, T) < 0),
        Ts>...>>,
    list<T>,
    apply<canonicalized, concatenate<select<(TYPE-ORDER(Ts, T) > 0),
        Ts>...>>
>
_canon<T, Ts...>;

static_assert(std::same_as<canonicalized<value_t<0>, value_t<-1>,
    value_t<-1>, value_t<1>>>, list<value_t<-1>, value_t<0>,
    value_t<1>>>>);

```

§ 12 Appendix C: __PRETTY_FUNCTION__ instability

This example is available at <https://godbolt.org/z/ojb9TnE99> .

Consider the following program, contributed by Barry Revzin:

```

#include <print>

enum class E;
template <E> struct C;
#ifdef DEFINED
enum class E { hi, gašper };
#endif

template <class T>
void show() {
    std::print("{}\n", __PRETTY_FUNCTION__);
}

int main() {
    show<C<E(0)>>>();
}

```

When compiled with `-std=c++23`, it yields the following output:

```
void show() [with T = C<(E)0>]
```

However, with `-std=c++23 -DDEFINED`, it produces a different one:

```
void show() [with T = C<E::hi>]
```

This makes external names of types incorporating enums as non-type template arguments have inconsistent between translation units.

§ 13 Appendix D: pros/cons slide from EWG Wroclaw

Implementation-defined or fully specified by the standard?

- Implementation defined:
 - Pro: ABIs already did all the work
 - Cons:
 - ABIs don't agree
 - Frontend doesn't know ABI for static analysis tools
 - Layering violation (Erich says that's already happening)
 - Compilers need to agree to have compatible ABI
 - Not necessarily self-consistent (name mangling uses compression) - Erich says it's probably ok
- Fully specified:
 - Pros:
 - Fully portable, including static analysis tools
 - Faster than mangling during constexpr evaluation (Erich says they can short-circuit)
 - (comparison does not require stringifying long symbol names, it short-circuits quickly)
 - Does not require the frontend to know the ABI (helps IDEs) (Erich says not true)
 - Cons:
 - Lots of work
 - Anonymous entities require a completely new-to-standard notion of a “declaration scope” with all the template arguments of all enclosing scopes

— We need to continue to specify ordering for every change to language entities

§ 14 Appendix ZZZ: This his how deep the rabbit hole goes

This section is left in the paper as a hint to the reader of how horrible specifying all of this would be for the language, without punting it to ABI specifications.

It is left here as the mess it was at the end of the attempt.

§ 14.1 Foreword

This section sets out an approach for defining an ordering of all compile-time entities; that is, the definition of `SORT_KEY(X)` for a given cv-qualified type `X`; and the comparison function between these sort-key-tuples.

Note to reviewers: *any* well-defined order will satisfy the design requirements; the chief design element here is whether we have achieved a well-ordering, not what exactly it is. The only other considration is whether we can evolve the order as we introduce new entities into the language.

§ 14.2 Approach

1. We define a **lowering to a sort-key-tuple** for every entity in the language.
2. The order is then defined on these *sort-key-tuples*.

The entities we must order are:

1. namespaces and modules
 1. the global namespace
 2. named namespaces
 3. the anonymous namespace
 4. modules
2. cv-qualified types
 1. scalar types
 2. array types
 3. class and union types
 4. class and union template specializations
 5. anonymous versions of all of the above

6. function types

3. templates

1. function templates

2. class templates

3. variable templates

4. concepts

5. type alias templates

6. deduction guides

4. partial template specializations

1. partial function template specializations

2. partial class template specializations

3. partial variable template specializations

5. parameters

1. runtime-parameters

2. type-parameters

3. constant-parameters

4. type-template parameters

5. concept-parameters

6. variable-template parameters

7. parameter packs

6. constants

1. integral constants

2. floating-point constants

3. pointer and reference constants

1. to functions

2. to named entities of static storage duration

3. to expressions

4. class-type constants

1. lambda literals

2. constants of compound types usable as template arguments

7. expressions

§ 14.3 Structure of sort-key-tuples

Every *sort-key-tuple* is of the form $(element \dots)$.

where an element is one of:

- *atom* (see [atoms](#))
- *name*
- *constant*
- *sort-key-tuple*

These tuples are then ordered lexicographically (ties broken in favor of shorter tuple), atoms first, then names, then constants, then other *sort-key-tuples*.

Let us name this transformation as `sort_key(entity)`.

The rest of the design is concerned with defining this transformation.

§ 14.4 decl-scopes

A sort-key for most entities is of the form

```
sort_key(x) = (sort_key(decl-scope(x)), x-dependent-elements...)
```

This section deals with the decl-scope of an entity.

To deal with anonymous types and templated entities like hidden friends, we have to observe more than just the namespace and class the entity is declared in; we must also observe anything at all that could contribute to the declaration, such as template parameters and arguments of partial and full template specializations of both class and function templates, alias templates, concepts, their parameter numbers (in case of default arguments), and so on. After that, we must make special allowances for entities that cannot be influenced by such surroundings, such as forward declarations of named classes.

§ 14.4.1 Every declaration introduces a decl-scope.

- Every declaration *of an anonymous entity* is ordered lexicographically within its decl-scope.
- Ties are not broken; there should be no conflicts, as all the context for a given *sort-key(x)* should be provided by the decl-scope of the declaration. (notably - all relevant parameters *and* their values should be part of the sort-key for a given declaration of an anonymous entity)
- Named entities' decl-scope is their normal declaration scope.

§ 14.5 What to do with undefined comparisons?

Some comparisons could be left to future revisions because of omissions or incomplete work. If a program /observes/ such a comparison, the program should be ill-formed (diagnostic required).

This is to enable fixing the issue in a future revision of the standard. It is the intention of the authors of this paper that such comparisons should be exceedingly difficult to observe in practice in normal code, since most comparisons should be tie-broken fairly early on.

§ 14.6 Modules

Every entity that belongs to a module has that module's full name as a string as the first part of its `sort_scope`. If an entity belongs to the global module, it has the atom *global-module* as its name.

§ 14.7 Namespaces

This section deals with namespaces and their sort-keys.

§ 14.7.1 The global namespace

Let x be the global namespace.

```
sort_scope(x) := ()  
sort_key(x) := (sort_scope(x), _namespace_, _global-namespace_)
```

§ 14.7.2 Named namespaces

Every namespace (except for the global namespace) has a parent namespace, which is its immediately enclosing namespace.

let x be some named namespace with the simple-name “namespace_name”.

```
sort_scope(x) := sort_key(parent_namespace)  
sort_key(x) := (sort_scope(x), _namespace_, "namespace_name")
```

§ 14.7.3 The anonymous namespace

The anonymous namespace sorts after all the other namespaces.

Let x be some anonymous namespace, with parent namespace `parent_namespace`.

```
sort_scope(x) := sort_key(parent_namespace)  
sort_key(x) := (sort_scope(x), _namespace_, _anonymous_)
```

This is ok, as there is only one anonymous namespace per namespace per compilation unit.

§ 14.7.4 Namespace examples

```
sort_key(::std::ranges) = (
    (
        ((), _namespace_, _global-namespace_), // sort_key(global-
        namespace)
        _namespace_, "std"
    ), // sort_key(::std)
    _namespace_, "ranges"
);
```

§ 14.8 cv-qualified types

§ 14.8.1 Qualifiers

Qualifiers are each assigned a score

```
&: 1
&&: 2
const: 3
volatile: 6
```

and ordering lowest-first after summing them.

Any implementation-defined qualifiers get a score that is 2x the largest score in the table; for instance `__restrict` gets 12.

Therefore, for an unqualified type `T`, the order of all possible qualified types would be:

```
0 T
1 T &
2 T &&
3 T const
4 T const &
5 T const &&
6 T volatile
7 T volatile &
8 T volatile &&
9 T const volatile
10 T const volatile &
11 T const volatile &&
```

This is accomplished by putting the qualifier sequence into the tuple just after the typename.

For a given type `T`, we define

```
qualifier_sequence(T const volatile &) := const volatile &
qualifier_sequence(T const volatile &&) := const volatile &&
```

so that we can put it into the `sort_key` tuple later.

§ 14.8.2 Scalar Types

All scalar types are built-in types, except for user-defined enumerations, which should be ordered as if they were class types.

Simple scalar types (everything in this section but function types and pointer types) are not ordered by “name” - they are ordered using the rules in the table. All atoms are still ordered before them, any name is ordered after them, as are all sort-key-tuples.

This is because some of the built-in types do not have names, only type aliases (such as `decltype(nullptr)`), and we do not order types by their aliases.

This causes any built-in scalar types to be ordered before any compound types.

In case of ties, built-in types with simple names shall be ordered before any nameless types.

In particular, scalar types shall be ordered as follows:

1. `void` comes first because it's not reifiable,
2. the type of `std::nullptr_t` as the first monostate
3. any other monostates, if added, sorted alphabetically by their common names (to be specified explicitly if added)
4. `bool` as the first bi-state
5. any other bi-states, if added, sorted alphabetically.
6. Raw-memory types (`char`, `signed char`, `unsigned char`) (`std::byte` is an enumeration in `std` so it falls under different rules)
7. Integral types in order of size, signed before unsigned (`short`, `unsigned short`, `int`, `unsigned int`, `long`, `unsigned long`, `long long`, `unsigned long long`, followed by any implementation-defined wider integral types like `__int128_t` etc.). Intersperse any implementation-defined built-in integral types as needed between the above following those rules.
8. Any remaining character types that are not type-aliases of any of the above, including unicode, according to the following rules: smallest first, unicode-specific variants after non-unicode variants.
9. Floating-point types, in order of size. In case of ties, `float`, `double` and `long double` come before any other floating point types of the same size. Any decimal-floating-point types come after binary-floating-point types; if multiple floating point types of the same bit-length exist, break ties by the bit-size of the exponent, lower first.
10. Implementation-defined vector-register types, ordered by the the integral type they consist of, and then by bit-size (example: `u8x16` from gcc's documentation).
11. Function types (internally ordered by rules in section [Function Types](#))
12. Pointer types (internally ordered by their pointee-type)

13. Pointer-to-member types (internally ordered by pointee-type)

§ 14.8.3 Array Types

TODO: this section is not complete.

Array types shall be ordered after scalar types but before class types.

Given a non-array type T

```
sort_key(T[]) := (sort_key(T), [])  
sort_key(T[n]) := (sort_key(T), [], n)
```

Multidimensional arrays, as an example:

```
sort_key(int[][n]) = (sort_key(int[]), [], n) = ((sort_key(int), []), [],  
n);
```

Notice:

```
sort_key(int[]) < sort_key(int[2])
```

because the shorter tuple wins a tie.

§ 14.8.4 Expressions

Example (courtesy of Lénárd Szolnoki)

```
template <auto x>  
struct S {};  
  
template <int x>  
void foo(S<2*x>) {}  
  
inline constexpr auto a = &foo<0>;  
  
template <int x>  
void foo(S<x>) {}  
  
inline constexpr void (*b)(S<0>) = &foo;  
  
inline constexpr auto x = TYPE-ORDER(S<x>, S<y>); // not equal
```

We must therefore sort arbitrary expressions; I propose we avoid that by making such comparisons, should they occur, ill-formed, until we are forced to define them.

§ 14.8.5 Lambdas

A lambda type is an anonymous type, which means it's lexicographically numbered within its `sort_scope`.

Example:

```
template <auto T> C {};
static_assert(type_order_v<
    C<[]{}>, // lambda-key: (sort_key(static_assert), _type_,
    _anonymous_, 1)
    C<[]{}> // lambda-key: (sort_key(static_assert), _type_,
    _anonymous_, 2)
    >
    == std::strong_order::less);
```

TODO: port stuff from bottom.

§ 14.9 Old design that needs to be revamped

§ 14.9.1 Example 1: `class foo` is declared in `struct bar`:

```
struct bar { class foo; }
```

```
sort_key(foo) = (sort_key(bar), sort_key(foo)) = ((type, bar), (type, foo,
))
```

This shall hold for any of the above named scopes.

§ 14.9.2 Example:

Given

```
namespace foo::bar {
    struct i;
}

namespace baz {
    struct j;
}
```

Then:

- `sort_key(foo::bar::i)` is `((namespace, foo), (namespace, bar), (type, i, ))`.
- `sort_key(baz::j)` is `((namespace, baz), (type, j, ))`

When compared, the result is that `baz::j < foo::bar::i`, since namespace `baz` precedes namespace `foo`.

§ 14.10 Atoms

The atoms of *key-tuples* are ordered as follows:

1. *global-module*
2. *global-namespace*
3. *anonymous*

4. kinds (see [kinds](#))
5. qualifiers (see [qualifiers](#))
6. `[]` (array of unknown bound)
7. `*` (pointer)
8. ellipsis `(... in f(...))`
9. parameter pack `(... in typename...)`
10. template parameter *auto*
11. template parameter *typename*
12. template parameter *T*

§ 14.11 identifiers

These are simple names (identifiers).

Examples: “pair”, “string”, “unique_ptr”, “std”, “hashmap”, “absl”.

§ 14.12 Kinds

There are the following kind tokens that can appear in *key-tuples*.

1. value
2. namespace
3. type
4. class template
5. type alias template
6. variable template
7. concept
8. function

Note: everything but “values” is pretty simple, but we haven’t dealt with values extensively yet with the R1 of this paper, though we should just defer to `<=>` and require a default strong structural ordering for values that may be template arguments.

§ 14.12.1 Identifiers

§ 14.12.1.1 Simple Names

Most names are strings that are valid (atomic) identifiers. Those are just themselves:

```
namespace foo::bar { struct baz; }
```

foo, bar and baz are such atomic identifiers.

§ 14.12.1.2 Unnamed entities

Unnamed entities are all assigned a key-tuple of their decl-scope and then numbered lexically, consecutively, starting with zero, with the counter being scoped to their decl-scope.

14.12.1.2.1 Decl-scopes

Decl-scopes are:

- namespaces
- class definitions
- union definitions
- function declarations
- function default argument definitions
- function definitions
- enumeration definitions
- type aliases
- pack expansions
- template parameter declarations
- definitions of template argument defaults
- class template declarations
- class template definitions
- definitions of partial class template specializations
- definitions of class template specializations
- function template declarations
- function template definitions
- function template default argument definitions
- type alias templates
- concept definitions
- default initializers for class members

14.12.1.2.2 Examples

Function declarations are scoped to the function itself.

Consider a lambda that appears as a default argument of a function template:

```
template <typename T>
void f(T x = []{ return T{0}; }());
//          ^^^^^^^^^^^^^^^^^^ this one
```

The *key-tuple* for `f<int>` is:

```
(function, (f, (type, int)), (type, void), ((type, int)))
```

The *key-tuple* for the lambda is:

```
((function, (f, (type, int)), (type, void), ((type, int))), (type,
(lambda, 0), )).
```

Note: because of the regular structure of *key-tuples*, such anonymous classes will compare greater than any entity that has a simple identifier, due to tuples comparing greater than atoms (which simple names are).

14.12.1.2.3 Lambda types

As unnamed entities - the types of lambdas are ordered first by where they are declared, then by declaration (lexical) order.

In effect, we assign them the name `(lambda, #)` where `#` is the count of other unnamed entities in the decl-scope.

```
namespace Banana {
    auto i = [](int) -> void {}; // 0th lambda instantiated in Banana
}

namespace Apple {
    auto i = [](float) -> int {}; // 0th lambda instantiated in Apple
    auto j = []() -> std::string {}; // 1st lambda instantiated in Apple
}
```

These would produce the following tuples:

```
sort_key(decltype(Banana::i)) = ((namespace, Banana), (type, (lambda, 0),
));
sort_key(decltype(Apple::i))  = ((namespace, Apple), (type, (lambda, 0),
));
sort_key(decltype(Apple::j))  = ((namespace, Apple), (type, (lambda, 1),
));
```

Note: the empty bit after the identifier is the empty qualifier pack.

14.12.1.2.4 Unnamed struct and union types

They are named, respectively, `(class, #)` and `(union, #)`.

Example: in a type alias like `typedef struct {} S;`, the unnamed struct type is decl-scoped to the alias declaration `S`.

§ 14.12.2 Namespaces

The `sort_key(namespace-name)` is `(namespace, identifier)`.

This means that namespaces are ordered alphabetically by comparing namespace names at the same rank. A namespace comes before any of its subnamespaces.

Example:

```
namespace outer1 {  
    struct i;  
}  
  
namespace outer2 {  
    namespace inner1 {  
        struct i;  
    }  
    namespace inner2 {  
        struct i;  
    }  
}
```

The order of the three structs w/ type `i` types shall be

```
sort_key(outer1::i)          <          sort_key(outer2::inner1::i)          <  
sort_key(outer2::inner2::i).
```

§ 14.12.3 Types

The `sort_key` of a type is `(type, <identifier>, <qualifiers>)`.

The `<identifier>` bit is a bit complicated, so let's deal with the qualifiers first.

Note: any decl-scopes the type is declared in are part of the parent *key-tuple*. The identifier portion is complicated because of possible template arguments for types that are template specializations.

§ 14.12.4 Ordering Class Types

§ 14.12.4.1 Ordering Simple Class Types

Class types shall be greater than scalar types.

Since we cannot redeclare two types with the same name, class types shall be ordered alphabetically.

```
struct Apple {};  
class Banana {};
```

```
struct Carrot {};
```

Would be ordered as Apple < Banana < Carrot

As such, we define sort key as:

```
sort_key(Apple) = (type, Apple, )
```

```
sort_key(Banana) = (type, Banana, )
```

```
sort_key(Carrot) = (type, Carrot, )
```

§ 14.12.4.2 Non Type Template Parameters

NTTPs shall first be ordered by their type, then their value.

Given:

```
template <auto T>
struct s {
    decltype(T) i = T;
};

s<1u> a;
s<1.0f> b;
```

```
sort_key(s<1u>) = ((type, (s, sort_key(1u))))
```

We can define sort_key of 1u as: `sort_key(1u) = ( sort_key(decltype(1u)), 1)`

s<1u> shall be ordered before s<1.0f>, as integral types come before floating point types.

NTTPs of the same type shall be lexicographically ordered by their scalar subobjects.
Meaning

```
struct F final {
    struct G final {
        int h;
        int i;
    } g;
    int j;
};

F f{{0,1}, 2};
F f2{{1,2}, 3};
```

```
sort_key(s<f>) < sort_key(s<f2>);
```

NTTPs of the same pointer or reference type shall be ordered by instantiation order.

§ 14.12.4.3 Ordering Class Template Specializations

Class templates shall be ordered by:

1. Class name, alphabetically.
2. Template arguments, applied lexicographically.

For example, given:

```
template <typename T, typename U>
struct Apple;

struct Banana;
struct Carrot;

Apple<Banana, Carrot>;
Apple<Banana, Banana>;
Apple<Carrot, Carrot>;
```

Note, `sort_key(<parameter>)...` will be used to denote a tuple where `sort_key` has been applied to all parameters.

For `void f(Foo, Bar)` `sort_key(<parameter>)...` would mean `(sort_key(Foo), sort_key(Bar))`

`sort_key` of a class template shall be defined as:

`sort_key(<class template>) = (type, (<name>, (sort_key(<parameter>)...)))`

So

`sort_key(Apple<Banana, Carrot>) = (type, (Apple, (sort_key(Banana), sort_key(Carrot))), )`

`sort_key(Apple<Banana, Carrot>) = (type, (Apple, ((type, Banana, ), (type, Carrot, )), )`

Note: the empty bit after the identifier is the empty qualifier pack.

The above would be ordered `sort_key(Apple<Banana, Banana>), sort_key(Apple<Banana, Carrot>), sort_key(Apple<Carrot, Carrot>).`

§ 14.12.4.4 Function Types

Function types shall be ordered by

1. Parameters, lexicographically.
2. Return type, if known from the declaration.

The `sort_key` of a function shall be defined as:

`sort_key(<function>) = (function, <name>, (sort_key(<parameter>)...), sort_key(<return type>))`

```
void foo(int i);
```

This function can be represented by: (function, (foo, (type, void), ((type, int))))

```
void foo(int)
void foo(int, double)
```

sort_key(void foo(int)) = (function, foo, (type, void), ((type, int)))

sort_key(void foo(int, double)) = (function, foo, (type, void), ((type, int), (type, double)))

So, the type of void foo(int) would precede the type of void foo(int, double)

§ 14.12.4.5 Member Function Types

Function types shall be ordered by

1. Return type
2. The type of the class it is a member of.
3. Parameters, lexicographically.

The sort key of a member function shall be defined as:

sort_key(<member function>) =

(function, (<name>, sort_key(<class>)), sort_key(<return type>), (sort_key(<parameter>)...)))

```
struct Foo {
    void bar(int i, float j);
};
```

sort_key(Foo::bar) =

(type, Foo,), (function, (bar, (type, Foo,)), (type, void), ((type, int,), (type, float,))))

§ 14.12.4.6 Variadic Function Types

Variadic function shall be ordered in a similar way. In a variadic function, the last argument is a variadic argument. A variadic argument shall be ordered immediately after its underlying type.

Given:

```
void foo(Foo);
void foo(Foo...);
```

In this case, the type of `void foo(Foo...)` is ordered immediately after the type of `void foo(Foo)`.

We can represent these as:

```
(function (type, void) (type, Foo, ))
```

```
(function (type, void) (type, Foo, ...))
```

§ 14.12.4.7 Parameter Packs

Parameter are ordered as class templates.

Given:

```
template<class... Types>
struct Tuple {};

class Foo {};
class Bar {};

Tuple<> t0;
Tuple<int> t1;
Tuple<Foo> t2;
Tuple<Bar> t3;
Tuple<Foo, Bar> t4;
```

would be ordered: `Tuple<>` < `Tuple<int>` < `Tuple<Bar>` < `Tuple<Foo>` < `Tuple<Foo, Bar>`

§ 14.12.4.8 Ordering Class Templates

Kinds of templates are ordered first by name, then by template arguments.

Given:

```
template <template <template<typename> class> class Template>
struct two{};

template <template <typename> class> struct one{};

template <typename> struct zero{};

zero<int> value0;
one<zero> value1;
two<one> value2;
```

These are represented by tuples:

```
sort_key(zero<int>) = (type, (zero, (type, int)))
```

```
sort_key(one<zero>) = (type, (one, (class_template, zero)))
```

```
sort_key(two<one>) = (type, (two, (class_template, one)))
```

§ 14.12.4.9 Variable Templates

Variable templates are ordered by name, then by template parameter.

```
sort_key(<variable_template>) = (variable_template, (<name>, (sort_key(<template_parameter>)...)))
```

```
template <typename F, typename S>
constexpr std::pair<F, S> pair_one_two = {1, 2};
```

the type of `pair_one_two<int, double>` can be represented as:

```
sort_key(pair_one_two<int, double>) = (variable_template, (pair_one_two, (type, int), (type, double)))
```

§ 14.12.4.10 Alias Templates

Alias templates are ordered alphabetically by name.

```
sort_key(<alias_template>) = (alias_template, <name>)
```

Given

```
template< class T >
using remove_cvref_t = typename remove_cvref<T>::type;
```

```
sort_key(remove_cvref_t) = (alias_template, remove_cvref_t)
```

§ 14.12.4.11 Concepts

Concepts are ordered in a similar manner to variable templates.

```
sort_key(<concept>) = (concept, (<name>, (sort_key(<template_parameter>)...)))
```

```
template <typename T, typename F = decltype([](T){})>
concept f = requires (T i, F f = [](T){}) {
    {f(i)} -> std::convertible_to<void>;
};
```

In order to order the type of the lambda declared in concept `f`, concept `f` must be comparable with other types.

Concepts shall be ordered first by name, then by template arguments.

```
sort_key(f<int>) = (concept, (f, (type, int), (lambda, 0)))
```

§ 15 References

[P1985R3] Gašper Ažman, Mateusz Pusz, Colin MacLean, Bengt Gustafsonn, Corentin Jabot. 2022-09-17. Universal template parameters.

<https://wg21.link/p1985r3>

[P2300R7] Eric Niebler, Michał Dominiak, Georgy Evtushenko, Lewis Baker, Lucian Radu Teodorescu, Lee Howes, Kirk Shoop, Michael Garland, Bryce Adelstein Lelbach. 2023-04-21. ``std::execution``.

<https://wg21.link/p2300r7>

[P2320R0] Andrew Sutton, Wyatt Childers, Daveed Vandevoorde. 2021-02-15. The Syntax of Static Reflection.

<https://wg21.link/p2320r0>

[P2841R1] Corentin Jabot, Gašper Ažman. 2023-10-14. Concept Template Parameters.

<https://wg21.link/p2841r1>