

aligned_accessor: An mdspan accessor expressing pointer overalignment

Document #: P2897R2
Date: 2024-07-15
Project: Programming Language C++ LEWG
Reply-to: Mark Hoemmen
<mhoemmen@nvidia.com>
Damien Lebrun-Grandie
<lebrungrandt@ornl.gov>
Nicolas Morales
<nmmoral@sandia.gov>
Christian Trott
<crtrott@sandia.gov>

Contents

- 1 Authors
- 2 Revision history
- 3 Purpose of this paper
- 4 Key features
- 5 Design discussion
 - 5.1 The accessor is not nestable
 - 5.2 Explicit constructor from `default_accessor`
 - 5.3 Why no explicit constructor from less to more alignment?
 - 5.4 We do not define an alias for aligned mdspan
 - 5.5 mdspan construction safety
 - 5.6 `is_sufficiently_aligned` is not `constexpr`
 - 5.7 Generalize `is_sufficiently_aligned` for all accessors?
 - 5.7.1 `is_sufficiently_aligned` is specific to `aligned_accessor`
 - 5.7.2 `detectably_invalid`: Generic validity check?
 - 5.7.3 Arguments against and for `detectably_invalid`
 - 5.7.4 Nonmember customization point design

- 5.7.5 We need both `is_sufficiently_aligned` and `detectably_invalid`
- 5.7.6 Alternative: Nonmember `is_sufficiently_aligned`
- 5.7.7 Do accessors need to check anything else?
- 5.7.8 Naming the function
- 5.7.9 Conclusions
- 5.8 Explicit conversions as the model for precondition-asserting conversions
 - 5.8.1 Example: conversion to `aligned_accessor`
 - 5.8.2 Status quo
 - 5.8.3 `mdspan` uses explicit conversions to assert preconditions
 - 5.8.4 Alternative: explicit cast function `naughty_cast`
 - 5.8.5 Conclusion: retain `mdspan`'s current design
- 6 Implementation
- 7 Example
- 8 References
- 9 Acknowledgments
- 10 Wording
 - 10.1 Add subsection  `[mdspan.accessor.aligned]` with the following
 - 10.2 Members `[mdspan.accessor.aligned.members]`
- 11 Wording for alternative nonmember `is_sufficiently_aligned` design
- 12 Appendix A: `detectably_invalid` nonmember function example
- 13 Appendix B: Implementation and demo

§ 1 Authors

- Mark Hoemmen (mhoemmen@nvidia.com) (NVIDIA)
- Damien Lebrun-Grandie (lebrungrandt@ornl.gov) (Oak Ridge National Laboratory)
- Nicolas Morales (nmmoral@sandia.gov) (Sandia National Laboratories)
- Christian Trott (crtrott@sandia.gov) (Sandia National Laboratories)

§ 2 Revision history

- Revision 0 (pre-Varna) to be submitted 2023-05-19
- Revision 1 (pre-Kona) to be submitted 2023-10-15

- Implement changes requested by LEWG review on 2023-10-10
 - Change `gcd` converting constructor Constraint to a Mandate
 - Add Example in the wording section that uses `is_sufficiently_aligned` to check the pointer overalignment precondition
 - Add Example in the wording section that uses `aligned_alloc` to create an overaligned allocation, to show that `aligned_accessor` exists as part of a system
 - Add an explicit constructor from `default_accessor`, so that users can type `aligned_mdspan y{x}` instead of `aligned_mdspan y{x.data_handle(), x.mapping() }`. Add an explanation in the design discussion section.
- Implement other wording changes
 - Add to `aligned_accessor`'s Mandates that `byte_alignment >= alignof(ElementType)` is true. This prevents construction of an invalid `aligned_accessor` object.
- Add more design discussion based on LEWG review on 2023-10-10
 - Explain why we do not include an `aligned_mdspan` alias
 - Explain `aligned_accessor` construction safety
- Revision 2 (post - St. Louis) to be submitted 2024-07-15
 - Implement required changes from LEWG review of R1 on 2024-06-28
 - Remove `constexpr` from `is_sufficiently_aligned`
 - Discuss optional suggestions from LEWG review of R1 on 2024-06-28
 - Add explicit converting constructor vs. named cast (“naughty_cast”) discussion
 - Add `detectably_invalid` discussion
 - Ask LEWG to consider the alternative design that makes `is_sufficiently_aligned` a nonmember function in `<bit>` instead of a member function of `aligned_accessor`, while LWG review of R2 proceeds concurrently
- P2389R2 was voted into the Working Draft at St. Louis, so replace use of `dextents` in examples with `dims`.
- Add non-wording section explaining why `aligned_accessor` has no explicit constructor from less to more alignment

- Add Compiler Explorer link with full implementation and demo
- Revision 3 (post - St. Louis) to be submitted 2024-07-15
 - Include updated feedback from David Sankel (see Acknowledgments) after his review of R2

§ 3 Purpose of this paper

We propose adding `aligned_accessor` to the C++ Standard Library. This class template is an `mdspan` accessor policy that uses `assume_aligned` to decorate pointer access. We think it belongs in the Standard Library for two reasons. First, it would serve as a common vocabulary type for interfaces that take `mdspan` to declare their minimum alignment requirements. Second, it extends to `mdspan` accesses the optimizations that compilers can perform to pointers decorated with `assume_aligned`.

`aligned_accessor` is analogous to the various `atomic_accessor_*` templates proposed by P2689. Both that proposal and this one start with a Standard Library feature that operates on a “raw” pointer (`assume_aligned` or the various `atomic_ref*` templates), and then propose an `mdspan` accessor policy that straightforwardly wraps the lower-level feature.

We had originally written `aligned_accessor` as an example in P2642, which proposes “padded” `mdspan` layouts. We realized that `aligned_accessor` was more generally applicable and that standardization would help the padded layouts proposed by P2642 reach their maximum value.

§ 4 Key features

- `offset_policy` is `default_accessor`
- `data_handle_type` is `ElementType*`
- Permitted implicit conversions
 - from `nonconst` to `const ElementType`,
 - from more overalignment to less overalignment, and
 - from overalignment to no overalignment (`default_accessor`)
- explicit converting constructor from `default_accessor` lets users assert overalignment
- `is_sufficiently_aligned` checks pointer alignment

The `offset_policy` alias is `default_accessor<ElementType>`, because even if a pointer `p` is aligned, `p + i` might not be.

The `data_handle_type` alias is `ElementType*`. It needs no further adornment, because alignment is asserted at the point of access, namely in the access function. Some implementations might have an easier time optimizing if they also apply some implementation-specific attribute to `data_handle_type` itself. Examples of such attributes include `__declspec(align_value(byte_alignment))` and `__attribute__((align_value(byte_alignment)))`. However, these attributes should not apply to the result of `offset`, for the same reason that `offset_policy` is `default_accessor` and not `aligned_accessor`.

The converting constructor from `aligned_accessor` is analogous to `default_accessor`'s constructor, in that it exists to permit conversion from `nonconst element_type` to `const element_type`. It additionally permits implicit conversion from more overalignment to less overalignment – something that we expect users may need to do. For example, users may start with `aligned_accessor<float, 128>`, because their allocation function promises 128-byte alignment. However, they may then need to call a function that takes an `mdspan` with `aligned_accessor<float, 32>`, which declares the function's intent to use 8-wide SIMD of `float`.

The `is_sufficiently_aligned` function checks whether a pointer has sufficient alignment to be used correctly with the class. This makes it easier for users to check preconditions, without needing to know how to cast a pointer to an integer of the correct size and signedness.

Per discussion below, we would like LEWG to consider the alternative design that removes `is_sufficiently_aligned` from `aligned_accessor` and adds it to the `<bit>` header as a separate nonmember function. We think LWG review of `aligned_accessor` can proceed concurrently. We present wording for this design alternative below.

§ 5 Design discussion

§ 5.1 The accessor is not nestable

We considered making `aligned_accessor` “wrap” any accessor type that meets the right requirements. For example, `aligned_accessor` could take the inner accessor as a template parameter, store an instance of it, and dispatch to its member functions. That would give users a way to apply multiple accessor “attributes” to their data handle, such as atomic access (see P2689) and overalignment.

We decided against this approach for three reasons. First, we would have no way to validate that the user's accessor type has the correct behavior. We could check that their accessor's `data_handle_type` is a pointer type, but we could not check that their accessor's access function actually dereferences the pointer. For instance, `access` might instead interpret the pointer as a file handle or a key into a distributed data store.

Second, even if the inner accessor's access function actually did return the result of dereferencing the pointer, the outer access function might not be able to recover the

effects of the inner access function, because `access` computes a reference, not a pointer. In order for `aligned_accessor`'s access function to get back that pointer, it would need to reach past the inner accessor's public interface. That would defeat the purpose of generic nesting.

Third, any way (not just this one) of nesting two generic accessors raises the question of order dependence. Even if it were possible to apply the effects of both the inner and outer accessors' access functions in sequence, it might be unpleasantly surprising to users if the effects depended on the order of nesting. A similar question came up in the “properties” proposal P0900, which we quote here.

Practically speaking, it would be considered a best practice of a high-quality implementation to ensure that a property's implementation of `properties::element_type_t` (and other traits) are invariant with respect to ordering with other known properties (such as those in the standard library), but with this approach it would be impossible to make that guarantee formal, particularly with respect to other vendor-defined and user-defined properties unknown to the property implementer.

For these reasons, we have made `aligned_accessor` stand-alone, instead of having it modify another user-provided accessor.

§ 5.2 Explicit constructor from `default_accessor`

LEWG's 2023-10-10 review of R0 pointed out that in R0, `aligned_accessor` lacks an explicit constructor from `default_accessor`. Having that constructor would make it easier for users to create an aligned `mdspan` from an unaligned `mdspan`. Making it explicit would prevent implicit conversion. Thus, we have decided to add this explicit constructor in R1.

Without the explicit constructor, users have two options for turning a nonaligned `mdspan` into an aligned `mdspan`. First, as in the following example, users could “take apart” the input nonaligned `mdspan` and use the pieces to construct an aligned `mdspan`, whose type they name completely.

```
void compute_with_aligned(
    std::mdspan<float, std::dims<2>, std::layout_left> matrix)
{
    const std::size_t byte_alignment = 4 * alignof(float);
    using aligned_matrix_t = std::mdspan<float, std::dims<2>,
        std::layout_left, std::aligned_accessor<float,
byte_alignment>>>;

    aligned_matrix_t aligned_matrix{matrix.data_handle(),
matrix.mapping()};
    // ... use aligned_matrix ...
}
```

Second, as in the following example, users could construct an `aligned_accessor` explicitly and use constructor template argument deduction (CTAD) to construct the aligned `mdspan` from its pieces.

```
void compute_with_aligned(  
    std::mdspan<float, std::dims<2>, std::layout_left> matrix)  
{  
    const std::size_t byte_alignment = 4 * alignof(float);  
  
    std::mdspan aligned_matrix{matrix.data_handle(),  
matrix.mapping(),  
    std::aligned_accessor<float, byte_alignment>{}};  
    // ... use aligned_matrix ...  
}
```

The first approach would likely be more common. This is because `mdspan` users commonly define their own type aliases for `mdspan`, with application-specific names that make code more self-documenting. The `aligned_matrix_t` definition above is an example.

Adding an explicit constructor from `default_accessor` lets users get the same effect more concisely, without needing to “take apart” the input `mdspan`.

```
void compute_with_aligned(std::mdspan<float, std::dims<2, int>,  
std::layout_left> matrix)  
{  
    const std::size_t byte_alignment = 4 * alignof(float);  
    using aligned_mdspan = std::mdspan<float, std::dims<2, int>,  
        std::layout_left, std::aligned_accessor<float,  
byte_alignment>>;  
  
    aligned_mdspan aligned_matrix{matrix};  
    // ... use aligned_matrix ...  
}
```

The explicit constructor does not decrease safety, in the sense that users were always allowed to convert from an `mdspan` with `default_accessor` to an `mdspan` with `aligned_accessor`. Before, users could perform this conversion by typing the following.

```
aligned_matrix_t aligned_matrix{matrix.data_handle(),  
matrix.mapping()};
```

Now, users can do the same thing with fewer characters.

```
aligned_matrix_t aligned_matrix{matrix};
```

§ 5.3 Why no explicit constructor from less to more alignment?

As explained in the previous section, `aligned_accessor` has an explicit converting constructor from `default_accessor` so that users can assert overalignment. It also has an

(implicit) converting constructor from another `aligned_accessor` with more alignment, to an `aligned_accessor` with less alignment. However, `aligned_accessor` does *not* have an explicit converting constructor from another `aligned_accessor` with *less* alignment, to an `aligned_accessor` with *more* alignment. Why not?

Consider the three typical use cases for `aligned_accessor`.

1. User knows an allocation's alignment at compile time.
2. User knows an allocation's alignment at run time, but not at compile time. For example, the value might depend on run-time detection of particular hardware features.
3. User doesn't know whether an allocation is overaligned. They might need to ask some system at run time, or check the pointer value themselves, in order to decide whether to call code that expects a particular alignment.

In Case (1), users would normally declare the maximum alignment. They would want to preserve this information at compile time as much as possible, by keeping the `aligned_accessor` `mdspan` with maximum compile-time alignment for the entire scope of its use. Users would only want implicit conversions to less alignment or `default_accessor` when calling functions whose parameter types encode these requirements.

Case (2) reduces to Case (3).

Case (3) reduces to Case (1). This works like any conversion from run-time type to compile-time type, with a fixed list of possible compile-time types (the alignments). As soon as a user's `mdspan` enters a scope where the alignment is known at compile time, the user would want to preserve that compile-time information and maximize the alignment for as large of a scope as possible.

None of these cases involve starting with more alignment, going to less (but still some) alignment, and then going back to more alignment again. Code that does that probably does not correctly use the types of function parameters to express its overalignment requirements. It's like code that uses `dynamic_cast` a lot. Users can still convert from less or more alignment by creating the result's `aligned_accessor` manually. However, we don't want to encourage this pattern, so we don't offer an explicit conversion for it.

§ 5.4 We do not define an alias for `aligned_mdspan`

In LEWG's 2023-10-10 review of R0, participants observed that this proposal's examples define an example-specific type alias for `mdspan` with `aligned_accessor`. They asked whether our proposal should include a *standard* alias `aligned_mdspan`. We do not *object* to such an alias, but we do not find it very useful, for the following reasons.

1. Users of `mdspan` commonly define their own type aliases whose names are meaningful for their applications.
2. It would not save much typing.

Examples may define aliases to make them more concise. One example in this proposal defines the following alias for an `mdspan` of `float` with alignment `byte_alignment`.

```
template<size_t byte_alignment>
using aligned_mdspan = std::mdspan<float, std::dims<1, int>,
std::layout_right, std::aligned_accessor<float,
byte_alignment>>;
```

This lets the example use `aligned_mdspan<32>` and `aligned_mdspan<16>`.

The above alias is specific to a particular example. A *general* version of alias would look like this.

```
template<class ElementType, class Extents, class Layout,
size_t byte_alignment>
using aligned_mdspan = std::mdspan<ElementType, Extents, Layout,
std::aligned_accessor<ElementType, byte_alignment>>;
```

This alias would save *some* typing. However, `mdspan` “power users” rarely type out all the template arguments. First, they can rely on CTAD to create `mdspans`, and `auto` to return them. Second, users commonly already define their own aliases whose names have an application-specific meaning. They define these aliases *once* and use them throughout the application. For instance, users might define the following.

```
template<class ElementType>
using vector_t = std::mdspan<ElementType,
std::dims<1>, std::layout_left>;
template<class ElementType>
using matrix_t = std::mdspan<ElementType,
std::dims<2>, std::layout_left>;

template<class ElementType, size_t byte_alignment>
using aligned_vector_t = std::mdspan<ElementType,
std::dims<1>, std::layout_left,
std::aligned_accessor<ElementType, byte_alignment>>;
template<class ElementType, size_t byte_alignment>
using aligned_matrix_t = std::mdspan<ElementType,
std::dims<2>, std::layout_left,
std::aligned_accessor<ElementType, byte_alignment>>;
```

Such users may never type the characters “`mdspan`” again. For this reason, while we do not object to an `aligned_mdspan` alias, we do not find the proliferation of aliases particularly ergonomic.

§ 5.5 `mdspan` construction safety

LEWG's 2023-10-10 review of R0 expressed concern that `mdspan`'s constructor has no way to check `aligned_accessor`'s alignment requirements. Users can call `aligned_accessor`'s `is_sufficiently_aligned(p)` static member function with a pointer `p` to check this themselves, before constructing the `mdspan`. However, `mdspan`'s constructor generally has no way to check whether its accessor finds the caller's data handle acceptable.

This is true for any accessor type, not just for `aligned_accessor`. It is a design feature of `mdspan` that accessors can be stateless. Most of them have no state. Even if they have state, they do not need to be constructed with or store the data handle. As a result, an accessor might not see a data handle until `access` or `offset` is called. Both of those member functions are performance critical, so they cannot afford an extra branch on every call. Compare to `vector::operator[]`, which has preconditions but is not required to perform bounds checks. Using exceptions in the manner of `vector::at` could reduce performance and would also make `mdspan` unusable in a freestanding or no-exceptions context.

Note that `aligned_accessor` does not introduce *additional* preconditions beyond those of the existing C++ Standard Library feature `assume_aligned`. In the words of one LEWG reviewer, `aligned_accessor` is not any more “pointy” than `assume_aligned`; it just passes the point through without “blunting” it.

Before submitting R0 of this paper, we considered an approach specific to `aligned_accessor`, that would wrap the pointer in a special data handle type. The data handle type would have an explicit constructor that takes a raw pointer, with a precondition that the raw pointer have sufficient alignment. The constructor would be explicit, because it would have a precondition. This design would force the precondition back to `mdspan` construction time. Users would have to construct the `mdspan` like this.

```
element_type* raw_pointer = get_pointer_from_somewhere();
using acc_type = aligned_accessor<element_type, byte_alignment>;
mdspan x{acc_type::data_handle_type{raw_pointer}, mapping,
acc_type{}};
```

We rejected this approach in favor of `is_sufficiently_aligned` for the following reasons.

1. Wrapping the pointer in a custom data handle class would make every access or offset call need to reach through the data handle's interface, instead of just taking the raw pointer directly. The access function, and to some extent also offset, need to be as fast as possible. Their performance depends on compilers being able to optimize through function calls. The authors of `mdspan` carefully balanced generality with function call depth and other code complexity factors that may hinder compilers from optimizing. Performance of `aligned_accessor` matters as much or even more than performance of `default_accessor`, because `aligned_accessor` exists to communicate optimization potential.
2. The alignment precondition would still exist. Requiring the data handle type to throw an exception if the pointer is not sufficiently aligned would make `mdspan`

unusable in a freestanding or no-exceptions context.

3. Users should not have to pay for unneeded checks. The two examples in the wording express the two most common cases. If users get a pointer from a function like `aligned_alloc`, then they already know its alignment, because they asked for it. If users are computing alignment at run time to dispatch to a more optimized code path, then they know alignment before dispatch. In both cases, users already know the alignment before constructing the `mdspan`.
4. The data handle is still a pointer, it's just a pointer with a constraint on its values. Users would reasonably expect to be able to use the result of `data_handle()` with existing interfaces that expect a raw pointer.

An LEWG poll on 2023-10-10, “[b]lock `aligned_accessor` progressing until we have a way of checking alignment requirements during `mdspan` construction,” resulted in no consensus. Attendance was 14.

Strongly Favor	Weakly Favor	Neutral	Weakly Against	Strongly Against
0	1	1	2	2

LEWG expressed an (unpolled) interest that we explore `mdspan` safety in subsequent work after the fall 2023 Kona WG21 meeting. LEWG asked us to explore safety in a way that is not specific to `aligned_accessor`. Part of that exploration is in the section below “Generalize `is_sufficiently_aligned` for all accessors?”. We plan further exploration of this topic elsewhere.

§ 5.6 `is_sufficiently_aligned` is not `constexpr`

LEWG reviewed R1 of this proposal at the June 2024 St. Louis WG21 meeting, and polled 1/10/0/0/1 (SF/F/N/A/SA) to remove `constexpr` from `is_sufficiently_aligned`. This is because it is not clear how to implement the function in a way that could ever be a constant expression. The straightforward cross-platform way to implement this would `bit_cast` the pointer to `uintptr_t`. However, `bit_cast` is not `constexpr` when converting from a pointer to an integer, per [\[bit.cast\] 3](#). Any `reinterpret_cast` similarly could not be a core constant expression, per [\[expr.const\] 5.15](#). One LEWG reviewer pointed out that some compilers have a built-in operation (e.g., Clang and GCC have `__builtin_bit_cast`) that might form a constant expression when `bit_cast` does not. On the other hand, the authors could not foresee a need for `is_sufficiently_aligned` to be `constexpr` and did not want to constrain implementations to use compiler-specific functionality.

§ 5.7 Generalize `is_sufficiently_aligned` for all accessors?

§ 5.7.1 `is_sufficiently_aligned` is specific to `aligned_accessor`

The `is_sufficiently_aligned` function exists so users can check the pointer's alignment precondition before constructing an `mdspan` with it. This precondition check is specific to `aligned_accessor`. Furthermore, the function has a precondition that the pointer points to a valid element. Standard C++ offers no way for users to check that. More importantly for `mdspan` users, Standard C++ offers no way to check whether a pointer and a layout mapping's `required_span_size()` form a valid range.

§ 5.7.2 `detectably_invalid`: Generic validity check?

During the June 2024 St. Louis WG21 meeting, one LEWG reviewer (please see Acknowledgments below) pointed out that code that is generic on the accessor type currently has no way to check whether a given data handle is valid. Specifically, given a `size_t` `size` (e.g., the `required_span_size()` of a given layout mapping), there is no way to check whether `[ 0, size )` forms an accessible range (see [\[mdspan.accessor.general\] 2](#)) of a given data handle and accessor. The reviewer suggested adding a new member function

```
bool detectably_invalid(data_handle_type handle, size_t size)
const noexcept;
```

to all `mdspan` accessors. This would return `true` if the implementation can show that `[ 0, size )` is *not* an accessible range for `handle` and the accessor, and `true` otherwise. The word “detectably” in the name would remind users that this is a “best effort” check. It might return `false` even if the handle is invalid or if `[ 0, size )` is not an accessible range. Also, it might return different values on different implementations, depending on their ability to check e.g., pointer range validity. The function would have the following design features.

- It must be a non-static member function, because in general, accessors may have state that determines validity of the data handle.
- It must be `const` because precondition-checking code should avoid observable side effects.
- It must be `noexcept` because precondition-checking code should not throw.

With such a function, users could write generic checked `mdspan` creation code like the following.

```
template<class LayoutMapping, class Accessor>
auto create_mdspan_with_check(
    typename Accessor::data_handle_type handle,
    LayoutMapping mapping,
    Accessor accessor)
{
    if (accessor.detectably_invalid(handle,
mapping.required_span_size())) {
        throw std::out_of_range("Invalid data handle and/or size");
    }
}
```

```
    return mdspan{handle, mapping, accessor};  
}
```

§ 5.7.3 Arguments against and for `detectably_invalid`

We didn't include this feature in the original `mdspan` design because most data handle types have no way to say with full accuracy whether a handle and size are valid. We didn't want to give users the false impression that a validity check was doing anything meaningful. Standard C++ has no way to check a raw pointer `T*` and a size, though some implementations such as CHERI C++ ([Davis 2019] and [Watson 2020]) and run-time profiling and debugging systems such as Valgrind do have this feature. We designed `mdspan` accessors to be able to wrap libraries that implement a partitioned global address space (PGAS) programming model for accessing remote data over a network. (See [P0009R18, Section 2.7, “Why custom accessors?”](#).) Such libraries include the one-sided communication interface in MPI (the Message Passing Interface for distributed-memory parallel programming) or NVSHMEM (NVIDIA's implementation of the SHMEM standard). Those libraries define their own data handle to represent remote data. For example, MPI uses an `MPI_Win` “window” object. NVSHMEM uses a C++ pointer to represent a “symmetric address” that points to an allocation from the “symmetric heap” (that is accessible to all participating parallel processes). Such libraries generally do not have validity checks for their handles.

On the other hand, a `detectably_invalid` function would let happen any checks that *could* happen. For instance, a hypothetical “GPU device memory accessor” (not proposed for the C++ Standard, but existing in projects like [RAPIDS RAFT](#)) might permit access to an allocation of GPU “device” memory from only GPU “device” code, not from ordinary “host” code. A common use case for GPU allocations is to allocate device memory in host code, then pass the pointer to device code for use there. Thus, it would be reasonable to create an `mdspan` in host code with that accessor. The accessor could use a CUDA run-time function like `cudaPointerGetAttributes` to check if the pointer points to valid GPU memory. Even `default_accessor` could have a simple check like this.

```
bool detectably_invalid(data_handle_type ptr, size_t size) const {  
    return ptr == nullptr && size != 0;  
}
```

§ 5.7.4 Nonmember customization point design

C++23 defines the generic interface of accessors through the accessor policy requirements [\[mdspan.accessor.reqmts\]](#). Adding `detectably_invalid` to these requirements would be a breaking change to C++23. Thus, generic code that wanted to call this function would need to fill in default behavior for both Standard accessors defined in C++23, and user-defined accessors that comply with the C++23 accessor requirements. The following `detectably_invalid` nonmember function (not proposed in this paper) shows one way users could do that. Please see Appendix A below for the full source code of a demonstration, along with a Compiler Explorer link. This demonstration shows that breaking backwards compatibility with C++23 is unnecessary, because users can

straightforwardly work around the lack of a `detectably_invalid` member function in C++23 - compliant accessors. Not standardizing this nonmember function work-around would also give users the freedom to fill in different default behavior. For example, some users may prefer to consider every (data handle, size) pair invalid unless proven otherwise, as a way to force use of custom accessors that have the ability to make accurate checks.

```
template<class Accessor>
concept has_detectably_invalid = requires(Accessor acc) {
    typename Accessor::data_handle_type;
    { std::as_const(acc).detectably_invalid(
        std::declval<typename Accessor::data_handle_type>(),
        std::declval<std::size_t>()
    ) } noexcept -> std::same_as<bool>;
};

template<class Accessor>
bool detectably_invalid(Accessor&& accessor,
    typename std::remove_cvref_t<Accessor>::data_handle_type handle,
    std::size_t size)
{
    if constexpr
(has_detectably_invalid<std::remove_cvref_t<Accessor>>) {
        return std::as_const(accessor).detectably_invalid(handle,
size);
    }
    else {
        return false;
    }
}
```

§ 5.7.5 We need both `is_sufficiently_aligned` and `detectably_invalid`

One could argue that if `aligned_accessor` had `detectably_invalid`, that would make `is_sufficiently_aligned` unnecessary. We disagree; we think `is_sufficiently_aligned` is useful by itself, whether or not `detectably_invalid` exists, for the following reasons.

1. Users will often want to check alignment separately from pointer range validity.
2. Checking alignment may be much less expensive than checking pointer range validity.

Regarding (1), we think the most common use case for `aligned_accessor`'s explicit converting constructor from `default_accessor` would be explicit construction of an `mdspan` with `aligned_accessor` from an `mdspan` with `default_accessor`. The latter exists, so the user has already asserted that the range formed by its data handle and `required_span_size()` is valid. Thus, the only thing the user would need to check would be whether the data handle is sufficiently aligned.

The same LEWG reviewer who suggested `detectably_invalid` had originally thought it would make `is_sufficiently_aligned` unnecessary. However, after reviewing R2 of

this paper, that reviewer changed their mind. They now agree with us that `is_sufficiently_aligned` is useful by itself. All their concerns would be addressed by making `is_sufficiently_aligned` a nonmember function, rather than a member function of `aligned_accessor`.

§ 5.7.6 Alternative: Nonmember `is_sufficiently_aligned`

The reviewer responded to our argument above by suggesting that we remove `is_sufficiently_aligned` from `aligned_accessor` and make it a separate nonmember function. This is a reasonable alternative suggestion. We think that LWG review of `aligned_accessor` can proceed concurrently while LEWG considers this alternative.

Into which header should this new function go? We think it belongs in `<bit>`, because it is fundamentally a bit arithmetic operation. There should be no reason why that could not be freestanding, like everything else in the `<bit>` header. Another option would be to put it in `<memory>` right after `assume_aligned`, and to mark it as freestanding (since `assume_aligned` is).

§ 5.7.7 Do accessors need to check anything else?

The only other thing an accessor’s user might want to check besides a (data handle, size) pair would be converting construction from another type of accessor. All `mdspan` components – extents, layout mappings, and accessors – implement conversions with preconditions via explicit constructors. (For more detail, please see the section below, “Explicit conversions as the model for precondition-asserting conversions.”) Accessors do *not* store their data handles, so the only reason to check whether converting construction is valid would be if the input or result accessor has separate run-time state. (Otherwise, the check could be a constraint or `static_assert`.) It’s rare for an accessor to need run-time state, so we don’t expect to need this feature in generic code. It would also be a separable addition from the feature of checking a data handle and size. Nevertheless, one could consider a design. We would favor just overloading `detectably_invalid` for accessors, as there would be no risk of ambiguity. Converting constructors only take one argument, so there would be no ambiguity between calling `detectably_invalid` with an accessor and calling it with a data handle and size.

§ 5.7.8 Naming the function

1. The function describes a property: “this (data handle, size) pair is not known to be invalid.” It’s an adjective (like “valid” or “is_valid”), not a verb (like “check” as in “check_valid”).
2. The function does not promise perfect accuracy. In the common case, it says whether it can *detect* whether the handle and size are `_in_valid`. Whether or not they are *valid* might be harder to say.
3. As discussed above, users may also want to check converting constructors from other accessor types. However, there would be no risk of ambiguity between that

and checking a data handle and size. Therefore, there's no need for the function's name to include the type of the thing being checked (e.g., "range").

4. Specifically, the function should not contain the word "pointer," because a data handle is not necessarily a pointer. Even if `data_handle_type` is a pointer type, a data handle might not necessarily be a pointer to the elements in the Standard C++ sense. For example, it might be some opaque handle that a library represents as a type alias of `void*`.

These points together suggest the name `detectably_invalid`.

§ 5.7.9 Conclusions

1. Adding `detectably_invalid` to the accessor requirements and existing Standard accessors in C++26 would be a breaking change to C++23. Nevertheless, users could write code that fills in reasonable behavior for C++23 accessors.
2. Few C++ implementations offer a way to check validity of a pointer range. Thus, users would experience `detectably_invalid` as mostly not useful for the common case of `default_accessor` and other accessors that access a pointer range.
3. Item (1) reduces the urgency of this suggestion for C++26. Item (2) reduces its potential to improve the `mdspan` user experience in a practical way. Therefore, we do not suggest adding `detectably_invalid` to the accessor requirements in this proposal. However, we do not discourage further work in separate proposals.
4. We would like LEWG to consider the alternative design that removes `is_sufficiently_aligned` from `aligned_accessor` and adds it to the C++ Standard Library as a separate nonmember function. We think LWG review of `aligned_accessor` can proceed concurrently.

§ 5.8 Explicit conversions as the model for precondition-asserting conversions

During the June 2024 St. Louis WG21 meeting, one LEWG reviewer asked about the explicit constructor from `default_accessor`. This constructor lets users assert that a pointer has sufficient alignment to be accessed by the `aligned_accessor`. The reviewer argued that this was an "unsafe" conversion, and wanted these "unsafe" conversions to be even more explicit than an explicit constructor: e.g., a new `*_cast` function template. We do not agree with this idea; this section explains why.

§ 5.8.1 Example: conversion to `aligned_accessor`

Suppose that some function that users can't change returns an `mdspan` of `float` with `default_accessor`, even though users know that the `mdspan` is overaligned to `8 * sizeof(float)` bytes. The function's parameter(s) don't matter for this example.


```
mdspan<float, dims<1>, layout_right, default_accessor<float>>
    overaligned_view(SomeParameters params);
```

Suppose also that users want to call some other function that they can't change. This function takes an `mdspan` of `float` with `aligned_accessor<float, 8>`. Its return type doesn't matter for this example.

```
SomeReturnType use_overaligned_view(
    mdspan<float, dims<1>, layout_right, aligned_accessor<float,
8>>);
```

§ 5.8.2 Status quo

How do users call `use_overaligned_view` with the object returned from `overaligned_view`? The status quo offers two ways. Both of them rely on `aligned_accessor<float, 8>`'s explicit converting constructor from `default_accessor<float>`.

1. Use `mdspan`'s explicit converting constructor.
2. Construct the new `mdspan` explicitly from its data handle, layout mapping, and accessor. (This is the ideal use case for CTAD, as an `mdspan` is nothing more than its data handle, layout mapping, and accessor.)

Way (1) looks like this.

```
auto x = overaligned_view(params);
auto result = use_overaligned_view(
    mdspan<float, dims<1>, layout_right,
    aligned_accessor<float, 8>>(x)
);
```

Way (2) looks like this. Note use of CTAD.

```
auto x = overaligned_view(params);
auto result = use_overaligned_view(
    mdspan{x.data_handle(), x.mapping(),
    aligned_accessor<float, 8>>(x.accessor())}
);
```

Which way is less verbose depends on `mdspan`'s template arguments. Both ways, though, force the user to name the type `aligned_accessor<float, 8>` explicitly. Users know that they have pulled out a sharp knife from the toolbox. It's verbose, it's nondefault, and it's a class with a short definition. Users can go to the specification, see `assume_aligned`, and know they are dealing with a low-level function that has a precondition.

§ 5.8.3 `mdspan` uses explicit conversions to assert preconditions

The entire system of `mdspan` components was designed so that

- conversions with preconditions happen through explicit conversions (mostly converting constructors); while
- conversions without preconditions happen through implicit conversions.

Changing this would break backwards compatibility with C++23. For example, one can see this with converting constructors for

- `extents` (for conversions from run-time to compile-time extents, or conversions from wider to narrower index type): [\[mdspan.extents.cons\]](#); and
- `layout_left::mapping`, and all the other layout mappings currently in the Standard that are not `layout_stride` or `layout_transpose` (for conversions from e.g., `layout_stride::mapping`, which assert that the strides are compatible): e.g., [\[mdspan.layout.left.cons\]](#).

This is consistent with C++ Standard Library class templates, in that construction asserts any preconditions. For example, if users construct a `string_view` or `span` from a pointer `ptr` and a size `size`, this asserts that the range `[ ptr, ptr + size )` is accessible.

§ 5.8.4 Alternative: explicit cast function `naughty_cast`

Everything we have described above is the status quo. What did the one LEWG reviewer want to see? They wanted all conversions with preconditions to use a “cast” function with an easily searchable name, analogous to `static_cast`. As a placeholder, we’ll call it “`naughty_cast`.” For the above `use_overaligned_view` example, the `naughty_cast` analog of Way (2) would look like this.

```
auto x = overaligned_view(params);
auto result = use_overaligned_view(
    mdspan{x.data_handle(), x.mapping(),
          naughty_cast<aligned_accessor<float, 8>>>(x.accessor())}
);
```

One could imagine defining `naughty_cast` of `mdspan` by `naughty_cast` of its components. This would enable an analog of Way (1).

```
auto x = overaligned_view(params);
auto result = use_overaligned_view(naughty_cast<
    mdspan<float, dims<1>, layout_right,
    aligned_accessor<float, 8>>>(x)
);
```

Another argument for `naughty_cast` besides searchability is to make conversions with preconditions “loud,” that is, easily seen in the code by human developers. However, the original Way (1) and Way (2) both are loud already in that they require a lot of extra code that spells out the result’s accessor type explicitly. The status quo’s difference in “volume” is implicit conversion

```
auto result = use_overaligned_view(x);
```

versus explicit construction.

```
auto result = use_overaligned_view(  
    mdspan{x.data_handle(), x.mapping(),  
        aligned_accessor<float, 8>(x)}};  
);
```

Adding `naughty_cast` to the latter doesn't make it much louder.

```
auto result = use_overaligned_view(  
    mdspan{x.data_handle(), x.mapping(),  
        naughty_cast<aligned_accessor<float, 8>>(x)}};  
);
```

There are other disadvantages to a `naughty_cast` design. The point of that design would be to remove or make non-public all the explicit constructors from `mdspan`'s components. That functionality would need to move somewhere. A typical implementation technique for a custom cast function is to rely on specializations of a struct with two template parameters, one for the input type and one for the output type of the cast. The `naughty_caster` struct example below shows how one could do that.

```
template<class Output, class Input>  
struct naughty_caster {};  
  
template<class Output, class Input>  
Output naughty_cast(const Input& input) {  
    return naughty_caster<Output, Input>::cast(input);  
}  
  
template<class OutputElementType, size_t ByteAlignment,  
    class InputElementType>  
    requires (is_convertible_v<InputElementType(*)[],  
        OutputElementType(*)[>)  
struct naughty_caster {  
    using output_type =  
        aligned_accessor<OutputElementType, ByteAlignment>;  
    using input_type = default_accessor<InputElementType>;  
  
    static output_type cast(const input_type&) {  
        return {};  
    }  
};
```

This technique takes a lot of effort and code, when by far the common case is that `cast` has a trivial body. For any accessors with state, it would almost certainly call for breaks of encapsulation, like making the `naughty_caster` specialization a friend of the input and/or output.

We emphasize that users are meant to write custom accessors. The intended typical author of a custom accessor is a performance expert who is not necessarily a C++ expert. It takes quite a bit of C++ experience to learn how to use encapsulation-breaking techniques safely; the lazy approaches all just expose implementation details or defeat the “safety”

that `naughty_cast` is supposed to introduce. Given that the main motivation of `naughty_cast` is safety, we shouldn't make it harder for users to write safe code.

More importantly, `naughty_cast` would obfuscate accessors. The architects of `mdspan` meant accessors to have to have a small number of “moving parts” and to define all those parts in a single place. Contrast `default_accessor` with the contiguous iterator requirements, for instance. The `naughty_cast` design would force custom accessors (and custom layouts) to define their different parts in different places, rather than all in one class. WG21 has moved away from this scattered design approach. For example, [P2855R1](#) (“Member customization points for Senders and Receivers”) changes `P2300` (`std::execution`) to use member functions instead of `tag_invoke`-based customization points.

§ 5.8.5 Conclusion: retain `mdspan`'s current design

For all these reasons, we do not support replacing `mdspan`'s current “conversions with preconditions are explicit conversions” design with a cast function design.

§ 6 Implementation

We have tested an implementation of this proposal with the [reference `mdspan` implementation](#). Appendix B below lists the source code of a full implementation.

§ 7 Example

```
template<size_t byte_alignment>
using aligned_mdspan =
    std::mdspan<float, std::dims<1, int>, std::layout_right,
std::aligned_accessor<float, byte_alignment>>;

// Interfaces that require 32-byte alignment,
// because they want to do 8-wide SIMD of float.
extern void vectorized_axpy(aligned_mdspan<32> y, float alpha,
aligned_mdspan<32> x);
extern float vectorized_norm(aligned_mdspan<32> y);

// Interfaces that require 16-byte alignment,
// because they want to do 4-wide SIMD of float.
extern void fill_x(aligned_mdspan<16> x);
extern void fill_y(aligned_mdspan<16> y);

// Helper functions for making overaligned array allocations.

template<class ElementType>
struct delete_raw {
    void operator()(ElementType* p) const {
        std::free(p);
    }
};
```

```

    template<class ElementType>
    using allocation = std::unique_ptr<ElementType[],
delete_raw<ElementType>>;

    template<class ElementType, std::size_t byte_alignment>
    allocation<ElementType> allocate_raw(const std::size_t
num_elements)
    {
        const std::size_t num_bytes = num_elements *
sizeof(ElementType);
        void* ptr = std::aligned_alloc(byte_alignment, num_bytes);
        return {ptr, delete_raw<ElementType>{}};
    }

    float user_function(size_t num_elements, float alpha)
    {
        // Code using the above two interfaces needs to allocate to the
max alignment.
        // Users could also query aligned_accessor::byte_alignment
        // for the various interfaces and take the max.
        constexpr size_t max_byte_alignment = 32;
        auto x_alloc = allocate_raw<float, max_byte_alignment>
(num_elements);
        auto y_alloc = allocate_raw<float, max_byte_alignment>
(num_elements);

        aligned_mdspan<max_byte_alignment> x(x_alloc.get());
        aligned_mdspan<max_byte_alignment> y(y_alloc.get());

        fill_x(x); // automatic conversion from 32-byte aligned to 16-
byte aligned
        fill_y(y); // automatic conversion again

        vectorized_axpy(y, alpha, x);
        return vectorized_norm(y);
    }

```

§ 8 References

- Davis et al., “CheriABI: Enforcing Valid Pointer Provenance and Minimizing Pointer Privilege in the POSIX C Run-time Environment,” ASPLOS ’19, April 2019, pp. 379 - 393. Available online [last accessed 2024-07-05]: <https://dl.acm.org/doi/10.1145/3297858.3304042>
- Watson et al., “CHERI C/C++ Programming Guide,” Technical Report UCAM-CL-TR-947, University of Cambridge Computer Laboratory, June 2020. Available online [last accessed 2024-07-05]: <https://doi.org/10.48456/tr-947>

§ 9 Acknowledgments

— For `detectably_invalid`, credit (with permission) to David Sankel (Adobe), `dsankel@adobe.com`

§ 10 Wording

Text in blockquotes is not proposed wording, but rather instructions for generating proposed wording. The  character is used to denote a placeholder section number which the editor shall determine.

In `[version.syn]`, add

```
#define __cpp_lib_aligned_accessor YYYYMMML // also in <mdspan>
```

Adjust the placeholder value as needed so as to denote this proposal's date of adoption.

To the Header `<mdspan>` synopsis `[mdspan.syn]`, after class `default_accessor` and before class `mdspan`, add the following.

```
// [mdspan.accessor.aligned], class template aligned_accessor
template<class ElementType, size_t byte_alignment>
class aligned_accessor;
```

At the end of `[mdspan.accessor.default]` and before `[mdspan.mdspan]`, add all the material that follows.

§ 10.1 Add subsection `[mdspan.accessor.aligned]` with the following

 Class template `aligned_accessor` `[mdspan.accessor.aligned]`

.1 Overview `[mdspan.accessor.aligned.overview]`

```
template<class ElementType, size_t the_byte_alignment>
struct aligned_accessor {
    using offset_policy = default_accessor<ElementType>;
    using element_type = ElementType;
    using reference = ElementType&;
    using data_handle_type = ElementType*;

    static constexpr size_t byte_alignment = the_byte_alignment;

    constexpr aligned_accessor() noexcept = default;

    template<class OtherElementType, size_t other_byte_alignment>
        constexpr aligned_accessor(
            aligned_accessor<OtherElementType, other_byte_alignment>)
noexcept;

    template<class OtherElementType>
        explicit constexpr aligned_accessor(
```

```

        default_accessor<OtherElementType>) noexcept;

    constexpr operator default_accessor<element_type>() const {
        return {};
    }

    constexpr reference access(data_handle_type p, size_t i) const
noexcept;

    constexpr typename offset_policy::data_handle_type
        offset(data_handle_type p, size_t i) const noexcept;

    static bool is_sufficiently_aligned(data_handle_type p);
};

```

1 *Mandates:*

- (1.1) — `byte_alignment` is a power of two, and
 - (1.2) — `byte_alignment >= alignof(ElementType)` is true.
- 2 `aligned_accessor` meets the accessor policy requirements.
- 3 `ElementType` is required to be a complete object type that is neither an abstract class type nor an array type.
- 4 Each specialization of `aligned_accessor` is a trivially copyable type that models semiregular.
- 5 $[0, n)$ is an accessible range for an object `p` of type `data_handle_type` and an object of type `aligned_accessor` if and only if $[p, p + n)$ is a valid range.

§ 10.2 Members [`mdspan.accessor.aligned.members`]

```

template<class OtherElementType, size_t other_byte_alignment>
constexpr aligned_accessor(
    aligned_accessor<OtherElementType, other_byte_alignment>)
noexcept {}

```

- 1 *Constraints:* `is_convertible_v<OtherElementType(*)[], element_type(*)[]>` is true.
- 2 *Mandates:* `gcd(other_byte_alignment, byte_alignment) == byte_alignment` is true.

```

template<class OtherElementType>
explicit constexpr aligned_accessor(
    default_accessor<OtherElementType>) noexcept {};

```

- 3 *Constraints:* `is_convertible_v<OtherElementType(*)[], element_type(*)[]>` is true.

```
constexpr reference access(data_handle_type p, size_t i) const
noexcept;
```

- 4 *Preconditions:* p points to an object X of a type similar ([**conv.qual**]) to element_type, where X has alignment byte_alignment ([**basic.align**]).

- 5 *Effects:* Equivalent to: return assume_aligned<byte_alignment>(p)[i];

```
constexpr typename offset_policy::data_handle_type
offset(data_handle_type p, size_t i) const noexcept;
```

- 6 *Preconditions:* p points to an object X of a type similar ([**conv.qual**]) to element_type, where X has alignment byte_alignment ([**basic.align**]).

- 7 *Effects:* Equivalent to: return p + i;

```
static bool is_sufficiently_aligned(data_handle_type p);
```

- 8 *Preconditions:* p points to an object X of a type similar ([**conv.qual**]) to element_type.

- 9 *Returns:* true if X has alignment at least byte_alignment, else false.

[*Example:* The following function compute uses is_sufficiently_aligned to check whether a given mdspan with default_accessor has a data handle with sufficient alignment to be used with aligned_accessor<float, 4 * sizeof(float)>. If so, the function dispatches to a function compute_using_fourfold_overalignment that requires fourfold overalignment of arrays, but can therefore use hardware-specific instructions, such as four-wide SIMD (Single Instruction Multiple Data) instructions. Otherwise, compute dispatches to a possibly less optimized function compute_without_requiring_overalignment that has no overalignment requirement.

```
extern void
compute_using_fourfold_overalignment(
    std::mdspan<float, std::dims<1>, std::layout_right,
    std::aligned_accessor<float, 4 * alignof(float)>> x);

extern void
compute_without_requiring_overalignment(
    std::mdspan<float, std::dims<1>> x);

void compute(std::mdspan<float, std::dims<1>> x)
{
    auto accessor = aligned_accessor<float, 4 * sizeof(float)>{};
    auto x_handle = x.data_handle();

    if (accessor.is_sufficiently_aligned(x_handle)) {
        compute_using_fourfold_overalignment(mdspan{x_handle,
x.mapping(), accessor});
    }
    else {
        compute_without_requiring_overalignment(x);
    }
}
```


–end example]

[Example: The following example shows how users can fulfill the preconditions of `aligned_accessor` by using existing C++ Standard Library functionality to create overaligned allocations. First, the `allocate_overaligned` helper function uses `aligned_alloc` to create an overaligned allocation.

```
template<class ElementType>
struct delete_with_free {
    void operator()(ElementType* p) const {
        std::free(p);
    }
};

template<class ElementType>
using allocation = std::unique_ptr<ElementType[],
delete_with_free<ElementType>>;

template<class ElementType, size_t byte_alignment>
allocation<ElementType> allocate_overaligned(const size_t
num_elements)
{
    const size_t num_bytes = num_elements * sizeof(ElementType);
    void* ptr = std::aligned_alloc(byte_alignment, num_bytes);
    return {ptr, delete_with_free<ElementType>{}};
}
```

Second, this example presumes that some third-party library provides functions requiring arrays of float to have 32-byte alignment.

```
template<size_t byte_alignment>
using aligned_mdspan = std::mdspan<float, std::dims<1, int>,
std::layout_right, std::aligned_accessor<float,
byte_alignment>>;

extern void vectorized_axpy(aligned_mdspan<32> y, float alpha,
aligned_mdspan<32> x);
extern float vectorized_norm(aligned_mdspan<32> y);
```

Third and finally, the user’s function `user_function` would begin by allocating “raw” overaligned arrays with `allocate_overaligned`. It would then create aligned `mdspan` with them, and pass the resulting `mdspan` into the third-party library’s functions.

```
float user_function(size_t num_elements, float alpha)
{
    constexpr size_t max_byte_alignment = 32;
    auto x_alloc = allocate_overaligned<float, max_byte_alignment>
(num_elements);
    auto y_alloc = allocate_overaligned<float, max_byte_alignment>
(num_elements);

    aligned_mdspan<max_byte_alignment> x(x_alloc.get());
    aligned_mdspan<max_byte_alignment> y(y_alloc.get());
}
```

```

// ... fill the elements of x and y ...

vectorized_axpy(y, alpha, x);
return vectorized_norm(y);
}

```

–end example]

§ 11 Wording for alternative nonmember `is_sufficiently_aligned` design

Text in blockquotes is not proposed wording, but rather instructions for generating proposed wording. The instructions here are expressed as a “diff” atop P2897R2, that is, as if the wording proposed in P2897R2 had already been applied to the Working Draft. The  character is used to denote a placeholder section number which the editor shall determine.

From the `aligned_accessor` synopsis in **[mdspan.accessor.aligned.overview]**, remove the member function static `bool is_sufficiently_aligned(data_handle_type p);`.

From the list of members of `aligned_accessor` in **[mdspan.accessor.aligned.members]**, remove static `bool is_sufficiently_aligned(data_handle_type p);` from after line 7 (“*Effects*: Equivalent to: `return p + i;`”), and remove line 8 (“*Preconditions*: ...”) and line 9 (“*Returns*: ...”). (Retain the *Example* using `is_sufficiently_aligned` that immediately follows the deleted lines.)

To the Header `<bit>` synopsis **[bit.syn]**, after `enum class endian` and before the `}` that closes namespace `std`, add the following.

```

// [bit.aligned], is_sufficiently_aligned
template<class ElementType, size_t byte_alignment>
bool is_sufficiently_aligned(ElementType*);

```

After **[bit.endian]**, add a new section **[bit.aligned]** and put in it all the material that follows.

```

template<class ElementType, size_t byte_alignment>
bool is_sufficiently_aligned(ElementType* p);

```

- ¹ *Preconditions*: `p` points to an object `X` of a type similar (**[conv.qual]**) to `ElementType`.
- ² *Returns*: true if `X` has alignment at least `byte_alignment`, else false.

§ 12 Appendix A: detectably_invalid nonmember function example

This section is nonnormative. This is the full source code with tests for the `detectably_invalid` nonmember function example above. Please see this [Compiler Explorer link](#) for a test with five different compilers: GCC 14.1, Clang 18.1.0, MSVC v19.40 (VS17.10), and nvc++ 24.5.

```
#include <cassert>
#include <concepts>
#include <cstdint>
#include <iostream>
#include <stdexcept>
#include <type_traits>
#include <utility>

template<class Accessor>
concept has_detectably_invalid = requires(Accessor acc) {
    typename Accessor::data_handle_type;
    { std::as_const(acc).detectably_invalid(
        std::declval<typename Accessor::data_handle_type>(),
        std::declval<std::size_t>()
    ) } noexcept -> std::convertible_to<bool>;
};

template<class Accessor>
bool detectably_invalid(Accessor&& accessor,
    typename std::remove_cvref_t<Accessor>::data_handle_type handle,
    std::size_t size)
{
    if constexpr
(has_detectably_invalid<std::remove_cvref_t<Accessor>>) {
        return std::as_const(accessor).detectably_invalid(handle,
size);
    }
    else {
        return false;
    }
}

struct A {
    using data_handle_type = float*;

    static bool detectably_invalid(data_handle_type ptr, std::size_t
size) noexcept {
        return ptr == nullptr && size != 0;
    }
};

struct B {
    using data_handle_type = float*;
};

struct C {
    using data_handle_type = float*;

    // This is nonconst, so it's not actually called.
    bool detectably_invalid(data_handle_type ptr, std::size_t size)
```

```

{
    throw std::runtime_error("C::detectably_invalid: uh oh");
}
};

struct D {
    using data_handle_type = float*;

    // This is const but not noexcept, so it's not actually called.
    bool detectably_invalid(data_handle_type ptr, std::size_t size)
const {
    throw std::runtime_error("D::detectably_invalid: uh oh");
}
};

int main()
{
    float* ptr = nullptr;

    assert(not detectably_invalid(A{}, ptr, 0));
    assert(detectably_invalid(A{}, ptr, 1));

    A a{};
    assert(not detectably_invalid(a, ptr, 0));
    assert(detectably_invalid(a, ptr, 1));

    const A a_c{};
    assert(not detectably_invalid(a_c, ptr, 0));
    assert(detectably_invalid(a_c, ptr, 1));

    assert(not detectably_invalid(B{}, ptr, 0));
    assert(not detectably_invalid(B{}, ptr, 1));

    // B doesn't know how to check pointer validity.

    assert(not detectably_invalid(B{}, ptr, 0));
    assert(not detectably_invalid(B{}, ptr, 1));

    B b{};
    assert(not detectably_invalid(b, ptr, 0));
    assert(not detectably_invalid(b, ptr, 1));

    const B b_c{};
    assert(not detectably_invalid(b_c, ptr, 0));
    assert(not detectably_invalid(b_c, ptr, 1));

    // If users make detectably_invalid nonconst or not noexcept,
    // the nonmember function falls back to a default
implementation.

    try {
        assert(not detectably_invalid(C{}, ptr, 0));
        assert(not detectably_invalid(C{}, ptr, 1));
    }
}

```

```

catch (const std::runtime_error& e) {
    std::cerr << "C{} threw runtime_error: " << e.what() << "\n";
}

try {
    const C c_c{};
    assert(not detectably_invalid(c_c, ptr, 0));
    assert(not detectably_invalid(c_c, ptr, 1));
}
catch (const std::runtime_error& e) {
    std::cerr << "const C threw runtime_error: " << e.what() <<
"\n";
}

try {
    C c{};
    assert(not detectably_invalid(c, ptr, 0));
    assert(not detectably_invalid(c, ptr, 1));
}
catch (const std::runtime_error& e) {
    std::cerr << "nonconst C threw runtime_error: " << e.what() <<
"\n";
}

try {
    assert(not detectably_invalid(D{}, ptr, 0));
    assert(not detectably_invalid(D{}, ptr, 1));
}
catch (const std::runtime_error& e) {
    std::cerr << "D{} threw runtime_error: " << e.what() << "\n";
}

try {
    const D d_c{};
    assert(not detectably_invalid(d_c, ptr, 0));
    assert(not detectably_invalid(d_c, ptr, 1));
}
catch (const std::runtime_error& e) {
    std::cerr << "const D threw runtime_error: " << e.what() <<
"\n";
}

try {
    D d{};
    assert(not detectably_invalid(d, ptr, 0));
    assert(not detectably_invalid(d, ptr, 1));
}
catch (const std::runtime_error& e) {
    std::cerr << "nonconst D threw runtime_error: " << e.what() <<
"\n";
}

std::cerr << "Made it to the end\n";
return 0;
}

```

§ 13 Appendix B: Implementation and demo

This [Compiler Explorer link](#) gives a full implementation of `aligned_accessor` and a demonstration. We show the full source code from that link here below.

```
#include <https://raw.githubusercontent.com/kokkos/mdspan/single-  
header/mdspan.hpp>  
#include <bit>  
#include <cassert>  
#include <cmath>  
#if defined(_MSC_VER)  
# include <cstdlib> // MSVC's _aligned_malloc  
#endif  
#include <exception>  
#include <functional>  
#include <memory>  
#include <numeric>  
#include <type_traits>  
  
namespace stdex = std::experimental;  
  
// P2389 (voted into C++ at June 2024 STL plenary)  
namespace std {  
    template<size_t Rank, class IndexType = size_t>  
    using dims = dextents<IndexType, Rank>;  
} // namespace std  
  
namespace {  
  
    template<class ElementType, std::size_t byte_alignment>  
    class aligned_accessor {  
    public:  
        static_assert(std::has_single_bit(byte_alignment),  
            "byte_alignment must be a power of two.");  
        static_assert(byte_alignment >= alignof(ElementType),  
            "Insufficient byte alignment for ElementType");  
  
        using offset_policy = stdex::default_accessor<ElementType>;  
        using element_type = ElementType;  
        using reference = ElementType&;  
        using data_handle_type = ElementType*;  
  
        constexpr aligned_accessor() noexcept = default;  
  
        template<  
            class OtherElementType,  
            std::size_t other_byte_alignment>  
        requires(std::is_convertible_v<  
            OtherElementType(*)[], element_type(*)[]>)  
        constexpr aligned_accessor(  
            aligned_accessor<OtherElementType, other_byte_alignment>  
            noexcept
```

```

{
    constexpr size_t the_gcd =
        std::gcd(other_byte_alignment, byte_alignment);
    static_assert(the_gcd == byte_alignment);
}

template<class OtherElementType>
requires(std::is_convertible_v<
    OtherElementType(*)[], element_type(*)[]>)
constexpr explicit aligned_accessor(
    stdex::default_accessor<OtherElementType>) noexcept
{}

constexpr
    operator stdex::default_accessor<element_type>() const
{
    return {};
}

constexpr reference
    access(data_handle_type p, size_t i) const noexcept
{
    return std::assume_aligned<byte_alignment>(p)[i];
}

constexpr typename offset_policy::data_handle_type
offset(data_handle_type p, size_t i) const noexcept {
    return p + i;
}

static bool
    is_sufficiently_aligned(data_handle_type p) noexcept
{
    return alignof(ElementType) >= byte_alignment &&
        std::bit_cast<uintptr_t>(p) % byte_alignment == 0;
};

template<size_t byte_alignment>
using aligned_mdspan =
    std::mdspan<float, std::dims<1, int>, std::layout_right,
        aligned_accessor<float, byte_alignment>>;

// Interfaces that require 32-byte alignment,
// because they want to do 8-wide SIMD of float.
void
vectorized_axpby(aligned_mdspan<32> y,
    float alpha, aligned_mdspan<32> x, float beta)
{
    assert(x.extent(0) == y.extent(0));
    for (int k = 0; k < x.extent(0); ++k) {
        y[k] = beta * y[k] + alpha * x[k];
    }
}

```

```

// 1-norm of the vector y
float vectorized_norm(aligned_mdspan<32> y)
{
    float one_norm = 0.0f;
    for (int k = 0; k < y.extent(0); ++k) {
        one_norm += std::fabs(y[k]);
    }
    return one_norm;
}

// Interfaces that require 16-byte alignment,
// because they want to do 4-wide SIMD of float.
void fill_x(aligned_mdspan<16> x) {
    for (int k = 0; k < x.extent(0); ++k) {
        x[k] = static_cast<float>(k + 2);
    }
}
void fill_y(aligned_mdspan<16> y) {
    for (int k = 0; k < y.extent(0); ++k) {
        y[k] = static_cast<float>(k - 1);
    }
}

// Helper functions for making overaligned array allocations.

template<class ElementType>
struct delete_raw {
    void operator()(ElementType* p) const {
        std::free(p);
    }
};

template<class ElementType>
using allocation =
    std::unique_ptr<ElementType[], delete_raw<ElementType>>;

template<class ElementType, std::size_t byte_alignment>
allocation<ElementType> allocate_raw(const std::size_t
num_elements)
{
    const std::size_t num_bytes = num_elements *
sizeof(ElementType);
    float* ptr = reinterpret_cast<float*>(
#ifdef _MSC_VER
        _aligned_malloc(byte_alignment, num_bytes)
#else
        std::aligned_alloc(byte_alignment, num_bytes)
#endif
    );
    return {ptr, delete_raw<ElementType>{}};
}

float user_function(size_t num_elements, float alpha, float beta)
{
    constexpr size_t max_byte_alignment = 32;

```



```

        // unique_ptr<float[], deleter_something>
        auto x_alloc = allocate_raw<float, max_byte_alignment>
(num_elements);
        auto y_alloc = allocate_raw<float, max_byte_alignment>
(num_elements);

        aligned_mdspan<max_byte_alignment> x(x_alloc.get(),
num_elements);
        aligned_mdspan<max_byte_alignment> y(y_alloc.get(),
num_elements);

        // expect aligned_mdspan<16>, get aligned_mdspan<32>
        fill_x(x); // automatic conversion from 32-byte aligned to 16-
byte aligned
        fill_y(y); // automatic conversion again

        // expect aligned_mdspan<32>, get aligned_mdspan<32>
        vectorized_axpby(y, alpha, x, beta);
        return vectorized_norm(y);
    }

} // namespace (anonymous)

int main(int argc, char* argv[])
{
    float result = user_function(10, 1.0f, -1.0f);
    // 3 + 3 + ... + 3 = 30
    assert(result == 30.0f);
    return 0;
}

```