

Document number: P3111R8.

Date: 2025-06-19.

Authors: Gonzalo Brito Gadeschi, Simon Cooksey, Daniel Lustig.

Reply to: Gonzalo Brito Gadeschi <gonzalob_at_nvidia.com (<http://nvidia.com>)>.

Audience: LWG.

Table of Contents

- Atomic Reduction Operations
 - Introduction
 - Motivation
 - Hardware Exposure
 - Performance
 - Functionality
 - Design
 - Implementation impact
 - Design Alternatives
 - API Naming
 - Definition Naming
 - Memory Ordering
 - Formalization
 - Wording
 - Forward progress
 - Atomic Modify-Write Operation APIs
 - No acquire sequences support

Changelog

- **R8:**
 - Update fp min/max to latest P3008 revision.

- **R7:** CWG saw this in Sofia and with the following changes made in R7 does not need to see it again:
 - Remove review notes.
 - Change section order to match standard.
 - Remove unnecessary wording added by P3008.
- **R6**
 - Fix typo: PPC64LE.
 - Add implementation recommendations on PPC64LE.
 - Add definition of atomic modify-write operation (previously atomic reduction operation) to wording of [atomics.order] (<https://eel.is/c++draft/atomics.order>).
 - Update definition name from "atomic reduction operation" to "atomic modify-write operation".
 - Cleanup how the functionality of P3008 is incorporated into the wording (it already was).
- **R5: post-Austria:**
 - Import `constexpr atomics` update from P3309 (<http://wg21.link/p3309>), i.e., new `store_key atomics` are `constexpr` in the same cases as `fetch_key`.
 - Conditionally import floating-point atomic min/max from P3008 (<http://wg21.link/p3008>).
 - Add feature testing macro.
 - Move `compare_store` into a separate paper: D3644R0 (<http://wg21.link/p3644>).
 - LEWG **POLL:** In P3111 Change the name from `reduce_*` to `store_*` and `atomic_reduce_*` to `atomic_store_*`

SF	F	N	A	SA
6	8	0	0	0

- Attendance: 14 (IP) + 6 (R)
 Author's Position: SF
 Outcome: Strong consensus in favour

- LEWG **POLL:** With the changed names forward P3111R4 to LWG and CWG for C++26

SF	F	N	A	SA
5	6	0	0	0

- Attendance: 14 (IP) + 6 (R)

Author's Position: SF

Outcome: Consensus in favour

- Change the name from `reduce_*` to `store_*` and `atomic_reduce_*` to `atomic_store_*` according to LEWG poll above.

- **R4: Austria:**

- EWG saw R3 in Hagenberg. **POLL:** P3111R3: Atomic Reduction Operations: Forward to LEWG and CWG for C++26

SF	F	N	A	SA
2	16	3	0	0

- Consensus

- **R3: pre Austria:**

- Refactor atomic reduction sequence replacements into tables.
- Provide alternative wording for atomic reduction sequence replacement using "as if" integer operations to provide same valid replacements as for integer reduction operations.
- Fixed FP reduction wording.

- **R2: post Wroclaw:** simplifies reduction sequence wording.

- **R1: post St. Louis '24**

- *Overview of changes:* incorporate SG6 and SG1 feedback into the proposal, which resolved all open questions. Restructure exposition about what's being proposed accordingly.
- [SG6 Review]:
 - Forward it; does not need to look at it again.
 - Agree on "generalized" reductions being the default and not providing version without.

- Do not use `GENERALIZED_SUM` to specify non-associative floating-point atomic reduction operations since it does not make sense for two values. Instead, attempt to add a `std::non_associative_add` operation that is exempted from the C standard clause in Annex F that makes re-ordering non-conforming. Recommended discussing this with CWG, which I did, and caused us to end up pursuing a different alternative (defining "reduction sequences") which is pursued in R1 instead.
- [SG1 Review]: Add wording for reduction sequences.
- [SG1 Review]: Add `compare_store` operation.
- [SG1 Review]: Update to handle `fenv` and `errno` for floating-point atomics.
- [SG1 Review]: Make "generalized" atomic floating-point reductions the default and do not provide fallback (beyond `fetch_<op>`).
- [SG1 Review]: Provide unsequenced support.
- [SG1 Review]: polls taken:
 1. Reduction operations are not a step, but infinite loops around reductions operations are not UB (practically implementations can assume reductions are finite)

SF	F	N	A	SA
3	9	1	0	0

2. The single floating point reduction operations should allow "fast math" (allow tree reductions), you can get precise results using `fetch_add/sub`

SF	F	N	A	SA
5	5	1	1	0

3. Any objection to approve the design of P3111R0: No objection.

- **R0: initial revision.**

Atomic Reduction Operations

Atomic Reduction Operations are atomic read-modify-write (RMW) operations (like `fetch_add`) that do not "fetch" the old value and are not reads for the purpose of building synchronization with acquire fences. This enables implementations to leverage hardware acceleration available in modern CPU and GPU architectures.

Furthermore, we propose to allow atomic memory operations that aren't reads in unsequenced execution, and to extend atomic arithmetic reduction operations for floating-point types with operations that assume floating-point arithmetic is associative.

Introduction

Concurrent algorithms performing atomic RMW operations that discard the old fetched value are very common in high-performance computing, e.g., finite-element matrix assembly, data analytics (e.g. building histograms), etc.

Consider the following parallel algorithm to build a histogram (full implementation (<https://clang.godbolt.org/z/zscrhEPYj>)):

```
1 // Example 0: Histogram.
2 span<unsigned> data;
3
4 array<atomic<unsigned>, N> buckets;
5 constexpr T bucket_sz = numeric_limits<T>::max() / (T)N;
6 unsigned nthreads = thread::hardware_concurrency();
7 for_each_n(execution::par_unseq, views::iota(0).begin(), nthreads,
8   [&](int thread) {
9     unsigned data_per_thread = data.size() / nthreads;
10    T* data_thread = data.data() + data_per_thread * thread;
11    for (auto e : span<T>(data_thread, data_per_thread))
12      buckets[e / bucket_sz].fetch_add(1, memory_order_relaxed);
13  });
```

This program has two main issues:

- **Correctness** (undefined behavior): The program should have used the `par` execution policy to avoid undefined behavior since the atomic operation is not:
 - Potentially concurrent and therefore exhibits a data-race in unsequenced contexts ([intro.execution]).

- Vectorization-safe [algorithms.parallel.defns#5]
(<https://eel.is/c++draft/algorithms.parallel.defns#5>) since it is specified to synchronize with other function invocations.

- **Performance:** Sophisticated compiler analysis required to optimize the above program for scalable hardware architectures with atomic reduction operations.

Atomic reduction operations address both shortcomings:

Before (https://clang.godbolt.org/z/xq8efq1WK)	After
<pre> 1 #include <algorithm> 2 #include <atomic> 3 #include <execution> 4 using namespace std; 5 using execution::par_unseq; 6 7 int main() { 8 size_t N = 10000; 9 vector<int> v(N, 0); 10 atomic<int> atom = 0; 11 for_each_n(par_unseq, 12 v.begin(), N, 13 [&](auto& e) { 14 // UB+SLOW: 15 atom.fetch_add(e); 16 }); 17 return atom.load(); 18 }</pre>	<pre> 1 #include <algorithm> 2 #include <atomic> 3 #include <execution> 4 using namespace std; 5 using execution::par_unseq; 6 7 int main() { 8 size_t N = 10000; 9 vector<int> v(N, 0); 10 atomic<int> atom = 0; 11 for_each_n(par_unseq, 12 v.begin(), N, 13 [&](auto& e) { 14 // OK+FAST 15 atom.store_add(e); 16 }); 17 return atom.load(); 18 }</pre>

This new operation can then be used in the Histogram Example (example 0), to replace the `fetch_add` with `store_add`.

Motivation

Hardware Exposure

The following ISAs provide Atomic Reduction Operation:

Architecture	Instructions
PTX	<code>red</code> (https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#parallel-synchronization-and-communication-instructions-red).
ARM	<code>LDADD RZ</code> (https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/LDADD--LDADDA--LDADDAL--LDADDL--Atomic-add-on-word-or-doubleword-in-memory-?lang=en), <code>STADD</code> (https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/STADD--STADDL--Atomic-add-on-word-or-doubleword-in-memory--without-return--an-alias-of-LDADD--LDADDA--LDADDAL--LDADDL-), <code>SWP RZ</code> (https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/SWP--SWPA--SWPAL--SWPL--Swap-word-or-doubleword-in-memory-?lang=en), <code>CAS RZ</code> (https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/CAS--CASA--CASAL--CASL--Compare-and-Swap-word-or-doubleword-in-memory-?lang=en).
x86-64	Remote Atomic Operations (RAO) (https://cdrdv2-public.intel.com/671368/architecture-instruction-set-extensions-programming-reference.pdf): <code>AADD</code> , <code>AAND</code> , <code>AOR</code> , <code>AXOR</code> .
RISC-V	None (note: AMOs are always loads and stores).
PPC64LE	None (note: <code>stdat</code> family of operations are always loads and stores).

Some of these instructions lack a destination operand (`red` (<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#parallel-synchronization-and-communication-instructions-red>), `STADD` (<https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/STADD--STADDL--Atomic-add-on-word-or-doubleword-in-memory--without-return--an-alias-of-LDADD--LDADDA--LDADDAL--LDADDL->), `AADD`). Others change semantics if the destination register used discards the result (Arm's `LDADD RZ` (<https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/LDADD--LDADDA--LDADDAL--LDADDL--Atomic-add-on-word-or-doubleword-in-memory-?lang=en>), `SWP RZ` (<https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/SWP--SWPA--SWPAL--SWPL--Swap-word-or-doubleword-in-memory-?lang=en>), `CAS RZ` (<https://developer.arm.com/documentation/ddi0602/2022-06/Base-Instructions/CAS--CASA--CASAL--CASL--Compare-and-Swap-word-or-doubleword-in-memory-?lang=en>)).

All ISAs provide the same semantics: these are not loads from the point-of-view of the Memory Model, and therefore do not participate in acquire sequences, but they do participate in release sequences:

- PTX Specification: `red` is "not a read operation" (<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#operation-types>)".
- Arm ARM: "where the destination register is WZR or XZR, are not regarded as doing a read for the purpose of a DMB LD barrier" (<https://developer.arm.com/documentation/ddi0487/latest>)".
- x86-64: "since they do not load data from memory into the processor." (<https://cdrdv2-public.intel.com/671368/architecture-instruction-set-extensions-programming-reference.pdf>)

These architectures provide both "relaxed" and "release" orderings for the reductions (e.g. `red.relaxed` / `red.release` , `STADD` / `STADDL`).

The Power ISA `stdat` family of instructions are intended to accelerate read-modify-write operations that benefit from being performed at memory. This is different from the intent of C++ atomic reduction operations, which is to avoid unnecessary synchronization caused by the use of read-modify-write operations and acquire fences. C++ implementations for the Power ISA should not use `stdat` to implement C++ atomic reduction operations to avoid low performance resulting from using these instructions for use cases they are not intended for. For more information, see Power ISA Version 3.1B (https://wiki.raptorcs.com/w/images/d/d3/OPF_PowerISA_v3.1B.pdf) Section 4.5 p.1080.

Performance

On hardware architectures that implement these as far atomics, the exposed latency of Atomic Reduction Operations may be as low as half that of " `fetch_<key>` " operations.

Example: on an NVIDIA Hopper H100 GPU, replacing `atomic.fetch_add` with `atomic.store_add` on the Histogram Example (Example 0) improves throughput by 1.2x.

Furthermore, non-associative floating-point atomic operations, like `fetch_add` , are required to read the "latest value", which sequentializes their execution. In the following example, the outcome `x == a + (b + c)` is not allowed, because either the atomic operation of thread0 happens-before that of thread1, or vice-versa, and floating-point arithmetic is not associative:

```
1 // Litmus test 2:
2 atomic<float> x = a;
3 void thread0() { x.fetch_add(b, memory_order_relaxed); }
4 void thread1() { x.fetch_add(c, memory_order_relaxed); }
```


Allowing the `x == a + (b + c)` outcome enables implementations to perform a tree-reduction, which improves complexity from $O(N)$ to $O(\log(N))$, at negligible higher amount of non-determinism which is already inherent to the use of atomic operations:

```
1 // Litmus test 3:
2 atomic<float> x = a;
3 x.store_add(b, memory_order_relaxed);
4 x.store_add(c, memory_order_relaxed);
5 // Sound to merge these two operations into one:
6 // x.store_add(b + c);
```

On GPU architectures, performing an horizontal reduction for then issuing a single atomic operation per thread group, reduces the number of atomic operation issued by up to the size of the thread group.

Functionality

Currently, all atomic memory operations are vectorization-unsafe (<https://eel.is/c++draft/algorithms.parallel.defns#5>) and therefore not allowed in element access functions of parallel algorithms when the `unseq` or `par_unseq` execution policies are used (see [\[algorithms.parallel.exec.5\]](#) (<https://eel.is/c++draft/algorithms.parallel#exec-5>) and [\[algorithms.parallel.exec.7\]](#) (<https://eel.is/c++draft/algorithms.parallel#exec-7>)). Atomic memory operations that "read" (e.g. `load`, `fetch_key`, `compare_exchange`, `exchange`, ...) enable building synchronization edges that block, which within `unseq` / `par_unseq` leads to deadlocks. N4070 (<https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4070.html>) solved this by tightening the wording to disallow any synchronization API from being called from within `unseq` / `par_unseq`.

Allowing Atomic Writes and Atomic Reduction Operations in unsequenced execution increases the set of concurrent algorithms that can be implemented in the lowest-common denominator of hardware that C++ supports. In particular, many hardware architectures that can accelerate `unseq` / `par_unseq` but cannot accelerate `par` (e.g. most non-NVIDIA GPUs), provide acceleration for atomic reduction operations.

We propose to make lock-free atomic operations that are not reads vectorization safe to enable calling them from unsequenced execution. Atomic operations that read remain vectorization-unsafe and therefore UB:

Design

For each `atomic_fetch_{OP}` in the `atomic<T>` and `atomic_ref<T>` class templates and their specializations, we introduce new `store_{OP}` member functions that return `void` :

```
1  template <class T>
2  struct atomic_ref {
3      T      fetch_add (T v, memory_order o) const noexcept;
4      void store_add(T v, memory_order o) const noexcept;
5  };
```

The `store_{OP}` APIs are vectorization safe if the atomic `is_always_lock_free` is true, i.e., they are allowed in Element Access Functions ([`algorithms.parallel.defns`] (https://eel.is/c++draft/algorithms.parallel.defns#def:element_access_functions)) of parallel algorithms:

```
for_each(par_unseq, ..., [&](auto old, ...) {
    assert(atom.is_lockfree()); // Only for lock-free atomics
    atom.store(42);              // OK: vectorization-safe
    atom.store_add(1);           // OK: vectorization-safe
    atom.fetch_add(1);           // UB: vectorization-unsafe
    atom.exchange(42);           // UB: vectorization-unsafe
    atom.compare_exchange_weak(old, 42); // UB: vectorization-unsafe
    atom.compare_exchange_strong(old, 42); // UB: vectorization-unsafe
    while (atom.load() < 42);    // UB: vectorization-unsafe
});
```

Furthermore, we specify non-associative floating-point atomic reduction operations to enable tree-reduction implementations that improve complexity from $O(N)$ to $O(\log(N))$ by minimally increasing the non-determinism which is already inherent to atomic operations. This allows producing the `x == a + (b + c)` outcome in the following example, which enables the optimization that merges the two `store_add` operations into a single one:

```
1  atomic<float> x = a;
2  x.store_add(b, memory_order_relaxed);
3  x.store_add(c, memory_order_relaxed);
4  // Sound to merge these two operations into one:
5  // x.store_add(b + c);
```

Applications that need the sequential semantics can use `fetch_add` instead.

Implementation impact

It is correct to conservatively implement `store_{0P}` as a call to `fetch_{0P}`. We evaluated the following implementations of unsequenced execution policies, which are not impacted:

- OpenMP `simd` (<https://www.openmp.org/spec-html/5.0/openmpsu42.html>) `pragma for unseq` and `par_unseq`, since OpenMP supports atomics within `simd` regions.
- `pragma clang loop` (<https://clang.llvm.org/docs/LanguageExtensions.html#extensions-for-loop-hint-optimizations>) is a hint.

Design Alternatives

Enable Atomic Reductions as `fetch_<key>` optimizations

Attempting to improve application performance by implementing compiler-optimizations to leverage Atomic Reduction Operations from `fetch_<key>` APIs whose result is unused has become a rite of passage for compiler engineers, e.g., GCC#509632 (<https://gcc.gnu.org/pipermail/gcc-patches/2018-October/509632.html>), LLVM#68428 (<https://github.com/llvm/llvm-project/issues/68428>), LLVM#72747 (<https://github.com/llvm/llvm-project/pull/72747>), ... Unfortunately, "simple" optimization strategies break backward compatibility in the following litmus tests (among others).

Litmus Test 0: from Will Deacon (<https://gcc.gnu.org/pipermail/gcc-patches/2018-October/509632.html>). Performing the optimization to replace the introduces the `y == 2 && r0 == 1 && r1 == 0` outcome:

```
1 void thread0(atomic_int* y,atomic_int* x) {
2     atomic_store_explicit(x,1,memory_order_relaxed);
3     atomic_thread_fence(memory_order_release);
4     atomic_store_explicit(y,1,memory_order_relaxed);
5 }
6
7 void thread1(atomic_int* y,atomic_int* x) {
8     atomic_fetch_add_explicit(y,1,memory_order_relaxed);
9     atomic_thread_fence(memory_order_acquire);
10    int r0 = atomic_load_explicit(x,memory_order_relaxed);
11 }
12
13 void thread2(atomic_int* y) {
14     int r1 = atomic_load_explicit(y,memory_order_relaxed);
15 }
```

Litmus Test 1: from Luke Geeson (<https://github.com/llvm/llvm-project/issues/68428#issue-1930595855>). Performing the optimization of replacing the exchange with a store introduces the `r0 == 0 && y == 2` outcome:

```
1 void thread0(atomic_int* y,atomic_int* x) {
2     atomic_store_explicit(x,1,memory_order_relaxed);
3     atomic_thread_fence(memory_order_release);
4     atomic_store_explicit(y,1,memory_order_relaxed);
5 }
6 void thread1(atomic_int* y,atomic_int* x) {
7     atomic_exchange_explicit(y,2,memory_order_release);
8     atomic_thread_fence(memory_order_acquire);
9     int r0 = atomic_load_explicit(x,memory_order_relaxed);
10 }
```

In some architectures, Atomic Reduction Operations can write to memory pages or memory locations that are not readable, e.g., MMIO registers on NVIDIA GPUs, and need a reliable programming model that does not depend on compiler-optimizations for functionality.

API Naming

SG1 and LEWG considered the following alternative atomic API names:

- `atomic.add` : breaks for logical operations (e.g. `atomic.and` , etc.).
- `atomic.reduce_add` : clearly state what this operation is for (reductions).
- `atomic.store_add` : clearly states that this operation is a "store" (and not a "load").
 - In Hagenberg LEWG decided to use this name

We considered providing separate version of non-associative floating-point atomic reduction operations with and without the support for tree-reductions, e.g., `atomic.store_add` (no tree-reduction support) and `atomic.store_add_generalized` (tree-reduction support), but decided against it because atomic operations are inherently non-deterministic, this relaxation only minimally impacts that, and `fetch_add` already provides (and has to continue to provide) the sequential semantics without this relaxation.

Definition Naming

The standard-internal name of these operations in the wording is under the purview of CWG. While SG1 and EWG proposed "atomic reduction operations", in preparation for CWG review several alternative names were considered (list below), and the name "atomic modify-write operation" was chosen.

Alternatives considered:

- atomic reduction operations:
 - PRO:
 - Some people find this functional connotation clear and approachable.
 - Name used by e.g. PTX.
 - CON:
 - Describes a particular problem distant from actual semantics.
 - Interaction with acquire fences not clear (but the `store_` in the APIs makes it clear).
 - The operation itself does not reduce anything; it's a chain of such operations that reduces values.
- atomic modify-write operations
 - PRO: implies lack of acquire-ness since acquire only go on reads.
 - CON: anything that does a store seems like it ought to be a "modify-write" operation.
- atomic modify-update operations: similar to atomic modify-write operations.
- "store read-modify-write atomics":
 - PRO:
 - Clear interaction with acquire.
 - Aligns with standard library API names (`store_<key>`).
- unordered atomics:
 - PRO: hints at tree reductions for `par_unseq` , where we simply guarantee no particular order of computation even on the atomic itself (for relaxed atomic

operations, we guarantee an order on the atomic, but not with respect to other memory).

- CON: different audiences use `unordered` to refer to widely different semantics.
 - LLVM IR uses the `unordered` atomic memory order to implement the semantics of Java non-atomic memory accesses to non-volatile shared memory variables. It is not strong enough for inter-thread synchronization, but gives well-defined semantics for data-races.
 - Networking programming models (e.g. SHMEM put/get) use `unordered` to refer to memory accesses that are not in program order (e.g. "unsequenced" accesses in C++ speak, i.e., accesses from the same thread that are not sequenced-before each other).

Memory Ordering

We choose to support `memory_order_relaxed` , `memory_order_release` , and `memory_order_seq_cst` .

We may need a note stating that replacing the operations in a reduction sequence is only valid as long as the replacement maintains the other ordering properties of the operations as defined in [intro.races]. Examples:

Litmus test 2:

```
1 // T0:
2 M1.store_add(1, relaxed);
3 M2.store(1, release);
4 M1.store_add(1, relaxed);
5
6 // T1:
7 r0 = M2.load(acquire);
8 r1 = M1.load(relaxed);
9
10 // r0 == 0 || r1 >= 1
```

Litmus test 3:

```
1 // T0
2 M1.store_add(1, seq_cst);
3 M2.store_add(1, seq_cst);
4 M1.store_add(1, seq_cst);
5
6 // T1
7 r0 = M2.load(seq_cst);
8 r1 = M1.load(seq_cst);
9
10 // r0 == 0 || r1 >= 1
```

Unresolved question: Are tree reductions only supported for `memory_order_relaxed` ? No, see litmus tests for `release` and `seq_cst`.

Formalization

Herd already support these for `STADD` on Arm, and the NVIDIA Volta Memory Model supports these for `red` and `multimem.red` on PTX. If we decide to pursue this exposure direction, this proposal would benefit from extending Herd's RC11 with reduction sequences for floating-point.

Wording

Forward progress

Modify [intro.progress#1] (<https://eel.is/c++draft/intro.progress#1>) as follows:

The implementation may assume that any thread will eventually do one of the following:

1. terminate,
2. invoke the function `std::this_thread::yield ([thread.thread.this])`,
3. make a call to a library I/O function,
4. perform an access through a volatile glvalue,
5. ~~perform an synchronization operation or an atomic operation~~atomic or synchronization operation other than an atomic modify-write operation [atomics.order]_(<https://eel.is/c++draft/atomics.order>), or
6. continue execution of a trivial infinite loop (`[stmt.iter.general]`).

[Note 1: This is intended to allow compiler transformations such as removal of empty loops, even when termination cannot be proven. — end note]

Modify [intro.progress#3] (<https://eel.is/c++draft/intro.progress#3>) as follows:

During the execution of a thread of execution, each of the following is termed an execution step:

1. termination of the thread of execution,
2. performing an access through a volatile glvalue, or
3. completion of a call to a library I/O function,~~a synchronization operation, or an atomic operation.~~or
4. completion of an atomic or synchronization operation other than an atomic modify-write operation [atomics.order]_(<https://eel.is/c++draft/atomics.order>).

Add to [algorithms.parallel.defns] (<https://eel.is/c++draft/algorithms.parallel.defns#5>):

A standard library function is vectorization-unsafe if it is specified to synchronize with another function invocation, or another function invocation is specified to synchronize with it, and if it is not a memory allocation or deallocation function or lock-free atomic modify-write operation [atomics.order] (<https://eel.is/c++draft/atomics.order>).

[Note 2: Implementations must ensure that internal synchronization inside standard library functions does not prevent forward progress when those functions are executed by threads of execution with weakly parallel forward progress guarantees. — end note]

[Example 2:

```
int x = 0;
std::mutex m;
void f() {
    int a[] = {1,2};
    std::for_each(std::execution::par_unseq, std::begin(a), std::end(a), [&](int i) {
        std::lock_guard<mutex> guard(m); // incorrect: lock_guard constructor call
        ++x;
    });
}
```

The above program may result in two consecutive calls to `m.lock()` on the same thread of execution (which may deadlock), because the applications of the function object are not guaranteed to run on different threads of execution. — end example]

Atomic Modify-Write Operation APIs

Add to [atomics.syn] (<https://eel.is/c++draft/atomics.syn>):

```

namespace std {
// [atomics.nonmembers], non-member functions
...

template<class T>
    void atomic_store_add(volatile atomic<T>*, _____ // freestanding
                        typename atomic<T>::difference_type) noexcept;

template<class T>
    constexpr void atomic_store_add(atomic<T>*, _____ // freestanding
                                   typename atomic<T>::difference_type) noexcept;

template<class T>
    void atomic_store_add_explicit(volatile atomic<T>*, _____ // freestanding
                                  typename atomic<T>::difference_type,
                                  memory_order) noexcept;

template<class T>
    constexpr void atomic_store_add_explicit(atomic<T>*, _____ // freestanding
                                             typename atomic<T>::difference_type,
                                             memory_order) noexcept;

template<class T>
    void atomic_store_sub(volatile atomic<T>*, _____ // freestanding
                        typename atomic<T>::difference_type) noexcept;

template<class T>
    constexpr void atomic_store_sub(atomic<T>*, _____ // freestanding
                                   typename atomic<T>::difference_type) noexcept;

template<class T>
    void atomic_store_sub_explicit(volatile atomic<T>*, _____ // freestanding
                                  typename atomic<T>::difference_type,
                                  memory_order) noexcept;

template<class T>
    constexpr void atomic_store_sub_explicit(atomic<T>*, _____ // freestanding
                                             typename atomic<T>::difference_type,
                                             memory_order) noexcept;

template<class T>
    void atomic_store_and(volatile atomic<T>*, _____ // freestanding
                        typename atomic<T>::value_type) noexcept;

template<class T>
    constexpr void atomic_store_and(atomic<T>*, _____ // freestanding
                                   typename atomic<T>::value_type) noexcept;

template<class T>
    void atomic_store_and_explicit(volatile atomic<T>*, _____ // freestanding
                                  typename atomic<T>::value_type,
                                  memory_order) noexcept;

template<class T>
    constexpr void atomic_store_and_explicit(atomic<T>*, _____ // freestanding
                                             typename atomic<T>::value_type,
                                             memory_order) noexcept;

template<class T>
    void atomic_store_or(volatile atomic<T>*, _____ // freestanding
                       typename atomic<T>::value_type) noexcept;

```



```

template<class T>
constexpr void atomic_store_min_explicit(atomic<T>*, _____ // freestanding
                                     typename atomic<T>::value_type,
                                     memory_order) noexcept;

...
}

```

Add to [atomics.ref.int] (<https://eel.is/c++draft/atomics.ref.int>):

```

namespace std {
  template<> struct atomic_ref<integral-type> {
    ...
    public:
      ...
      constexpr void store_add(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_sub(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_and(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_or(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_xor(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_max(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
      constexpr void store_min(value_type,
                              memory_order = memory_order::seq_cst) const noexcept;
  }
}

```

- Drafting Note: `key` in `store_key` below needs to be italics

```

constexpr void store_key(value_type operand,
                        memory_order order = memory_order_seq_cst) const noexcept

```

- Preconditions: `order` is `memory_order_relaxed`, `memory_order_release`, or `memory_order_seq_cst`.
- Effects: Atomically replaces the value referenced by `*ptr` with the result of the computation applied to the value referenced by `*ptr` and the given `operand`. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([atomics.order] (<https://eel.is/c++draft/atomics.order>)).

- Remarks: Except for store_max and store_min, for signed integer types, the result is as if *ptr and parameters were converted to their corresponding unsigned types, the computation performed on those types, and the result converted back to the signed type. [Note 1: There are no undefined results arising from the computation. — end note]. For store_max and store_min, the maximum and minimum computation is performed as if by max and min algorithms ([alg.min.max]), respectively, with *ptr and the first parameter as the arguments.

Add to [atomics.ref.float] (<https://eel.is/c++draft/atomics.ref.float>):

```
namespace std {
    template<> struct atomic_ref<floating-point-type> {
        ...
    public:
        ...
        constexpr void store_add(value_type,
                                memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_sub(value_type,
                                memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_max(value_type, memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_min(value_type, memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_fmaximum(value_type, memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_fminimum(value_type, memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_fmaximum_num(value_type, memory_order = memory_order::seq_cst) const noexcept;
        constexpr void store_fminimum_num(value_type, memory_order = memory_order::seq_cst) const noexcept;
    };
}
```

P3008 modified [atomics.ref.float.4] (<https://eel.is/c++draft/atomics#ref.float-4>) as follows:

The following operations perform arithmetic computations. The correspondence among key, operator, and computation is specified in Table 153 except for the keys max, min, fmaximum, fminimum, fmaximum_num, and fminimum_num, which are specified below.

- Drafting Note: key in store_key below needs to be italics

```
constexpr void store_key(value_type operand,
                        memory_order order = memory_order::seq_cst) const noexcept;
```

- Preconditions: order is memory_order_relaxed, memory_order_release, or memory_order_seq_cst.

- Effects: Atomically replaces the value referenced by `*ptr` with the result of the computation applied to the value referenced by `*ptr` and the given operand. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([`atomics.order`])(<https://eel.is/c++draft/atomics.order>)).
- Remarks: If the result is not a representable value for its type ([`expr.pre`]), the result is unspecified, but the operations otherwise have no undefined behavior. Atomic arithmetic operations on floating-point-type should conform to the `numeric_limits<floating-point-type>` traits associated with the floating-point type ([`limits.syn`]). The floating-point environment ([`cfenv`]) for atomic arithmetic operations on floating-point-type may be different than the calling thread's floating-point environment. The arithmetic rules of floating-point atomic modify-write operations may be different from operations on floating-point types or atomic floating-point types.
[Note: Tree reductions are permitted for atomic modify-write operations. – end note]
- Remarks:
 - For `store fmaximum` and `store fminimum`, the maximum and minimum computation is performed as if by `fmaximum` and `fminimum`, respectively, with `*ptr` and the first parameter as the arguments.
 - For `store fmaximum num` and `store fminimum num`, the maximum and minimum computation is performed as if by `fmaximum num` and `fminimum num`, respectively, with `*ptr` and the first parameter as the arguments.
 - For `store max` and `store min`, the maximum and minimum computation is performed as if by `fmaximum num` and `fminimum num`, respectively, with `*ptr` and the first parameter as the arguments, except that:
 - If both arguments are NaN, an unspecified NaN value is stored at `*ptr`.
 - If exactly one argument is a NaN, either the other argument or an unspecified NaN value is stored at `*ptr`, it is unspecified which.
 - If the arguments are differently signed zeros, which of these values is stored at `*ptr` is unspecified.
- Recommended practice: The implementation of `store max` and `store min` should treat negative zero as smaller than positive zero.

Add to [`atomics.ref.pointer`] (<https://eel.is/c++draft/atomics.ref.generic#atomics.ref.pointer>):

```

namespace std {
  template<class T> struct atomic_ref<T*> {
    ...
  public:
    ...
    constexpr void store_add(difference_type,
                             memory_order = memory_order::seq_cst) const noexcept;
    constexpr void store_sub(difference_type,
                             memory_order = memory_order::seq_cst) const noexcept;
    constexpr void store_max(value_type,
                             memory_order = memory_order::seq_cst) const noexcept;
    constexpr void store_min(value_type,
                             memory_order = memory_order::seq_cst) const noexcept;
  };
}

```

- Drafting Note: `key` in `store_key` below and "see above" need to be italics and not codefont.

```

constexpr void store_key(see above operand,
                        memory_order order = memory_order::seq_cst) const noexcept;

```

- Mandates: `T` is a complete object type.
- Preconditions: `order` is `memory_order_relaxed`, `memory_order_release`, or `memory_order_seq_cst`.
- Effects: Atomically replaces the value referenced by `*ptr` with the result of the computation applied to the value referenced by `*ptr` and the given operand. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([atomics.order] (<https://eel.is/c++draft/atomics.order>)).
- Remarks: The result may be an undefined address, but the operations otherwise have no undefined behavior.

For `store_max` and `store_min`, the maximum and minimum computation is performed as if by `max` and `min` algorithms ([alg.min.max]), respectively, with `*ptr` and the first parameter as the arguments.

[Note 1: If the pointers point to different complete objects (or subobjects thereof), the `<` operator does not establish a strict weak ordering (Table 29, [expr.rel]). — end note]

Add to [atomics.types.int] (<https://eel.is/c++draft/atomics.types.int>):

- Drafting Note: all `*-type` below should be in italic.

```
namespace std {
    template<> struct atomic<integral-type> {
        ...

        void store_add(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_add(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_sub(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_sub(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_and(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_and(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_or(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_or(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_xor(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_xor(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_max(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_max(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;
        void store_min(integral-type,
                     memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_min(integral-type,
                     memory_order = memory_order::seq_cst) noexcept;

    };
}
```

```
void store_key(integral-type operand,
               memory_order order = memory_order_seq_cst) volatile noexcept;
constexpr void store_key(integral-type operand,
               memory_order order = memory_order_seq_cst) noexcept;
```

- Constraints: For the volatile overload of this function, is_always_lock_free is true.

- Preconditions: `order` is `memory_order_relaxed`, `memory_order_release`, or `memory_order_seq_cst`.
- Effects: Atomically replaces the value pointed to by `this` with the result of the computation applied to the value pointed to by `this` and the given operand. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([atomics.order] (<https://eel.is/c++draft/atomics.order>)).
- Remarks: Except for `store_max` and `store_min`, for signed integer types the result is as if the value pointed to by `this` and parameters were converted to their corresponding unsigned types, the computation performed on those types, and the result converted back to the signed type.

[Note: There are no undefined results arising from the computation. — end note]

For `store_max` and `store_min`, the maximum and minimum computation is performed as if by `max` and `min` algorithms ([alg.min.max]), respectively, with the value pointed to by `this` and the first parameter as the arguments.

Add to [atomics.types.float] (<https://eel.is/c++draft/atomics.types.float>):

```

namespace std {
    template<> struct atomic<floating-point-type> {
        ...

        void store_add(floating-point-type,
                       memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_add(floating-point-type,
                                 memory_order = memory_order::seq_cst) noexcept;
        void store_sub(floating-point-type,
                       memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_sub(floating-point-type,
                                 memory_order = memory_order::seq_cst) noexcept;
        void store_max(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_max(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
        void store_min(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_min(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
        void store_fmaximum(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_fmaximum(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
        void store_fminimum(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_fminimum(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
        void store_fmaximum_num(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_fmaximum_num(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
        void store_fminimum_num(floating-point-type, memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_fminimum_num(floating-point-type, memory_order = memory_order::seq_cst) noexcept;
    };
}

```

P3008 modified [atomics.types.float] (<https://eel.is/c++draft/atomics#types.float-4>) as follows:

The following operations perform arithmetic addition and subtraction computations. The correspondence among key, operator, and computation is specified in Table 153 except for the keys `max`, `min`, `fmaximum`, `fminimum`, `fmaximum_num`, and `fminimum_num`, which are specified below.

```

void store_key(floating-point-type operand,
               memory_order order = memory_order::seq_cst) volatile noexcept;
constexpr void store_key(floating-point-type operand,
                         memory_order order = memory_order::seq_cst) noexcept;

```

- Constraints: For the `volatile` overload of this function, `is always lock free` is true.
- Preconditions: `order` is `memory_order relaxed`, `memory_order release`, or `memory_order seq_cst`.

- Effects: Atomically replaces the value pointed to by `this` with the result of the computation applied to the value pointed to by `this` and the given operand. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([`atomics.order`]) (<https://eel.is/c++draft/atomics.order>)).
- Remarks: If the result is not a representable value for its type ([`expr.pre`]) the result is unspecified, but the operations otherwise have no undefined behavior. Atomic arithmetic operations on `floating-point-type` should conform to the `numeric_limits<floating-point-type>` traits associated with the floating-point type ([`limits.syn`]). The floating-point environment ([`cfenv`]) for atomic arithmetic operations on `floating-point-type` may be different than the calling thread's floating-point environment. The arithmetic rules of floating-point atomic modify-write operations may be different from operations on floating-point types or atomic floating-point types.
[Note: Tree reductions are permitted for atomic modify-write operations. – end note]
- Remarks:
 - For `store fmaximum` and `store fminimum`, the maximum and minimum computation is performed as if by `fmaximum` and `fminimum`, respectively, with the value pointed to by `this` and the first parameter as the arguments.
 - For `store fmaximum num` and `store fminimum num`, the maximum and minimum computation is performed as if by `fmaximum num` and `fminimum num`, respectively, with the value pointed to by `this` and the first parameter as the arguments.
 - For `store max` and `store min`, the maximum and minimum computation is performed as if by `fmaximum num` and `fminimum num`, respectively, with the value pointed to by `this` and the first parameter as the arguments, except that:
 - If both arguments are NaN, an unspecified NaN value replaces the value pointed to by `this`.
 - If exactly one argument is a NaN, either the other argument or an unspecified NaN value replaces the value pointed to by `this`, it is unspecified which.
 - If the arguments are differently signed zeros, which of these values replaces the value pointed to by `this` is unspecified.
- Recommended practice: The implementation of `store max` and `store min` should treat negative zero as smaller than positive zero.

Add to [`atomics.types.pointer`] (<https://eel.is/c++draft/atomics.types.pointer>):

```

namespace std {
    template<class T> struct atomic<T*> {
        ...
        void store_add(ptrdiff_t,
            memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_add(ptrdiff_t,
            memory_order = memory_order::seq_cst) noexcept;
        void store_sub(ptrdiff_t,
            memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_sub(ptrdiff_t,
            memory_order = memory_order::seq_cst) noexcept;
        void store_max(T*,
            memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_max(T*,
            memory_order = memory_order::seq_cst) noexcept;
        void store_min(T*,
            memory_order = memory_order::seq_cst) volatile noexcept;
        constexpr void store_min(T*,
            memory_order = memory_order::seq_cst) noexcept;
    };
}

```

Drafting note: add the following to the end of the atomics.types.pointer.

```

void store_key(see above operand,
    memory_order order = memory_order::seq_cst) volatile noexcept;
constexpr void store_key(see above operand,
    memory_order order = memory_order::seq_cst) noexcept;

```

- Constraints: For the `volatile` overload of this function, `is always lock free` is true.
- Mandates: `T` is a complete object type.
[Note 1: Pointer arithmetic on `void*` or function pointers is ill-formed. — end note]
- Effects: Atomically replaces the value pointed to by `this` with the result of the computation applied to the value pointed to by `this` and the given operand. Memory is affected according to the value of `order`. These operations are atomic modify-write operations ([atomics.order] (<https://eel.is/c++draft/atomics.order>)).
- Remarks: The result may be an undefined address, but the operations otherwise have no undefined behavior.

For `store_max` and `store_min`, the maximum and minimum computation is performed as if by `max` and `min` algorithms ([alg.min.max]), respectively, with the value pointed to by `this` and the first parameter as the arguments.

[Note 2: If the pointers point to different complete objects (or subobjects thereof), the < operator does not establish a strict weak ordering (Table 29, [expr.rel]). — end note]

No acquire sequences support

Modify [atomics.fences] (<https://eel.is/c++draft/atomics.fences>) as follows:

33.5.11 Fences[atomics.fences]

1. This subclause introduces synchronization primitives called fences. Fences can have acquire semantics, release semantics, or both. A fence with acquire semantics is called an acquire fence. A fence with release semantics is called a release fence.
2. A release fence A synchronizes with an acquire fence B if there exist atomic operations X and Y, where Y is not an atomic modify-write operation ([atomics.order] (<https://eel.is/c++draft/atomics.order>)), both operating on some atomic object M, such that A is sequenced before X, X modifies M, Y is sequenced before B, and Y reads the value written by X or a value written by any side effect in the hypothetical release sequence X would head if it were a release operation.
3. A release fence A synchronizes with an atomic operation B that performs an acquire operation on an atomic object M if there exists an atomic operation X such that A is sequenced before X, X modifies M, and B reads the value written by X or a value written by any side effect in the hypothetical release sequence X would head if it were a release operation.
4. An atomic operation A that is a release operation on an atomic object M synchronizes with an acquire fence B if there exists some atomic operation X on M such that X is sequenced before B and reads the value written by A or a value written by any side effect in the release sequence headed by A.

Add the following to [atomics.order] (<https://eel.is/c++draft/atomics.order>) after the definition of read-modify-write operations (<https://eel.is/c++draft/atomics.order#10>):

1. An *atomic modify-write operation* is an atomic read-modify-write operation with weaker synchronization requirements as specified in [atomic.fences] (<https://eel.is/c++draft/atomics.fences>).

[Note: The intent is for atomic modify-write operations to be implemented using mechanisms that are not ordered, in hardware, by the implementation of acquire fences. No other semantic or hardware property (e.g., that the mechanism is a far atomic operation) is implied. - end note]

Add `__cpp_lib_atomic_reductions` version macro to `<version>` synopsis [version.syn] (<https://eel.is/c++draft/version.syn#2>):

`#define cpp_lib_atomic_reductions L // freestanding, also in <atomic>`