

# Atomic stores and object lifetimes

Doc. No: P3181R0

Authors: Hans Boehm, Dave Claussen, David Goldblatt

Contact: Hans Boehm (hboehm@google.com)

Audience: SG1

Date: Apr. 15, 2024

Target: C++26

## Abstract

The current object lifetime rules insist that the last update to an object happens-before destruction of the object. This is arguably too conservative in ways that the authors found surprising and undesirable. Most likely it is possible to relax the rule in a way that better aligns it with programmer intuition. Should we? Are there implementation gotchas?

[ This is essentially the draft version we discussed in Tokyo, for the record. There seemed to be general consensus to proceed, and that we need carefully worked out wording. There was no real consensus on direction for the wording. There is clearly work remaining to be done here. ]

## Introduction

The intent of the standard is that all accesses to an object should happen during its lifetime, i.e. construction should happen before the access and the access should happen before destruction.<sup>1</sup>

Now consider the case in which the final access is in Thread A, and it is a store to an `atomic<bool>`. Thread B performs an acquire load on the same atomic and then destroys the object, after which it gets deallocated, or goes out of scope.

According to the standard, this is fine if the store in A is a release-store. It is not fine if the store to A is a release fence, followed by a relaxed store, since in that case it is the fence, not the store, that synchronizes with the acquire-load. In the release-store case, the store happens-before the deletion; in the fence case it does not.

---

<sup>1</sup> This paper will ignore accesses during construction and destruction, which is an important, but separable issue.

For something like the standard message passing litmus test the difference between “fence; relaxed-store” is usually stronger than “release-store”, and the difference often doesn’t matter. Here it is the other way around.

## Detailed example

We can write the version with the release-store as follows. Here we also include a third thread C that subsequently happens to allocate the same memory used for the previously described load and store. The third thread, though not very relevant from the standard’s perspective, is important for discussing possible resulting implementation issues.

Thread A:

```
fence(release);           // irrelevant for present purposes.
a1: a->store(1, relaxed);
```

Thread B:

```
b1: r1 = a->load(acquire); // sees 1.
b2: delete a;
```

Reallocating thread C:

```
c1: b = new atomic<int>(0); // gets the same location as a.
c2: r2 = b->load(relaxed);  // Should not see the 1 from thread A.
```

[res.on.objects]p2 states:

“If an object of a standard library type is accessed, and ... the access does not happen before the end of the object’s lifetime, the behavior is undefined unless otherwise specified.

[Note 2 : This applies even to objects such as mutexes intended for thread synchronization. — end note]”

The allocator must ensure that b2 happens-before c1. We need a1 happens-before c2, which just doesn’t follow. If we somehow promised that a1 did happen-before b1 would effectively turn `memory_order::relaxed` into `memory_order::release`, which is clearly unacceptable.

The current rules are backed by some possible intuition: [ Warning: This is a bit of a stretch. ] Consider a distributed shared-memory implementation that may nondeterministically send asynchronous updates, and has release operations send visible changes to acquire operations. (A release fence would ensure that all later acquire operations see changes.) b1 could happen to see a1’s update, but there is no guarantee the reallocating thread sees it before the allocation, since there is no communicating release-acquire pair involving the child thread.

## Stronger semantics are probably still fine

On the other hand, this can't really happen on implementations that compile the non-atomic initializing store in c1 in the same way as they would a relaxed atomic store. In that case, the modification order / coherence rules in the standard prohibit the above behavior.

We can roughly model the preceding example as:

Thread A:

```
a1: a.store(1, relaxed);
```

Thread B:

```
b1: r1 = a.load(relaxed); // sees 1
b2: b.store(1, release);  // Somewhere inside the delete implementation.
```

Thread C: (reallocating thread)

```
c1: r2 = b.load(acquire); // Somewhere inside allocator implementation.
c2: a.store(0, non-atomic); // initialization of new object, kind of
c3: a.load();              // Sees 1?
```

For now, let's treat c2 as a relaxed store.

If c2 follows a1 in the modification order, then, by [intro.races]p16, this can't happen, since b1 happens before c3, and thus c3 must see a later store.

So we would have to have a1 follow c2 in the modification order. But that would violate [intro.races]p17: "If a value computation A [b1] of an atomic object M [a] happens before an operation B [c2] that modifies M [a], then A[b1] takes its value from a side effect X on M, where X precedes B [c2] in the modification order of M."

This seems to be a fairly convincing argument that the problematic outcome in which the load at c2 sees 1 isn't possible on common implementations.

The only small fly in the ointment here is the non-atomic store at c2. The question then is: What implementations are there that implement a non-atomic store differently from a relaxed one? Any sort of compiler reordering seems implausible here, since in reality b2 and c1 must communicate a's address. The only architecture I know of that by design required stronger ordering for relaxed stores was Itanium. There were some errata for other old processors that sort of had this effect, but we don't know of anything modern enough to be a serious constraint. And we're not convinced that even those old processors would actually pose a problem here; it's just that the above safety argument does not directly apply.

## Why this matters

An important property of synchronization primitives generally is that they be able to "synchronize their own destruction". As an example: if an object has an embedded mutex protecting it, as well as a refcount, it should be valid for a thread to acquire the mutex, decrement the refcount, see that it became zero, release the mutex, and destroy the object.

(C++ and posix mutexes work this way, but historically some mutex implementations did not -- if your platform implements mutexes by embedding an OS handle alongside an atomic word there are very difficult to manage races in between unlocking a mutex and waking any waiters).

If fences aren't strong enough to make this work, "relaxed atomic + fence" suddenly loses the ability to provide this guarantee. That hurts the memory model sanity generally: "inserting seq\_cst fences before and after relaxed accesses gives seq\_cst semantics" is very intuitive, and true except for this edge case (related "very obviously should be true" claims break too, like "it's correct to replace store(release) with fence(release); store(relaxed)"). Fences end up with a weird safety restriction that's not motivated by the underlying hardware.

For example, a reference count decrement operation might be written as

```
int count = x->refcount.fetch_sub(1, acq_rel);
if (count == 1) delete x;
```

This becomes subtly incorrect when rewritten as

```
atomic_thread_fence(release);
int count = x->refcount.fetch_sub(1, relaxed);
atomic_thread_fence(acquire);
if (count == 1) delete x;
```

When you're writing some bit of synchronization you don't necessarily know how your caller will use the synchronization you provide, so you can't assume it *won't* be used to synchronize destruction. When you're using someone else's primitive, you don't know how they implemented their synchronization guarantees, and shouldn't have to in order to be able to safely destroy it.

## Towards a solution

The exact wording here is TBD, but the approach would be to relax the rule that the last access to an object happens-before the end of the object lifetime. One suggested starting point for the wording is:

"If there are atomic operations X and Y, where X modifies some atomic object M, Y reads the value in the hypothetical release sequence X would head if X were a release operation, and Y happens before the end of M's lifetime, then X's access of M is exempt from the requirement in [res.on.objects] that the access happens before the end of M's lifetime."

We suspect this could safely be generalized further, by only insisting that X be coherence-ordered before Y, rather than insisting that Y reads the value written by X. But it is unclear this provides enough added utility to make up for the added wording issues and implementation risk.

The original, more restrictive, version could probably also be phrased as relaxing the previously quoted [res.on.objects]p2 to a more nicely phrased version of:

"If an object of a standard library type is accessed, and ... the access does not happen before the end of the object's lifetime and, if the access is an update, no evaluation both observes the update and happens before the end of the object's lifetime, the behavior is undefined unless otherwise specified."

[Note 2 : This applies even to objects such as mutexes intended for thread synchronization. — end note]"