

Making `std::basic_istream::ignore` less surprising

- - [Revision History](#)
 - [Introduction](#)
 - [Discussion](#)
 - [Possible solutions](#)
 - [Change `ignore` to handle negative `delim` values](#)
 - [Split `ignore` into two functions and change `delim` to `char\_type`](#)
 - [Add an overload of `ignore` that takes a `char\_type`](#)
 - [As above, but constrain the new overload to prevent ambiguities](#)
 - [Do nothing](#)
 - [Proposed wording](#)
 - [Acknowledgements](#)
 - [References](#)

Revision History

- Initial revision

Introduction

`std::istream::ignore(n, delim)` has surprising behaviour if `delim` is a `char` with a negative value. We can remove the surprise and make code more robust.

Discussion

A negative `delim` value passed to `std::istream::ignore` can never match any character in the stream's input sequence, because the comparison is done using `traits_type::eq_int_type(rdbuf()->sgetc(), delim)` and `sgetc()` never returns negative values (except at EOF). `sgetc()` extracts a character `c` from the input sequence and then calls `traits_type::to_int_type(c)` to convert it to `int_type`, which produces a non-negative value.

The `std::char_traits<char>::to_int_type(char c)` function converts the character `c` to a non-negative integer, as if by `(int)(unsigned char)c`. This allows the value `(int_type)-1` to be reserved for EOF without worrying about whether `char` is signed and whether `(char)-1` can equal EOF. But it means that any code dealing with raw `char` values from the stream must consistently use `to_int_type` to convert the (possibly negative) `char` value into a non-negative `int_type` so that all characters are represented in the same form and can be compared like-for-like.

So if a negative delimiter `char` can never match, this means that users should call `std::cin.ignore(n, std::char_traits<char>::to_int_type(c))` or `std::cin.ignore(n, (unsigned char)c)` in case `char` is signed on their platform, and `c` happens to have the most significant bit set, i.e., is a negative value. For example, on a typical x86_64 Linux system where `char` is signed, `std::cin.ignore(std::numeric_limits<std::streamsize>::max(), '\x80')` will always discard all input up to EOF, even if `\x80` is present in the input.

In generic code that works with streams of either `char` or `wchar_t` you need to use the more verbose `to_int_type` form rather than casting to `unsigned char`, because casting a `wchar_t` to `unsigned char` would be wrong. That's unfortunate, because `std::char_traits<wchar_t>::to_int_type` doesn't have the same "cast to unsigned" behaviour, and so passing a `wchar_t` directly to `std::wistream::ignore` without conversion works correctly. But to work with both `char` and `wchar_t`, generic code must assume the worst and defend against the signed `char` trap.

Even in non-generic code that only works with `char`, you still need to remember that the trap exists, and remember to avoid it. It's rare that I see anybody get that right for `std::isalpha(c)` et al, and I wasn't even aware of the need to do it for `ignore` until this week.

I suspect that most users are not aware of the need to use `to_int_type` here, which means that the very useful `ignore` function is surprisingly fragile. It's also quite ugly to have to cast or deal with the traits type directly in a high-level `istream` API like `ignore`, which is not a low-level `streambuf` member function. The other unformatted input functions that take a delimiter (`get` and `getline`) take a `char_type`, so they are agnostic about whether `char` is signed or unsigned, and any conversion to `int_type` is done by the stream, not expected to be done by the caller.

Possible solutions

Change `ignore` to handle negative `delim` values

We could modify the spec for `basic_istream::ignore` so that negative values (except for -1 which must be reserved for EOF) are automatically fed through `to_int_type` to "clean" them, so that they're in the same domain as the values returned by `sgetc()`.

The GNU C library takes this approach for its `<ctype.h>` functions, which are specified to take `int`, but which have undefined behaviour if the argument is not an `unsigned char` value or EOF. So `isalpha('\x80')` has undefined behaviour if `char` is signed. But with Glibc, it works. Values in the range [-128,-1) are handled as if converted to `unsigned char` automatically, so that negative `char` values don't produce undefined behaviour. Obviously it's still possible to misuse those functions, e.g., by passing an `int` value outside the range of `char` or `unsigned char`, e.g., `isalpha(1000)`. But the apparently simple `isalpha(c)` for a `char` value isn't undefined just because `c` happens to be a negative value of a signed type.

Taking the same approach for `ignore` would remove the trap for cases like `cin.ignore(n, '\x80')`, making it behave as `cin.ignore(n, std::char_traits<char>::to_int_type('\x80'))`. However, it would also "fix" cases like `cin.ignore(n, -10)` which are less likely to be correct. There would be no way to tell the difference between a `char` with the value -10 and an `int` with the value -10, but the latter seems odd, and possibly a bug. Depending how we specified it, this solution might also give defined behaviour to `cin.ignore(n, +1000)` and `cin.ignore(n, -9999)` which are just nonsense.

The other downside of this solution is that it only fixes negative delimiters less than or equal to -2, because -1 still has to be reserved to mean EOF. So on a platform where `char` is signed, `cin.ignore(n, '\xfe')` would work, but `cin.ignore(n, '\xff')` would not, because that value is `traits_type::eof()`. On a platform where `char` is unsigned, both work. So this solution removes *most* surprises, but not all, and doesn't have portable guarantees. Users should really still use `in.ignore(n, std::char_traits<char>::to_int_type(c))` to work with arbitrary delimiter characters on arbitrary platforms.

Split `ignore` into two functions and change `delim` to `char_type`

We could change `delim` to `char_type` and make `ignore` convert that to `int_type` internally by using `to_int_type`.

```
basic_istream& ignore(streamsize n = 1, int_type = traits_type::eof());  
basic_istream& ignore(streamsize n, char_type delim);
```

This would preserve the same behaviour for `in.ignore(n)`, i.e., ignore up to `n` characters or up to EOF, whichever happens first. But it would break explicit uses of `eof` as the second argument, e.g., `in.ignore(n, traits_type::eof())`. This would implicitly convert the `eof()` value to `char_type` and treat it as a delimiter. If `(char)-1` happens to be present in the next `n` characters of the input sequence, it would match and we would stop ignoring too soon. This would break too much code, and we can do better.

Add an overload of `ignore` that takes a `char_type`

We don't need to change the existing `ignore`, we can just add an overload that does the correct conversion to `int_type`:

```
basic_istream& ignore(streamsize n, char_type delim)  
{ return ignore(n, traits_type::to_int_type(delim)); }
```

This does exactly what users expect when calling `ignore(n, c)` with a `char_type` argument (even `(char)-1`), with no alarms and no surprises. The behaviour is entirely consistent for signed or unsigned `char` and always matches the given `delim` if it occurs in the input sequence.

The downside of this solution is that calls that pass a delimiter that is neither `int_type` nor `char_type` become ambiguous, e.g. `std::cin.ignore(n, 1ULL)` is valid today but would become ambiguous with this new overload. Arguably, that's a good thing. What is this code trying to do? What if it passes a value that doesn't even fit in `int_type`, e.g. `numeric_limits<int_type>::max()+1LL`? Maybe it's good for such calls to not compile.

Since the problem only exists for `char` istreams and not `wchar_t`, this new overload could be specified as present only when `char_type` is `char`. That would avoid introducing any ambiguities for `std::wistream`.

As above, but constrain the new overload to prevent ambiguities

```
basic_istream& ignore(streamsize n, same_as<char_type> auto delim)  
{ return ignore(n, traits_type::to_int_type(delim)); }
```

This will only be selected by overload resolution when called with an argument of type `char_type`. For all other argument types, the existing overload will be selected and if the argument isn't of type `int_type` it will implicitly convert to `int_type`, exactly as happens today.

This doesn't change the meaning of any existing code, except for calls with negative `char_type` values which would start to work as users probably expected them to all along. The downside is the additional complexity of using a constrained overload, which would need to be emulated with SFINAE if vendors wanted to backport this fix to older standards modes.

Personally, I think the non-constrained overload is the best option, and that the cases which become ambiguous should probably be fixed to clarify what they're intending to do. Making the conversion to `int_type` explicit (and using `to_int_type` if appropriate) would probably be an improvement.

Do nothing

Boo! Users deserve better. Well, most of them. Some don't.

Proposed wording

Modify the class synopsis in [istream.general] as shown:

```
// [istream.unformatted], unformatted input
streamsize gcount() const;
int_type get();
basic_istream& get(char_type& c);
basic_istream& get(char_type* s, streamsize n);
basic_istream& get(char_type* s, streamsize n, char_type delim);
basic_istream& get(basic_streambuf<char_type, traits>& sb);
basic_istream& get(basic_streambuf<char_type, traits>& sb, char_type delim);

basic_istream& getline(char_type* s, streamsize n);
basic_istream& getline(char_type* s, streamsize n, char_type delim);

basic_istream& ignore(streamsize n = 1, int_type delim = traits::eof());
basic_istream& ignore(streamsize n, char_type delim);
int_type peek();
basic_istream& read(char_type* s, streamsize n);
streamsize readsome(char_type* s, streamsize n);
```

Modify [istream.unformatted] as shown:

```
basic_istream& ignore(streamsize n = 1, int_type delim = traits::eof());
```

-25- *Effects*: Behaves as an unformatted input function (as described above). After constructing a sentry object, extracts characters and discards them. Characters are extracted until any of the following occurs:

- (25.1) — `n != numeric_limits<streamsize>::max()` ([numeric.limits]) and `n` characters have been extracted so far;
- (25.2) — end-of-file occurs on the input sequence (in which case the function calls `setstate(eofbit)`, which may throw `ios_base::failure` ([iostate.flags]));
- (25.3) — `traits::eq_int_type(traits::to_int_type(c), delim)` for the next available input character `c` (in which case `c` is extracted).

[*Note 1*: The last condition will never occur if `traits::eq_int_type(delim, traits::eof())`. — *end note*]

-26- *Returns*: `*this`.

```
basic_istream& ignore(streamsize n = 1, char_type delim);
```

```
-?- Effects: Equivalent to: return ignore(n, traits type::to_int_type(delim));
```

Acknowledgements

Thanks to Iain Sandoe and Ulrich Drepper for inspiring this.

References

[N4971](#), Working Draft - Programming Language -- C++, Thomas Köppe, 2023.