

P3351R4

views :: scan

Published Proposal, 2026-04-27

This version:

<https://wg21.link/P3351R4>

Author:

[Yihe Li](#)

Audience:

LEWG

Project:

ISO/IEC 14882 Programming Languages — C++, ISO/IEC JTC1/SC22/WG21

Target:

C++29

Source:

github.com/Mick235711/wg21-papers/blob/main/P3351/P3351R4.bs

Issue Tracking:

[GitHub Mick235711/wg21-papers](#)

Table of Contents

- 1 Revision History**
 - 1.1 R4 (2026-05 pre-Brno mailing)
 - 1.2 R3 (2025-10 pre-Kona mailing)
 - 1.3 R2 (2025-01 pre-Hagenberg Mailing)
 - 1.4 R1 (2024-10 pre-Wrocław Mailing)
 - 1.5 R0 (2024-07 post-St. Louis Mailing)
- 2 Motivation**
- 3 Design**

- 3.1 What Is Being Proposed?
- 3.2 Background
- 3.3 Prior Art
- 3.4 Alternative Names
- 3.5 Left or Right Fold?
 - 3.5.1 SG9 Review, Wrocław (2024-11)
- 3.6 More Convenience Aliases
 - 3.6.1 SG9 Review, Hagenberg (2025-02)
- 3.7 Range Properties
 - 3.7.1 Reference and Value Type
 - 3.7.2 Category
 - 3.7.2.1 SG9 Review, Hagenberg (2025-02)
 - 3.7.3 Common
 - 3.7.4 Sized
 - 3.7.5 Const-Iterable
 - 3.7.6 Borrowed
 - 3.7.6.1 SG9 Review, Hagenberg (2025-02)
- 3.8 Feature Test Macro
- 3.9 Freestanding
- 3.10 Future Extensions

4 Implementation Experience

5 Questions To Resolve

- 5.1 Questions for SG9
 - 5.1.1 SG9 Review, Hagenberg (2025-02)
- 5.2 Questions for LEWG

6 Wording

- 6.1 17.3.2 Header `<version>` synopsis [version.syn]
- 6.2 25.2 Header `<ranges>` synopsis [ranges.syn]
- 6.3 25.7.❖ Scan view [range.scan]
- 6.4 25.7.❖.1 Overview [range.scan.overview]
- 6.5 25.7.❖.2 Class template `scan_view` [range.scan.view]
- 6.6 25.7.❖.3 Class template `scan_view::iterator` [range.scan.iterator]

7 Poll Results

- 7.1 SG9 Review, Wrocław (2024-11)
- 7.2 SG9 Review, Hagenberg (2025-02)
- 7.3 SG9 Review, Kona (2025-11)

References

Normative References

Informative References

This paper proposes the `views::scan` range adaptor, which takes a range and a function that takes the current element *and* the current state as parameters. Basically, `views::scan` is a lazy view version of `std::inclusive_scan`, or `views::transform` with a stateful function.

The `views::scan` adaptor is classified as a Tier 1 item in the Ranges plan for C++26 ([\[P2760R1\]](#)).

§ 1. Revision History

§ 1.1. R4 (2026-05 pre-Brno mailing)

- In Kona (2025-11), SG9 sees [\[P3351R3\]](#) and forwarded it to LEWG ([polls](#)). Audience is altered accordingly.
- Wording changes since [\[P3351R3\]](#) based on SG9 feedback:
 - Removed impossible requirements on `forward_range` in iterator's postfix `operator++`.
 - Changed `bool IsInit` parameter to an `enum class` for clarity
 - Actually implement borrowed range semantics
- Removed no longer relevant sections from the paper

§ 1.2. R3 (2025-10 pre-Kona mailing)

- Retargeted to C++29, given time constraints.
- Removed `partial_sum` alias as per SG9 guidance in Hagenberg (2025-02).
- Add some material in the transversal category section, including discussion on stashing iterators, move-only types, and the possible choices between unconditionally input or conditionally forward for `scan_view`.

- Add some more discussion on the feasibility of making `scan_view` conditionally borrowed.
- Add implementation experience in Beman Project.
- Rebase onto latest draft [\[N5014\]](#), which incorporates changes induced by [\[P2846R6\]](#).
 - Since [\[LWG4189\]](#) is now accepted, the freestanding section is obsolete and is modified accordingly.
- Editorial changes since [\[P3351R2\]](#):
 - Add poll results and comments from SG9 review in Hagenberg (2025-02).
 - Removed question on `operator-`; apparently only random access ranges need to provide that in the STL (see `filter_view`, for example).
 - Move naming questions to LEWG.

§ 1.3. R2 (2025-01 pre-Hagenberg Mailing)

- Redesigned the entire paper, following SG9's feedback in Wrocław (2024-11). Significant changes since [\[P3351R1\]](#) are:
 - `views::scan` is now equivalent to `std::inclusive_scan`, both in terms of argument sequence (`rng`, `func`, `init`) and also the semantics.
 - `views::prescan` is removed as it did not gain enough support and have no prior experience in the STL.
 - The prior arts section is redesigned:
 - Added STL algorithm to prior arts section, and also reformat the table.
 - More languages like Kotlin and Haskell are added.
 - After the parameter sequence change, conflict between range and initial value is no longer an issue, thus removing this workaround which previously prevent the with-initial-value version to overload with the normal version. As a result, now `scan` and `partial_sum` supports an optional initial value.
 - A section is added to explore the relationship with parallel range algorithms in general.
 - Group the wording question into different groups based on audience.
- Added poll results section and several review results sections.
- Added some minor correction and discussion on reference type.
- Added some discussion on naming conflicts with [\[P1729R5\]](#).

- Several wording fixes:
 - Fixed the mismatch template arguments for `scan_view` in synopsis.
 - Add *italics* for exposition-only concepts.
- Preemptively add `reserve_hint` members in anticipation of [P2846R6]’s acceptance.
- Rebase onto latest draft [N5001].

§ 1.4. R1 (2024-10 pre-Wrocław Mailing)

- Several wording fixes:
 - Added the missing `noexcept` to `end()`.
 - Refactored the constraints of `scan_view` out into its own concept, such that it tests for both assignability from `range_reference_t<R>` and the invoke result of `f`.
 - Replace `regular_invocable` with `invocable`.
 - Pass by move in `scan_view`’s constructor and when invoking the function.
- Rebase onto latest draft [N4988].

§ 1.5. R0 (2024-07 post-St. Louis Mailing)

- Initial revision.

§ 2. Motivation

The motivation for this view is given in [P2760R1] and quoted below for convenience:

If you want to take a range of elements and get a new range that is applying `f` to every element, that's `transform(f)`. But there are many cases where you need a `transform` to that is stateful. That is, rather than have the input to `f` be the current element (and require that `f` be `regular_invocable`), have the input to `f` be both the current element and the current state.

For instance, given the range `[1, 2, 3, 4, 5]`, if you want to produce the range `[1, 3, 6, 10, 15]` - you can't get there with `transform`. Instead, you need to use `scan` using `+` as the binary operator. The special case of `scan` over `+` is `partial_sum`.

One consideration here is how to process the first element. You might want `[1, 3, 6, 10, 15]` and you might want `[0, 1, 3, 6, 10, 15]` (with one extra element), the latter could be called a `prescan`.

This adaptor is also present in `ranges-v3`, where it is called `views::partial_sum` with the function parameter defaulted to `std::plus{}`. However, as [P2760R1] rightfully pointed out, `partial_sum` is probably not suitable for a generic operation like this that can do much more than just calculating partial sum, similar to how `accumulate` is not a suitable name for a general fold. Therefore, the more generic `scan` name is chosen to reflect the nature of this operation. (More discussion on the naming are present in the later sections.)

§ 3. Design

§ 3.1. What Is Being Proposed?

```
std::vector vec{1, 2, 3, 4, 5, 4, 3, 2, 1};
std::println("{} ", vec | views::scan(std::plus{})); // [1, 3, 6, 10, 15, 19, 22, 25, 26]
std::println("{} ", vec | views::scan(std::plus{}, 10)); // [11, 13, 16, 20, 25, 30, 35, 40, 41]
std::println("{} ", vec | views::scan(ranges::max)); // [1, 2, 3, 4, 5, 5, 4, 3, 2]
std::println("{} ", vec | views::scan(ranges::max, 3)); // [3, 3, 3, 4, 5, 5, 4, 3, 2]
```

§ 3.2. Background

Scan is a pretty common operation in functional programming. On its simplest form, scanning a sequence `x0, x1, x2, ...` gives you back the sequence `x0, x0 + x1, x0 + x1 + x2, ...`, essentially doing a prefix sum. When generalizing `+` to an arbitrary binary function, scan become a form of "partial" fold, in which fold-like operation is performed but with intermediate result preserved.

There are two forms of scan defined in a mathematical sense: *inclusive* scan and *exclusive* scan. An inclusive scan includes *i*-th term of the original sequence when computing the *i*-th term of the resulting sequence, while the exclusive scan does not. Inclusive scan may optionally include an initial seed *s* that will be used to compute the first output term, while exclusive scan requires an initial seed as the first output term directly. When summarized in a table, both forms of scan looks like this:

Original	<i>x</i> 0	<i>x</i> 1	<i>x</i> 2	<i>x</i> 3
Inclusive	<i>x</i> 0	<i>f</i> (<i>x</i> 0, <i>x</i> 1)	<i>f</i> (<i>f</i> (<i>x</i> 0, <i>x</i> 1), <i>x</i> 2)	<i>f</i> (<i>f</i> (<i>f</i> (<i>x</i> 0, <i>x</i> 1), <i>x</i> 2), <i>x</i> 3)
Inclusive With Seed	<i>f</i> (<i>s</i> , <i>x</i> 0)	<i>f</i> (<i>f</i> (<i>s</i> , <i>x</i> 0), <i>x</i> 1)	<i>f</i> (<i>f</i> (<i>f</i> (<i>s</i> , <i>x</i> 0), <i>x</i> 1), <i>x</i> 2)	<i>f</i> (<i>f</i> (<i>f</i> (<i>f</i> (<i>s</i> , <i>x</i> 0), <i>x</i> 1), <i>x</i> 2), <i>x</i> 3)
Exclusive	<i>s</i>	<i>f</i> (<i>s</i> , <i>x</i> 0)	<i>f</i> (<i>f</i> (<i>s</i> , <i>x</i> 0), <i>x</i> 1)	<i>f</i> (<i>f</i> (<i>f</i> (<i>s</i> , <i>x</i> 0), <i>x</i> 1), <i>x</i> 2)

We can see clearly that both forms of scan are easily interconvertible, as shifting the inclusive scan result right once, prepend the initial seed, and drop the last term gives you the exclusive scan result. Another critical characteristic of exclusive scan is that they disregard the last input term to preserve the length of the sequence.

In light of this, a third, more general scan semantics was implemented by some languages (like Haskell's `scanl`): Prepend the initial seed, but does not drop the last term. This is called `prescan` semantics in previous revisions of this proposal, and has the notable characteristic that the resulting sequence will be one longer than the input. When applied to the sequence `[1, 2, 3, 4, 5]` and using addition as the binary function, the three scan semantics gives the following results:

Original	<code>[1, 2, 3, 4, 5]</code>
Inclusive	<code>[1, 3, 6, 10, 15]</code>
Inclusive With Seed 10	<code>[11, 13, 16, 20, 25]</code>
Exclusive With Seed 10	<code>[10, 11, 13, 16, 20]</code>
<code>prescan</code> With Seed 10	<code>[10, 11, 13, 16, 20, 25]</code>

C++ STL provides `std::partial_sum` and `std::inclusive_scan` (latter since C++17) with the inclusive semantics, and `std::exclusive_scan` (since C++17) with the exclusive semantics. Previous revisions of this proposal proposes both the inclusive and the `prescan` semantics, while the current revision only proposes the inclusive semantics. As can be seen in the Prior Arts sections below, inclusive semantics is the most commonly used and provided semantic in most languages, and other

two semantics can be easily constructed by manipulating the inclusive result with existing range adaptors:

```
exclusive_scan(rng, func, init) = concat(single(init), scan(rng, func, init))
prescan(rng, func, init) = concat(single(init), scan(rng, func, init))
```

Therefore, the author thinks that providing only the inclusive semantics is sufficient.

§ 3.3. Prior Art

The scan adaptor/algorithm had made an appearance in many different libraries and languages:

- range-v3 has a `views::partial_sum` adaptor that don't take initial seeds, but takes arbitrary function parameter (defaults to `+`).
 - range-v3 also has an `views::exclusive_scan` adaptor that implements exclusive scan semantics. Its initial seed parameter is mandatory, but also takes an optional function parameter.
- STL has `std::partial_sum` algorithm in `<numeric>` since C++98, and `std::[in,ex]clusive_scan` algorithms in `<numeric>` since C++17.
- `tl::ranges` also has an implementation of `views::partial_sum`.
- Python has an `itertools.accumulate` algorithm that optionally takes initial seeds (with a named argument), and takes arbitrary function parameter (defaults to `+`).
 - Furthermore, NumPy provides `cumsum` algorithm that returns a partial sum (without function parameter), and an `ufunc.accumulate` algorithm that performs arbitrary scans with arbitrary function parameter that have no default. Neither of those algorithms takes an initial seed.
- Rust has an `iter.scan` adaptor that takes initial seeds and arbitrary function parameters (no defaults provided). Kotlin also provides a `scan` member that works similarly.
- Haskell provides two variants of left scan: `scanl1` and `scanl`. The former performs an inclusive scan, but does not permit an initial seed. The latter performs a `prescan`-like operation that preserve both the initial seed provided and the last fold result, resulting in a one-larger range than passed in.

Summarized in a table (without additional note, all functions provide inclusive scan semantics):

Library	Signature	Parameter	Function	With Default	Initial Seed
Lazy Adaptors					
range-v3 / <code>tl::ranges</code>	<code>partial_sum</code>	<code>(rng[, fun])</code>	✓	✓	✗
	<code>exclusive_scan</code> (Exclusive semantics)	<code>(rng, init[, fun])</code>	✓	✓	✓
Proposed	<code>scan</code>	<code>(rng, func[, init])</code>	✓	✗	✓
Eager Algorithms					
STL	<code>std::partial_sum</code>	<code>(rng, out[, func])</code>	✓	✓	✗
	<code>std::inclusive_scan</code>	<code>(rng, out[, func[, init]])</code>	✓	✓	✓
	<code>std::exclusive_scan</code> (Exclusive semantics)	<code>(rng, out, init[, func])</code>	✓	✓	✓
Python <code>itertools</code>	<code>accumulate</code>	<code>(rng[, fun[, init=init]])</code>	✓	✓	✓
NumPy	<code>cumsum</code>	<code>(rng)</code>	✗	N/A	✗
	<code><func>.accumulate</code>	<code>(rng)</code>	✓	✗	✗
Rust	<code>iter.scan</code>	<code>(init, func)</code>	✓	✗	✓
Kotlin	<code><rng>.scan</code>	<code>(init, func)</code>	✓	✗	✓
Haskell	<code>scanl1</code>	<code>(func, rng)</code>	✓	✗	✗
	<code>scanl</code> (<code>prescan</code> semantics)	<code>(func, init, rng)</code>	✓	✗	✓

§ 3.4. Alternative Names

The author thinks `views::scan` is a pretty good name, which have prior example in `std::[in,ex]clusive_scan` and in Rust/Kotlin, and is a generic enough name that will not cause confusion.

However, one possible naming conflict recently has arisen: `std::scan` ([P1729R5]). Here, the name `scan` refers to a completely different concept: scanning user's input, and `std::scan` is proposed as

an upgraded version of the `scanf` family and uses `std::format`-like syntax to parse variables from a given string input:

```
auto result = std::scan<std::string, int>("answer = 42", "{} = {}");
const auto& [key, value] = result->values();
// key == "answer", value == 42
```

As noted above, this interpretation of the name `scan` also has precedence in the standard library: the `scanf` family, so this meaning also makes some sense.

Notice that `std::scan` and `views::scan` do not conflict with each other since these two names are in different namespaces. However, it may be of interest to LEWG to rename `views::scan` to a different name in order to avoid user confusion about two unrelated facilities being named the same. In light of this, the following list of alternative names is provided for consideration:

Alternative names considered for `views::scan`: (in order of decreasing preference)

- `views::partial_fold` (suggested by [P2214R2], and makes the connection with `ranges::fold` clear): Pretty good name, since `scan` is just a `fold` with intermediate state saved. However, the correct analogy is actually `ranges::fold_left_first`, and `ranges::fold_left` actually corresponds to an exclusive scan, which the author felt may cause some confusion.
- `views::inclusive_scan`: Makes the equivalence with the STL numeric algorithm clear, and have no conflict potential with `scanf` anymore. However, the author felt that this name is a bit too long.
- `views::accumulate`, `views::fold`, `views::fold_left`: The output is a range instead of a number, so using the same name probably is not accurate. The latter two also suffer from the same displaced correspondence problem.

Overall, the author doesn't find any of these options particularly intriguing and opts to propose the original naming of `scan`. This does not actually conflict with `std::scan`, and besides, the standard library is already adopting two meanings of `scan` by introducing both `scanf` and `std::[in,ex]clusive_scan`, so it shouldn't be that much of a confusion risk to the users.

Another consideration worth pointing out is the consistency with the Range-ified versions of the existing `<numeric>` algorithms. [P3732R2] proposed `ranges::{in,ex}clusive_scan` as the newer version of the existing `std::{in,ex}clusive_scan` algorithms; if consistency in naming is desired, `views::inclusive_scan` may be a more appropriate name. However, the author leans towards naming the new algorithm `ranges::scan` to gain consistency instead, to provide convenience to the users.

§ 3.5. Left or Right Fold?

Theoretically, there are two possible direction of scanning a range:

```
// rng = [x1, x2, x3, ...]
scan_left(rng, f, i) // [x1, f(x1, x2), f(f(x1, x2), x3), ...]
scan_right(rng, f, i) // [x1, f(x2, x1), f(x3, f(x2, x1)), ...]
```

Both are certainly viable, which begs the question: Should we provide both?

On the one hand, `ranges::fold` provided both the left and the right fold version, despite the fact that right fold can be simulated by reversing the range and the function parameter order. However, here, the simulation is even easier: just reversing the order of the function parameters will turn a left scan to a right scan.

Furthermore, all of the mentioned prior arts perform left scan, and it is hard to come up with a valid use case of right scan that cannot be easily covered by left scan. Therefore, the author only proposes left scan in this proposal.

§ 3.5.1. SG9 Review, Wrocław (2024-11)

SG9 agreed that only left scan is needed, since a right scan can be easily simulated by flipping the arguments of the function provided. The fundamental difference between `fold` and `scan` here is that `fold` requires both reversing the range *and* flipping the arguments when changing from left to right, while `scan` only requires the latter.

§ 3.6. More Convenience Aliases

Obviously, there are some common aliases that can be provided on top of `scan`. The four most useful aliases are:

```
std::vector<int> vec{3, 4, 6, 2, 1, 9, 0, 7, 5, 8}
partial_sum(vec) // [3, 7, 13, 15, 16, 25, 25, 32, 37, 45]
partial_product(vec) // [3, 12, 72, 144, 144, 1296, 0, 0, 0, 0]
running_min(vec) // [3, 3, 3, 2, 1, 1, 0, 0, 0, 0]
running_max(vec) // [3, 4, 6, 6, 6, 9, 9, 9, 9, 9]
```

Which are the results of applying `std::plus{}`, `std::multiplies{}`, `ranges::min`, and `ranges::max` as the functor to be passed to `views::scan`, respectively.

Looking at the results, these are certainly very useful aliases. However, as stated above, all of these four aliases can be achieved by simply passing the corresponding function object (currently, `min/max` doesn't have a function object, but `ranges::min` and `ranges::max` become useable after [P3136R1] was adopted in Wrocław (2024-11)) as the function parameter, so I'm not sure it is worth the hassle of specification. Therefore, currently this proposal do not include these aliases, but the author is happy to add them should SG9/LEWG request.

R2 of this paper originally proposed the `partial_sum` alias, as there is certainly a stronger case for it than the case for `product`, `min` and `max`:

- All existing implementation of scan-like algorithm defaults the function argument to `+` whenever there is a default.
- `std::partial_sum` exists
- Partial sum is one of the most studied and used concept in programming, arguably even more useful than running `min/max`.

It was also considered whether `partial_sum` should be a standalone adaptor alias, or should it be folded into `scan` as the default function parameter. On the one hand, range-v3's `views::partial_sum` does this folding by using a function parameter defaults to `std::plus{}`. On the other hand, the author thinks that this will definitely cause some confusion to the users, due to the fact that a generic name like `scan` defaults to `std::plus` as its function parameter:

```
vec | views::scan // what should this mean?
```

This is similar to the scenario encountered by `ranges::fold` ([P2322R6]), where despite the old algorithm `std::accumulate` took `std::plus` as the default function parameter, `ranges::fold` still choose to not have a default. The author felt that the same should be done for `views::scan` (i.e. no default function), and introduced a separate `views::partial_sum` alias that more clearly convey the intent.

§ 3.6.1. SG9 Review, Hagenberg (2025-02)

However, the guidance and comments received during SG9 review of [P3351R2] in Hagenberg (2025-02) suggested that the `partial_sum` alias is not needed. Instead, forcing users to specify `std::plus{}` explicitly convey the intent of the operation more clearly. Therefore, the suggested aliases was dropped.

§ 3.7. Range Properties

The proposed `scan_view` is a range adaptor, i.e. can be piped into.

§ 3.7.1. Reference and Value Type

Consider the following:

```
std::vector<double> vec{1.0, 1.5, 2.0};  
vec | views::scan(std::plus{}, 1);
```

Obviously, we expect the result to be `[2.0, 3.5, 5.5]`, not `[2, 3, 5]`, therefore for `scan` we cannot just use the initial seed's type as the resulting range's reference/value type.

There are two choices we can make: (for input range type `Rng`, function type `F` and initial seed type `T`)

1. Just use the reference/value type of the input range. This is consistent with `std::partial_sum`. (range-v3 also chose this approach, using the input range's value type as the reference type of `scan_view`)
2. Be a bit clever, and use `decay_t<invoke_result_t<F&, T, ranges::range_reference_t<Rng>>>`. In other words, the return type of `func(init, *rng.begin())`.

Note that the second option don't really covers all kind of functions, since it is entirely plausible that `func` will change its return type in every invocation. However, this should be enough for nearly all normal cases, and is the decision chosen by [\[P2322R6\]](#) for `ranges::fold_left[_first]`. For `scan` without an initial seed, this approach can also work, by returning the return type of `func(*rng.begin(), *rng.begin())`.

Although the second choice is a bit complex in design, it also avoids the following footgun:

```
std::vector<int> vec{1, 4, 2147483647, 3};  
vec | views::scan(std::plus{}, 0L);
```

With the first choice, the resulting range's value type will be `int`, which would result in UB due to overflow. With the second choice the resulting value type will be `long` which is fine. (Unfortunately, if you omit the initial seed here it will still be UB, and that cannot be fixed.)

The second choice also enables the following use case:

```
// Assumes that std::to_string also has an overload for std::string that just
// concatenates two strings
std::vector<int> vec{1, 2, 3};
vec | views::scan(
    [](const auto& a, const auto& b) { return std::to_string(a) + std::to_string(b); },
    "2"
) // ["21", "212", "2123"]
```

Given that `ranges::fold_left[_first]` chose the second approach, the author thinks that the second approach is the correct one to pursue. In other words, the value type of `scan_view` will be `decay_t<invoke_result_t<F&, T, ranges::range_reference_t<Rng>>>`, and the reference type will be `const value_type&`.

§ 3.7.2. Category

Always input.

The resulting range cannot be bidirectional, since the function parameter cannot be applied in reverse. A future extension may enable a `scan` with both forward and backward function parameter, but that is outside the scope of this paper.

However, making `scan_view` forward also have its consequences. To support ranges with move-only element types, `scan_view::iterator::operator*` must return a *reference* to the underlying accumulator, which makes it a stashing iterator (an iterator that returns a reference to something within itself). [\[iterator.concept.forward\]/3](#) specifies that stashing iterators are only allowed to be input, so we have a difficult choice:

- Either, accept that `scan_view` is input only, and returns a reference in `operator*`.
- Or, make `scan_view` at most forward for range with copyable element types, but copy the underlying accumulator every time in `operator*`.

Note that ranges-v3's `partial_sum_view` does not support move-only element types at all, and thus copy the accumulator every time to support forward iterators. `tl::ranges` on the other hand incorrectly marked a stashing iterator implementation as forward.

§ 3.7.2.1. SG9 Review, Hagenberg (2025-02)

None of the two options are particularly appealing, but the author and SG9 ultimately thinks that copying the accumulator every time you dereference the iterator is too surprising as a hidden cost. Therefore, the author decided to make `scan_view` input only.

Theoretically we can use some type traits similar to `is_trivially_copyable` to determine a criteria for when to copy the accumulator, but the author feels that depending a fundamental property of the iterator on such difference is hard to understand for the users.

§ 3.7.3. Common

Never.

This is consistent with range-v3's implementation. The reasoning is that each iterator needs to store both the current position and the current partial sum (generalized), so that the `end()` position's iterator is not readily available in O(1) time.

§ 3.7.4. Sized

If and only if the input range is sized. The reported size is always equal to the input range's size.

Also, given that [P2846R6] is adopted for C++26 in Hagenberg (2025-02), `scan_view` also adapted that proposal by adding `reserve_hint()` members who forwards the underlying view's `reserve_hint()` result appropriately.

§ 3.7.5. Const-Iterable

Similar to `views::transform`, if and only if the input range is const-iterable and `func` is `const`-invocable.

§ 3.7.6. Borrowed

If and only if the input range is borrowed, and the function parameter models `tidy-func`.

Currently, `views::transform` is never borrowed, but after [P3117R1] determined suitable criteria for storing the function parameter in the iterator, it can be conditionally borrowed. Therefore, the same

can be applied to `scan_view` to make it conditionally borrowed by storing the function and the initial value inside the iterator.

§ 3.7.6.1. SG9 Review, Hagenberg (2025-02)

SG9 had consensus that `scan_view` should be conditionally borrowed similar to `transform_view`. The wording is thus modified to be rebased on [P3117R1]’s changes.

§ 3.8. Feature Test Macro

This proposal added a new feature test macro `__cpp_lib_ranges_scan`.

§ 3.9. Freestanding

[P1642R11] included nearly everything in `<ranges>` into freestanding, and [LWG4189] (accepted at Hagenberg (2025-02)) make the few leftover cases such as `views::cache_latest` freestanding too. Currently, everything in `<ranges>` except `ranges::elements_of`, `views::istream` and `ranges::[basic_]istream_view` is marked freestanding, and due to the blanket `// mostly freestanding` new additions into the header will be automatically made freestanding, which is the desired behavior for `views::scan`.

However, [LWG4189] failed to mark `__cpp_lib_ranges_enumerate` as freestanding, which the author assumes is an oversight (and submitted [LWG4286] to rectify). This paper thus marks its feature-test macro as freestanding as usual.

§ 3.10. Future Extensions

Several future extensions to the proposed `scan` adaptor are possible and seen in other languages. These are outside the scope of this proposal and are only listed here as a record for inspiration.

One obvious extension is to provide a variant that allows functions that operate on indexes, similar to Kotlin’s `scanIndexed` `member`:

```
>>> scanIndexed(["a", "b", "c"], "start", lambda index, acc, elem: acc + f'
["start, 0: a", "start, 0: a, 1: b", "start, 0: a, 1: b, 2: c"]
```


This is already achievable in C++ by using `views::enumerate` and manually unpacking the tuple elements inside the function. A direct `scan_indexed` adaptor would be more convenient but ultimately does not provide new functionalities that cannot be achieved using existing tools, so it is not proposed.

Another extension concerns the parallelism potential for the scan algorithm. At first glance, scan seems to be a completely serial algorithm that cannot be paralleled, as each term depends on the previous intermediate result. However, by iteratively performing the scanned function, the algorithm can actually be paralleled and can be completed in logarithmic time regarding to the number of elements when parallel resources are abundant. See [the Wikipedia page](#) for details.

In light of this, concerns are raised about the proposed `scan` adaptor in that it does not permit this kind of reorder for paralleled execution:

```
// Previous: two passes, two launches, each paralleled
ranges::inclusive_scan(par, rng, rng);
ranges::transform(par, rng, some_func);

// After: one pass but actually serial
ranges::transform(par, rng | views::scan(std::plus{}), some_func);
```

This is due to the fact that `scan_view` is at most forward, and thus, all the scanning operations can only be performed in sequence.

However, the author is unable to find a satisfactory solution to this problem. With the C++ iterator model, efficient paralleled algorithm execution is **only** possible with random access iterator/range inputs, as referring to an arbitrary point in the iteration space at a constant time is often an integral part of the algorithm. Although C++17 parallel algorithms generally only require forward iterators, in practice, most implementations (like libstdc++ and libc++) will *de facto* require random access iterators to enable parallelism and silently fall back to serial execution with lower categories. Existing third-party parallel libraries like OpenMP and oneTBB also nearly always require random access iterators in their algorithms.

Yet, it is impossible to make `scan_view` a random access range. Doing so will not only require a way to get the previous results (probably through an additional reversal functor) but also requires efficient computation of the N-th element in the output range, which by definition requires N execution of the function and thus is impossible to be made efficient. The parallel algorithm for scan requires a substantially different iterator/adaptor model that allows efficient and arbitrary chunking of inner operations, which the C++ iterator model does not permit.

This is a similar problem faced by `views::filter`. On its own, `filter_view` is at most bidirectional regardless of inputs, which makes it impossible to be executed efficiently using standard parallel algorithms and most third-party libraries as it is never random access. However, a natural parallel algorithm exists for `filter_view`: just filter whatever element you are processing in parallel and exit the current thread/worker if the filter returns `false`. Unfortunately, this is not expressible as the iterator advancement is strictly tied to the filter operation.

There are several possible mitigation strategies for this problem, but none of them are complete:

- Add named algorithms. For instance, given that `filter_view` is not parallelized, C++17 provided a parallel `copy_if` to provide the desired semantics. However, algorithms are not composable (which is the whole reason that range adaptors are introduced), and using parallel STL algorithms like `copy_if` and `inclusive_scan` in lieu of adaptors leads to suboptimal performance (due to `copy_if` needs some synchronization for writing the output while `filter` does not) and the multi-launch problem described above.
- Add named algorithms for compositions. For instance, to solve the transform + scan multi-launch problem, STL directly provides `transform_{in,ex}clusive_scan`, whose parallelized version can perform two operations in one launch. However, compositions are infinite, and STL cannot provide every composition as named algorithms (the above scan + transform operation, for instance, is not provided).
- Special-case the range adaptors. For instance, parallelized `ranges::transform` can identify the `scan_view` being passed in and utilize `base()` and functor members to do the parallelization. However, this solution is brittle, cannot handle complex pipelines, and requires substantial work for all algorithms.

Ultimately, the C++20 range adaptor is not a good abstraction for parallelism, and complex pipelines using adaptors are inherently only able to be executed serially. The parallelism problem is not unique to `scan_view`, and substantial changes to both the adaptor and iterator model (perhaps somehow separate the advancement and compute stages) are required to enable sufficient parallelism opportunities even for long pipeline compositions, which are best explored in a separate proposal. Fortunately, at least the current range parallel algorithm proposal [P3179R9] wisely restricted the input to be random access, so the above scan + transform invocation will result in a compile error instead of silently degrading to serial performance.

§ 4. Implementation Experience

The author implemented this proposal in `beman.scan_view` as part of Beman Project. No significant obstacles are observed.

§ 5. Questions To Resolve

§ 5.1. Questions for SG9

- Should the function parameter have a default? - Resolved ✓
- Should more convenience aliases be added? - Resolved ✓

§ 5.1.1. SG9 Review, Hagenberg (2025-02)

The poll results indicated that the function parameter should not have a default, and no more convenience aliases should be added.

§ 5.2. Questions for LEWG

- Should the naming be `views::inclusive_scan` to avoid conflict?
- `regular_invocable` or `invocable`? Currently the wording uses the latter (following `range-v3`), but `transform_view` used the former.

§ 6. Wording

The wording below is based on [N5032], and assumes that [P3117R1]'s `tidy_func` concept is already applied on top.

Wording notes for LWG and editor:

- Currently, the three views' definition reside just after `views::transform`'s definition and synopsis, due to the author's perception that they are pretty similar.
- The exposition-only concept `scannable` and `scannable-impl` is basically the same as `indirectly-binary-left-foldable` and `indirectly-binary-left-foldable-impl`; the only difference is that `F` is required to be move constructible instead of copy constructible.

§ 6.1. 17.3.2 Header `<version>` synopsis [version.syn]

In this clause's synopsis, insert a new macro definition in a place that respects the current alphabetical order of the synopsis, and substituting `20XXYYL` by the date of adoption.

```
#define __cpp_lib_ranges_scan 20XXYYL // freestanding, also in <ranges>
```

§ 6.2. 25.2 Header `<ranges>` synopsis [ranges.syn]

Modify the synopsis as follows:

```
// [...]
namespace std::ranges {
    // [...]
    // [range.transform], transform view
    template<input_range V, move_constructible F>
        requires view<V> && is_object_v<F> &&
            regular_invocable<F&, range_reference_t<V>> &&
            can_reference_invoke_result_t<F&, range_reference_t<V>>>
    class transform_view;

    namespace views { inline constexpr unspecified transform = unspecified; }

    // [range.scan], scan view
    enum class scan_view_kind : bool { unseeded, seeded };

    template<input_range V, move_constructible F, move_constructible T, scan
    requires see below
    class scan_view;

    template<input_range V, move_constructible F, move_constructible T, scan
    constexpr bool enable_borrowed_range<scan_view<V, F, T, K>> =
    enable_borrowed_range<V> && tidy_func<F>;

    namespace views {
    inline constexpr unspecified scan = unspecified;
    }

    // [range.take], take view
    template<view> class take_view;
```

```

template<class T>
    constexpr bool enable_borrowed_range<take_view<T>> =
        enable_borrowed_range<T>;

namespace views { inline constexpr unspecified take = unspecified; }
// [...]
}

```

Editor's Note: Add the following subclause to 25.7 Range adaptors [range.adaptors], after 25.7.9 Transform view [range.transform]

§ 6.3. 25.7. Scan view [range.scan]

§ 6.4. 25.7. .1 Overview [range.scan.overview]

1. `scan_view` presents a view that accumulates the results of applying a transformation function to the current state and each element.
2. The name `views::scan` denotes a range adaptor object ([range.adaptor.object]). Given subexpressions `E`, `F` and `G`, the expressions `views::scan(E, F)` and `views::scan(E, F, G)` are expression-equivalent to `scan_view(E, F)` and `scan_view(E, F, G)`, respectively.

[Example 1:

```

vector<int> vec{1, 2, 3, 4, 5};
for (auto&& i : std::views::scan(vec, std::plus{})) {
    std::print("{} ", i); // prints 1 3 6 10 15
}
for (auto&& i : std::views::scan(vec, std::plus{}, 10)) {
    std::print("{} ", i); // prints 11 13 16 20 25
}

```

-- end example]

§ 6.5.25.7.2 Class template `scan_view` [range.scan.view]

```

namespace std::ranges {
    template<typename V, typename F, typename T, typename U>
        concept scannable-impl = // exposition only
            movable<U> &&
            convertible_to<T, U> && invocable<F&, U, range_reference_t<V>> &&
            assignable_from<U&, invoke_result_t<F&, U, range_reference_t<V>>>;

    template<typename V, typename F, typename T>
        concept scannable = // exposition only
            invocable<F&, T, range_reference_t<V>> &&
            convertible_to<invoke_result_t<F&, T, range_reference_t<V>>,
                decay_t<invoke_result_t<F&, T, range_reference_t<V>>>> &&
            scannable-impl<V, F, T,
                decay_t<invoke_result_t<F&, T, range_reference_t<V>>>>;

    template<input_range V, move_constructible F, move_constructible T, scan_
        requires view<V> && is_object_v<F> && is_object_v<T> && scannable<V, F>
    class scan_view : public view_interface<scan_view<V, F, T, K>> {
    private:
        // [range.scan.iterator], class template scan_view::iterator
        template<bool> struct iterator; // exposition only

        V base_ = V(); // exposition only
        movable_box<F> fun_; // exposition only
        movable_box<T> init_; // exposition only

    public:
        scan_view() requires default_initializable<V> && default_initializable<
        constexpr explicit scan_view(V base, F fun) requires (K == scan_view_k
        constexpr explicit scan_view(V base, F fun, T init) requires (K == scan

        constexpr V base() const & requires copy_constructible<V> { return base
        constexpr V base() && { return std::move(base_); }

        constexpr iterator<false> begin();
        constexpr iterator<true> begin() const
            requires range<const V> && scannable<const V, const F, T>;

        constexpr default_sentinel_t end() const noexcept { return default_sen

        constexpr auto size() requires sized_range<V> { return ranges::size(ba

```

```
constexpr auto size() const requires sized_range<const V>
{ return ranges::size(base_); }

constexpr auto reserve_hint() requires approximately_sized_range<V>
{ return ranges::reserve_hint(base_); }

constexpr auto reserve_hint() const requires approximately_sized_range<const V>
{ return ranges::reserve_hint(base_); }
};

template<class R, class F>
    scan_view(R&&, F) -> scan_view<views::all_t<R>, F, range_value_t<R>>;
template<class R, class F, class T>
    scan_view(R&&, F, T) -> scan_view<views::all_t<R>, F, T, scan_view_kind_t>
}
```

```
constexpr explicit scan_view(V base, F fun) requires (K == scan_view_kind::single)
```

1. *Effects*: Initializes `base_` with `std::move(base)` and `fun_` with `std::move(fun)`.

```
constexpr explicit scan_view(V base, F fun, T init) requires (K == scan_view_kind::triple)
```

2. *Effects*: Initializes `base_` with `std::move(base)`, `fun_` with `std::move(fun)`, and `init_` with `std::move(init)`.

```
constexpr iterator<false> begin();
```

3. *Effects*: Equivalent to: `return iterator<false>{*this, ranges::begin(base_)};`

```
constexpr iterator<true> begin() const
    requires range<const V> && scannable<const V, const F, T>;
```

4. *Effects*: Equivalent to: `return iterator<true>{*this, ranges::begin(base_)};`

§ 6.6.25.7.3 Class template `scan_view::iterator` [range.scan.iterator]

```

namespace std::ranges {
    template<input_range V, move_constructible F, move_constructible T, scan_
        requires view<V> && is_object_v<F> && is_object_v<T> && scannable<V, F>
    template<bool Const>
    class scan_view<V, F, T, K>::iterator {
    private:
        using Parent = maybe-const<Const, scan_view>; // exposition only
        using Base = maybe-const<Const, V>;           // exposition only
        using Func = maybe-const<Const, F>;           // exposition only
        using ResultType = decay_t<invoke_result_t<Func&, T, range_reference_t<V>>>;

        struct Holder { // exposition only
            sentinel_t<Base> end_ = sentinel_t<Base>(); // exposition only
            movable-box<F> fun_;                        // exposition only
            movable-box<T> init_;                       // exposition only
        };
        using HolderType = conditional_t<tidy-func<F>, Holder, Parent*>; // exp

        constexpr sentinel_t<Base> get-end() const { // exposition only
            if constexpr (tidy-func<F>)
                return parent_.end_;
            else
                return std::ranges::end(parent_>base_);
        }
        constexpr Func& get-fun() { // exposition only
            if constexpr (tidy-func<F>)
                return *parent_.fun_;
            else
                return *parent_>fun_;
        }
        constexpr T& get-init() { // exposition only
            if constexpr (tidy-func<F>)
                return *parent_.init_;
            else
                return *parent_>init_;
        }
        static constexpr HolderType init(Parent& parent) { // exposition only
            if constexpr (tidy-func<F>)
                return {end(parent.base_), movable-box<F>{}, parent.init_};
            else
                return addressof(parent);
        }
    };
}

```



```

}

iterator_t<Base> current_ = iterator_t<Base>(); // exposition only
HolderType parent_ = {}; // exposition only
movable_box<ResultType> sum_; // exposition only

public:
    using iterator_concept = input_iterator_tag;
    using value_type = ResultType;
    using difference_type = range_difference_t<Base>;

    iterator() requires default_initializable<iterator_t<Base>> = default;
    constexpr iterator(Parent& parent, iterator_t<Base> current);
    constexpr iterator(iterator<!Const> i)
        requires Const && convertible_to<iterator_t<V>, iterator_t<Base>>;

    constexpr const iterator_t<Base>& base() const & noexcept { return current_; }
    constexpr iterator_t<Base> base() && { return std::move(current_); }

    constexpr const value_type& operator*() const { return *sum_; }

    constexpr iterator& operator++();
    constexpr void operator++(int);

    friend constexpr bool operator==(const iterator& x, const iterator& y)
        requires equality_comparable<iterator_t<Base>>;
    friend constexpr bool operator==(const iterator& x, default_sentinel_t)
        requires equality_comparable<iterator_t<Base>>;
};

```

```

constexpr iterator(Parent& parent, iterator_t<Base> current);

```

1. *Effects*: Initializes `current_` with `std::move(current)` and `parent_` with `init(parent)`. Then, equivalent to:

```

if (current_ == get_end()) return;
if constexpr (K == scan_view_kind::seeded) {
    sum_ = invoke(get_fun(), get_init(), *current_);
}

```

```

} else {
    sum_ = *current_;
}

```

```

constexpr iterator(iterator<!Const> i)
    requires Const && convertible_to<iterator_t<V>, iterator_t<Base>>;

```

2. *Effects*: Initializes `current_` with `std::move(i.current_)`, `parent_` with `std::move(i.parent_)`, and `sum_` with `std::move(i.sum_)`.

```

constexpr iterator& operator++();

```

3. *Effects*: Equivalent to:

```

if (++current_ != get_end()) {
    sum_ = invoke(get_fun(), std::move(*sum_), *current_);
}
return *this;

```

```

constexpr void operator++(int);

```

4. *Effects*: Equivalent to `++*this`.

```

friend constexpr bool operator==(const iterator& x, const iterator& y)
    requires equality_comparable<iterator_t<Base>>;

```

5. *Effects*: Equivalent to: `return x.current_ == y.current_;`

```

friend constexpr bool operator==(const iterator& x, default_sentinel_t);

```

6. *Effects*: Equivalent to: `return x.current_ == x.get_end();`

§ 7. Poll Results

§ 7.1. SG9 Review, Wrocław (2024-11)

Poll 1: We want a view with a convenient spelling for the semantics of `std::inclusive_scan` without initial value.

Strongly Favor	Favor	Neutral	Against	Strongly Against
1	5	3	0	0

- **Attendance:** 10
- **Author's Position:** Favor
- **Outcome:** Consensus.

Poll 2: We want a view with a convenient spelling for the semantics of `std::inclusive_scan` with initial value.

Strongly Favor	Favor	Neutral	Against	Strongly Against
1	7	1	0	0

- **Attendance:** 10
- **Author's Position:** Favor
- **Outcome:** Consensus. Stronger than Poll 1.

Poll 3: We want a view with a convenient spelling for the semantics of `std::exclusive_scan` (which would always require an initial value).

Strongly Favor	Favor	Neutral	Against	Strongly Against
0	3	6	1	0

- **Attendance:** 10
- **Author's Position:** Against
- **Outcome:** Weak consensus.

Poll 4: We want a view with a convenient spelling for the semantics of `views::prescan` as proposed in the paper ([P3351R1]) (`std::exclusive_scan` plus the final sum).

Strongly Favor	Favor	Neutral	Against	Strongly Against
1	2	4	2	0

- **Attendance:** 10

- **Author's Position:** Favor
- **Outcome:** No consensus either way. Author's choice.

§ 7.2. SG9 Review, Hagenberg (2025-02)

Poll 1: We think `scan_view` should always be an `input_range` because (among other reasons) we think that copying the value every time we perform a dereference is an unacceptable hidden cost.

Strongly Favor	Favor	Neutral	Against	Strongly Against
1	3	4	1	1

- **Attendance:** 13
- **Author's Position:** Neutral
- **Outcome:** Very weak consensus. Author's choice.

Poll 2: We think `scan_view` should be conditionally borrowed when the function is a tidy function (in the [P3117R1] sense).

Strongly Favor	Favor	Neutral	Against	Strongly Against
4	1	3	0	0

- **Attendance:** 13
- **Author's Position:** Neutral
- **Outcome:** Consensus.

Poll 3: Eliminate `partial_sum` as a short-hand for `scan` with `plus`, and `scan` should not have a defaulted function argument.

Strongly Favor	Favor	Neutral	Against	Strongly Against
3	5	3	0	0

- **Attendance:** 13
- **Author's Position:** Neutral
- **Outcome:** Consensus.

§ 7.3. SG9 Review, Kona (2025-11)

Poll 1: Replace the `bool IsInit` template parameter with some kind of enum similar to the `std::ranges::subrange_kind` enum of `std::ranges::subrange`.

Strongly Favor	Favor	Neutral	Against	Strongly Against
2	3	2	0	0

- **Attendance:** 8
- **Author's Position:** Favor
- **Outcome:** Strong consensus in favor.

Poll 2: Instead of having an templated `scan_view` templated on having an initializer, provide two separate classes `scan_view` and `scan_view_with_init` or something like that.

Strongly Favor	Favor	Neutral	Against	Strongly Against
2	0	2	3	0

- **Attendance:** 8
- **Author's Position:** Neutral
- **Outcome:** No consensus.

Poll 3: Allow an implementation to make the iterator not dependent on the `IsInit` parameter.

Strongly Favor	Favor	Neutral	Against	Strongly Against
1	2	3	0	0

- **Attendance:** 8
- **Author's Position:** Neutral
- **Outcome:** Weak consensus in favor.

Poll 4: Forward [\[P3351R3\]](#) `views::scan` with the changes above to LEWG for inclusion in C++29.

Strongly Favor	Favor	Neutral	Against	Strongly Against
4	3	0	0	0

- **Attendance:** 8
- **Author's Position:** Strongly Favor
- **Outcome:** Strong consensus in favor.

§ References

§ Normative References

[N5032]

Thomas Köppe. *Working Draft, Standard for Programming Language C++*. 15 December 2025. URL: <https://wg21.link/n5032>

[P2760R1]

Barry Revzin. *A Plan for C++26 Ranges*. 15 December 2023. URL: <https://wg21.link/p2760r1>

[P3117R1]

Zach Laine, Barry Revzin, Jonathan Müller. *Extending Conditionally Borrowed*. 15 December 2024. URL: <https://wg21.link/p3117r1>

[P3351R1]

Yihe Li. *views::scan*. 9 October 2024. URL: <https://wg21.link/p3351r1>

[P3351R2]

Yihe Li. *views::scan*. 12 January 2025. URL: <https://wg21.link/p3351r2>

[P3351R3]

Yihe Li. *views::scan*. 4 October 2025. URL: <https://wg21.link/p3351r3>

§ Informative References

[LWG4189]

Hewill Kang. *cache_latest_view should be freestanding*. WP. URL: <https://wg21.link/lwg4189>

[LWG4286]

Yihe Li. *Some more feature-test macros for fully freestanding features are not marked freestanding*. WP. URL: <https://wg21.link/lwg4286>

[N4988]

Thomas Köppe. *Working Draft, Programming Languages — C++*. 5 August 2024. URL: <https://wg21.link/n4988>

[N5001]

Thomas Köppe. *Working Draft, Programming Languages — C++*. 17 December 2024. URL: <https://wg21.link/n5001>

[N5014]

Thomas Köppe. *Working Draft, Standard for Programming Language C++*. 5 August 2025. URL: <https://wg21.link/n5014>

[P1642R11]

Ben Craig. *Freestanding Library: Easy [utilities], [ranges], and [iterators]*. 1 July 2022. URL: <https://wg21.link/p1642r11>

[P1729R5]

Elias Kosunen, Victor Zverovich. *Text Parsing*. 15 October 2024. URL: <https://wg21.link/p1729r5>

[P2214R2]

Barry Revzin, Conor Hoekstra, Tim Song. *A Plan for C++23 Ranges*. 18 February 2022. URL: <https://wg21.link/p2214r2>

[P2322R6]

Barry Revzin. *ranges::fold*. 22 April 2022. URL: <https://wg21.link/p2322r6>

[P2846R6]

Corentin Jabot. *reserve_hint: Eagerly reserving memory for not-quite-sized lazy ranges*. 15 February 2025. URL: <https://wg21.link/p2846r6>

[P3136R1]

Tim Song. *Retiring niebloids*. 18 November 2024. URL: <https://wg21.link/p3136r1>

[P3179R9]

Ruslan Arutyunyan, Alexey Kukanov, Bryce Adelstein Lelbach. *C++ parallel range algorithms*. 29 May 2025. URL: <https://wg21.link/p3179r9>

[P3732R2]

Ruslan Arutyunyan, Mark Hoemmen, Alexey Kukanov, Bryce Adelstein Lelbach, Abhilash Majumder. *Numeric Range Algorithms*. 26 March 2026. URL: <https://wg21.link/p3732r2>