

Annotations for Reflection

Document #: P3394R2 [[Latest](#)] [[Status](#)]
Date: 2025-03-14
Project: Programming Language C++
Audience: CWG, LWG
Reply-to: Wyatt Childers
<wcc@edg.com>
Dan Katz
<dkatz85@bloomberg.net>
Barry Revzin
<barry.revzin@gmail.com>
Daveed Vandevoorde
<daveed@edg.com>

Contents

[1 Revision History](#)

[2 Introduction](#)

[3 Motivating Examples](#)

[3.1 Command-Line Argument Parsing](#)

[3.2 Test Parametrization](#)

[3.3 Serialization](#)

[4 Proposal](#)

[4.1 Why not Attributes?](#)

[4.2 Library Queries](#)

[4.3 Additional Syntactic Constraints](#)

[4.4 Implementation Experience](#)

[4.5 Other Directions We Are Exploring](#)

[5 Wording](#)

[5.1 Language](#)

[5.2 Library](#)

[6 References](#)

§ 1 Revision History

Since [\[P3394R1\]](#):

- loosened the P2996 scope restrictions around injected declarations: allow the definition of an entity to produce injected declaration of itself
- reduced library API down to just `is_annotation`, `annotations_of`, `annotations_of_with_type` (a binary function that returns a vector<info>), and `annotate`.

Since [\[P3394R0\]](#):

— added wording

§ 2 Introduction

Ever since writing [\[P1240R0\]](#) ([Scalable Reflection in C++](#)), but more so since [\[P2996R0\]](#) ([Reflection for C++26](#)), we have been requested to add a capability to annotate declarations in a way that reflection can observe. For example, Jeremy Ong presented compelling arguments in a post to the [SG7 reflector](#). Corentin Jabot also noticed the need while P1240 was evolving and wrote [\[P1887R0\]](#) ([Typesafe Reflection on attributes](#)), which proposes syntax not entirely unlike what we present here.

In early versions of P2996 (and P1240 before that), a workaround was to encode properties in the template arguments of alias template specializations:

```
template <typename T, auto... Annotations>
using Noted = T;

struct C {
    Noted<int, 1> a;
    Noted<int*, some, thing> b;
};
```

It was expected that something like `type_of(^C::a)` would produce a reflection of `Noted<int, 1>` and that can be taken apart with metafunctions like `template_arguments_of` — which both preserves the type as desired (a is still an `int`) and allows reflection queries to get at the desired annotations (the `1`, `some`, and `thing` in this case).

There are problems with this approach, unfortunately:

- It doesn't work for all contexts where we might want to annotate declarations (e.g., enumerators).
- It doesn't directly express the intent.
- It turns out that providing access to aliases used in the declaration of reflected entities raises difficult questions and P2996 is therefore likely going to drop the ability to access that information (with the intention of resurrecting the capability with a later paper once specification and implementation challenges have been addressed).

In this paper, we propose simple mechanisms that more directly support the ability to annotate C++ constructs.

§ 3 Motivating Examples

We'll start with a few motivating examples for the feature. We'll describe the details of the feature in the subsequent section.

These examples are inspired from libraries in other programming languages that provide some mechanism to annotate declarations.

§ 3.1 Command-Line Argument Parsing

Rust's [clap](#) library provides a way to add annotations to declarations to help drive how the parser is declared. We can now [do the same](#):

```
struct Args {
    [[=clap::Help("Name of the person to greet")]]
    [[=clap::Short, =clap::Long]]
    std::string name;

    [[=clap::Help("Number of times to greet")]]
    [[=clap::Short, =clap::Long]]
    int count = 1;
};

int main(int argc, char** argv) {
    Args args = clap::parse<Args>(argc, argv);

    for (int i = 0; i < args.count; ++i) {
        std::cout << "Hello " << args.name << '\n';
    }
}
```

Here, we provide three types (Short, Long, and Help) which help define how these member variables are intended to be used on the command-line. This is implemented on top of [Lyra](#).

When run:

```
$ demo -h
USAGE:
  demo [-?|-h|--help] [-n|--name <name>] [-c|--count <count>]

Display usage information.

OPTIONS, ARGUMENTS:
  -?, -h, --help
  -n, --name <name>      Name of the person to greet
  -c, --count <count>    Number of times to greet

$ demo -n wg21 --count 3
Hello wg21
Hello wg21
Hello wg21
```

While Short and Long can take explicit values, by default they use the first letter and whole name of the member that they annotate.

The core of the implementation is that `parse<Args>` loops over all the non-static data members of `Args`, then finds all the `clap`-related annotations and invokes them:

```
template <typename Args>
auto parse(int argc, char** argv) -> Args {
    Args args;
    auto cli = lyra::cli();

    // ...

    template for (constexpr info M : nonstatic_data_members_of(^Args))
    {
        auto id = std::string id(identifier_of(mem));
        auto opt = lyra::opt(args.[:M:], id);

        template for (constexpr info A : annotations_of(M)) {
            if constexpr (parent_of(type_of(A)) == ^clap) {
                // for those annotations that are in the clap namespace
                // invoke them on our option
                extract<[:type_of(A):]>(A).apply_annotation(opt, id);
            }
        }

        cli.add_argument(opt);
    }

    // ...
};
```

So, for instance, Short would be implemented like this:

```
namespace clap {
    struct ShortArg {
        // optional isn't structural yet but let's pretend
        optional<char> value;

        constexpr auto operator()(char c) const -> ShortArg {
            return {.value=c};
        };

        auto apply_annotation(lyra::opt& opt, std::string_view id) const
        -> void {
            char first = value.value_or(id[0]);
            opt[std::string("-") + first];
        }
    };
};
```

```
inline constexpr auto Short = ShortArg();
}
```

Overall, a fairly concise implementation for an extremely user-friendly approach to command-line argument parsing.

§ 3.2 Test Parametrization

The pytest framework comes with a decorator to [parametrize](#) test functions. We can now do [the same](#):

```
namespace N {
    [[=parametrize({
        Tuple{1, 1, 2},
        Tuple{1, 2, 3}
    })]]
    void test_sum(int x, int y, int z) {
        std::println("Called test_sum(x={}, y={}, z={})", x, y, z);
    }

    struct Fixture {
        Fixture() {
            std::println("setup fixture");
        }

        ~Fixture() {
            std::println("teardown fixture");
        }

        [[=parametrize({Tuple{1}, Tuple{2}})]]
        void test_one(int x) {
            std::println("test one({})", x);
        }

        void test_two() {
            std::println("test two");
        }
    };
}

int main() {
    invoke_all<^N>();
}
```

When run, this prints:

```
Called test_sum(x=1, y=1, z=2)
Called test_sum(x=1, y=2, z=3)
setup fixture
test one(1)
teardown fixture
```

```

setup fixture
test one(2)
teardown fixture
setup fixture
test two
teardown fixture

```

Here, `parametrize` returns a value that is some specialization of `Parametrize` (which is basically an array of tuples, except that `std::tuple` isn't structural so the implementation rolls its own).

The rest of the implementation looks for all the free functions named `test_*` or nonstatic member functions of class types that start with `test_*` and invokes them once, or with each parameter, depending on the presence of the annotation. That looks like this:

```

constexpr auto parametrization_of(std::meta::info M) -> std::meta::info
{
    for (auto a : annotations_of(M)) {
        auto t = type_of(a);
        if (has_template_arguments(t) and template_of(t) ==
            ^^Parametrize) {
            return a;
        }
    }
    return std::meta::info();
}

template <std::meta::info M, class F>
void invoke_single_test(F f) {
    constexpr auto A = parametrization_of(M);

    if constexpr (A != std::meta::info()) {
        // if we are parametrized, pull out that value
        // and for each tuple, invoke the function
        // this is basically calling std::apply on an array of tuples
        constexpr auto Params = extract<[:type_of(A):]>(A);
        for (auto P : Params) {
            P.apply(f);
        }
    } else {
        f();
    }
}

template <std::meta::info Namespace>
void invoke_all() {
    template for (constexpr std::meta::info M : members_of(Namespace)) {
        if constexpr (is_function(M) and
            identifier_of(M).starts_with("test_")) {
            invoke_single_test<M>([:M:]);
        } else if constexpr (is_type(M)) {
            template for (constexpr std::meta::info F :
                nonstatic_member_functions_of(M)) {
                if constexpr (identifier_of(F).starts_with("test_")) {

```

```

        invoke_single_test<F>([&](auto... args){
            typename [:M:] fixture;
            fixture.[:F:] (args...);
        });
    }
}
}
}
}
}

```

§ 3.3 Serialization

Rust's [serde](#) library is a framework for serialization and deserialization. It is easy enough with reflection to do member-wise serialization. But how do you opt into that? An annotation provides a cheap mechanism of [doing just that](#) (built on top of [Boost.Json](#)):

```

struct [[=serde::derive]] Point {
    int x, y;
};

```

Allowing:

```

// prints {"x":1,"y":2}
std::cout << boost::json::value_from(Point{.x=1, .y=2});

```

But opting in is just the first thing you might want to do with serialization. You might also, for instance, want to change how fields are serialized. `serde` provides a lot of attributes to do so. The easiest to look at is `rename`, which uses the provided string instead of the name of the non-static data member:

```

struct [[=serde::derive]] Person {
    [[=serde::rename("first name")]] std::string first;
    [[=serde::rename("last name")]] std::string last;
};

```

Which leads to:

```

// prints {"first name":"Peter","last name":"Dimov"}
std::cout << boost::json::value_from(Person{.first="Peter",
    .last="Dimov"});

```

The implementation for these pieces is fairly straightforward. We provide an opt-in for the value conversion function when the `serde::derive` annotation is present. In that case, we walk all the non-static data members and write them into the `boost::json::value` output. If a `serde::Rename` annotation is present, we use that instead of the data member's name:

```

namespace serde {
    inline constexpr struct{} derive{};
    struct rename { char const* field; };
}

namespace boost::json {
    template <class T>
        requires (has_annotation(^T, serde::derive))
    void tag_invoke(value_from_tag const&, value& v, T const& t) {
        auto& obj = v.emplace_object();
        template for (constexpr auto M : nonstatic_data_members_of(^T))
        {
            constexpr auto field = []{
                std::optional<std::meta::info> res;
                for (std::meta::info r : annotations_of_with_type(M,
                    ^serde::rename)) {
                    if (res and *res != r) {
                        throw "oops";
                    }
                    res = r;
                }

                return res
                    .transform([](std::meta::info r){
                        return std::string_view(extract<serde::rename>
                            (r).field);
                    })
                    .value_or(identifier_of(M));
            }();

            obj[field] = boost::json::value_from(t.[:M:]);
        }
    }
}

```

You can imagine extending this out to support a wide variety of other serialization-specific attributes that shouldn't otherwise affect the C++ usage of the type. For instance, a [more complex approach](#) additionally supports the `skip_serializing_if` annotation while first collecting all `serde` annotations into a struct.

§ 4 Proposal

The core idea is that an *annotation* is a compile-time value that can be associated with a construct to which attributes can appertain. Annotation and attributes are somewhat related ideas, and we therefore propose a syntax for the former that builds on the existing syntax for the latter.

At its simplest:

```

struct C {
    [[=1]] int a;
}

```



```
};
```

Syntactically, an annotation is an attribute of the form `= expr` where `expr` is a *constant-expression* (which syntactically excludes, e.g., *comma-expression*) to which the glvalue-to-prvalue conversion has been applied if the expression wasn't a prvalue to start with.

Currently, we require that an annotation has structural type because we're going to return annotations through `std::meta::info`, and currently all reflection values must be structural.

§ 4.1 Why not Attributes?

Attributes are very close in spirit to annotations. So it made sense to piggy-back on the attribute syntax to add annotations. Existing attributes are designed with fairly open grammar and they can be ignored by implementations, which makes it difficult to connect them to user code. Given a declarations like:

```
[[nodiscard, gnu::always_inline]]  
[[deprecated("don't use me")]]  
void f();
```

What could reflecting on `f` return? Because attributes are ignorable, an implementation might simply ignore them. Additionally, there is no particular value associated with any of these attributes that would be sensible to return. We're limited to returning either a sequence of strings. Or, with [P3294](#), token sequences.

But it turns out to be quite helpful to preserve the actual values without requiring libraries to do additional parsing work. Thus, we need to distinguish annotations (whose values we need to preserve and return back to the user) from attributes (whose values we do not). Thus, we looked for a sigil introducing a general expression.

Originally, the plus sign (+) was considered (as in P1887), but it is not ideal because a prefix `+` has a meaning for some expressions and not for others, and that would not carry over to the attribute notation. A prefix `=` was found to be reasonably meaningful in the sense that the annotation “equals” the value on the right, while also being syntactically unambiguous. We also discussed using the reflection operator (`^^`) as an introducer (which is attractive because the annotation ultimately comes back to the programmer as a reflection value), but that raised questions about an annotation itself being a reflection value (which is not entirely improbable).

As such, this paper proposes annotations as distinct from attributes, introduced with a prefix `=`.

§ 4.2 Library Queries

We propose the following set of library functions to work with annotations:

```
namespace std::meta {  
    consteval bool is_annotation(info);
```

```

    consteval      vector<info>      annotations_of(info      item);
    // (1)
    consteval      vector<info> annotations_of_with_type(info item, info type);
    // (2)

    consteval info annotate(info item,
                           info value,
                           source_location loc =
    source_location::current());
}

```

`is_annotation` checks whether a particular reflection represents an annotation.

We provide two functions to retrieve all the annotations of a particular item:

1. `annotations_of(item)` returns all the annotations on item.
2. `annotations_of_item_with_type(item, type)` returns all the annotations a on item such that `type_of(a) == type`.

In earlier revisions, these were overloads (both spelled `annotations_of`), but starting in R2 the second function was renamed to add clarity on what the second parameter actually means. There were previously additional functions proposed for ergonomic reasons, but those were removed during the [Hagenberg meeting](#).

And lastly, `annotate` provides the ability to programmatically add an annotation to a declaration.

§ 4.3 Additional Syntactic Constraints

Annotations can be repeated:

```

[[=42, =42]] int x;
static_assert(annotations_of(^x).size() == 2);

```

Annotations spread over multiple declarations of the same entity accumulate:

```

[[=42]] int f();
[[=24]] int f();
static_assert(annotations_of(^f).size() == 2);

```

Annotations follow appertenance rules like attributes, but shall not appear in the *attribute-specifier-seq* of a *type-specifier-seq* or an *empty-declaration*:

```

struct [[=0]] S {}; // Okay: Appertains to S.
[[=42]] int f();    // Okay: Appertains to f.
int f[[=0]] ();     // Ditto.

```

```
int [[=24]] f();      // Error: Cannot appertain to int.  
[[=123]];            // Error: No applicable construct.
```

To avoid confusion, annotations are not permitted after an *attribute-using-prefix*. So this is an error:

```
[[using clang: amdgpu_waves_per_eu, =nick("weapon")]]  
int select_footgun(int);
```

Instead, use:

```
[[using clang: amdgpu_waves_per_eu]] [[=nick("weapon")]]  
int select_footgun(int);
```

§ 4.4 Implementation Experience

The core language feature and the basic query functions have been implemented in the EDG front end and in Bloomberg's P2996 Clang fork (with option `-freflection-latest`), both available on Compiler Explorer.

§ 4.5 Other Directions We Are Exploring

As evidenced in the motivating examples earlier, there is a lot of value in this proposal even in this simple form. However, there is more to consider when it comes to annotations.

This proposal right now lets us unconditionally add an annotation to a type:

```
struct [[=X]] Always;
```

But it does not let us conditionally add an annotation to a type:

```
template <class T>  
struct /* X only for some T */ Sometimes;
```

Or to really generalize annotations. For instance, in the clap example earlier, our example showed usage with `clap::Short` and `clap::Long`. What if somebody wants to compose these into their own annotation that attaches both `clap::Short` and `clap::Long` to a declaration?

More broadly, there is clear value in having an annotation be able to be invoked by the declaration itself. Doing so allows the two uses above easily enough. An interesting question, though, is whether this callback (syntax to be determined) is invoked at the *beginning* of the declaration or at the *end* of the declaration. For annotations on classes, this would be before the class is complete or after the class is complete. Before completeness allows the class to observe the annotation during instantiation. After

completeness allows the annotation callback to observe properties of the type. In some sense, Herb Sutter's [\[P0707R4\] \(Metaclasses: Generative C++\)](#) was adding annotations on classes, invoked on class completeness, that allow mutation of the class.

One concrete, simpler example. We can, with this proposal as-is, create a Debug annotation that a user can add to their type and a specialization of `std::formatter` for all types that have a Debug annotation [as follows](#):

```
template <auto V> struct Derive { };
template <auto V> inline constexpr Derive<V> derive;

inline constexpr struct{} Debug;

template <class T> requires (has_annotation(^T, derive<Debug>))
struct std::formatter<T> {
    // ...
};

struct [[=derive<Debug>]] Point {
    int x;
    int y;
};

int main() {
    auto p = Point{.x=1, .y=2};
    // prints p=Point{.x=1, .y=2}
    std::println("p={}", p);
}
```

This *works*, but it's not really the ideal way of doing it. This could still run into potential issues with ambiguous specialization of `std::formatter`. Better would be to allow the Debug annotation to, at the point of completion of `Point`, inject an explicit specialization of `std::formatter`. This would rely both on the ability for the annotation to be called back and language support for such injection (see [\[P3294R2\] \(Code Injection with Token Sequences\)](#)).

There are still open questions as to how to handle such callbacks. Does an annotation that gets called back merit different syntax from an annotation that doesn't? Can it mutate the entity that it is attached to? How do we name the potential callbacks? Should the callback be registered implicitly (e.g., if C of type X with member `X::annotate_declaration(...)` appears, that member is automatically a callback invoked when an entity is first declared with an annotation of type X) or explicitly (e.g., calling `annotated_declaration_callback(^X, X_handler)` would cause `X_handler(...)` to be invoked when an entity is first declared with an annotation of type X).

§ 5 Wording

The wording is relative to [\[P2996R10\]](#).

§ 5.1 Language

Change 6.8.2 [\[basic.fundamental\]](#) to add “annotation” to the list of reflection kinds:

- 17-1 A value of type `std::meta::info` is called a *reflection*. There exists a unique *null reflection*; every other reflection is a representation of
- a value with structural type ([temp.param]),
 - an object with static storage duration,
 - a variable,
 - a structured binding,
 - a function,
 - an enumerator,
 - an annotation ([dcl.attr.grammar]),
 - a type,
 - a *typedef-name*,
 - a class member,
 - a bit-field,
 - a primary class template, function template, primary variable template, alias template, or concept,
 - a namespace or namespace alias,
 - a base class specifier, or
 - a description of a declaration of a non-static data member.

Update annotation equality in 7.6.10 [\[expr.eq\]](#)/5+:

- 5+ If both operands are of type `std::meta::info`, comparison is defined as follows:
- (5+.1) — If one operand is a null reflection value, then they compare equal if and only if the other operand is also a null reflection value.
 - (5+.2) — Otherwise, if one operand represents a value, then they compare equal if and only if the other operand represents a value that is template-argument-equivalent (13.6 [\[temp.type\]](#)).
 - (5+.x) — Otherwise, if one operand represents an annotation, then they compare equal if and only if the other operand represents the same annotation.
 - (5+.3) — Otherwise, if one operand represents an object, then they compare equal if and only if the other operand represents the same object.
 - (5+.4) — Otherwise, if one operand represents an entity, then they compare equal if and only if the other operand represents the same entity.

- (5+.5) — Otherwise, if one operand represents a direct base class relationship, then they compare equal if and only if the other operand represents the same direct base class relationship.
- (5+.6) — Otherwise, both operands 0_1 and 0_2 represent data member descriptions. The operands compare equal if and only if the data member descriptions represented by 0_1 and 0_2 compare equal (11.4.1 [\[class.mem.general\]](#)).

Loosen the scope restrictions in 7.7 [\[expr.const\]/29](#) to allow definitions to inject declarations of the entity that they are defining.

- 29 Let C be a *consteval-block-declaration*, the evaluation of whose corresponding expression produces an injected declaration D ([meta.reflection.define.aggregate]). **The If D is a definition, then the** scope of D shall not enclose C . The program is ill-formed if a scope S encloses exactly one of C or D where S is
 - (29.1) — a function parameter scope, or
 - (29.2) — a class scope.

Extend the grammar in 9.12.1 [\[dcl.attr.grammar\]](#):

- 1 Attributes **and annotations** specify additional information for various source constructs such as types, variables, names, blocks, or translation units.

```

attribute-specifier-seq:
    attribute-specifier-seqopt attribute-specifier

attribute-specifier:
    [ [ attribute-using-prefixopt attribute-list ] ]
    alignment-specifier

alignment-specifier:
    alignas ( type-id ...opt )
    alignas ( constant-expression ...opt )

attribute-using-prefix:
    using attribute-namespace :
```

```

- attribute-list:
-   attributeopt
-   attribute-list , attributeopt
-   attribute ...
-   attribute-list , attribute ...
+ attribute-or-annotation-list:
+   attribute-or-annotationopt
+   attribute-or-annotation-list , attribute-or-annotationopt
+   attribute-or-annotation ...
+   attribute-or-annotation-list , attribute-or-annotation ...
```

```

+ attribute-or-annotation:
+   attribute
+   annotation

attribute:
  attribute-token attribute-argument-clauseopt

+ annotation:
+   = constant-expression

```

- 2 If an *attribute-specifier* contains an *attribute-using-prefix*, the *attribute-list* following that *attribute-using-prefix* shall not contain *either* an *attribute-scoped-token* *or an annotation* and every *attribute-token* in that *attribute-list* is treated as if its identifier were prefixed with *N::*, where *N* is the *attribute-namespace* specified in the *attribute-using-prefix*.

Adjust the restriction on ellipses in 9.12.1 [\[dcl.attr.grammar\]](#)/4:

- 4 In an *attribute-list*, an ellipsis may ~~appear only~~ *only appear following an attribute* if that *attribute*'s specification permits it. An ~~attribute~~ *attribute-or-annotation* followed by an ellipsis is a pack expansion. An *attribute-specifier* that contains no *attributes* has no effect. The order in which the *attribute-tokens* appear in an *attribute-list* is not significant. *The constant-expression in each annotation shall have structural type ([temp.param]). [...]*
- 5 Each *attribute-specifier-seq* is said to appertain to some entity or statement, identified by the syntactic context where it appears ([stmt.stmt], [dcl.dcl], [dcl.decl]). If an *attribute-specifier-seq* that appertains to some entity or statement contains an *attribute* or *alignment-specifier* that is not allowed to apply to that entity or statement, the program is ill-formed. If an *attribute-specifier-seq* appertains to a friend declaration ([class.friend]), that declaration shall be a definition. *The constant-expression in each annotation in an attribute-specifier-seq shall have structural type ([temp.param]) [Note 1: Reflections representing annotations can be retrieved with functions like std::meta::annotations_of ([meta.reflection.annotation]) — end note].*

Change the pack expansion rule in 13.7.4 [\[temp.variadic\]](#)/5.9:

- 5 [...] Pack expansions can occur in the following contexts:
- (5.9) — In an *attribute-list* ([dcl.attr.grammar]); the pattern is an ~~attribute~~ *attribute-or-annotation*.

§ 5.2 Library

Add to the `<meta>` synopsis:

```
namespace std::meta {
    // ...
    consteval bool is_bit_field(info r);
    consteval bool is_enumerator(info r);
+   consteval bool is_annotation(info);
    // ..

    // [meta.reflection.annotation], annotation reflection
+   consteval vector<info> annotations_of(info item);
+   consteval vector<info> annotations_of_with_type(info item, info type);
+
+   consteval info annotate(info item, info value, source_location loc =
        source_location::current());
}
```

Add `is_annotation` to `[meta.reflection.queries]`:

```
consteval bool is_enumerator(info r);
+ consteval bool is_annotation(info r);
```

- 1 *Returns:* `true` if `r` represents an enumerator or annotation, respectively. Otherwise, `false`.

Update `source_location_of` in `[meta.reflection.names]`:

```
consteval source_location source_location_of(info r);
```

- 7 *Returns:* If `r` represents a value, a type other than a class type or an enumeration type, the global namespace, or a data member description, then `source_location{}`. Otherwise, an implementation-defined `source_location` value.
- 8 *Recommended practice:* If `r` represents an annotation that was added with a call to `std::meta::annotate`, then implementations should return a value corresponding to the `source_location` argument passed to that function. If `r` represents an entity with a definition that is reachable from the evaluation context, a value corresponding to a definition should be returned.

Update the meanings of `type_of` and `value_of` in `[meta.reflection.queries]`:

```
consteval info type_of(info r);
```

- 35 *Constant When:* `r` represents a value, object, variable, function that is not a constructor or destructor, enumerator, annotation, non-static data member, bit-field, base class specifier, or description of a declaration of a non-static data member.

36 *Returns:* If *r* represents an entity, object, or value, then a reflection of the type of what is represented by *r*. Otherwise, if *r* represents an annotation, then the type of the annotated value. Otherwise, if *r* represents a direct base class relationship, then a reflection of the type of the direct base class. Otherwise, for a data member description (*T*, *N*, *A*, *W*, *NUA*) (11.4.1 [\[class.mem.general\]](#)), a reflection of the type *T*.

```
consteval info value_of(info r);
```

39 *Constant When:* *r* is a reflection representing

(39.1) — either an object or variable, usable in constant expressions from a point in the evaluation context ([*expr.const*]), whose type is a structural type ([*temp.type*]),

(39.2) — an annotation,

(39.3) — an enumerator, or

(39.4) — a value.

40 *Returns:*

(40.1) — If *r* is a reflection of an object *o*, or a reflection of a variable which designates an object *o*, then a reflection of the value held by *o*. The reflected value has type *type_of(o)*, with the cv-qualifiers removed if this is a scalar type

(40.2) — Otherwise, if *r* is a reflection of an enumerator or an annotation, then a reflection of the value of the enumerator or annotation, respectively.

(40.3) — Otherwise, *r*.

And update `extract` in `[meta.reflection.extract]`, specifically *extract-val*:

```
template <class T>
consteval T extract-val(info r); // exposition only
```

41 Let *U* be the type of the value, annotation, or enumerator that *r* represents.

42 *Constant When:*

(42.1) — *U* is a pointer type, *T* and *U* are similar types ([*conv.qual*]), and `is_convertible_v<U, T>` is `true`,

(42.2) — *U* is not a pointer type and the cv-unqualified types of *T* and *U* are the same, or

(42.3) — *U* is a closure type, *T* is a function pointer type, and the value *r* represents is convertible to *T*.

43 *Returns:* the value, annotation, or enumerator *V* represented by *r*, converted to *T*.

And to the new section `[meta.reflection.annotation]`:

1 Annotations can be used to add additional information to declarations for future introspection.

```
consteval vector<info> annotations_of(info item);
```

2 *Constant When:* `dealias(item)` represents a type, variable, function, or a namespace.

3 *Returns:* A vector containing the following reflections:

- For each declaration *D* of the entity represented by *item* that precedes either some point in the evaluation context (`[expr.const]`) or a point immediately following the *class-specifier* of a class for which such a point is in a complete-class context,
- For each *annotation A* in each *attribute-specifier-seq* appertaining to *D*,
- A reflection *R* representing *A*, whose value is the result of the corresponding *constant-expression*.

For any two reflections R_1 and R_2 returned, if the annotation represented by R_1 precedes the annotation represented by R_2 , then R_1 is sequenced before R_2 in the returned vector.

[*Example 1:*

```
struct Option { bool value; };

[[=1]] void f();
[[=2, =3]] void g();
[[=4]] void g();

struct C {
    [[=Option{true}]] int a;
    [[=Option{false}]] int b;
};

static_assert(annotations_of(^f).size() == 1);
static_assert(annotations_of(^g).size() == 3);
static_assert(extract<int>(annotations_of(^g)[0]) == 2);
static_assert(extract<int>(annotations_of(^g)[1]) == 3);
static_assert(extract<int>(annotations_of(^g)[2]) == 4);

static_assert(extract<Option>(annotations_of(^C::a)[0]).value);
static_assert(!extract<Option>(annotations_of(^C::b)[0]).value);
```

— *end example*]

```
consteval vector<info> annotations_of_with_type(info item, info type);
```

4 *Constant When:* `annotations_of(item)` is constant and `dealias(type)` is a reflection representing a complete type.

5 *Returns:* A vector containing each element, *e*, of `annotations_of(item)`, in order, such that `dealias(type_of(e)) == dealias(type)`.

```
constexpr info annotate(info item, info value, source_location loc =
    source_location::current());
```

6 *Constant When:*

(6.1) — `dealias(item)` represents a class type, variable, function, or a namespace; and

(6.2) — `value` represents a value.

7 *Effects:* Produces an injected declaration (`[expr.const]`) at location `loc` redeclaring the entity represented by `dealias(item)`. That injected declaration is annotated by `value` and its locus is immediately following the manifestly constant-evaluated expression currently under evaluation.

8 *Returns:* `dealias(item)`.

§ 6 References

[P0707R4] Herb Sutter. 2019-06-17. Metaclasses: Generative C++.

<https://wg21.link/p0707r4>

[P1240R0] Andrew Sutton, Faisal Vali, Daveed Vandevoorde. 2018-10-08. Scalable Reflection in C++.

<https://wg21.link/p1240r0>

[P1887R0] Corentin Jabot. 2019-10-06. Typesafe Reflection on attributes.

<https://wg21.link/p1887r0>

[P2996R0] Barry Revzin, Wyatt Childers, Peter Dimov, Andrew Sutton, Faisal Vali, Daveed Vandevoorde. 2023-10-15. Reflection for C++26.

<https://wg21.link/p2996r0>

[P2996R10] Barry Revzin, Wyatt Childers, Peter Dimov, Andrew Sutton, Faisal Vali, Daveed Vandevoorde, and Dan Katz. 2025-02-26. Reflection for C++26.

<https://wg21.link/p2996r10>

[P3294R2] Barry Revzin, Andrei Alexandrescu, Daveed Vandevoorde. 2024-10-15. Code Injection with Token Sequences.

<https://wg21.link/p3294r2>

[P3394R0] Daveed Vandevoorde, Wyatt Childers, Dan Katz,. 2024-10-14. Annotations for Reflection.

<https://wg21.link/p3394r0>

[P3394R1] Daveed Vandevoorde, Wyatt Childers, Dan Katz, Barry Revzin. 2025-01-13. Annotations for Reflection.

<https://wg21.link/p3394r1>