

# P3477R0

## There are exactly 8 bits in a byte

Published Proposal, 2024-10-15

**This version:**

<http://wg21.link/P3477r0>

**Author:**

[JF Bastien](#) (Woven by Toyota)

**Audience:**

EWG, LEWG, SG22

**Project:**

ISO/IEC 14882 Programming Languages — C++, ISO/IEC JTC1/SC22/WG21

**Source:**

[github.com/jfbastien/papers/blob/master/source/P3477r0.bs](https://github.com/jfbastien/papers/blob/master/source/P3477r0.bs)

---

## Table of Contents

### 1 Rationale

### 2 Impact on C

### 3 Wording

#### 3.1 Language

#### 3.2 Library

### References

Informative References

## § 1. Rationale

C has the `CHAR_BIT` macro which contains the implementation-defined number of bits in a byte, without restrictions on the value of this number. C++ imports this macro as-is. Many other macros and character traits have values derived from `CHAR_BIT`. While this was historically relevant in computing's early days, modern hardware has overwhelmingly converged on the assumption that a byte is 8 bits. This document proposes that C++ formally mandates that a byte is 8 bits.

Mainstream compilers already support this reality:

- GCC [sets a default value of 8](#) but [no upstream target changes the default](#);
- LLVM sets `__CHAR_BIT__` to 8; and
- MSVC defines `CHAR_BIT` to 8.

We can find vestigial support, for example GCC dropped [dsp16xx in 2004](#), and [1750a in 2002](#). Search [the web for more evidence](#) finds a few GCC out-of-tree ports which do not seem relevant to modern C++.

[\[POSIX\]](#) has mandated this reality since POSIX.1-2001 (or IEEE Std 1003.1-2001), saying:

As a consequence of adding `int8_t`, the following are true:

- A byte is exactly 8 bits.
- `CHAR_BIT` has the value 8, `SCHAR_MAX` has the value 127, `SCHAR_MIN` has the value -128, and `UCHAR_MAX` has the value 255.

Since the POSIX.1 standard explicitly requires 8-bit char with two's complement arithmetic, it is easier for application writers if the same two's complement guarantees are extended to all of the other standard integer types. Furthermore, in programming environments with a 32-bit long, some POSIX.1 interfaces, such as `rand48()`, cannot be implemented if long does not use a two's complement representation.

To add onto the reality that POSIX chose in 2001, C++20 has only supported two's complement storage since [\[P0907r4\]](#), and C23 has followed suit.

The overwhelming support for 8-bit bytes in hardware and software platforms means that software written for non-8-bit byte is incompatible with software written for 8-bit bytes, and vice versa. C and C++ code targeting non-8-bit byte are incompatible dialects of C and C++.

[Wikipedia quotes](#) the following operating systems as being currently POSIX compliant (and therefore supporting 8-bit bytes):

- AIX
- HP-UX
- INTEGRITY
- macOS
- OpenServer
- UnixWare
- VxWorks
- vz/OS

And many others as being formerly compliant, or mostly compliant.

Even StackOverflow, the pre-AI repository of our best knowledge (after Wikipedia), [gushes with enthusiasm about non-8-bit byte architectures](#), and asks [which exotic architecture the committee cares about](#).

Why bother? A few reasons:

- The complexity of supporting non-8-bit byte architectures sprinkles small but unnecessary burden in quite a few parts of language and library;
- Compilers and toolchains are required to support edge cases that do not reflect modern usage;
- New programmers are easily confused and find C++'s exotic tastes obtuse;
- Some seasoned programmers joyfully spend time supporting non-existent platforms "for portability" if they are unwise;
- [Our language looks silly](#), solving problems that nobody has.

One reason not to bother: there still are processors with non-8-bit bytes. The question facing us is: are they relevant to modern C++? If we keep supporting the current approach where Bytes<sup>R</sup>Us, will developers who use these processors use the new versions of C++?

A cut-the-baby-in-half alternative is to mandate that `CHAR_BIT % 8 == 0`. Is that even making anything better? Only if the Committee decides to keep supporting Digital Signal Processors (DSPs) and other processors where `CHAR_BIT` is not 8 but is a multiple of 8.

This paper cannot succeed without mentioning the PDP-10 (though noting that PDP-11 has 8-bit bytes), and the fact that some DSPs have 24-bit or 32-bit words treated as "bytes." These architectures made sense in their era, where word sizes varied and the notion of a byte wasn't standardized. Today, nearly every general-purpose and embedded system adheres to the 8-bit byte model. The question isn't

whether there are still architectures where bytes aren't 8-bits (there are!) but whether these care about modern C++... and whether modern C++ cares about them.

## § 2. Impact on C

This proposal explores whether C++ is relevant to architectures where bytes are not 8 bits, and whether these architectures are relevant to C++. The C committee might reach a different conclusion with respect to this language. Ideally, both committees would be aligned. This paper therefore defers to WG14 and the SG22 liaison group to inform WG21.

## § 3. Wording

### § 3.1. Language

Edit [**intro.memory**] as follows:

The fundamental storage unit in the ++ memory model is the byte. A byte is **8 bits, which is** at least large enough to contain the ordinary literal encoding of any element of the basic character set literal character set and the eight-bit code units of the Unicode UTF-8 UTF-8 encoding form and is composed of a contiguous sequence of bits, the number of which is bits in a byte. The least significant bit is called the low-order bit; the most significant bit is called the high-order bit. The memory available to a C++ program consists of one or more sequences of contiguous bytes. Every byte has a unique address.

The number of bits in a byte is reported by the macro `CHAR_BIT` in the header `climits`.

## § 3.2. Library

Edit [**climits.syn**] as follows:

```
// all freestanding
#define CHAR_BIT see below8
#define SCHAR_MIN see below-127
#define SCHAR_MAX see below128
#define UCHAR_MAX see below255
#define CHAR_MIN see below
#define CHAR_MAX see below
#define MB_LEN_MAX see below
#define SHRT_MIN see below
#define SHRT_MAX see below
#define USHRT_MAX see below
#define INT_MIN see below
#define INT_MAX see below
#define UINT_MAX see below
#define LONG_MIN see below
#define LONG_MAX see below
#define ULONG_MAX see below
#define LLONG_MIN see below
#define LLONG_MAX see below
#define ULLONG_MAX see below
```

The header `climits` defines all macros the same as the C standard library header `limits.h`, except as noted above.

Except for `CHAR_BIT` and `MB_LEN_MAX`, a macro referring to an integer type `T` defines a constant whose type is the promoted type of `T`.

Edit [**cstdint.syn**] as follows:

The header `cstdint` supplies integer types having specified widths, and macros that specify limits of integer types.

```
int8_t
int16_t
int32_t
int64_t
int_fast8_t
int_fast16_t
int_fast32_t
int_fast64_t
int_least8_t
int_least16_t
int_least32_t
int_least64_t
intmax_t
intptr_t
uint8_t
uint16_t
uint32_t
uint64_t
uint_fast8_t
uint_fast16_t
uint_fast32_t
uint_fast64_t
uint_least8_t
uint_least16_t
uint_least32_t
uint_least64_t
uintmax_t
uintptr_t

// all freestanding
namespace std {
    using int8_t          = signed integer type; // optional
    using int16_t         = signed integer type; // optional
    using int32_t         = signed integer type; // optional
    using int64_t         = signed integer type; // optional
    using intN_t          = see below;           // optional

    using int_fast8_t     = signed integer type;
    using int_fast16_t    = signed integer type;
    using int_fast32_t     = signed integer type;
    using int_fast64_t    = signed integer type;
```

```

using int_fastN_t      = see below;           // optional

using int_least8_t     = signed integer type;
using int_least16_t    = signed integer type;
using int_least32_t    = signed integer type;
using int_least64_t    = signed integer type;
using int_leastN_t     = see below;           // optional

using intmax_t         = signed integer type;
using intptr_t        = signed integer type; // optional

using uint8_t          = unsigned integer type; // optional
using uint16_t         = unsigned integer type; // optional
using uint32_t         = unsigned integer type; // optional
using uint64_t         = unsigned integer type; // optional
using uintN_t          = see below;           // optional

using uint_fast8_t     = unsigned integer type;
using uint_fast16_t    = unsigned integer type;
using uint_fast32_t    = unsigned integer type;
using uint_fast64_t    = unsigned integer type;
using uint_fastN_t     = see below;           // optional

using uint_least8_t    = unsigned integer type;
using uint_least16_t   = unsigned integer type;
using uint_least32_t   = unsigned integer type;
using uint_least64_t   = unsigned integer type;
using uint_leastN_t    = see below;           // optional

using uintmax_t        = unsigned integer type;
using uintptr_t        = unsigned integer type; // optional
}

#define INTN_MIN        see below
#define INTN_MAX        see below
#define UINTN_MAX       see below

#define INT_FASTN_MIN   see below
#define INT_FASTN_MAX   see below
#define UINT_FASTN_MAX  see below

#define INT_LEASTN_MIN  see below
#define INT_LEASTN_MAX  see below

```

```

#define UINT_LEASTN_MAX    see below

#define INTMAX_MIN         see below
#define INTMAX_MAX         see below
#define UINTMAX_MAX        see below

#define INTPTR_MIN         see below           // optional
#define INTPTR_MAX         see below           // optional
#define UINTPTR_MAX        see below           // optional

#define PTRDIFF_MIN        see below
#define PTRDIFF_MAX        see below
#define SIZE_MAX           see below

#define SIG_ATOMIC_MIN     see below
#define SIG_ATOMIC_MAX     see below

#define WCHAR_MIN          see below
#define WCHAR_MAX          see below

#define WINT_MIN           see below
#define WINT_MAX           see below

#define INTN_C(value)      see below
#define UINTN_C(value)     see below
#define INTMAX_C(value)    see below
#define UINTMAX_C(value)   see below

```

The header defines all types and macros the same as the C standard library header `stdint.h` , except that none of the types nor macros are optional because bytes are 8 bits .

All types that use the placeholder *N* are optional when *N* is not 8, 16, 32, or 64. The exact-width types `intN_t` and `uintN_t` for *N* = 8, 16, 32, and 64 are also optional; however, if an implementation defines integer types with the corresponding width and no padding bits, it defines the corresponding typedef-names. Each of the macros listed in this subclause is defined if and only if the implementation defines the corresponding typedef-name.

The macros `INTN_C` and `UINTN_C` correspond to the typedef-names `int_leastN_t` and `uint_leastN_t`, respectively.

Within [**localization**], remove the 4 *mandates* clauses specifying:

```
CHAR_BIT == 8 is true
```

## § References

### § Informative References

#### [P0907r4]

JF Bastien. *Signed Integers are Two's Complement*. 6 October 2018. URL:  
<https://wg21.link/p0907r4>

#### [POSIX]

*The Open Group Base Specifications Issue 8*. 2024. URL:  
<https://pubs.opengroup.org/onlinepubs/9799919799/>