

P3564R0: Make the concurrent forward progress guarantee usable in bulk

Mark Hoemmen, Bryce Adelstein Lelbach, and Michael Garland
(NVIDIA)

SG1, WG21

Hagenberg, Austria, 2025-02

“Make bulk make progress”

- Concurrent forward progress is currently useless for bulk
 - Forward progress guarantee relates to distinct execution agents
 - Nothing specifies when different $f(k)$ invocations run on distinct agents
 - ➔ Parallel is the strongest guarantee that generic code can assume
- Fix is mostly wording
 - Specify that each of bulk's $f(k)$ run on distinct execution agents
 - “Agent” is a legal fiction that need not correspond to “thread”
 - Make it ill-formed to use default (sequential) bulk with scheduler that promises concurrent forward progress
 - Let bulk report error via error channel if it can't fulfill promise for given number of agents

Original intent: 1 execution agent per iteration

- SG1, Wrocław 2024 review of P3481R0, unanimous consent
 - “We need a version of bulk that creates an execution agent per iteration”
- P2300R0
 - Later relaxed bulk wording, only to permit default (sequential) bulk
- P2181R1 (“Correcting the design of bulk execution”)
 - Clarifying P0443R14 bulk’s forward progress wording

Why expose concurrent forward progress?

- Concurrent is effortful to implement; don't throw it away
 - Bulk users more likely to write algorithms with explicit synchronization
 - Concurrent makes e.g., sort more efficient (e.g., fewer bulk launches)
- Parallel programming models expose it as much as possible
 - OpenMP: Concurrent across different “progress units”
 - CUDA
 - Threads in a block: can wait on each other
 - Blocks: concurrent not default, but can ask for it
 - HPX: Can ask for a “parallel section” that permits blocking sync

Run on N agents != parallel algorithm

- Parallel algorithm (e.g., `for_each`) has 2 parts
 - Make available N distinct execution agents
 - Distribute work items (loop iterations) to agents
- Programming models separate them
 - OpenMP: “`#pragma omp parallel`” vs. loop & distribution directives
 - CUDA kernel launch vs. parallel algorithms (Thrust, `stdpar`)
 - ScaLAPACK, High Performance Fortran (HPF)
 - Process grid / `PROCESSORS` directive, vs.
 - Data distributions / `DISTRIBUTE` (et al.) directives
- Programming models let users know & control distribution

Current wording not useful for implementing performant generic parallel algorithms

- Run on ??? number of agents
- No control over number of agents
 - Does large N oversubscribe hardware?
 - Do I need to distribute work items?

Let bulk fail if concurrent + N too large

- Implementations may not be able to promise concurrent forward progress for all possible N
- Already true of popular programming models, e.g., OpenMP
- Implementations can handle this in 2 ways
 - If recoverable: Error channel (CHANGE)
 - If not: terminate() (already permitted)