



`std::array` is a wrapper for an array!

Document number:

[P3737R2](#)

Date:

2025-11-20

Audience:

LEWG

Project:

ISO/IEC 14882 Programming Languages — C++, ISO/IEC JTC1/SC22/WG21

Reply-to:

Jan Schultke <janschultke@gmail.com>

GitHub Issue:

wg21.link/P3737/github

Source:

github.com/Eisenwave/cpp-proposals/blob/master/src/array.cow



ABSTRACT: The `std::array` class template is implemented as a simple wrapper type for a "C-style array". However, its specification in the standard is considerably more permissive and should be simplified.

§ Contents

- 1 **Revision history**
 - 1.1 Changes since R1
 - 1.2 Changes since R0
- 2 **Introduction**
 - 2.1 What the standard says
 - 2.2 What the standard does not say
- 3 **Motivation**
 - 3.1 Isn't this a waste of time?
- 4 **Design considerations**
 - 4.1 Zero-length `std::array` status quo



4.1.1	Conclusion
4.2	Trivial copyability of zero-length arrays
4.3	Double-brace initialization for zero-length arrays
4.3.1	Conclusion
4.4	Problematic iterator requirements for zero-length arrays
4.4.1	Conclusion
5	Impact on implementations
6	Wording
6.1	[array.overview]
6.2	[array.zero]
7	References

§ 1. Revision history

§ 1.1. Changes since R1

- Fixed a typo in §4.4. Problematic iterator requirements for zero-length arrays
- Updated discussion and wording based on [LWG4276] being accepted in Kona

§ 1.2. Changes since R0

- Formulated an overall design strategy in §4. Design considerations.
- Removed any changes related to `array<T, 0>::front()` and `array<T, 0>::back()`. Instead, R1 assumes that [LWG4276] will be accepted.
- Excluded bit-field members from appearing in `std::array`.
- Made trivial copyability and trivial default construction of `std::array<T, 0>` unconditional; see §4.2. Trivial copyability of zero-length arrays.
- Rebased §6. Wording on [N5014] and made minor wording tweaks.
- Minor editorial changes throughout the paper.

§ 2. Introduction

The `std::array` class template has established itself as a de-facto replacement for "builtin arrays" or "C-style arrays" in many code bases. This also means that it is frequently taught to novice programmers, with an explanation along the lines of:

”

`std::array` is just a wrapper for a C-style array:

```
template<class T, size_t N>
struct array {
    T __array[N];
    // ...
};
```



While this explanation is not correct for zero-length `std::array`s, it does match how the template is implemented in every standard library for `N != 0`, and there is very little reason not to implement it in this obvious fashion.

§ 2.1. What the standard says

The actual specification of `std::array` is not so simple, and is a combination of multiple constraints on the implementation:

- It is a class template, a contiguous container ([array.overview] paragraph 1), and a reversible container (with an exception; see [array.overview] paragraph 3), and it meets some requirements of a sequence container.
- It can be list-initialized with up to `N` elements whose types are convertible to `T` ([array.overview] paragraph 2). This is obviously not exhaustive; initialization with `{}` or `{other_array}` should also be possible.
- It is a structural type if `T` is a structural type ([array.overview] paragraph 4), and therefore, also a literal class type in that event.

Additionally, while this does not strictly specify anything about the layout, some helper functions in the standard library de-facto rely on it. Take [array.creation] for example:

”

```
template<class T, size_t N>
constexpr array<remove_cv_t<T>, N> to_array(T (&a)[N]);
```

Mandates: `is_array_v<T>` is `false` and `is_constructible_v<remove_cv_t<T>, T&>` is `true`.

Preconditions: `T` meets the `Cpp17CopyConstructible` requirements.

Returns: `{{ a[0], ..., a[N - 1] }}`

The use of double braces in the *Returns* specification would be nonsensical if `std::array` was not "a wrapper for an array".

§ 2.2. What the standard does not say

Notably, there are quite a couple of guarantees that are absent.



EXAMPLE: It would be compliant to implement `std::array` as follows:



```

struct alignas(1024) malice_and_evil {
    constexpr malice_and_evil() { }
    malice_and_evil(const malice_and_evil&) { }
};

template<class T, size_t N>
struct array {
    T __array[N];
    malice_and_evil evil;
};

```

Such an implementation technically satisfies all the requirements for `std::array`, but

- copying the array would not be a constant expression,
- `std::array` would not be trivially copyable for any type, and
- its size and alignment would be much greater than that of `T[N]`.



EXAMPLE: An even more insane implementation would be:

```

struct gobbler {
    constexpr gobbler() = default;
    constexpr gobbler(auto&&) {}
};

template<class T, size_t N>
struct array {
    gobbler __gobblers[N];
    T __array[N]; // necessary to satisfy contiguous container requirement.
    // ...
};

```

Since `[array]` never states what effect list-initialization has for a `std::array`, and even `std::to_array` is just stated to return the result of *some expression*, nothing suggests that `begin()` would give us an iterator to `x` after initializing like `std::array<T, 1> {{x}}`. All list-initialization could be "gobbled up".

§ 3. Motivation

It seems like the vagueness in the specification serves no practical purpose; it is unclear what implementations could do with the additional freedom, other than pranking their users. It would be beneficial to the C++ community if the simplified explanation in §2. [Introduction](#) was what the standard actually said.

A stricter specification would provide additional useful guarantees such as `std::array<T, N>` being trivially copyable when `T` is trivially copyable. This is relevant to use cases like `std::bit_cast<std::array<std::byte, sizeof(x)>>(x)`, which *technically* rely on implementation details, not on standard behavior.

§ 3.1. Isn't this a waste of time?



While it could be argued that only a malicious implementation would violate our user expectations as in §2.2. [What the standard does not say](#) and it is therefore time-wasting to restrict `std::array` any further, it would be unusual for WG21 to shy away from standardizing universally existing practice and to recommend users to rely on non-standard implementation details, simply because those implementation details are widespread. If the remaining implementation freedom can only be used for evil, perhaps we should not grant it.

§ 4. Design considerations

The overall design strategy is to standardize the "greatest common denominator" of `libc++` and `libstdc++`. No breaking changes should be made; existing practice should be standardized.

While the specification for arrays of nonzero length is rather obvious, it is unclear how many guarantees we want to provide for zero-length arrays. For example, should `std::array<std::string, 0>` be trivially copyable, even though `std::string` is not?

Within [\[array.zero\]](#), there are some long-standing issues going back to 2012. LWG has visited this subclause many times in [\[LWG2157\]](#), but never fully completed a solution. This work has been absorbed mostly unmodified into §6. [Wording](#).

§ 4.1. Zero-length `std::array` status quo

The zero-length case is also where we see some implementation divergence in size and alignment of the array. The following table shows how major standard libraries implement zero-length `std::array`.

Library	Implementation	Size	Trivially copyable	Assignable
MSVC STL	contains T if T is default-constructible, otherwise <code>struct{}</code>	<code>sizeof(T)</code> or 1	depends on T	depends on T
libstdc++	contains <code>struct{}</code>	1	always	always
libc++	contains (possibly const) <code>__empty[sizeof(T)]</code>	<code>sizeof(T)</code>	always	depends on T



BUG: The MSVC STL implementation is non-compliant. Despite `std::array<T, 0>` being a *zero-length* container with *no* elements, it actually holds one element (and call its constructors and destructors) as long as T is default-constructible.

This bug has been reported, and fixing it would break ABI ([\[STL5583\]](#)).

§ 4.1.1. Conclusion

Generally speaking, it is desirable if a zero-length `std::array` behaves as similarly to a regular `std::array` of the same element type. `libc++` is the only implementation that does this well. The "greatest common denominator" between these implementations (excluding MSVC STL) should be standardized, which is:

- `std::array<T, 0>` is trivially copyable.
- `std::array<T, 0>` is assignable if `T` is.
- `std::array<T, 0>` has size and alignment at most that of `T`.
- `std::array<T, 0>` is not an empty class.

An implementation satisfying these requirements would look something like:

```
struct __empty {};  
  
template<class T>  
struct array<T, 0> {  
    // const __empty if T is const  
    using __empty_type = copy_cv<T, __empty>;  
  
    // No alignas for libstdc++.  
    alignas(T) __empty_type arr;  
  
    // ...  
};
```

§ 4.2. Trivial copyability of zero-length arrays

Both `libc++` and `libstdc++` provide a trivially copyable and trivially default-constructible `std::array<T, 0>`.

While it could be argued that these properties should only be guaranteed when `T` has them too, there is no clear benefit to the implementation or wording to make these guarantees conditional, and we are simply standardizing existing practice with this design.

§ 4.3. Double-brace initialization for zero-length arrays

Note that we need to make double-brace initialization like `std::array<int, 0>{{{}}` valid to make generic programming easier. It is plausible that we perform this when expanding an empty pack like: `std::array<int, sizeof...(args)>{{ args... }}`.

Making this valid requires either some non-static data member, or a base class. An empty base class would make the array as a whole an empty class, and this would be an ABI break, so it is out of the question.

§ 4.3.1. Conclusion

Standardize the existing practice of having a non-static data member which enables double-brace initialization.

§ 4.4. Problematic iterator requirements for zero-length arrays



[array.zero] paragraph 2 specifies:

” In the case that `N == 0`, `begin()` == `end()` == unique value. The return value of `data()` is unspecified.

Firstly, it is unclear whether this "unique value" is meant to be unique per object, unique for each invocation, etc.

Secondly, this requirement was never implemented by any compiler and it is too late to fix now. Note that MSVC STL, libc++, and libstdc++ all use `T*` as an iterator type. Considering that, a possible implementation looks like:

```
template<class T>
struct zero_length_array {
    union U {
        char c;
        T i;
        U() = default;
    } u;

    constexpr const T* begin() const noexcept { return &u.i; }
    constexpr const T* end() const noexcept { return begin(); }
    // ...
};

constexpr zero_length_array<int> a{}; // OK
static_assert(a.begin() == a.end()); // OK
```

However, this would require `std::array<T, 0>` to be at least one `T` large, and it is only a single byte large for libstdc++. Changing the size of the type would break ABI. The only way to conjure up a `T*` out of thin air would be to use `reinterpret_cast`, but that would not work in constant expressions.

§ 4.4.1. Conclusion

Delete the uniqueness requirement. Without specifying anything special for zero-length arrays, it still acts as an empty range, and `begin() == end()` is `true`, which is all we really need.

§ 5. Impact on implementations

For `std::array`s of nonzero length, every implementation already complies with the proposed changes. For zero-length `std::array`s, libstdc++ and libc++ already comply with the proposed changes.

The MSVC STL implements zero-length `std::array`s incorrectly and does not comply with C++26 (or any prior standard), and fixing this would require an ABI break. See ([STL5583], §4.1.

Zero-length `std::array` status quo). If that ABI break ever happens, the implementation can simply be aligned with what is being proposed here.



§ 6. Wording

The following changes are relative to [N5014] with the changes in [LWG4276] applied.

§ 6.1. [array.overview]

Change [array.overview] paragraph 1 as follows:



The header `<array>` defines a class template for storing fixed-size sequences of objects. ~~An array is a contiguous container ([container.reqmts]).~~ An ~~instance~~ object of type `array<T, N>` stores `N` elements of type `T`, so that `size()` ~~== `N` is an invariant~~ always equals `N`.

Change [array.overview] paragraph 2 as follows:



An array is an aggregate ([dcl.init.aggr]) ~~that can be list-initialized with up to `N` elements whose types are convertible to `T`~~ with no base classes and no bit-field members. A specialization `array<T, N>` has a single public non-static data member of type "array of `N` `T`" if `N` is nonzero; otherwise the contents are specified in [array.zero].

[Note: An array is trivially copyable, trivially default-constructible, standard-layout, and a structural type if `T` has these properties, respectively. — end note]

Change [array.overview] paragraph 3 as follows:



An array meets all of the requirements of a container ([container.reqmts]), of a contiguous container, and of a reversible container ([container.rev.reqmts]), except that a ~~default-constructed array object~~ default-initialized or value-initialized object of type `array<T, N>` is not empty if `N > 0`. An array meets some of the requirements of a sequence container ([sequence.reqmts]). Descriptions are provided here only for operations on array that are not described in one of these tables, and for operations where there is additional semantic information.

Delete [array.overview] paragraph 4:



~~`array<T, N>` is a structural type ([term.structural.type]) if `T` is a structural type. Two values `a1` and `a2` of type `array<T, N>` are template-argument-equivalent ([temp.type]) if and only if each pair of corresponding elements in `a1` and `a2` are template-argument-equivalent.~~

Change [array.overview] paragraph 5 as follows:



```
namespace std {
    template<class T, size_t N>
    struct array {
        // non-static data members
        T arr[N]; // exposition-only

        // types
        using value_type = T;
        using pointer = T*;
        [...]
    };
}
```



§ 6.2. [array.zero]

Delete all paragraphs within [array.zero]:



- 1 array shall provide support for the special case `N == 0`.
- 2 In the case that `N == 0`, `begin() == end() == unique value`. The return value of `data()` is unspecified.
- 3 Member function `swap()` shall have a non-throwing exception specification.

Insert new paragraphs within [array.zero]:



- 1 A specialization `array<T, 0>` does not have an *arr* data member. Instead, it has a non-static data member of unspecified, trivially copyable, standard-layout, empty aggregate *U* with no base classes and with the same cv-qualification as *T*. The size and alignment of *U* is an implementation-defined choice between 1 and the size and alignment of *T*.
 - 2 The value representation of a specialization `array<T, 0>` is empty.
 - 3 A specialization `array<T, 0>` has identical member functions as the primary template, except that
 - the `begin`, `end`, `cbegin`, `cend`, `rbegin`, `rend`, and `data` member functions return value-initialized results, and
 - the `fill` and `swap` member functions have no effect and have a non-throwing exception specification.
- [Note: Any call to `front`, `back`, or `operator[]` violates a hardened precondition ([structure.specifications]). — end note]

§ 7. References

[N5014] *Thomas Köppe*. Working Draft, Programming Languages — C++ (2025-08-05)
<https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2025/n5014.pdf>



[LWG2157] *Daryle Walker*. How does `std::array<T,0>` initialization work when T is not default-constructible? (2012-05-08)
<https://cplusplus.github.io/LWG/issue2157>

[LWG4276] *Jan Schultke*. `front()` and `back()` are not hardened for zero-length `std::arrays` (2025-06-08)
<https://cplusplus.github.io/LWG/issue4276>

[STL5583] *Jan Schultke*. `<array>`: `std::array<T,0>` calls constructors and destructors of T (2025-06-13)
<https://github.com/microsoft/STL/issues/5583>